

Distributed Sagas

A Protocol for Coordinating Microservices

Caitie McCaffrey

Distributed Systems Engineer



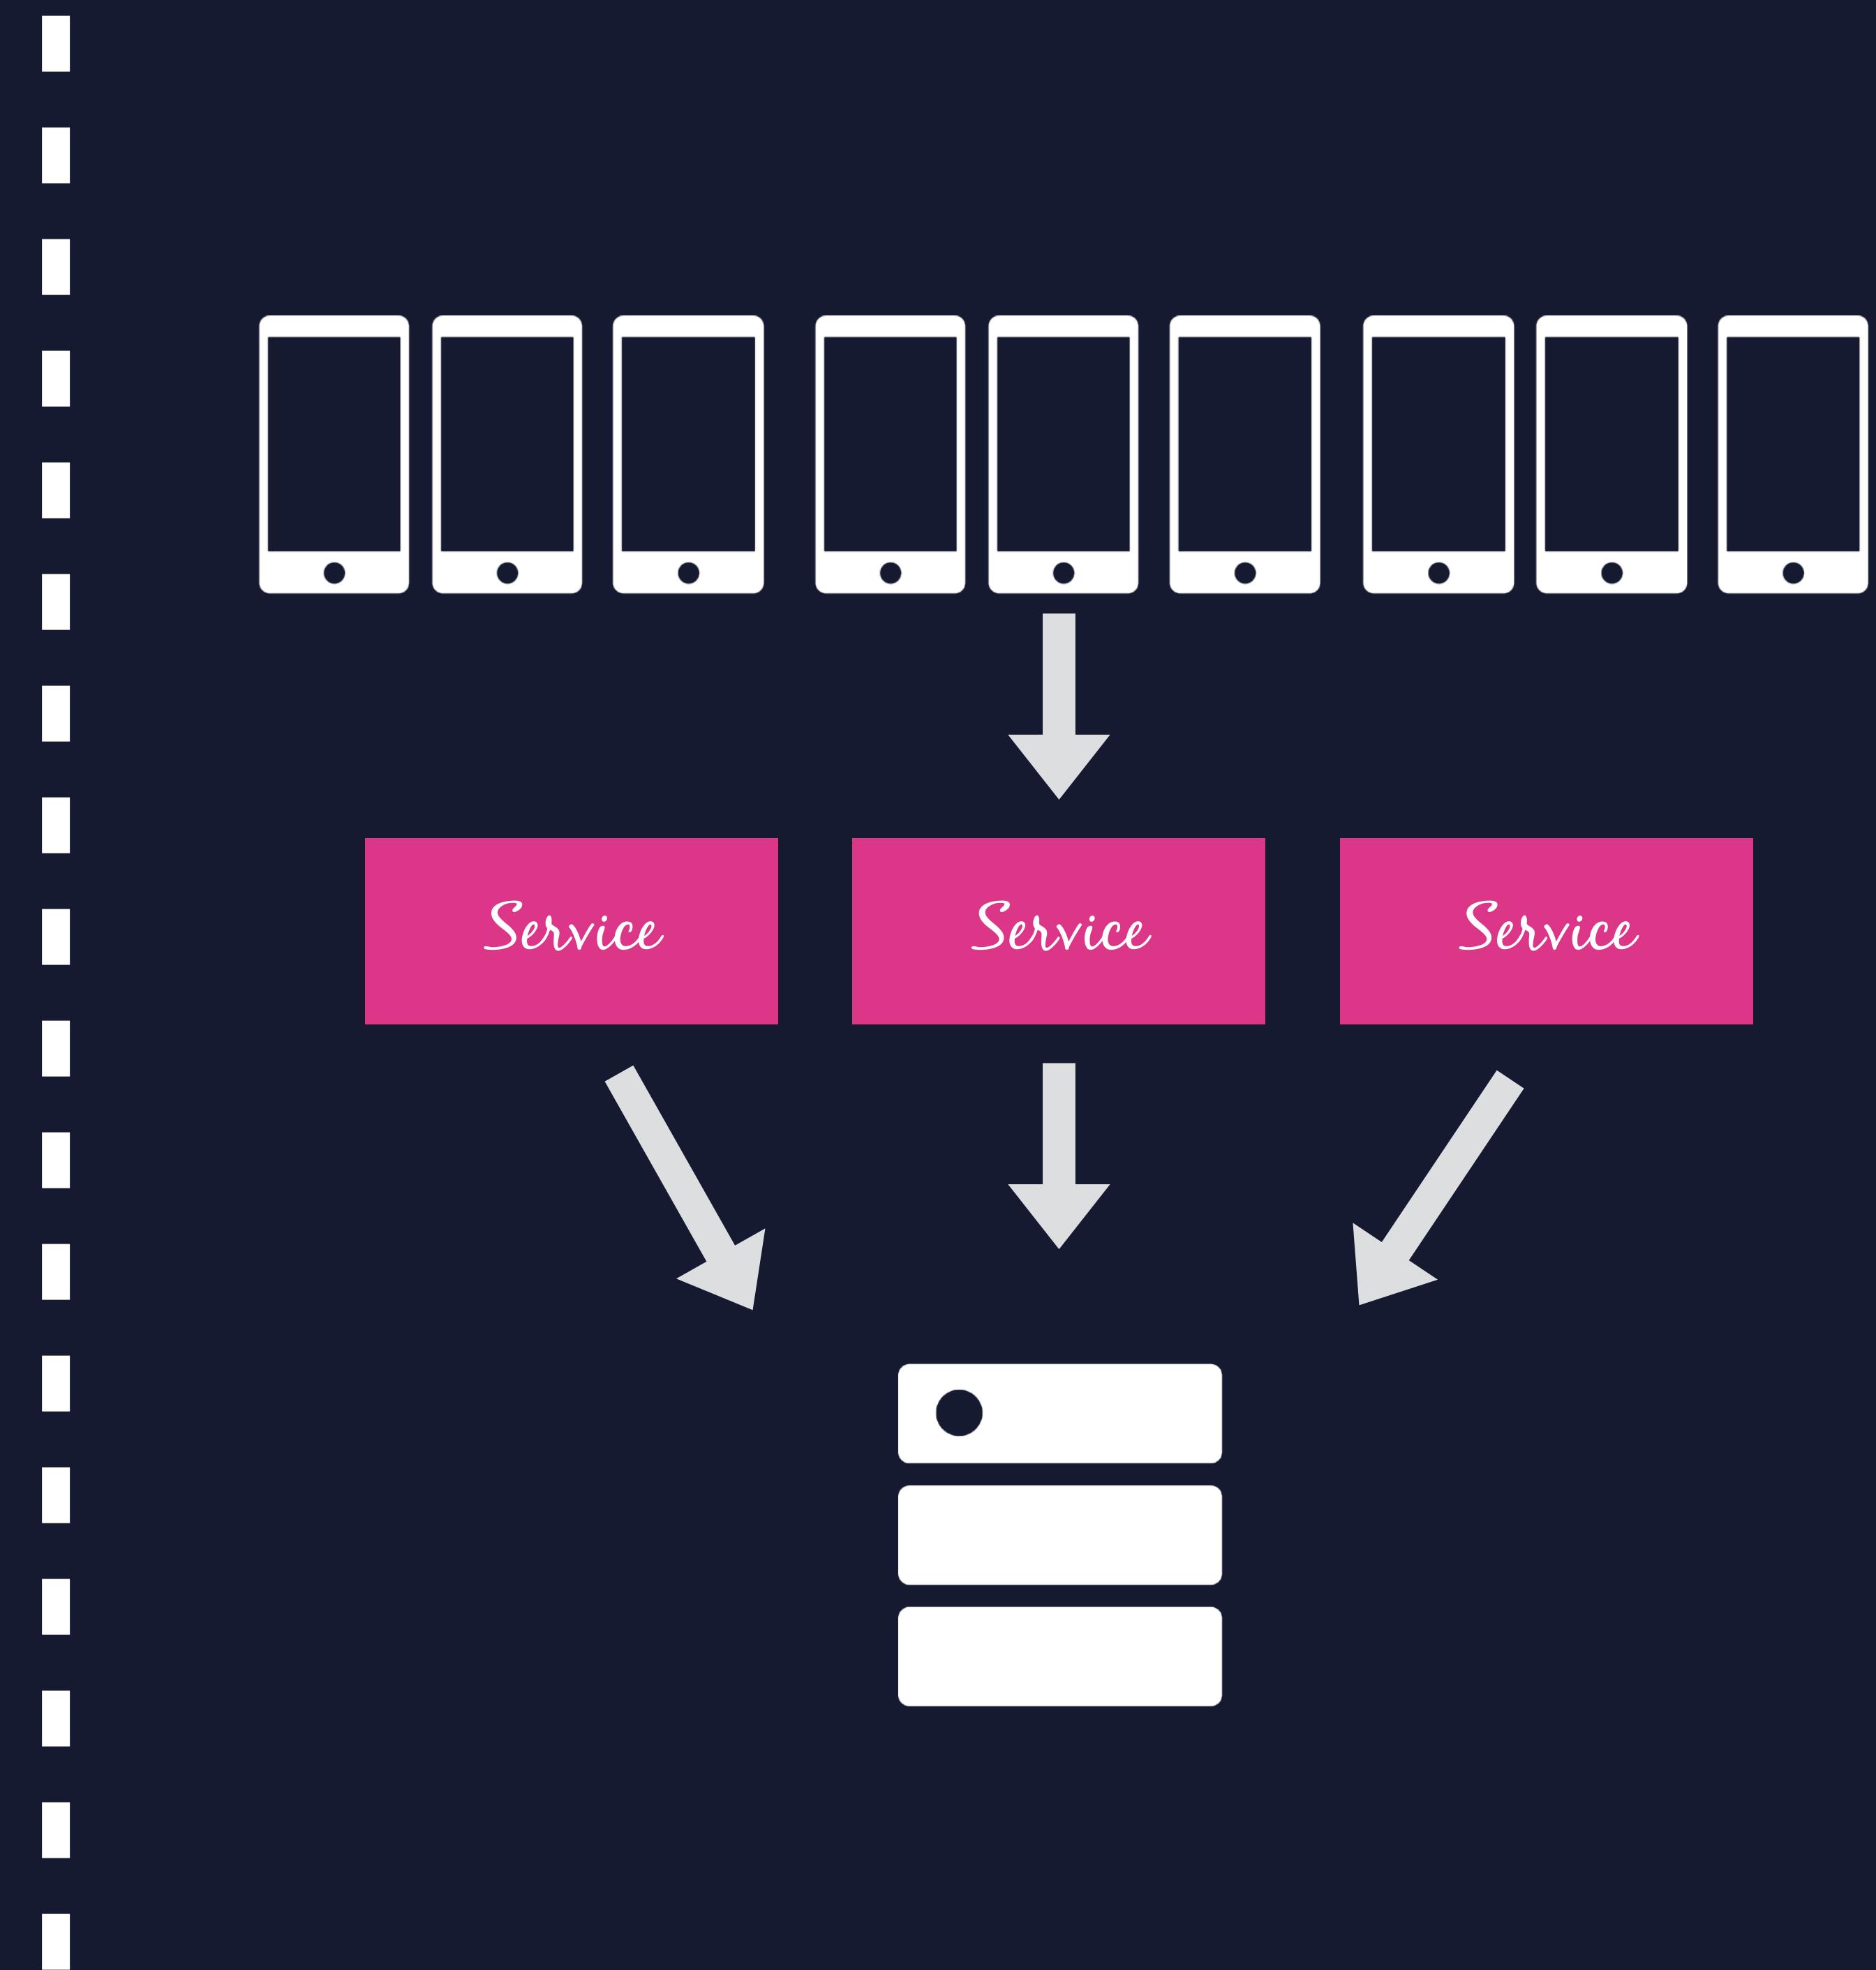
@caitie



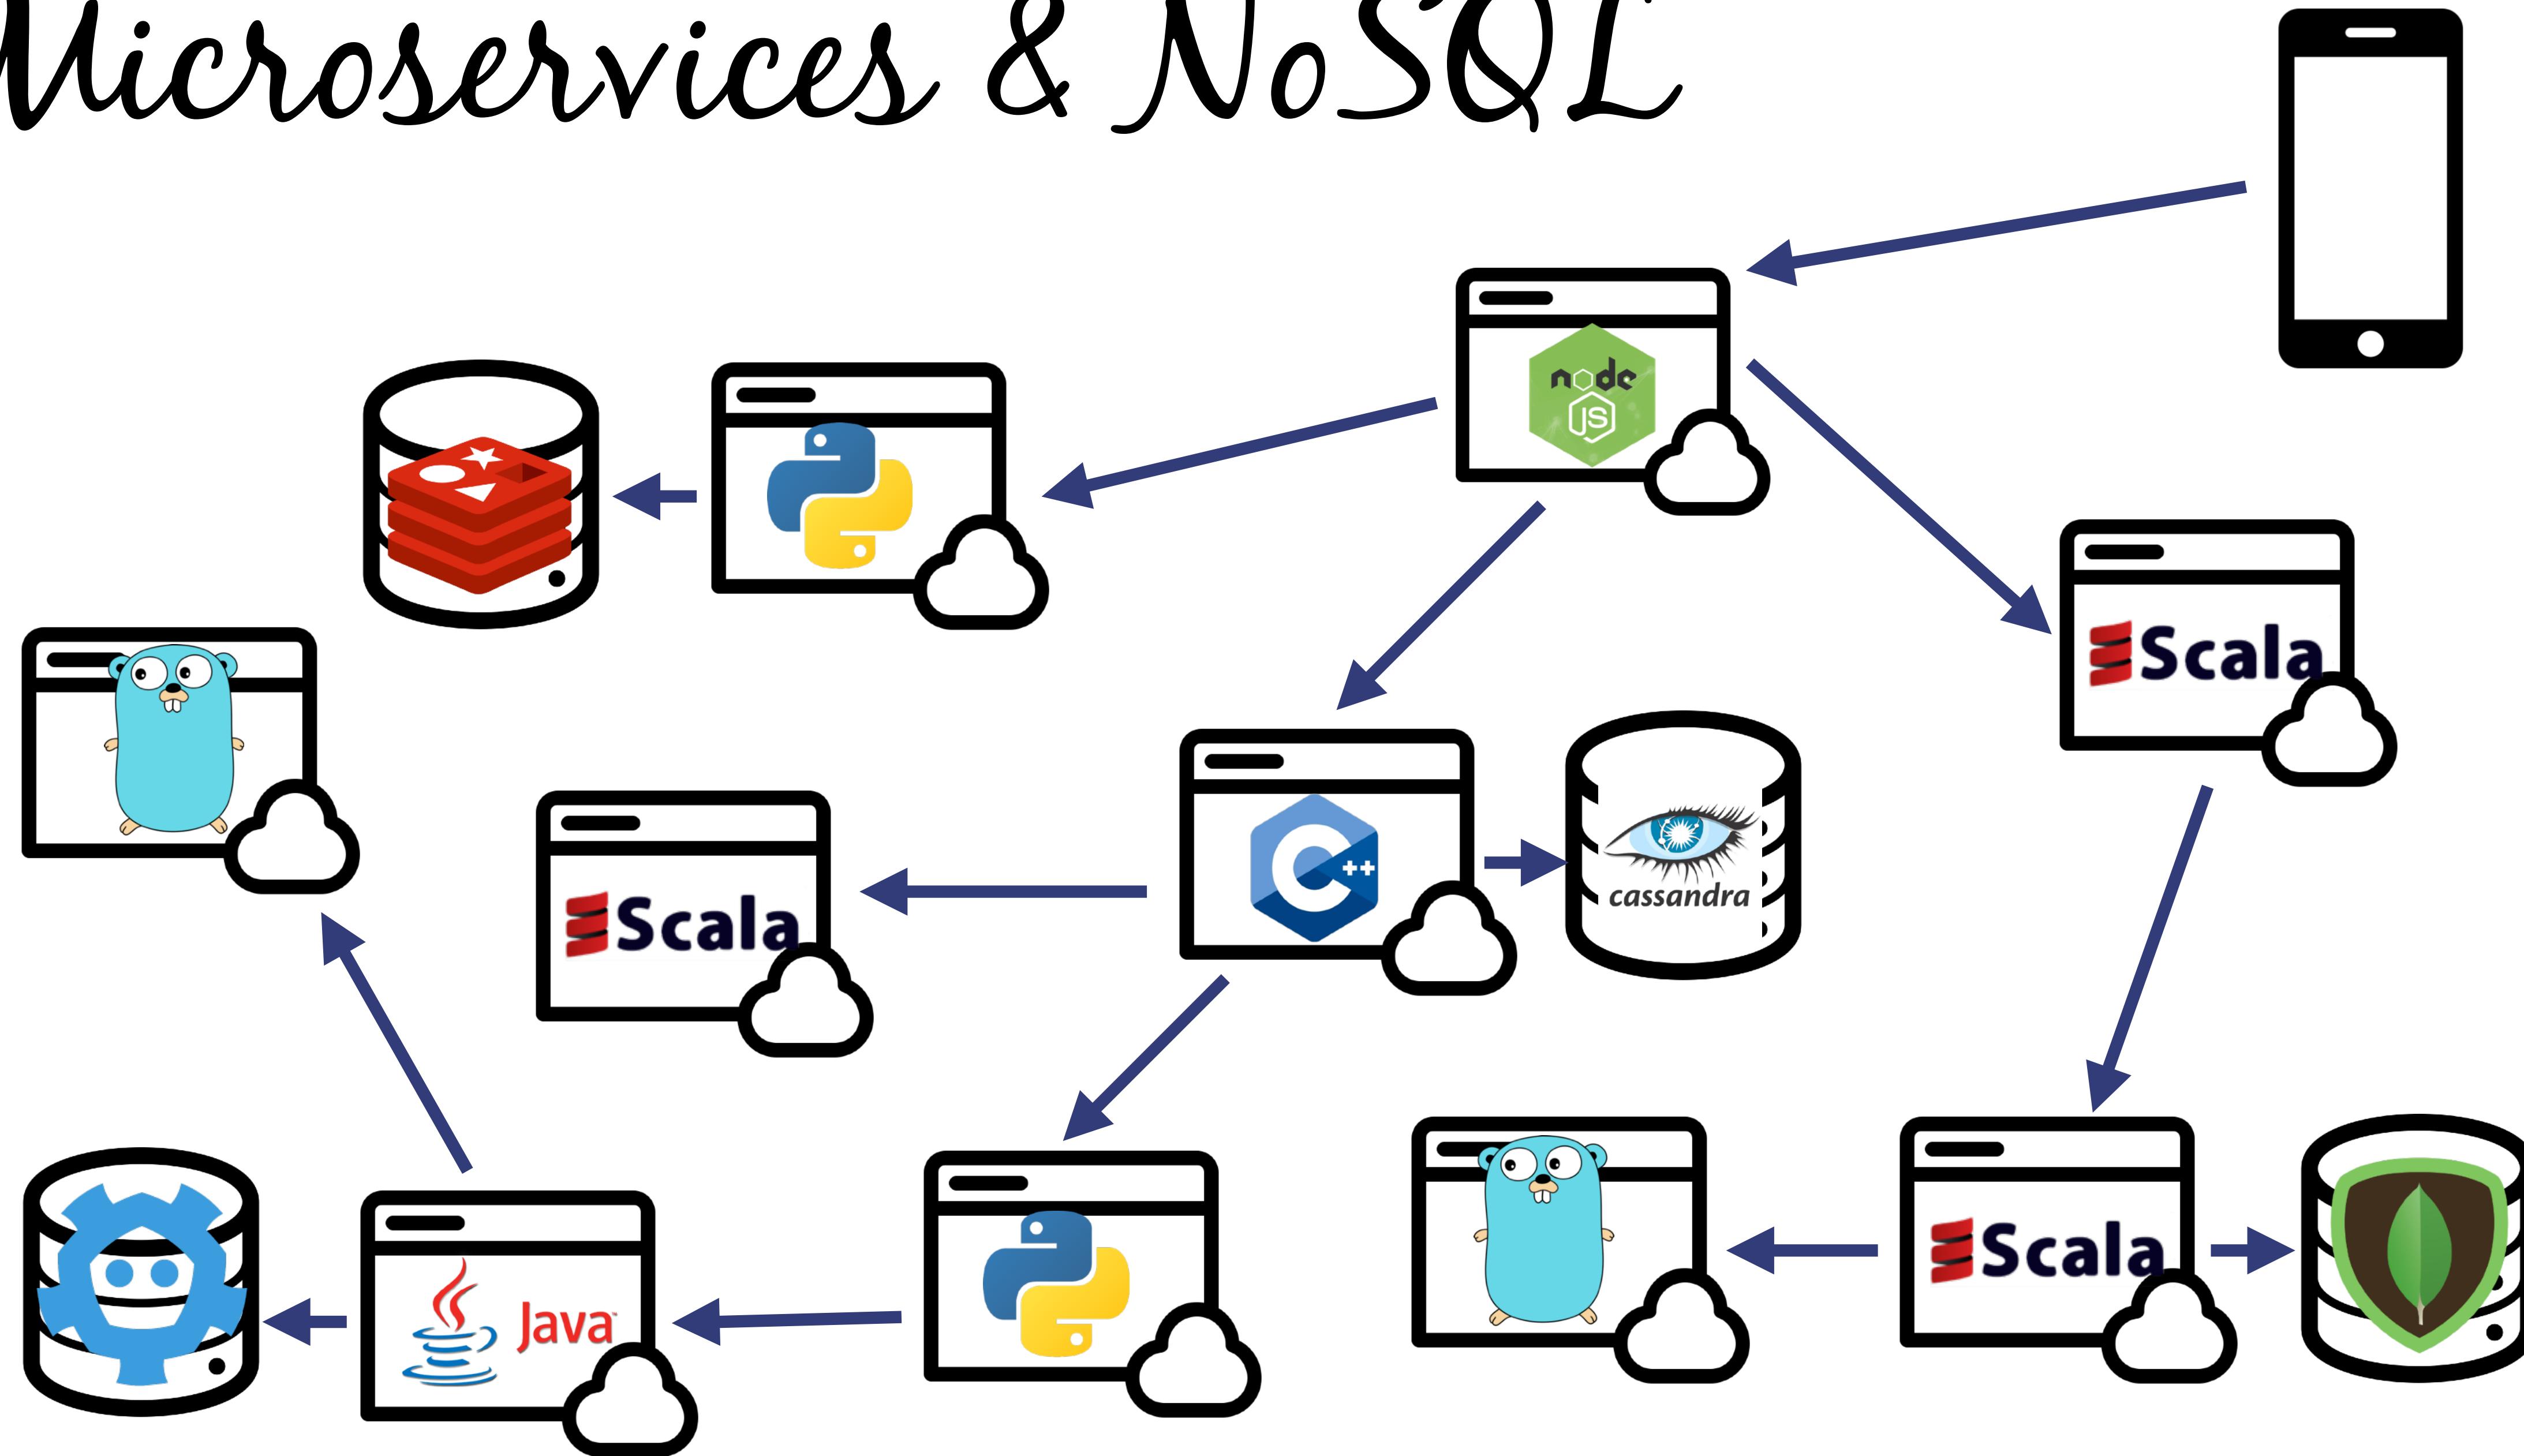
caitiem.com



Monoliths



Microservices & NoSQL



2015

Feral Concurrency Control: An Empirical Investigation of Modern Application Integrity

Peter Bailis, Alan Fekete[†], Michael J. Franklin, Ali Ghodsi, Joseph M. Hellerstein, Ion Stoica
UC Berkeley and [†]University of Sydney

ABSTRACT

The rise of data-intensive “Web 2.0” Internet services has led to a range of popular new programming frameworks that collectively embody the latest incarnation of the vision of Object-Relational Mapping (ORM) systems, albeit at unprecedented scale. In this work, we empirically investigate modern ORM-backed applications’ use and disuse of database concurrency control mechanisms. Specifically, we focus our study on the common use of *feral*, or application-level, mechanisms for maintaining database integrity, which, across a range of ORM systems, often take the form of declarative correctness criteria, or invariants. We quantitatively analyze the use of these mechanisms in a range of open source applications written using the Ruby on Rails ORM and find that feral invariants are the most popular means of ensuring integrity (and, by usage, are over 37 times more popular than transactions). We evaluate which of these feral invariants actually ensure integrity (by usage, up to 86.9%) and which—due to concurrency errors and lack of database support—may lead to data corruption (the remainder), which we experimentally quantify. In light of these findings, we present recommendations for database system designers for better supporting

Rails is interesting for at least two reasons. First, it continues to be a popular means of developing responsive web application front-end and business logic, with an active open source community and user base. Rails recently celebrated its tenth anniversary and enjoys considerable commercial interest, both in terms of deployment and the availability of hosted “cloud” environments such as Heroku. Thus, Rails programmers represent a large class of consumers of database technology. Second, and perhaps more importantly, Rails is “opinionated software” [41]. That is, Rails embodies the strong personal convictions of its developer community, and, in particular, David Heinemeier Hansson (known as DHH), its creator. Rails is particularly opinionated towards the database systems that it tasks with data storage. To quote DHH:

“I don’t want my database to be clever! . . . I consider stored procedures and constraints vile and reckless destroyers of coherence. No, Mr. Database, you can not have my business logic. Your procedural ambitions will bear no fruit and you’ll have to pry that logic from my dead, cold object-oriented hands . . . I want a single layer of cleverness: My domain model.” [55]

2015

Feral Concurrency Control: An Empirical Investigation of Modern Application Integrity

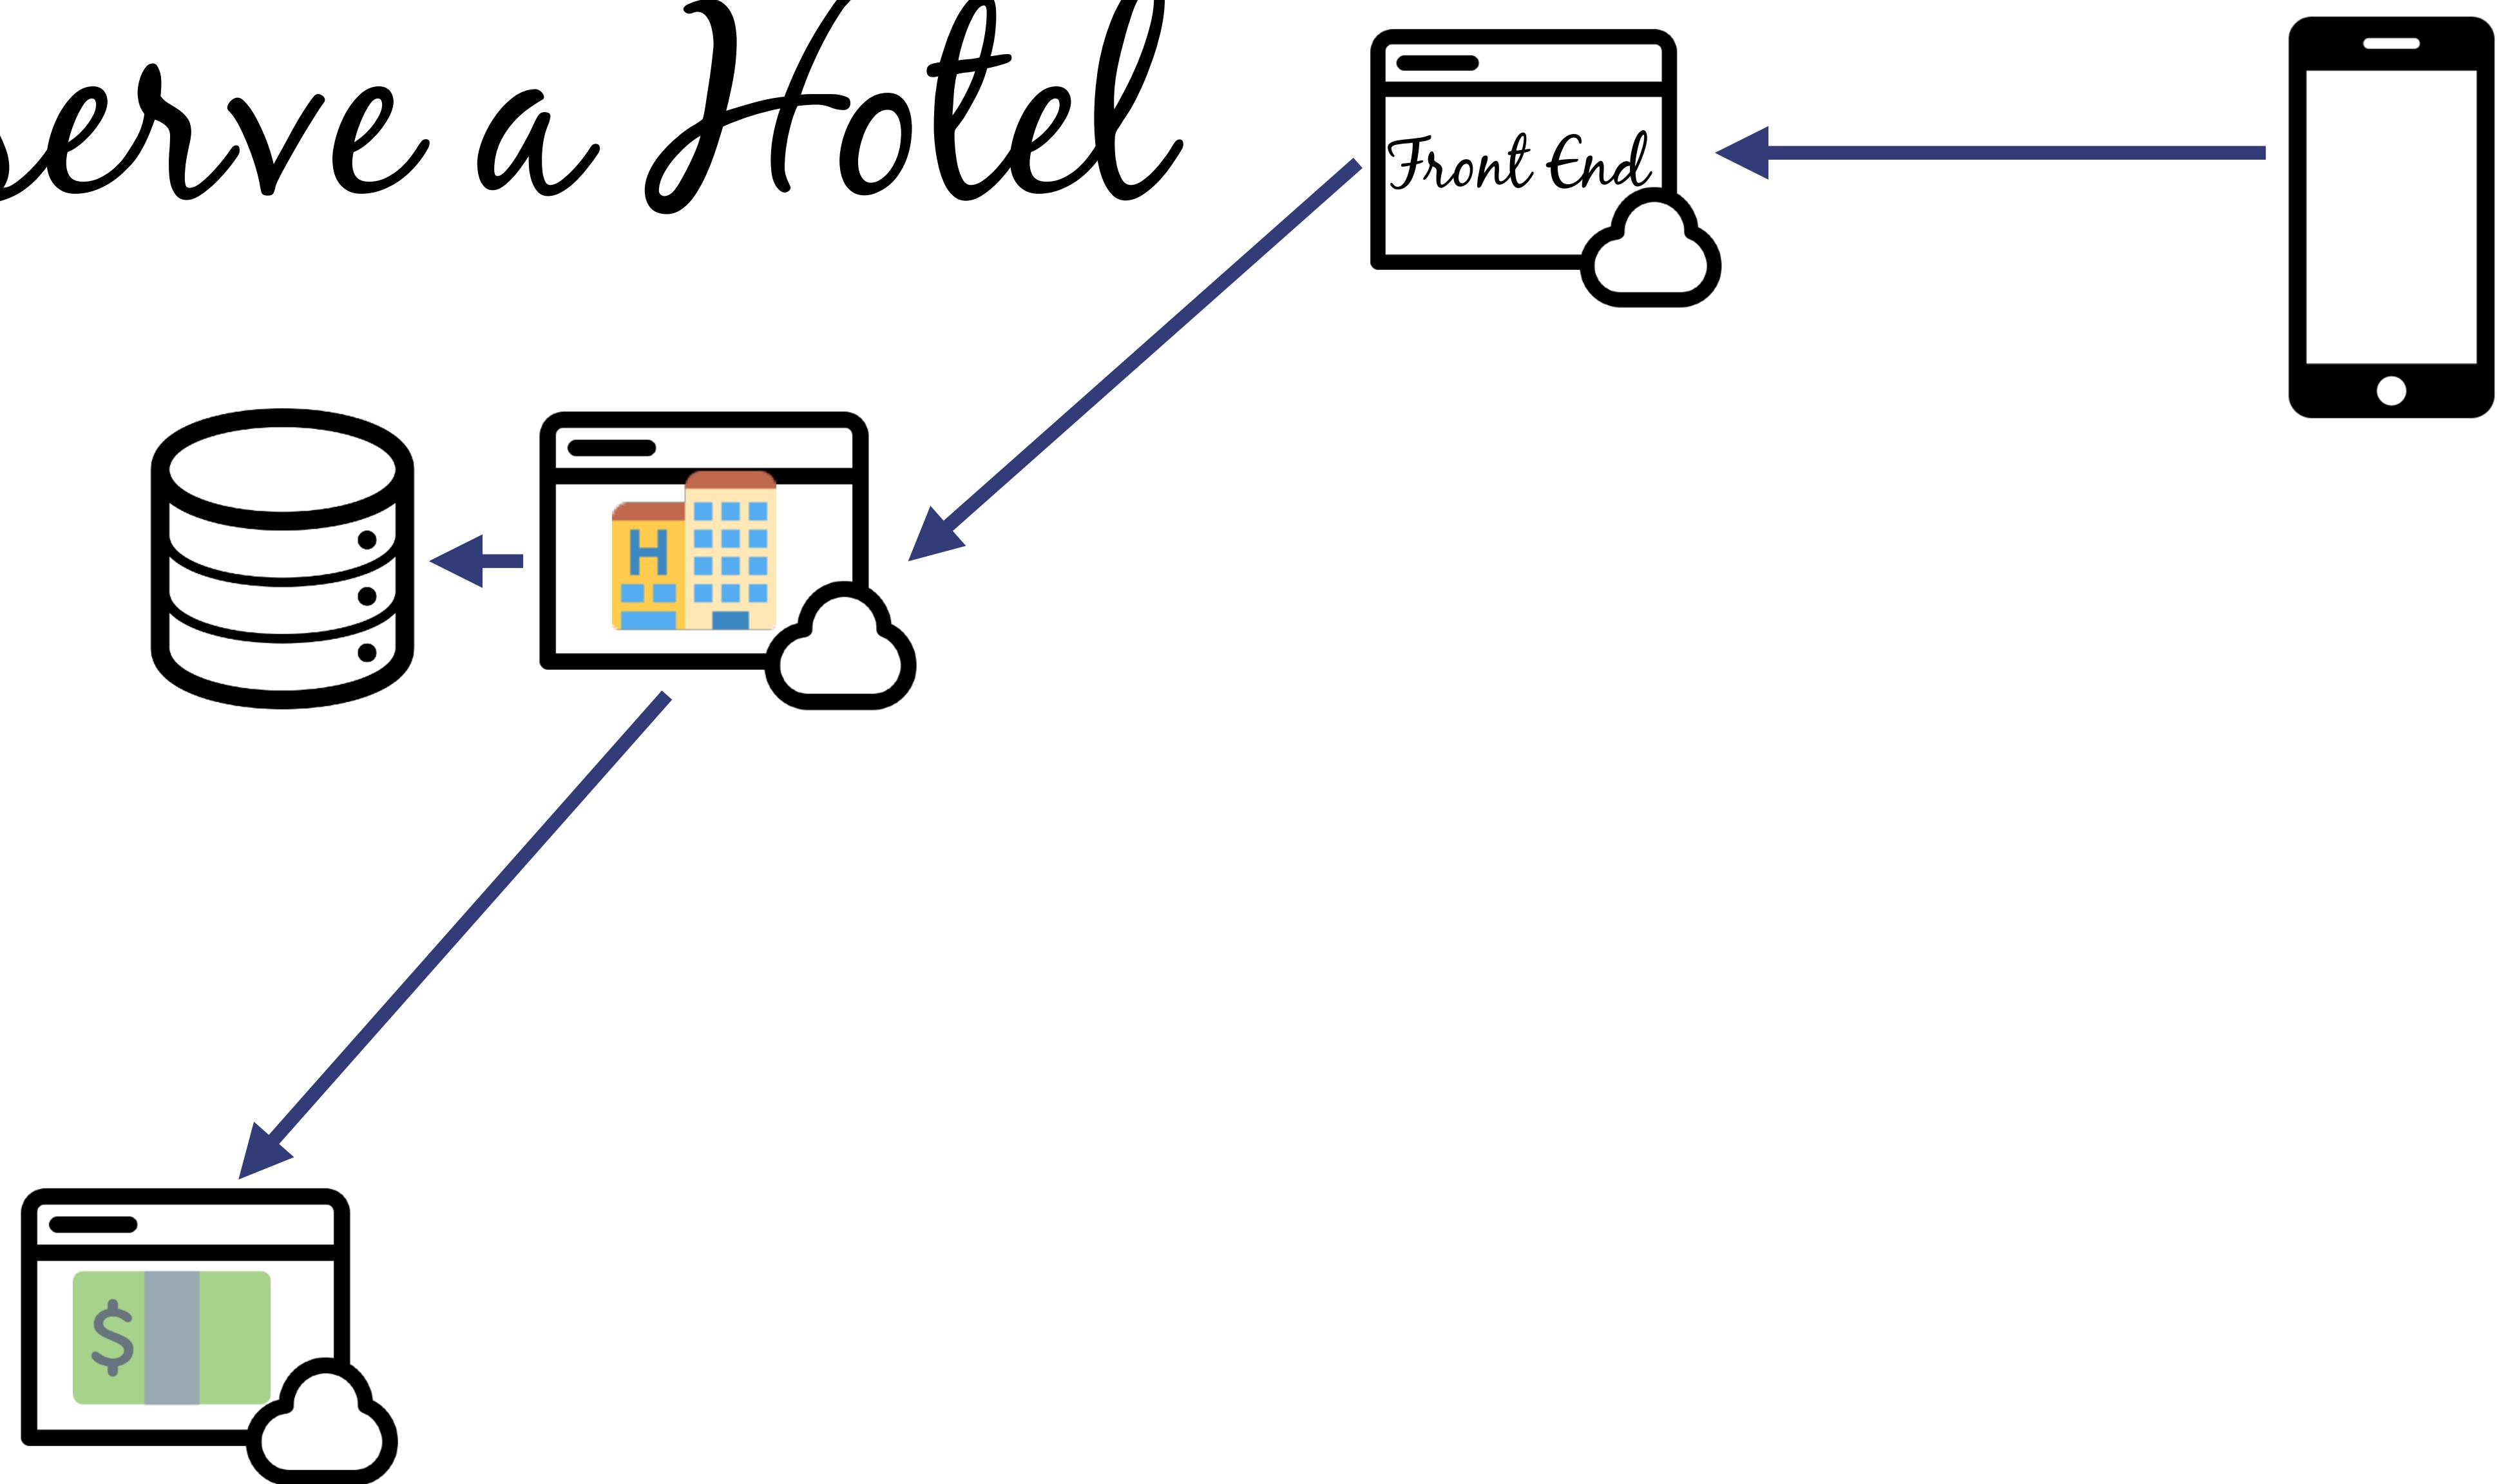
Peter Bailis, Alan Fekete[†], Michael J. Franklin, Ali Ghodsi, Joseph M. Hellerstein, Ion Stoica
UC Berkeley and [†]University of Sydney

“Application-level Mechanisms for maintaining database integrity”

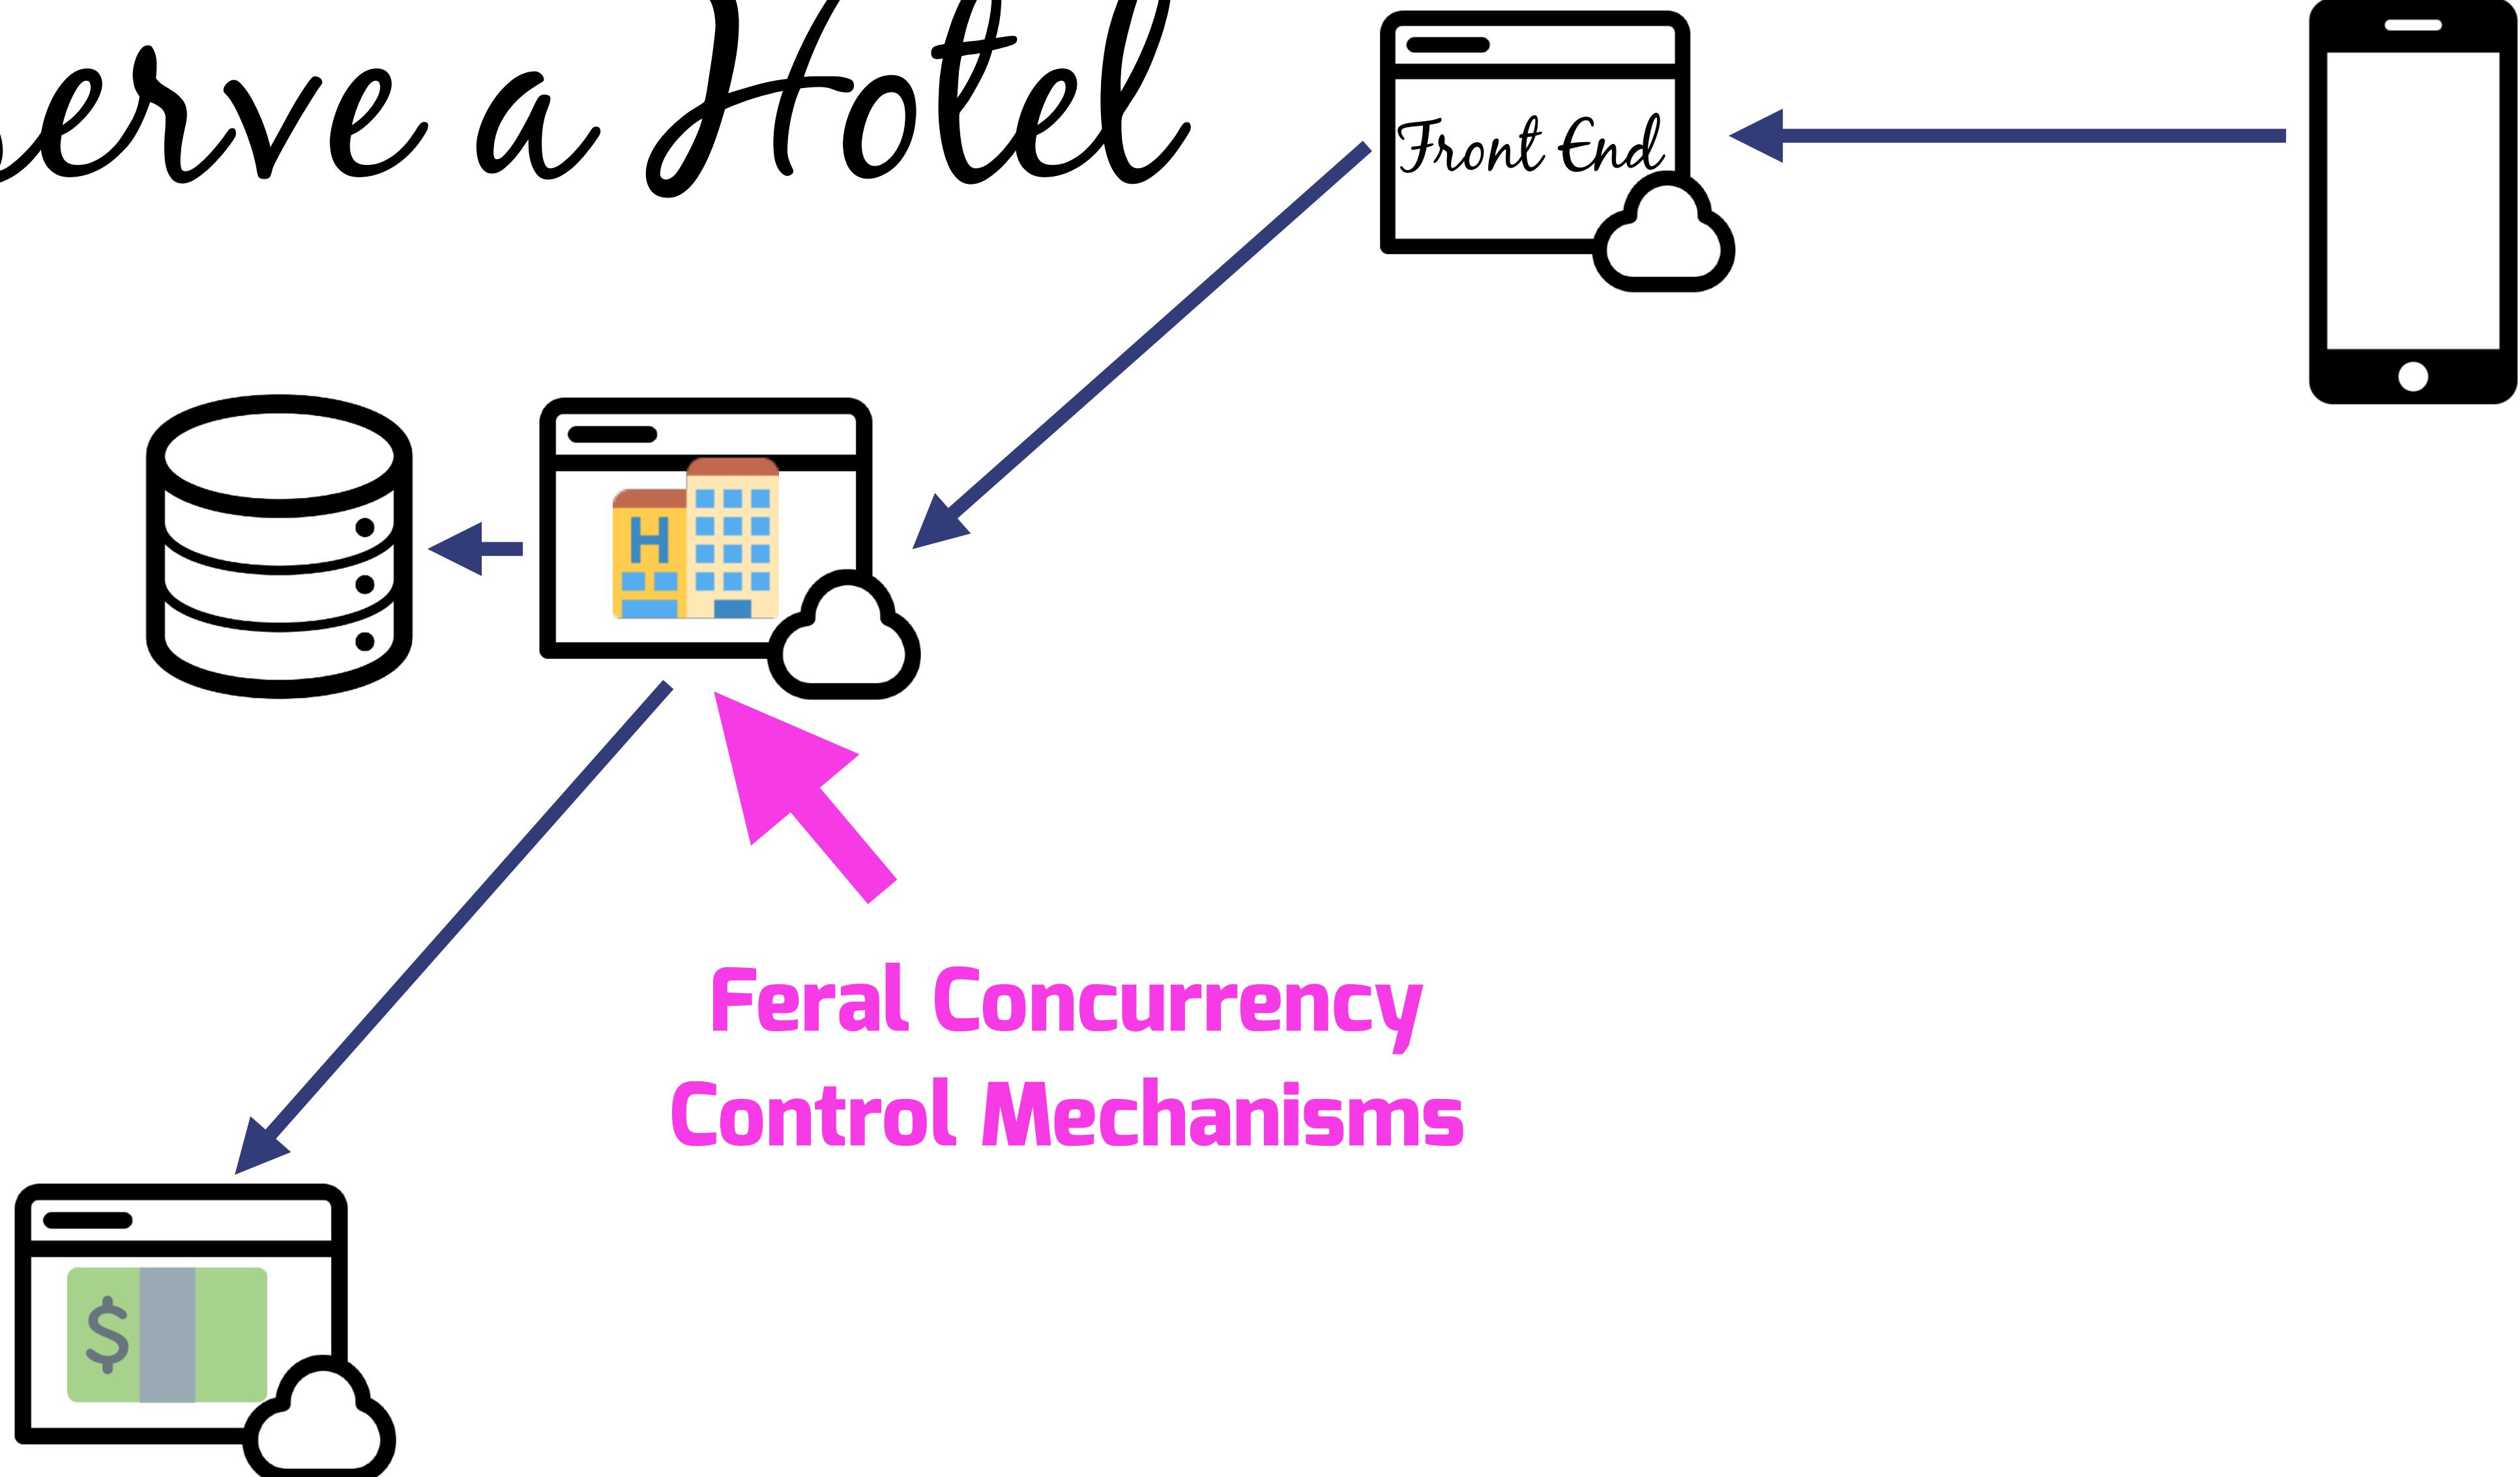
86.9%) and which—due to concurrency errors and lack of database support—may lead to data corruption (the remainder), which we experimentally quantify. In light of these findings, we present recommendations for database system designers for better supporting

destroyers of coherence. NO, MR. Database, you can not have my business logic. Your procedural ambitions will bear no fruit and you'll have to pry that logic from my dead, cold object-oriented hands ...I want a single layer of cleverness: My domain model.” [55]

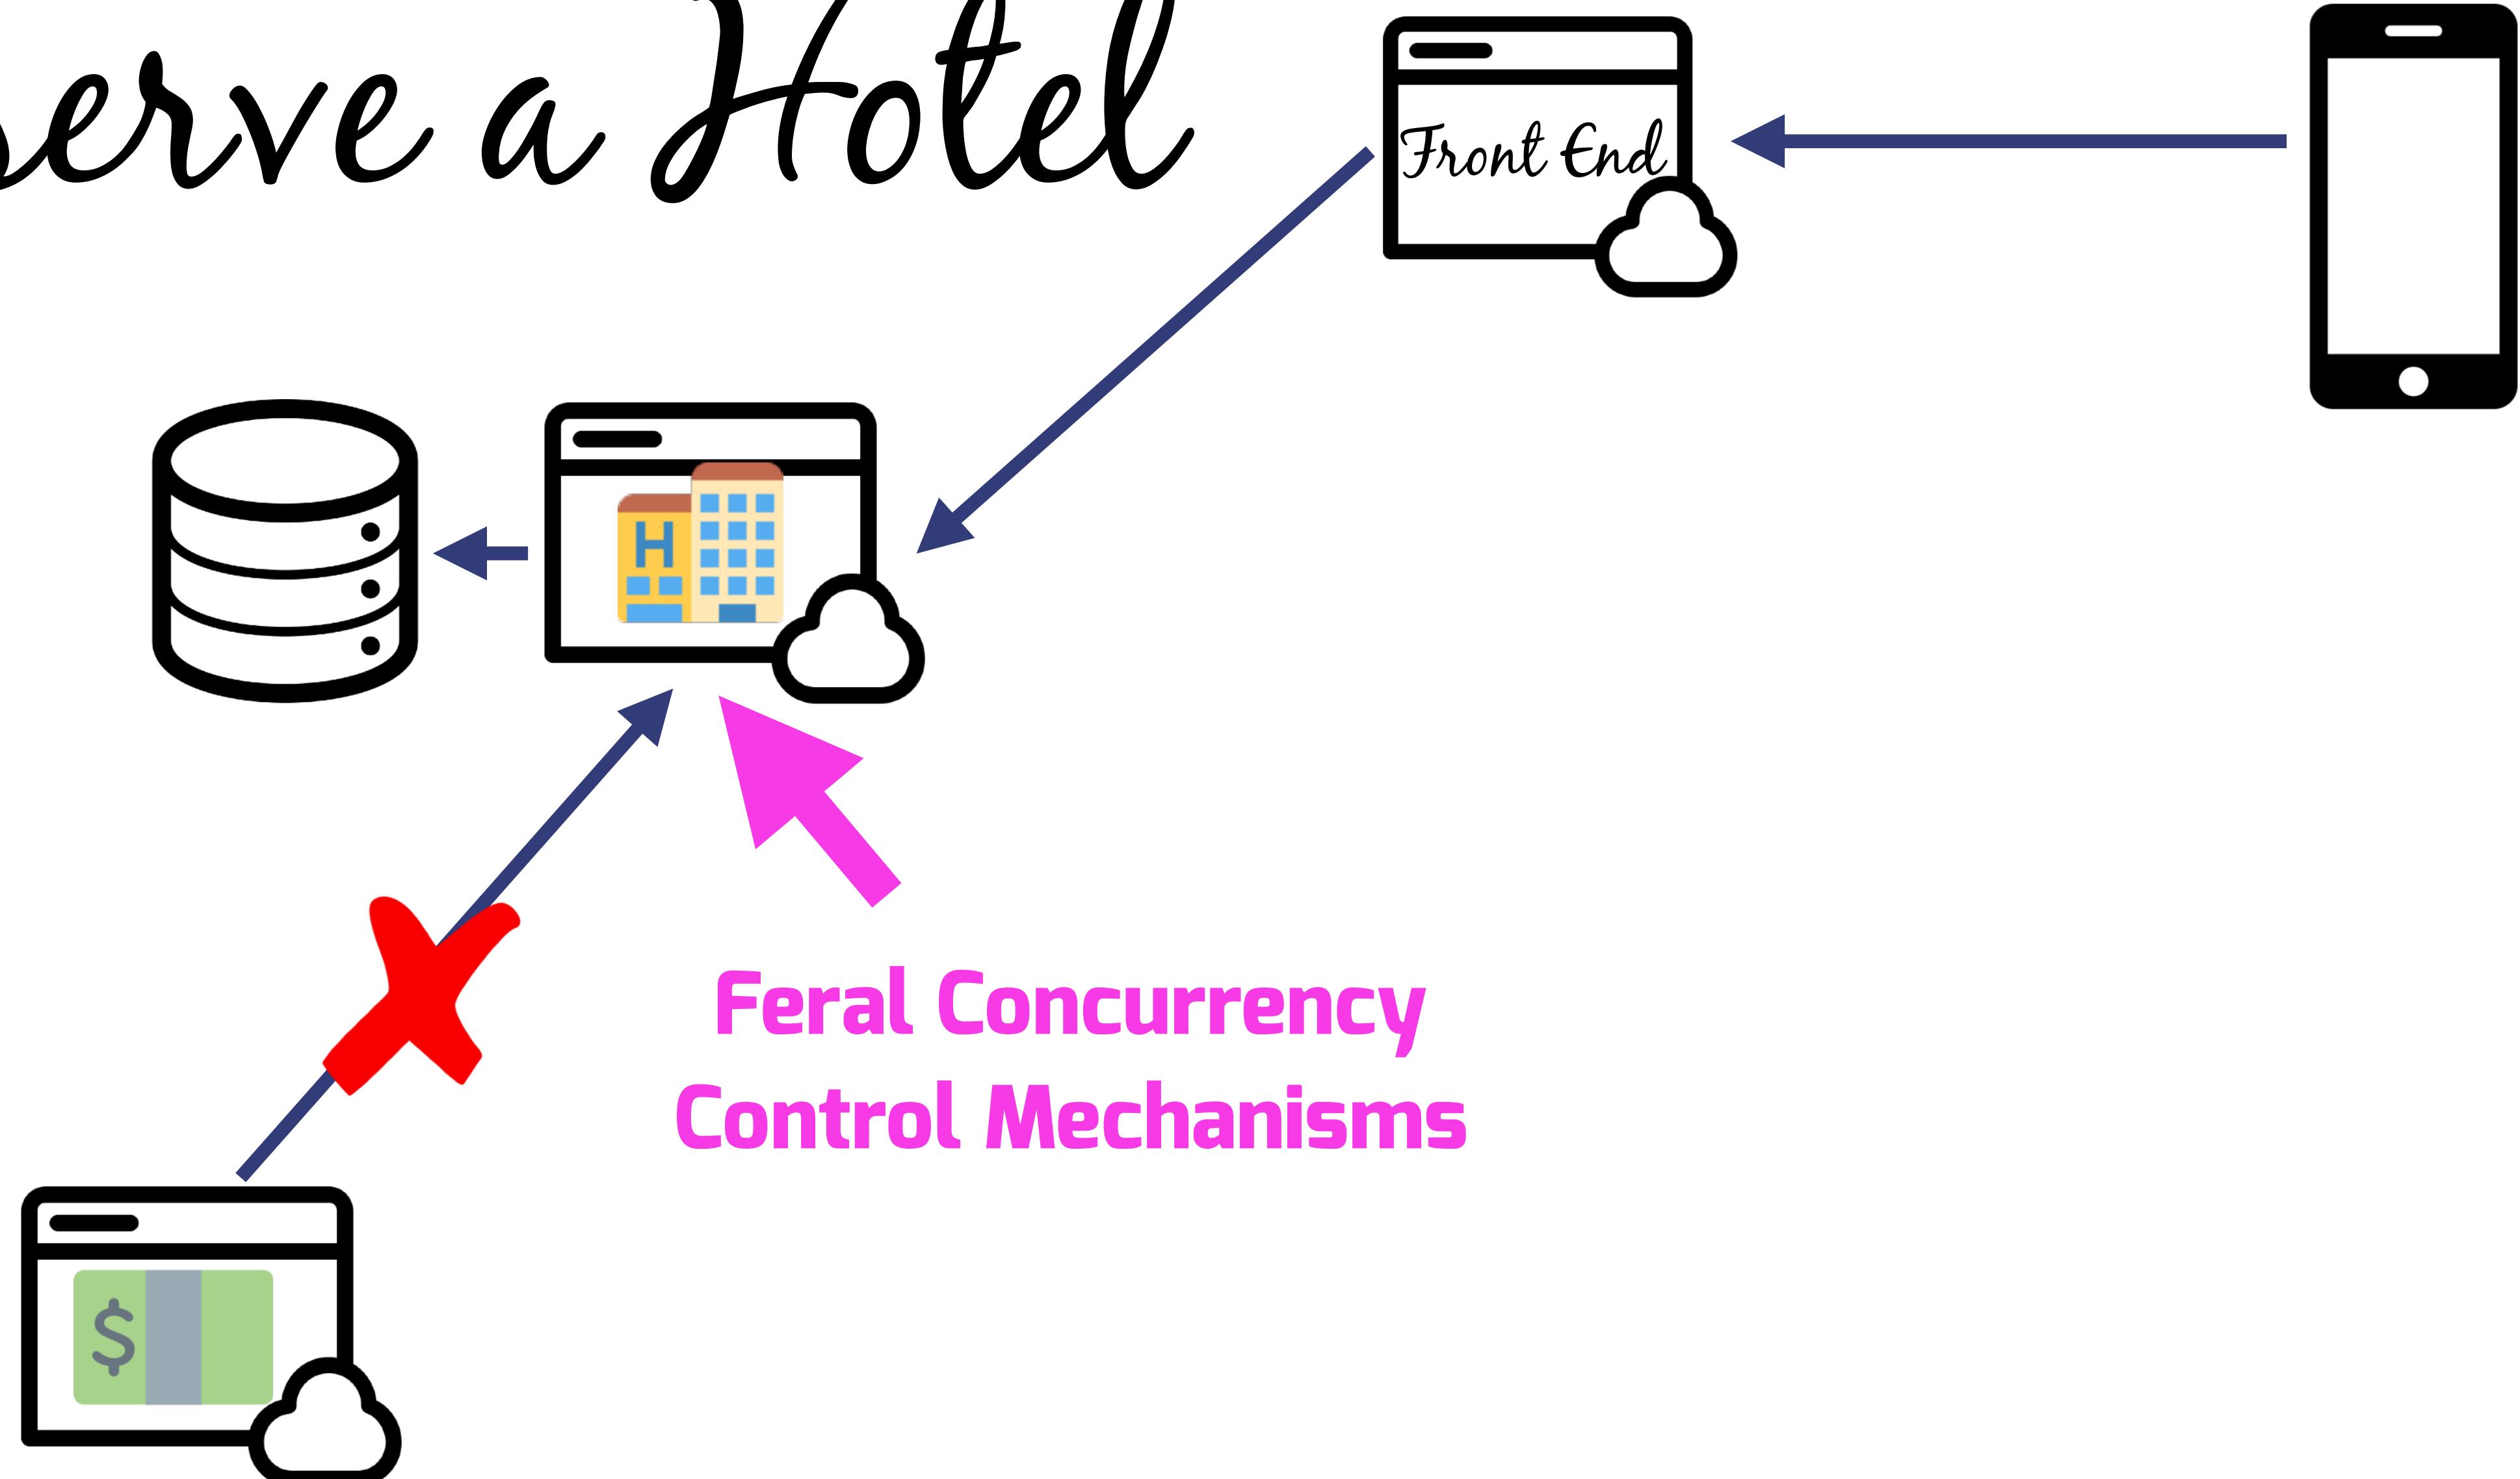
Reserve a Hotel



Reserve a Hotel

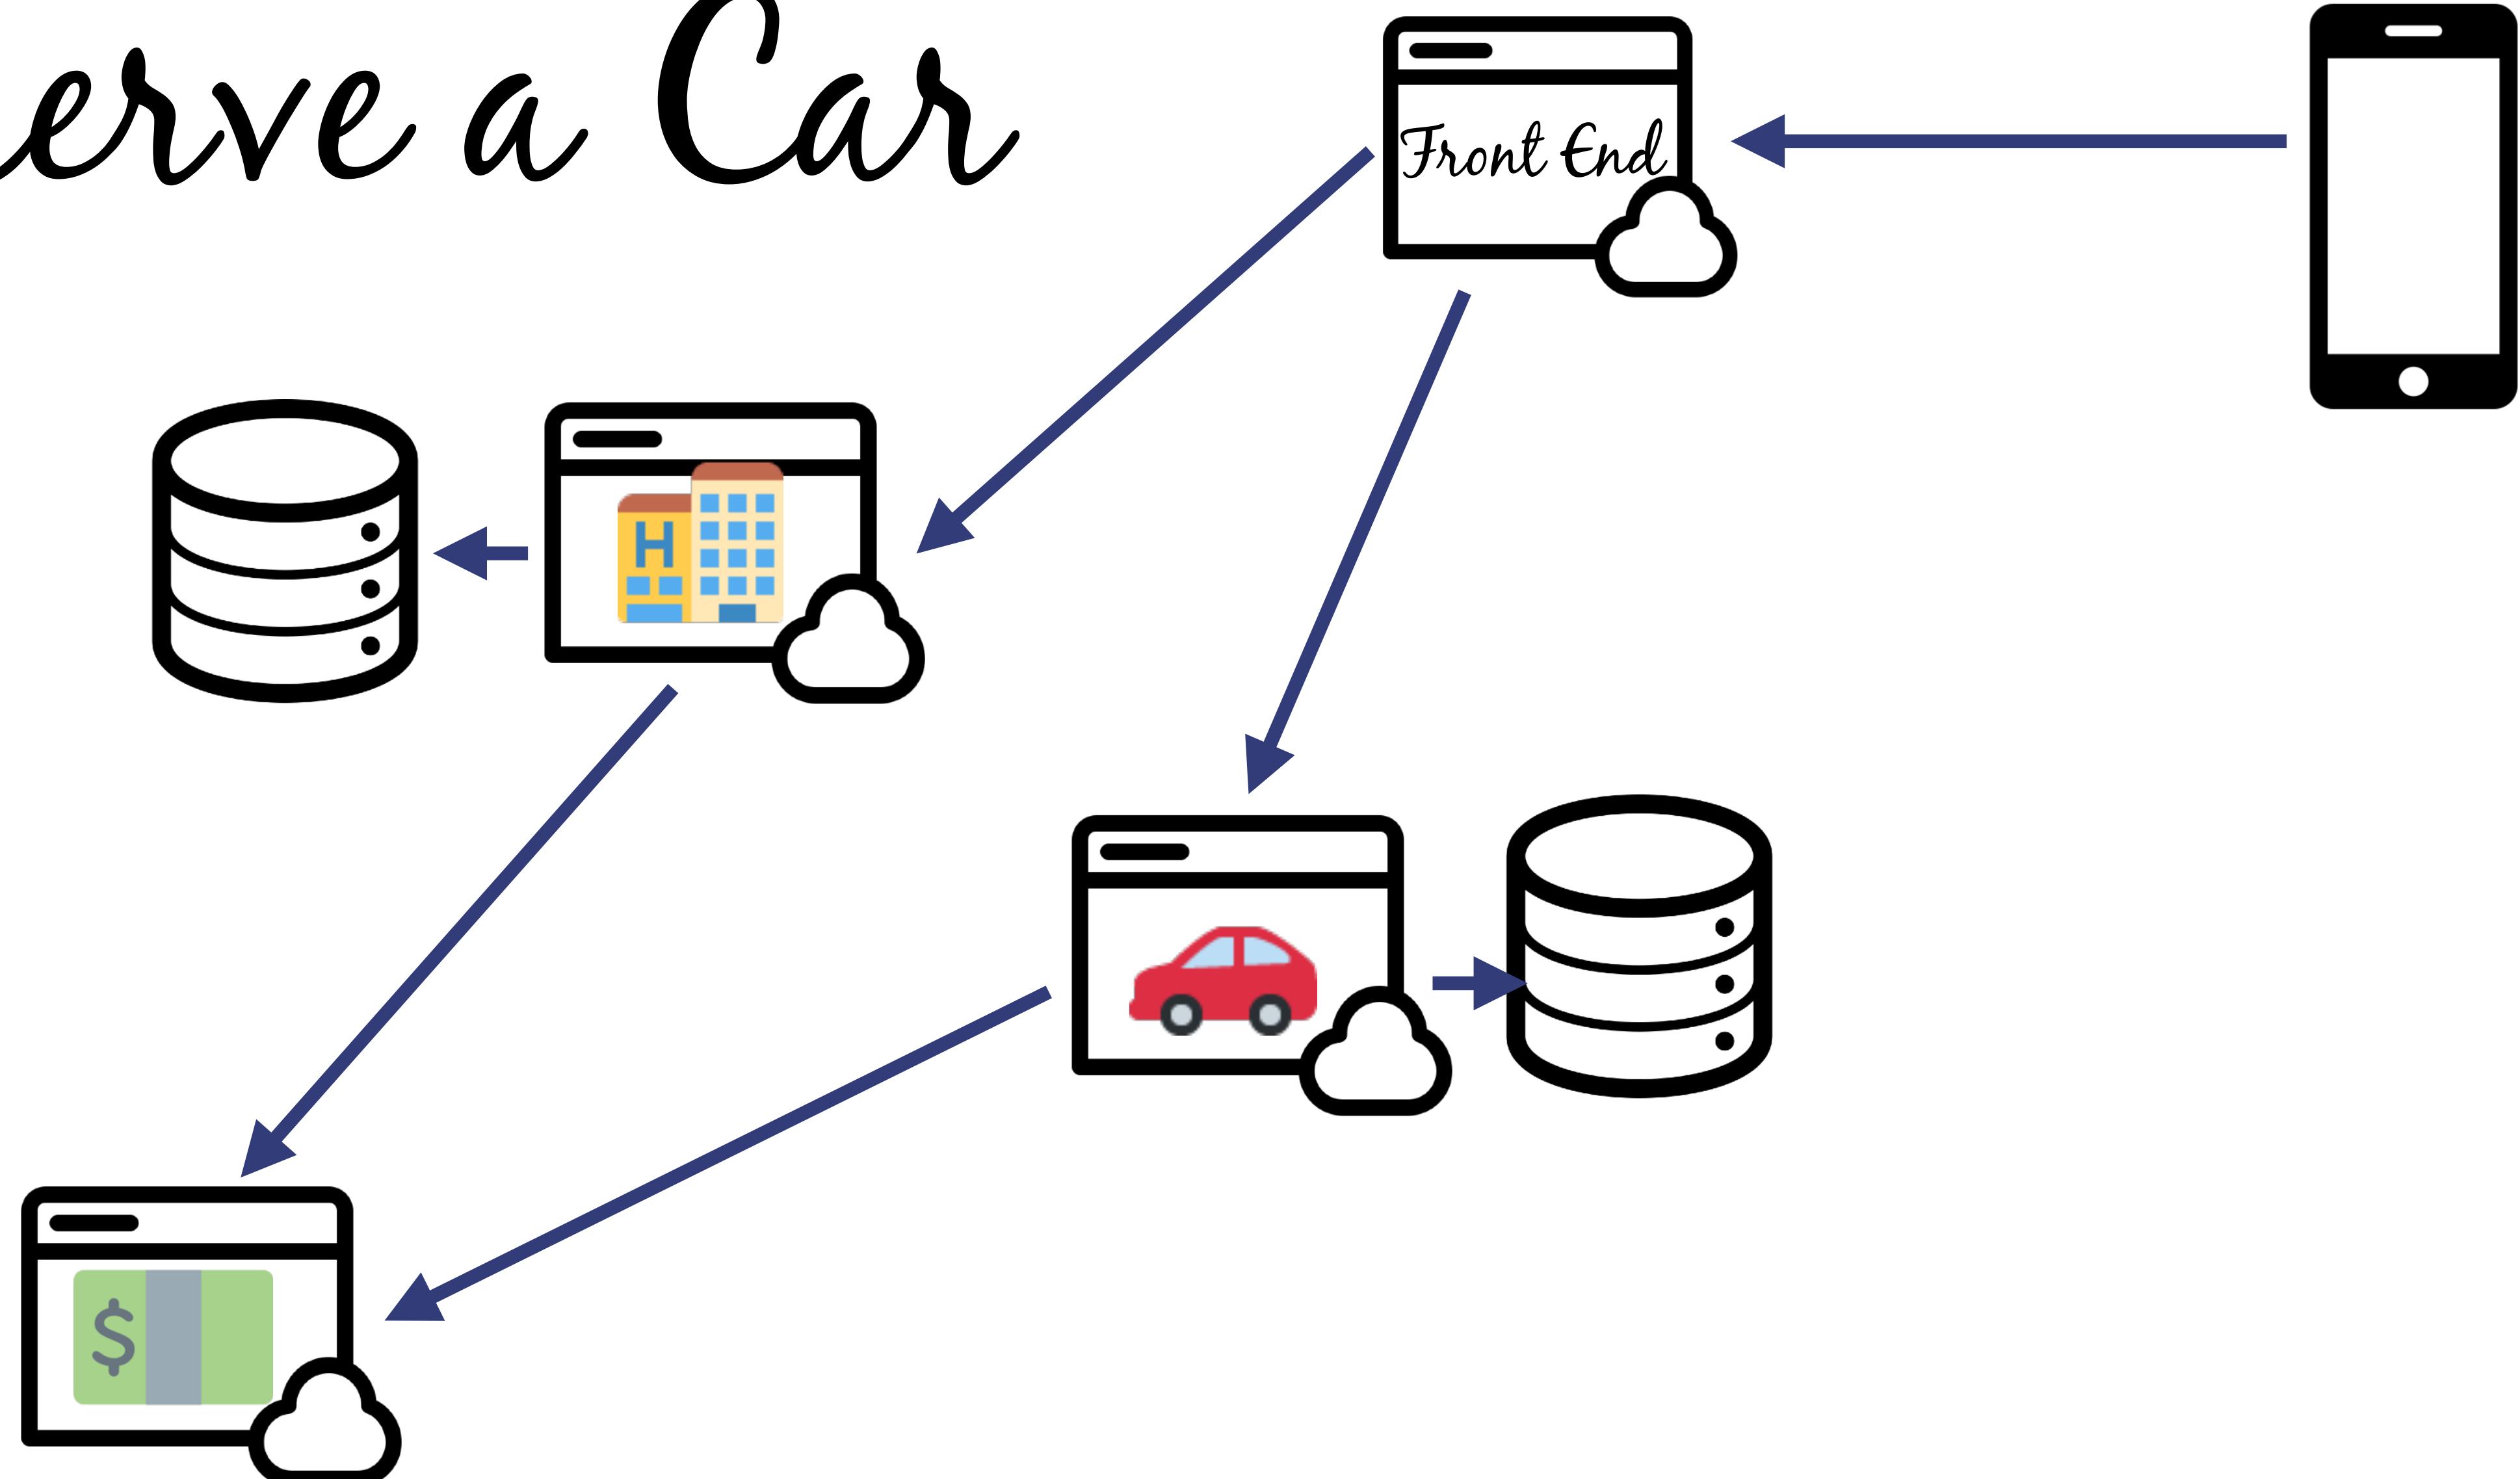


Reserve a Hotel

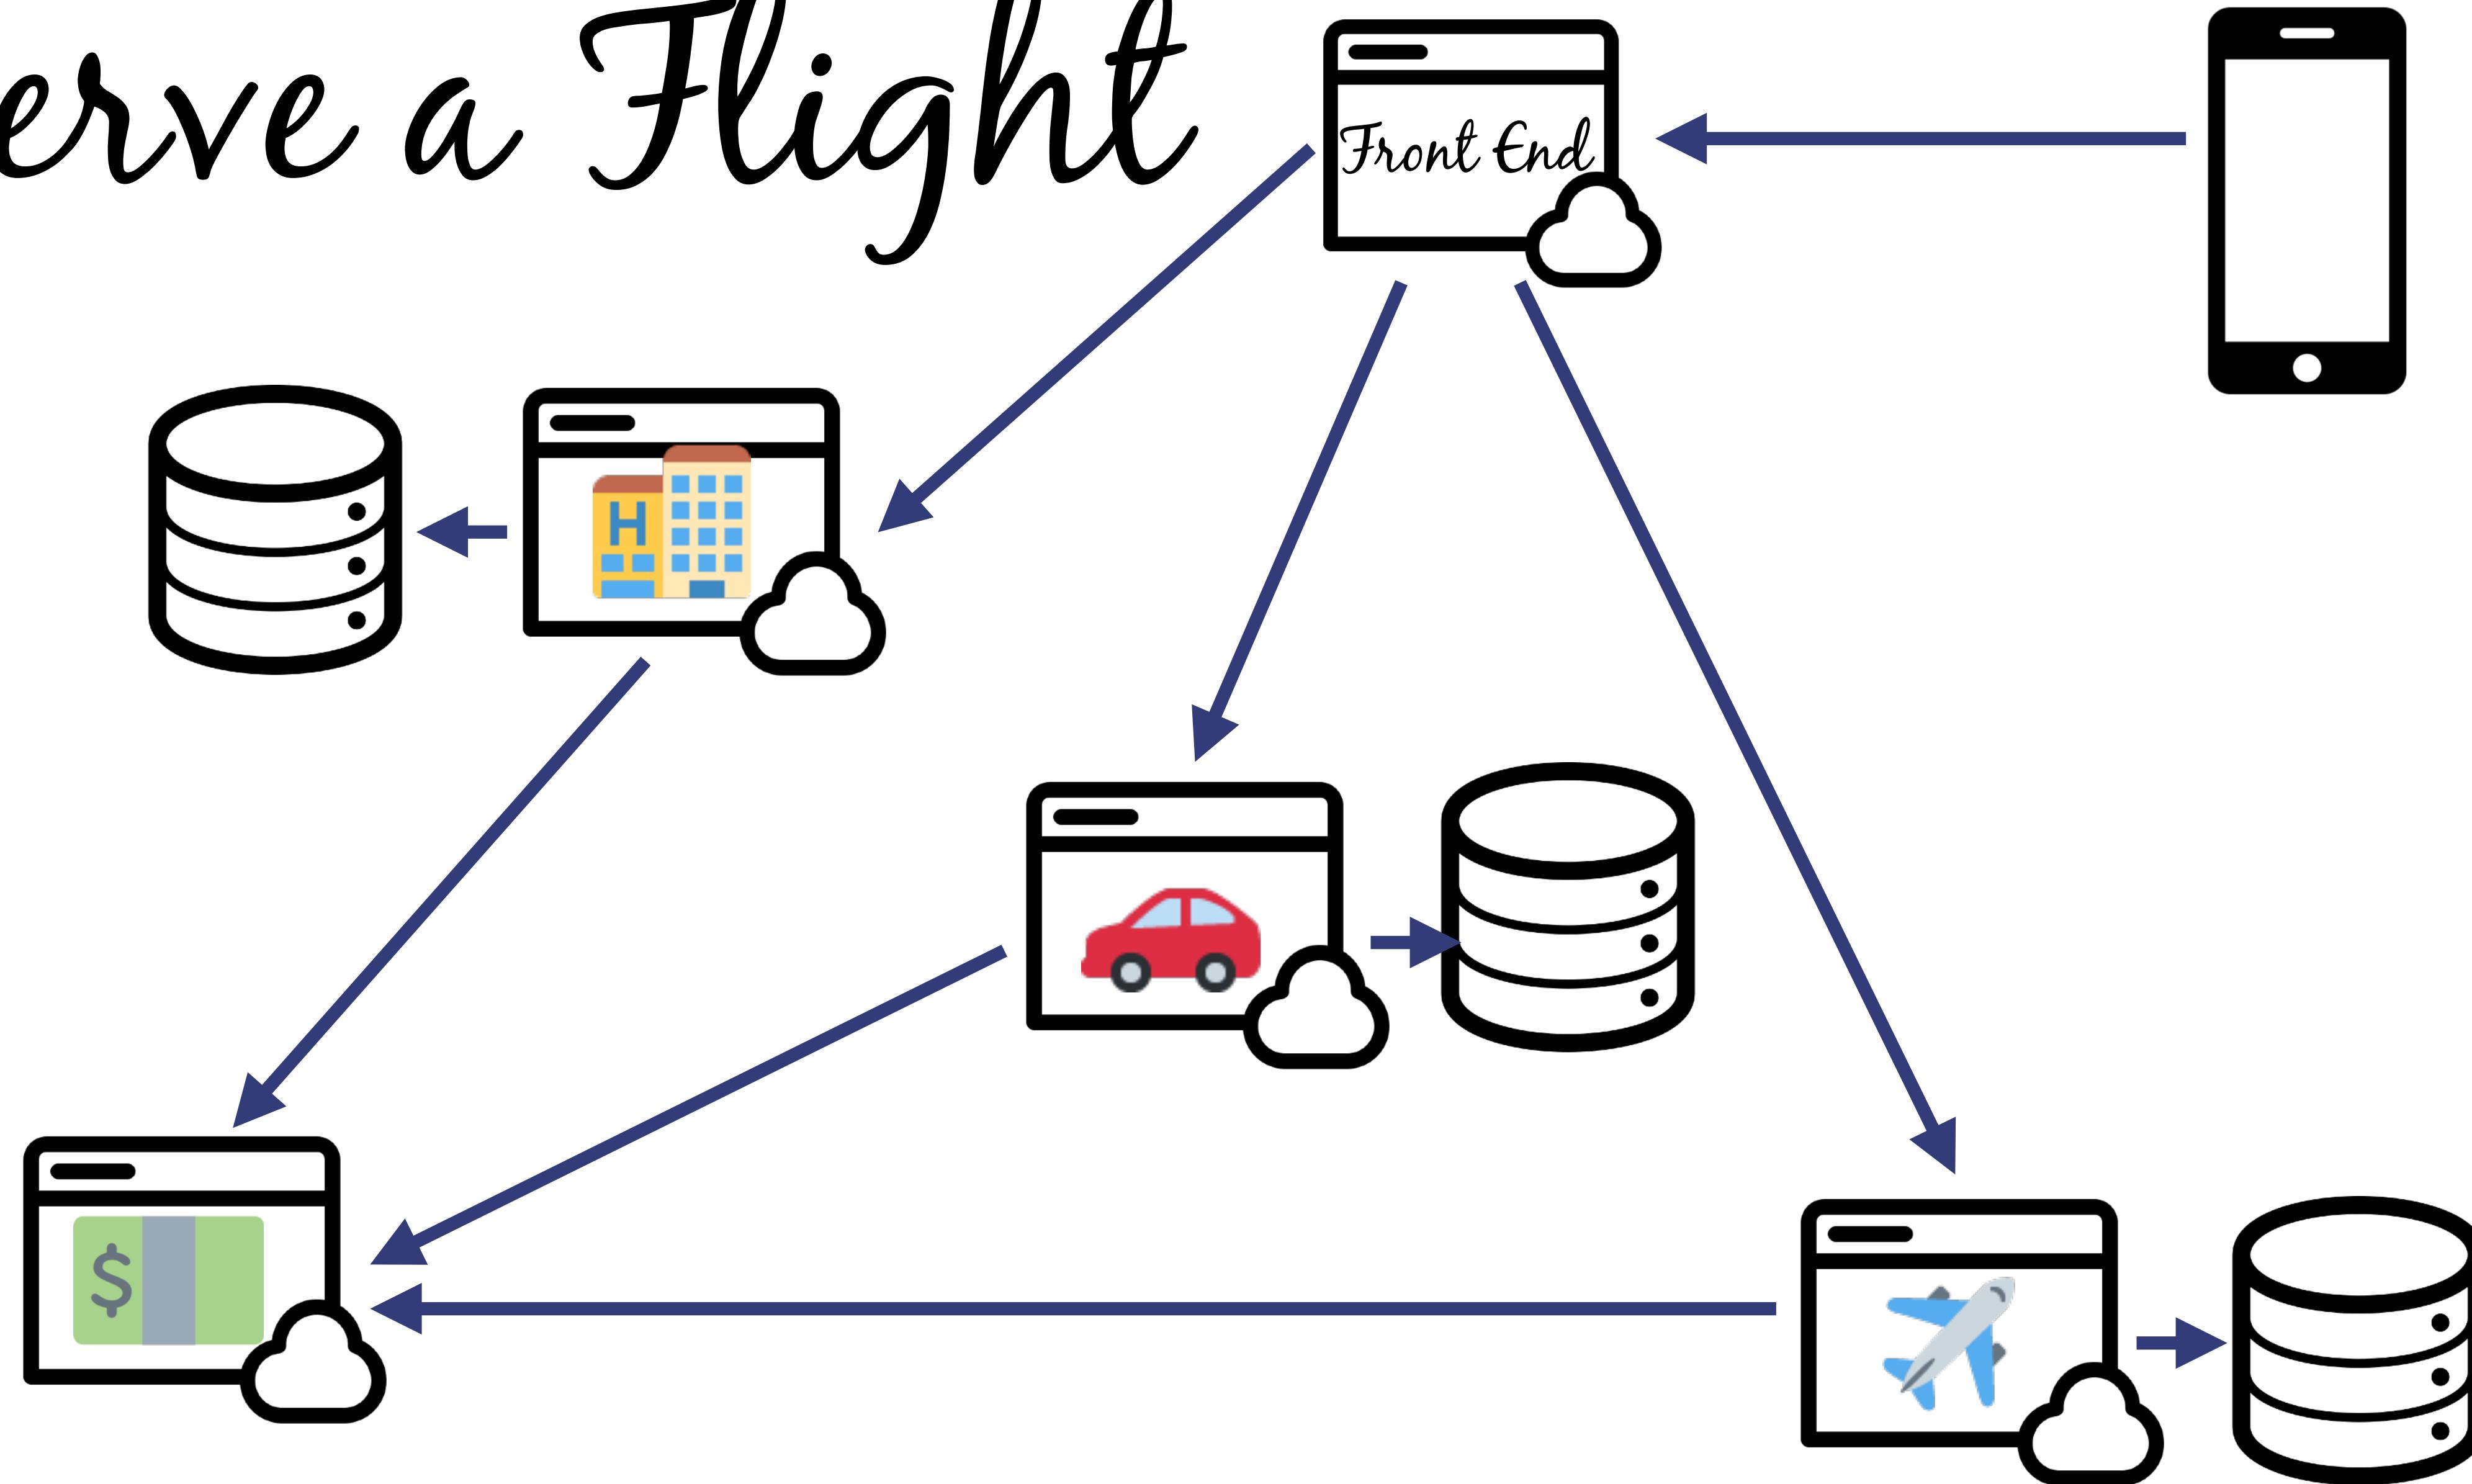


Feral Concurrency Control Mechanisms

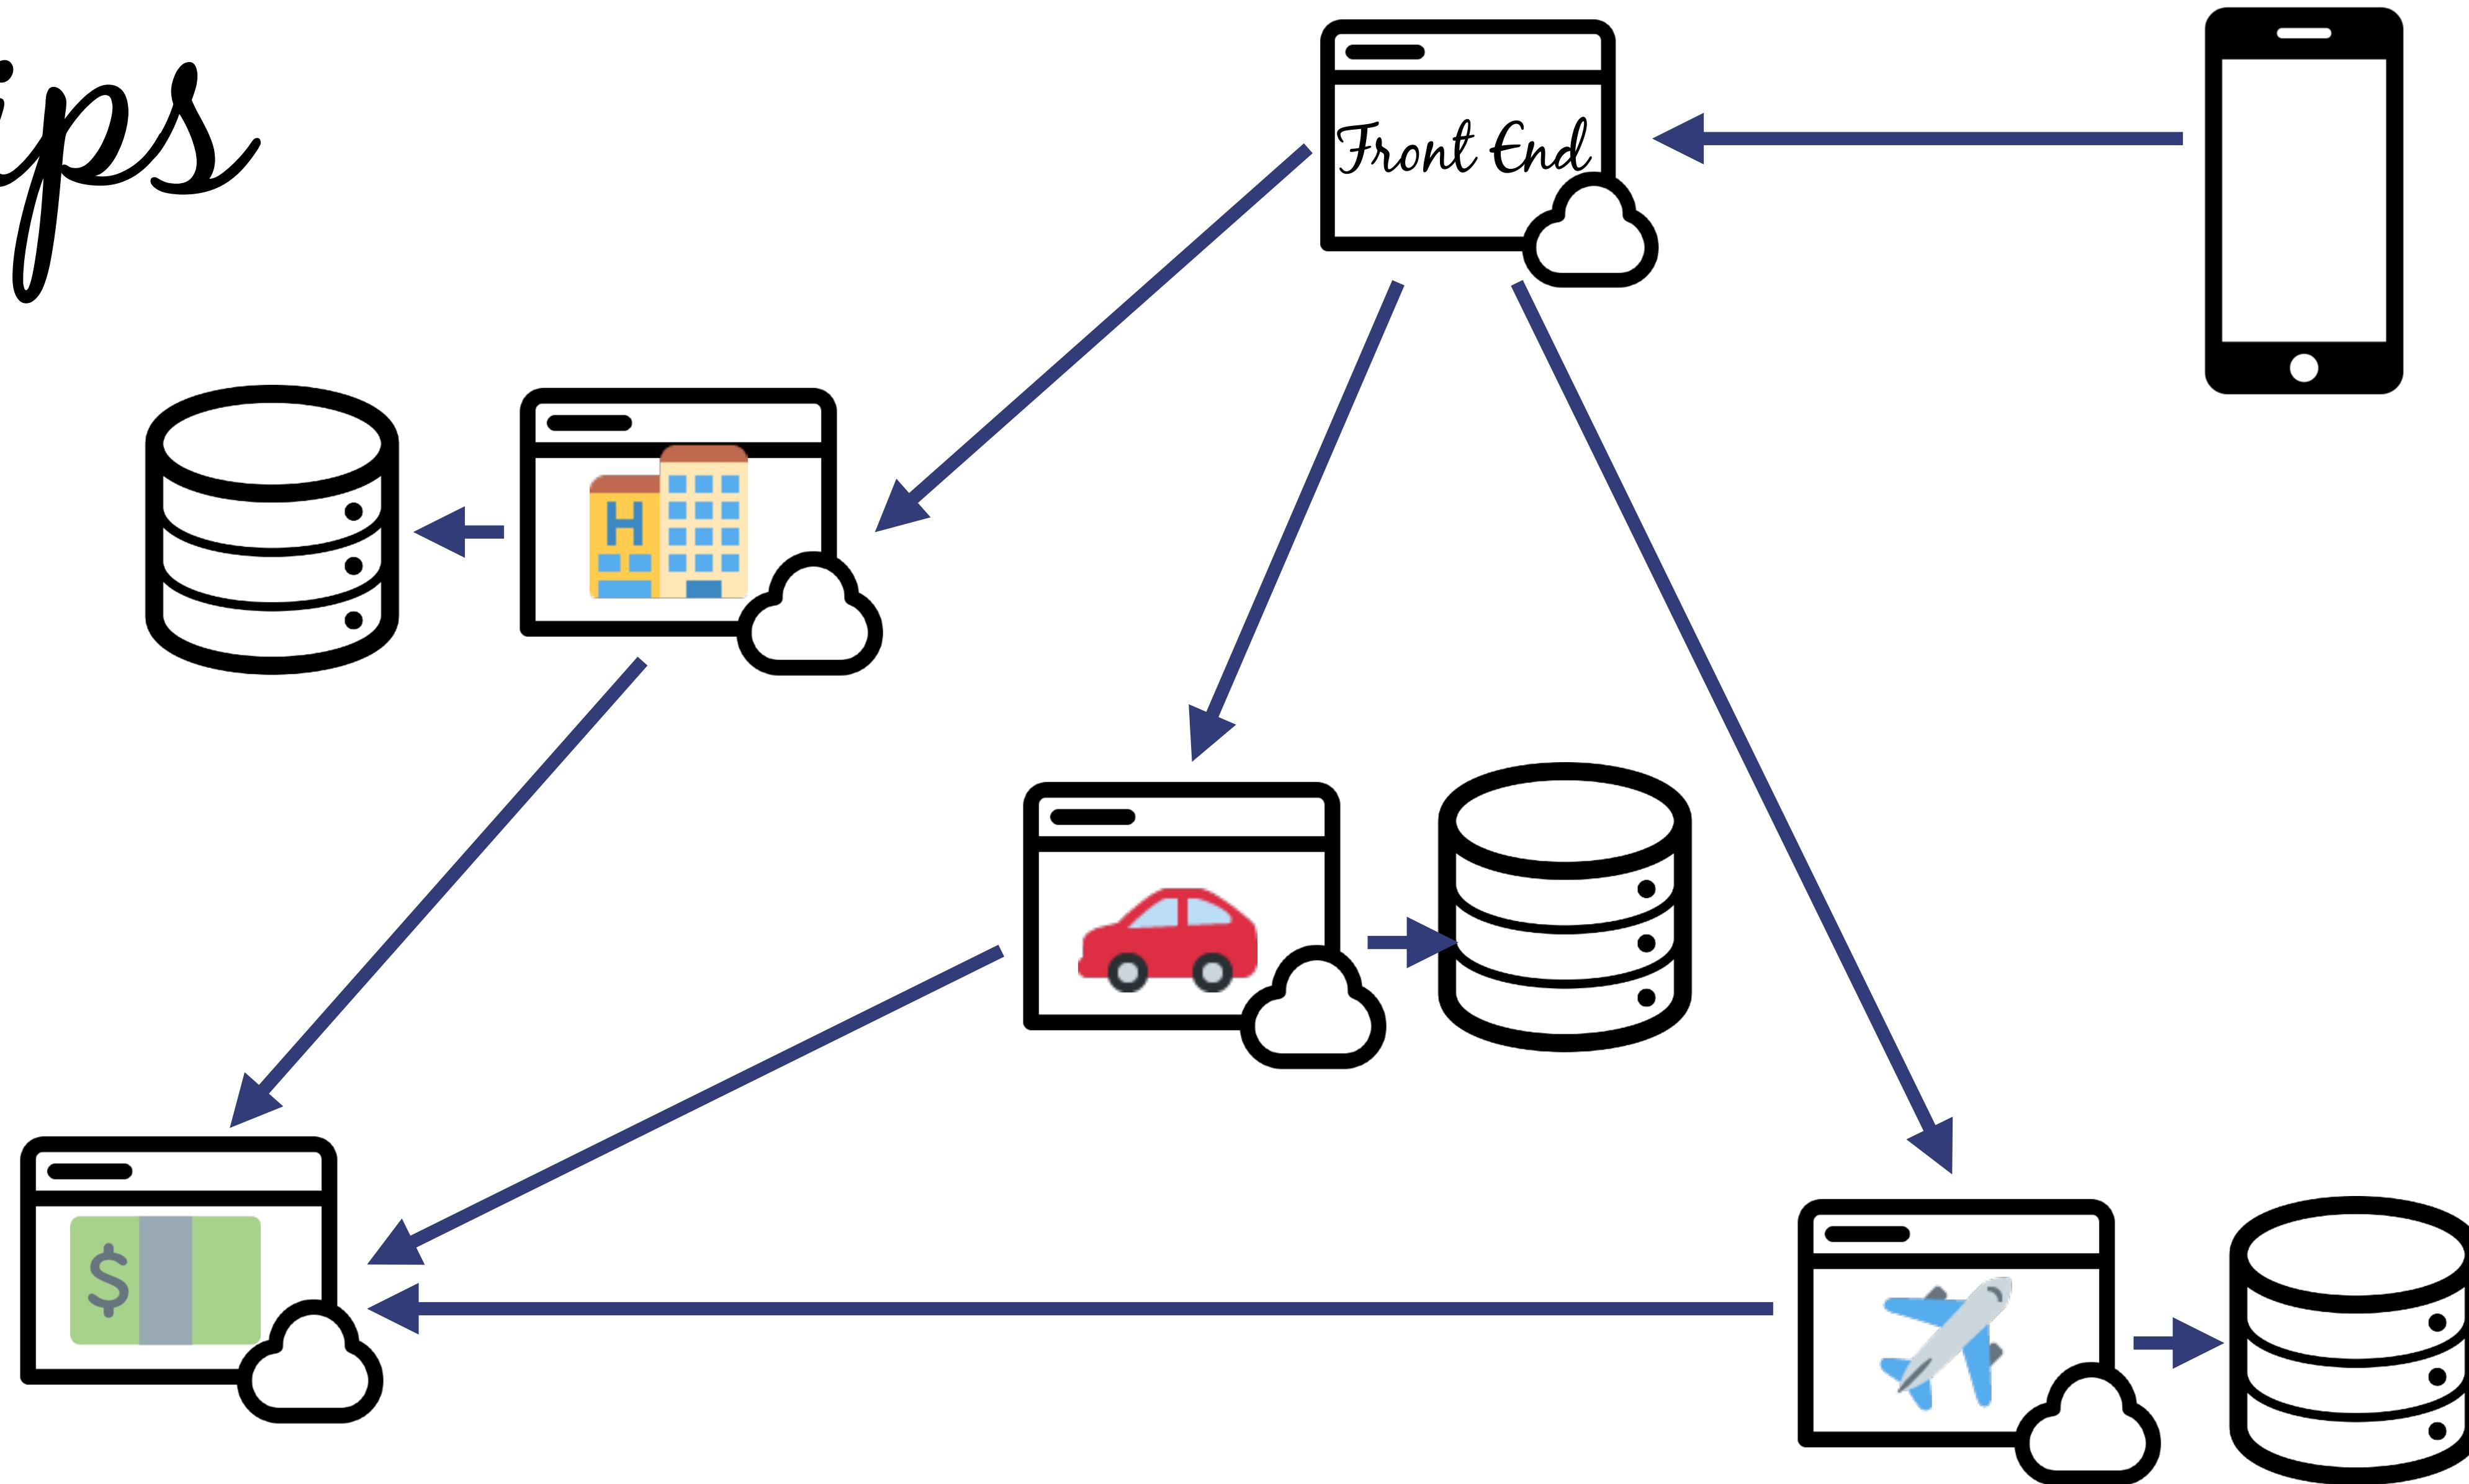
Reserve a Car



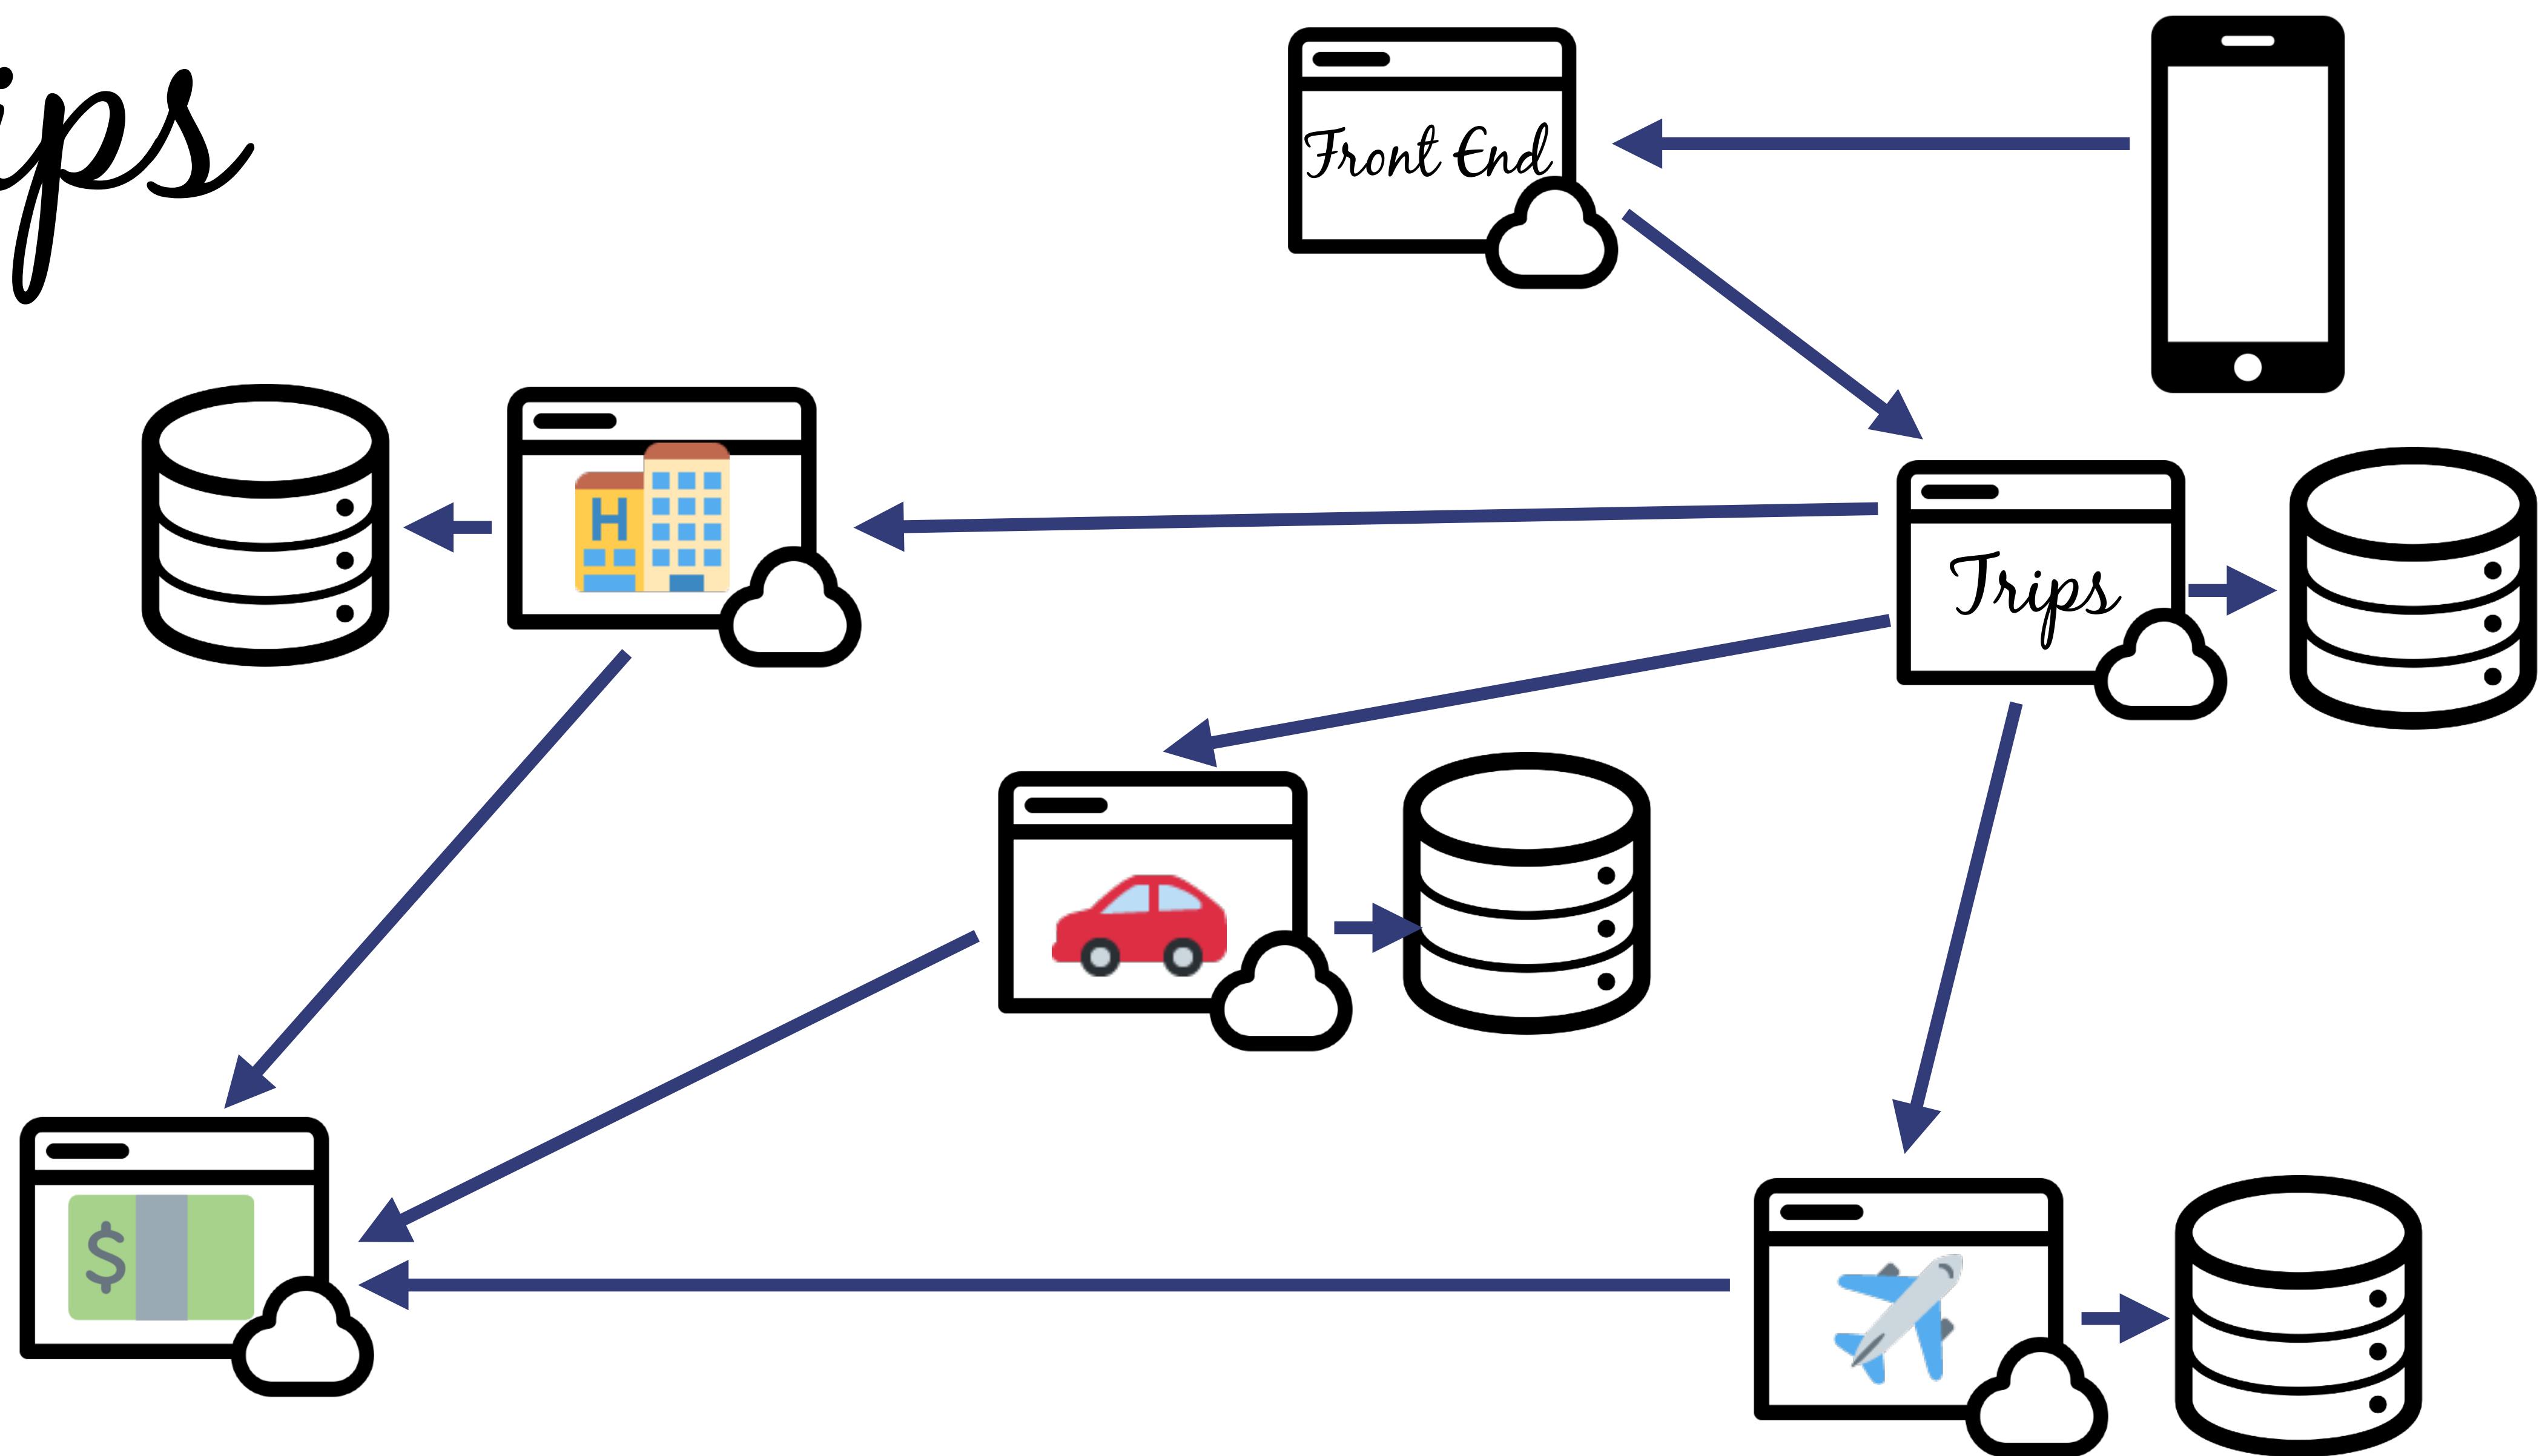
Reserve a Flight



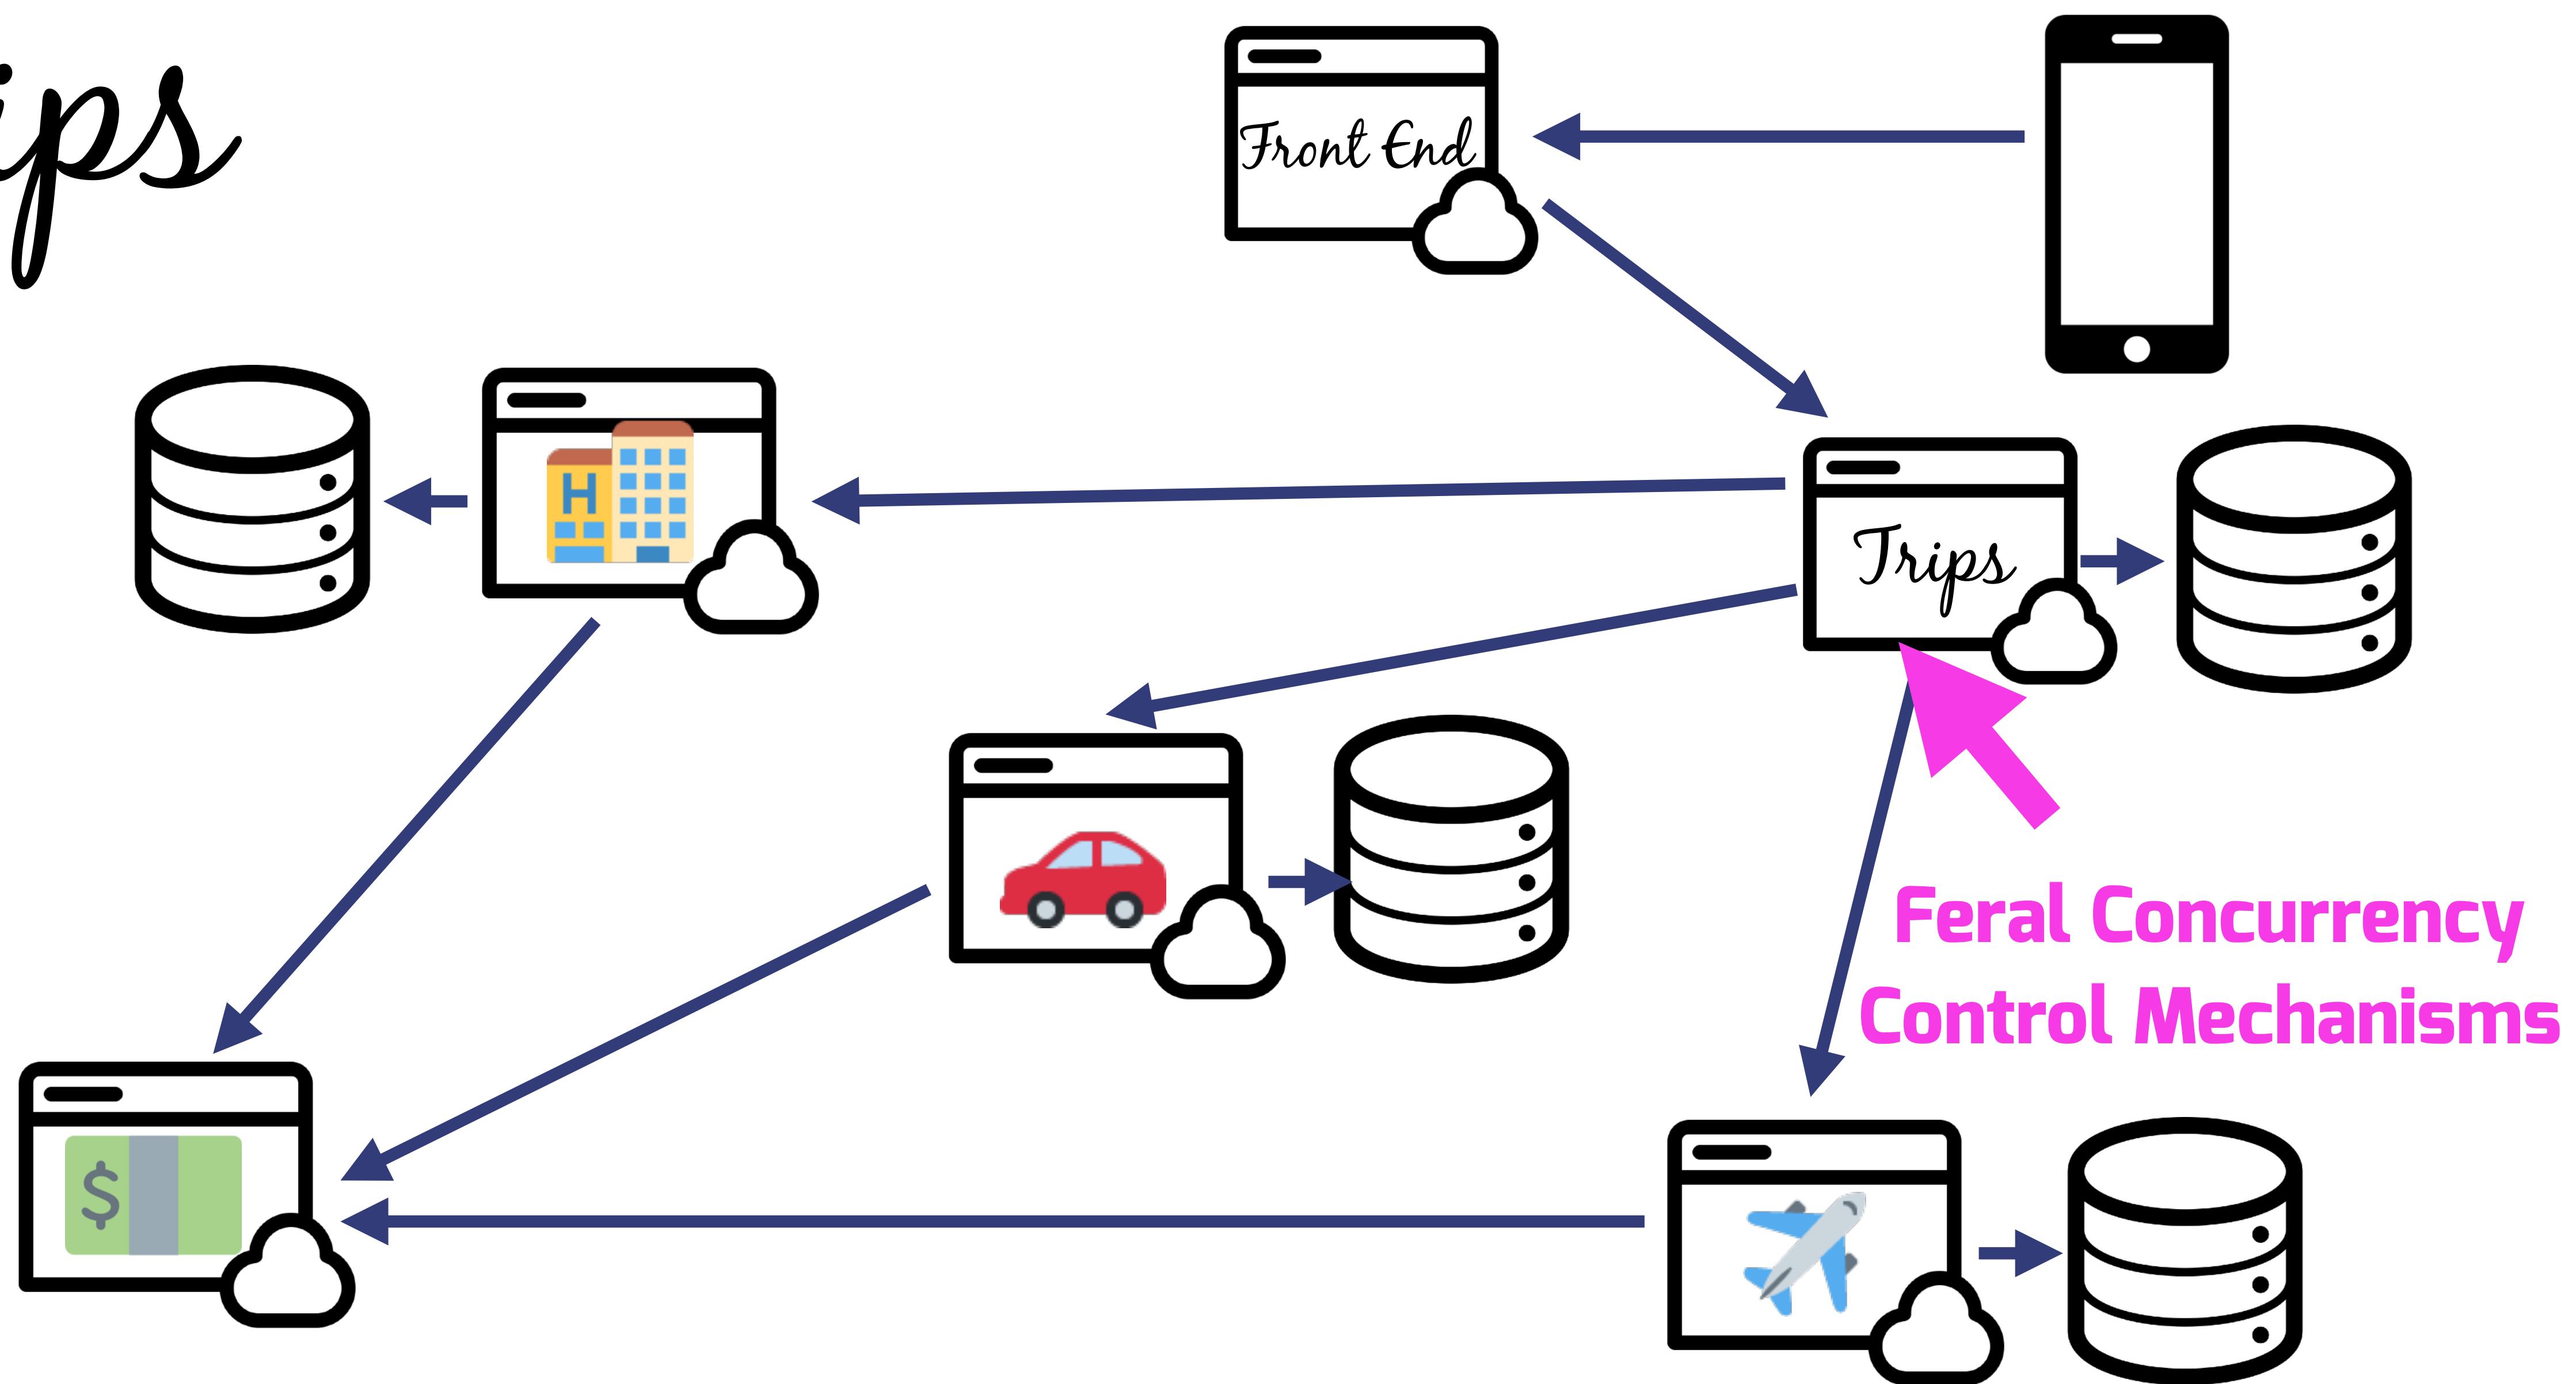
Trips



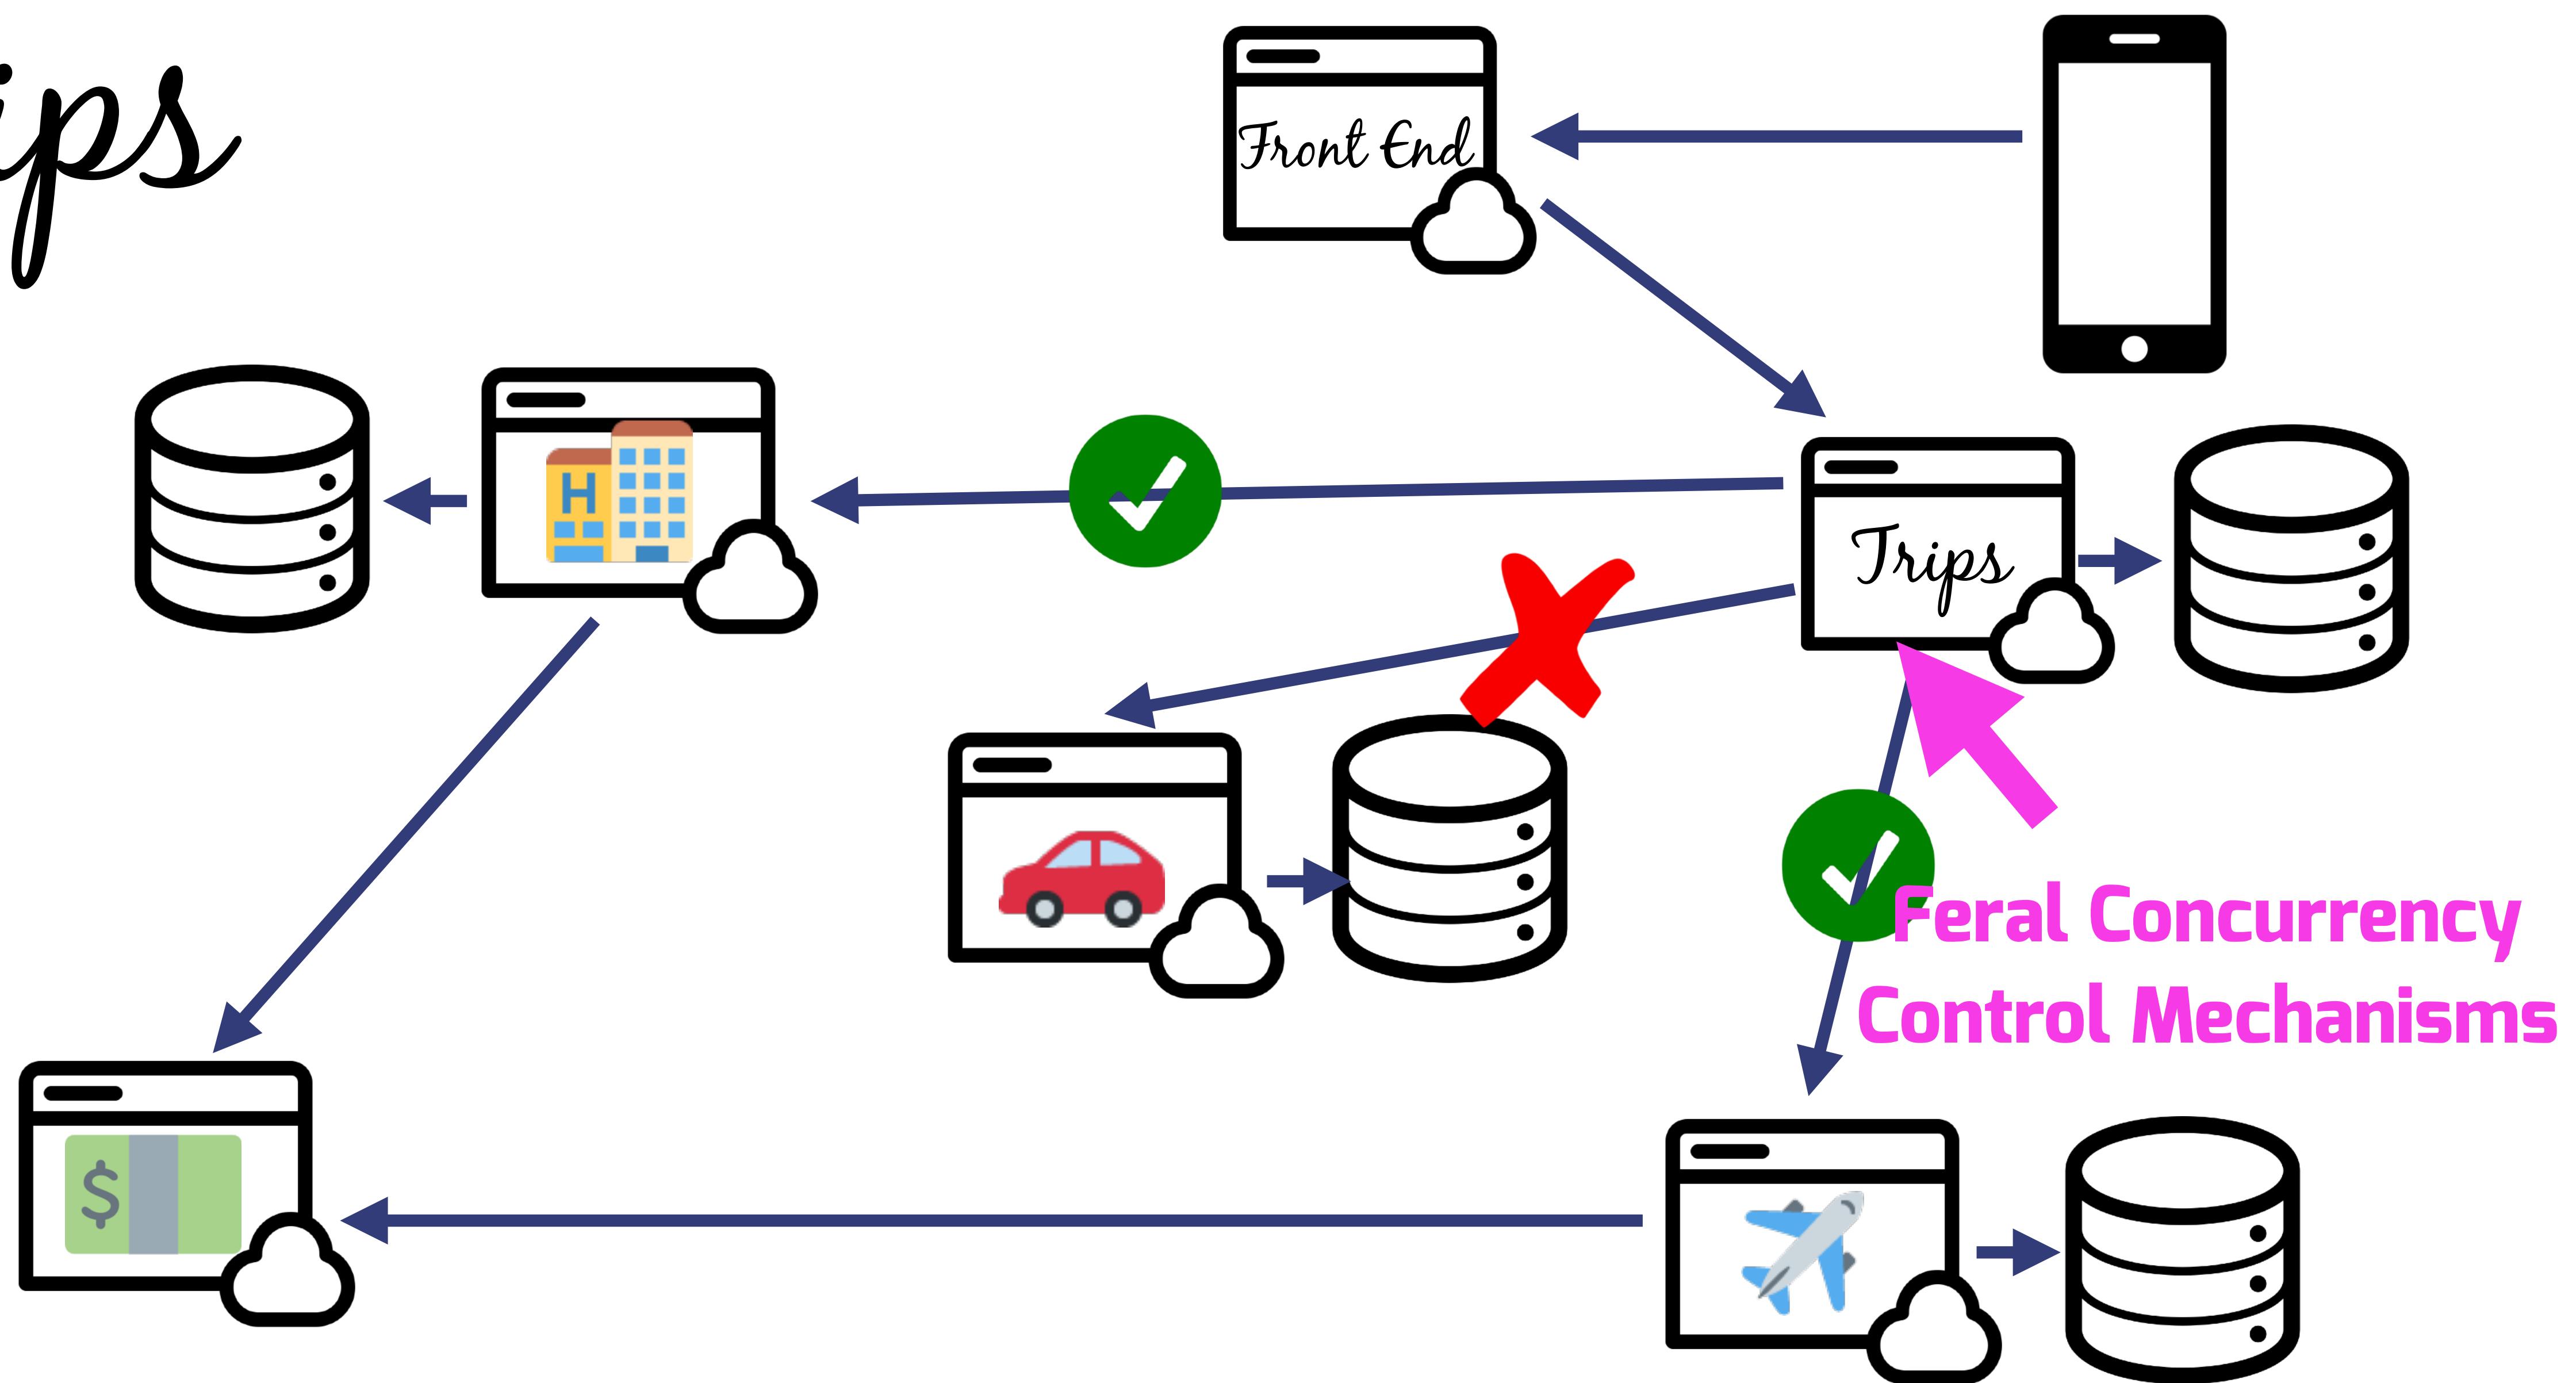
Trips



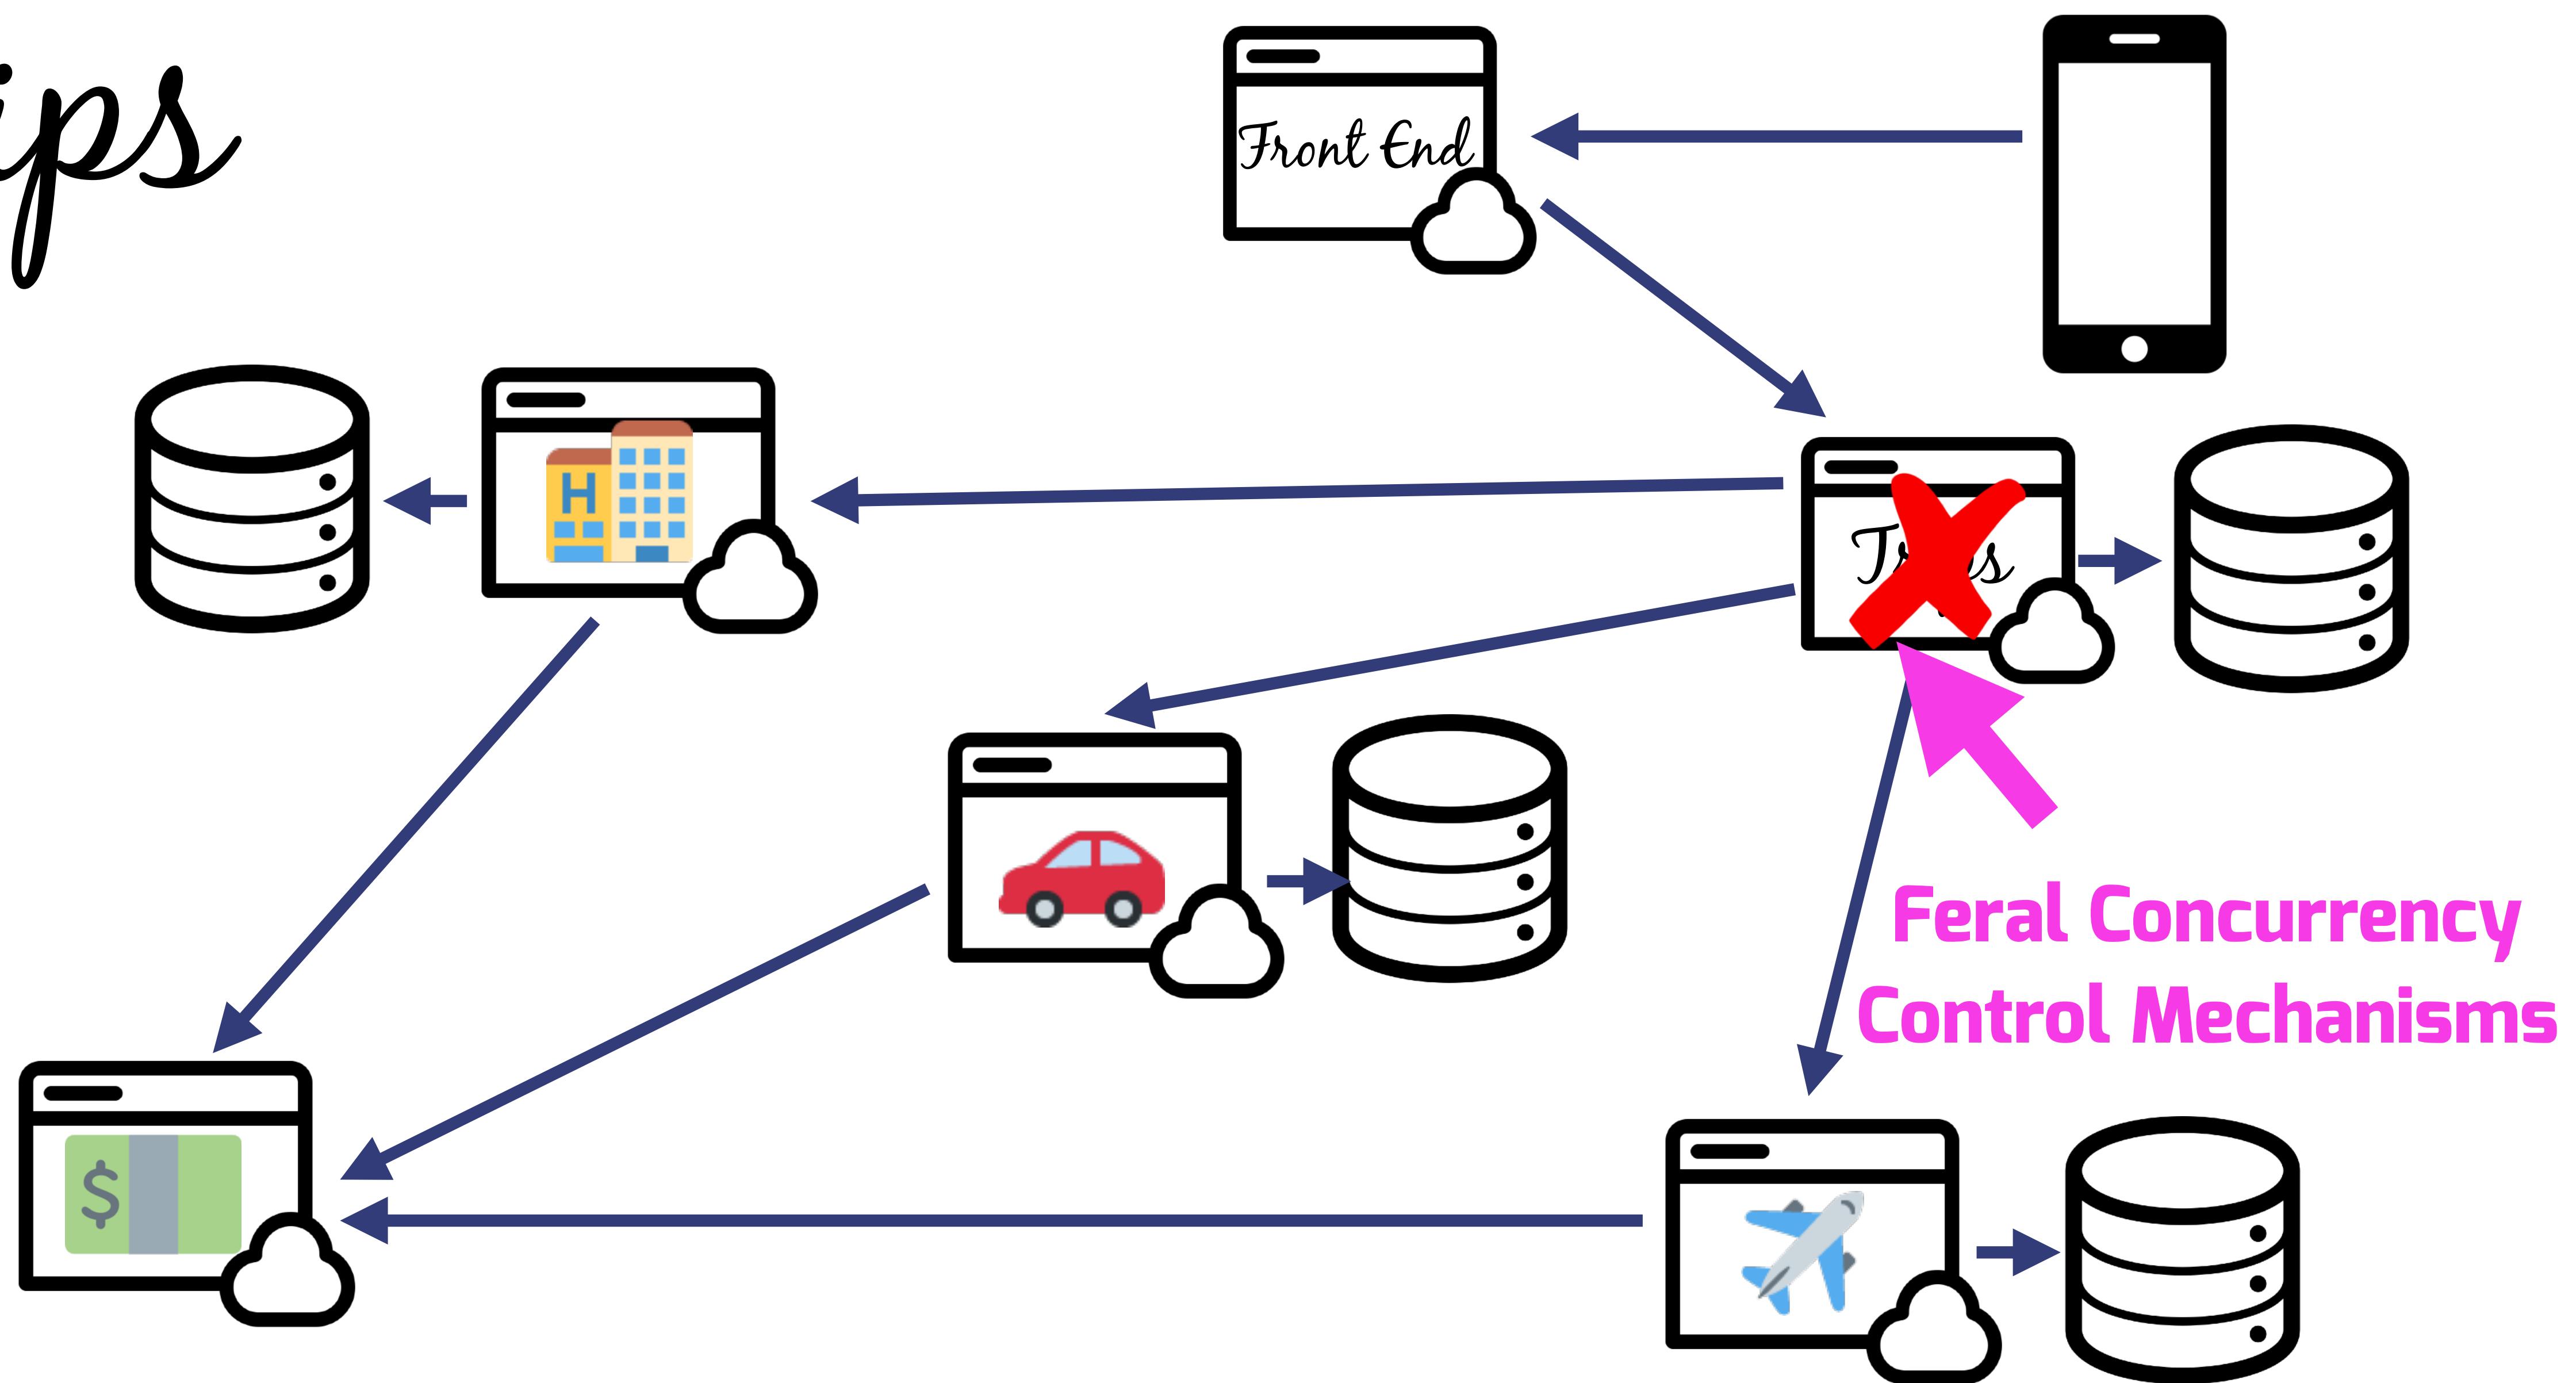
Trips



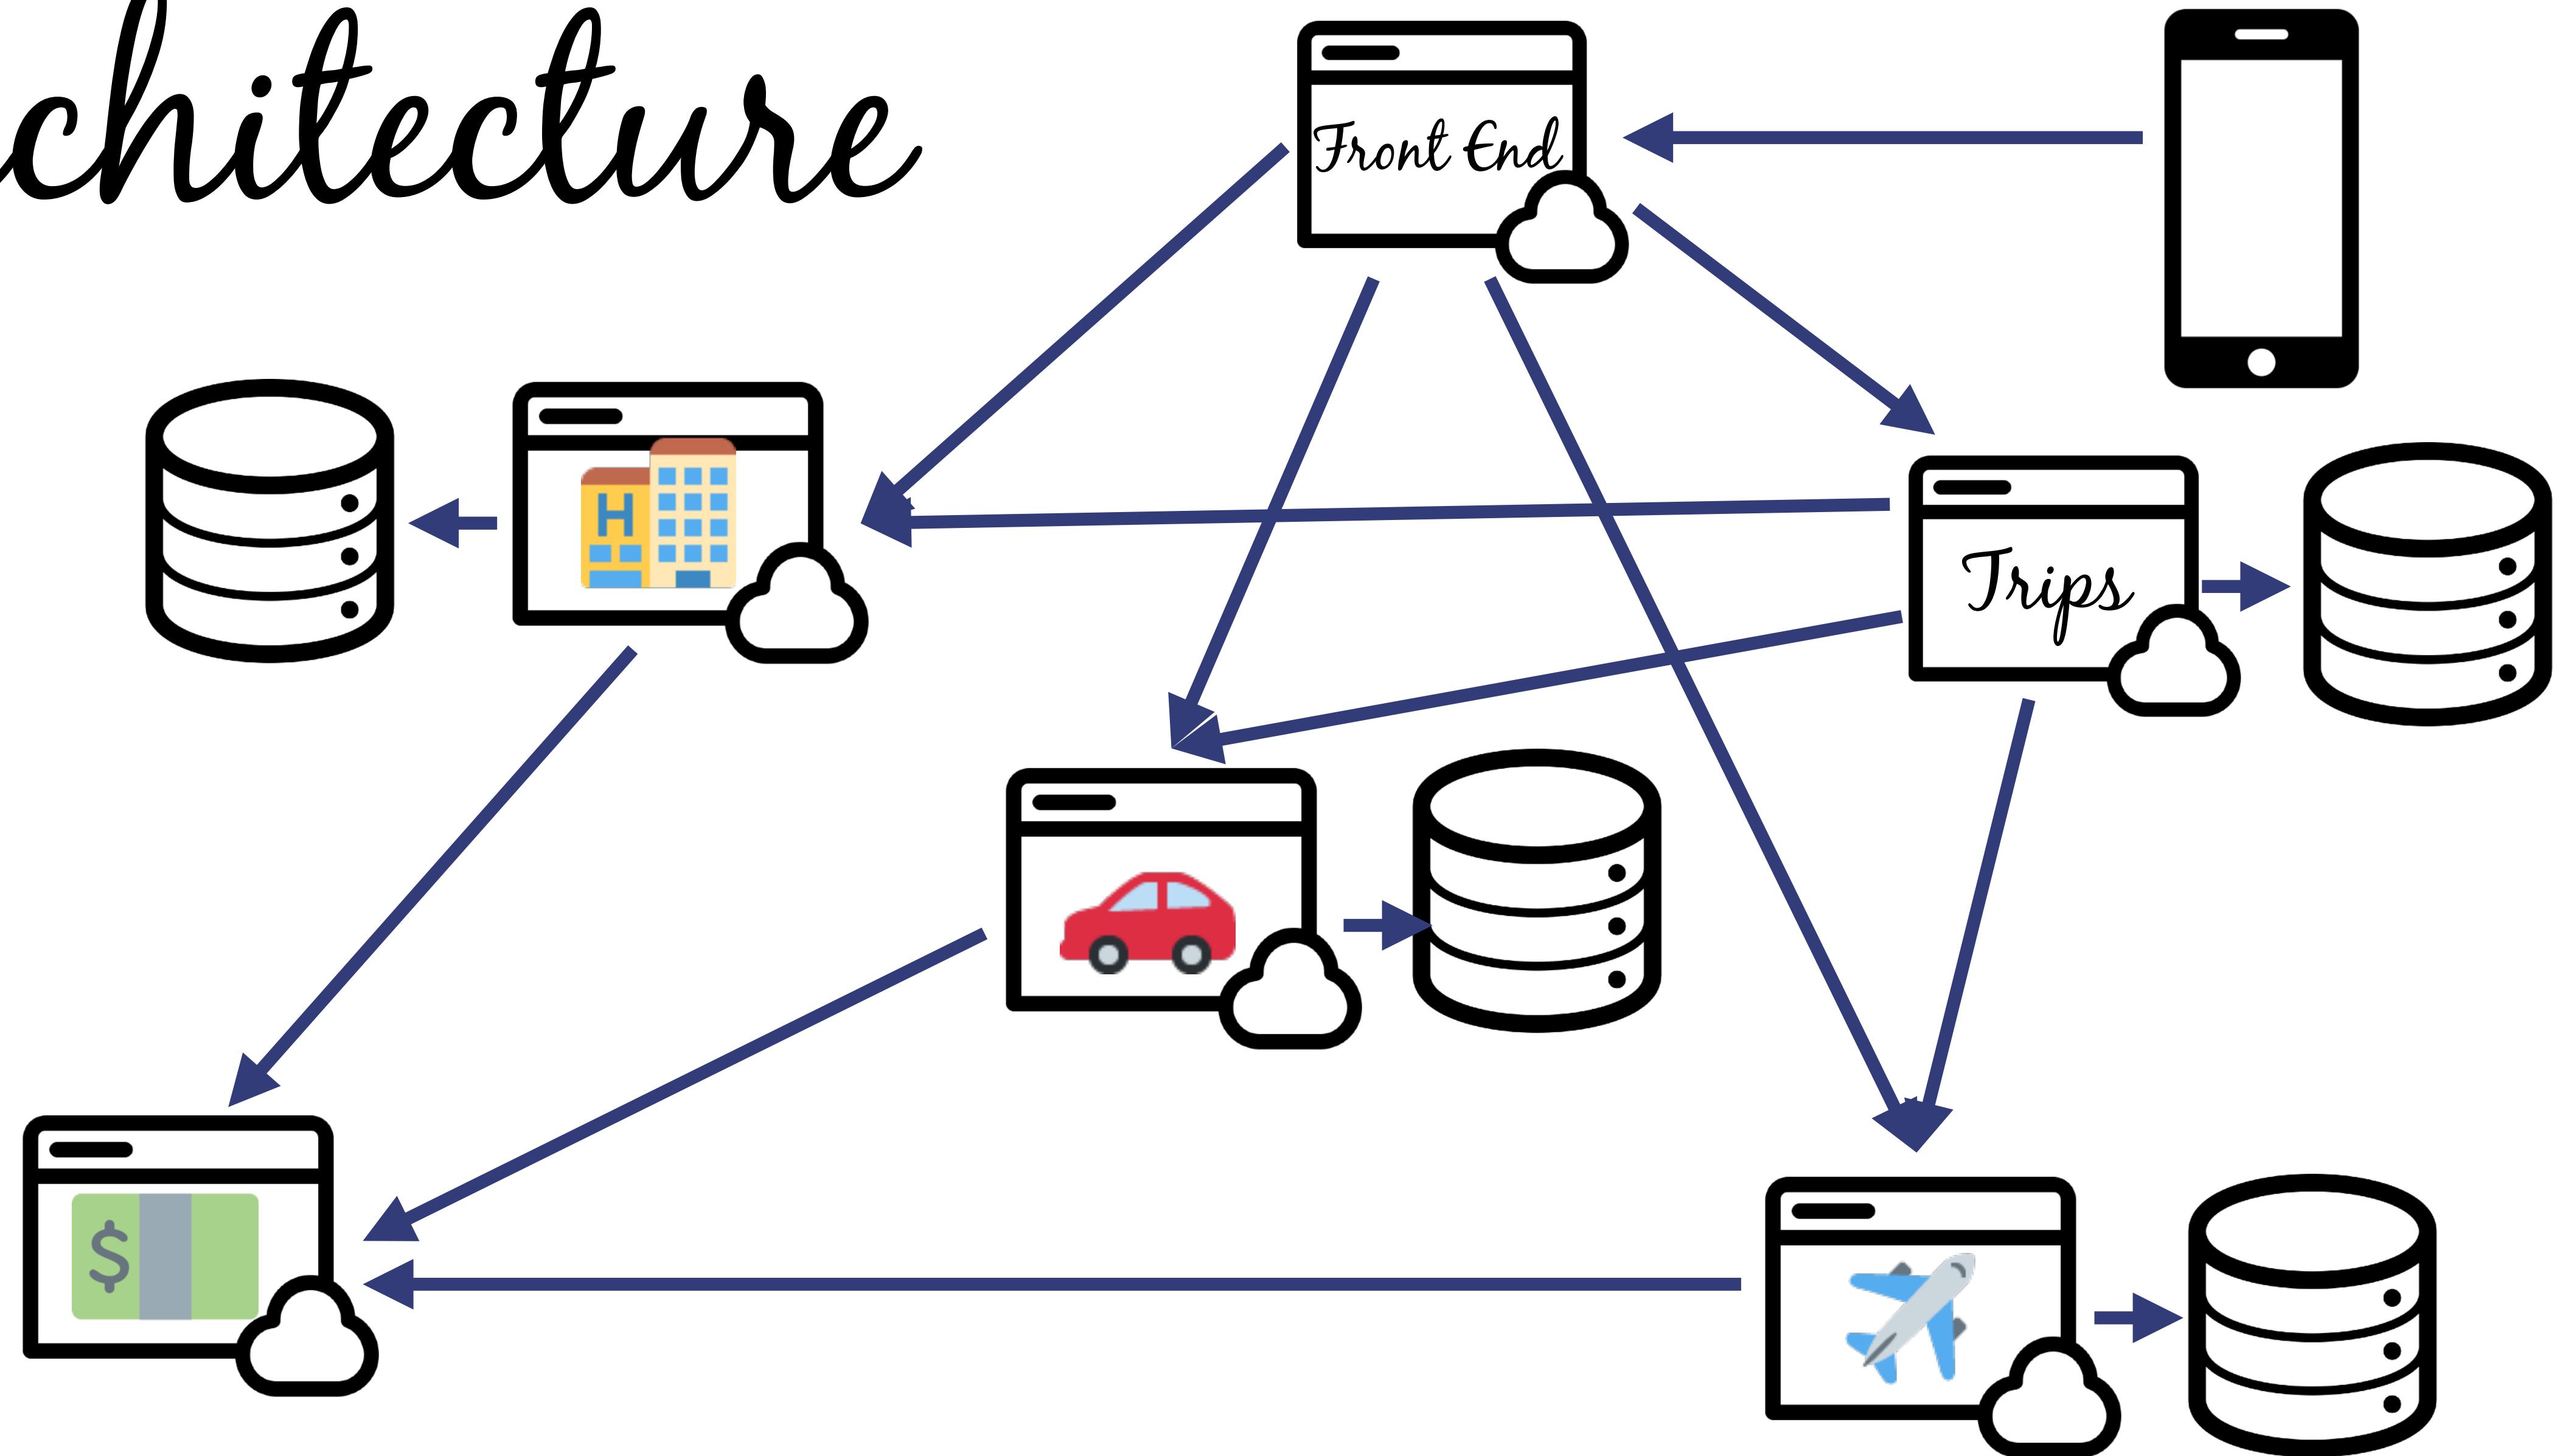
Trips



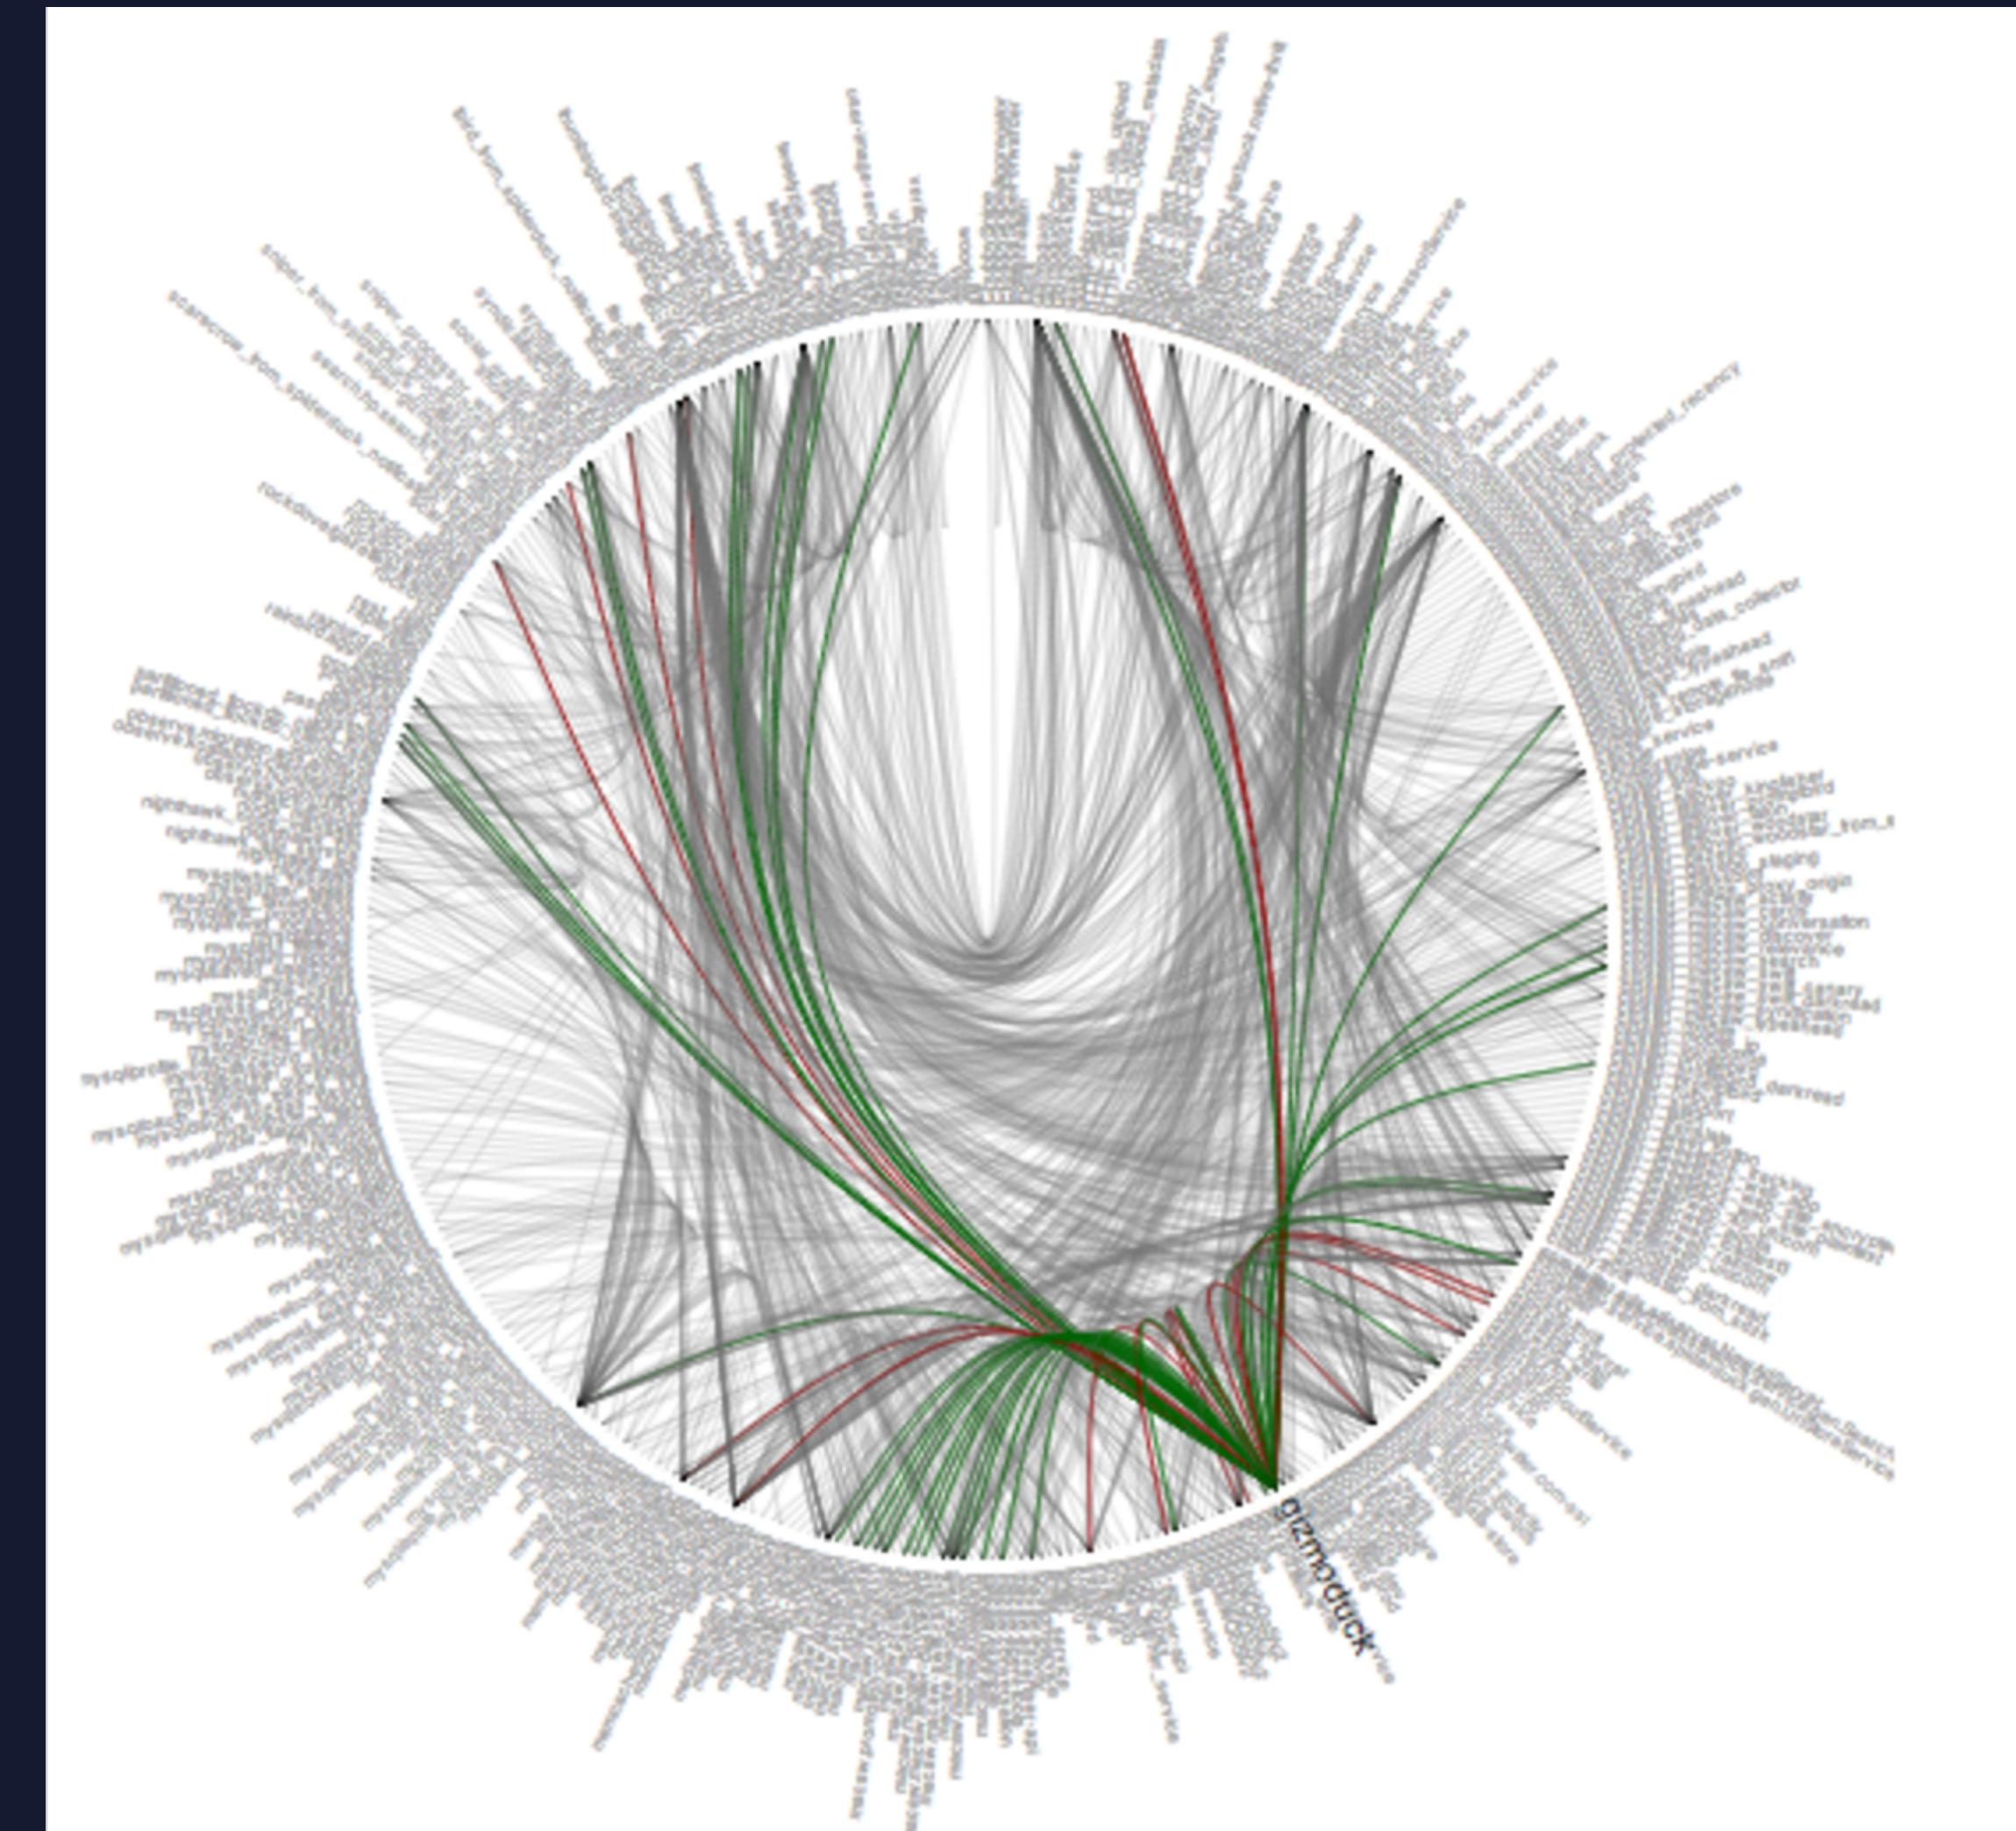
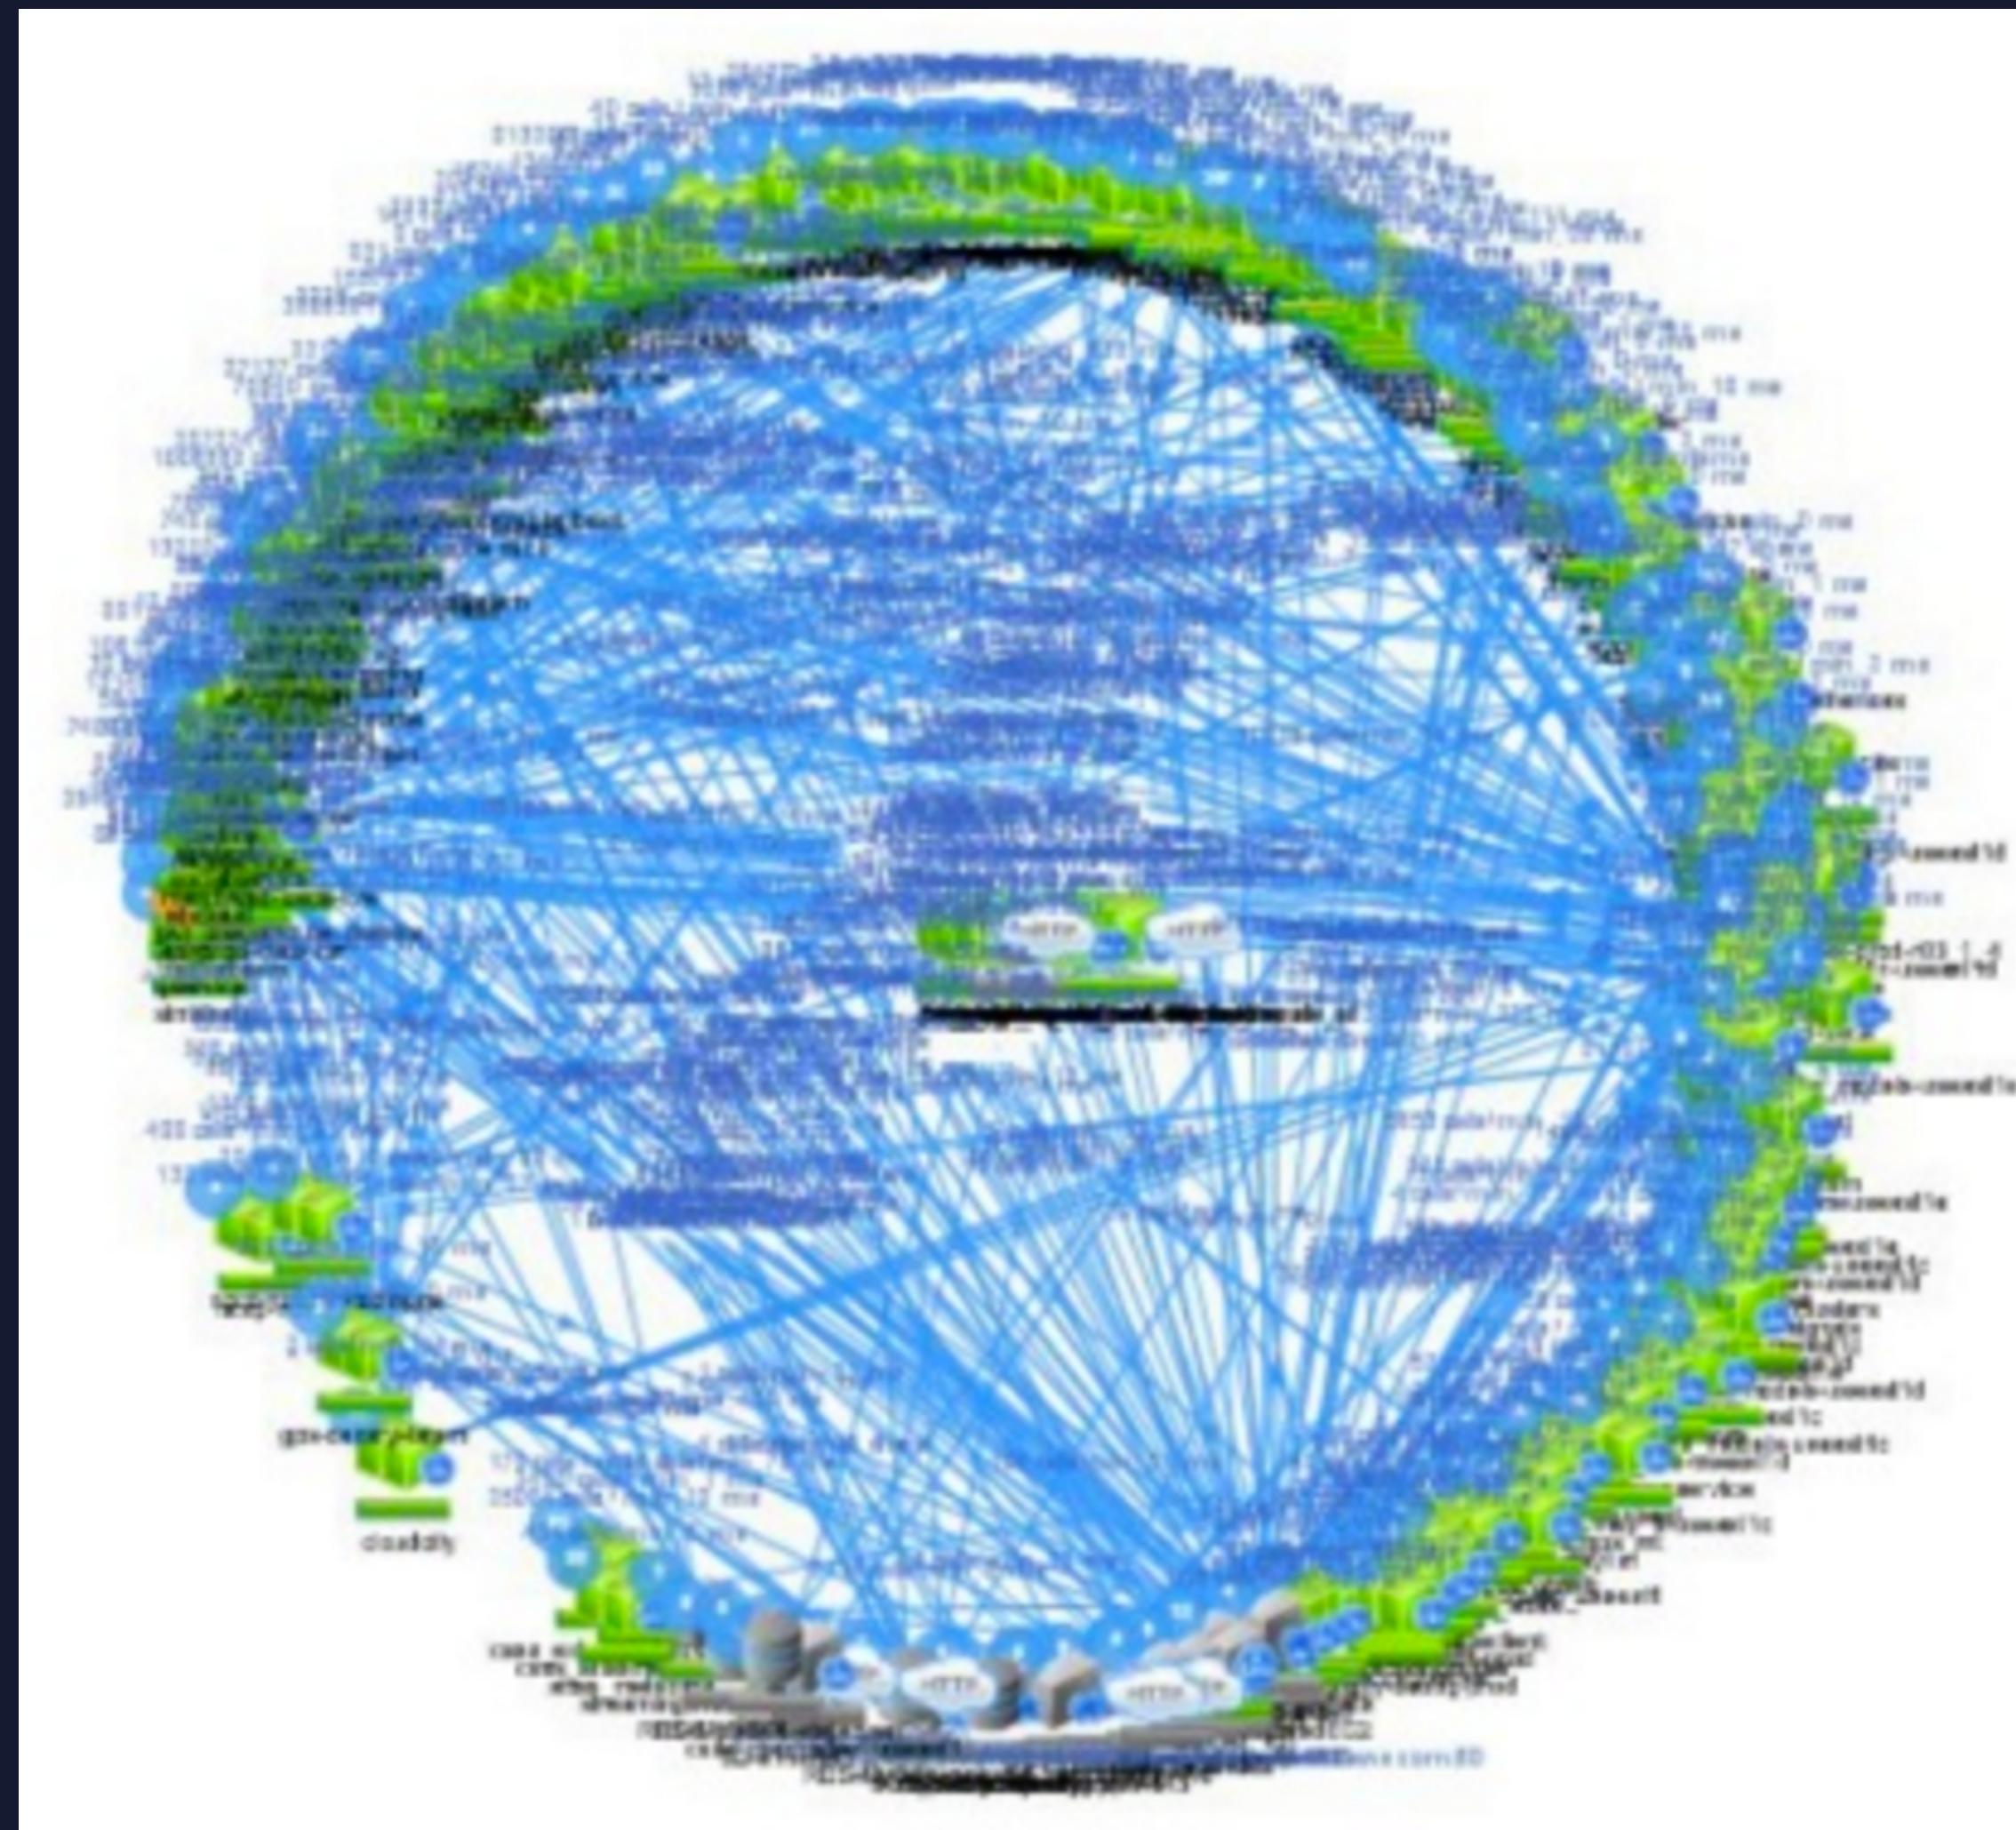
Trips



Architecture



Death Star Architectures



Can We Do Better
than Feral Concurrency Control?

Spanner: Google’s Globally-Distributed Database

James C. Corbett, Jeffrey Dean, Michael Epstein, Andrew Fikes, Christopher Frost, JJ Furman, Sanjay Ghemawat, Andrey Gubarev, Christopher Heiser, Peter Hochschild, Wilson Hsieh, Sebastian Kanthak, Eugene Kogan, Hongyi Li, Alexander Lloyd, Sergey Melnik, David Mwaura, David Nagle, Sean Quinlan, Rajesh Rao, Lindsay Rolig, Yasushi Saito, Michal Szymaniak, Christopher Taylor, Ruth Wang, Dale Woodford

Google, Inc.

Abstract

Spanner is Google’s scalable, multi-version, globally-distributed, and synchronously-replicated database. It is the first system to distribute data at global scale and support externally-consistent distributed transactions. This paper describes how Spanner is structured, its feature set, the rationale underlying various design decisions, and a novel time API that exposes clock uncertainty. This API and its implementation are critical to supporting external consistency and a variety of powerful features: non-blocking reads in the past, lock-free read-only transactions, and atomic schema changes, across all of Spanner.

1 Introduction

tency over higher availability, as long as they can survive 1 or 2 datacenter failures.

Spanner’s main focus is managing cross-datacenter replicated data, but we have also spent a great deal of time in designing and implementing important database features on top of our distributed-systems infrastructure. Even though many projects happily use Bigtable [9], we have also consistently received complaints from users that Bigtable can be difficult to use for some kinds of applications: those that have complex, evolving schemas, or those that want strong consistency in the presence of wide-area replication. (Similar claims have been made by other authors [37].) Many applications at Google have chosen to use Megastore [5] because of its semi-relational data model and support for synchronous repli-

2012

Spanner: Google’s Globally-Distributed Database

James C. Corbett, Jeffrey Dean, Michael Epstein, Andrew Fikes, Christopher Frost, JJ Furman, Sanjay Ghemawat, Andrey Gubarev, Christopher Heiser, Peter Hochschild, Wilson Hsieh,

“Spanner is Google’s scalable, multi-version, globally distributed, and synchronously-replicated database”

and its implementation are critical to supporting external consistency and a variety of powerful features: non-blocking reads in the past, lock-free read-only transactions, and atomic schema changes, across all of Spanner.

1 Introduction

that Bigtable can be difficult to use for some kinds of applications: those that have complex, evolving schemas, or those that want strong consistency in the presence of wide-area replication. (Similar claims have been made by other authors [37].) Many applications at Google have chosen to use Megastore [5] because of its semi-relational data model and support for synchronous repli-

Challenges to Adopting Stronger Consistency at Scale

Phillipe Ajoux*, Nathan Bronson*, Sanjeev Kumar*, Wyatt Lloyd†*, Kaushik Veeraraghavan*

*Facebook, †University of Southern California

Abstract

There have been many recent advances in distributed systems that provide stronger semantics for geo-replicated data stores like those underlying Facebook. These research systems provide a range of consistency models and transactional abilities while demonstrating good performance and scalability on experimental workloads. At Facebook we are excited by these lines of research, but fundamental and operational challenges currently make it infeasible to incorporate these advances into deployed systems. This paper describes some of these challenges with the hope that future advances will address them.

1 Introduction

Facebook is a social network that connects 1.35 billion people [19]. Facebook’s social graph reflects its users, their relationships, the content they create, and the actions they take. The social graph is large and constantly growing and changing. Data from this graph is stored in several geo-replicated data stores and tightly inte-

such as Search and News Feed. No single data placement strategy can efficiently serve all workloads in a heavily sharded system, so many of these services choose a specialized sharding function and maintain their own data store, caches, and indexes.

The biggest barrier to providing stronger consistency guarantees in an environment like Facebook’s is that the consistency mechanism must **integrate consistency across many stateful services**. All of the scaling mechanisms described above lead to extra copies of data. Even if each of our caches and independent services were linearizable [24], inconsistencies would still be present in the aggregated result. Inter-service tracking is complicated by services that store data derived from a query to a lower layer, making object-level tracking insufficient. While these challenges are acute for Facebook at its present scale, we first encountered them when Facebook was much smaller. Solutions to these problems will benefit the growing class of applications whose implementation relies on sharding and separation into stateful services.

Another fundamental challenge is that a general consistency mechanism at scale must **tolerate high query**

Challenges to Adopting Stronger Consistency at Scale

Phillipe Ajoux*, Nathan Bronson*, Sanjeev Kumar*, Wyatt Lloyd†*, Kaushik Veeraraghavan*

*Facebook, †University of Southern California

“The biggest barrier...is that consistency mechanisms must integrate across many stateful services”

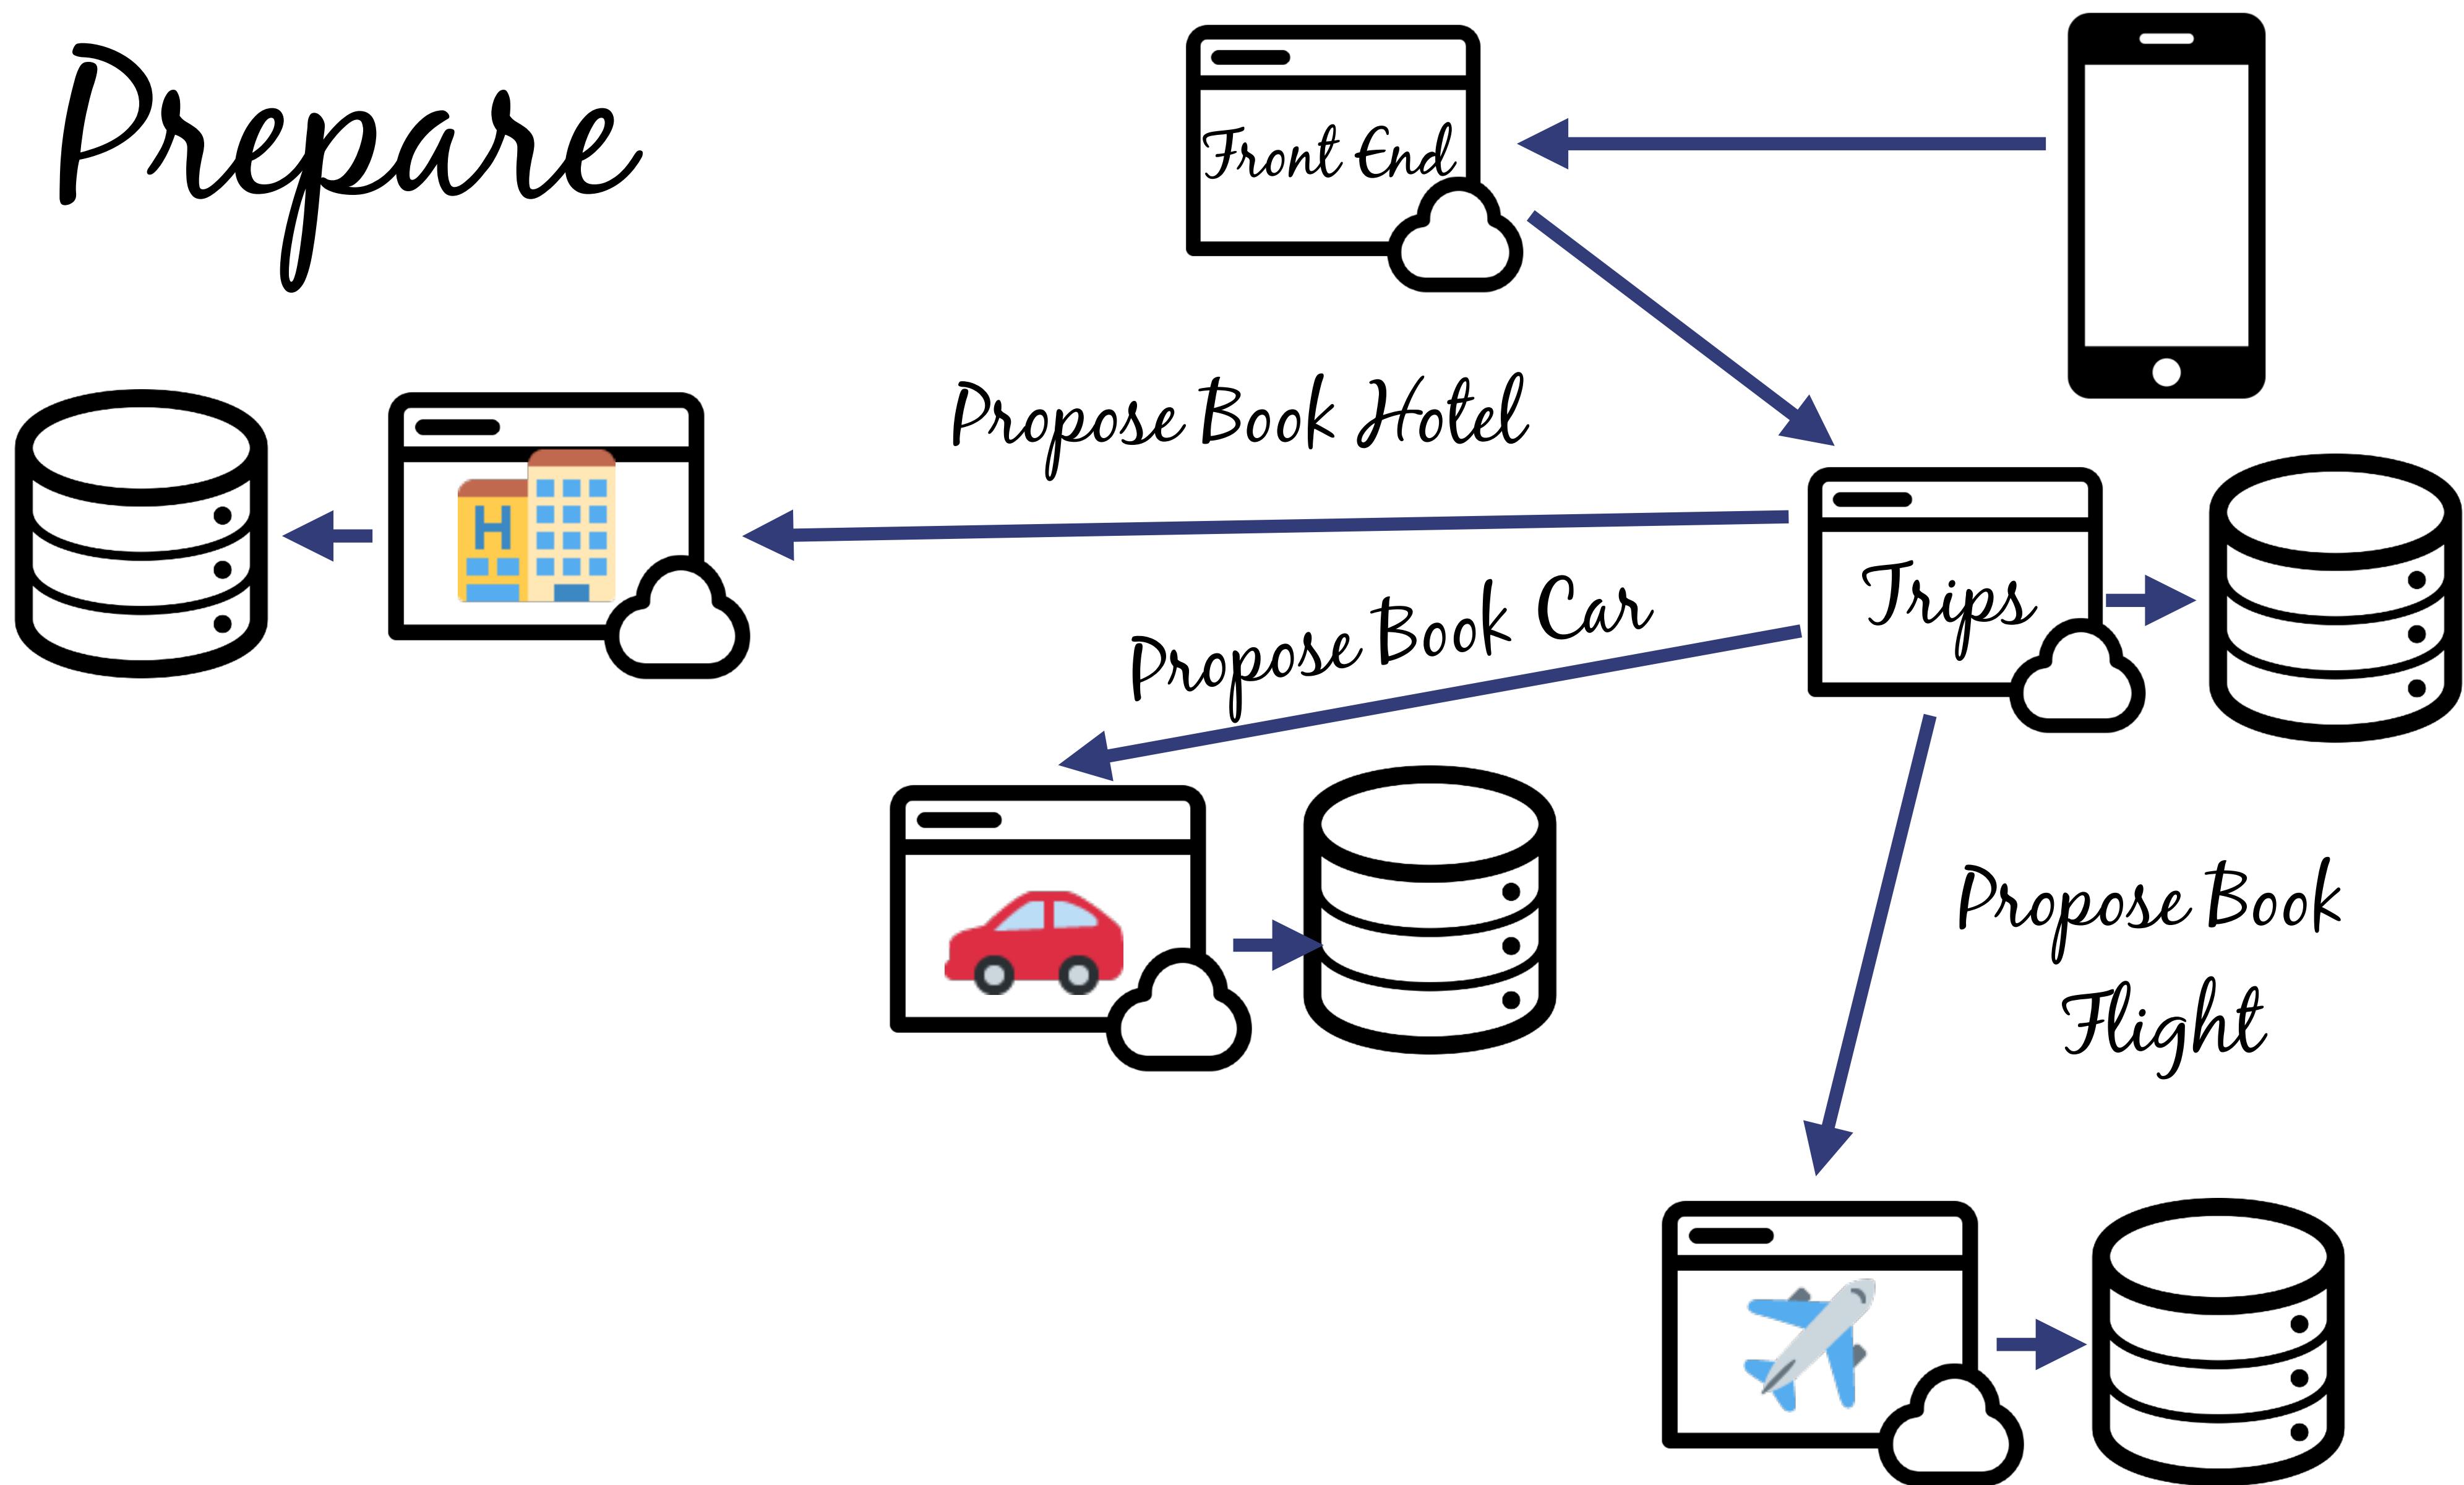
people [19]. Facebook’s social graph reflects its users, their relationships, the content they create, and the actions they take. The social graph is large and constantly growing and changing. Data from this graph is stored in several geo-replicated data stores and tightly inte-

grated with other data stores to support the needs of the system. This is a challenge for consistency mechanisms, as they must ensure that data is consistent across all these stores.

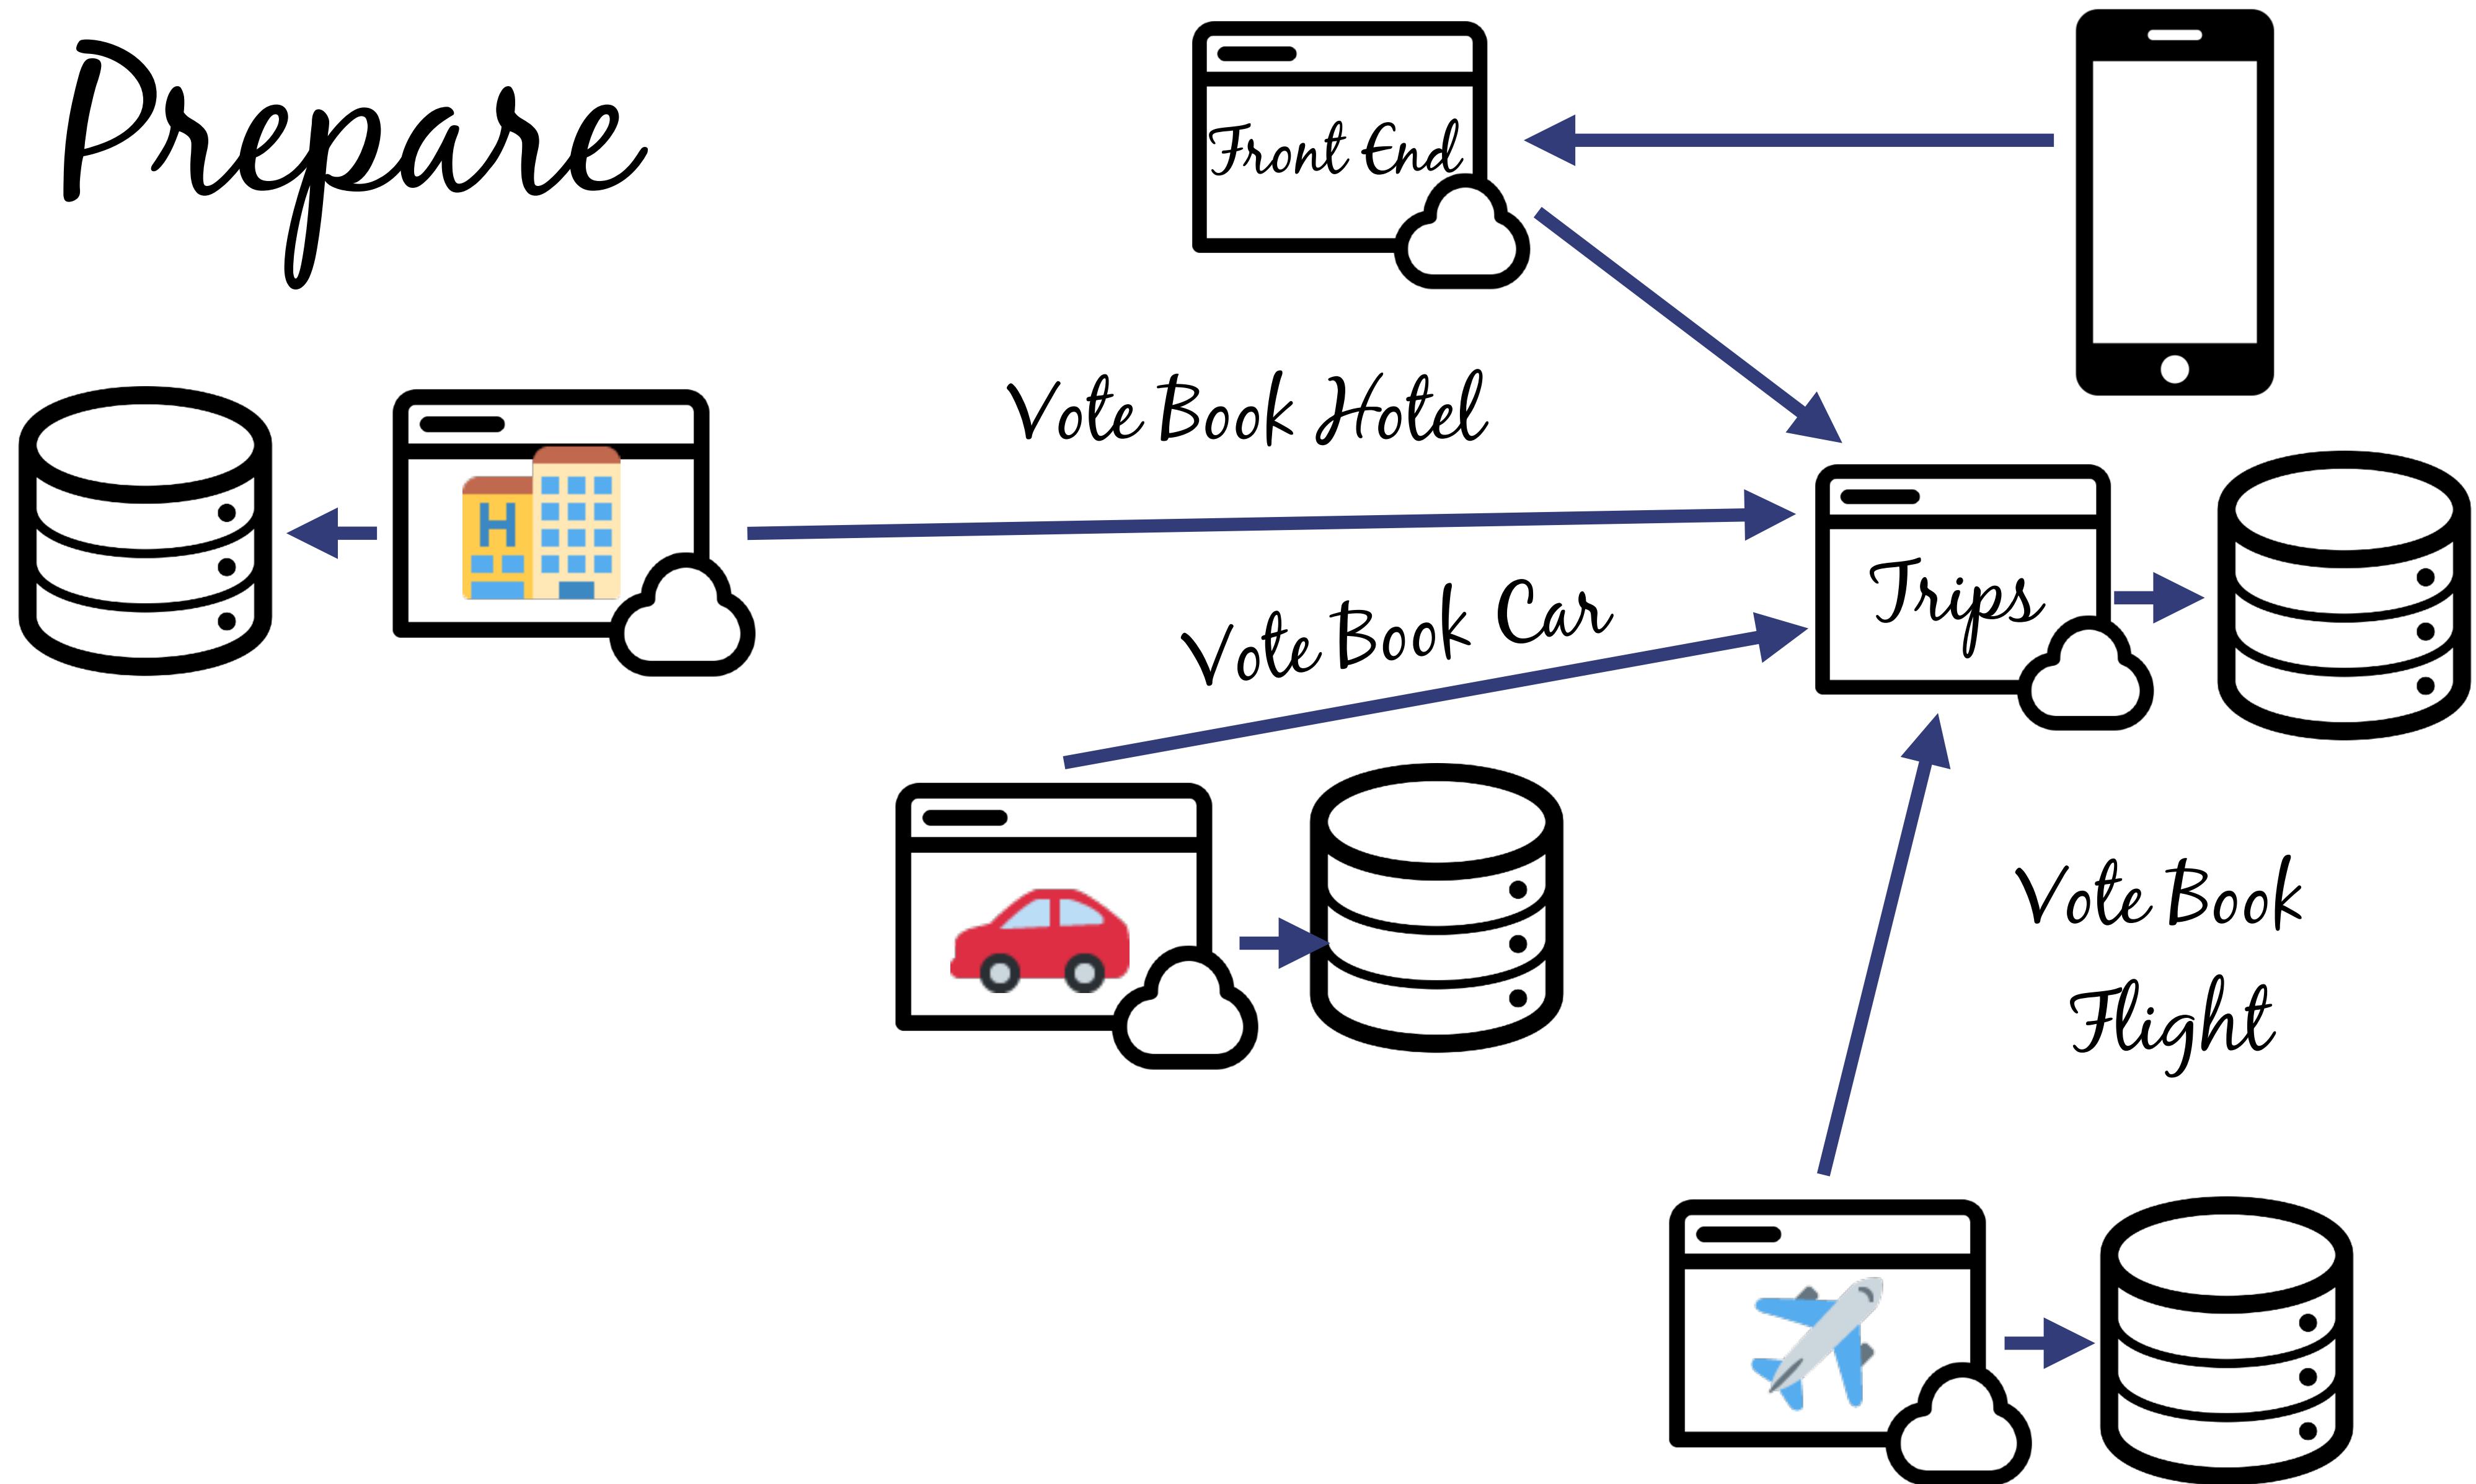
Another fundamental challenge is that a general consistency mechanism at scale must tolerate high query

Two Phase Commit

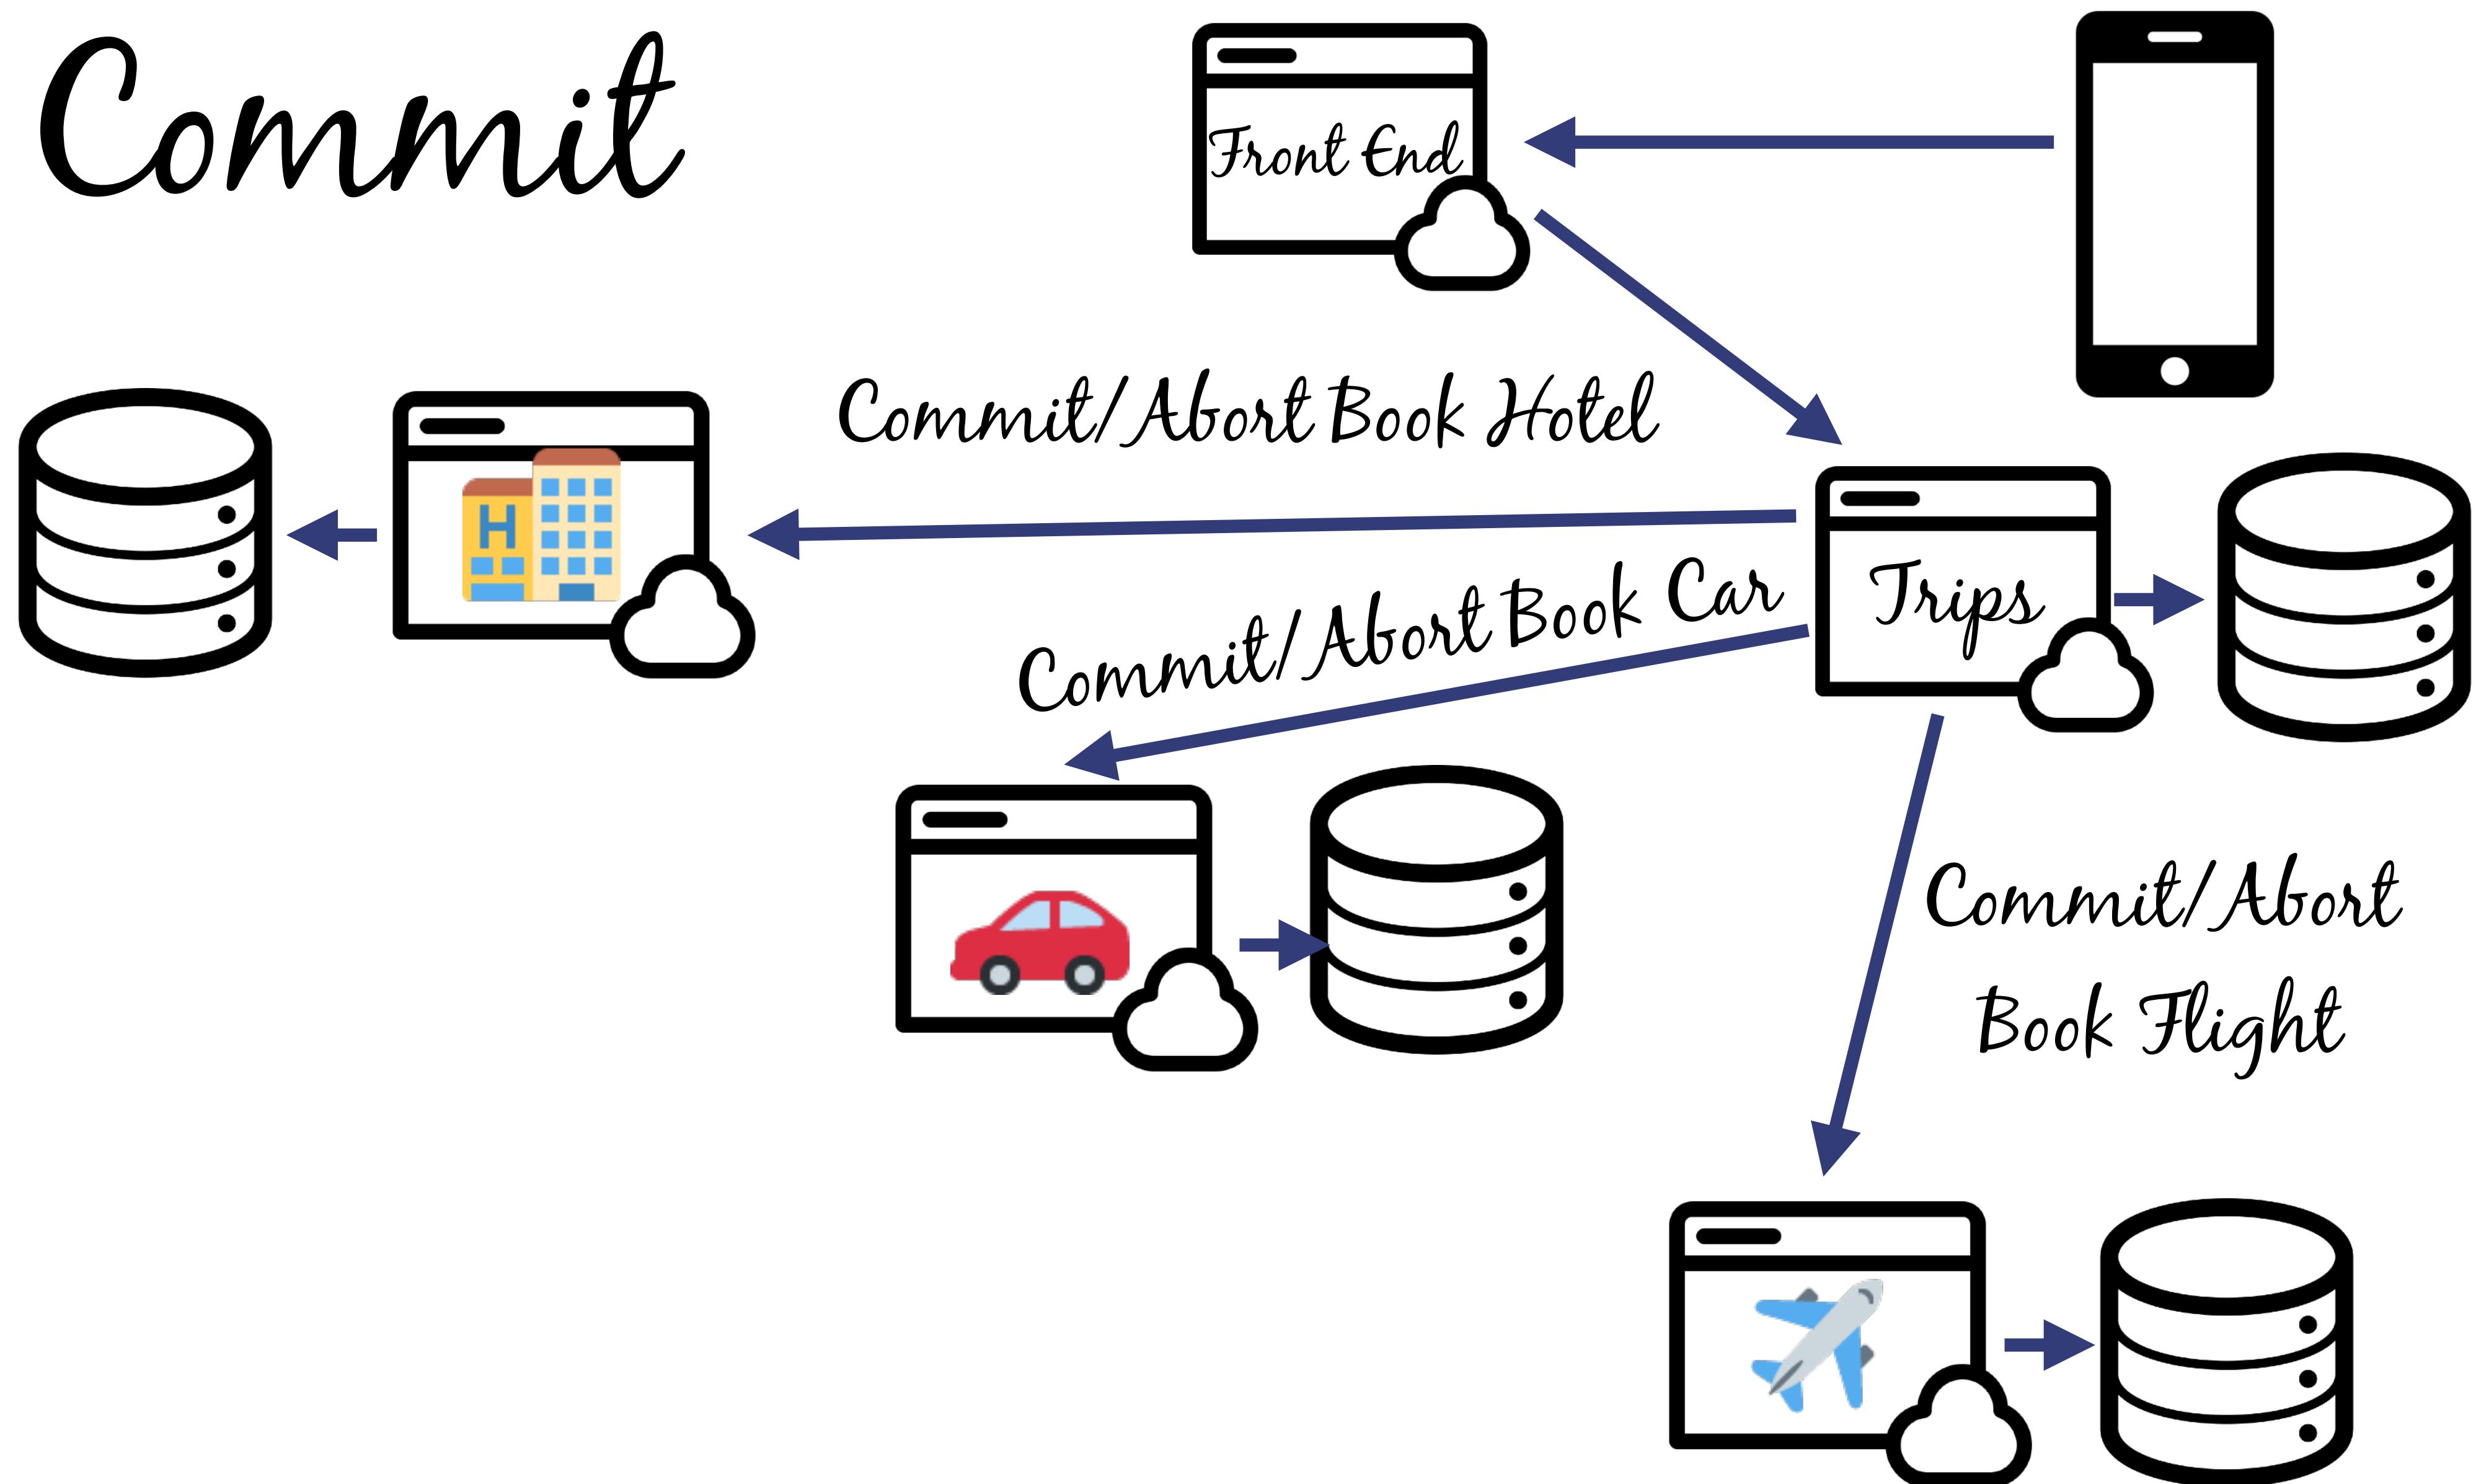
2PC: Prepare



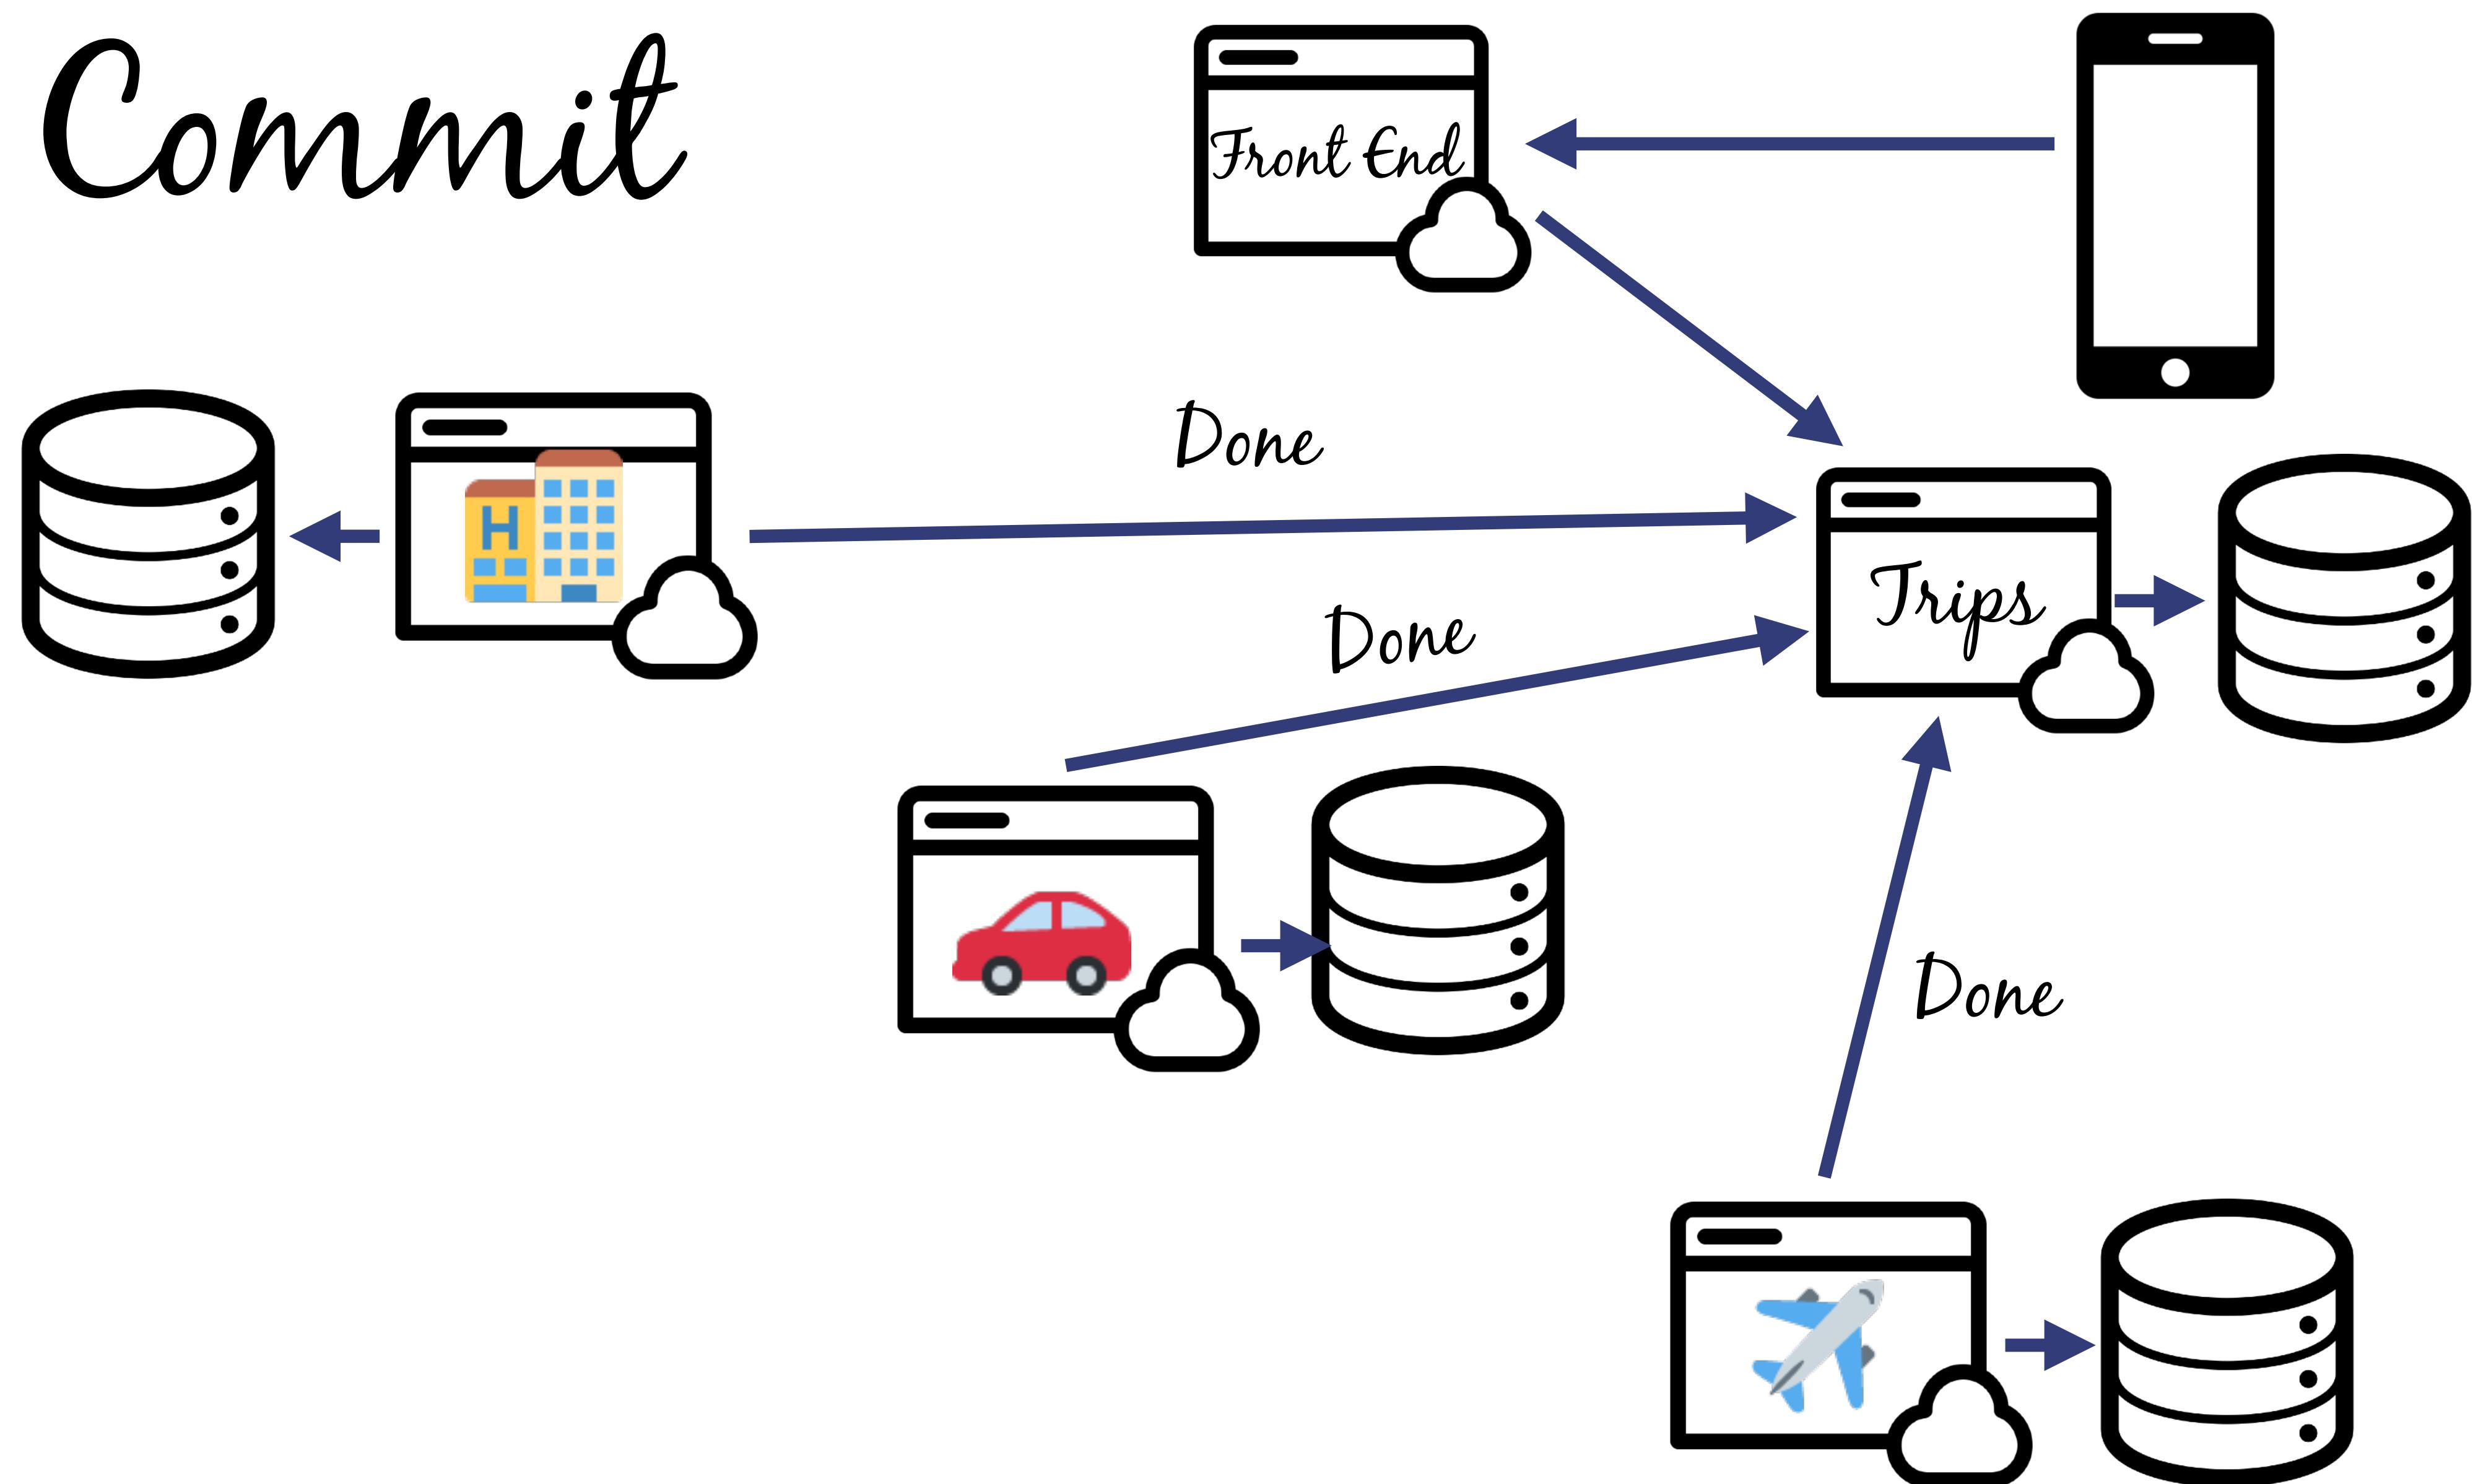
2PC: Prepare



2PC: Commit



2PC: Commit



2PC:

Doesn't Scale

- $O(N^2)$ Messages in the worst case
- Coordinator is a Single Point of Failure
- Reduced Throughput

Distributed Sagas

A Protocol for Coordinating Microservices

SAGAS

*Hector Garcia-Molina
Kenneth Salem*

Department of Computer Science
Princeton University
Princeton, N J 08544

Abstract

Long lived transactions (LLTs) hold on to database resources for relatively long periods of time, significantly delaying the termination of shorter and more common transactions. To alleviate these problems we propose the notion of a saga. A LLT is a saga if it can be written as a sequence of transactions that can be interleaved with other transactions. The database management system guarantees that either all the transactions in a saga are successfully completed or compensating transactions are run to amend a partial execution. Both the concept of saga and its implementation are relatively simple, but they have the potential to improve performance significantly. We analyze the various implementation issues related to sagas, including how they can be run on an existing system that does not directly support them. We also discuss techniques for dealing with LLTs that are short

the majority of other transactions either because it accesses many database objects, it has lengthy computations, it pauses for inputs from the users, or a combination of these factors. Examples of LLTs are transactions to produce monthly account statements at a bank, transactions to process claims at an insurance company, and transactions to collect statistics over an entire database [Gray81a].

In most cases, LLTs present serious performance problems. Since they are transactions, the system must execute them as atomic actions, thus preserving the consistency of the database [Date81a, Ullm82a]. To make a transaction atomic, the system usually locks the objects accessed by the transaction until it commits, and this typically occurs at the end of the transaction. As a consequence, other transactions wishing to access the LLT's objects suffer a long delay. If LLTs are long because they

1987

SAGAS

*Hector Garcia-Molina
Kenneth Salem*

Department of Computer Science

“Sagas are Long Lived Transactions [in a single relational Database]”

have the potential to improve performance significantly. We analyze the various implementation issues related to sagas, including how they can be run on an existing system that does not directly support them. We also discuss techniques for detecting LLT bugs that can lead to

saction atomic, the system usually locks the objects accessed by the transaction until it commits, and this typically occurs at the end of the transaction. As a consequence, other transactions wishing to access the LLT’s objects suffer a long delay. If LLT are long-running, then

1987

SAGAS

*Hector Garcia-Molina
Kenneth Salem*

“A Saga is a Long Lived Transaction that can be written as a sequence of transactions that can be interleaved.

All transactions in the sequence complete successfully or compensating transactions are ran to amend a partial execution.”

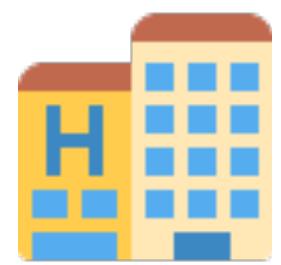
tation issues related to sagas, including how they can be run on an existing system that does not directly support them. We also discuss techniques for detecting an LLT that has not yet

commits, and this typically occurs at the end of the transaction. As a consequence, other transactions wishing to access the LLT’s objects suffer a long delay. If LLTs are long-running, then

Distributed Sagas

A Protocol for Coordinating Microservices

A Distributed Saga is a Collection of Requests



Book Hotel



Book Car



Book Flight



Charge Money

A Distributed Saga is a Collection of Requests



Book Hotel



Book Car

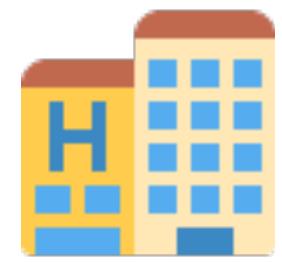


Book Flight



Charge Money

and Compensating Requests



Cancel Hotel



Cancel Car

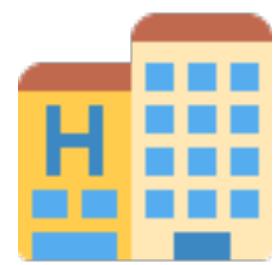


Cancel Flight



Refund Money

A Distributed Saga is a Collection of Requests



Book Hotel



Book Car



Book Flight



Charge Money

and Compensating Requests



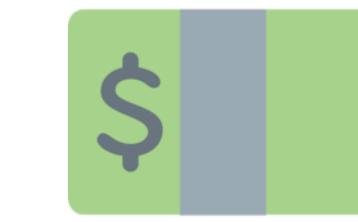
Cancel Hotel



Cancel Car



Cancel Flight

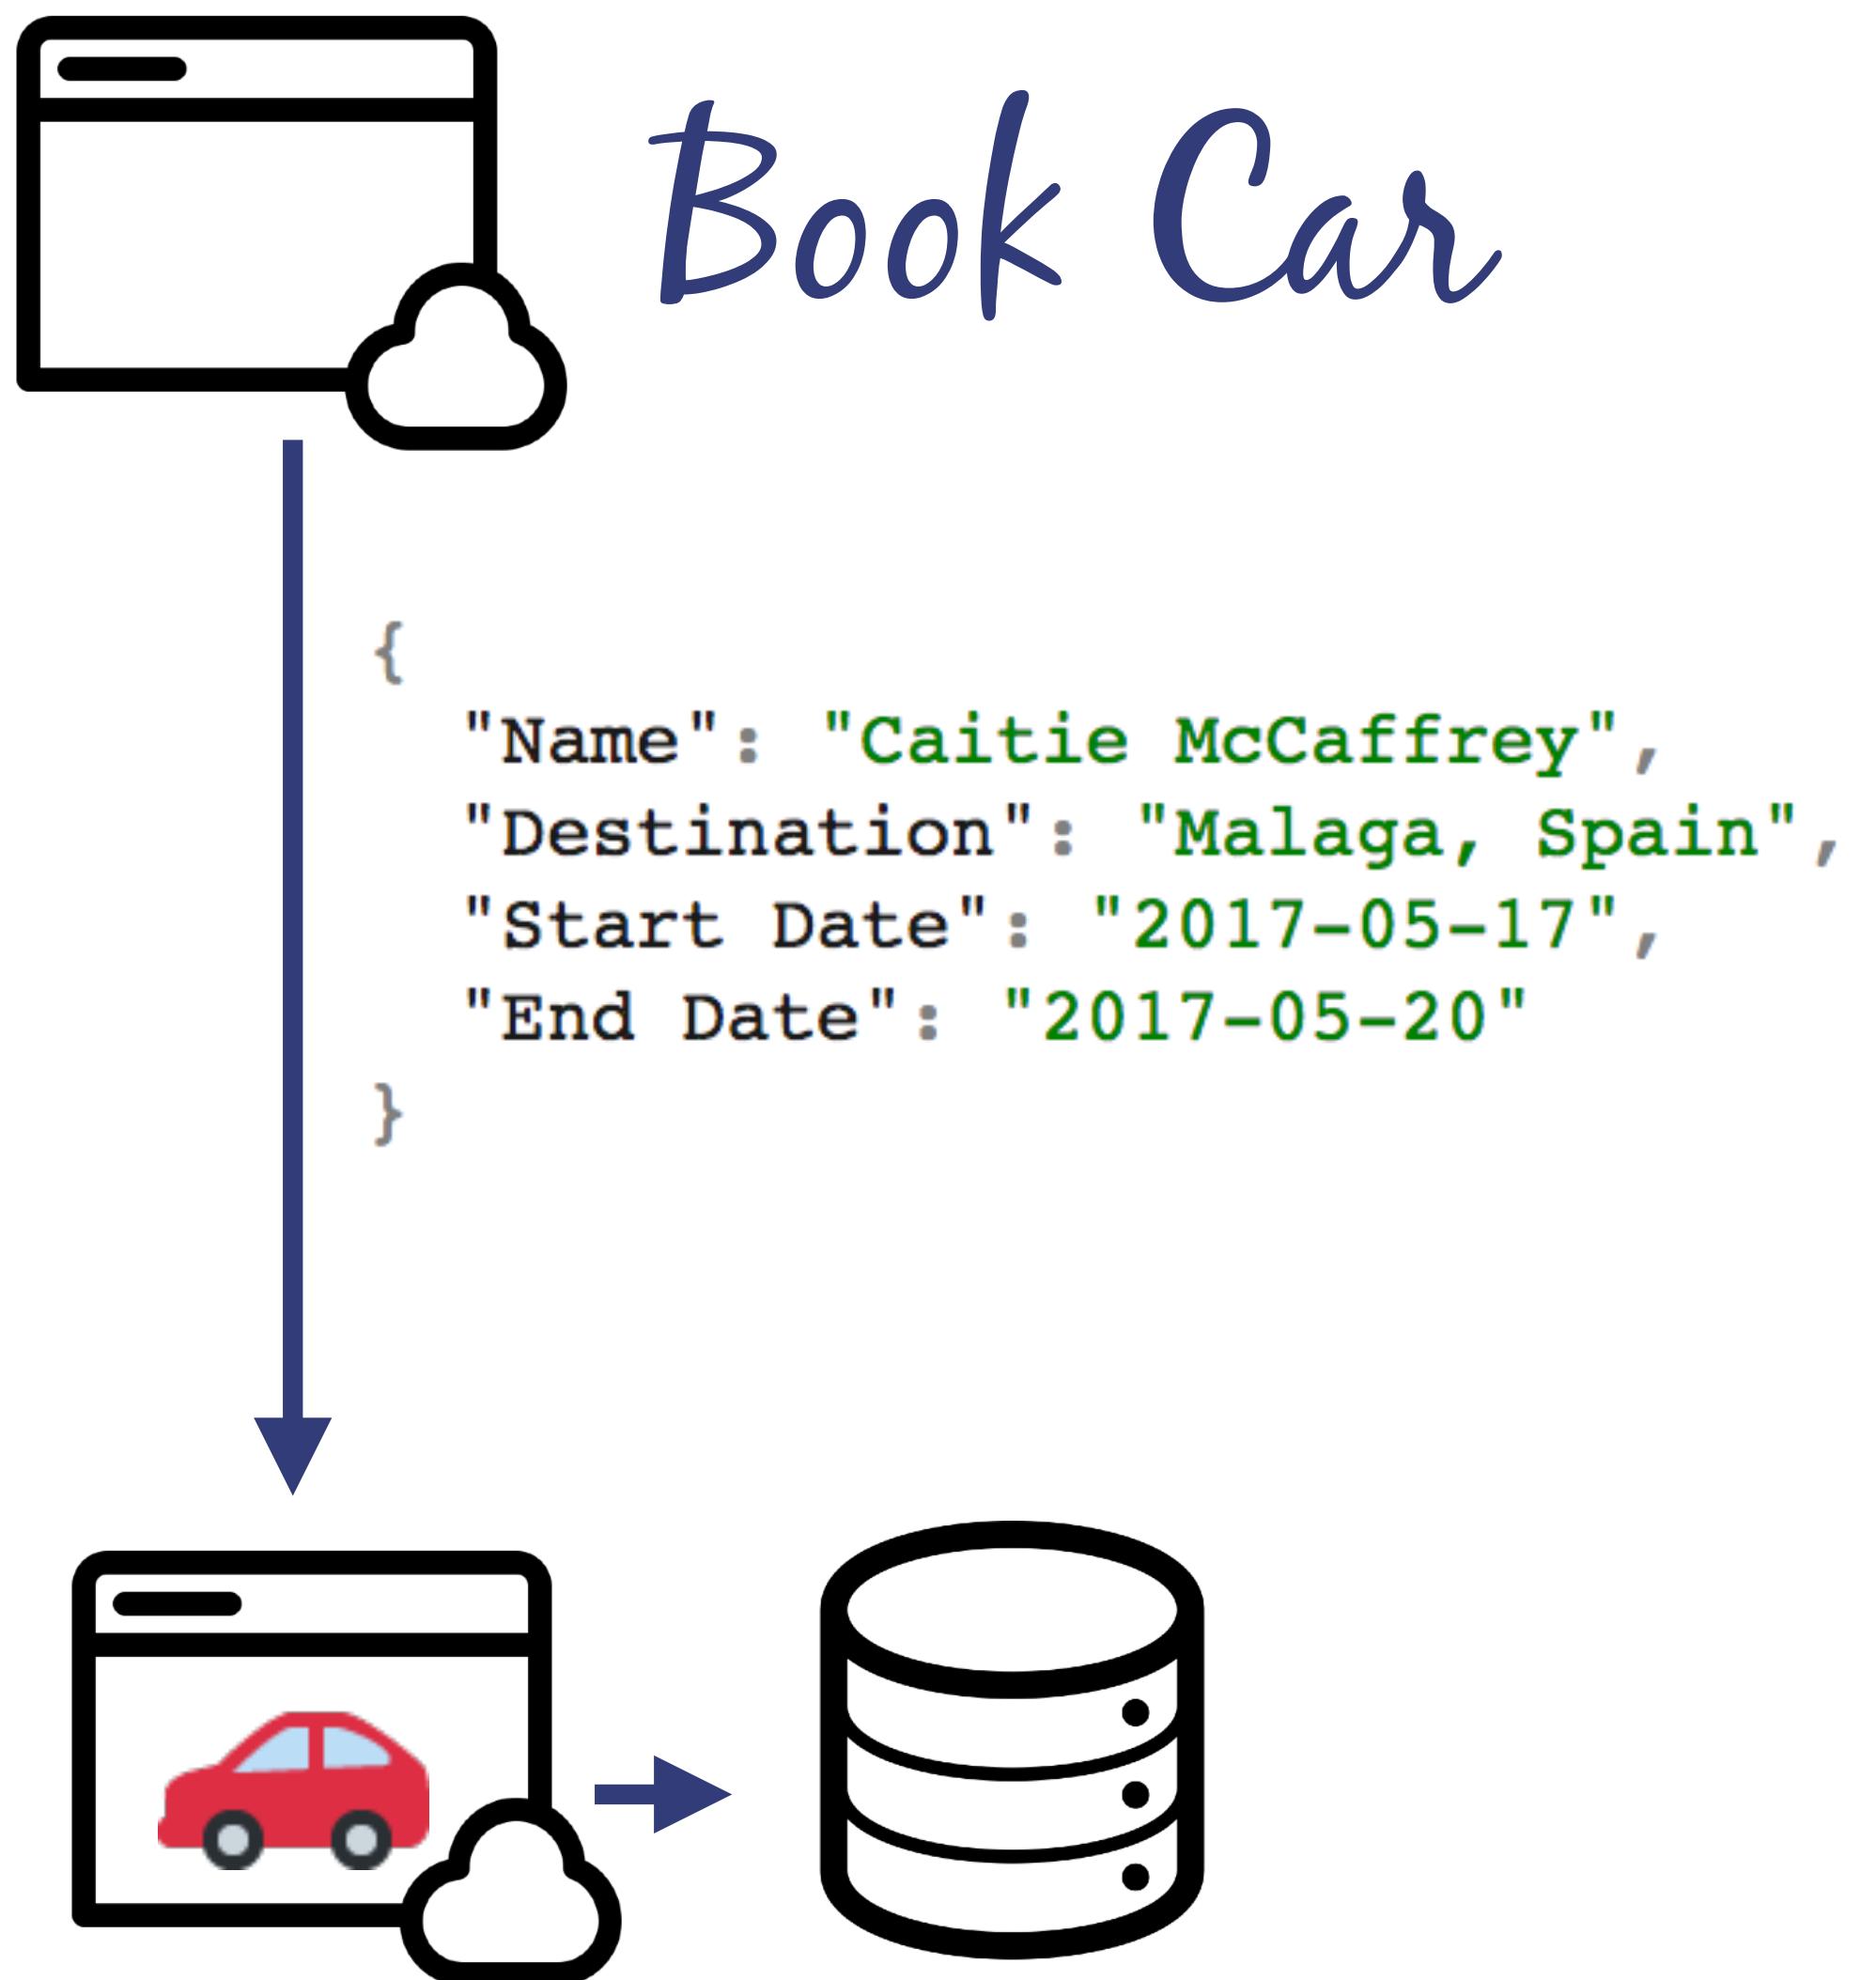


Refund Money

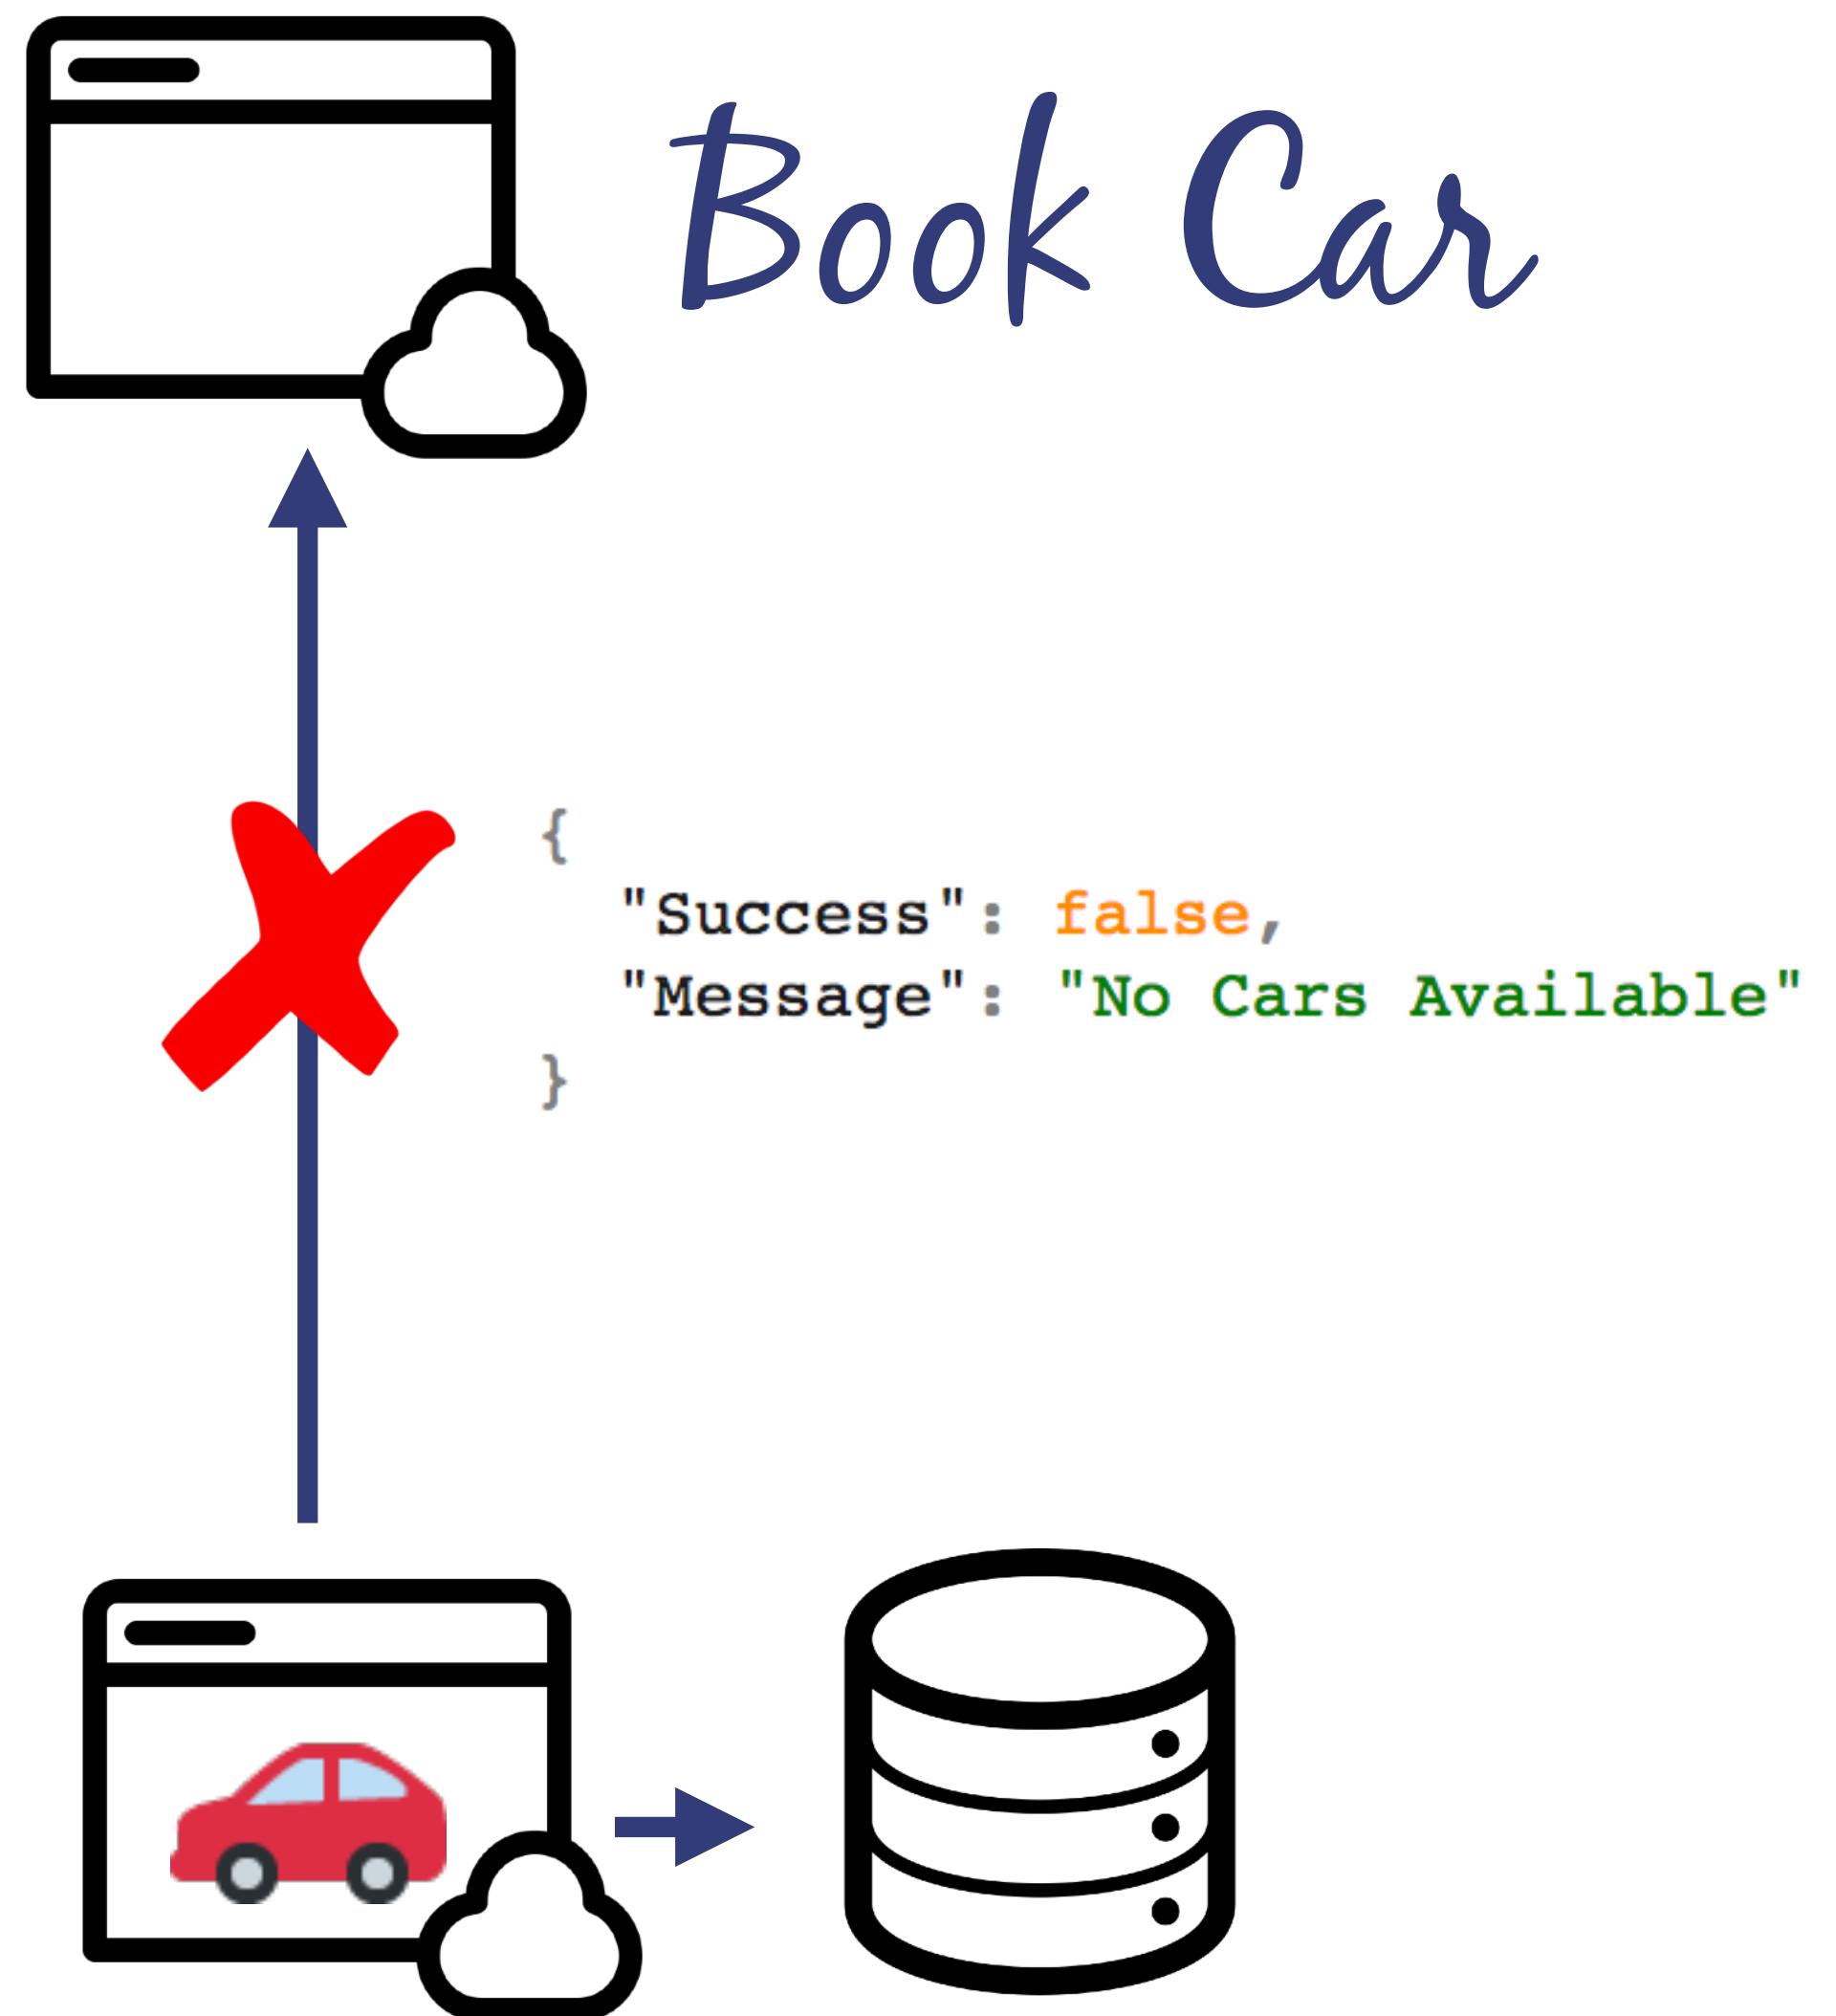
that represent a single business level action

Distributed Saga *Requests*

Requests Can Abort



Requests Can Abort



Requests

Must Be Idempotent



```
{  
  "Name": "Caitie McCaffrey",  
  "Destination": "Malaga, Spain",  
  "Start Date": "2017-05-17",  
  "End Date": "2017-05-20"  
}
```

Requests

Must Be Idempotent



```
{  
  "Name": "Caitie McCaffrey",  
  "Destination": "Malaga, Spain",  
  "Start Date": "2017-05-17",  
  "End Date": "2017-05-20"  
}
```

Requests Must Be Idempotent



```
{  
  "Success": "true",  
  "Confirmation Number": "ABC456"  
}
```

Requests

Must Be Idempotent

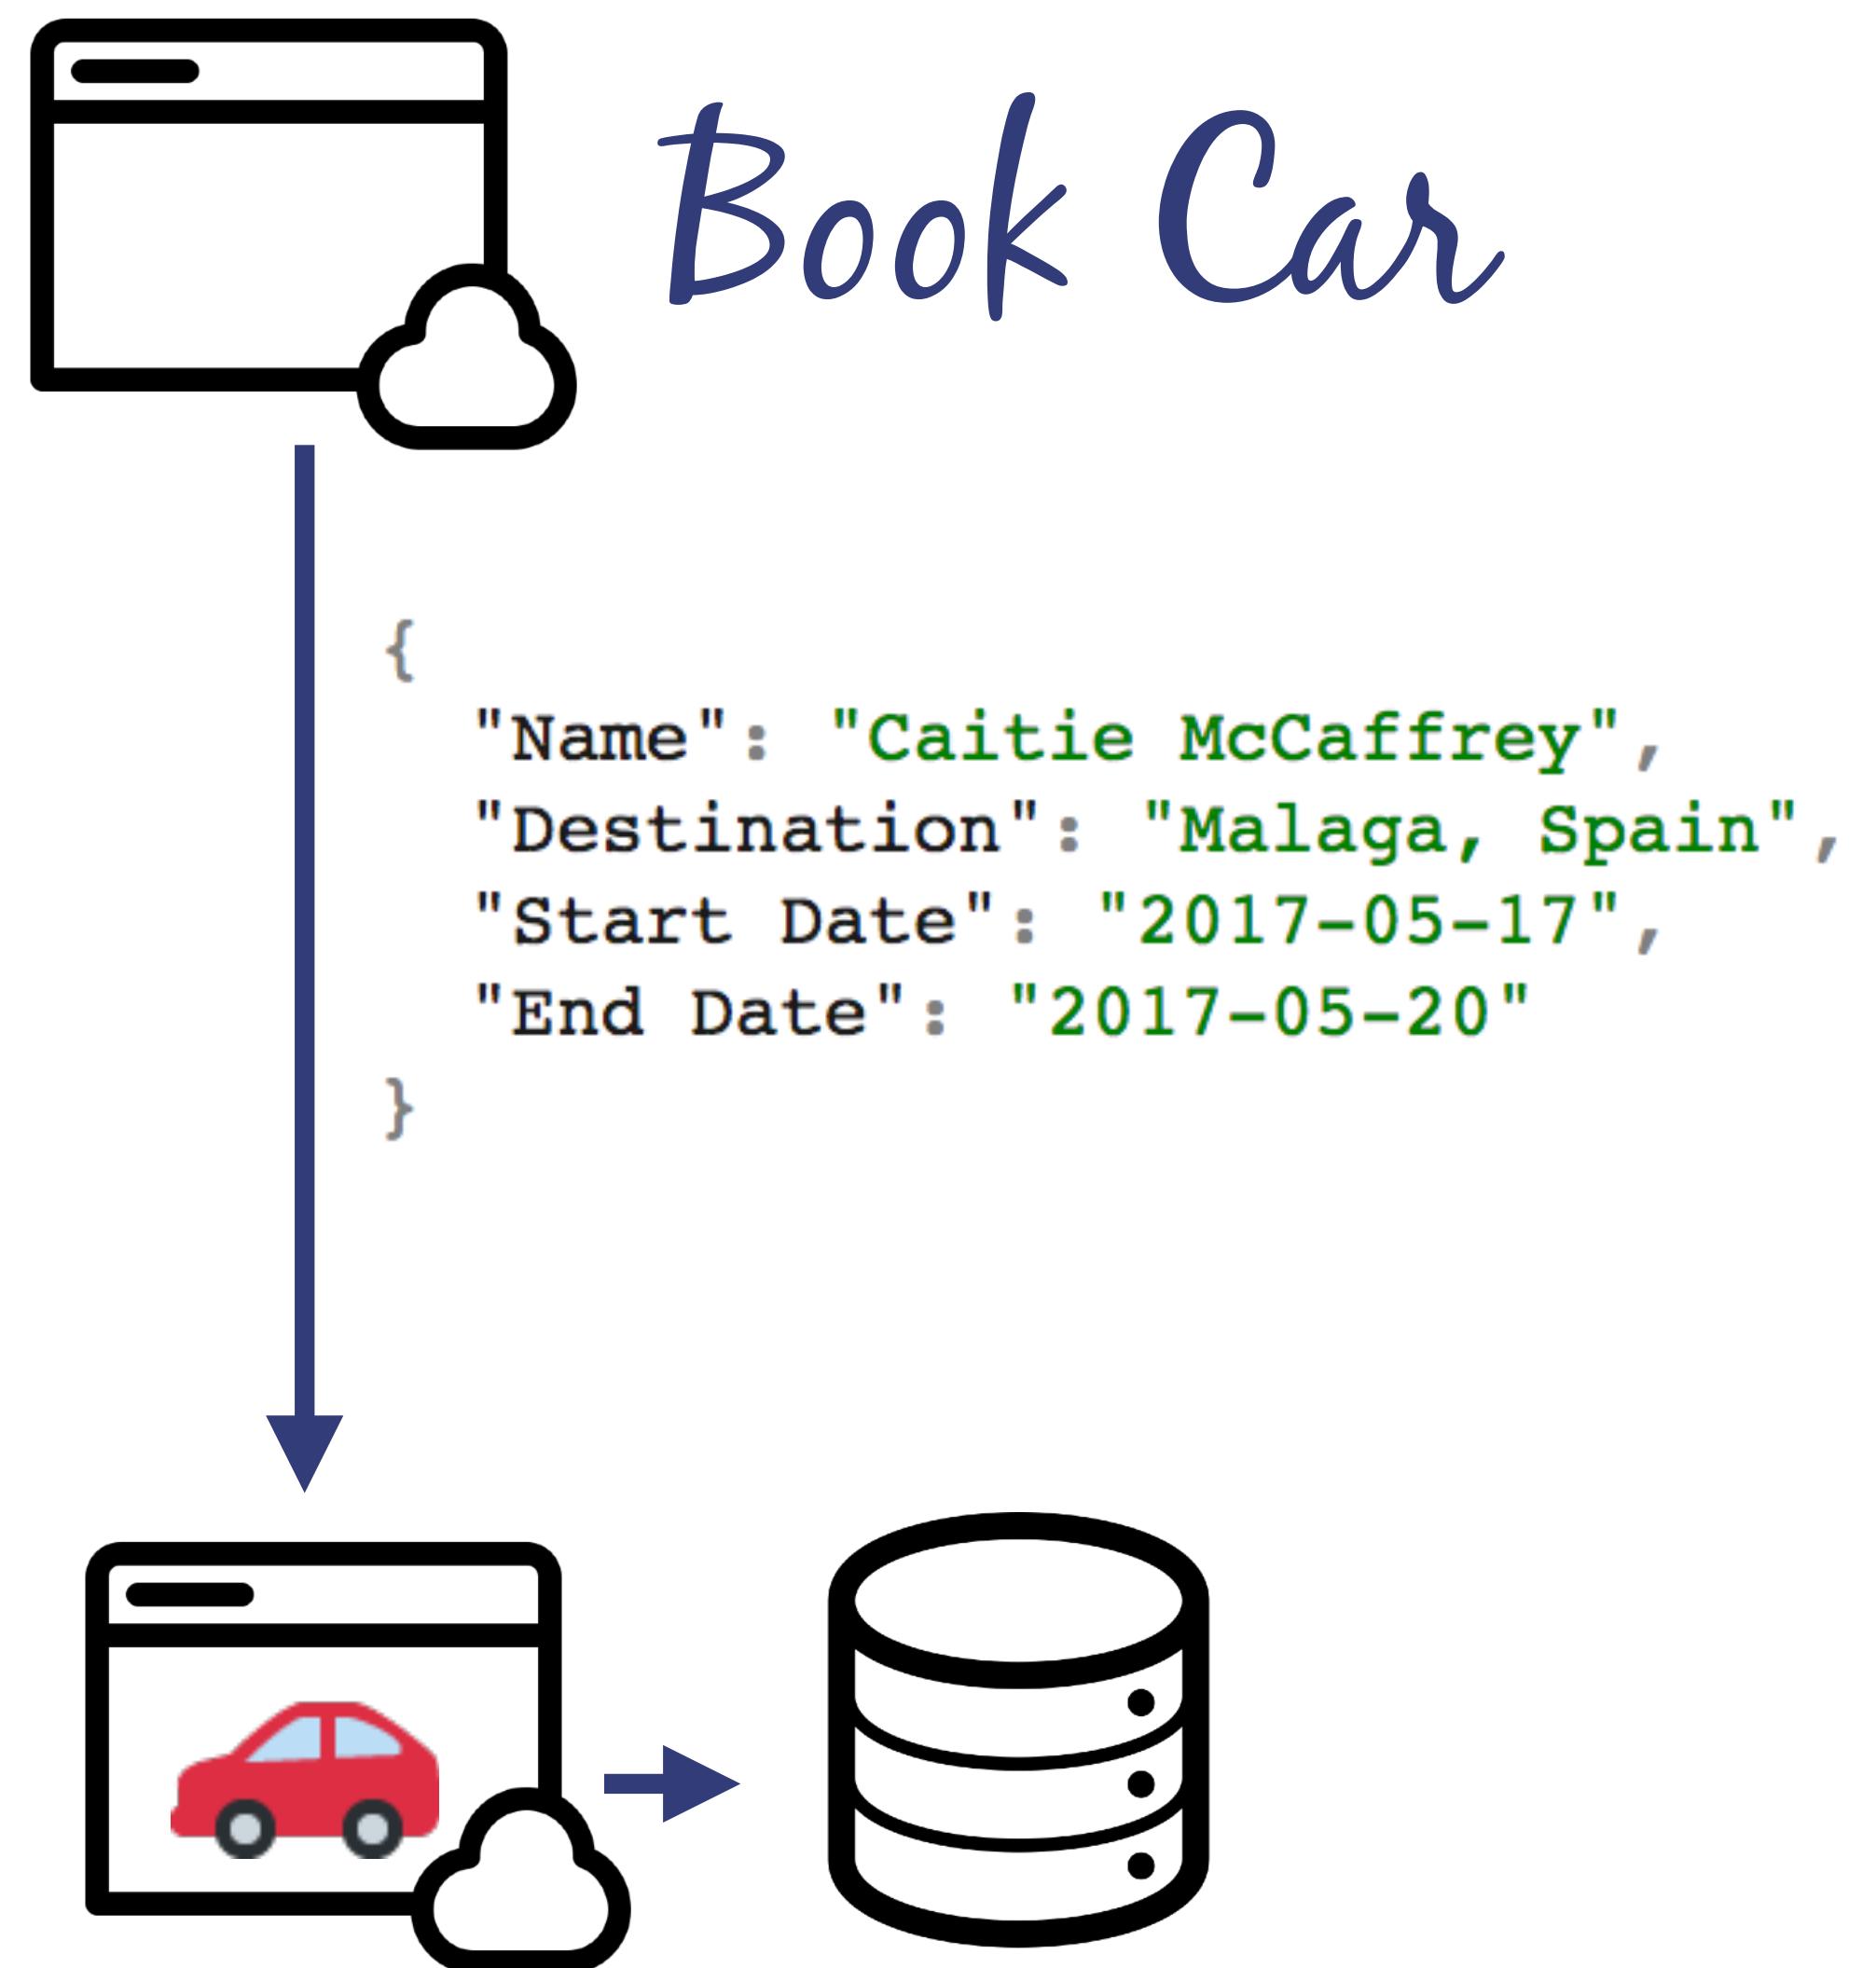
```
{  
  "Name": "Caitie McCaffrey",  
  "Destination": "Malaga, Spain",  
  "Start Date": "2017-05-17",  
  "End Date": "2017-05-20"  
}
```



```
{  
  "Success": "true",  
  "Confirmation Number": "ABC456"  
}
```

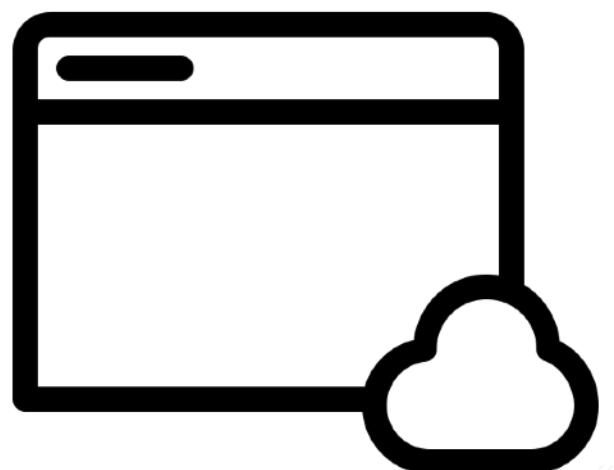
Requests

Must Be Idempotent



Requests

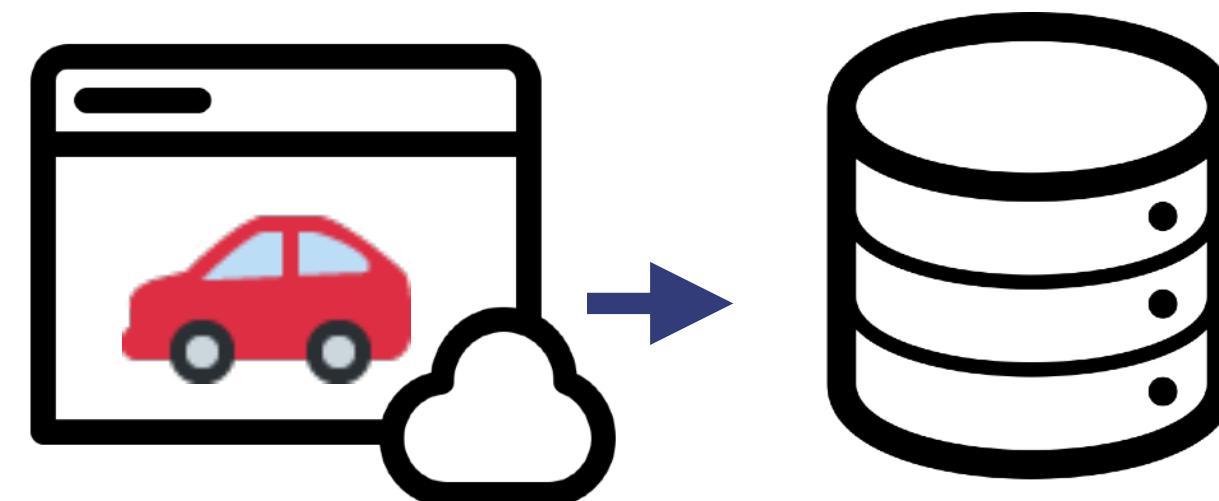
Must Be Idempotent



Book Car

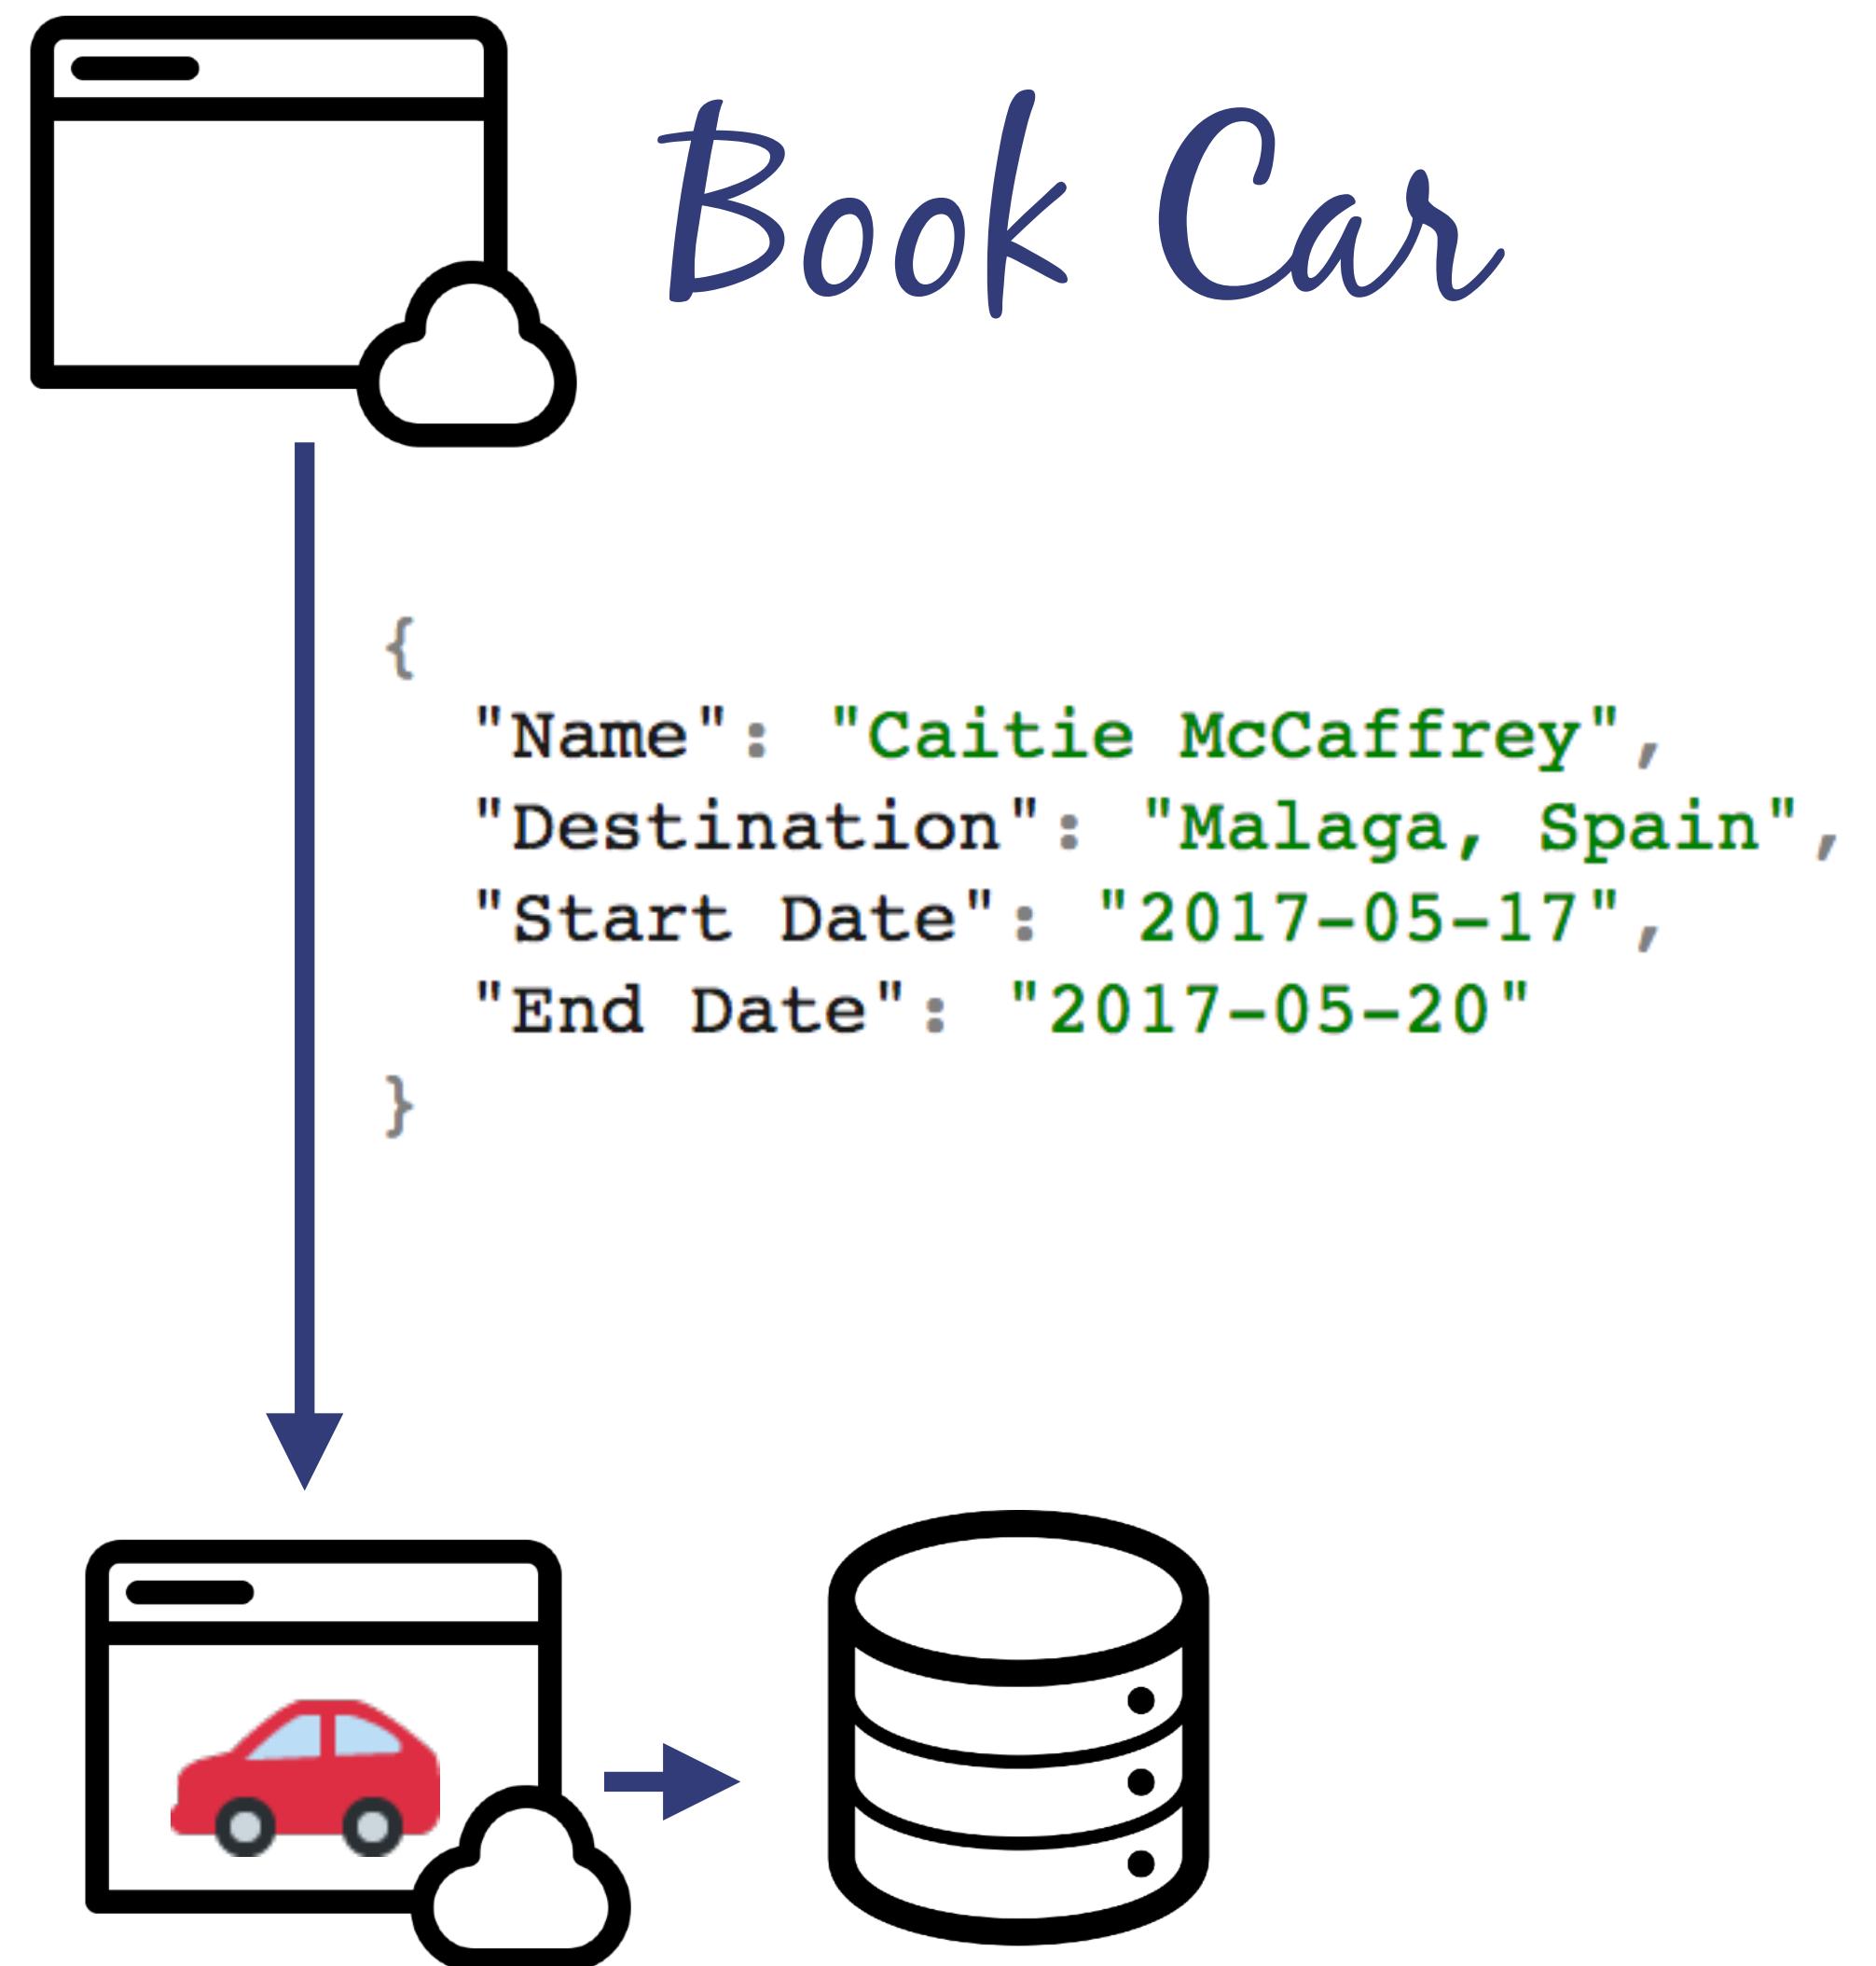


Timeout



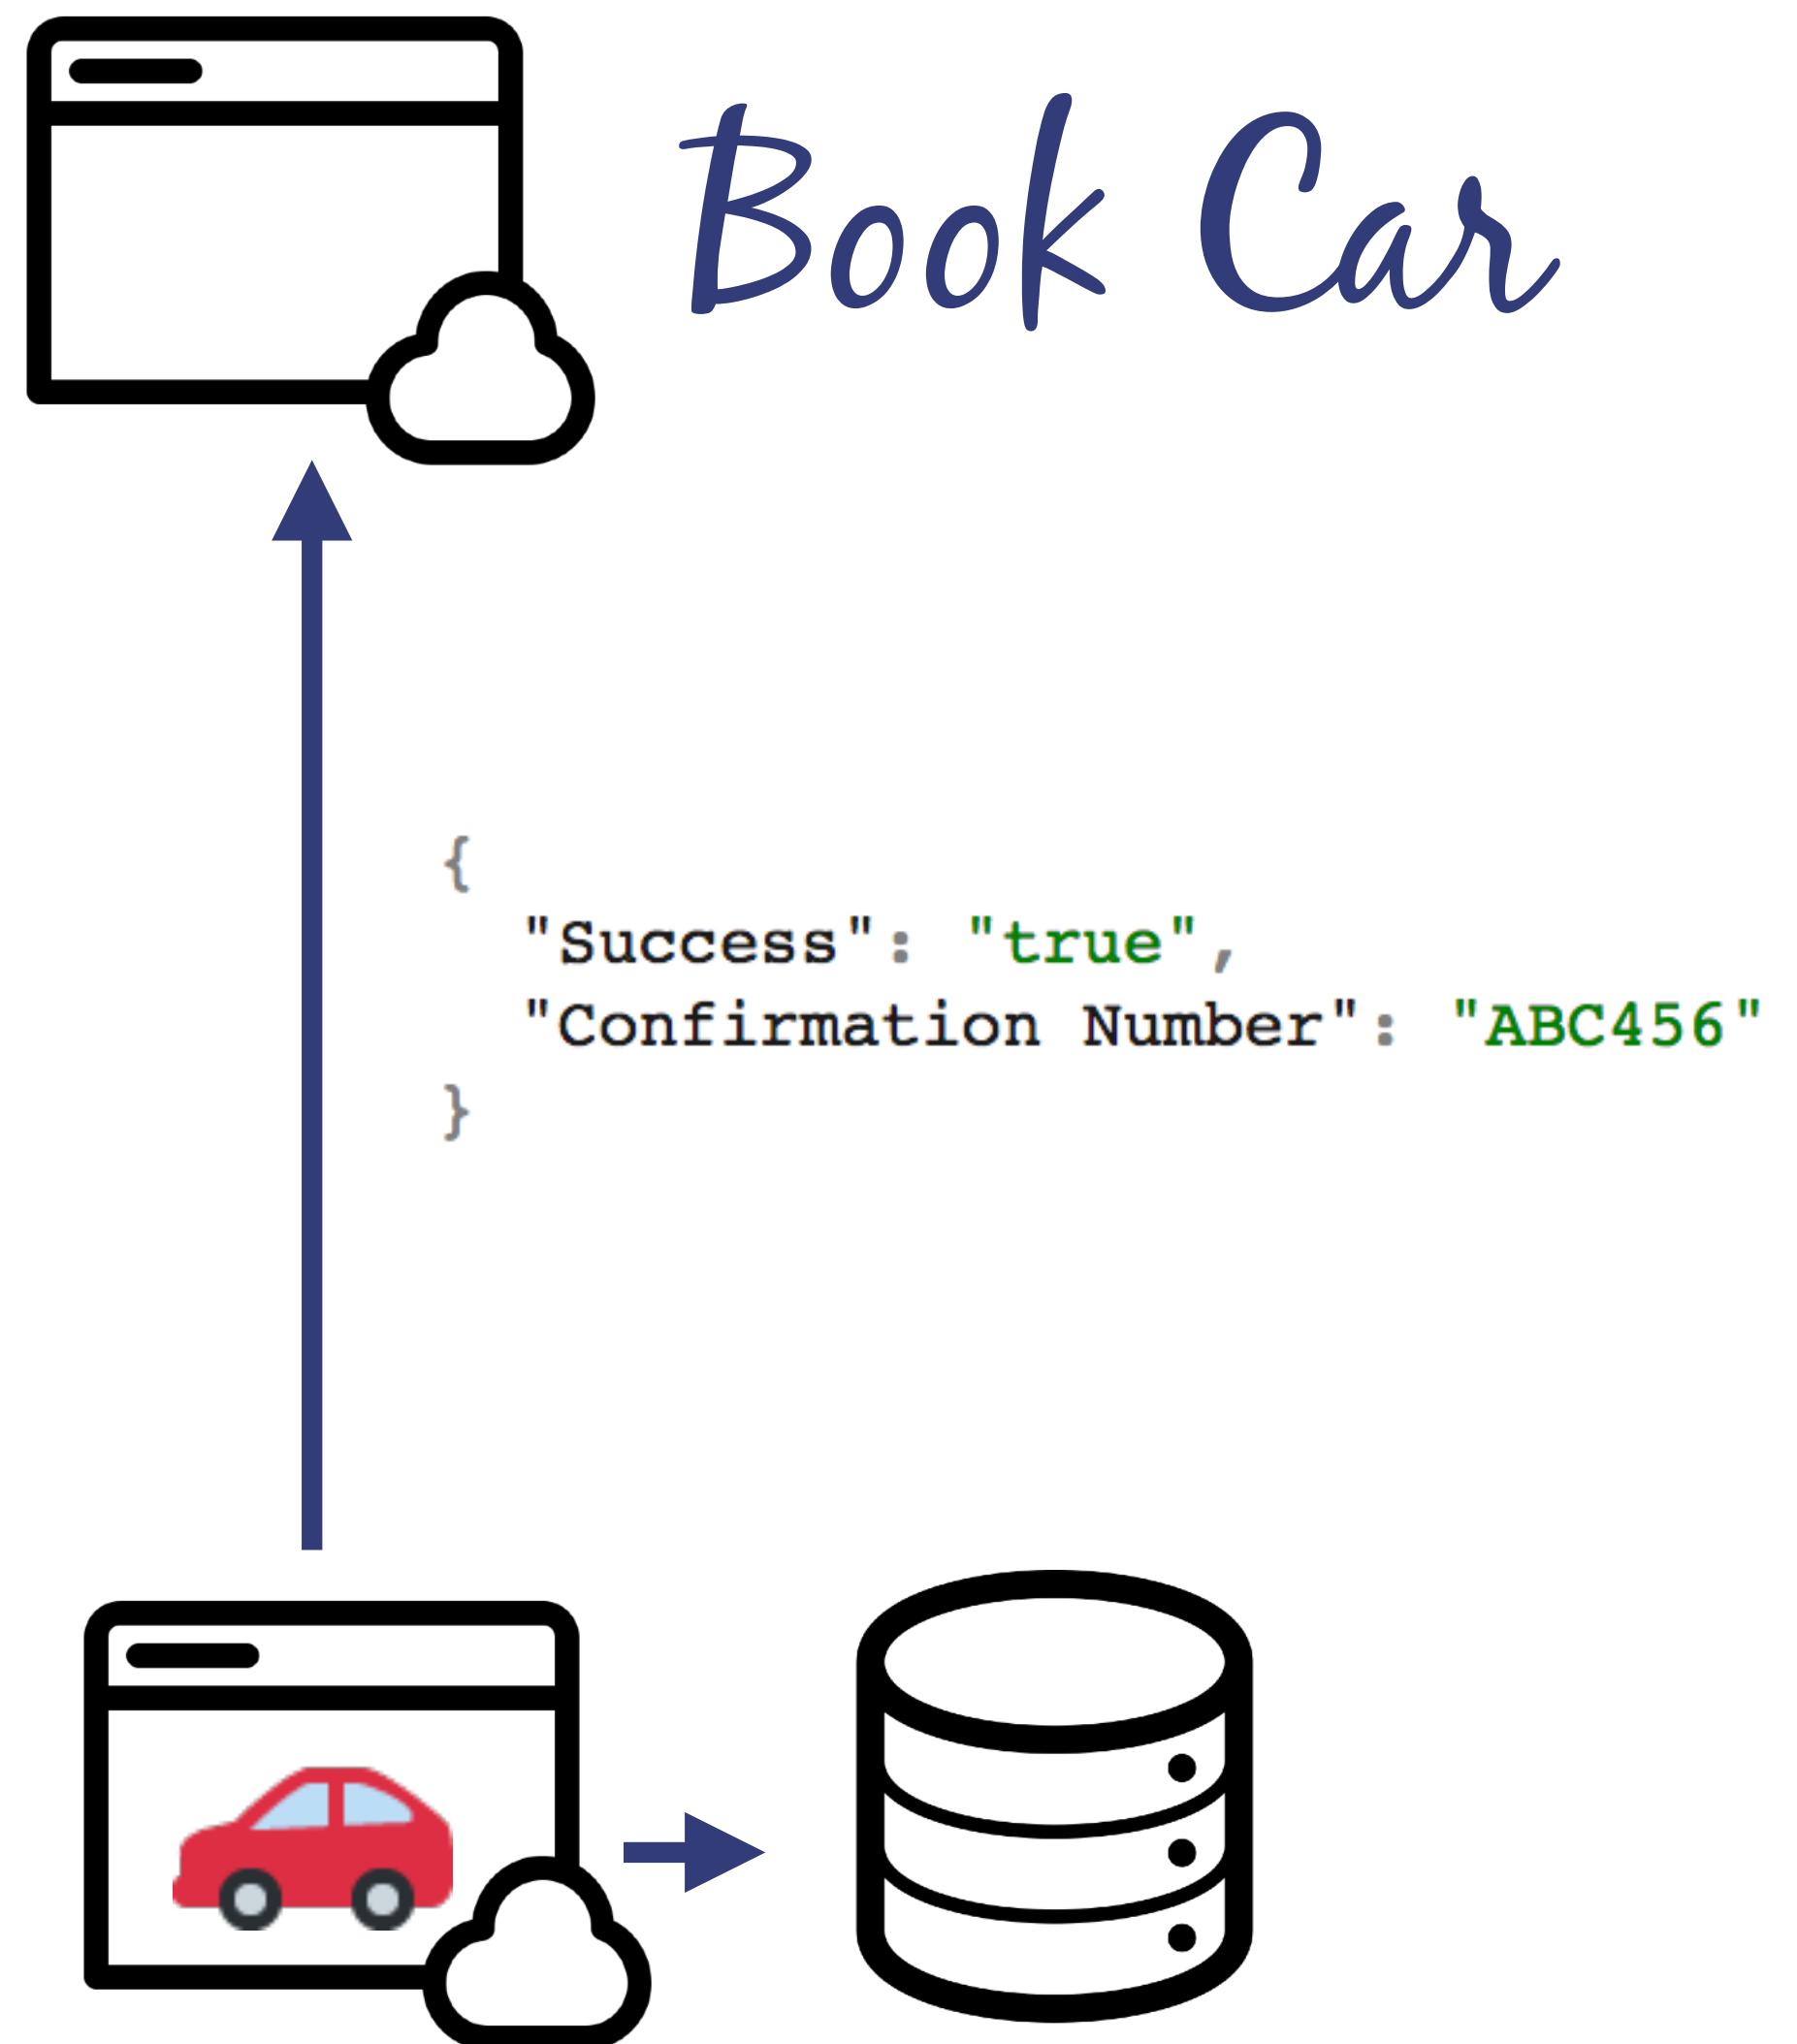
Requests

Must Be Idempotent



Requests

Must Be Idempotent



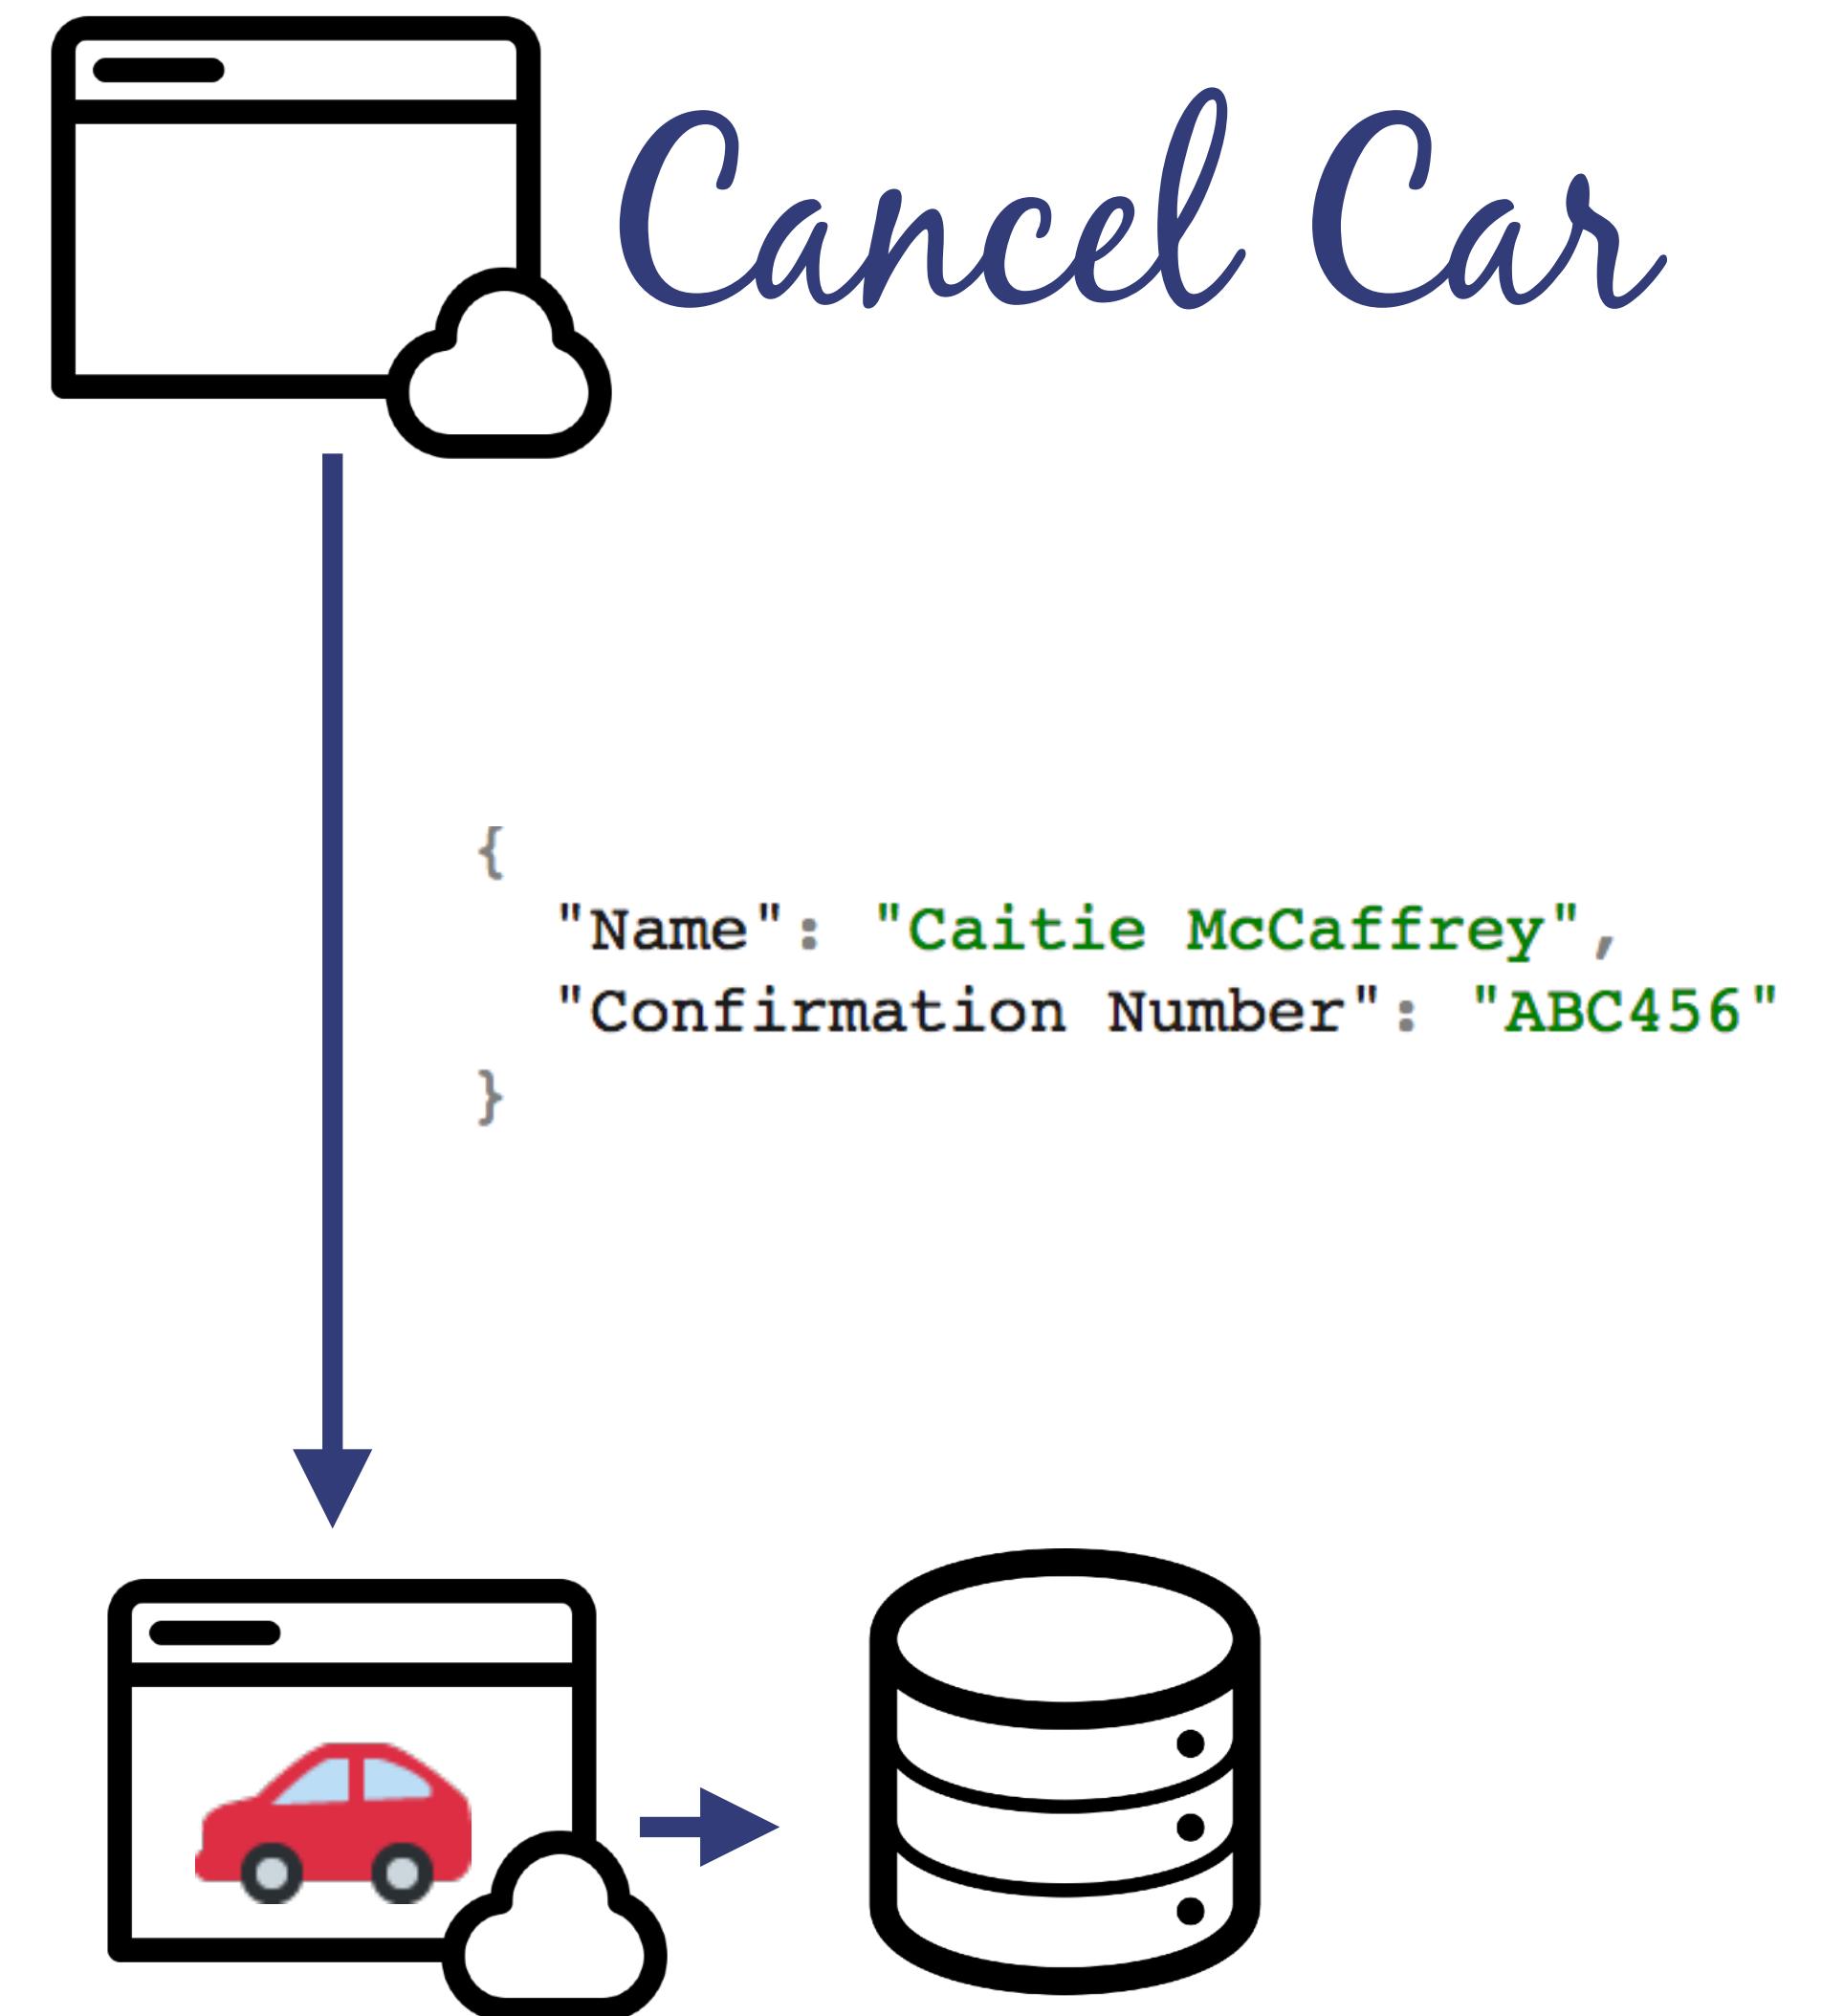
Distributed Saga *Compensating Requests*

Compensating Requests

Semantically undoes the effect of a request

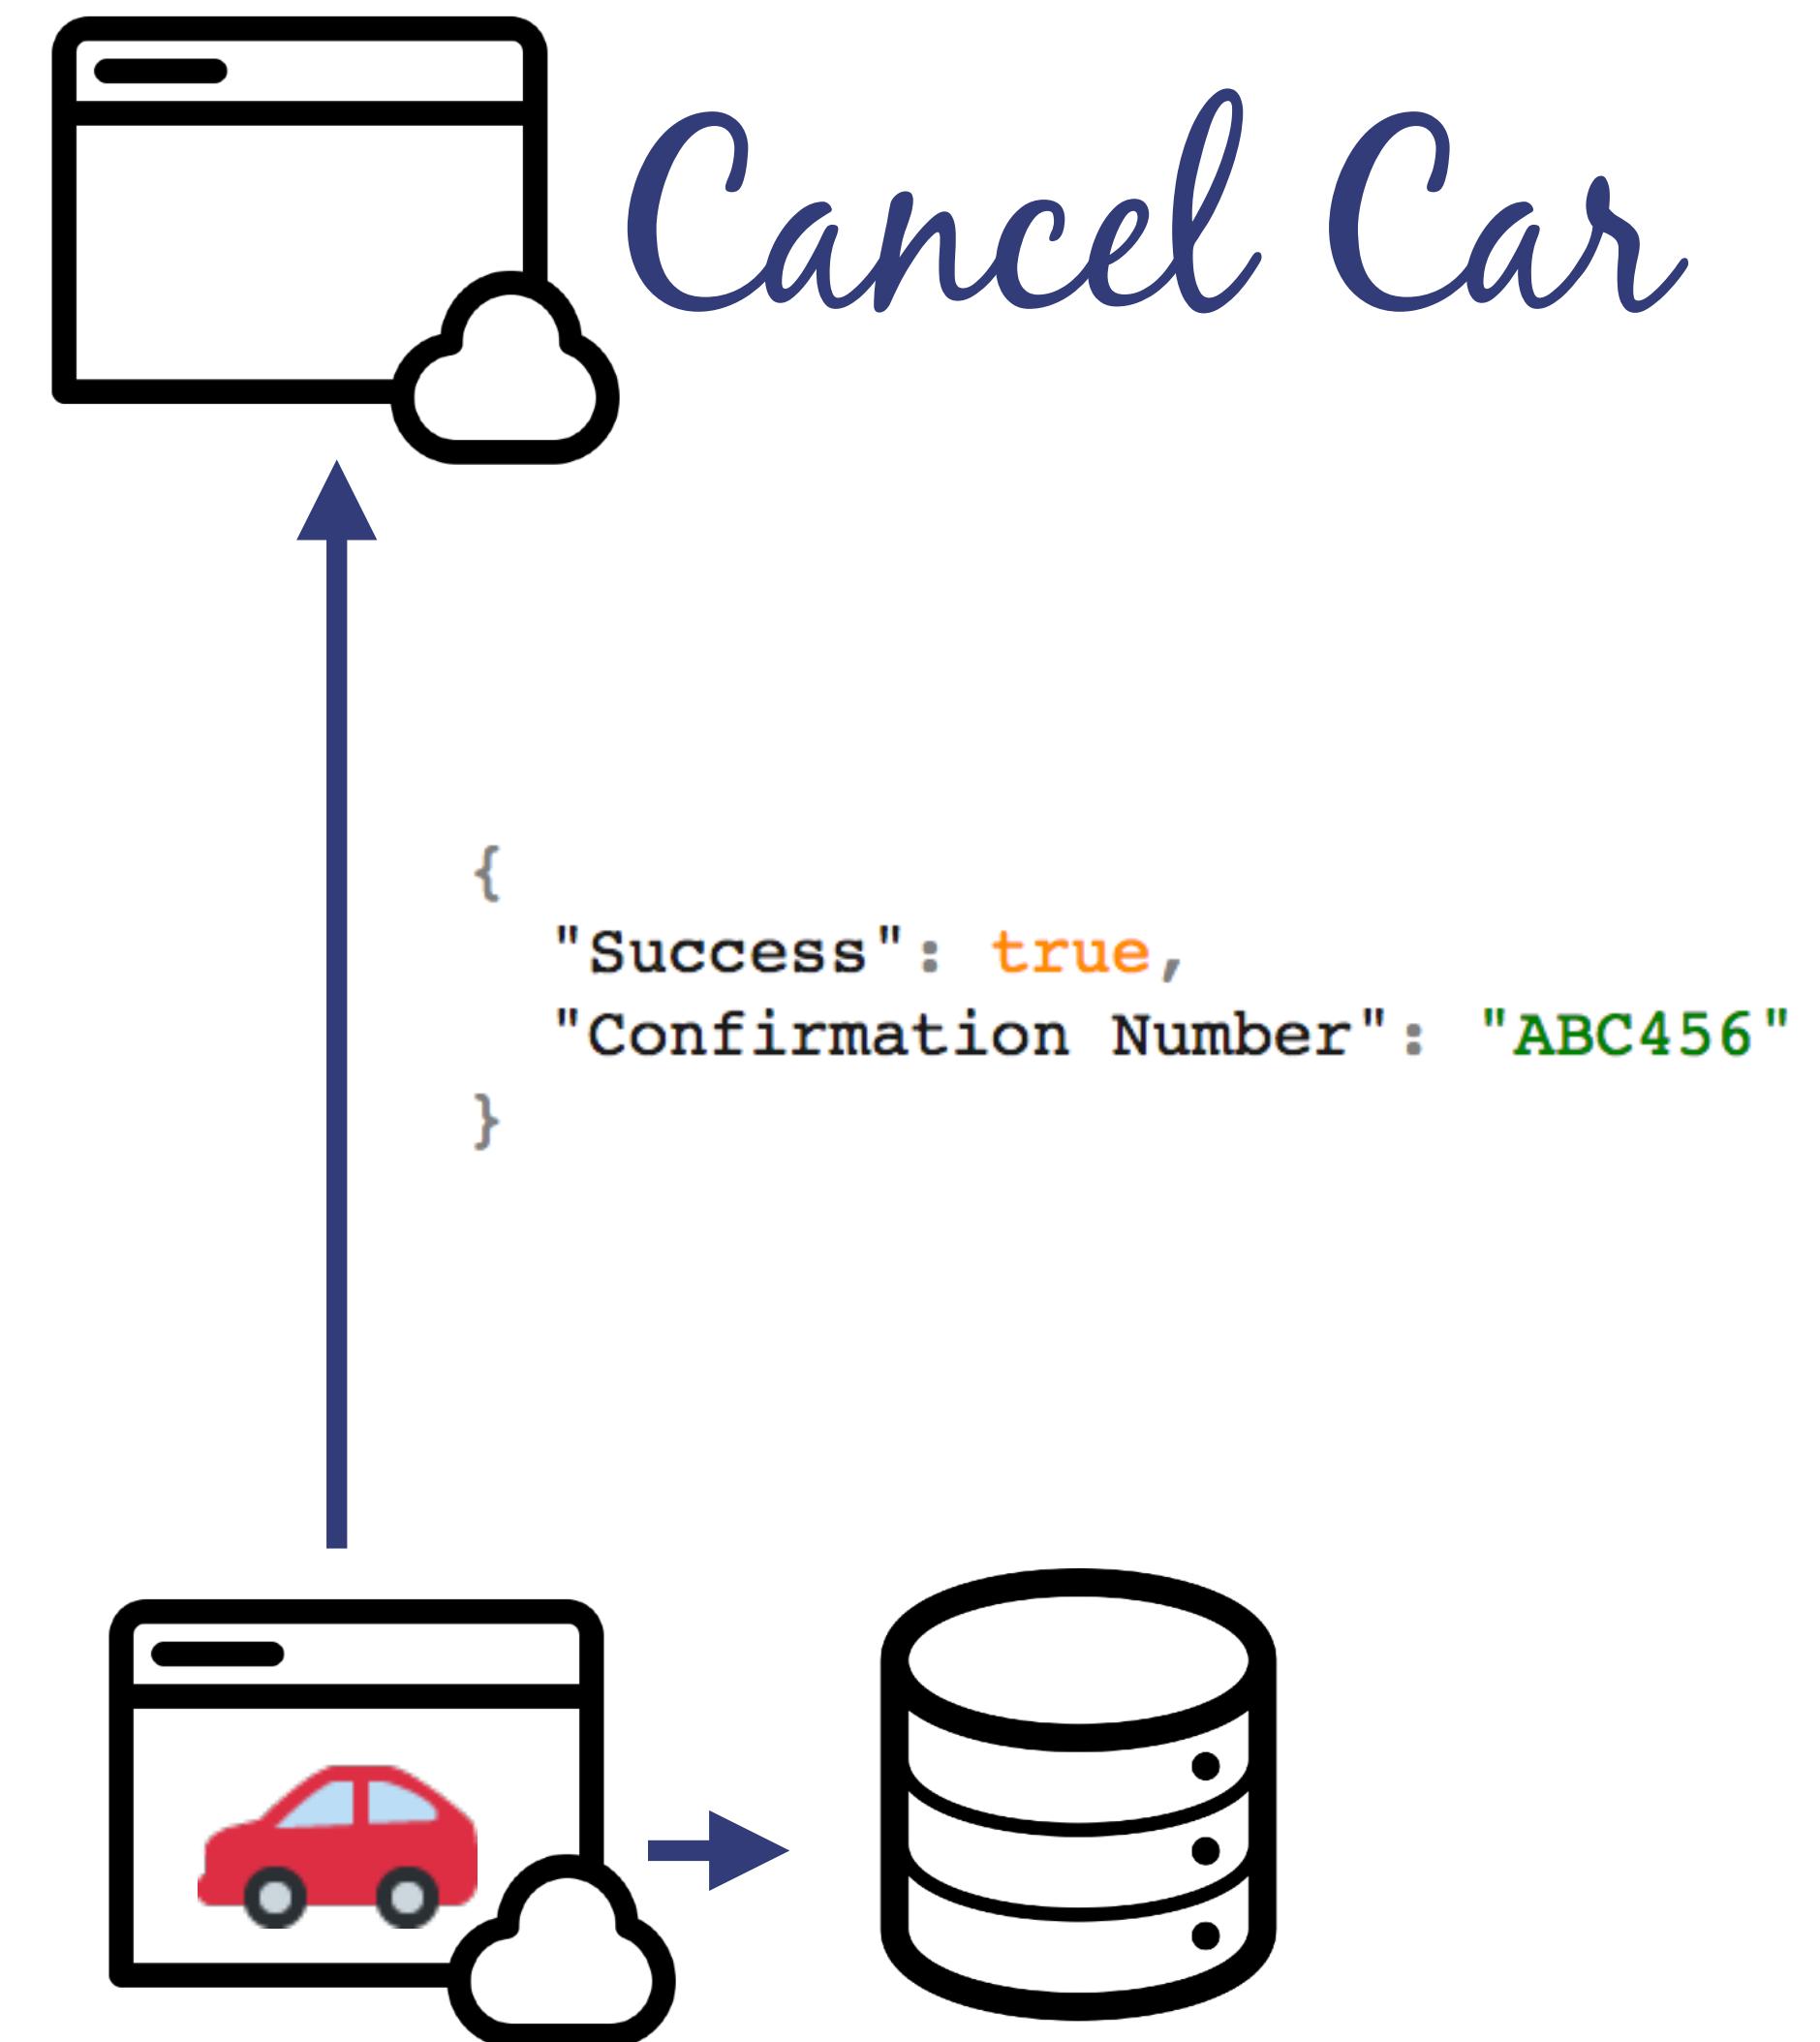
Compensating Requests

Cannot Abort



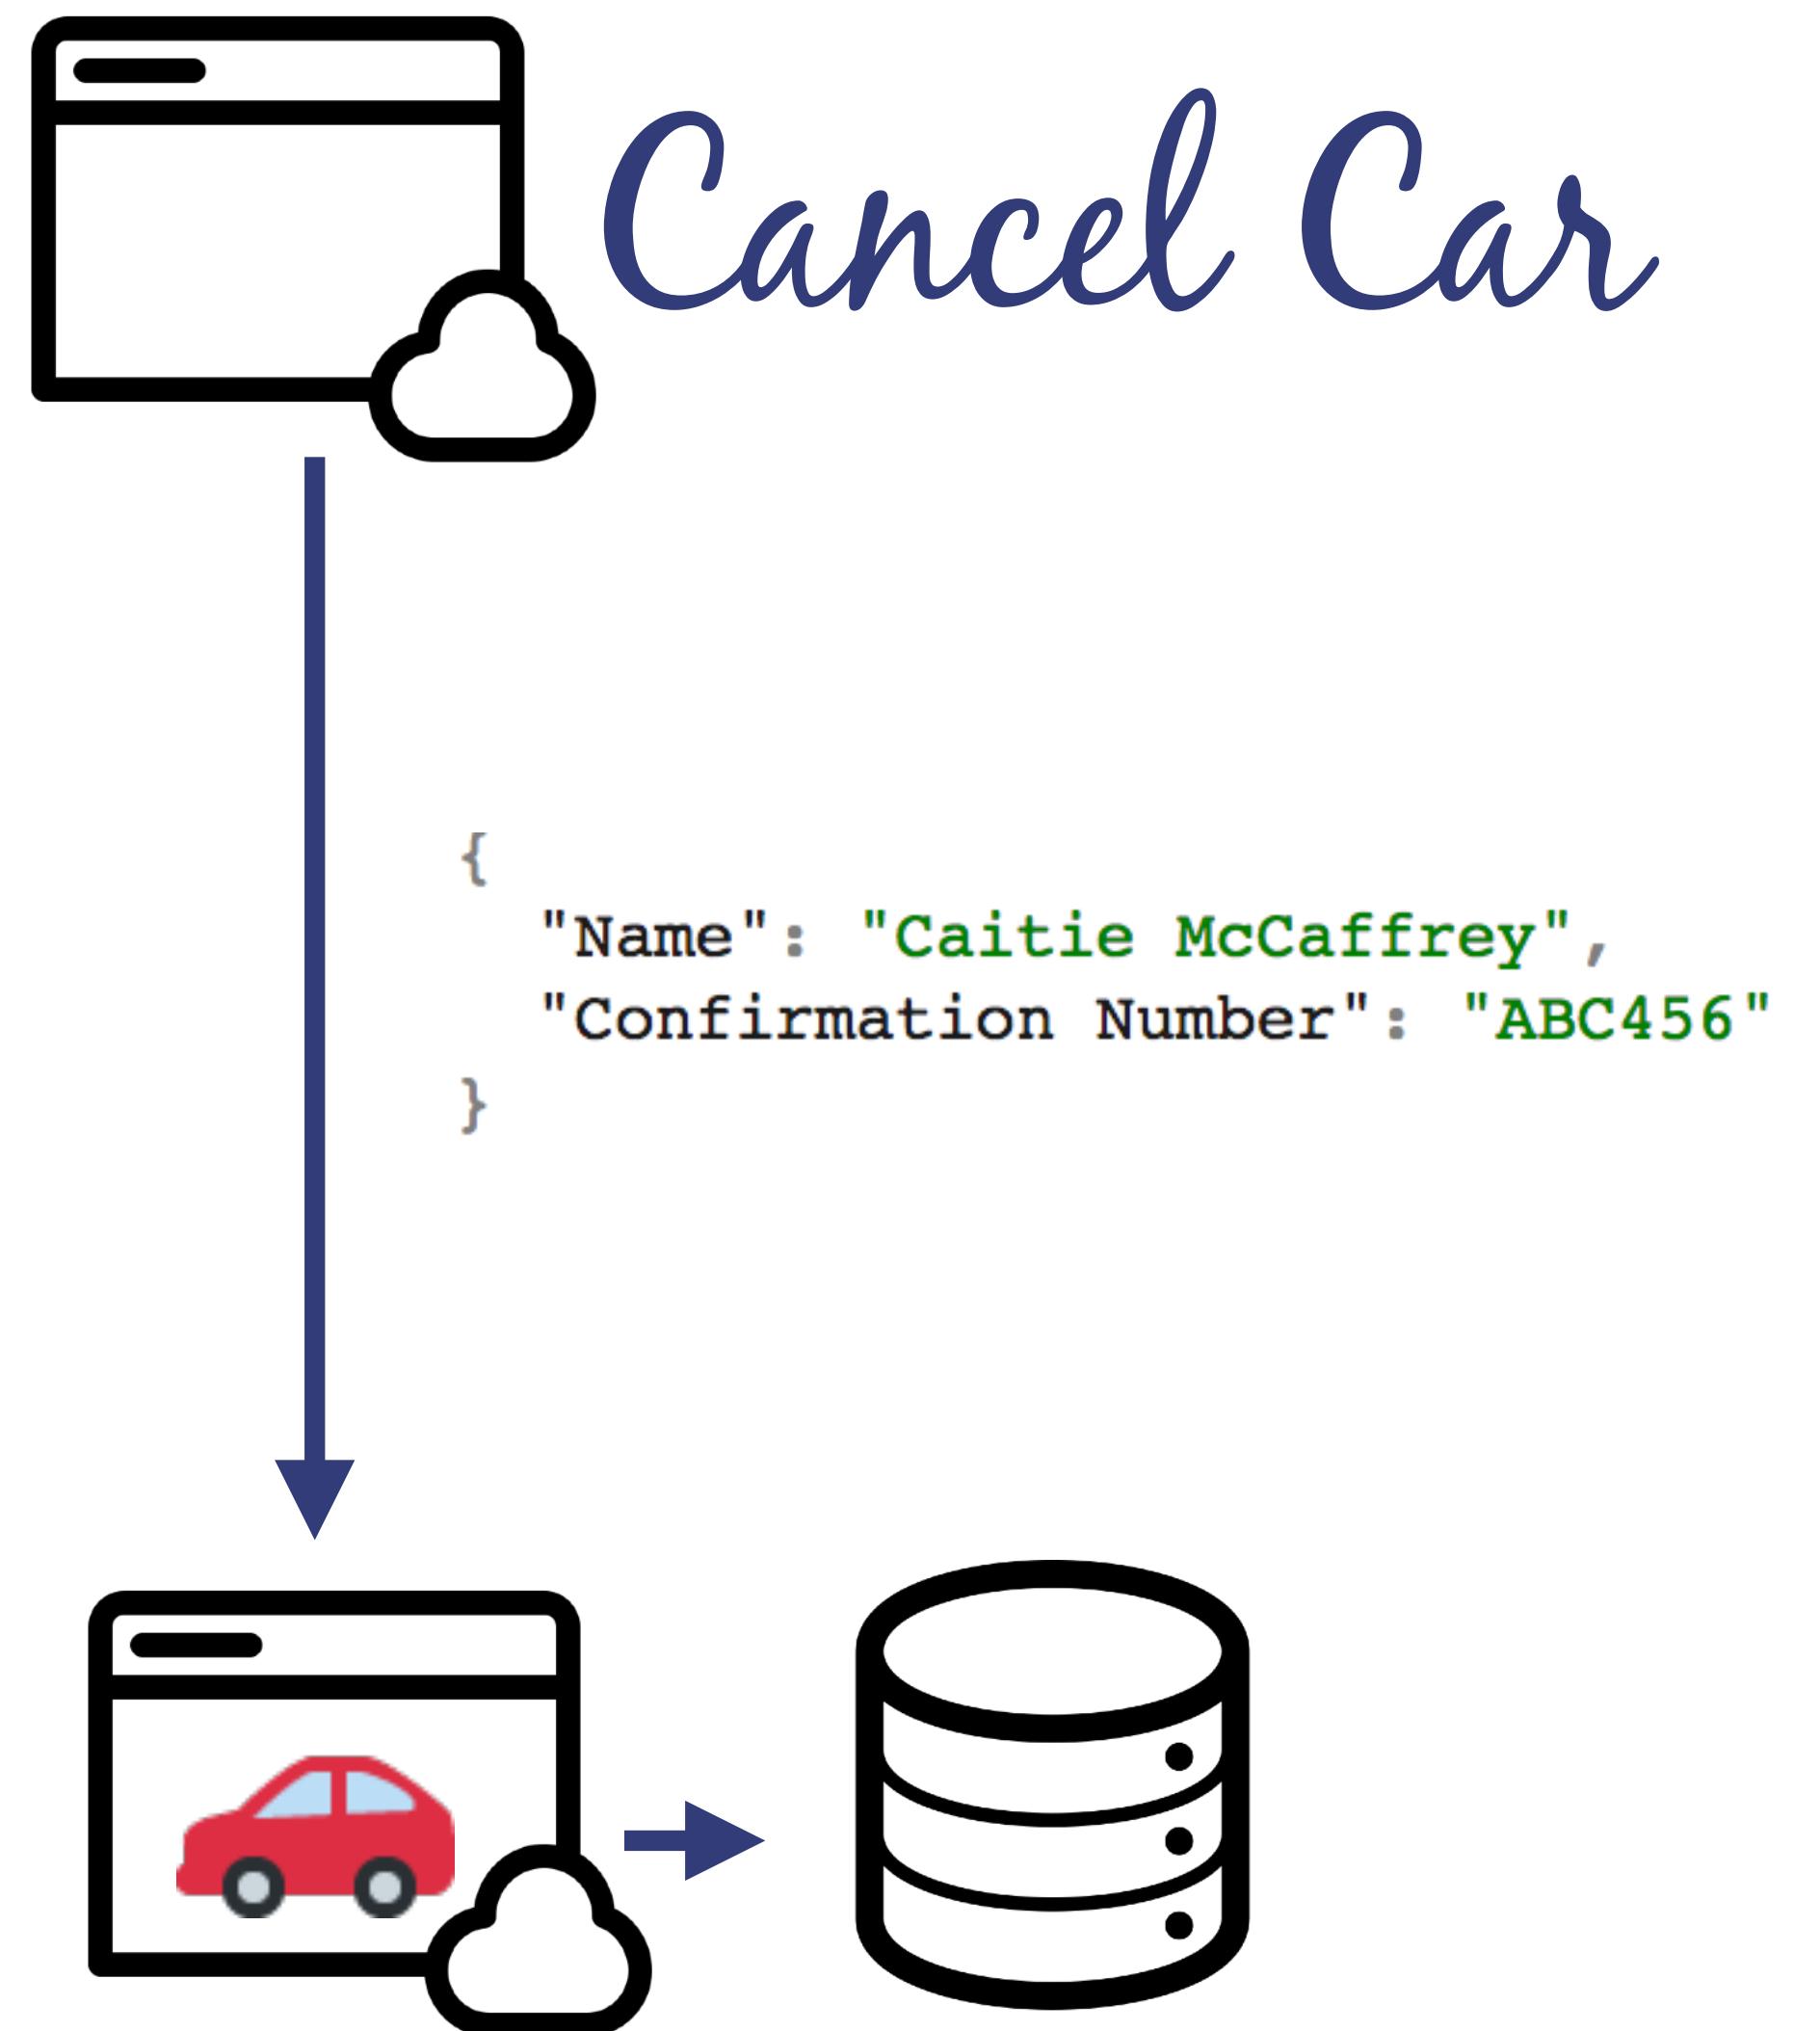
Compensating Requests

Can Not Abort



Compensating Requests

Must Be Idempotent



Compensating Requests

Must Be Commutative
with Requests



Book Car



Cancel Car

is the same as



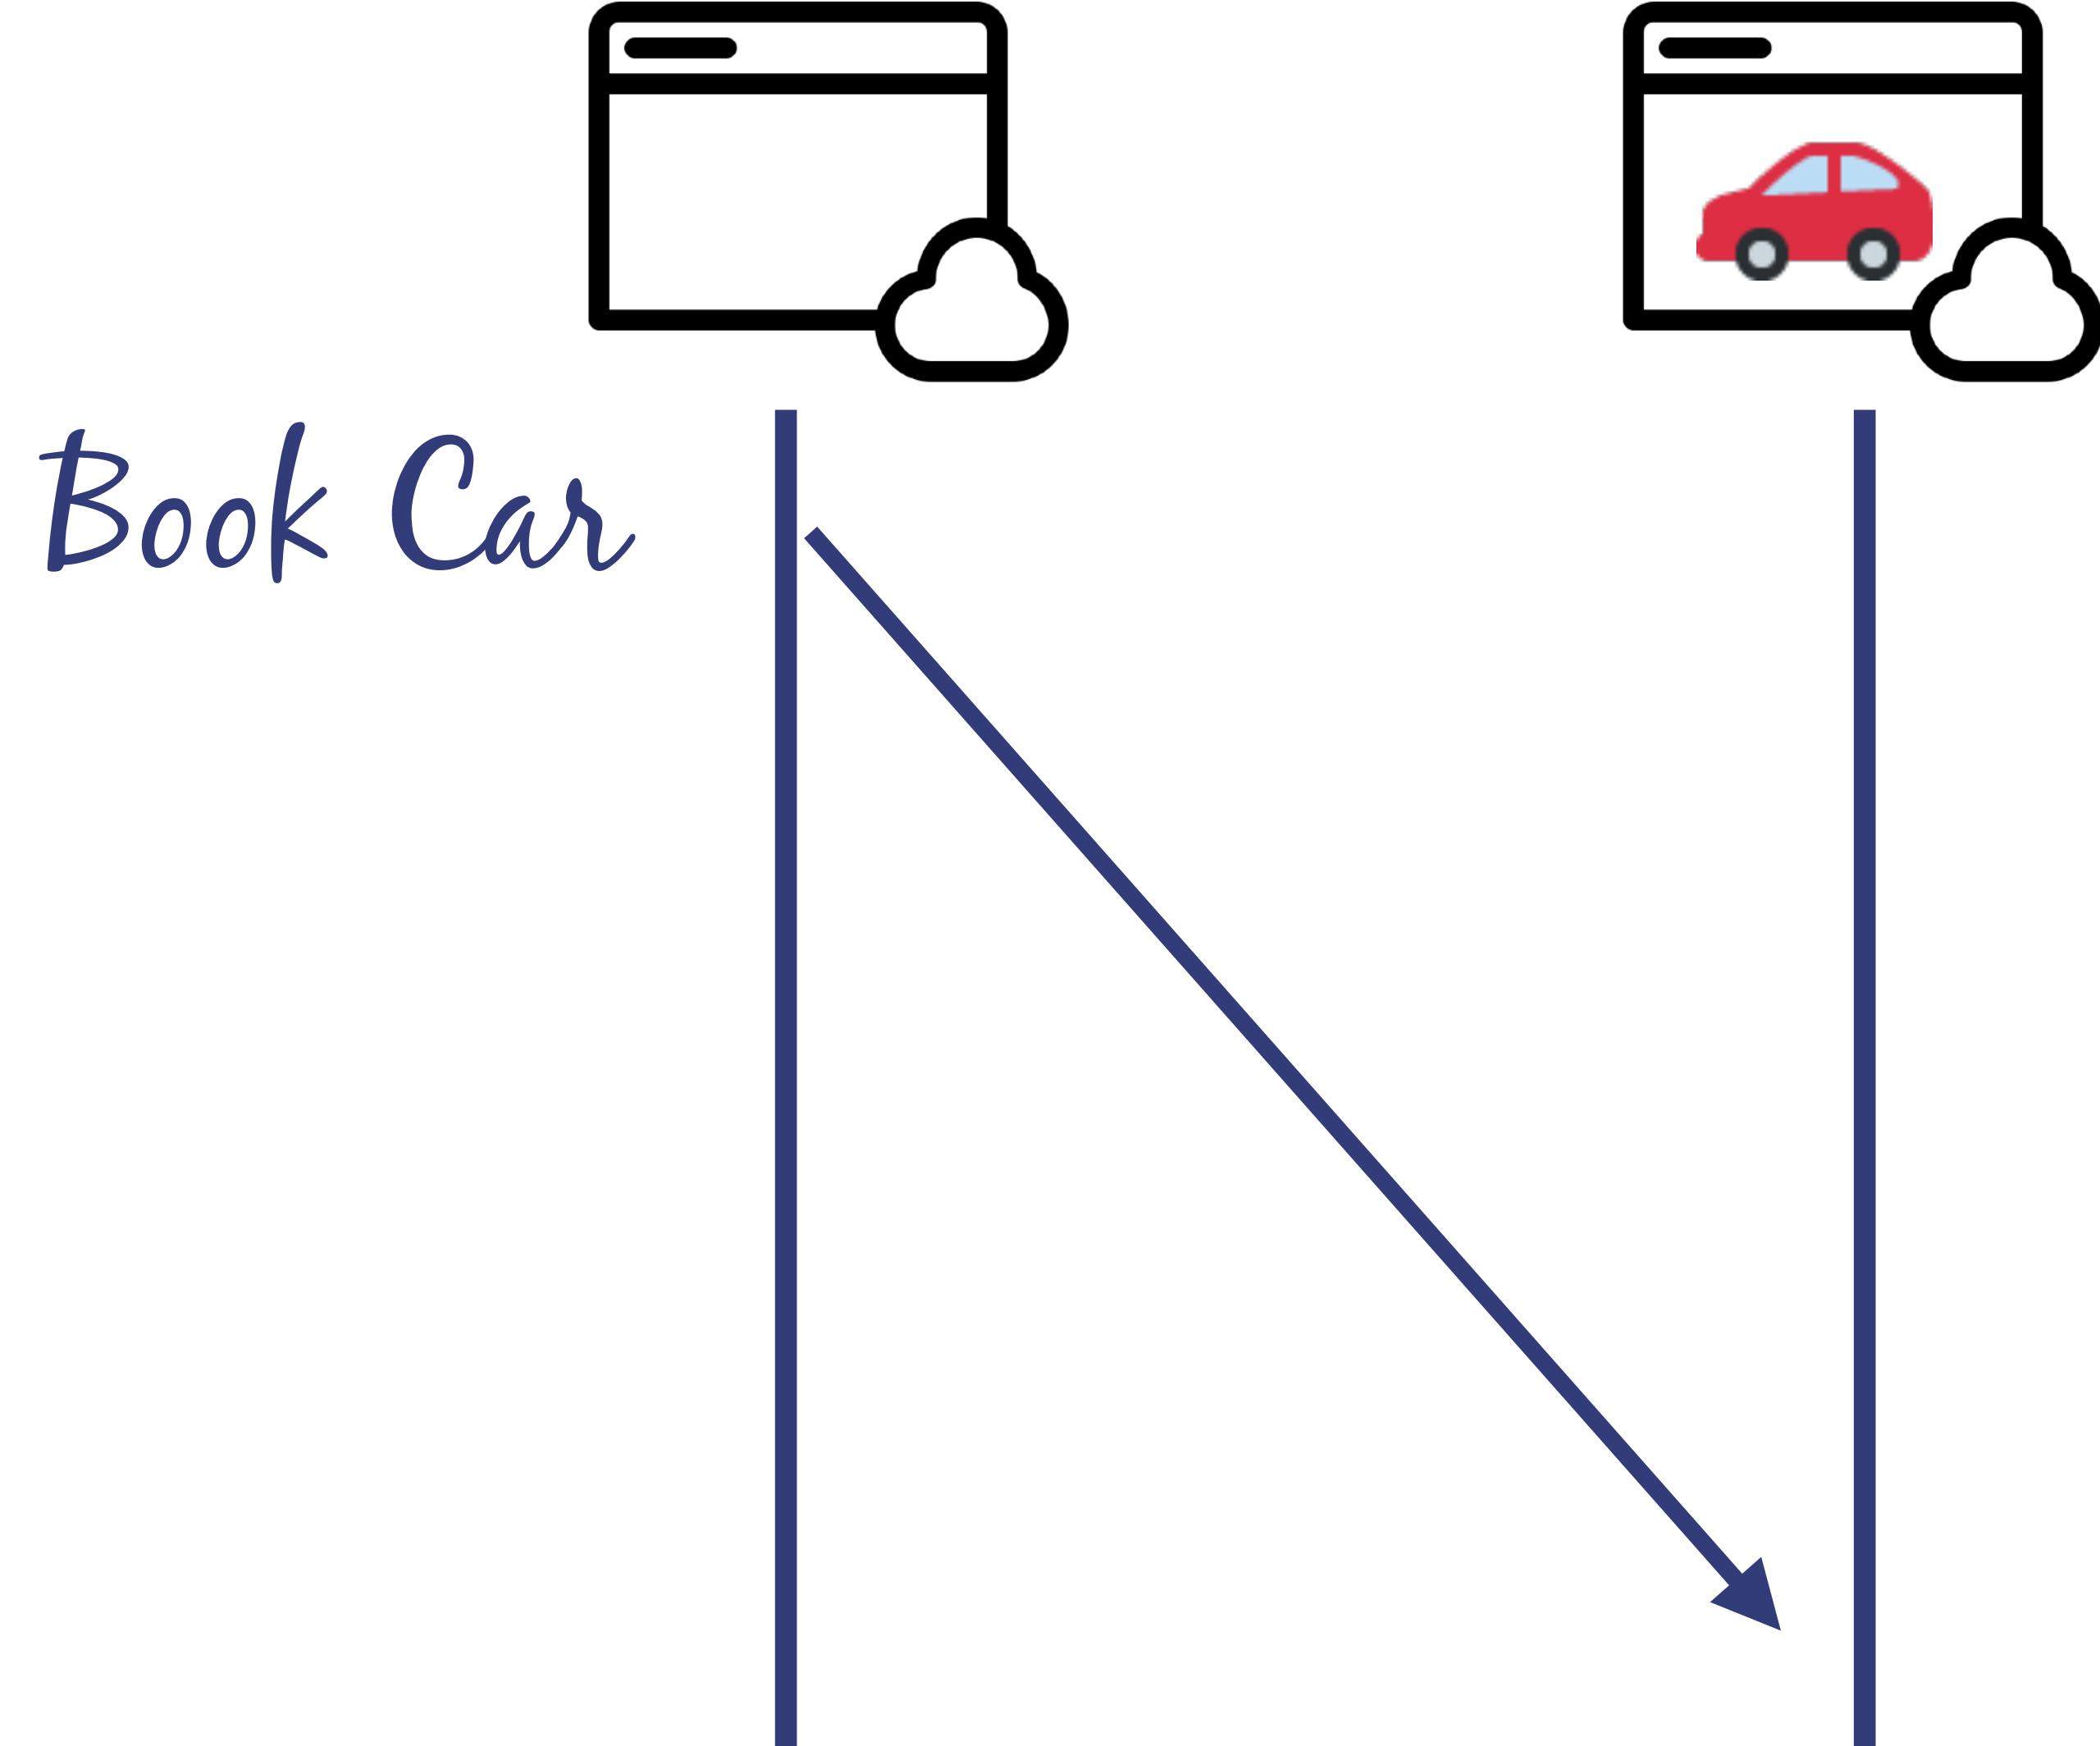
Cancel Car



Book Car

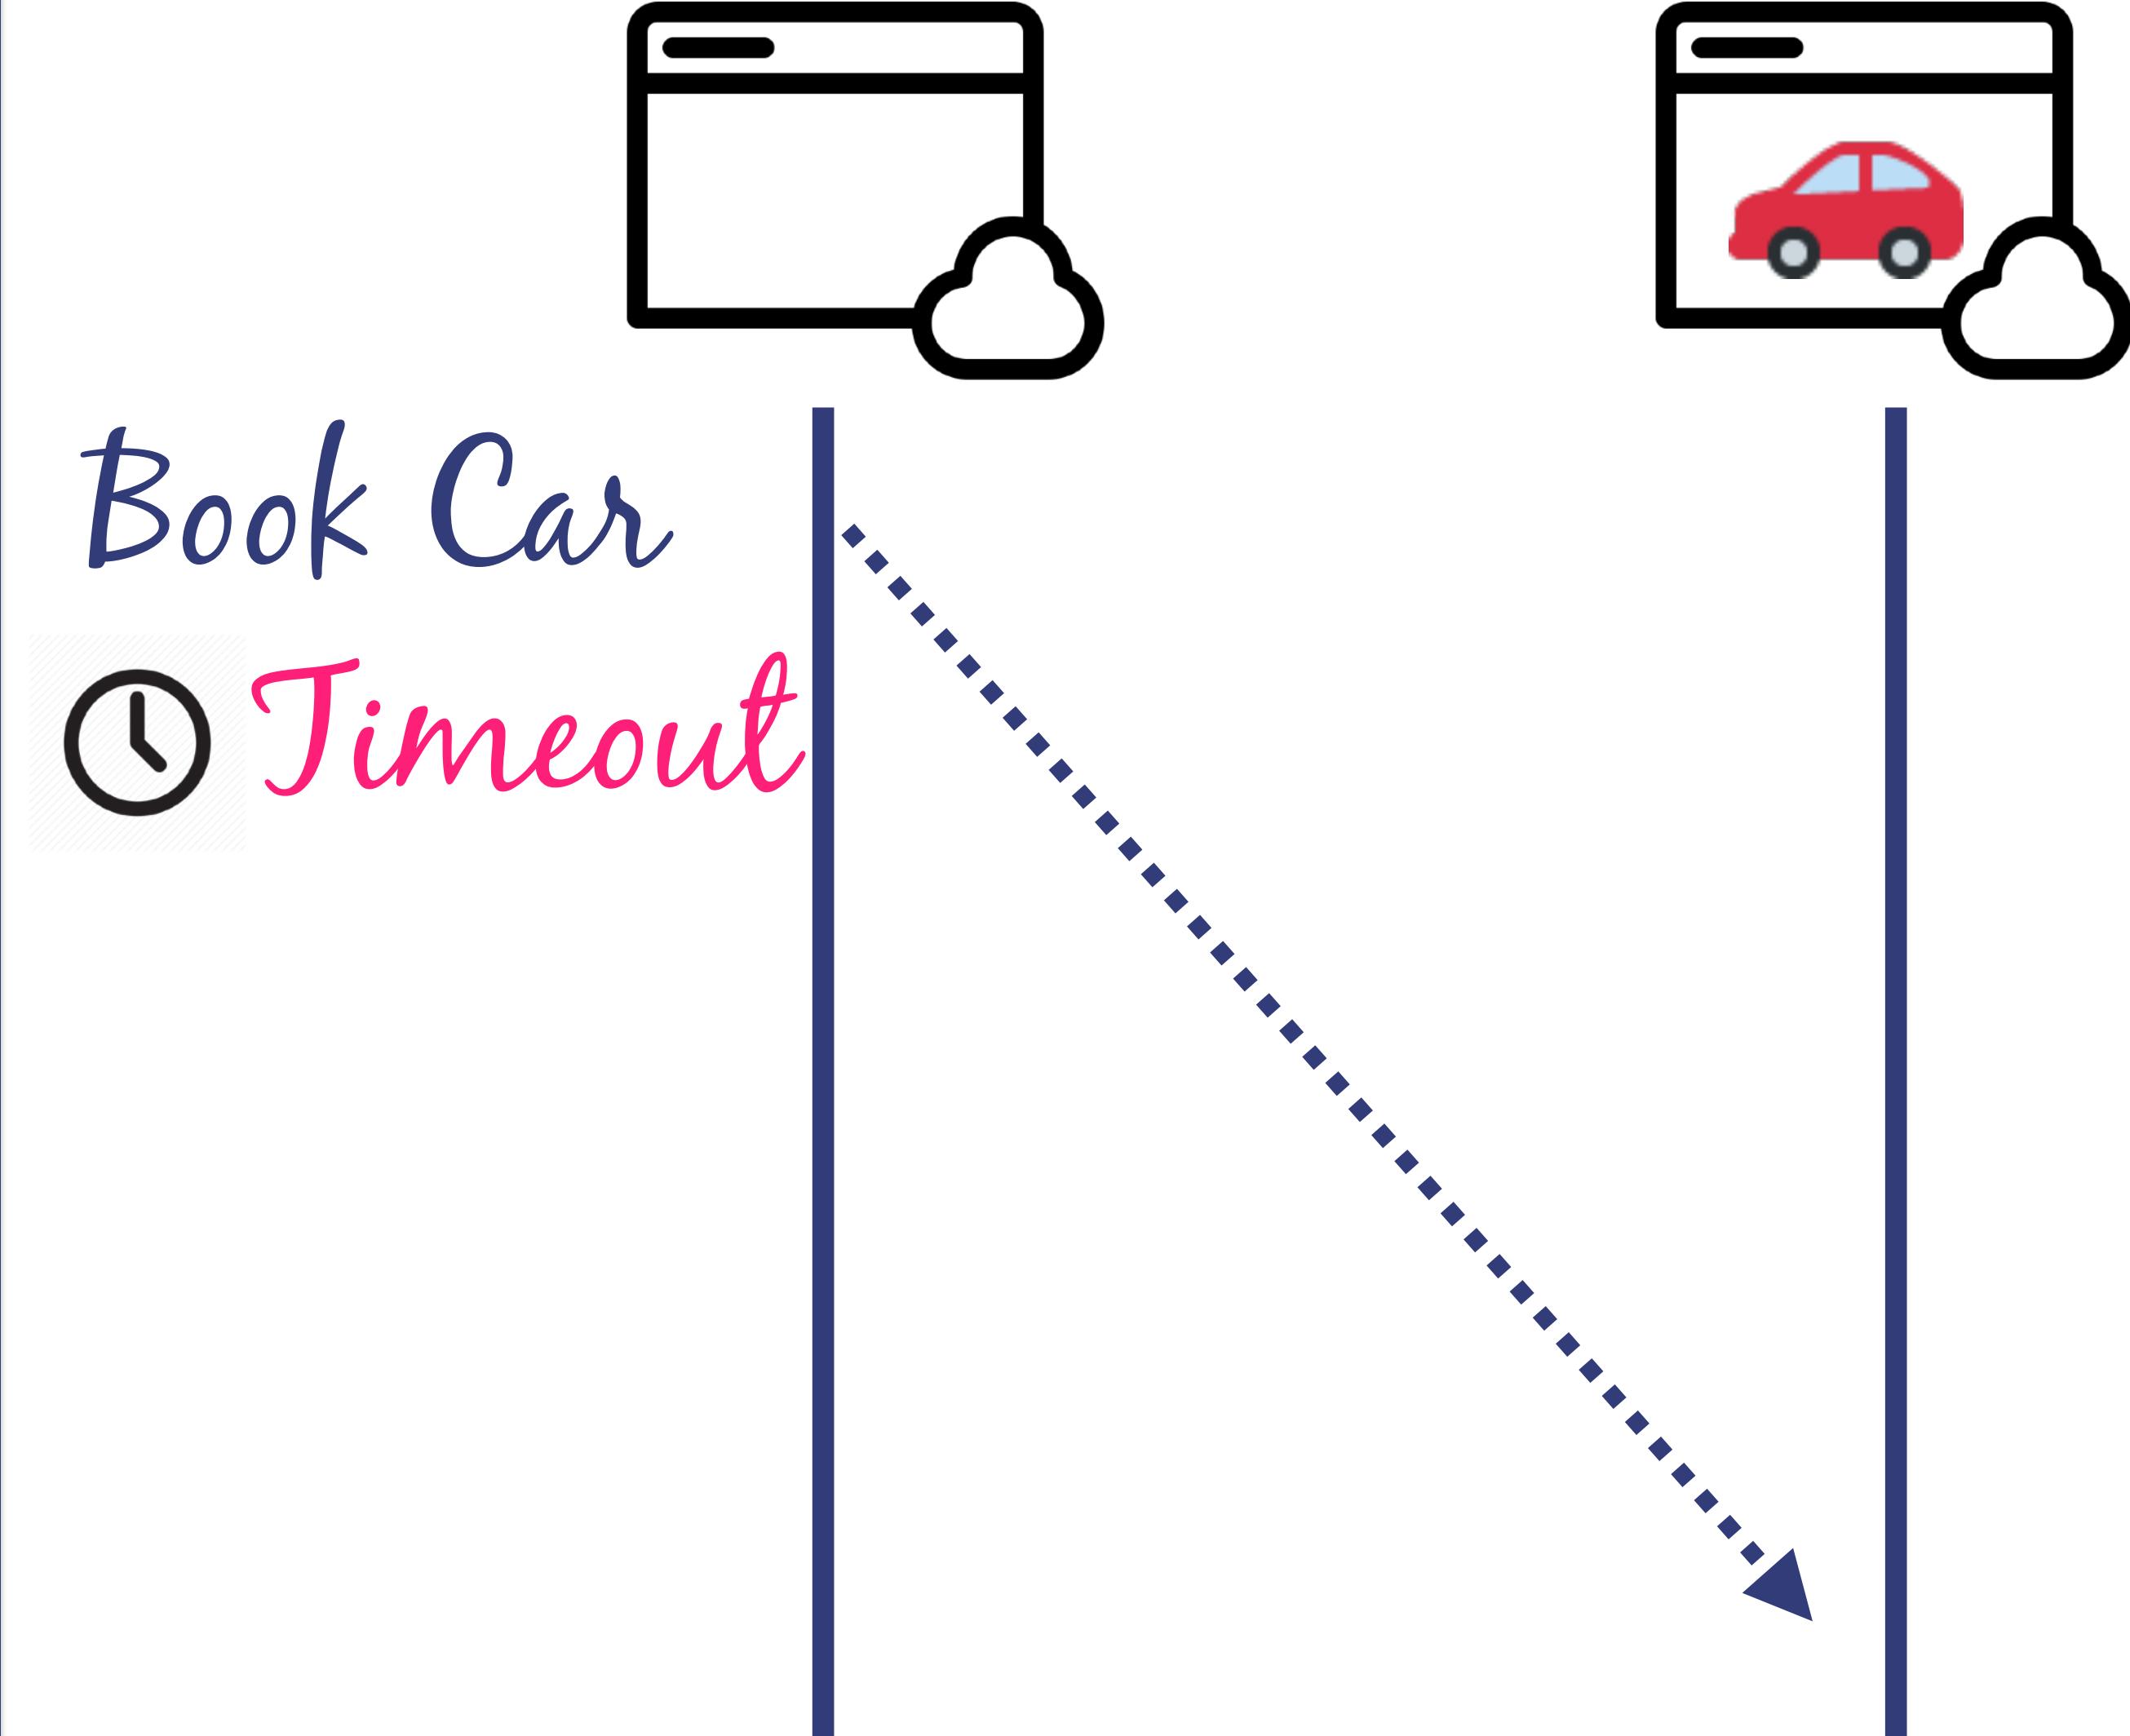
Compensating Requests

Must Be Commutative
with Requests



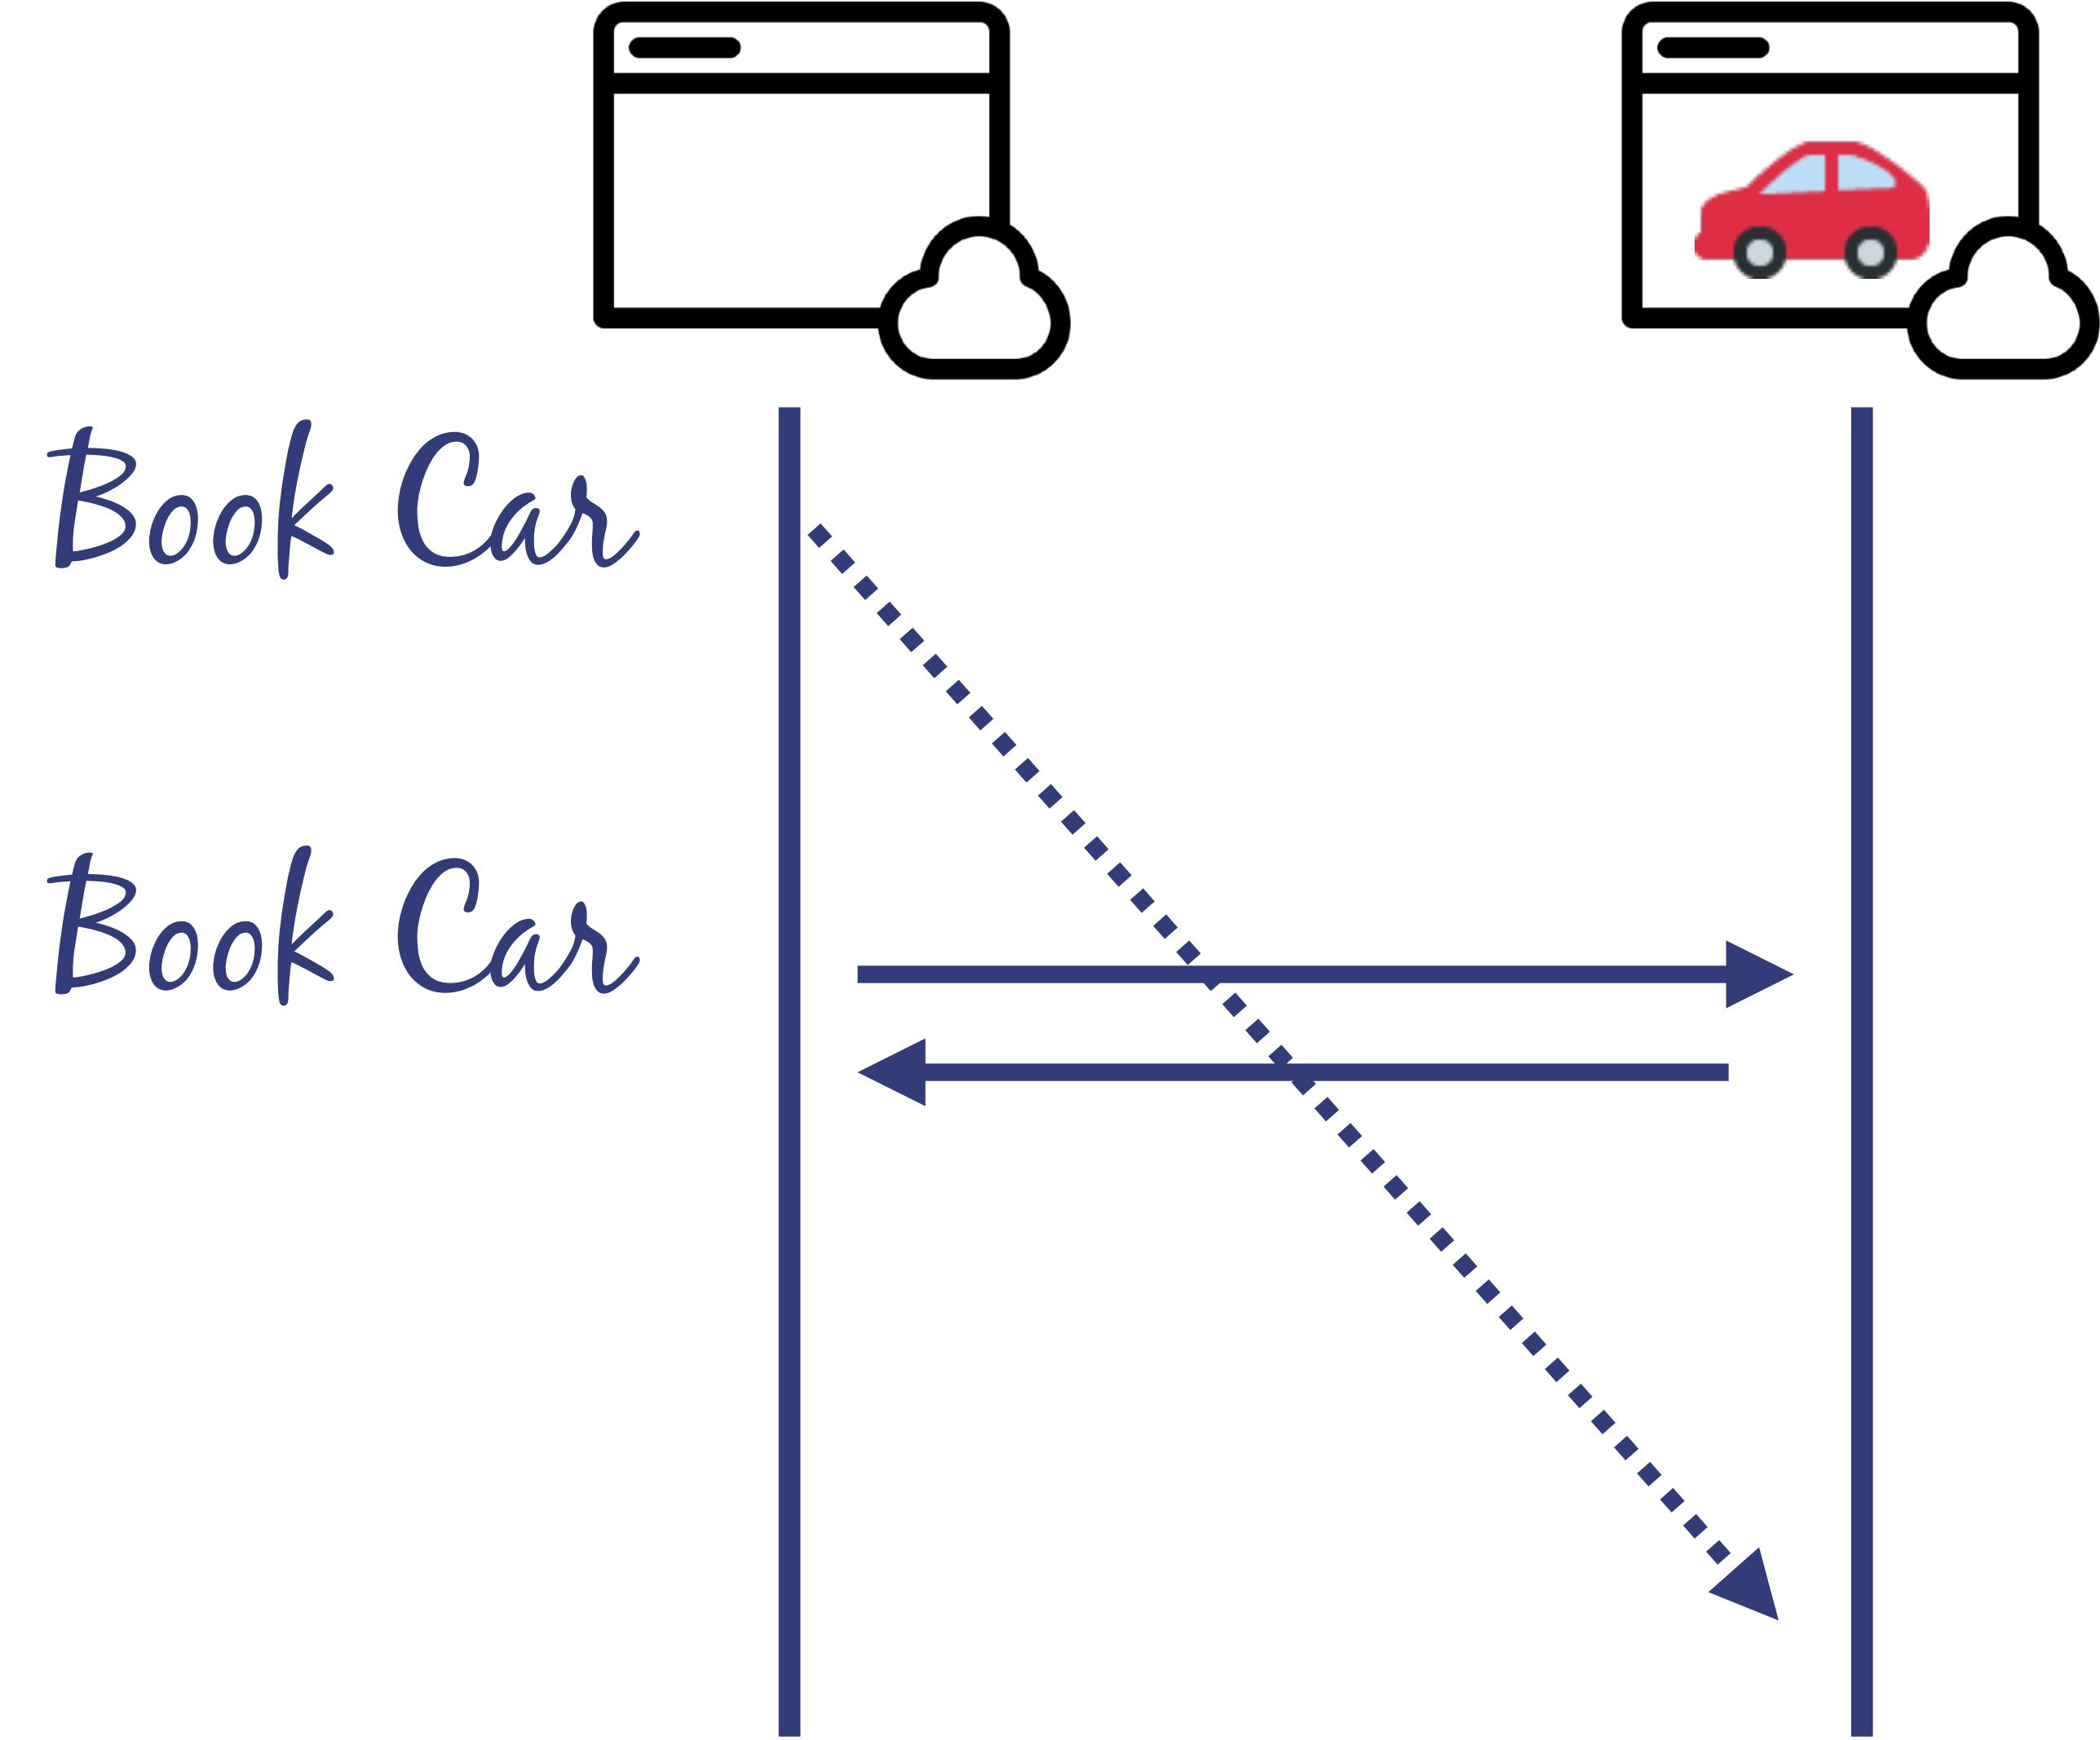
Compensating Requests

Must Be Commutative
with Requests



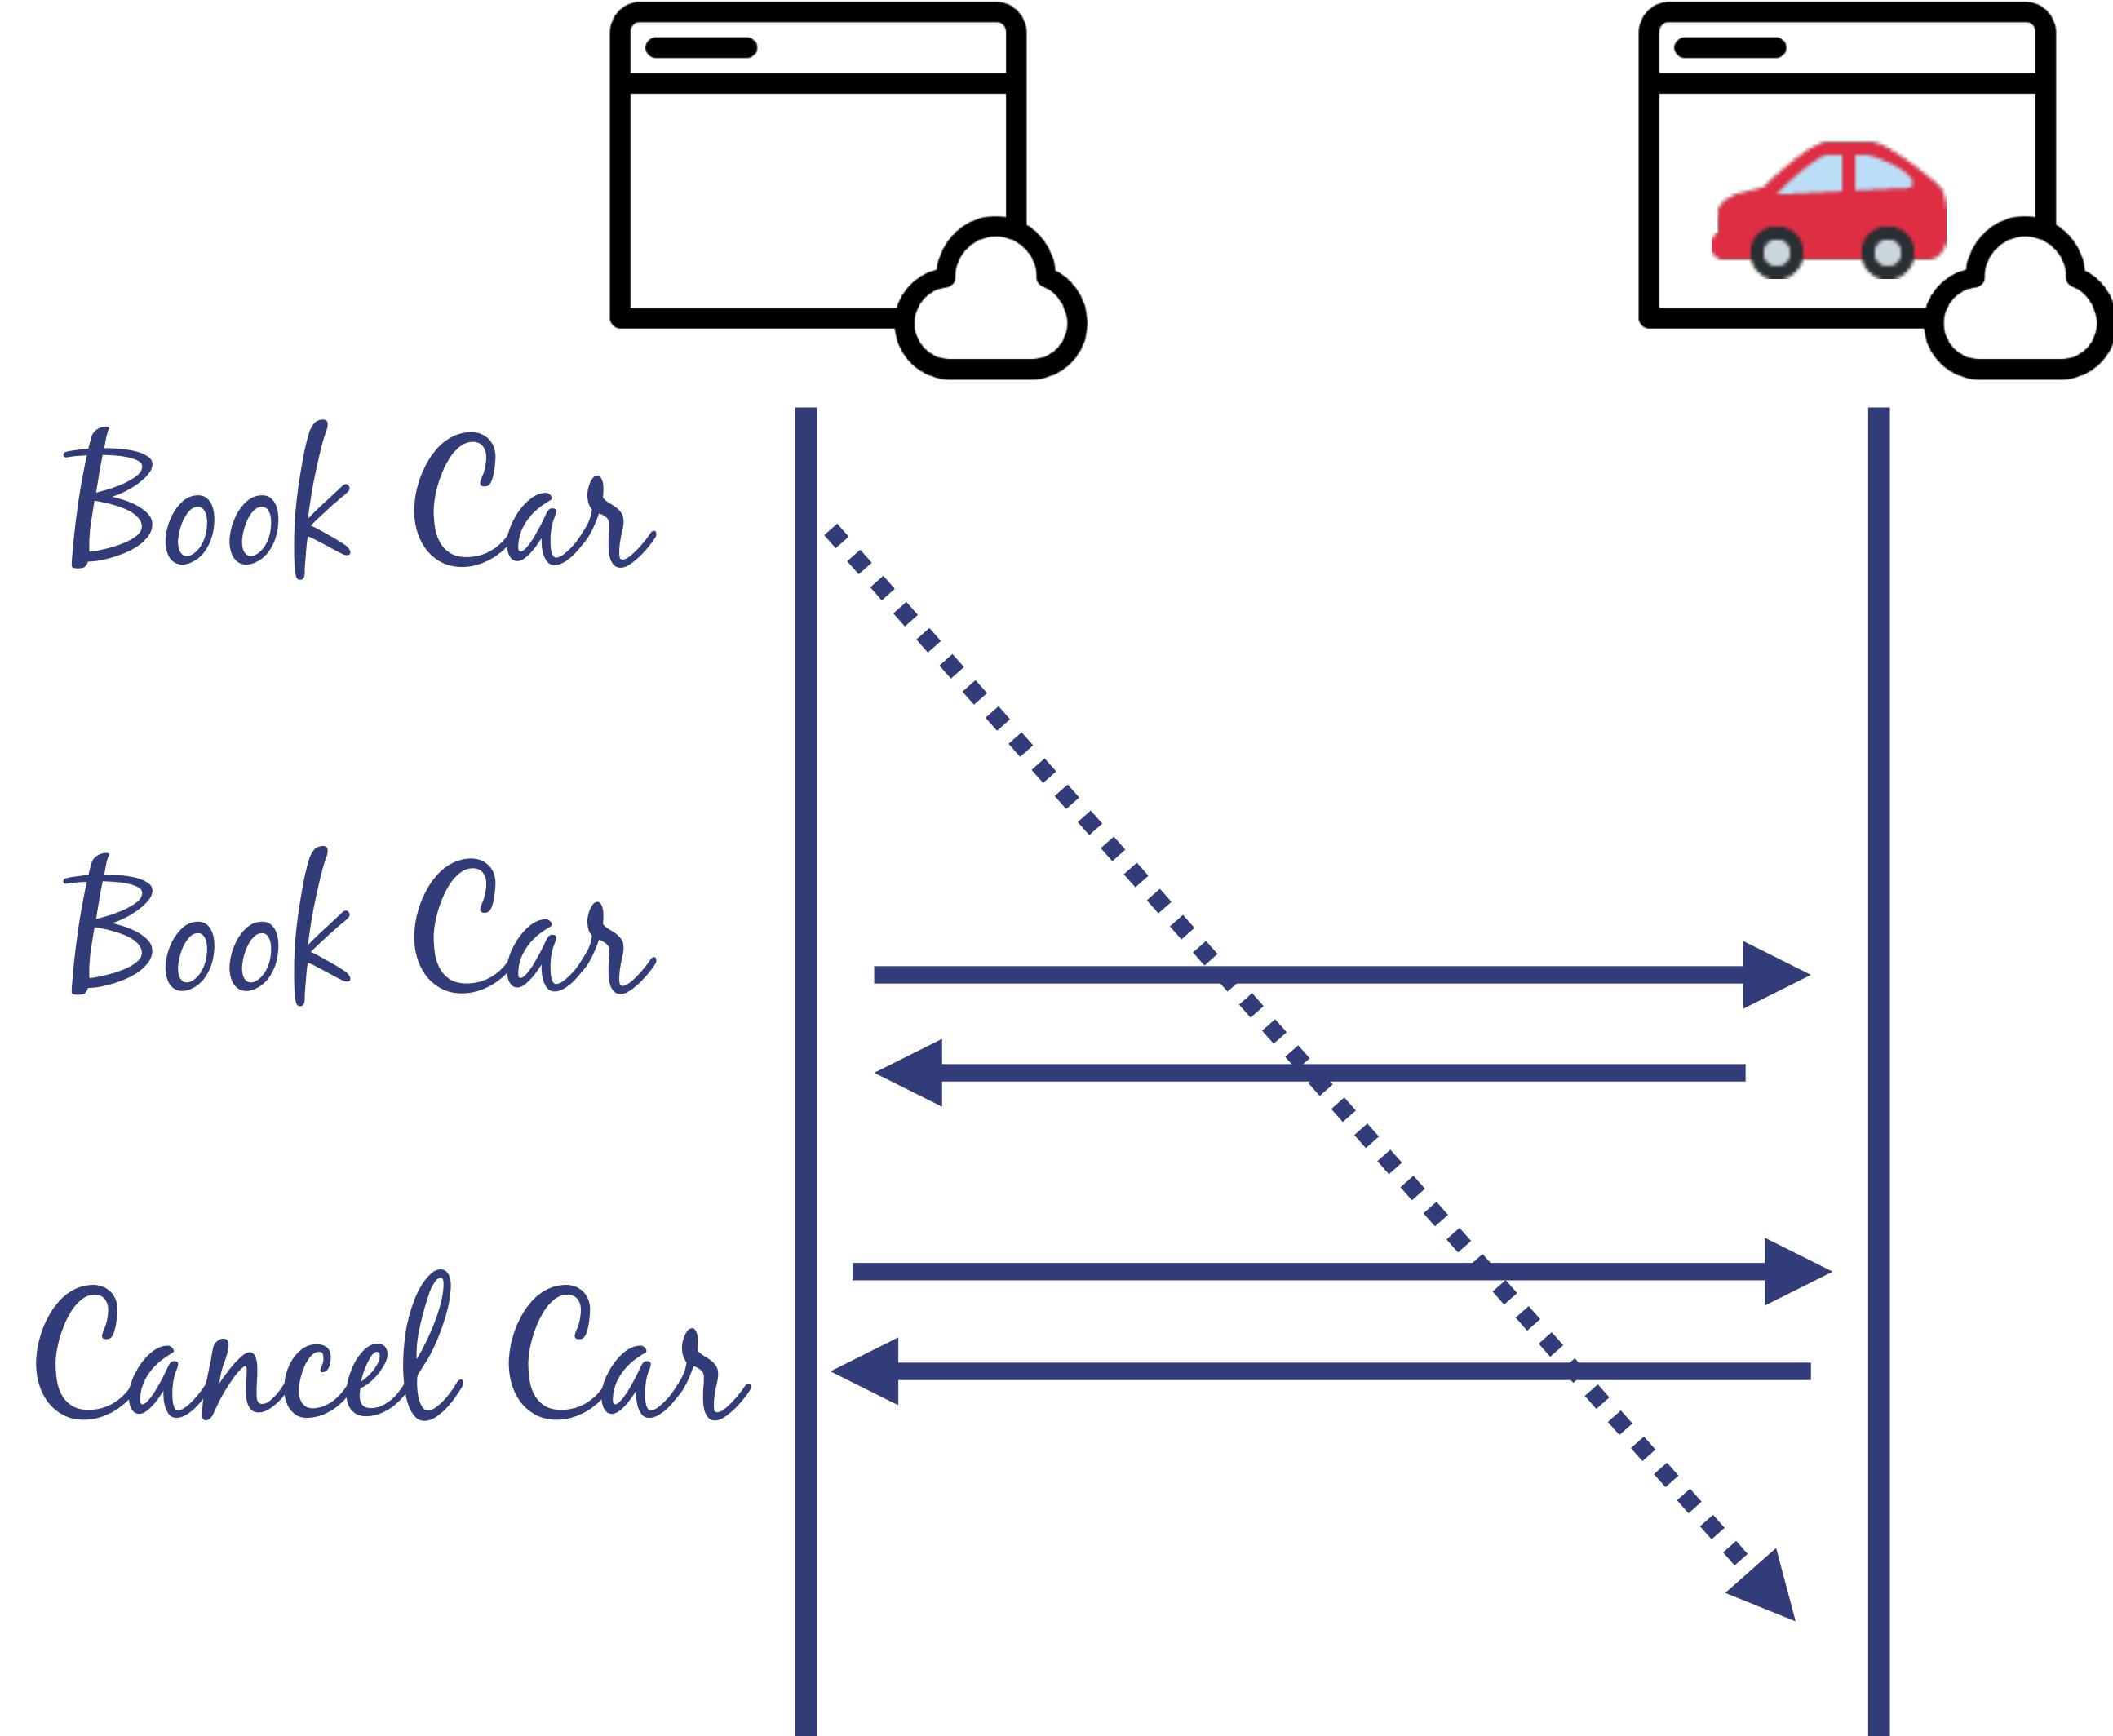
Compensating Requests

Must Be Commutative
with Requests



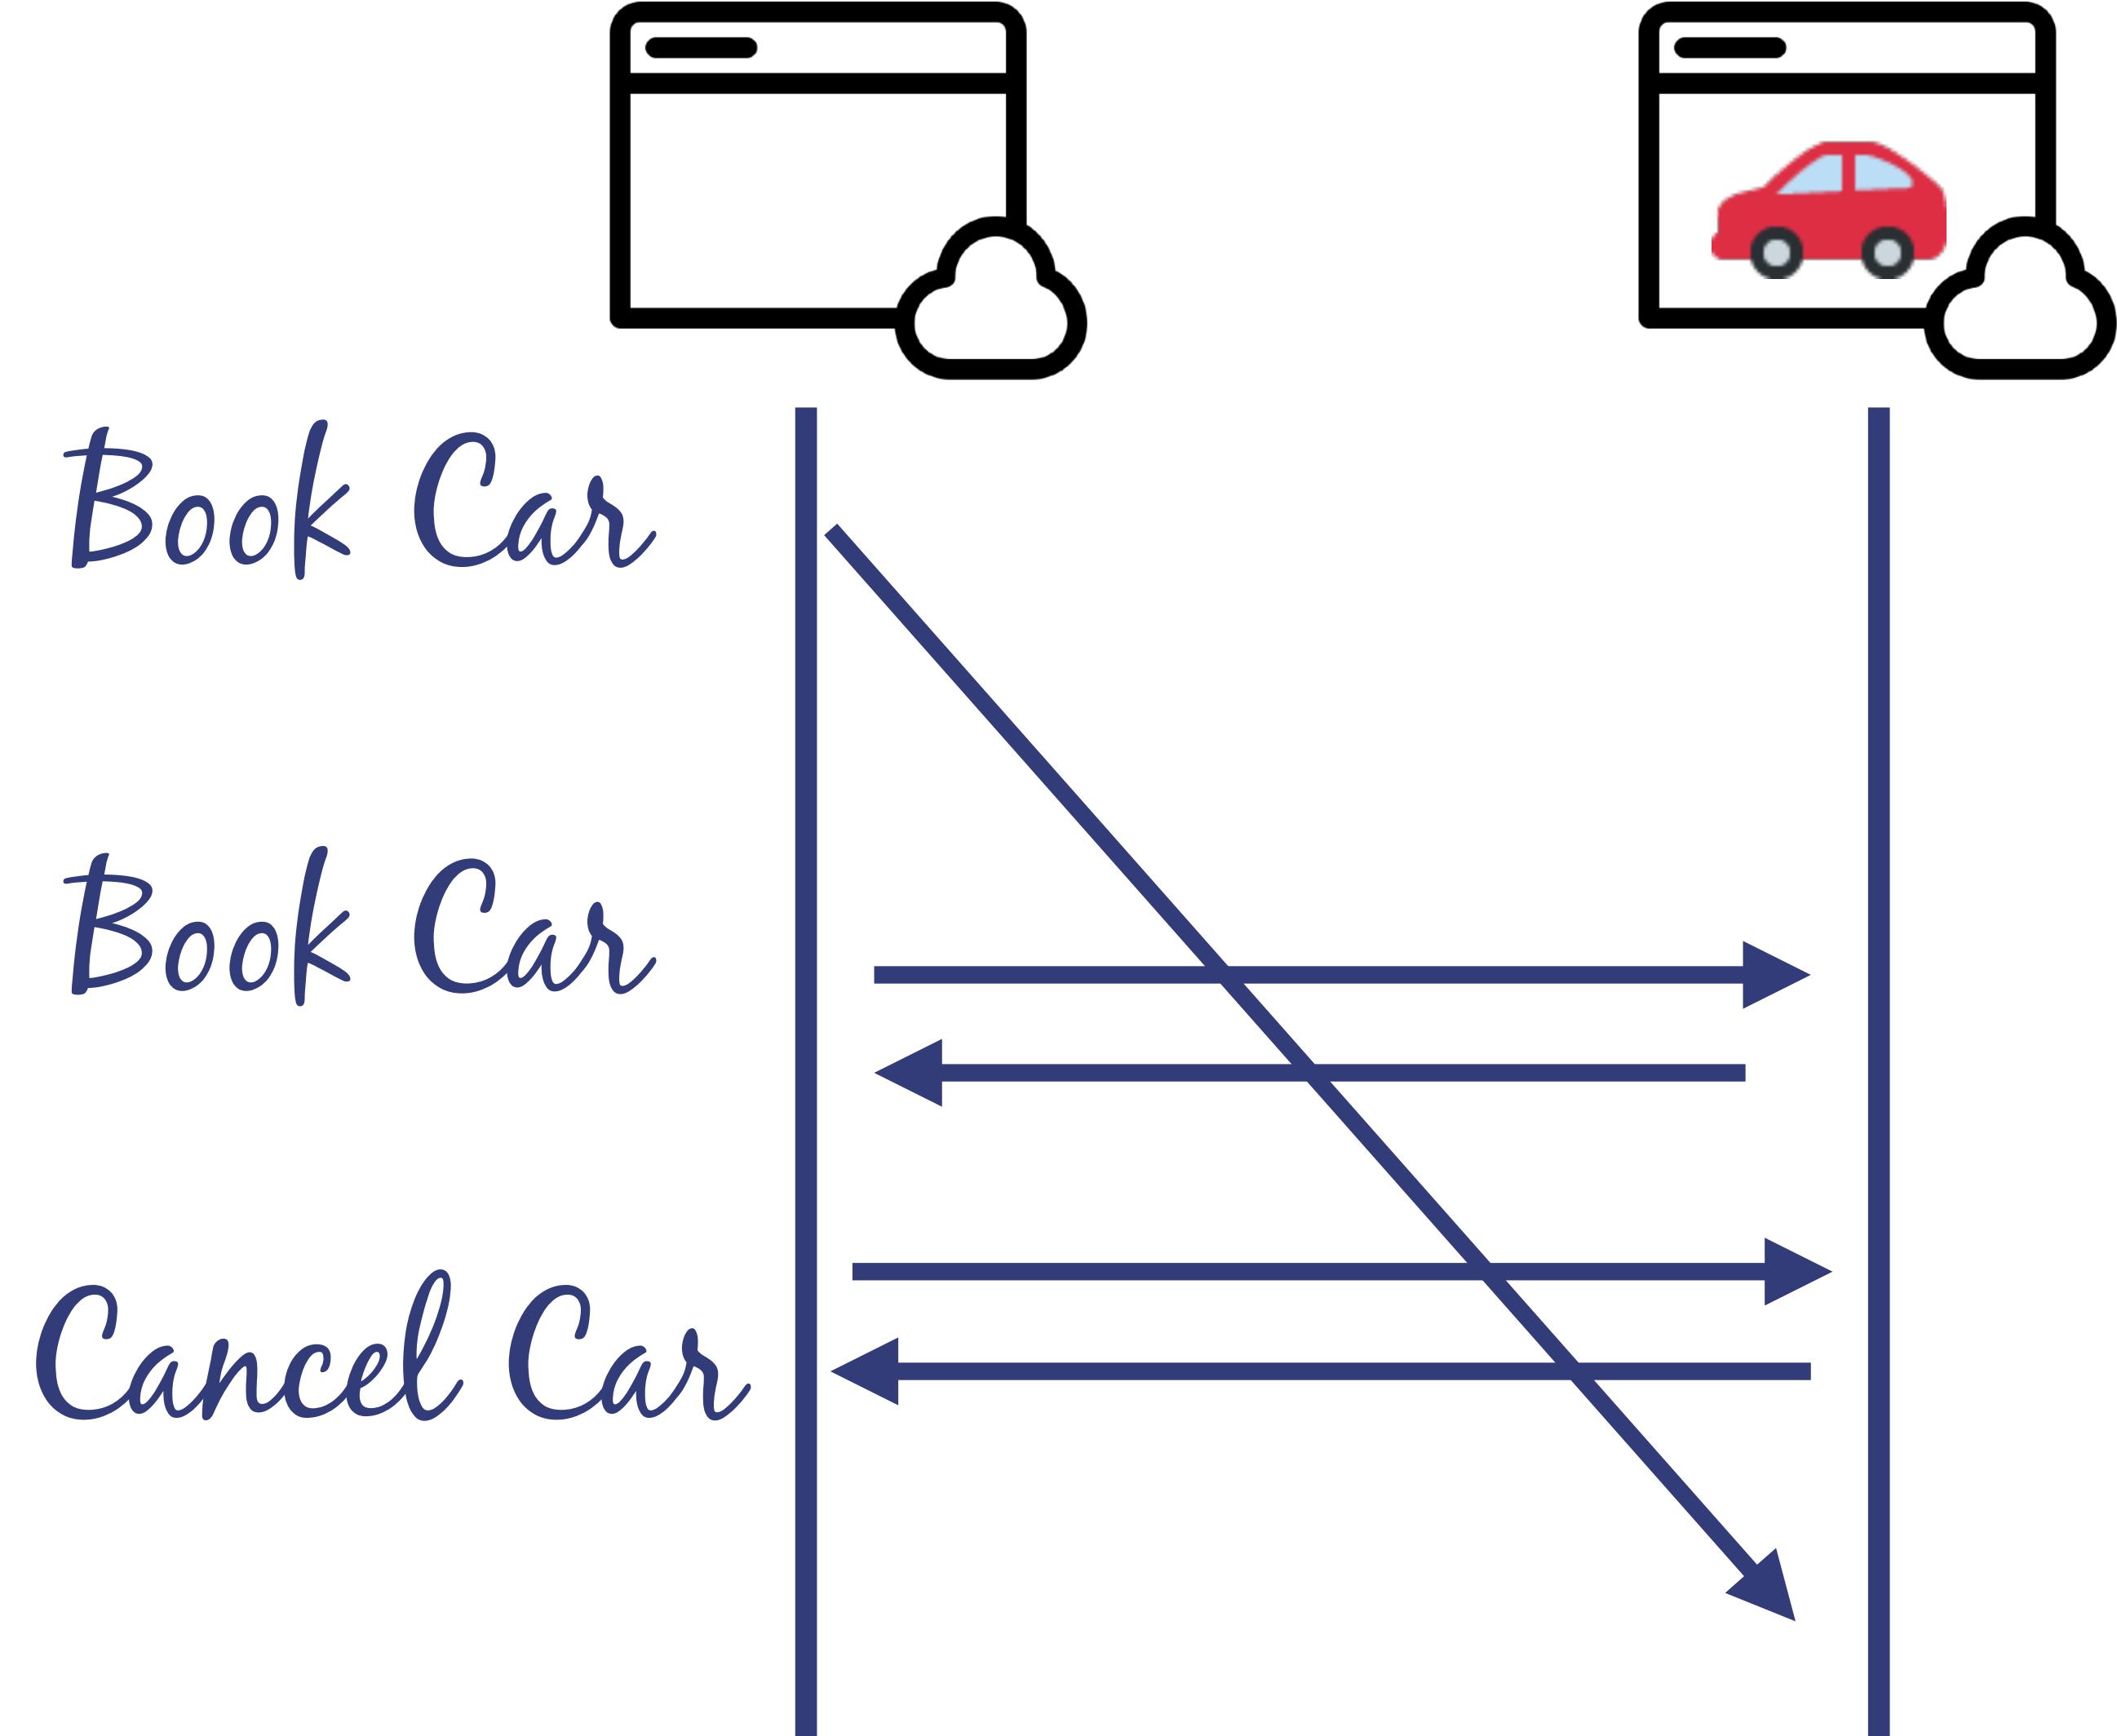
Compensating Requests

Must Be Commutative with Requests



Compensating Requests

Must Be Commutative with Requests



Requests

Idempotent

Can Abort

Compensating Requests

Idempotent

Commutative

Can Not Abort

Distributed Saga Guarantee

All requests were completed successfully



Book Hotel



Book Car



Book Flight



Charge Money

Distributed Saga Guarantee

All requests were completed successfully



Book Hotel



Book Car



Book Flight



Charge Money

Or a subset of requests and the corresponding compensating requests were executed



Book Hotel



Book Car



Cancel Hotel



Cancel Car

Distributed Saga Guarantee

No Atomicity

No Isolation

Distributed Saga Guarantee

No Atomicity

No Isolation



Visible before
Saga Completes



Book Hotel



Book Car

Distributed Saga Guarantee

All requests were completed successfully



Book Hotel



Book Car



Book Flight



Charge Money

Or a subset of requests and the corresponding compensating requests were executed



Book Hotel



Book Car



Cancel Hotel

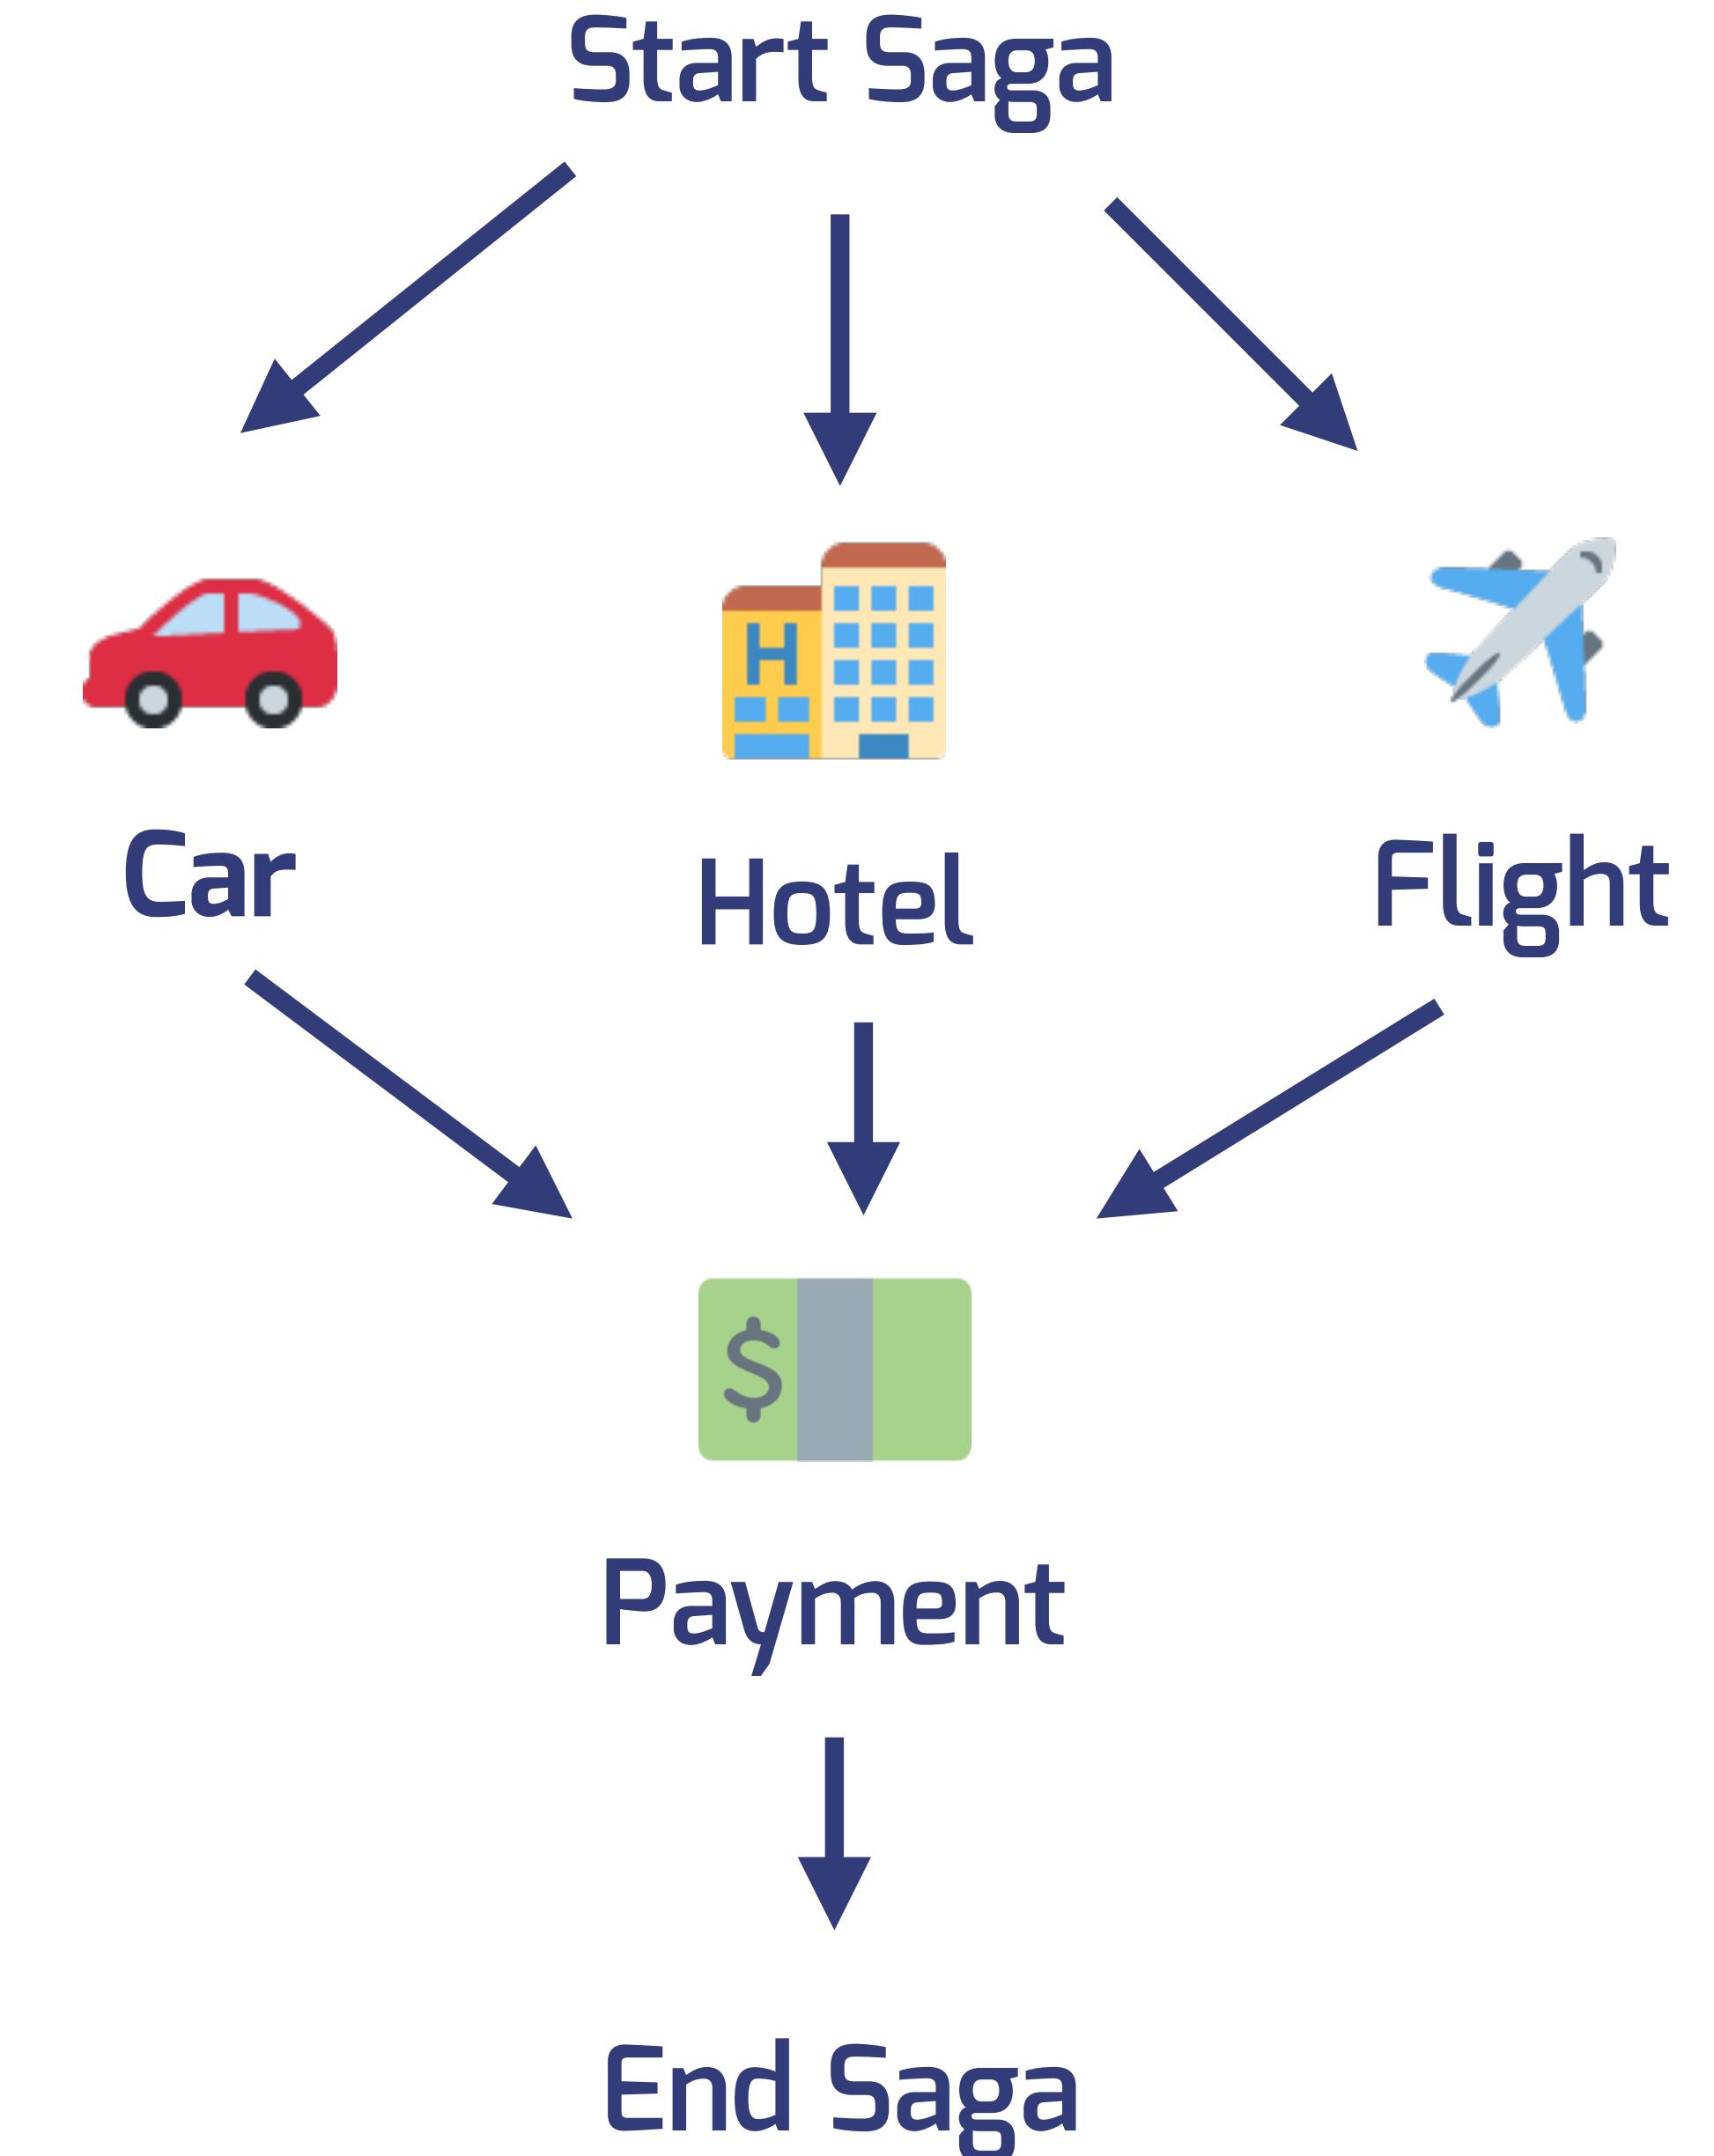


Cancel Car

defining a

Distributed Saga

Distributed Saga Directed Acyclic Graph



Distributed Saga Vertex

Name: Hotel

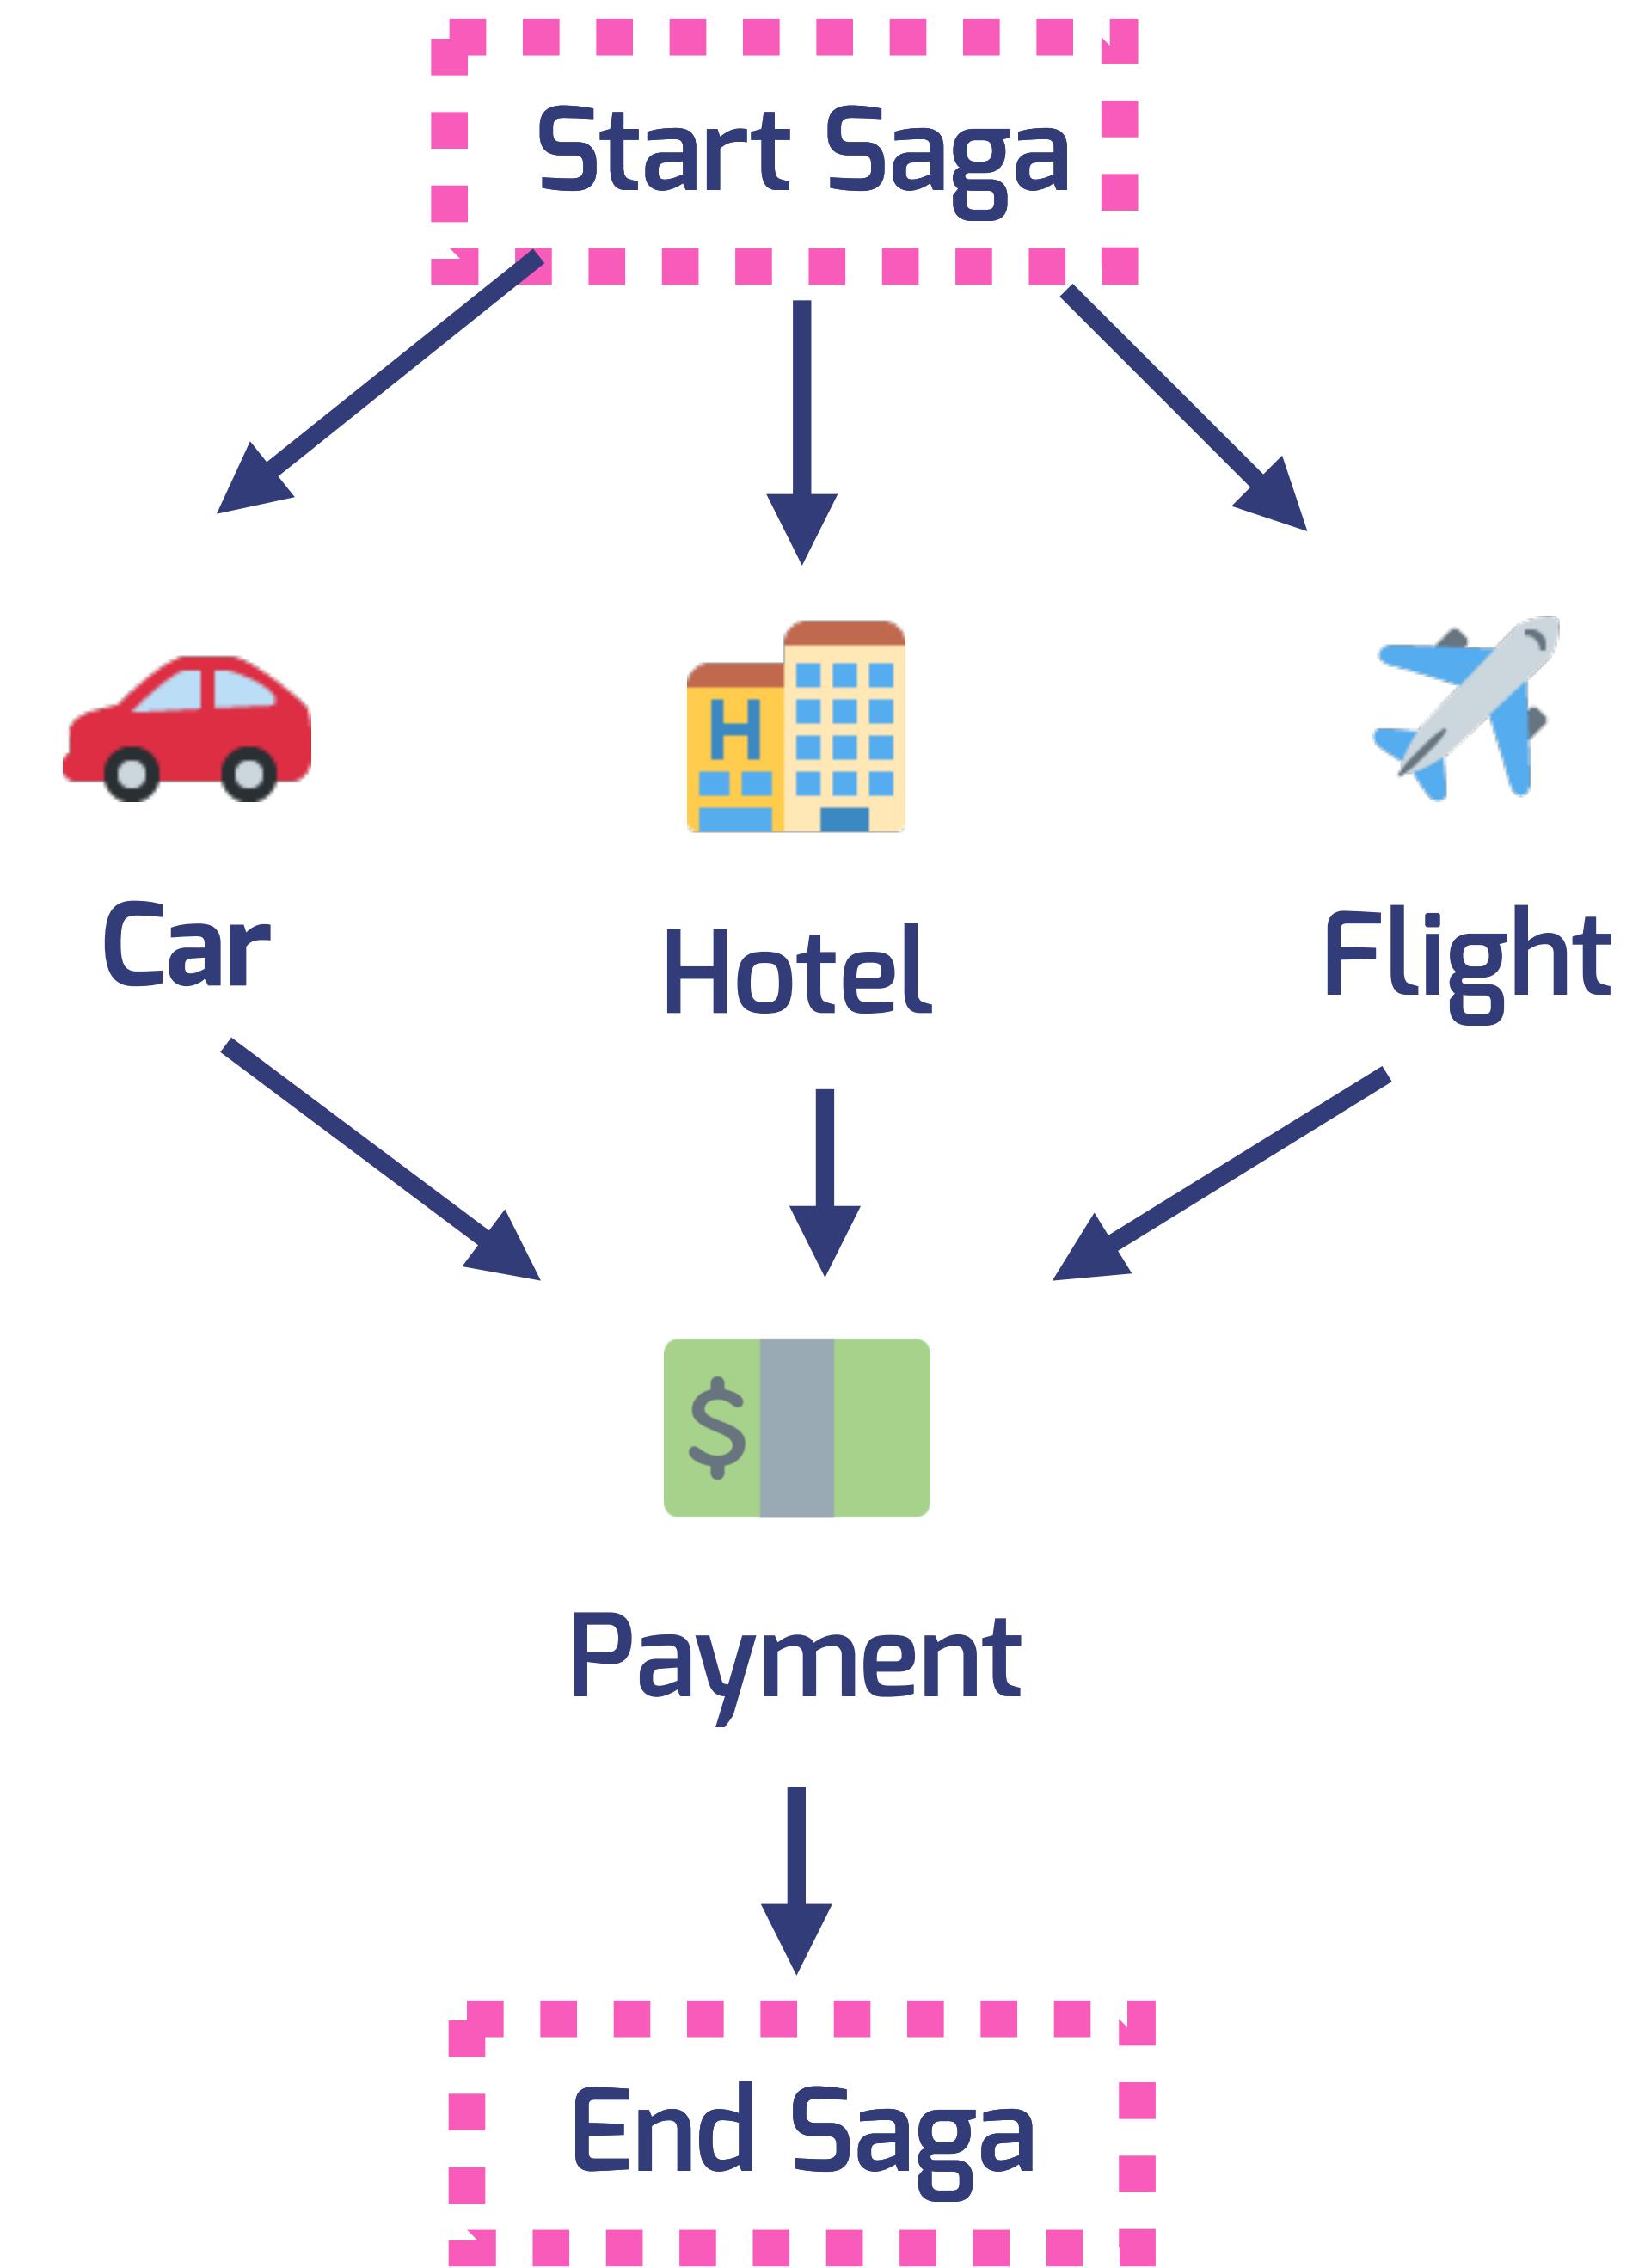


Request: Book Hotel

Compensating Request: Cancel Hotel

Status: Not Completed

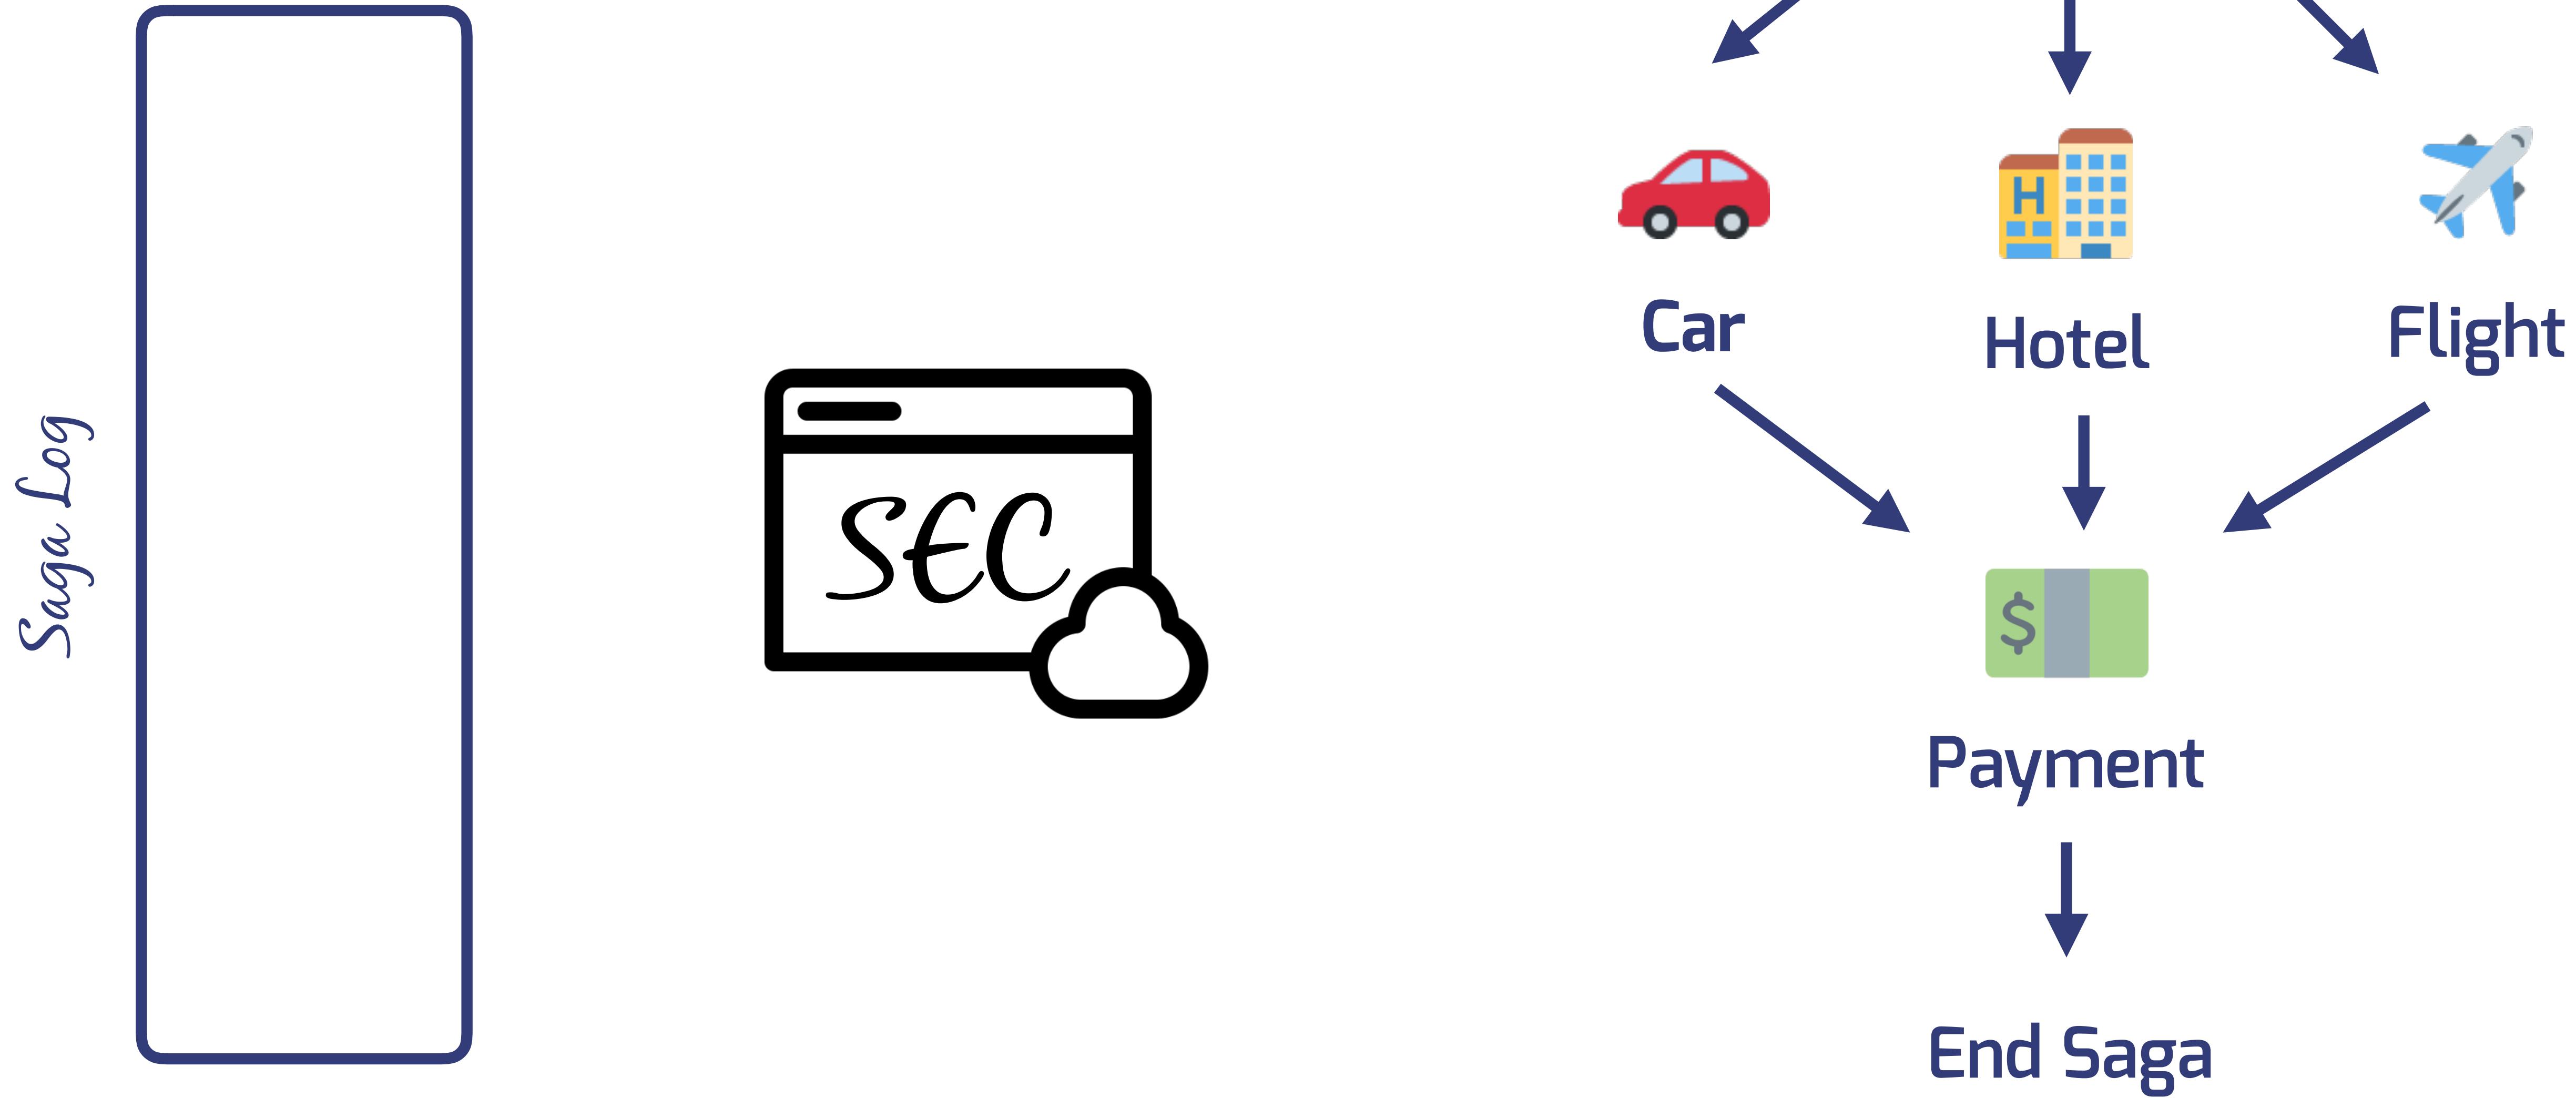
Distributed Saga Start & End Vertices



Distributed Saga Log

fault-tolerant & highly available

Saga Execution Coordinator

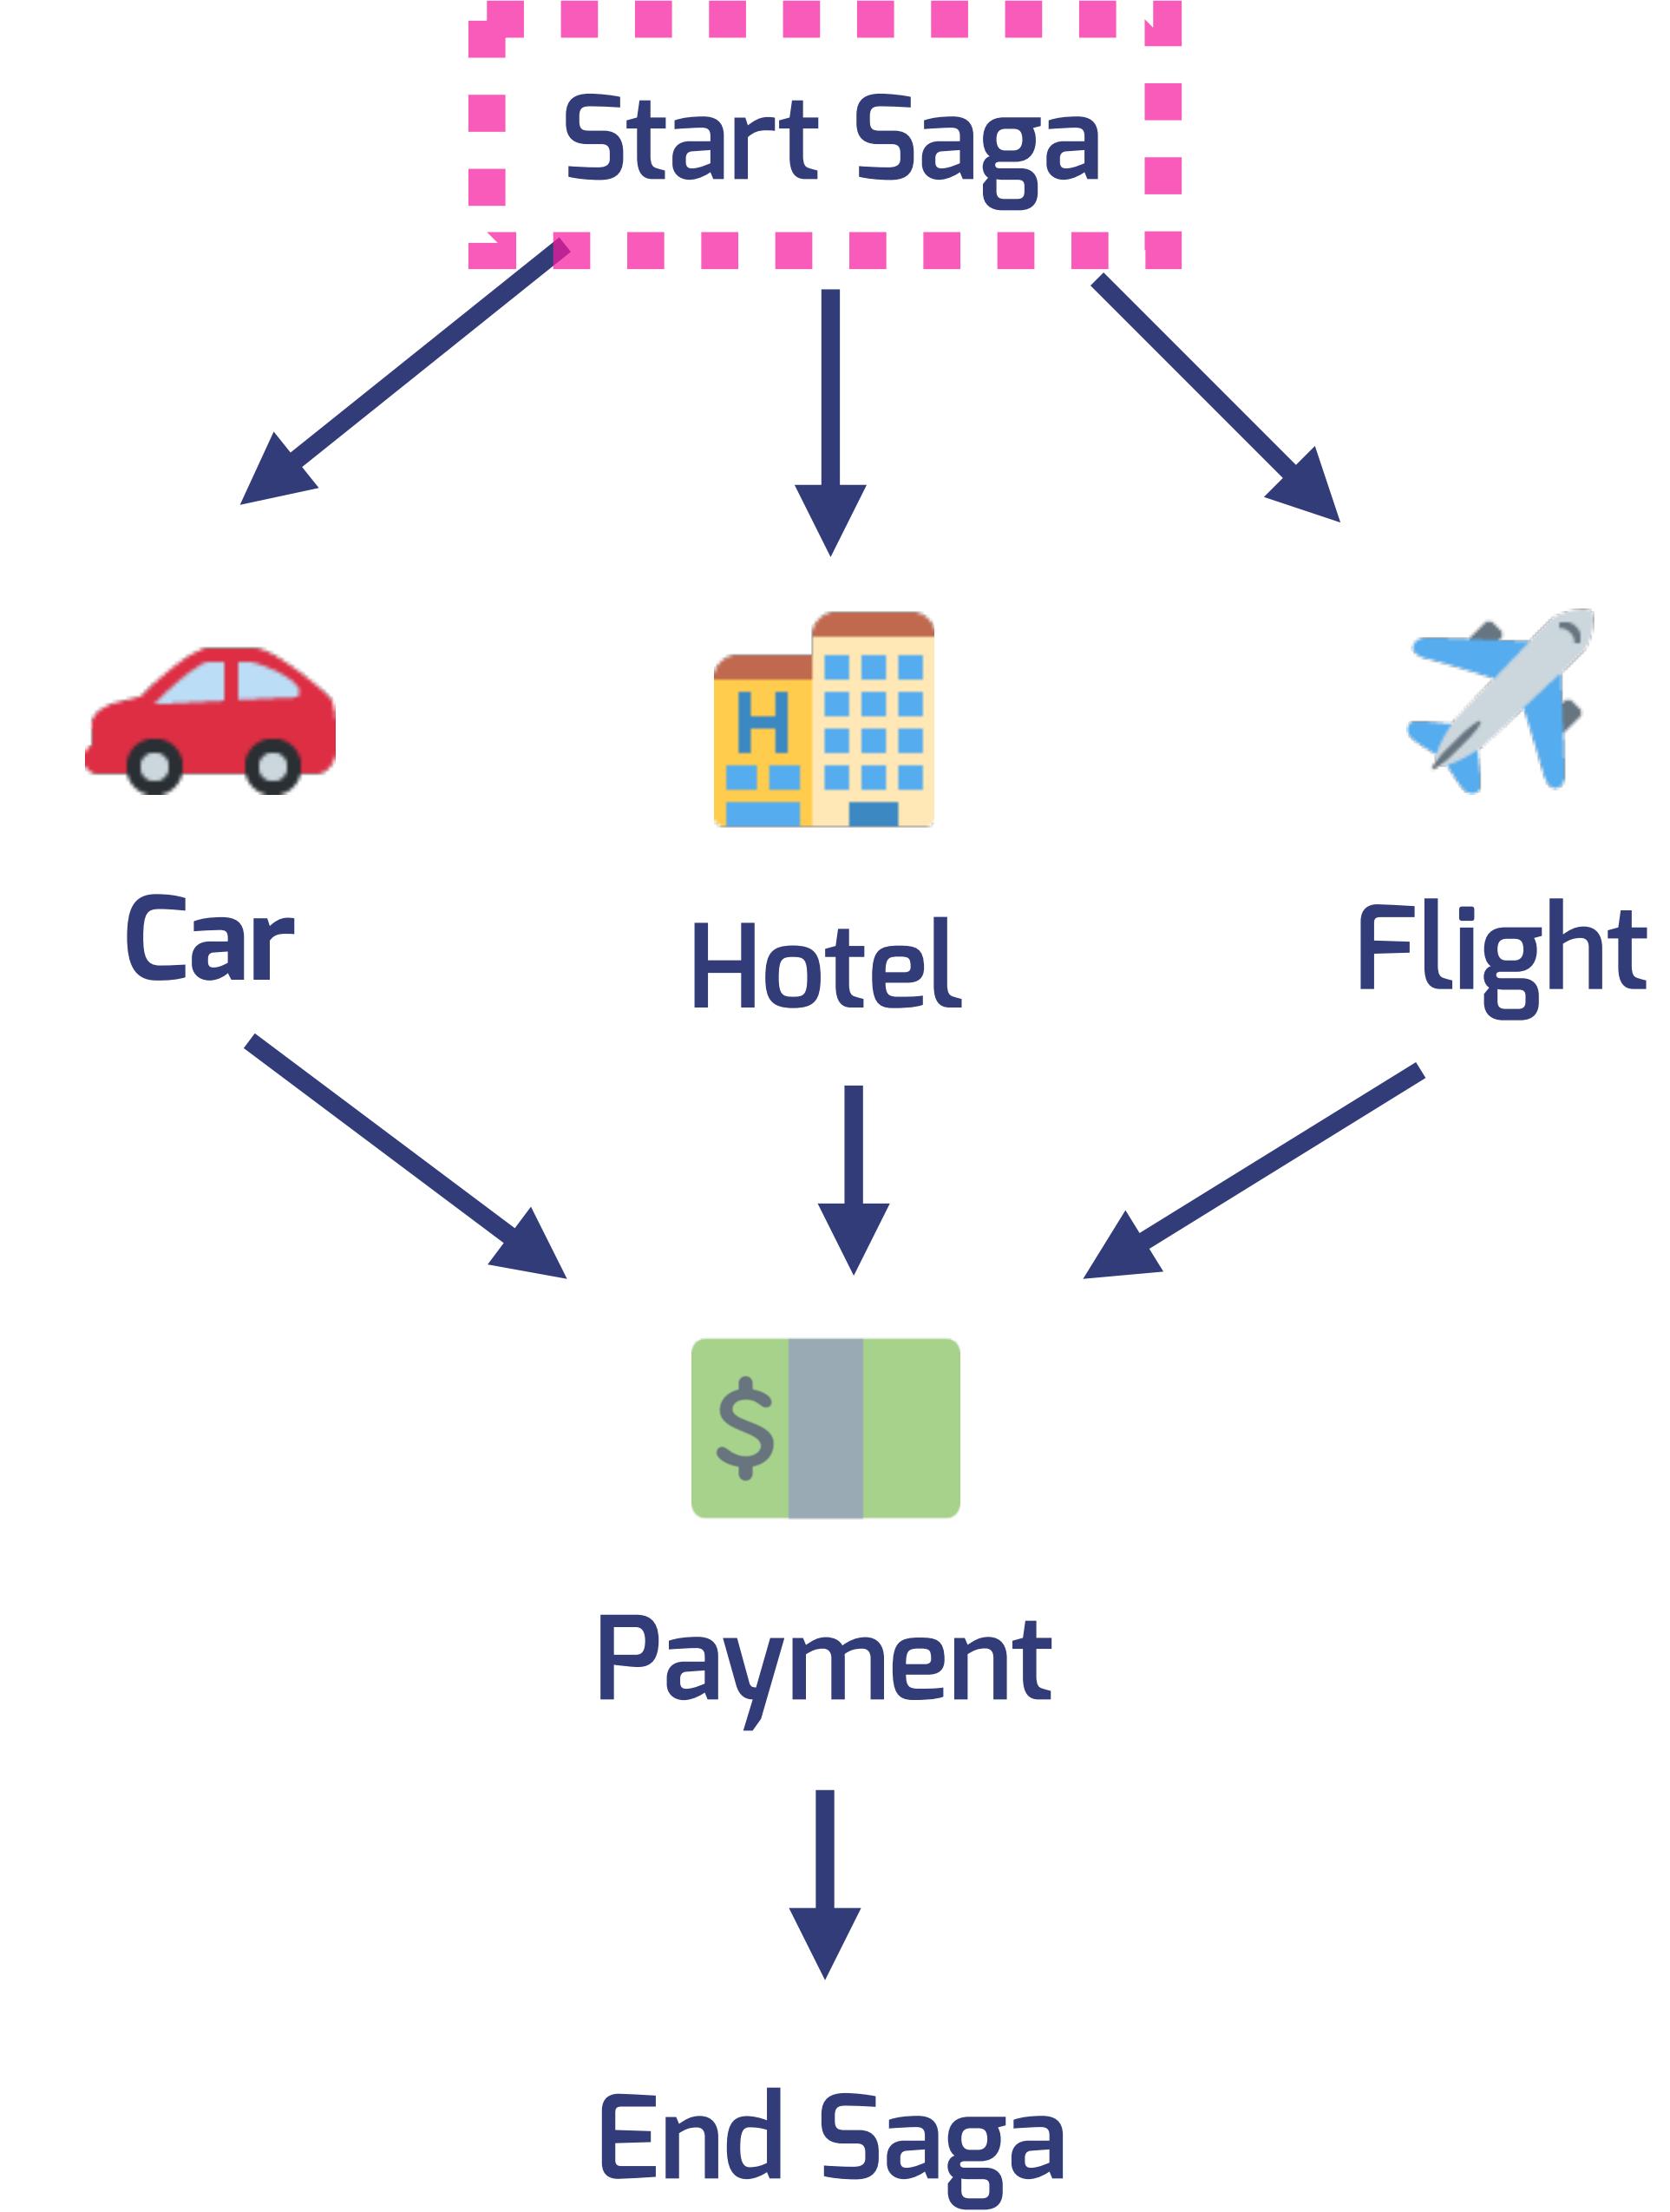


Executing a
Distributed Saga

Saga Log

Book Trip Request

```
{  
  "Name": "Caitie McCaffrey",  
  "Destination": "Malaga, Spain",  
  "Start Date": "2017-05-17",  
  "End Date": "2017-05-20",  
  "Payment Token":  
    "Tm90IG15IHJ1YWwgY3JlZG10IGNhcmQgaW5mbvA6KQ==",  
  "Price": "2500USD"  
}
```



Saga Log

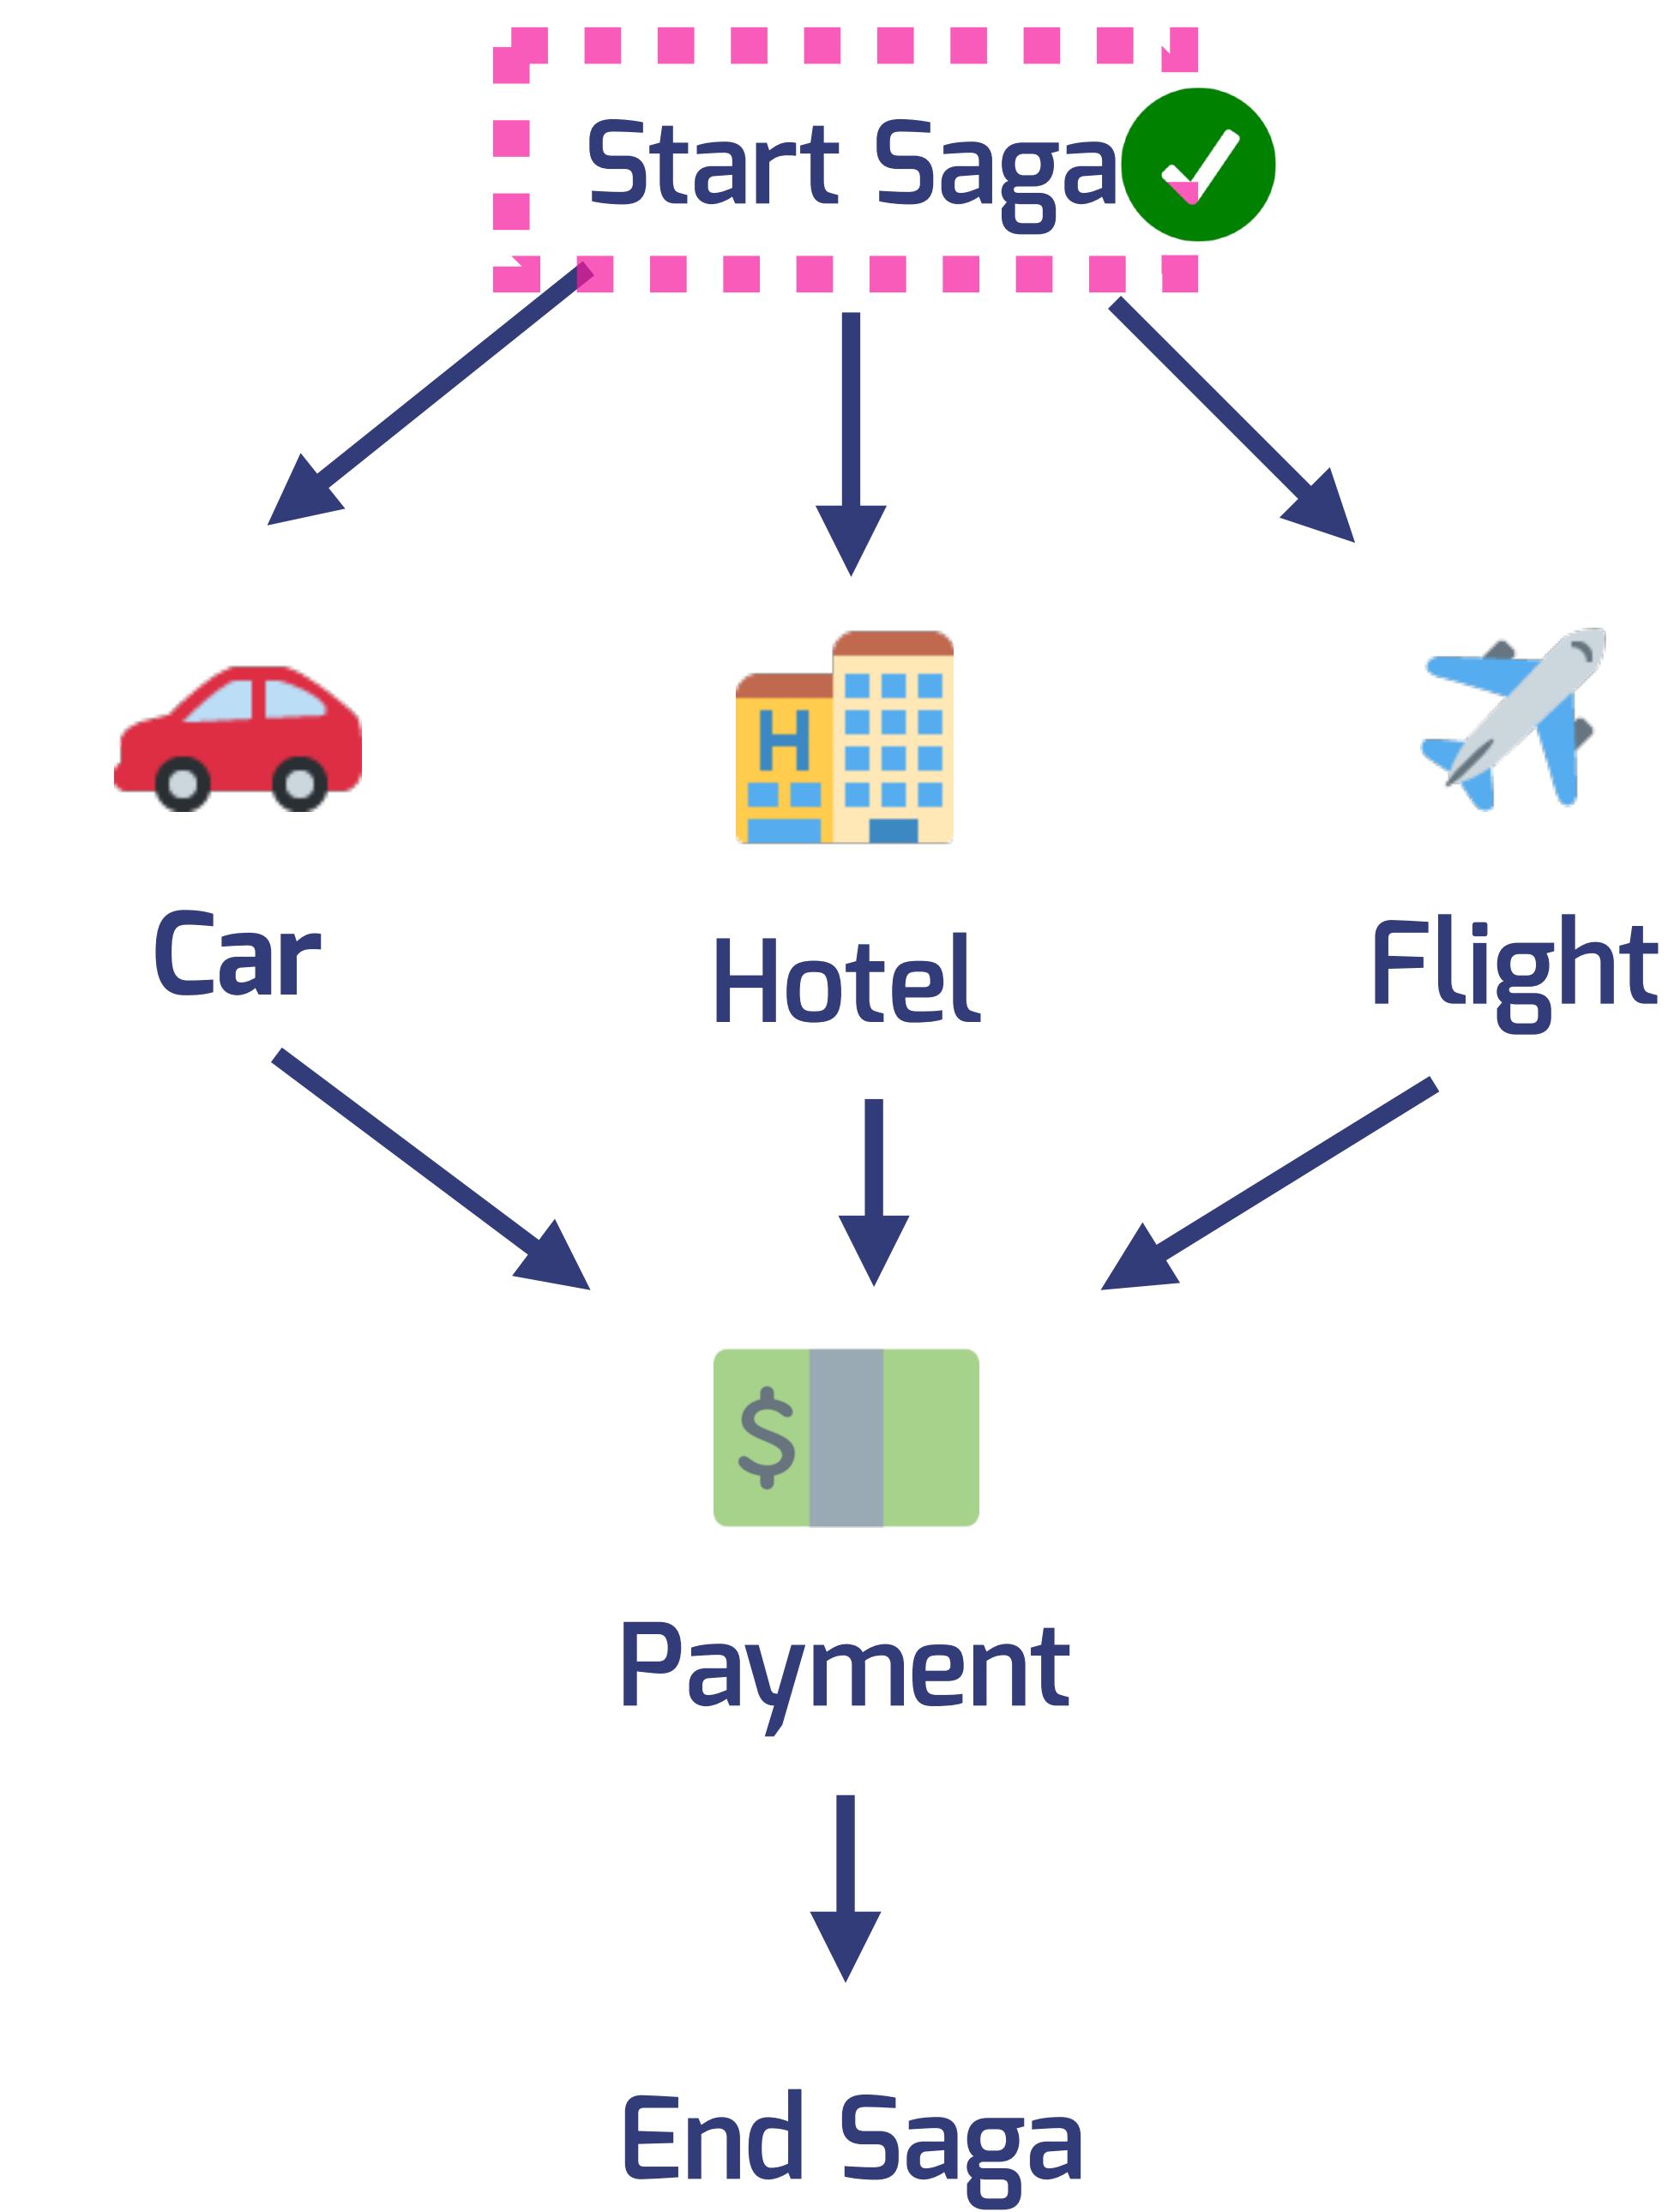
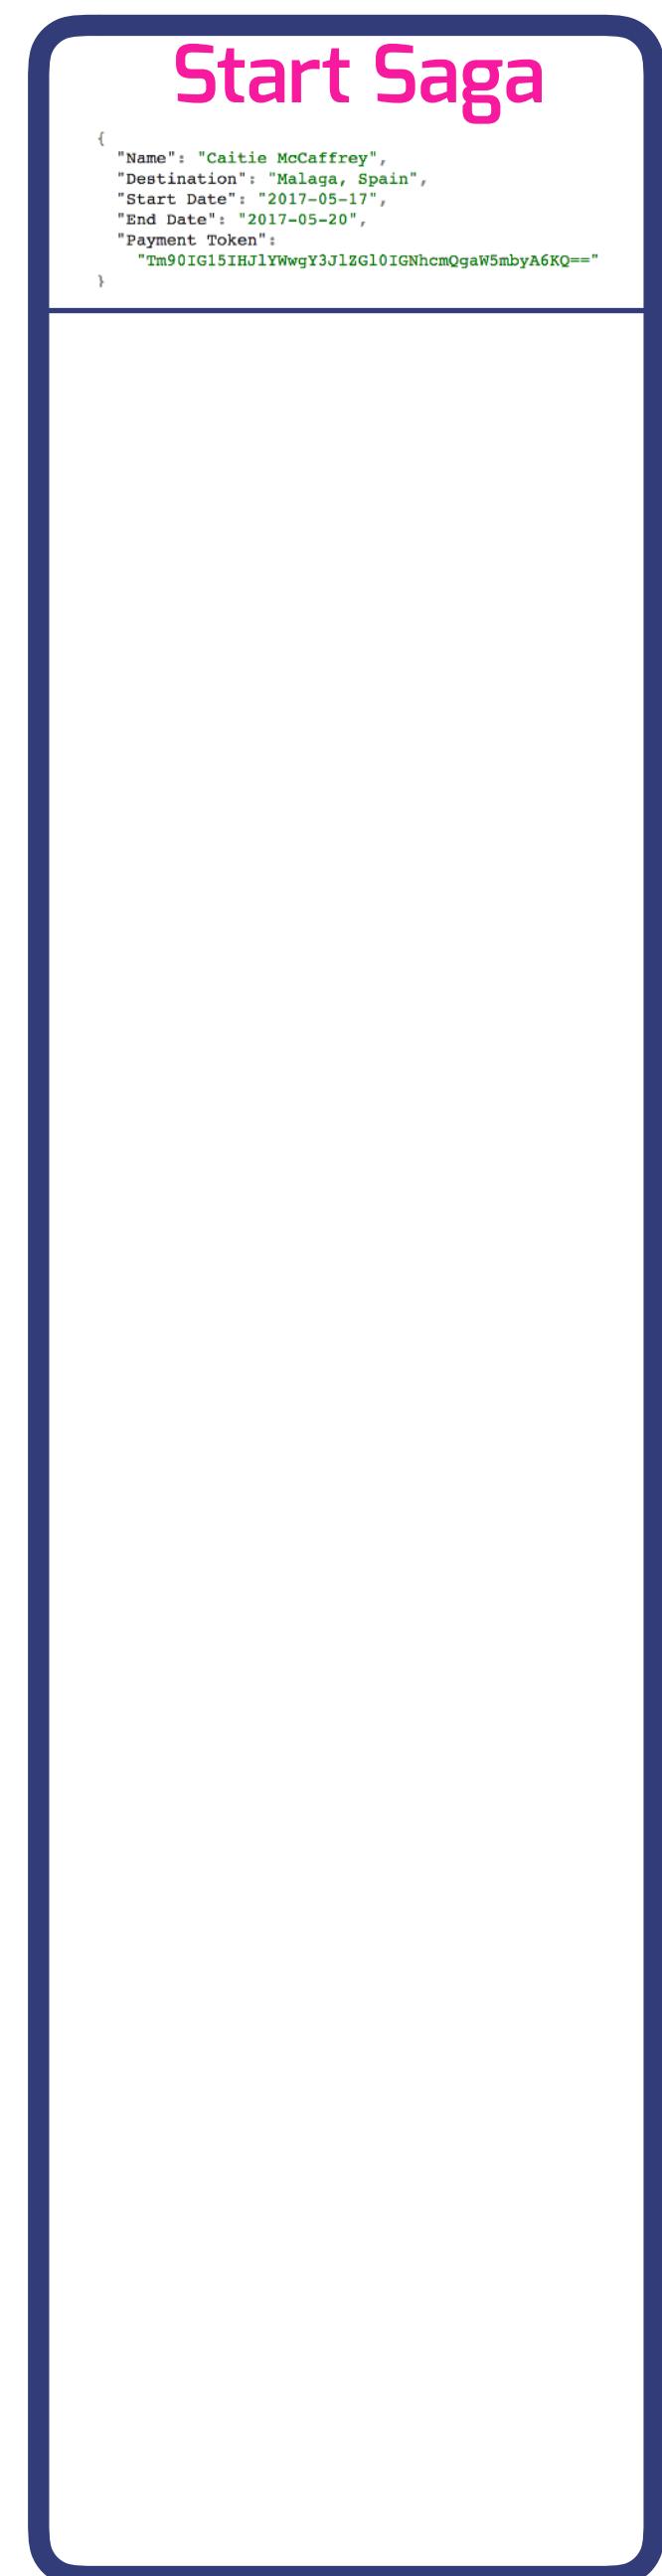
Book Trip Request

```
{  
  "Name": "Caitie McCaffrey",  
  "Destination": "Malaga, Spain",  
  "Start Date": "2017-05-17",  
  "End Date": "2017-05-20",  
  "Payment Token":  
    "Tm90IG15IHJ1YWwgY3JlZG10IGNhcmQgaW5mbvA6KQ==",  
  "Price": "2500USD"  
}
```

Start Saga



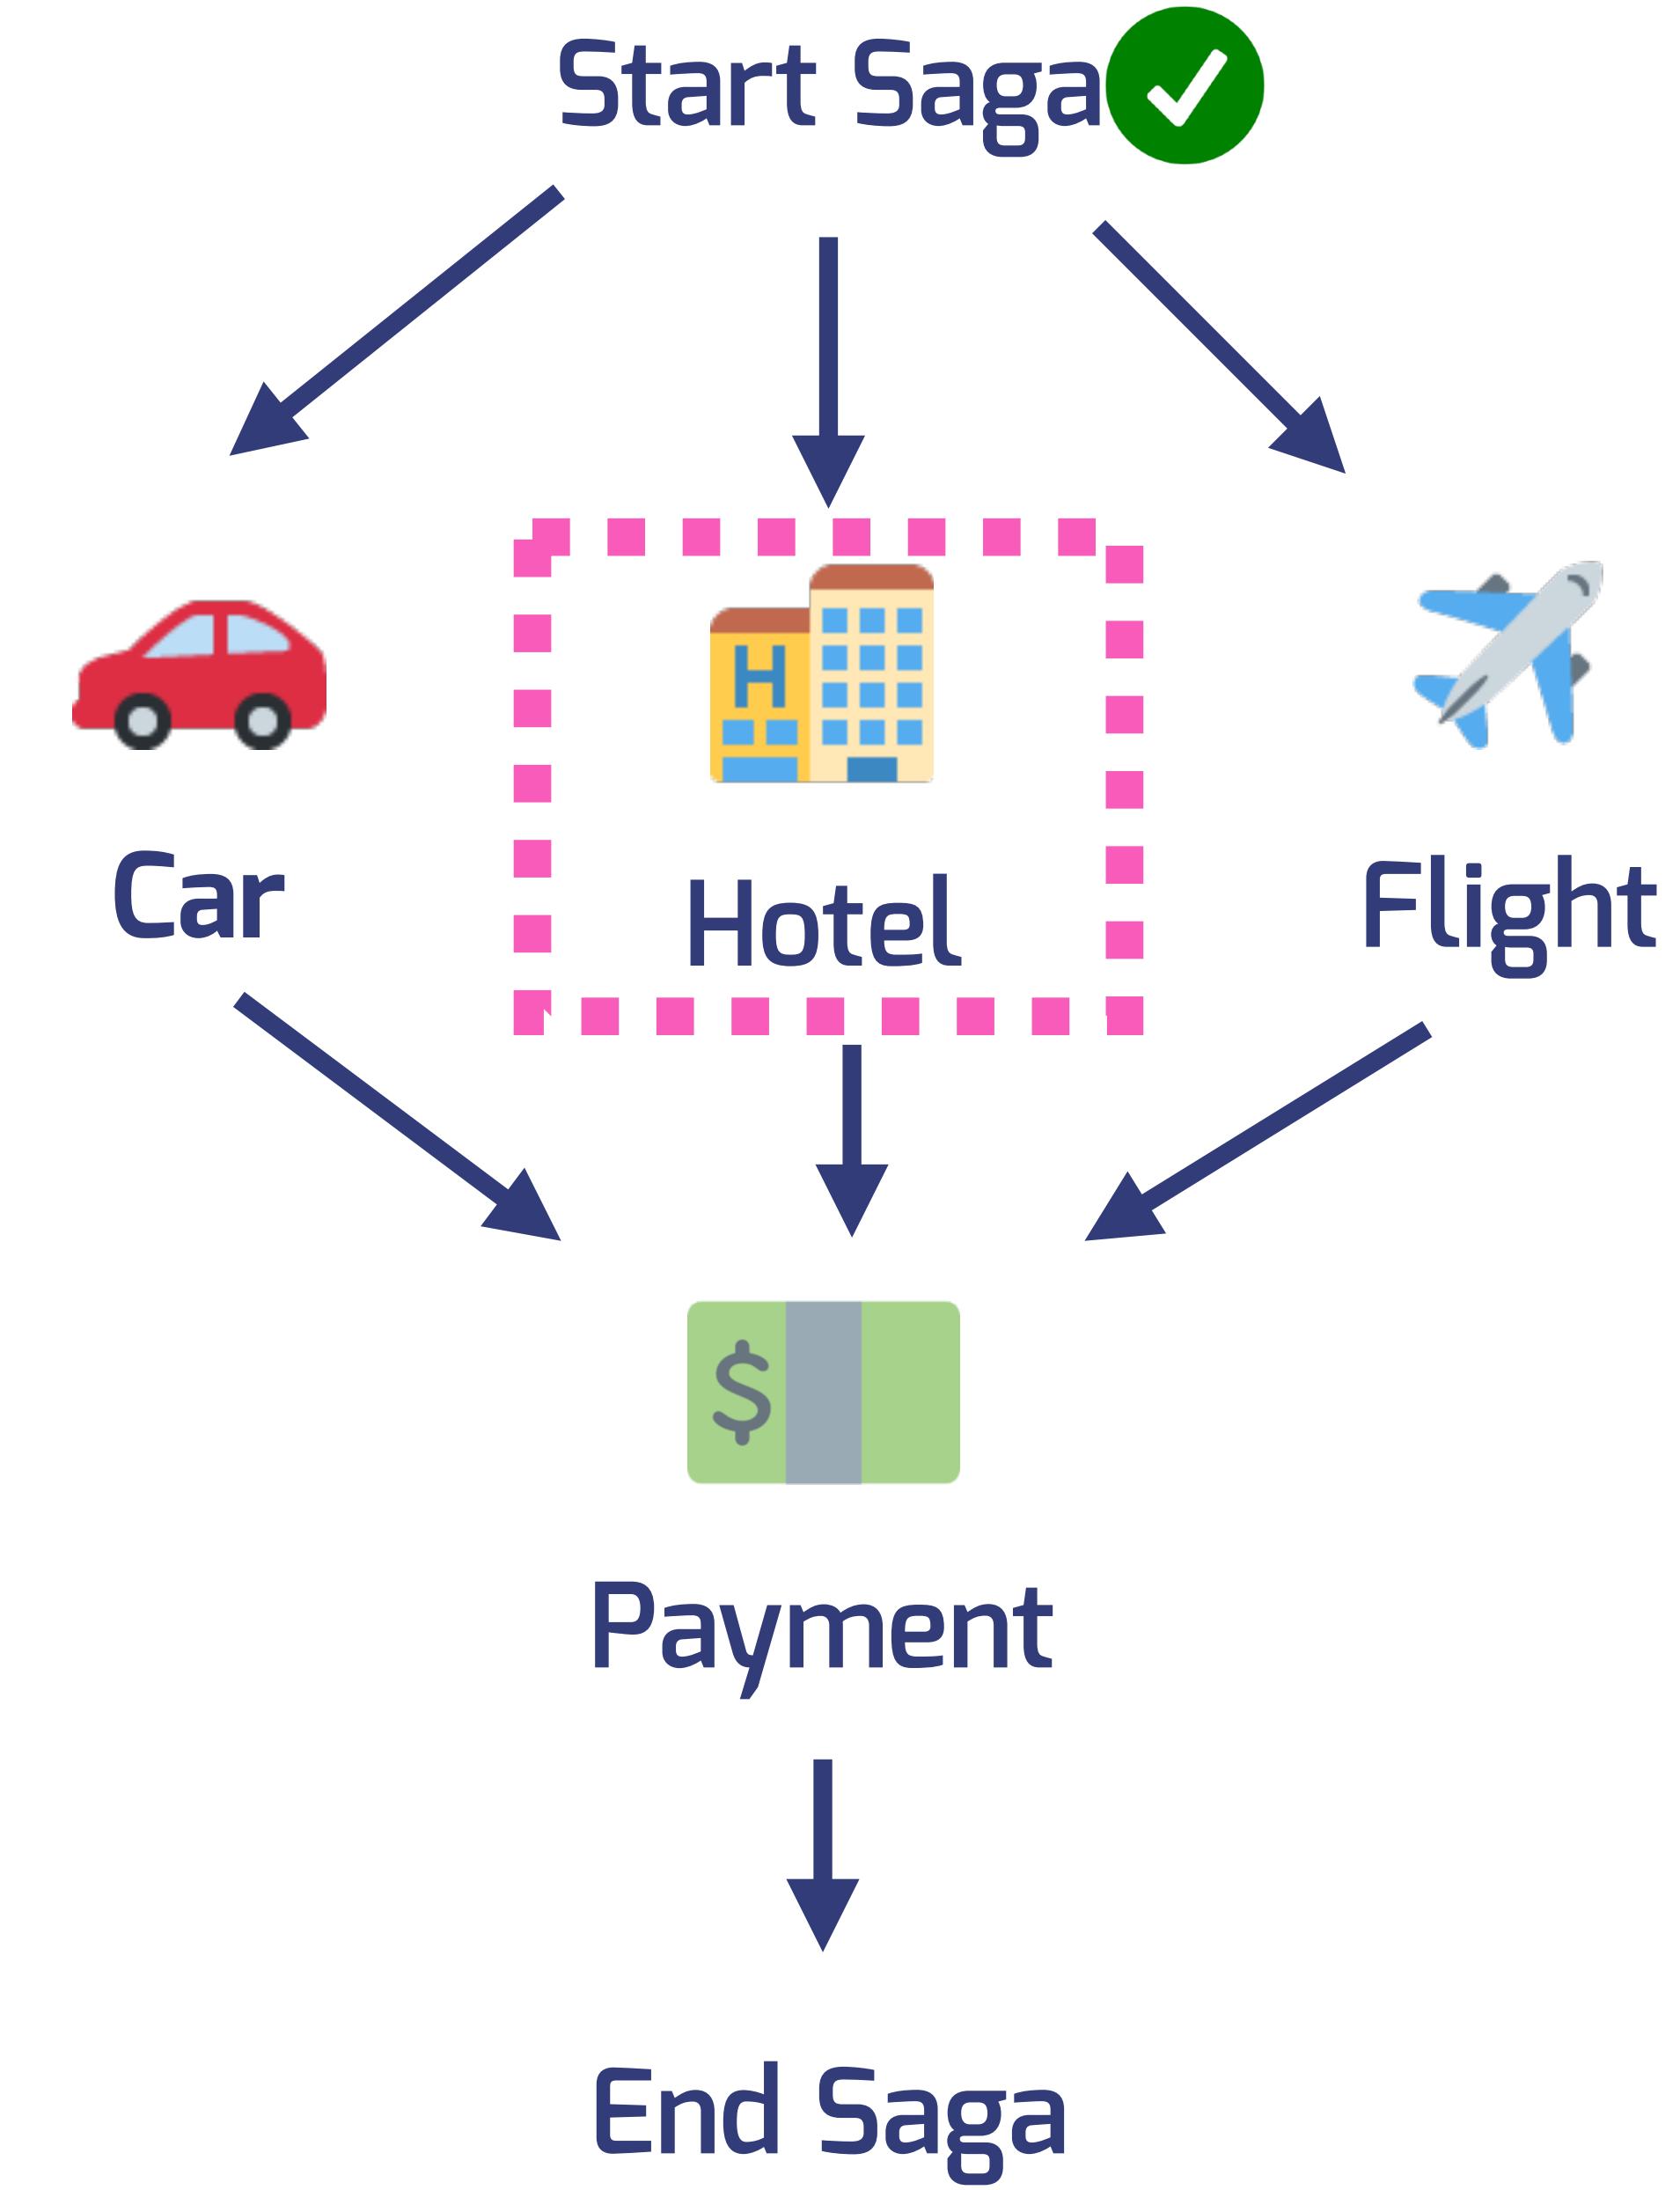
Saga Log

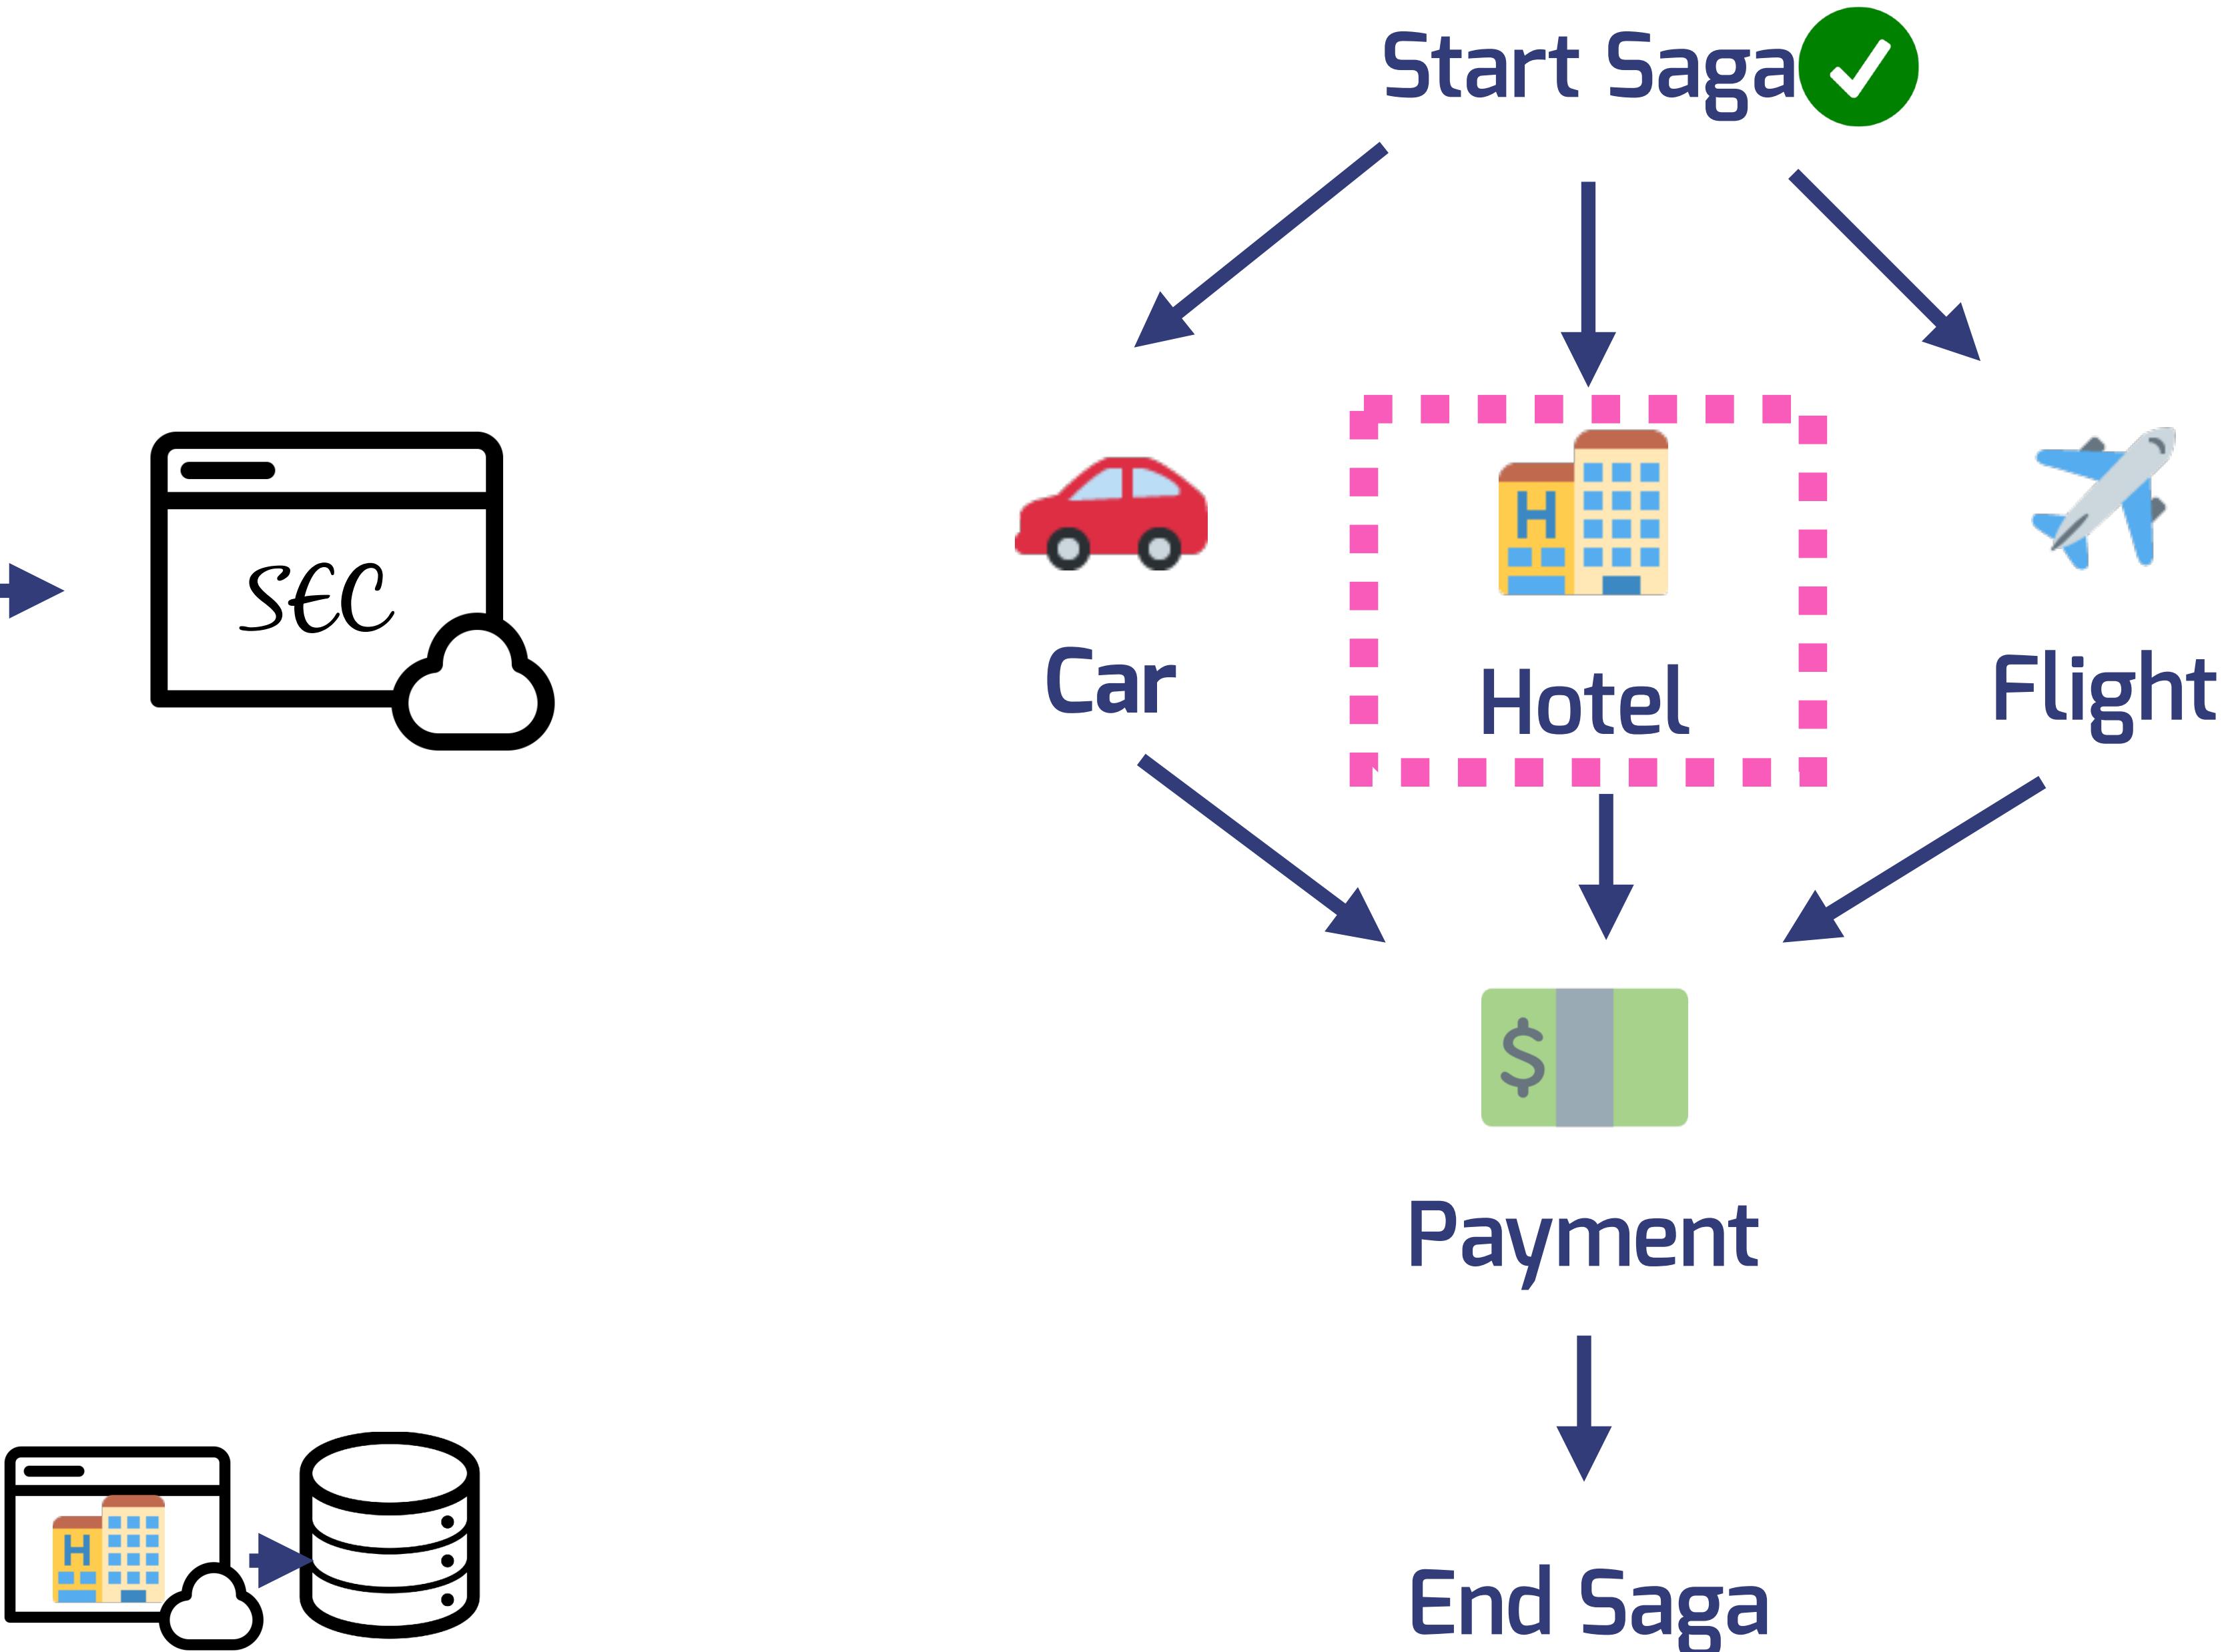
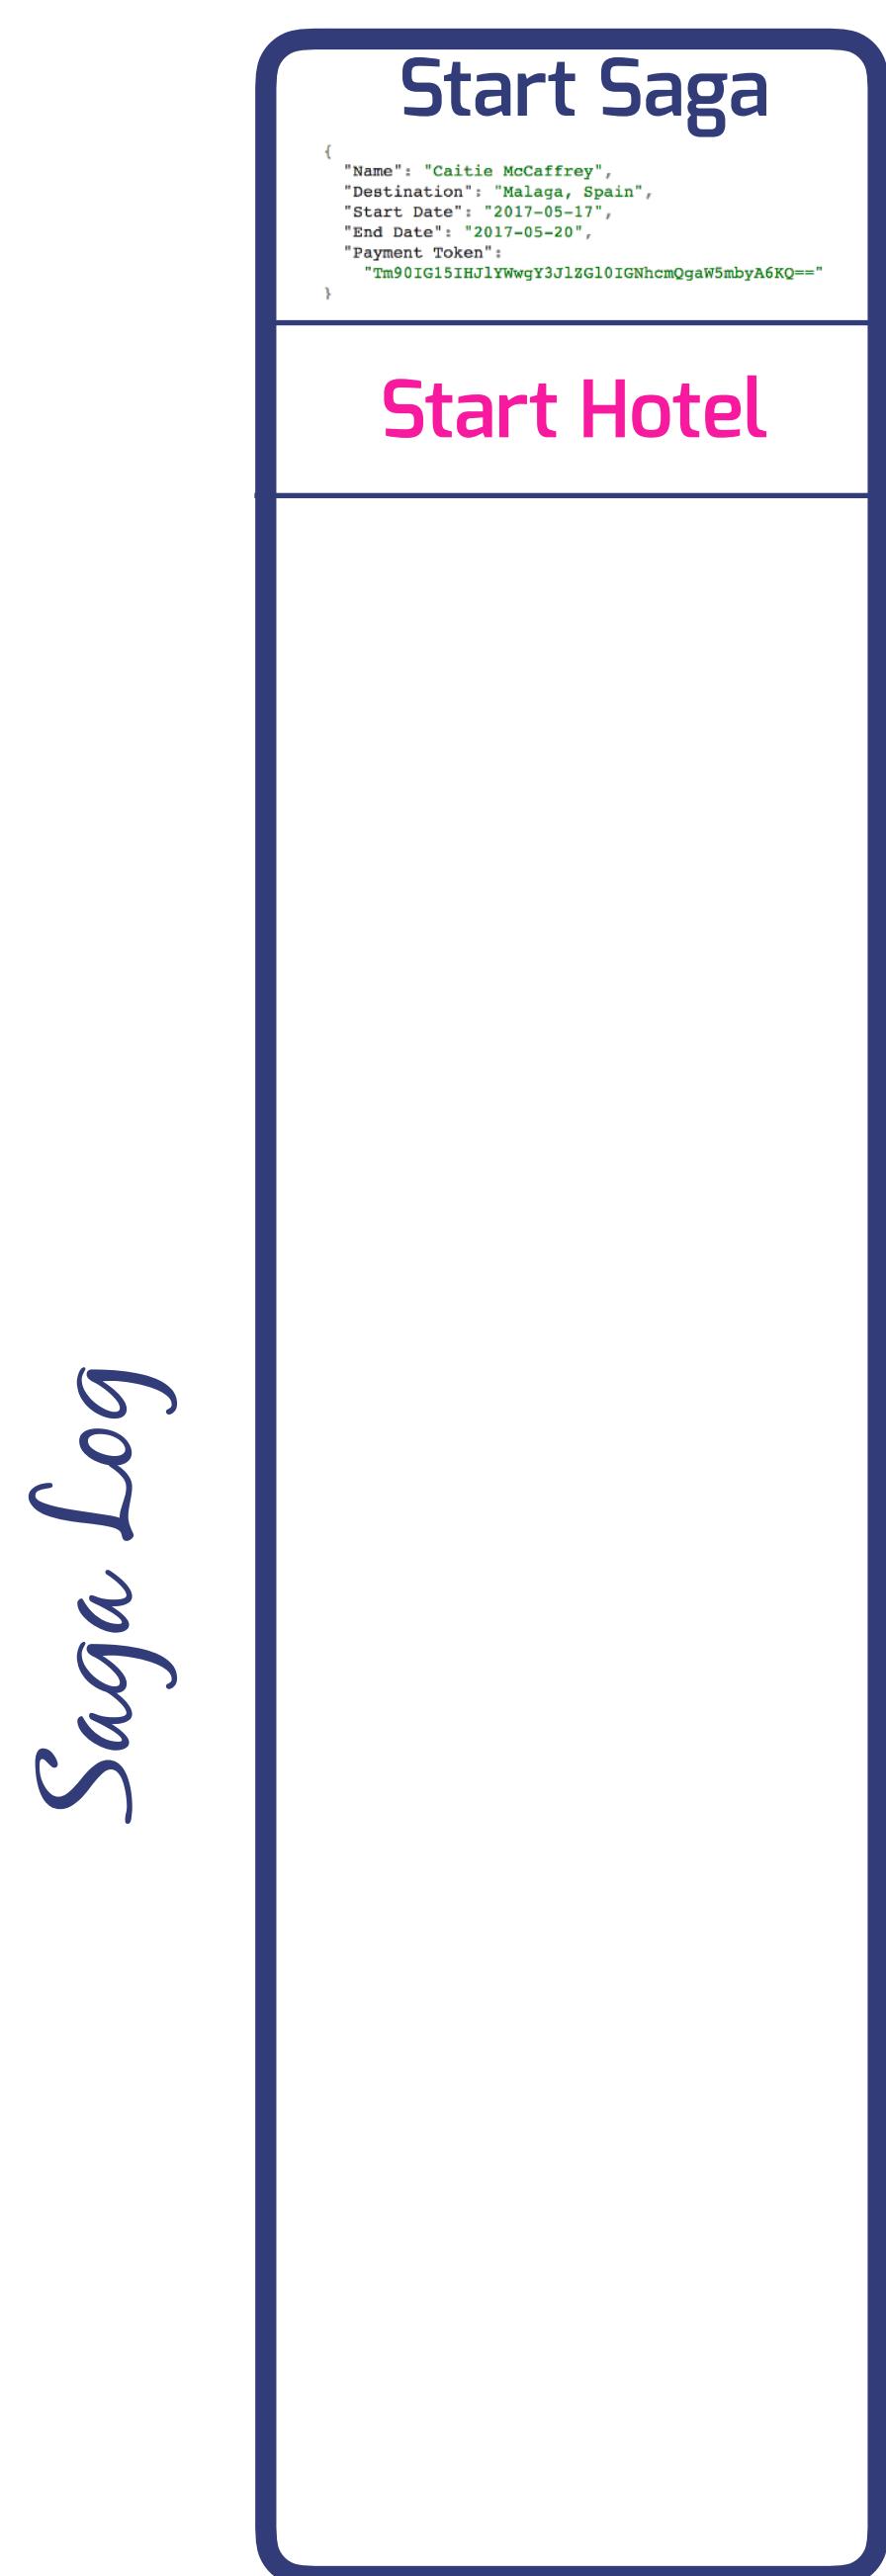


Saga Log

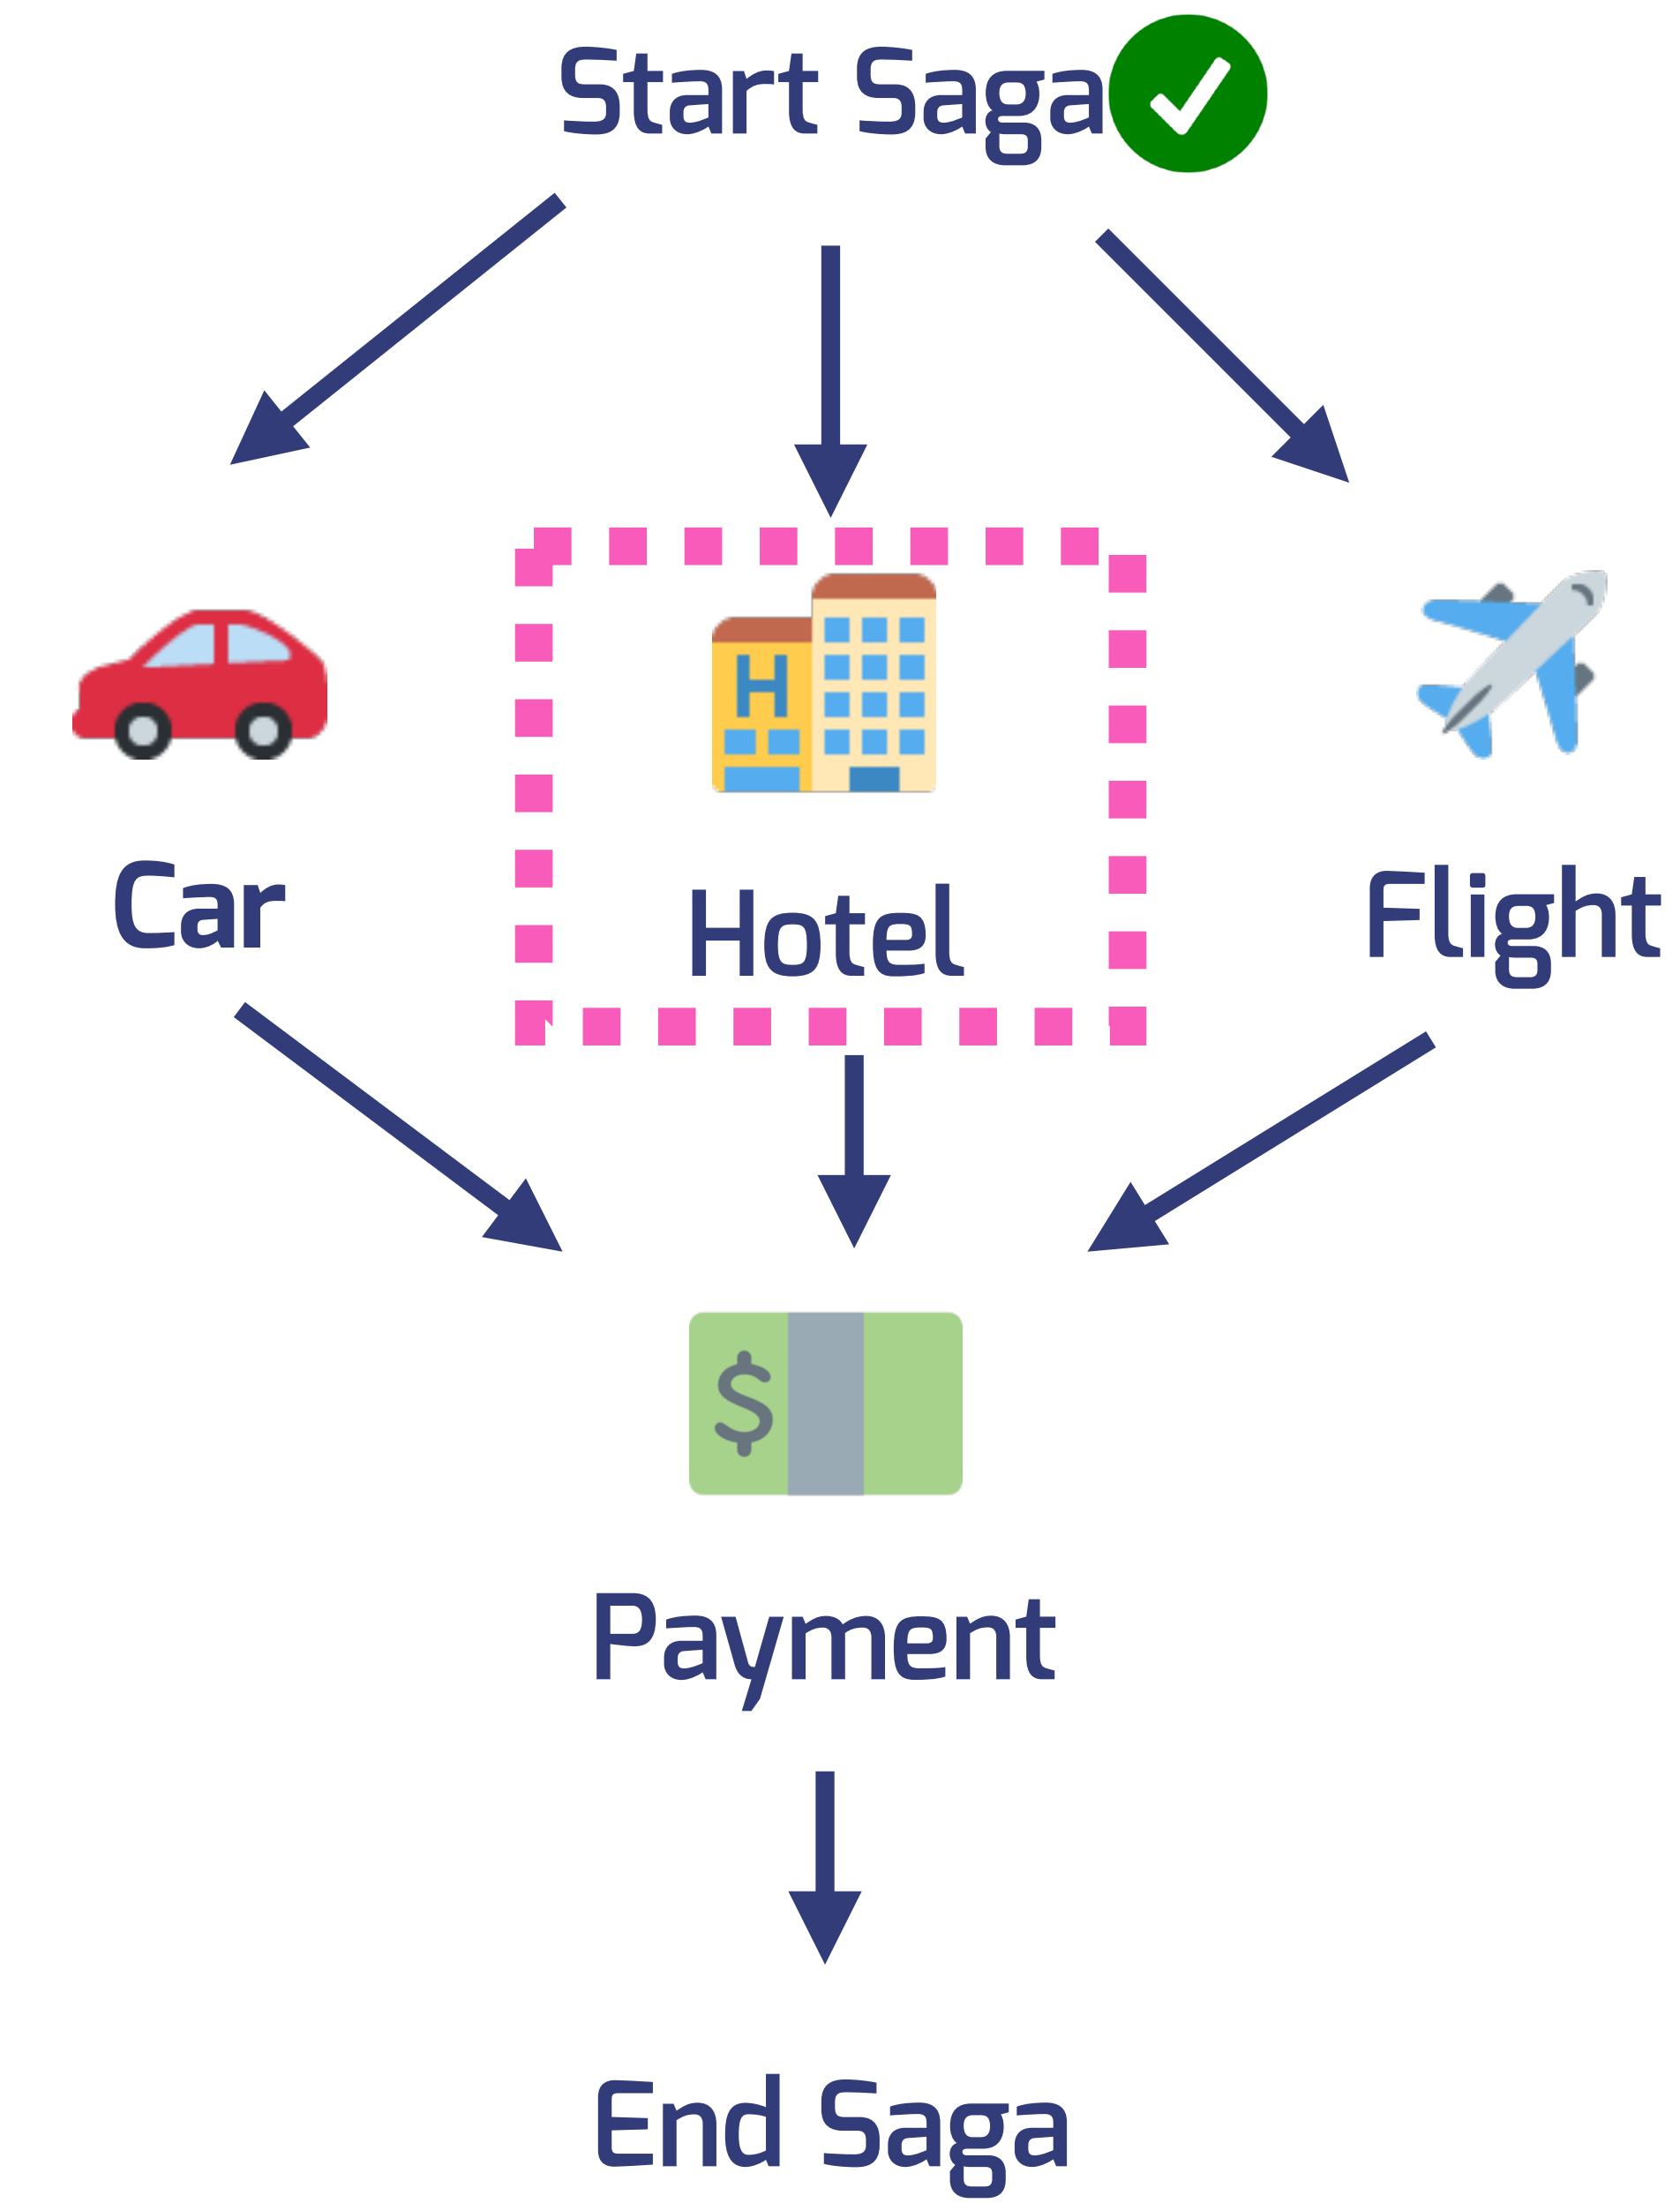
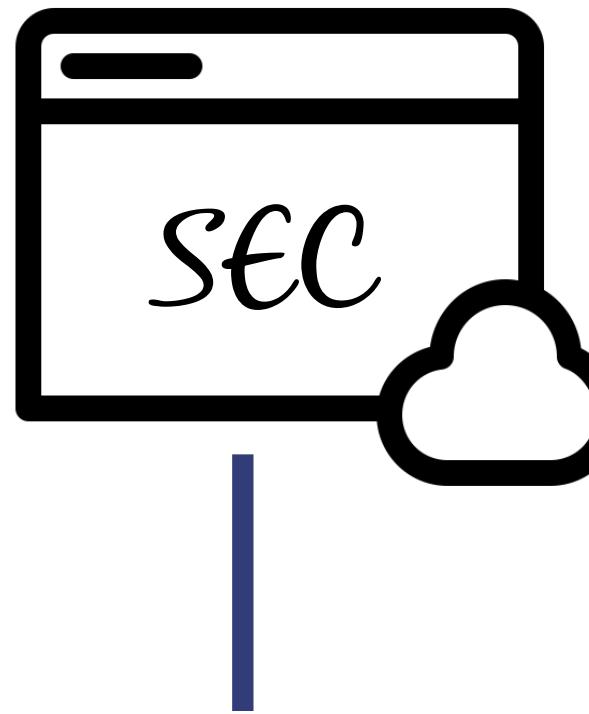
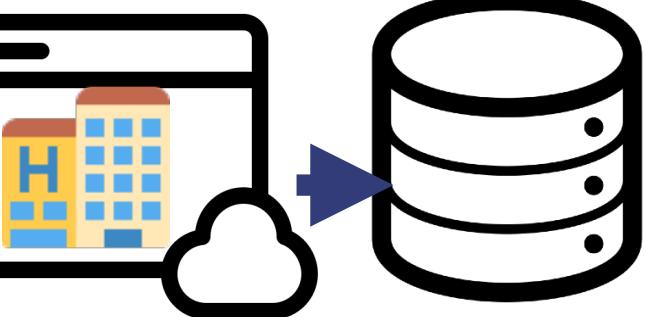


Start Hotel

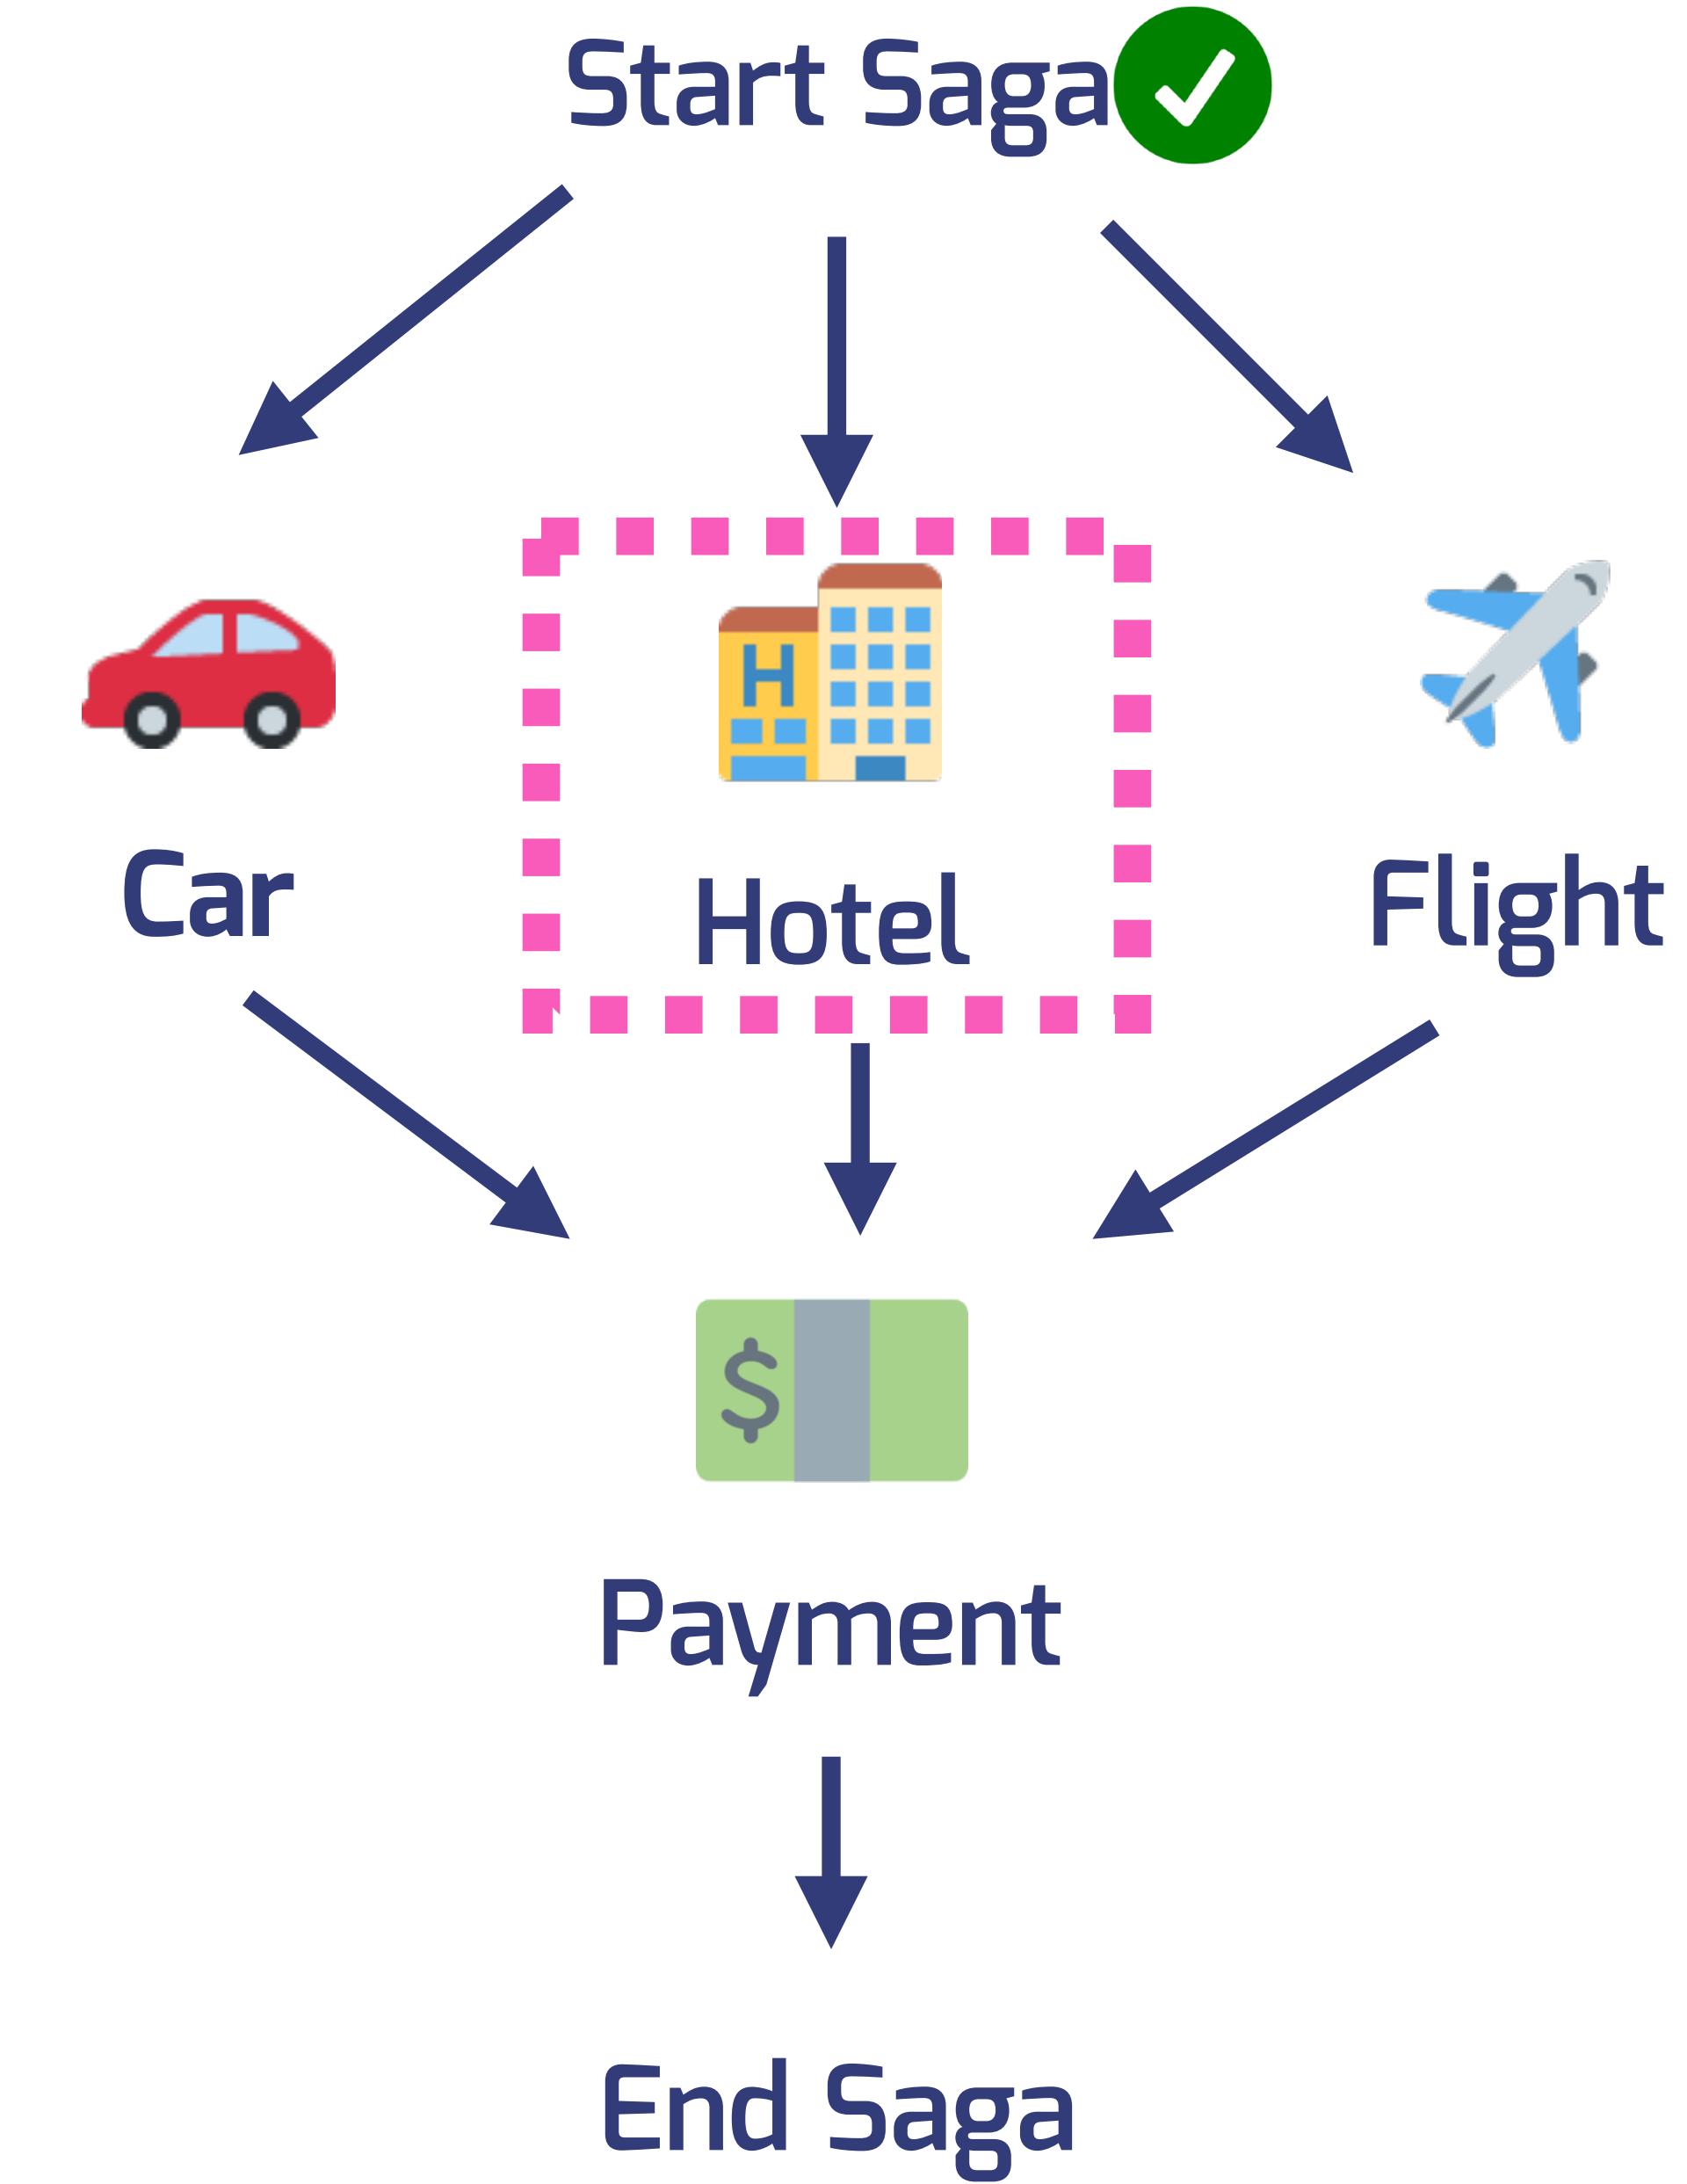
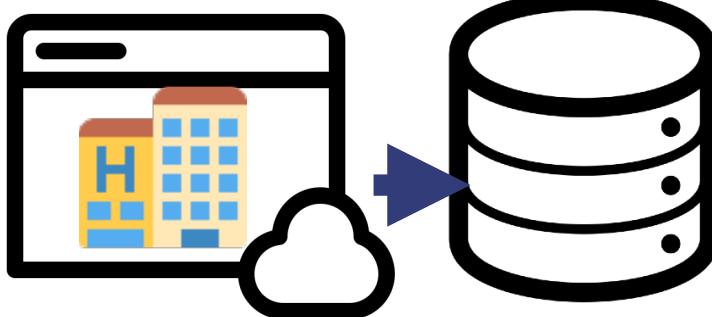


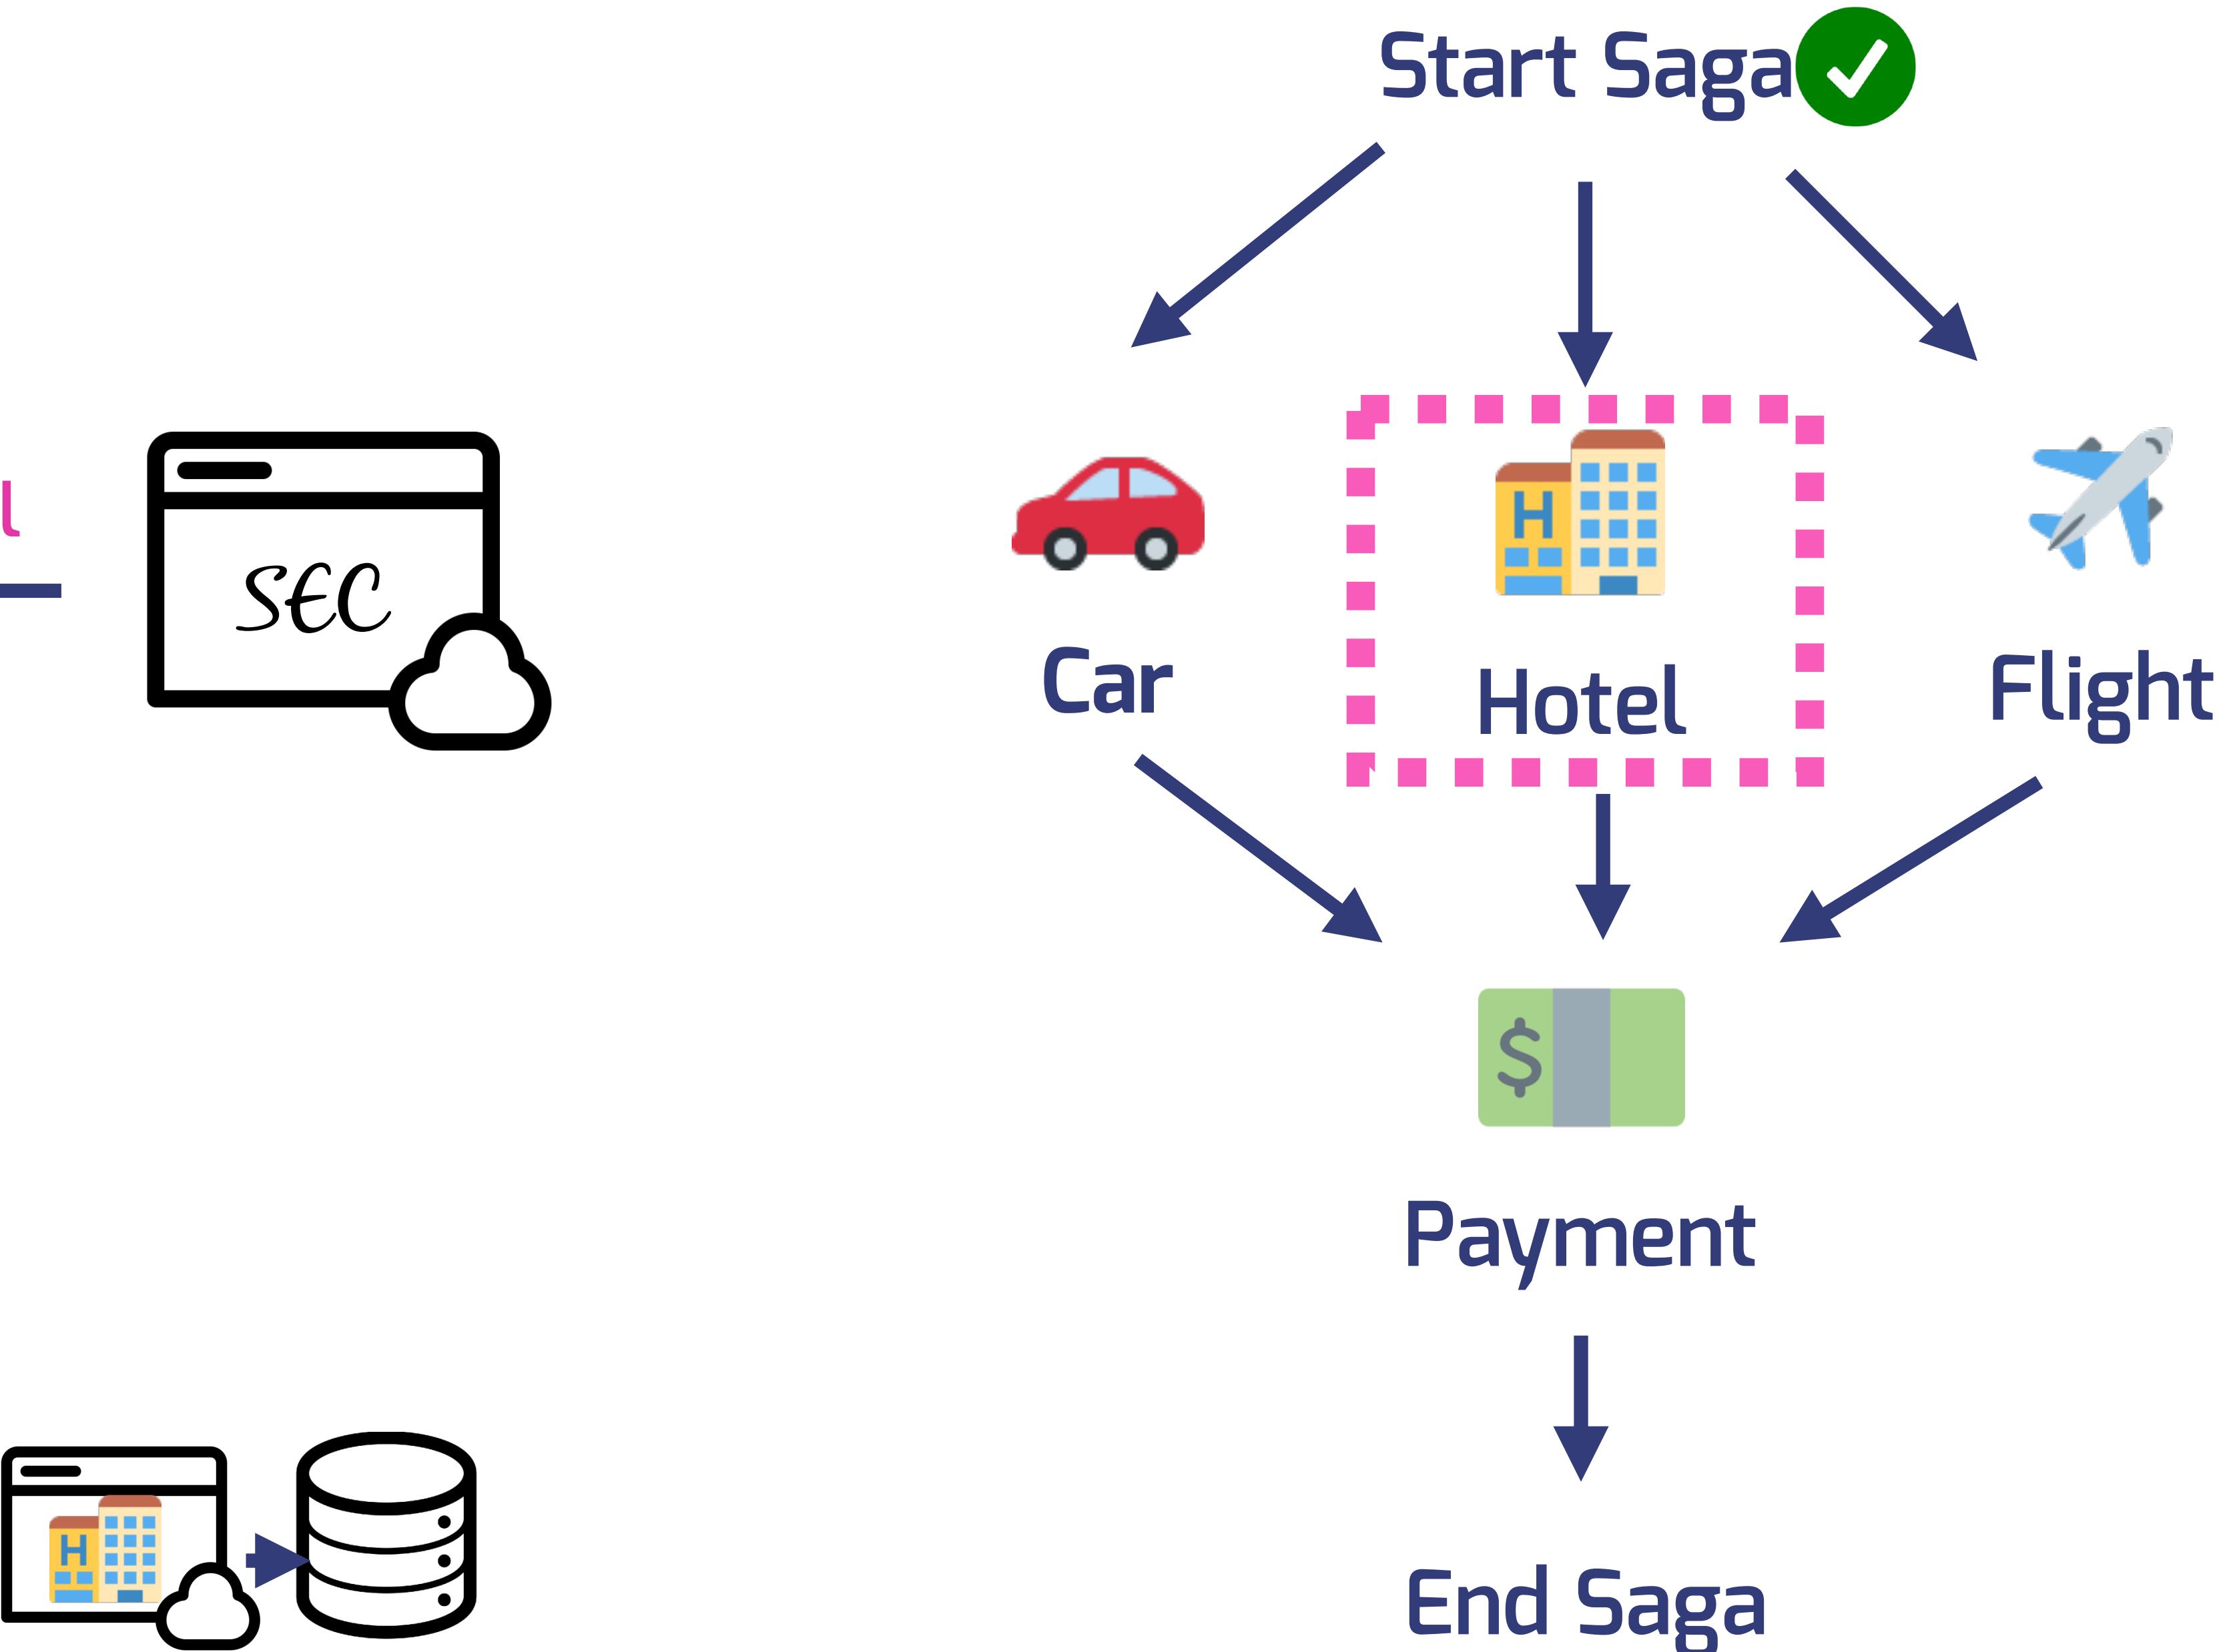
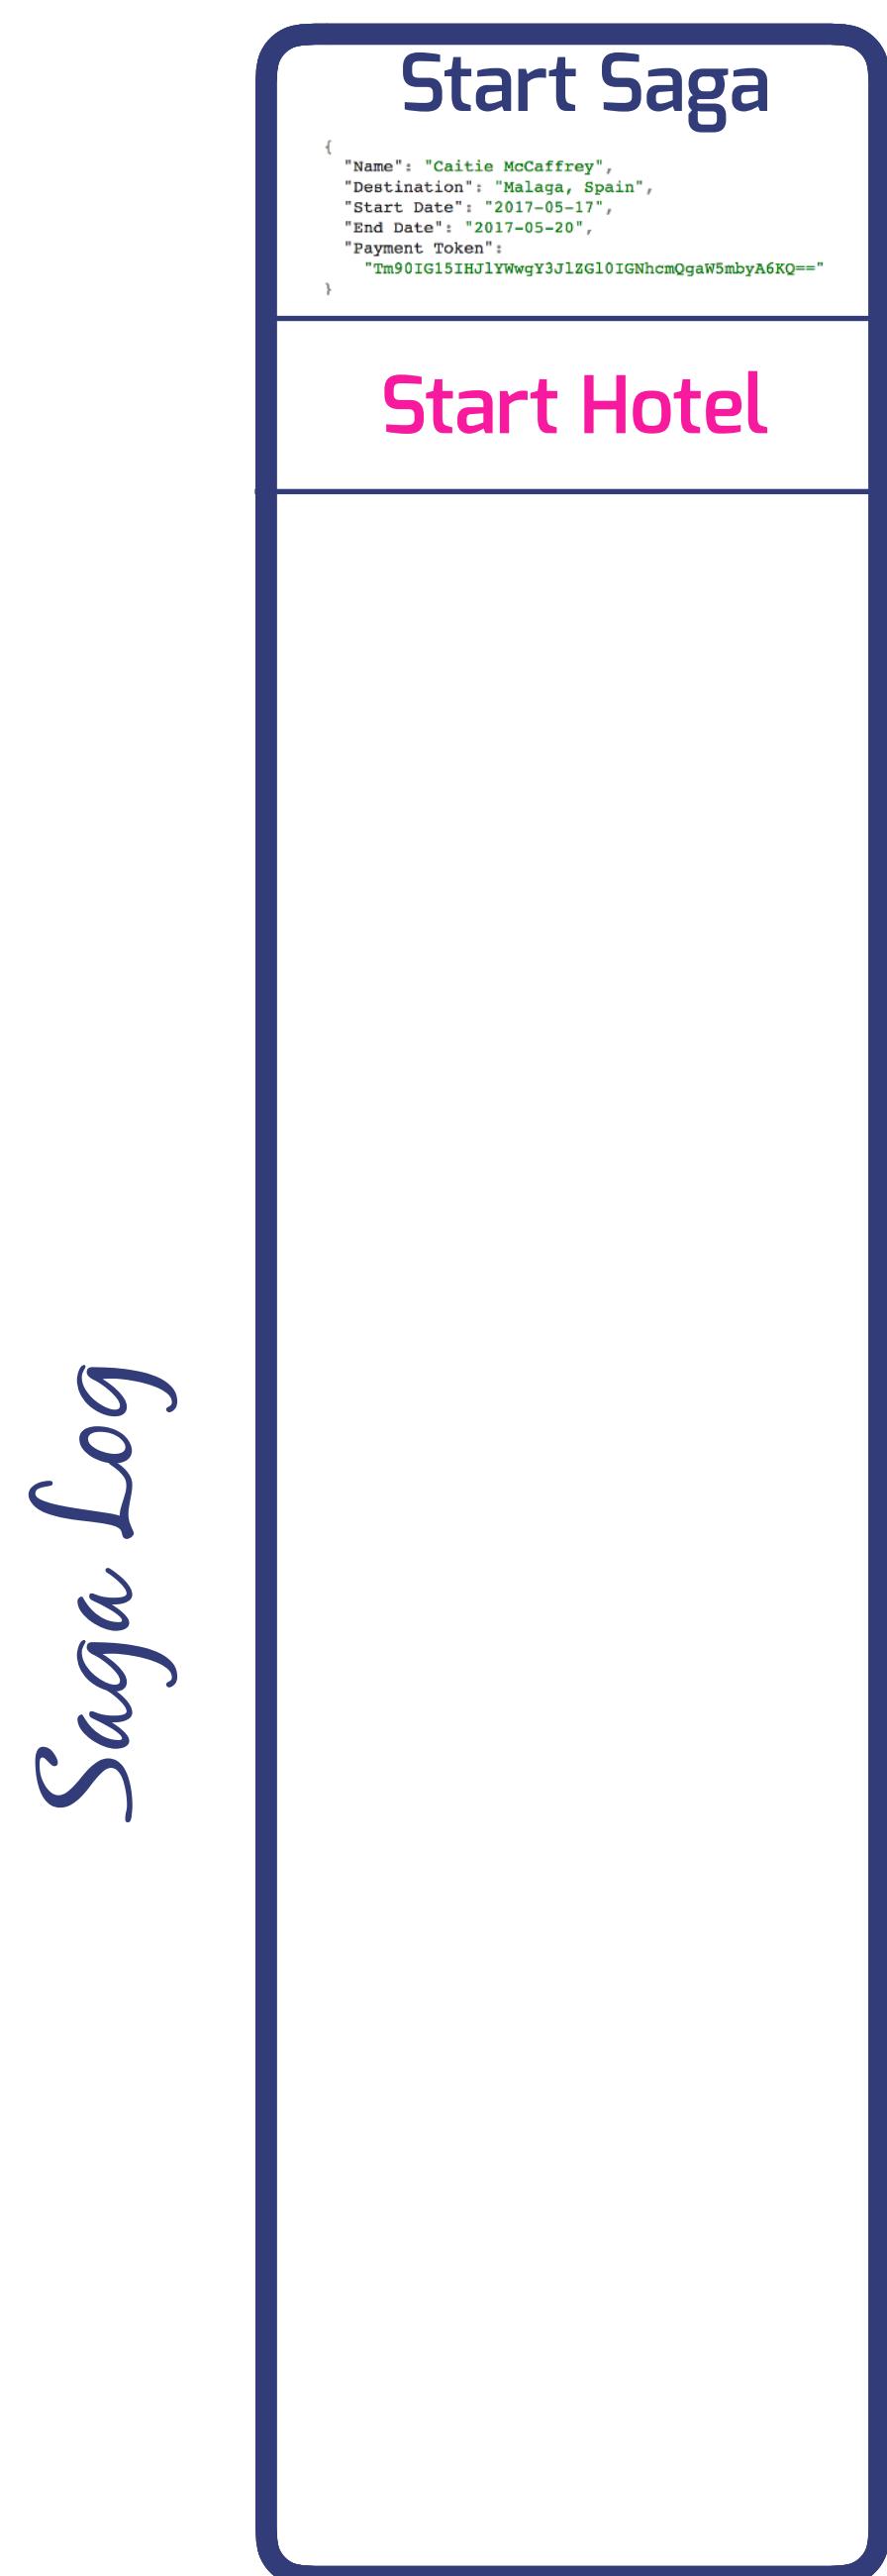


Saga Log

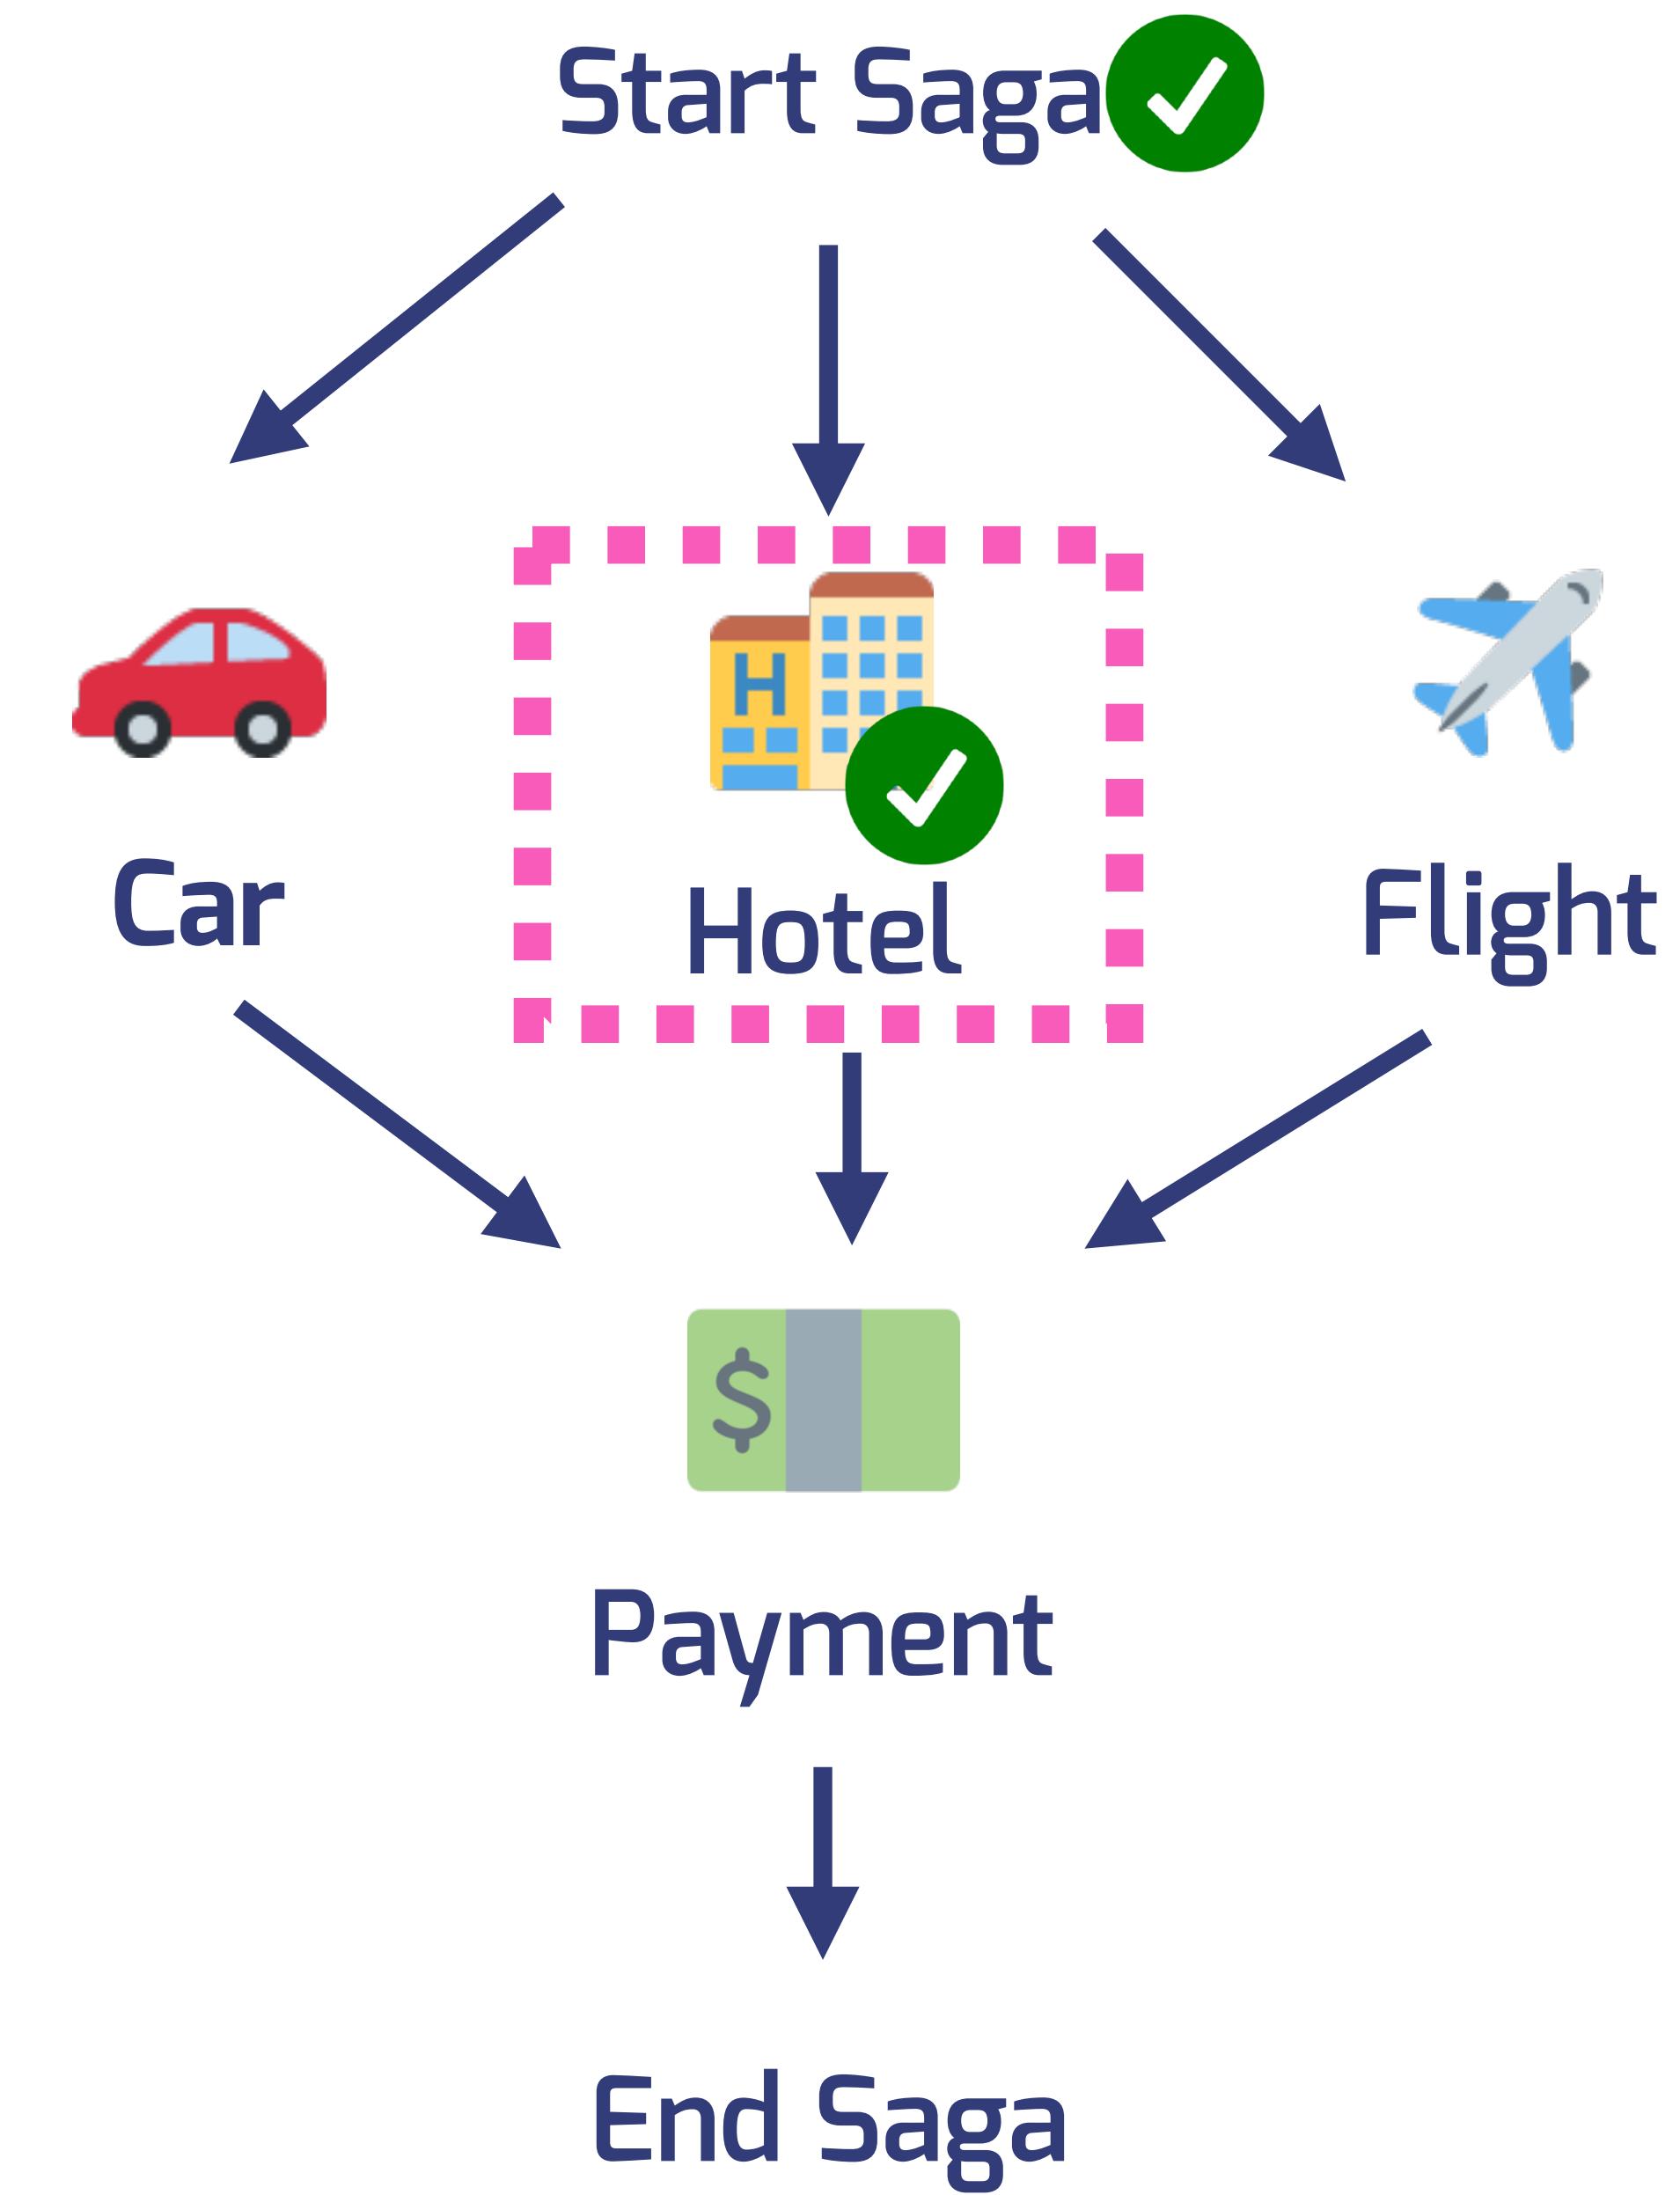
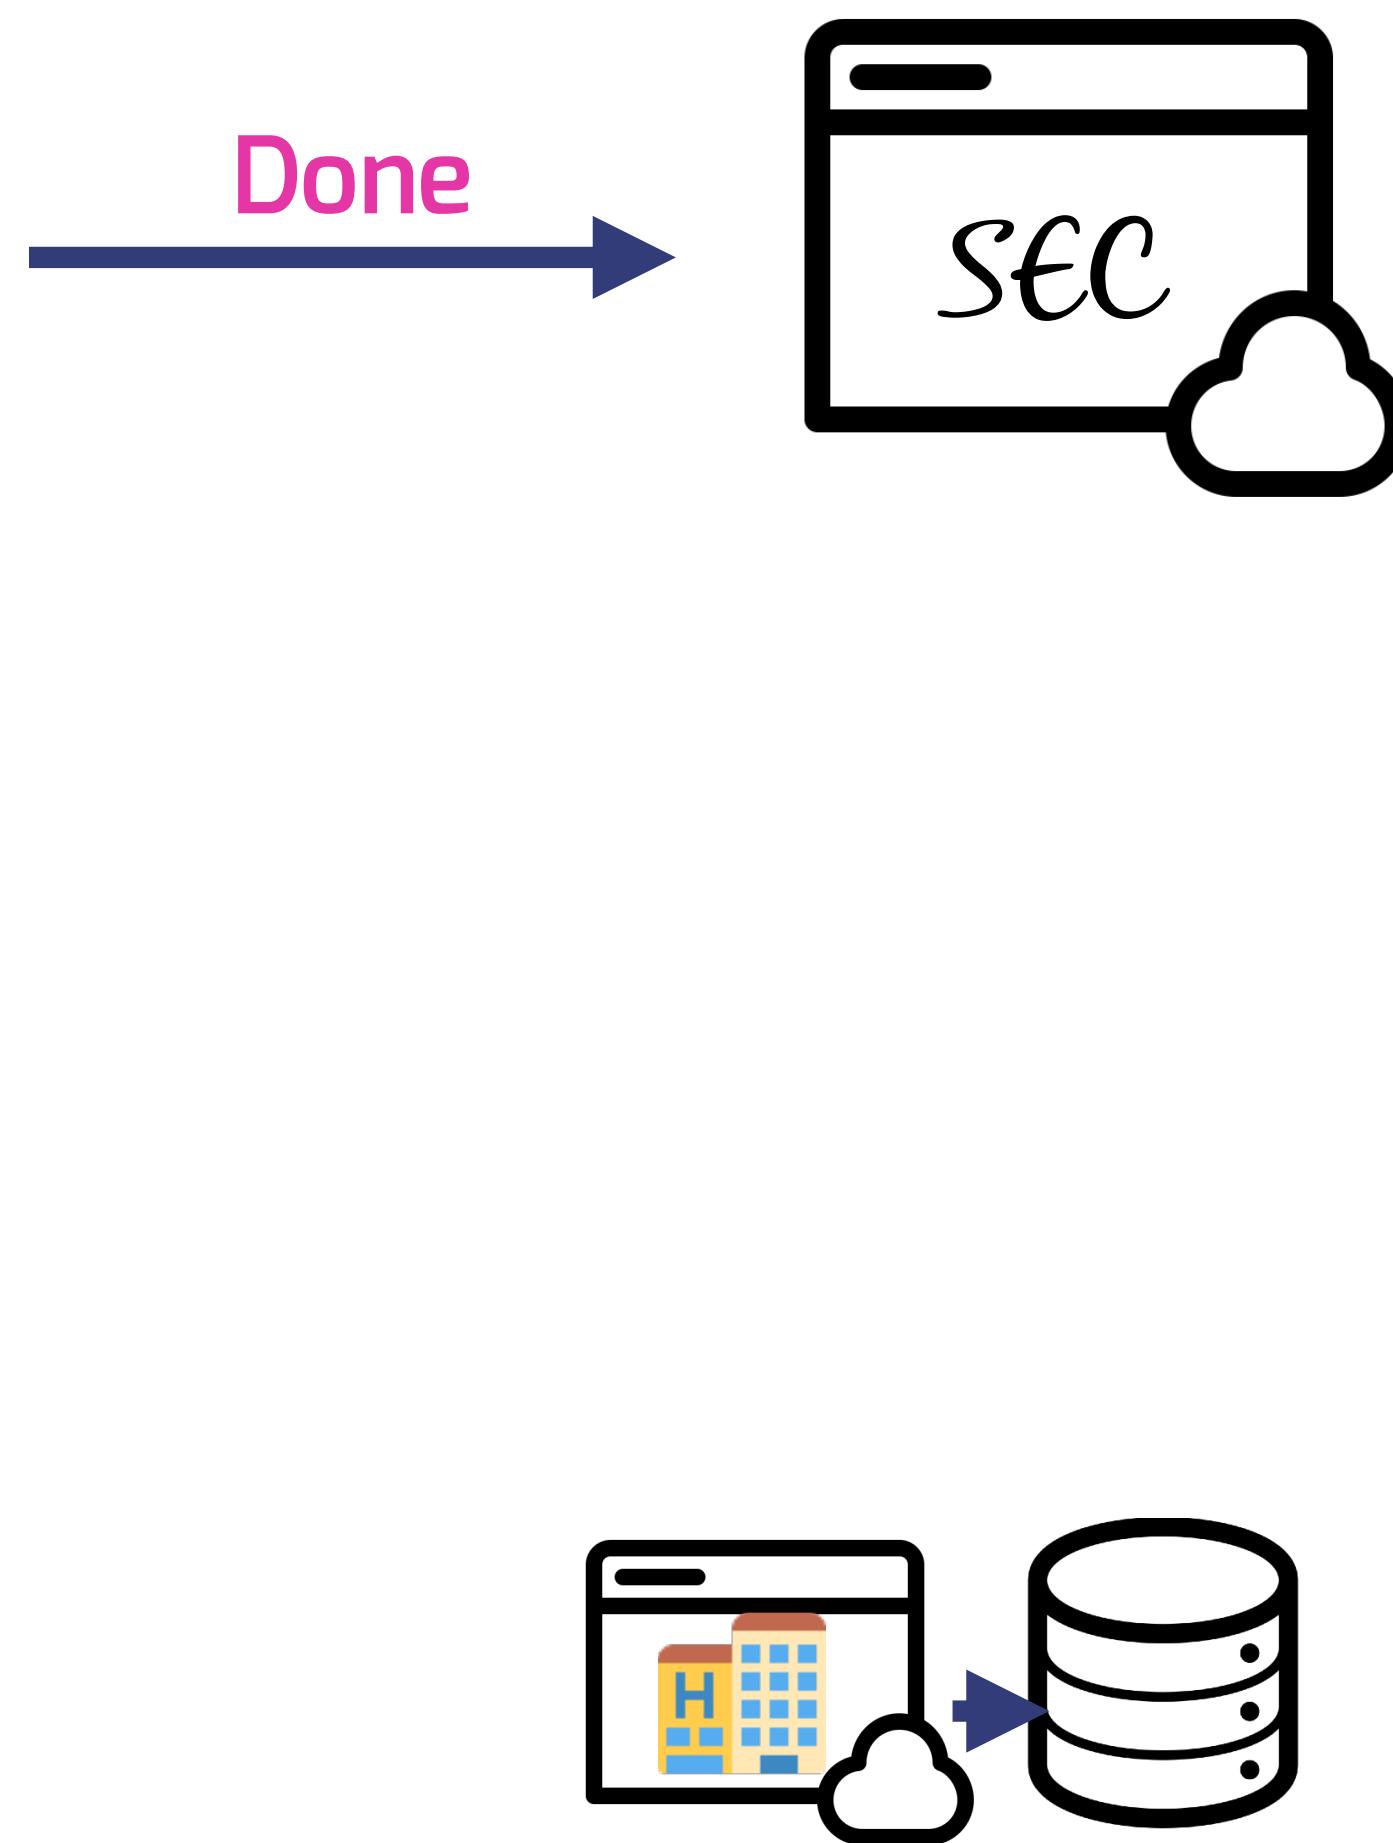


Saga Log





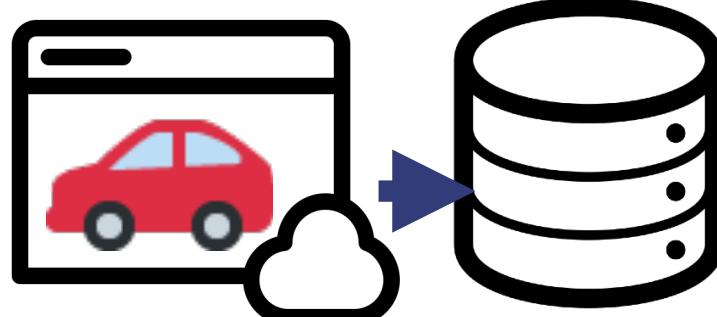
Saga Log



Saga Log



Start Car



Start Saga



Car



Hotel

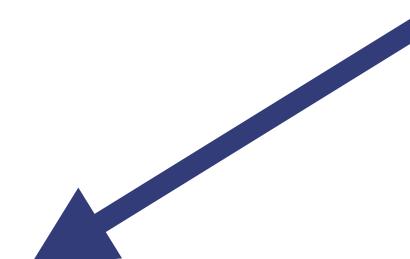
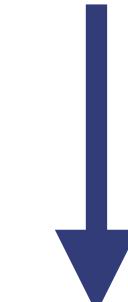
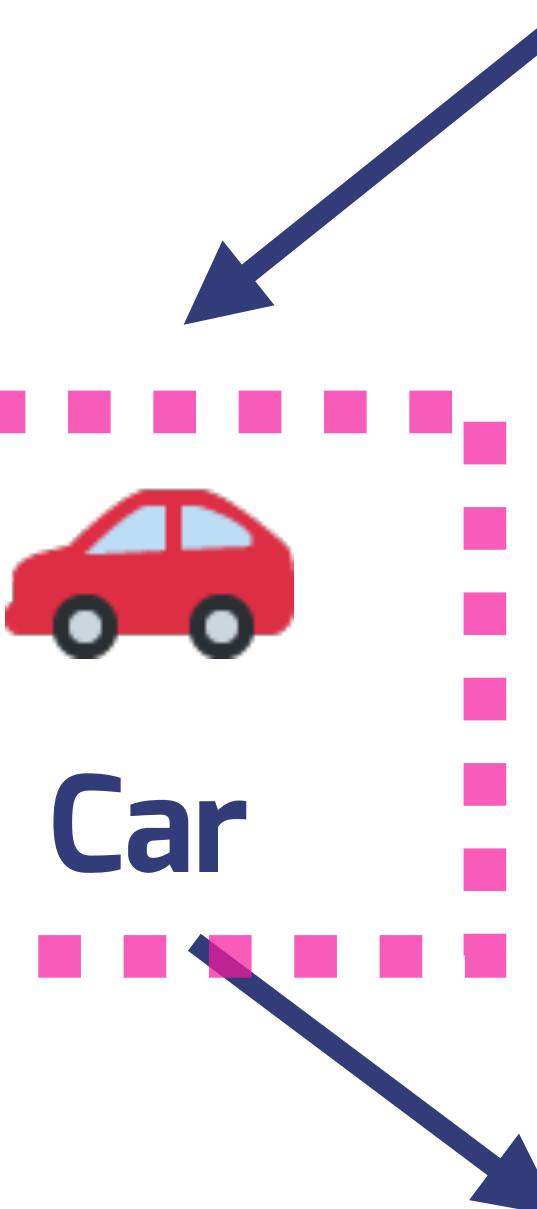


Flight

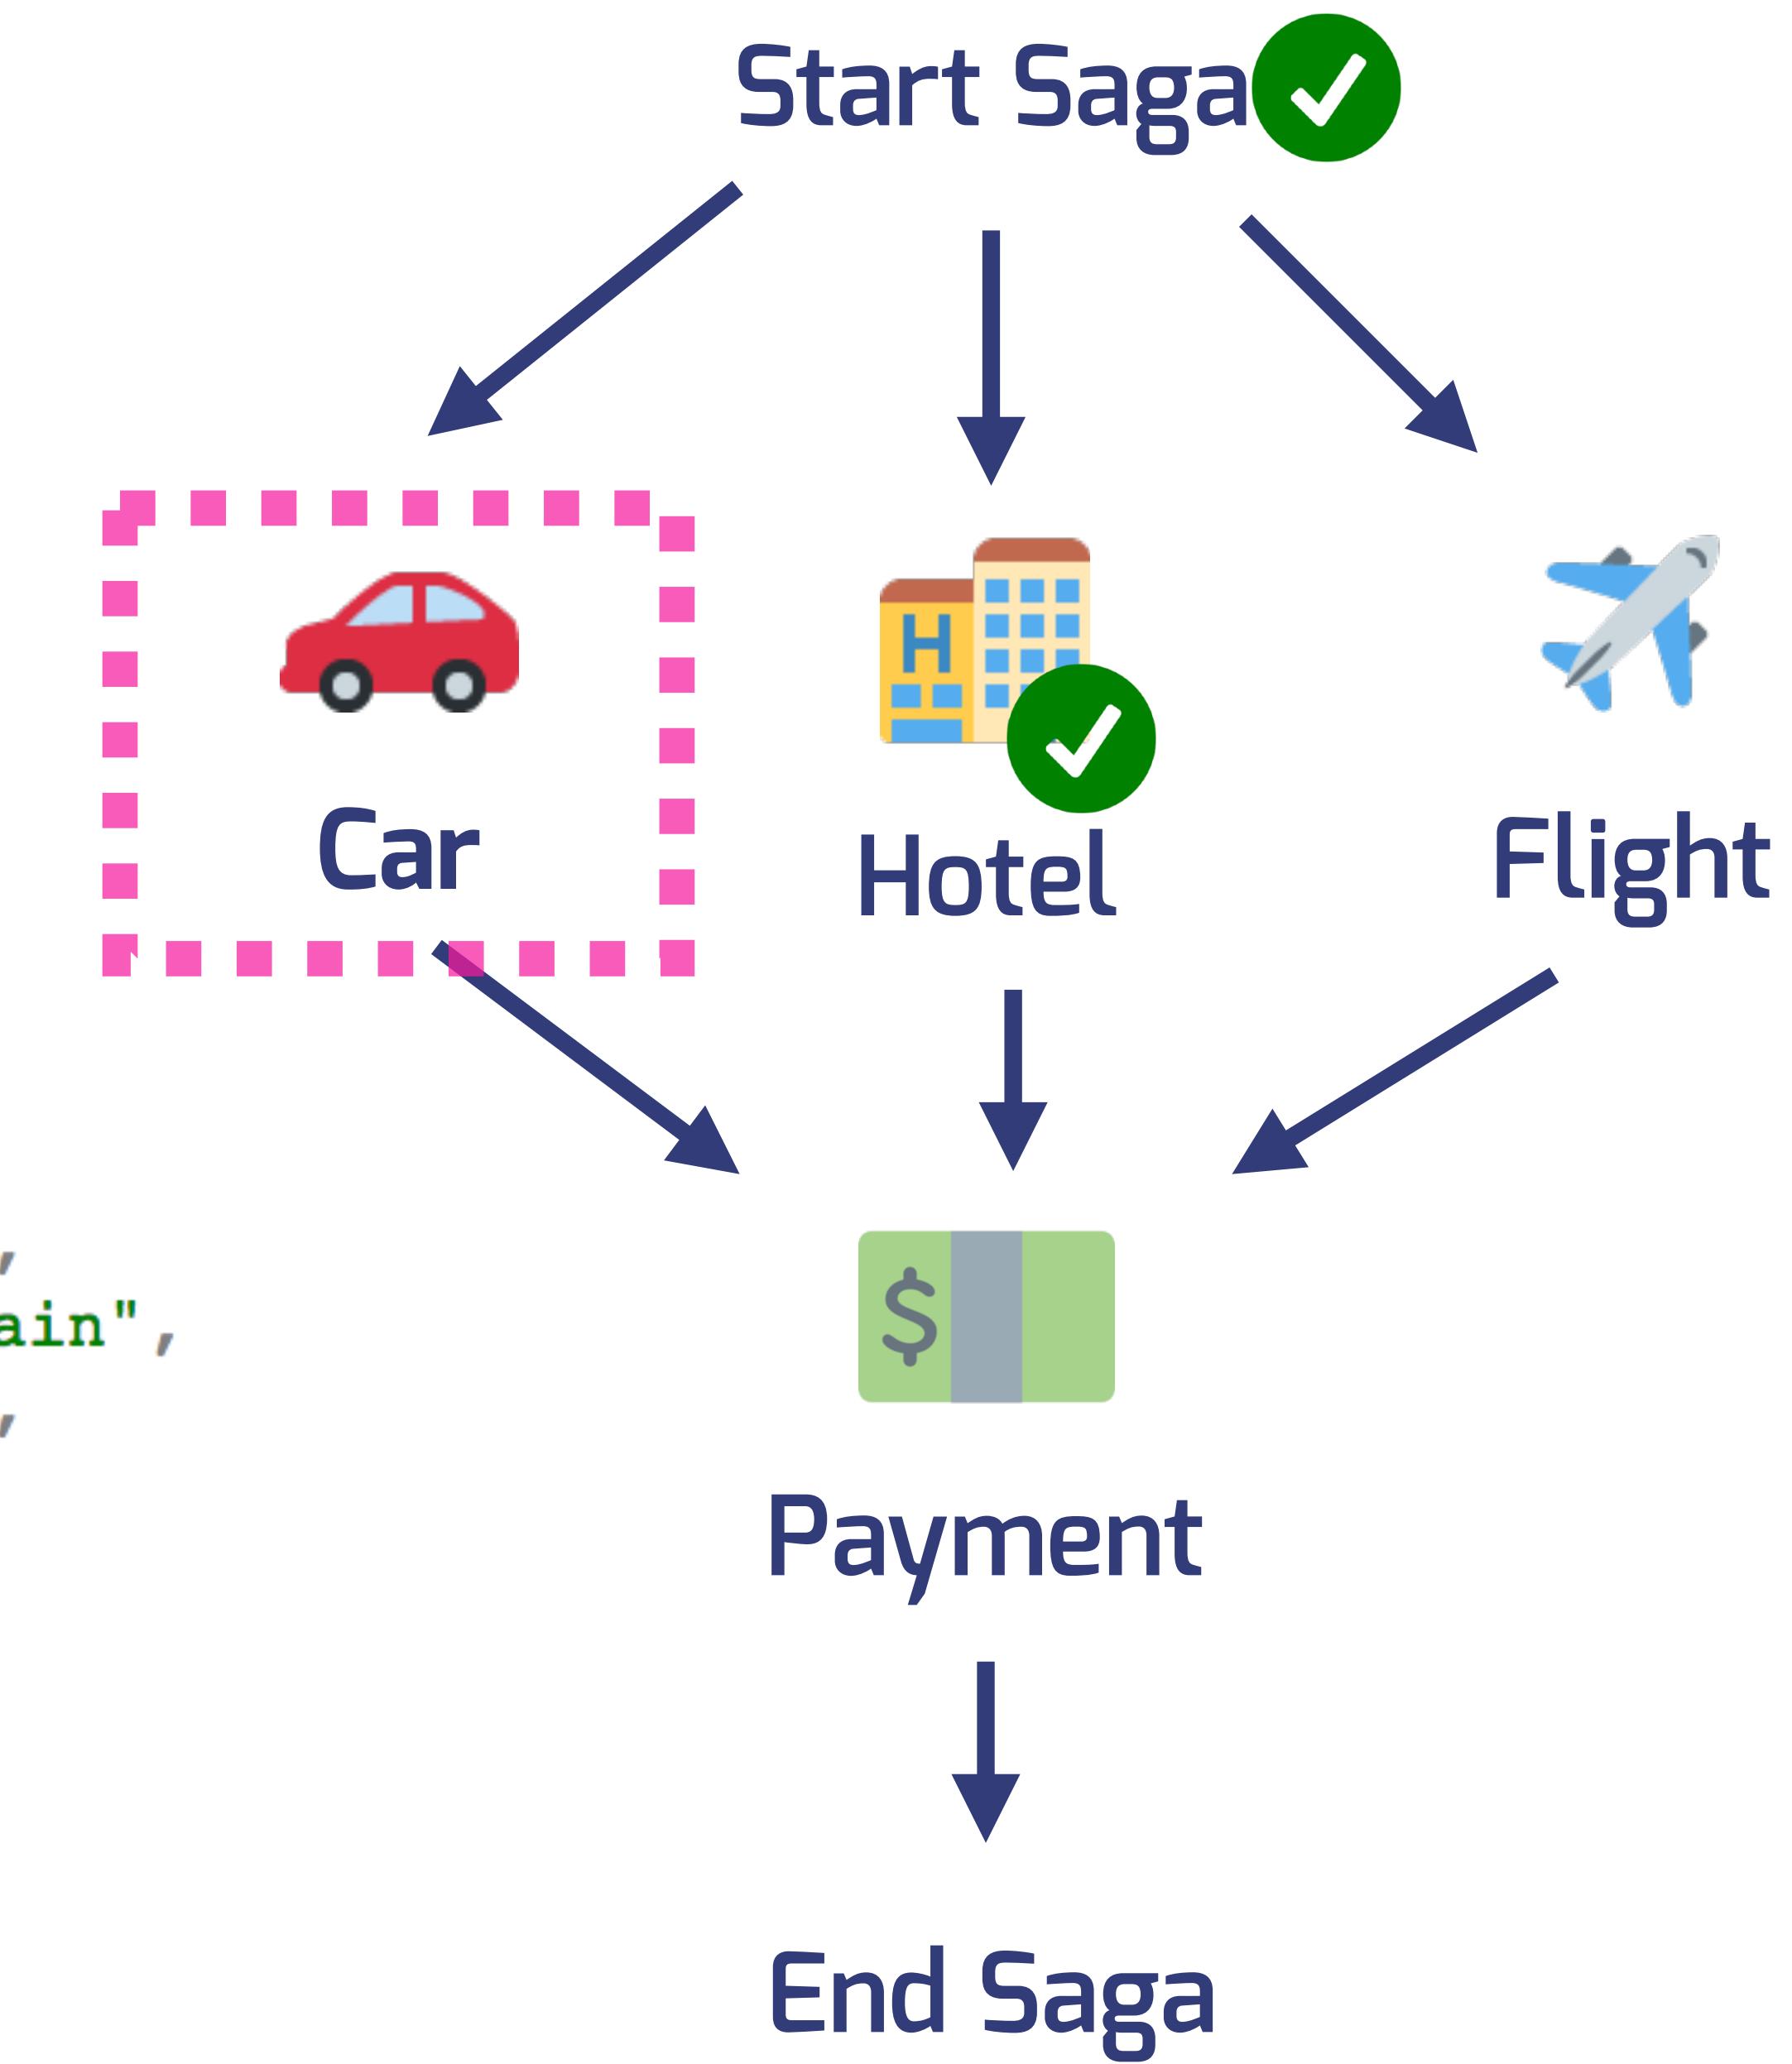
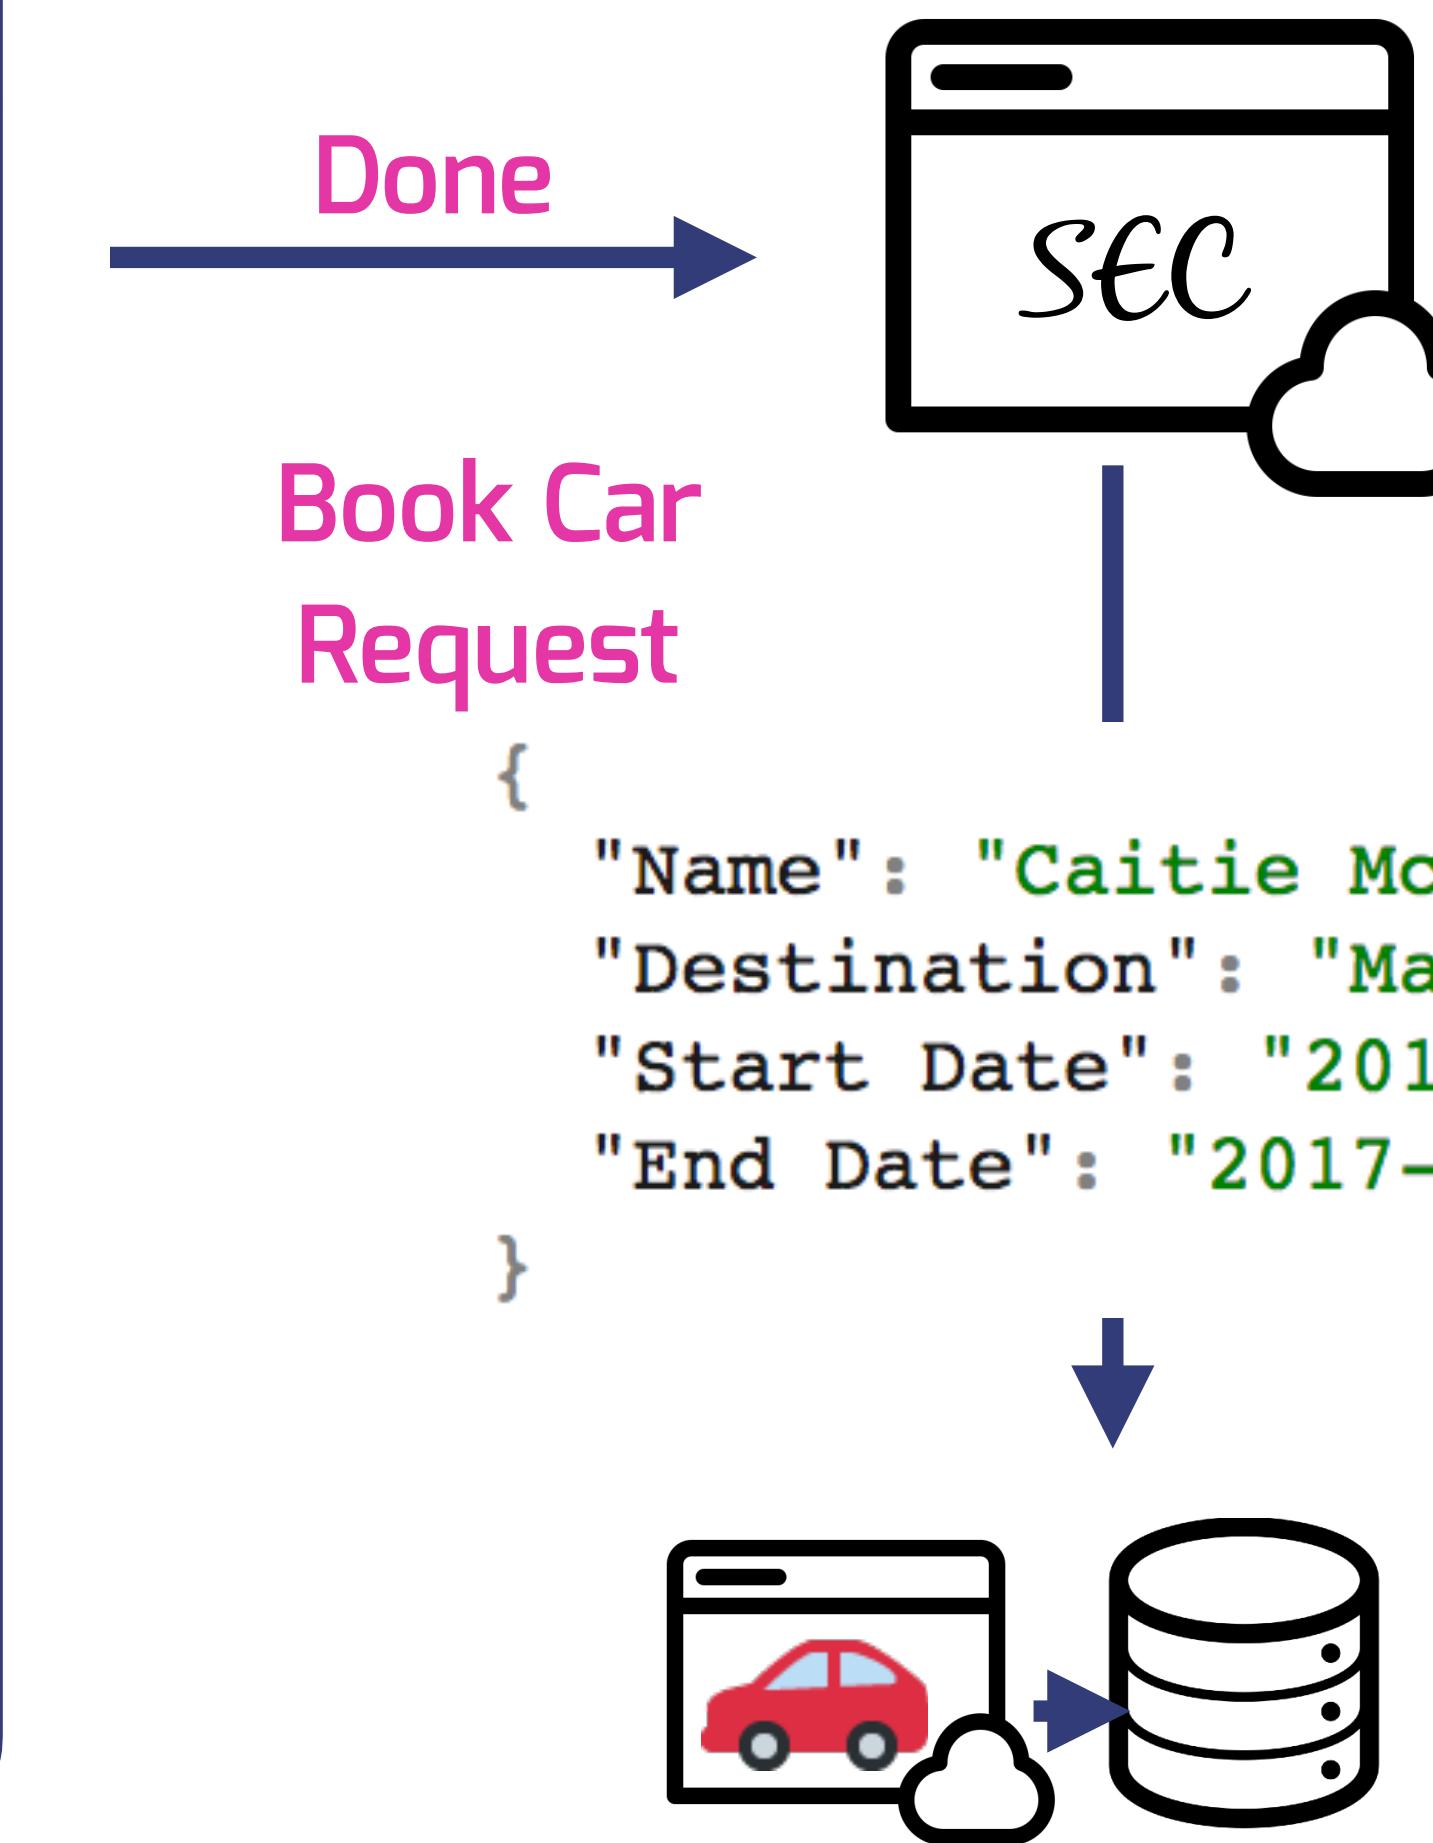


Payment

End Saga



Saga Log



Saga Log

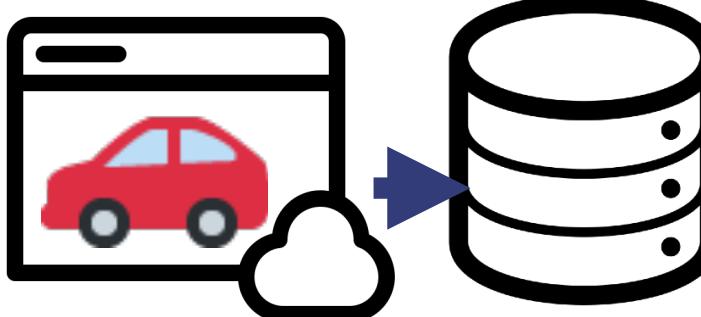


End Car



Book Car Response

```
{ "Success": "true", "Confirmation Number": "ABC456" }
```



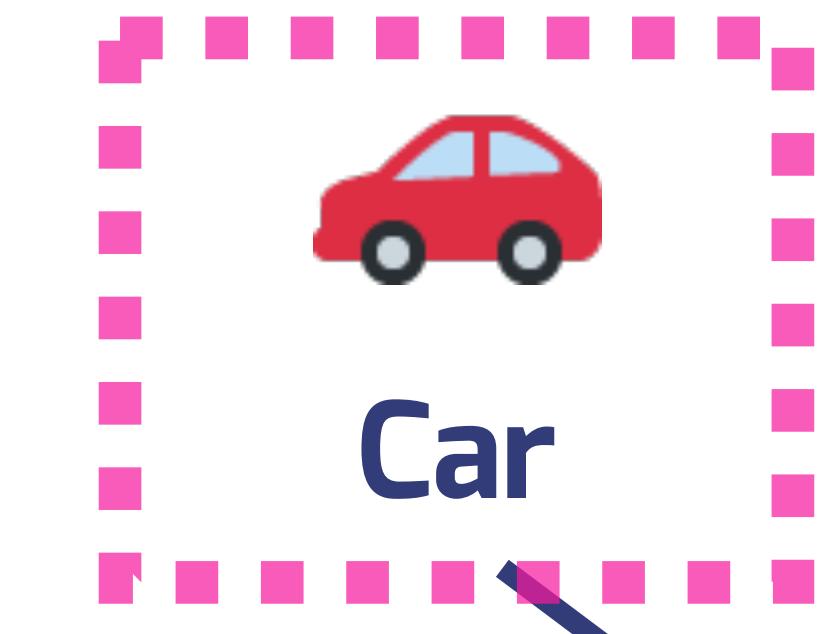
Start Saga



Hotel



Flight

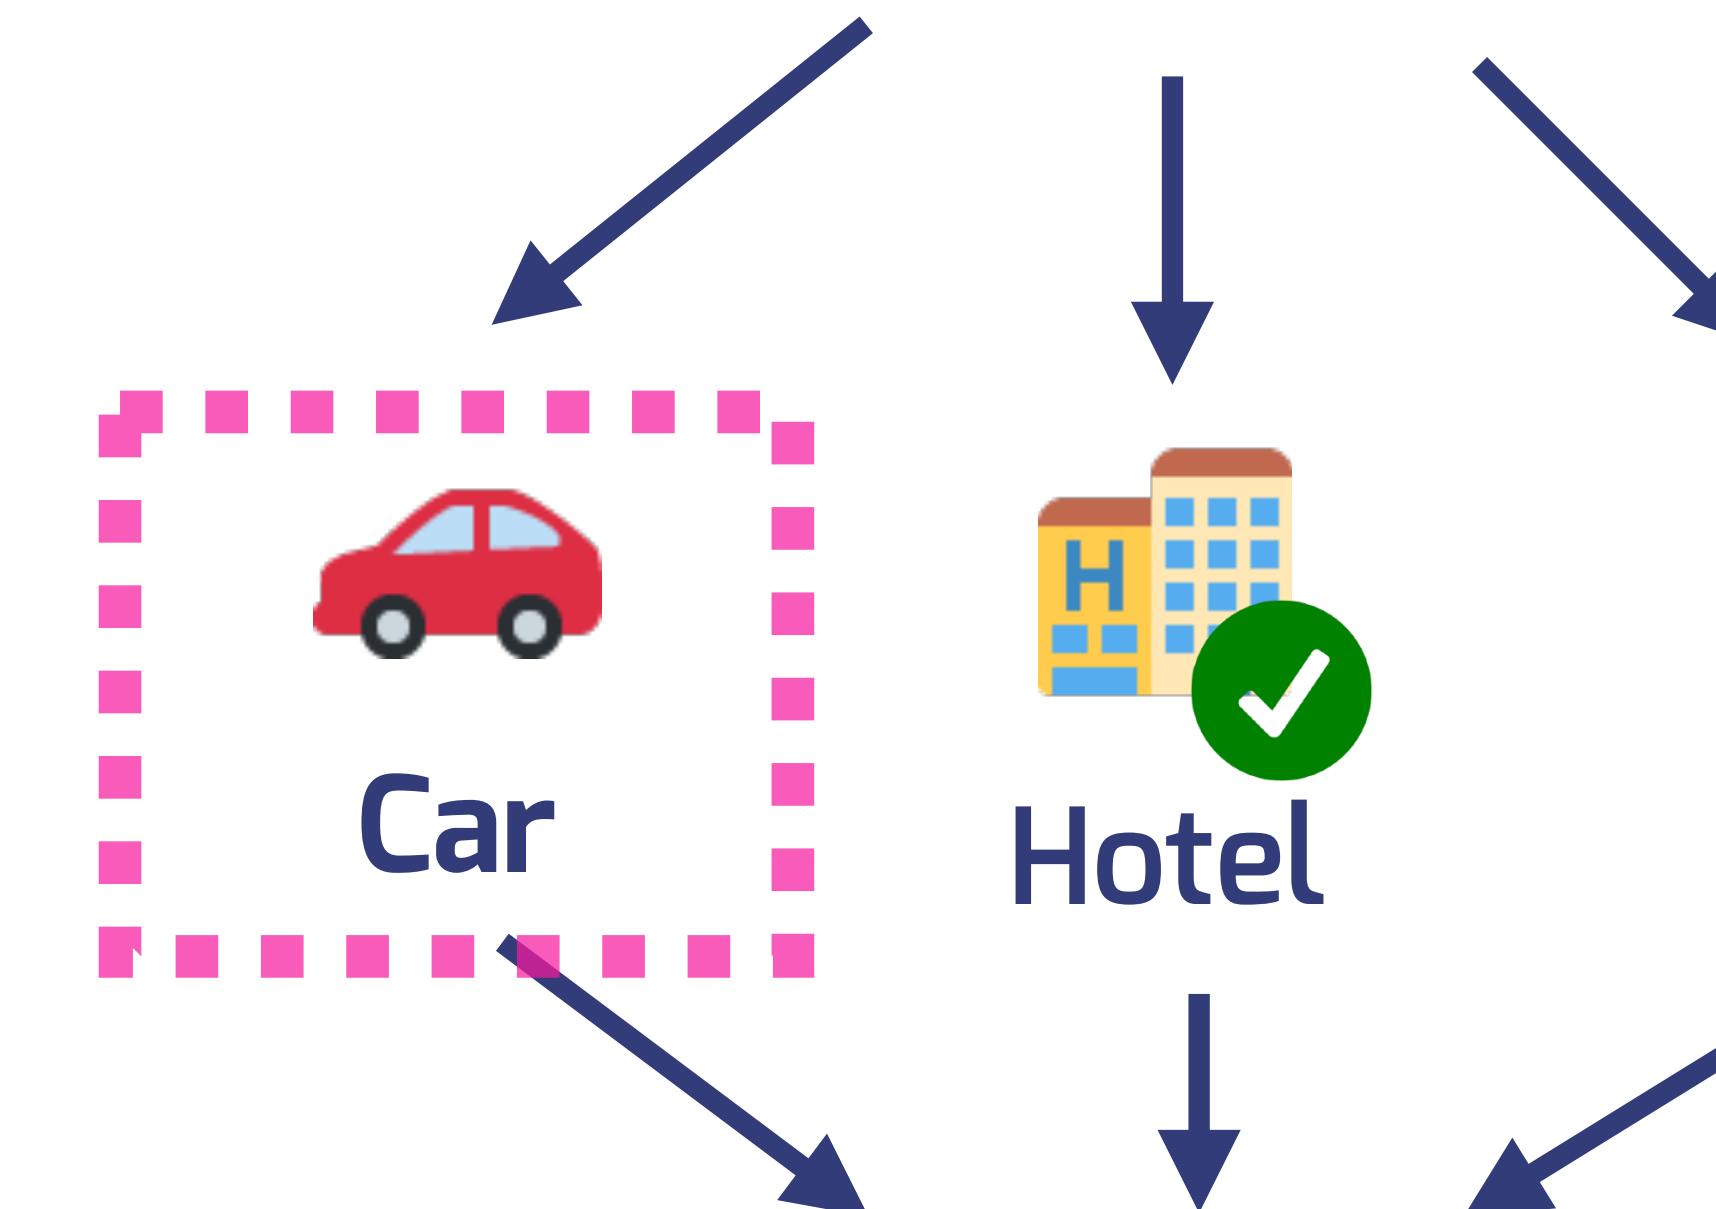


Car

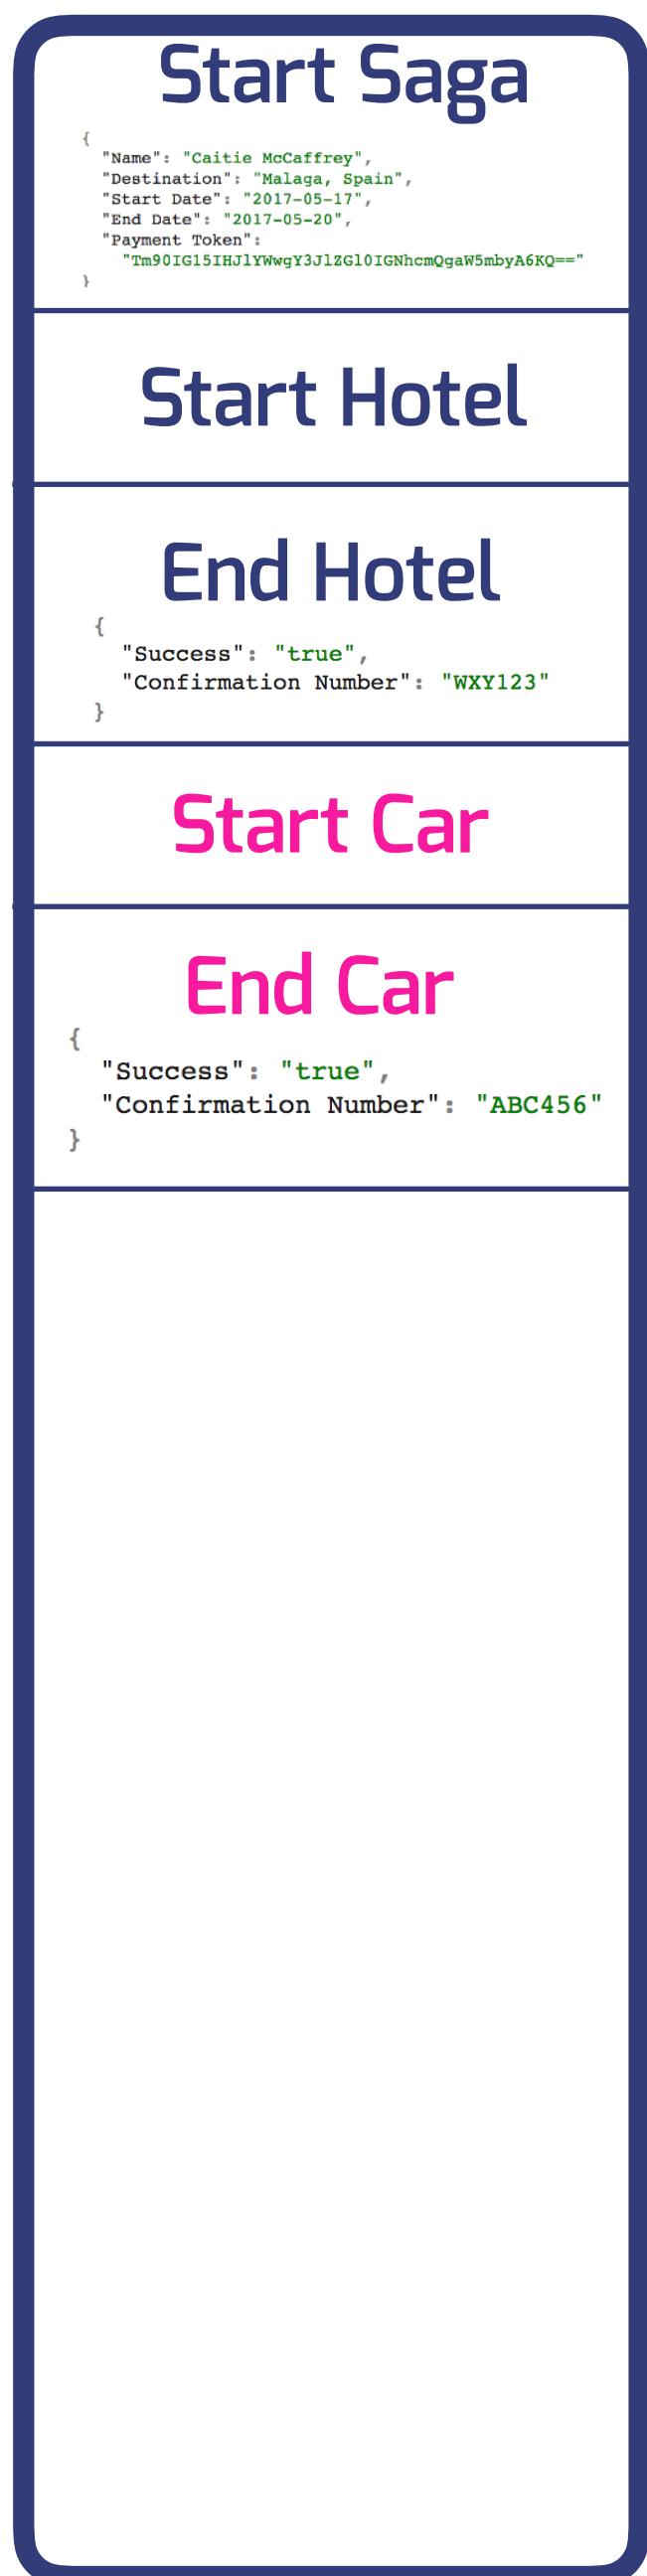


Payment

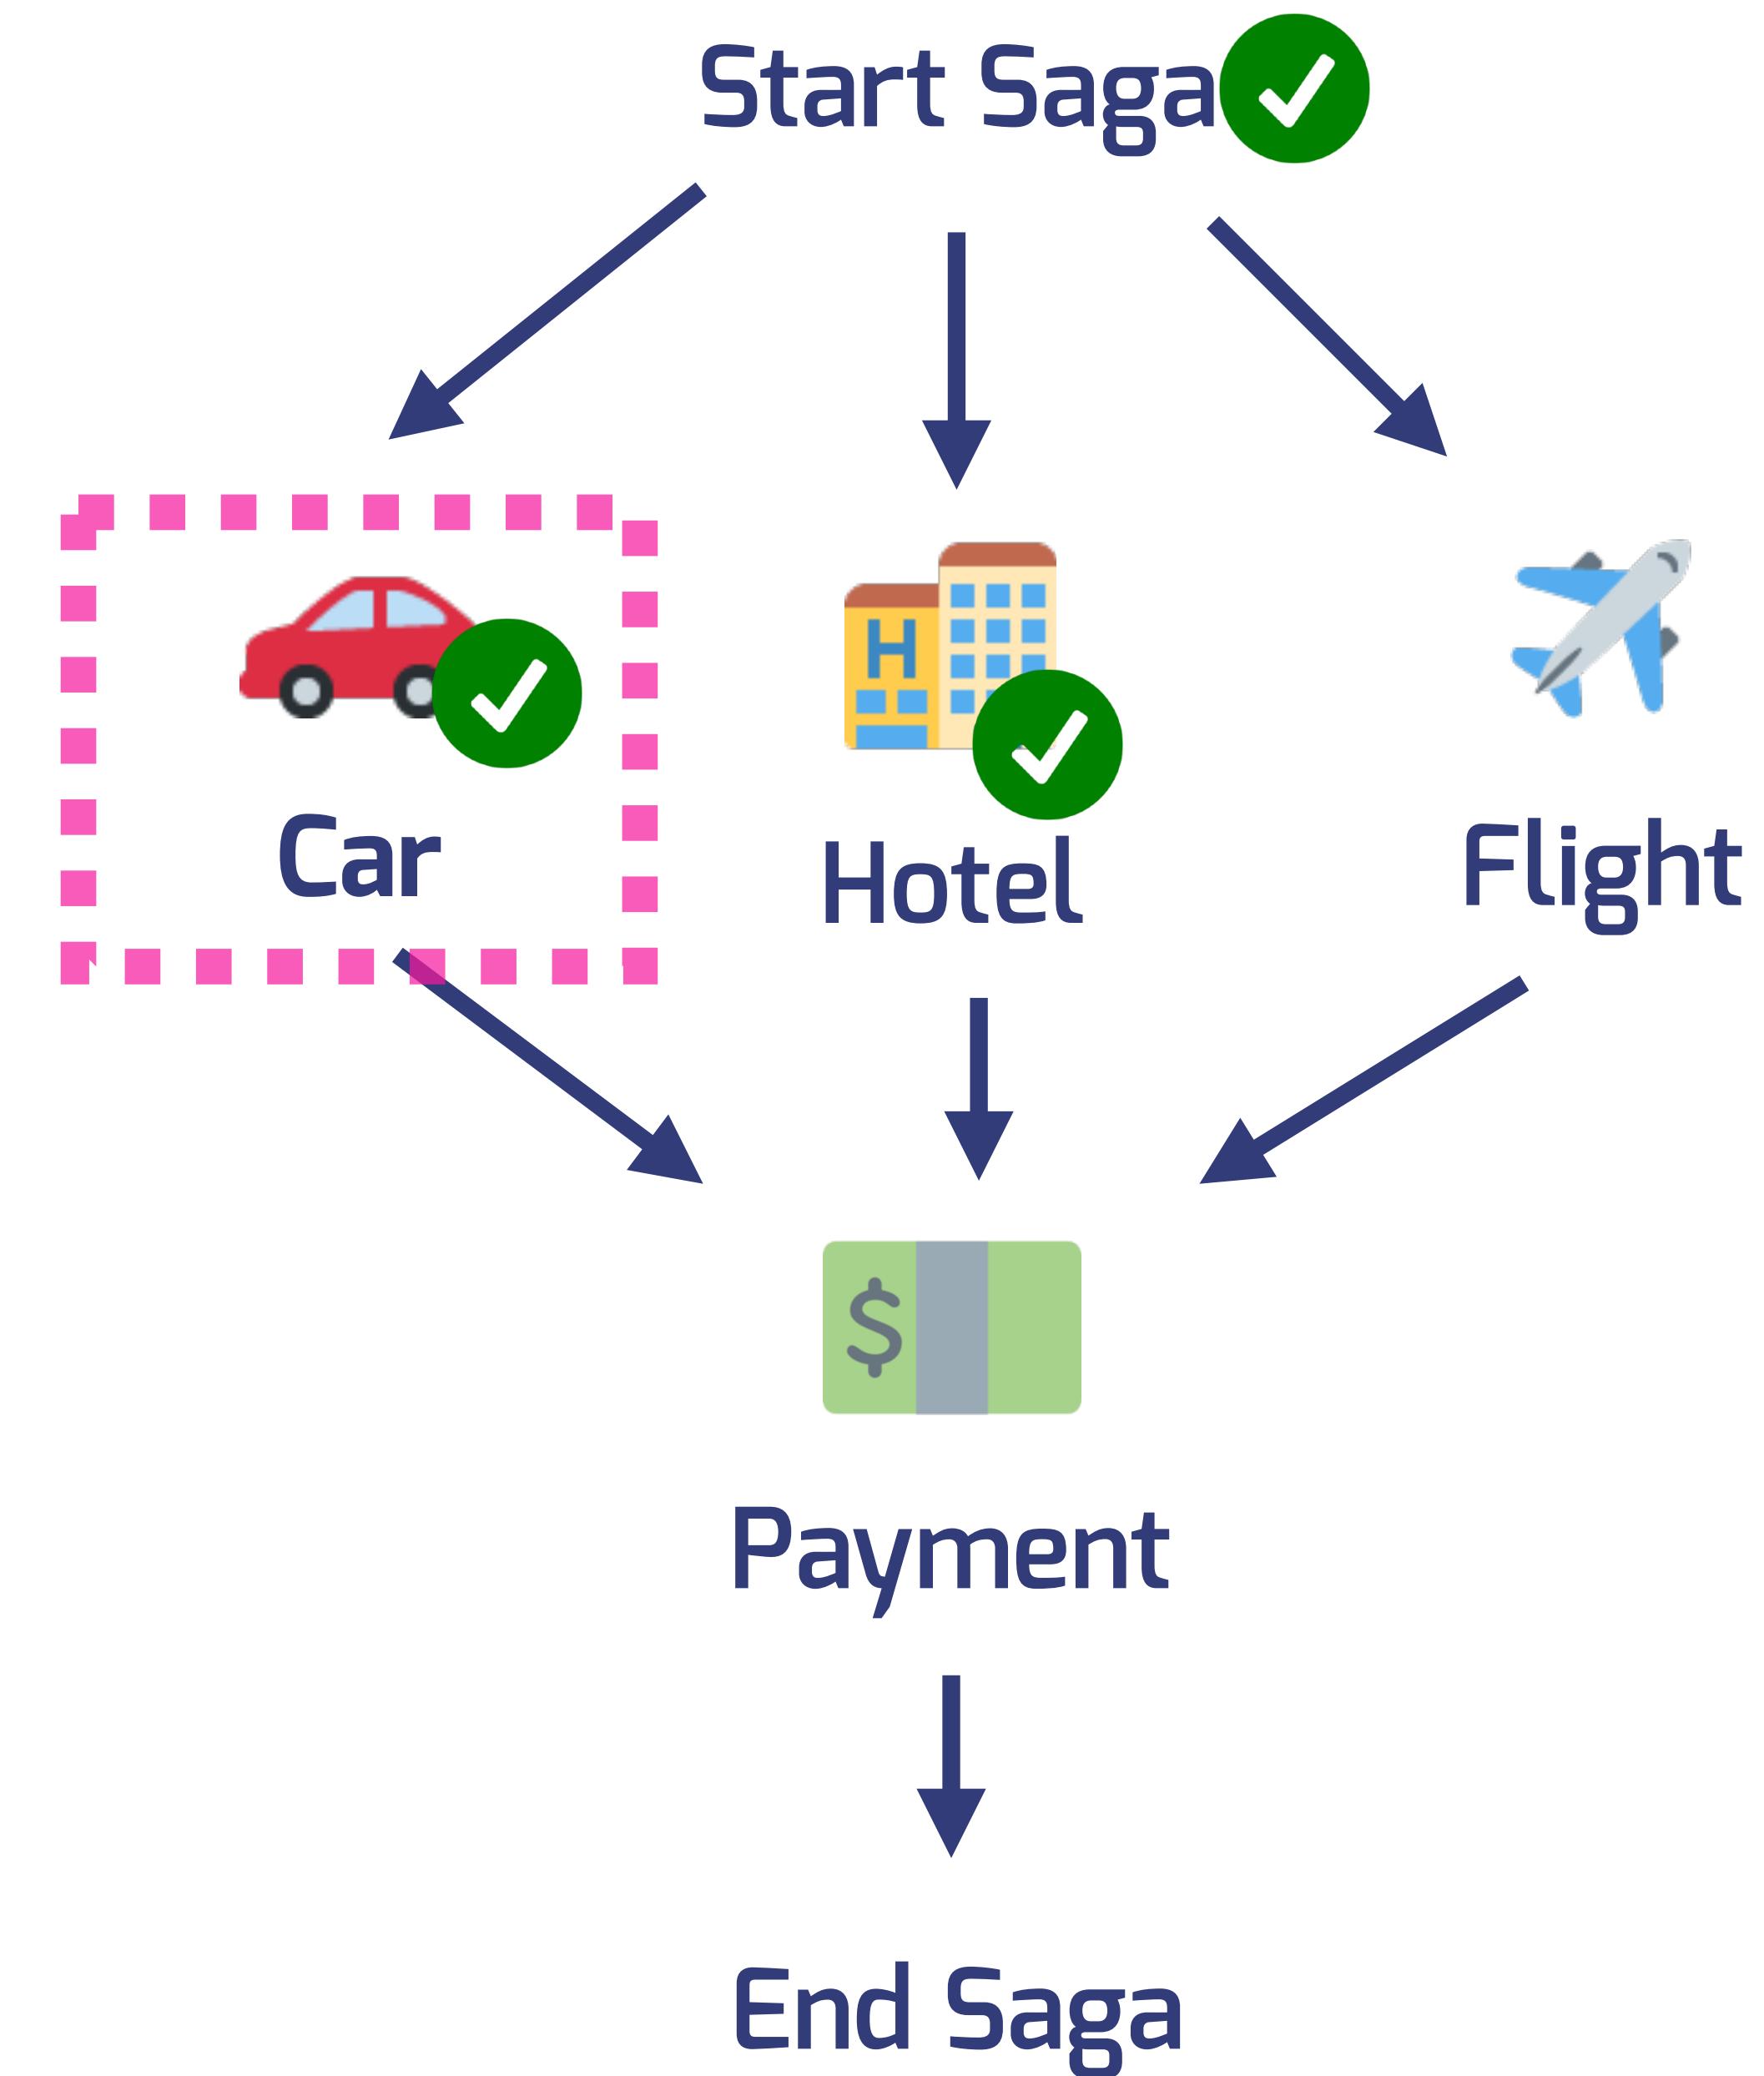
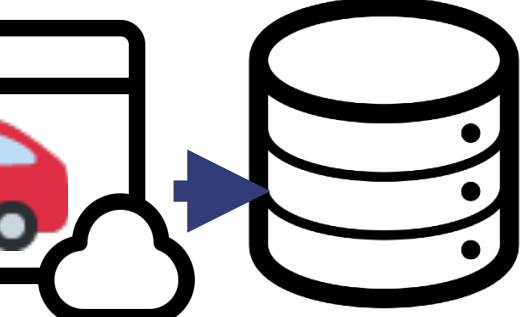
End Saga



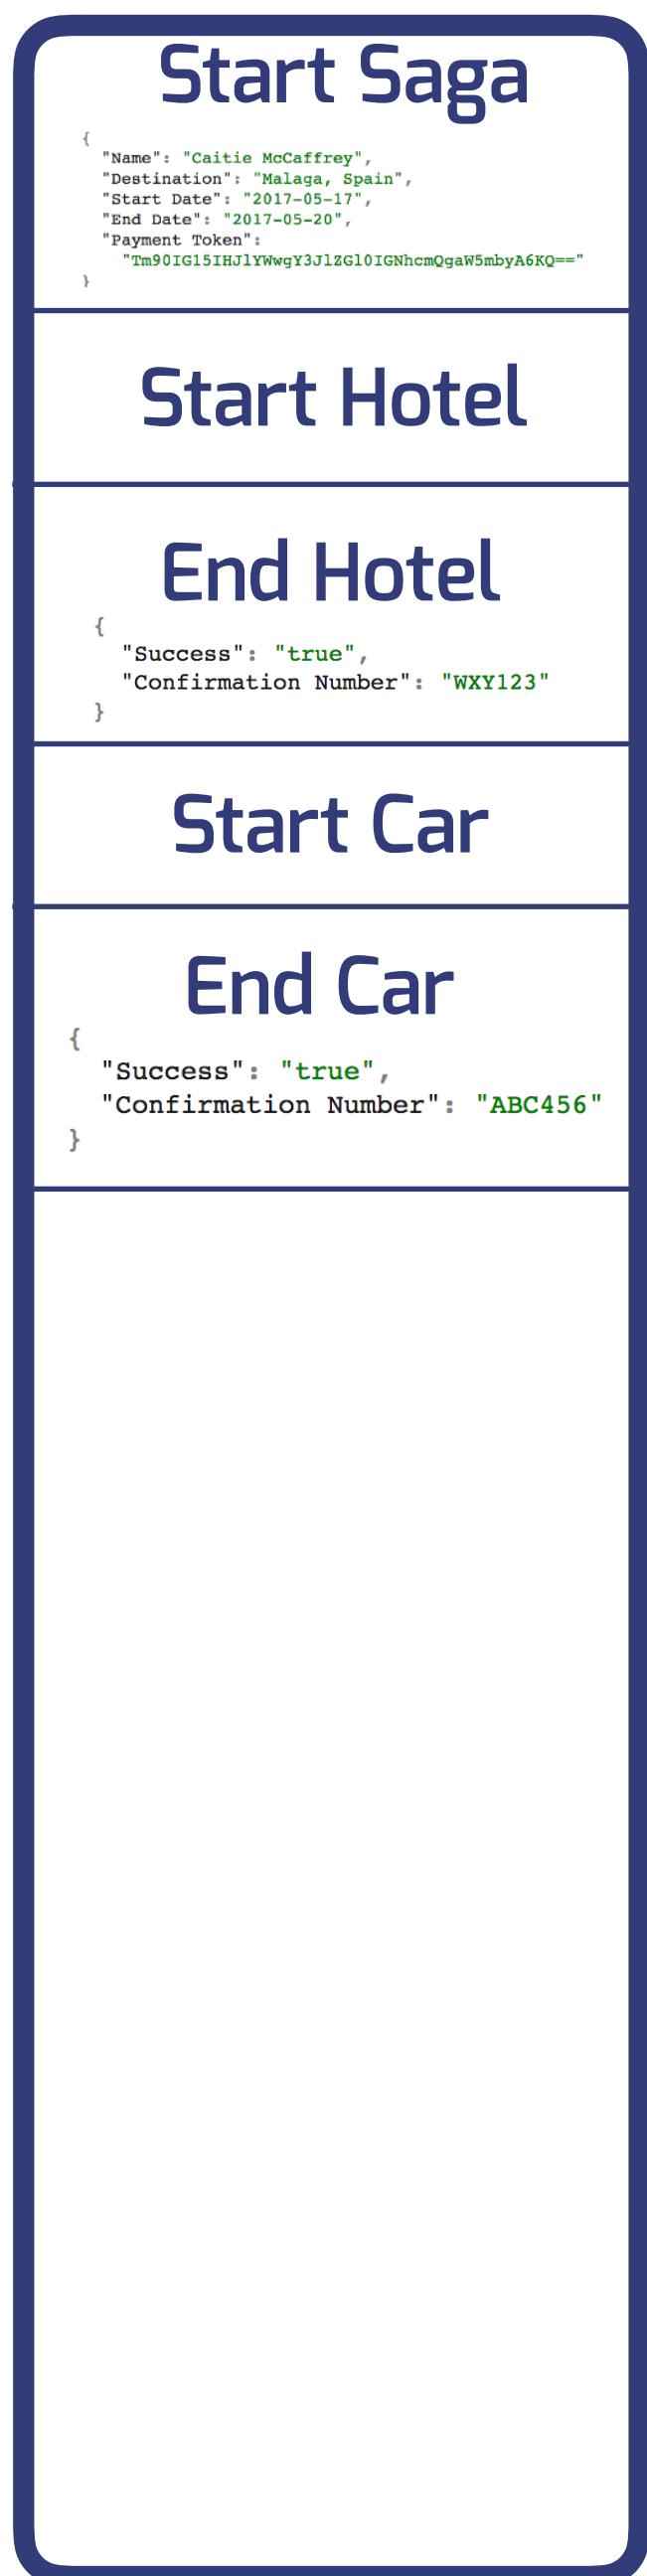
Saga Log



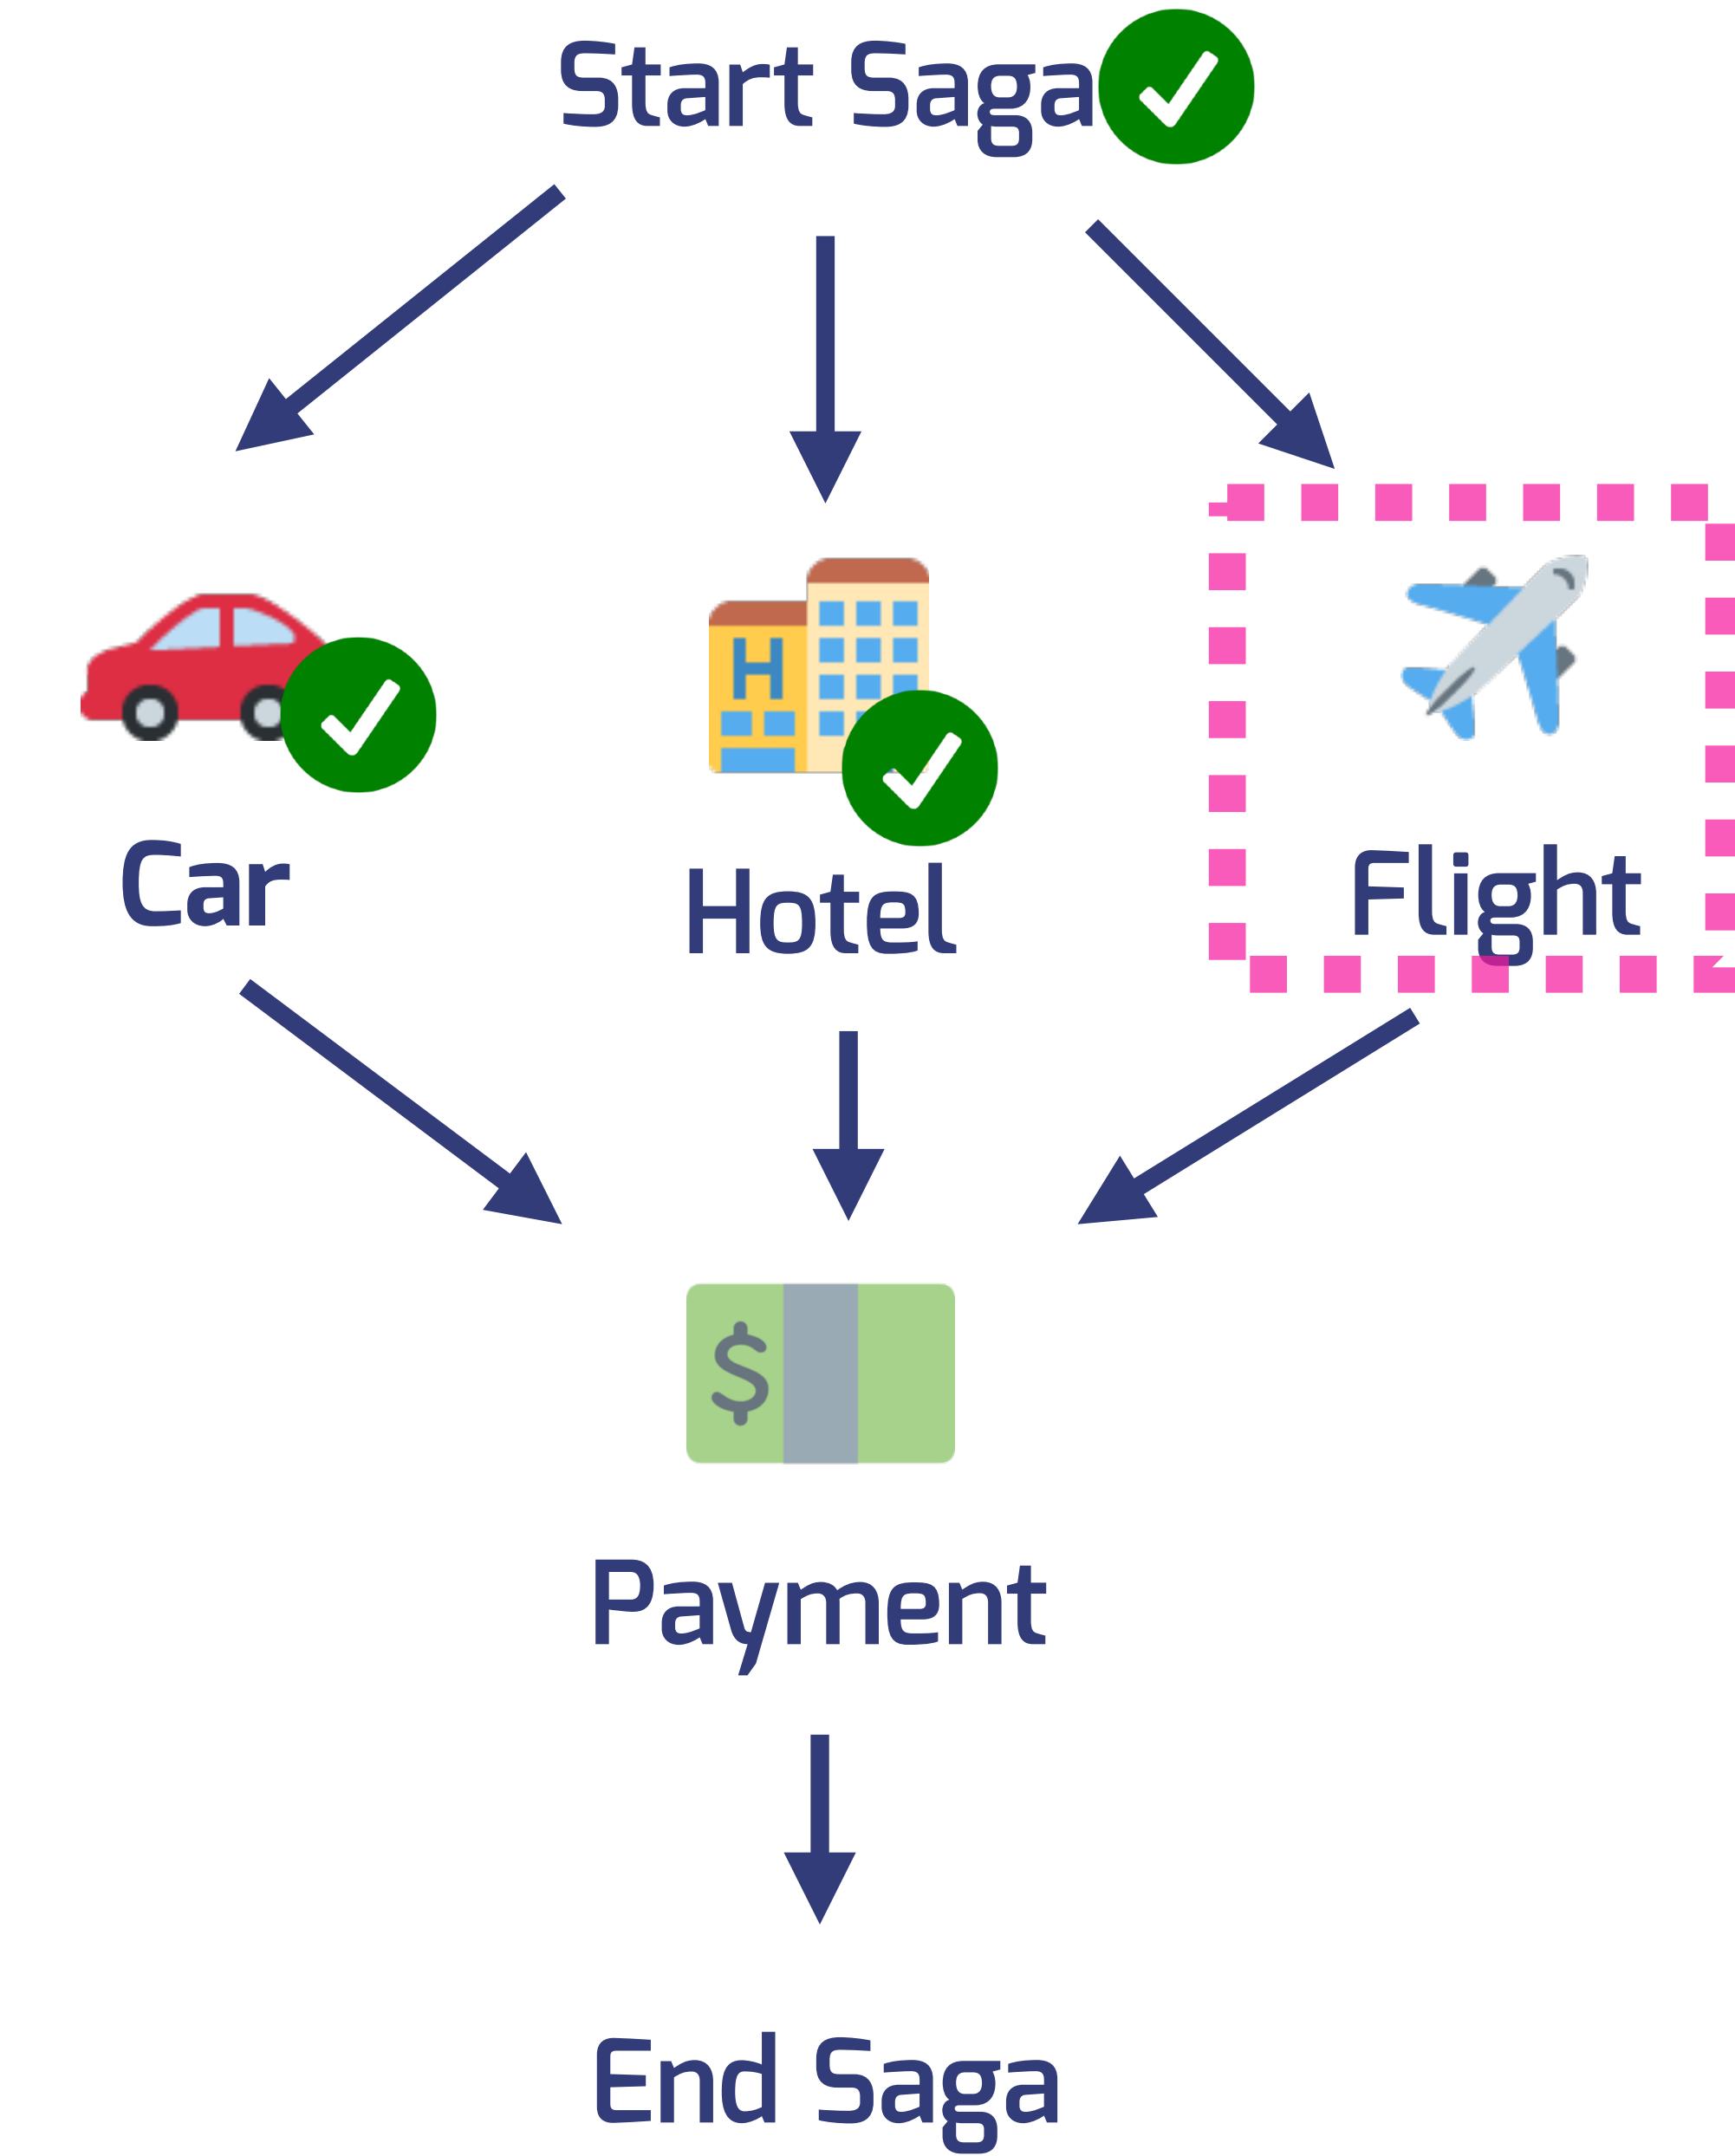
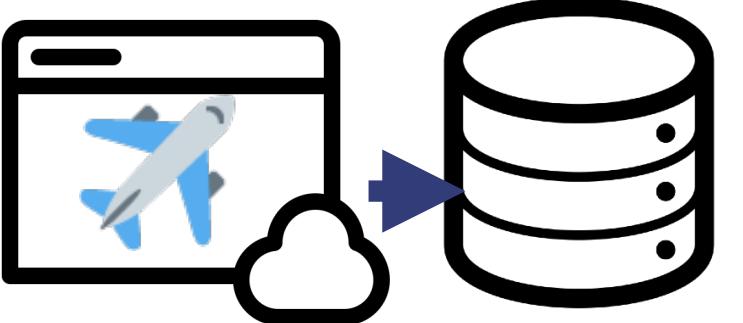
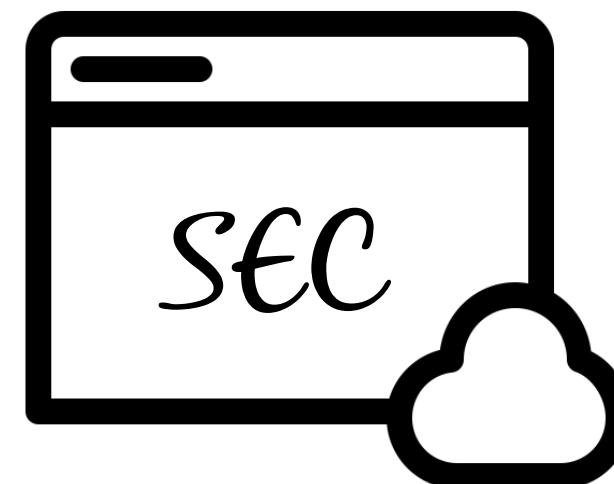
Done

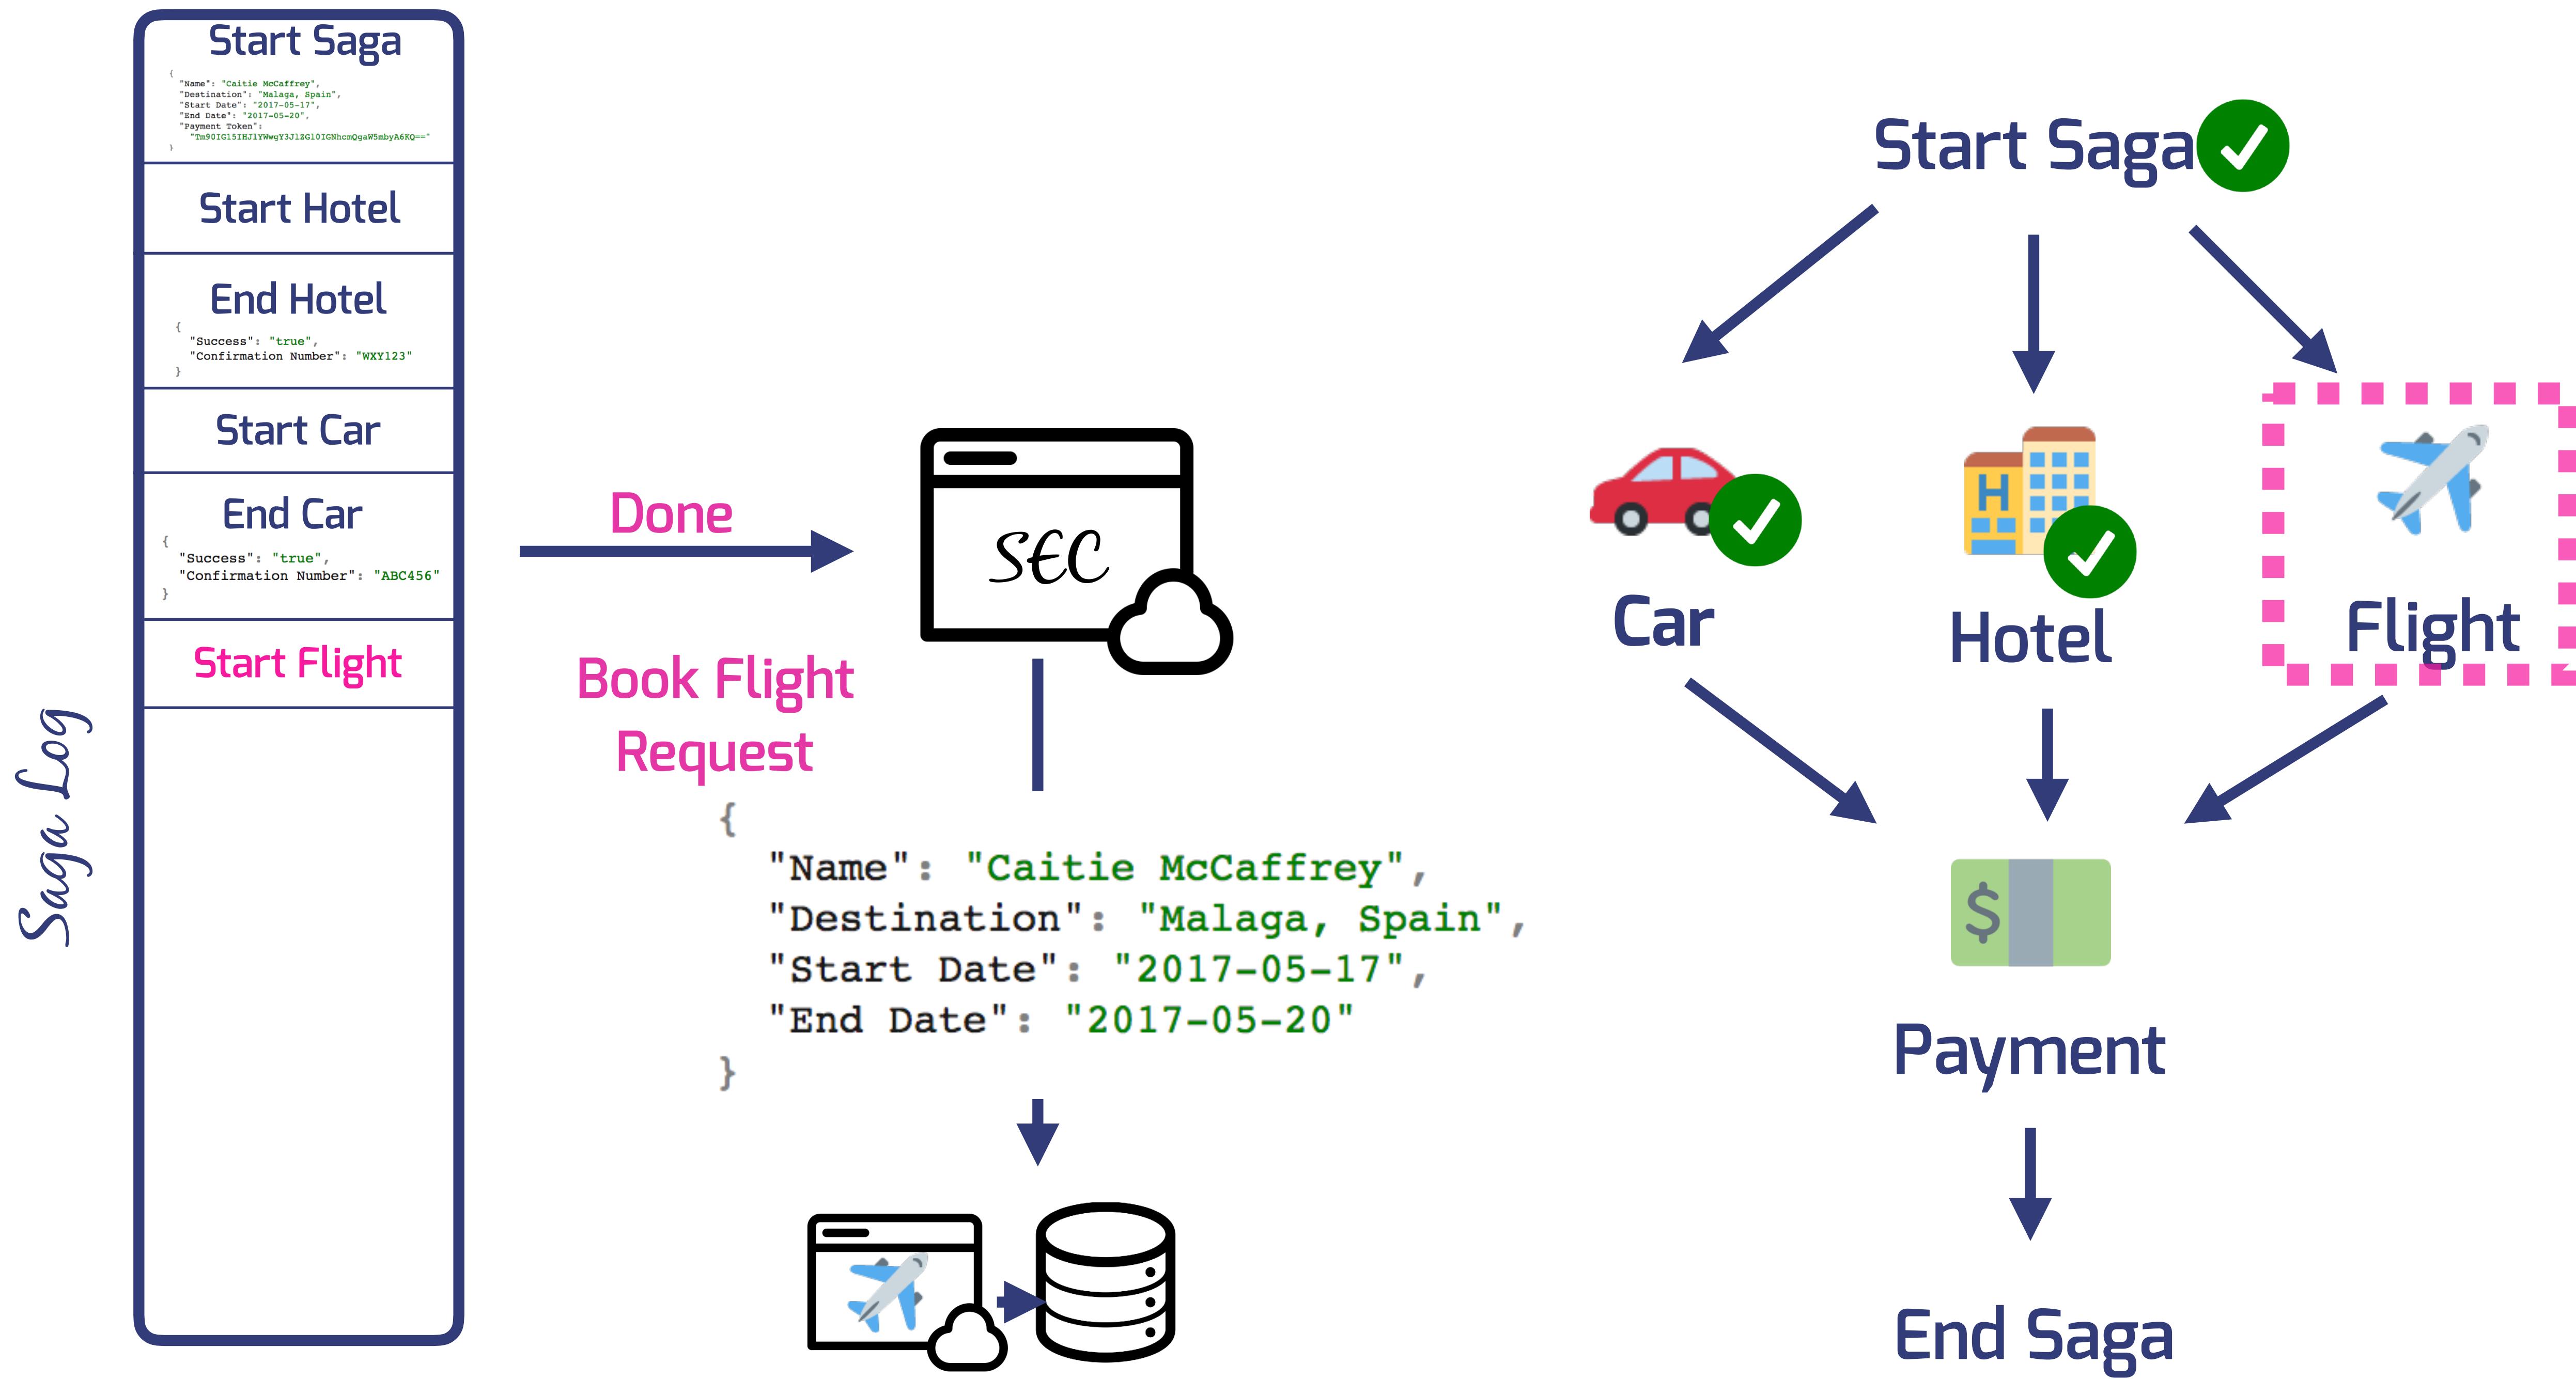


Saga Log

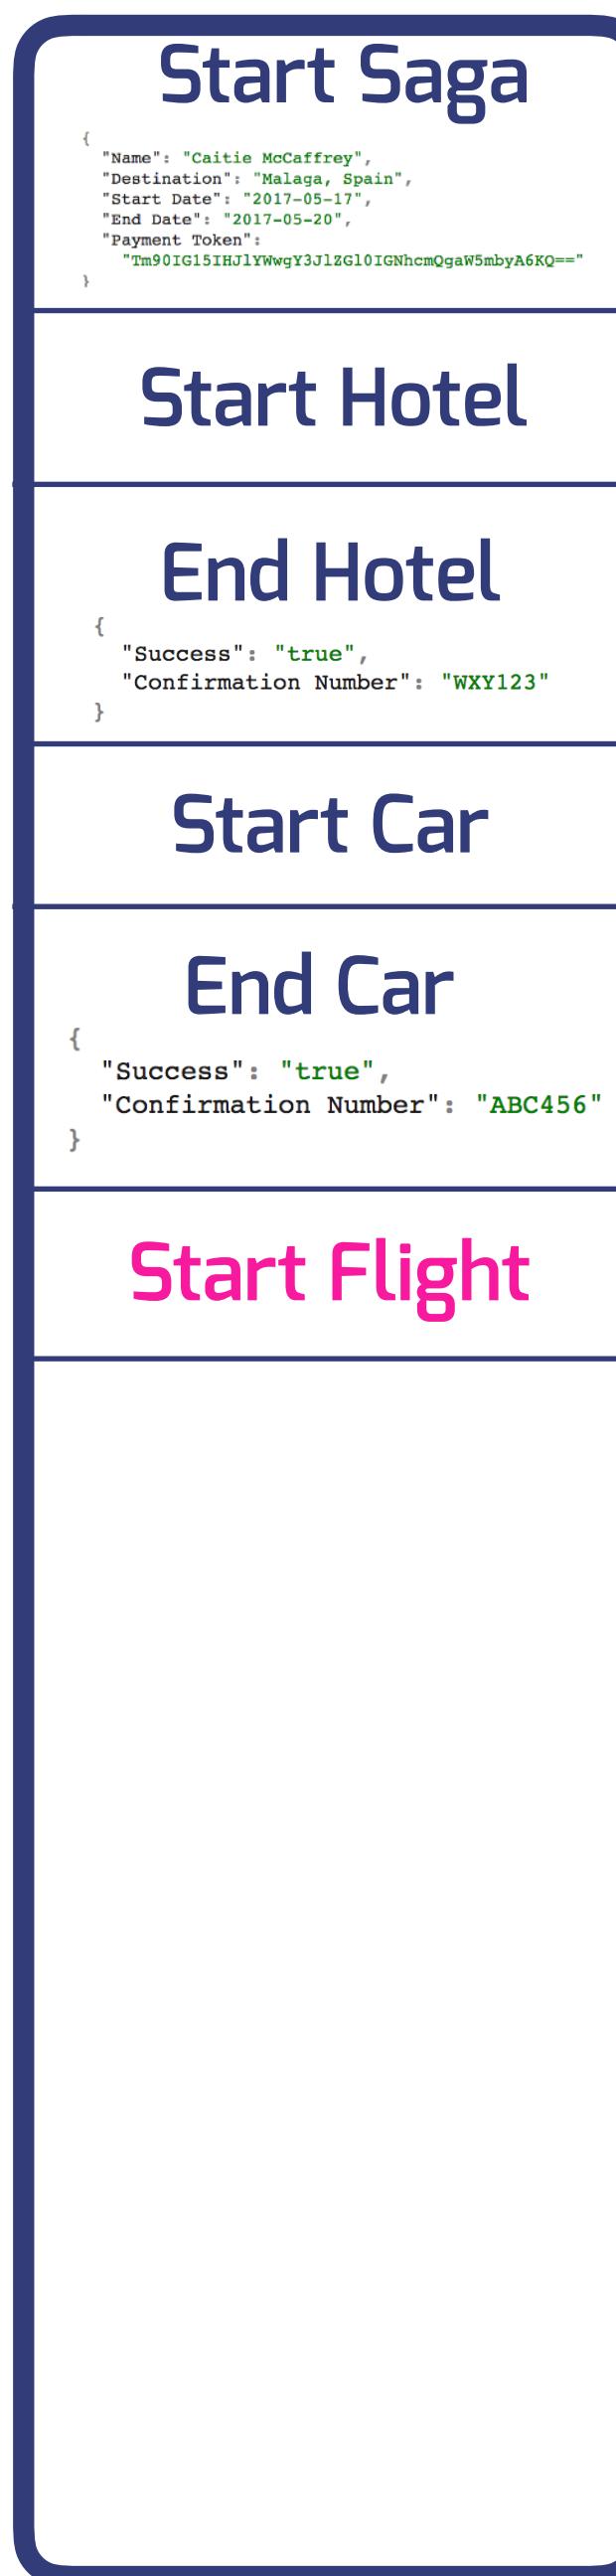


Start Flight





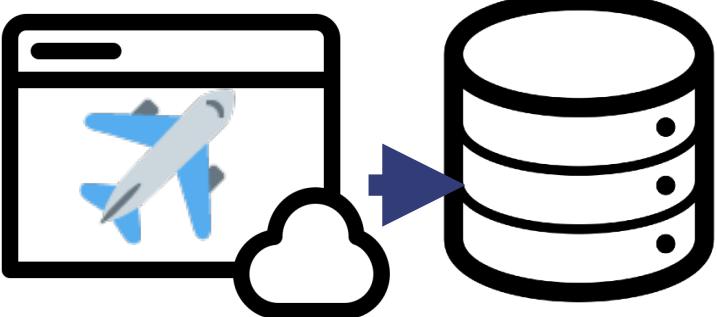
Saga Log



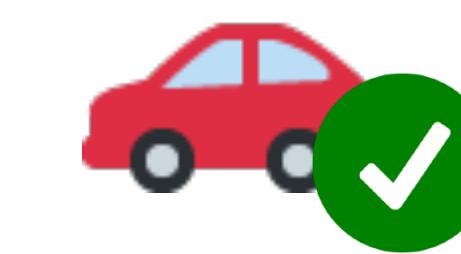
End Flight

Book Flight Response

```
{ "Success": "true", "Confirmation Number": "789QPZ" }
```



Start Saga



Car



Hotel



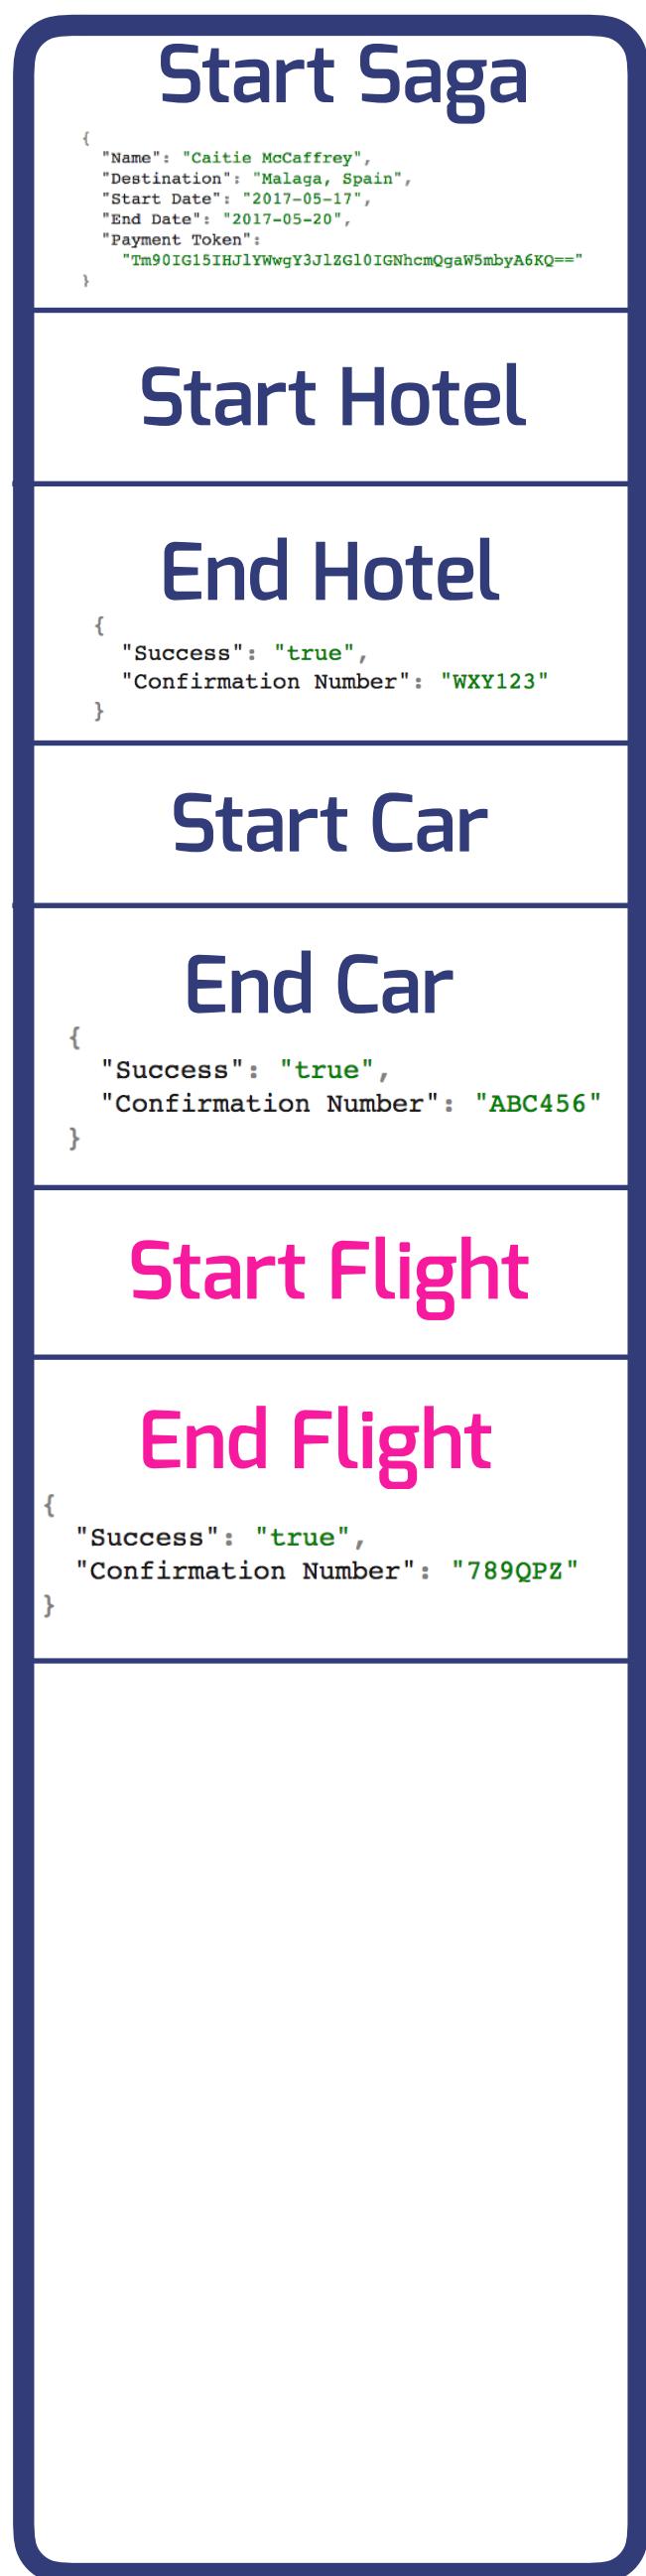
Flight

Payment

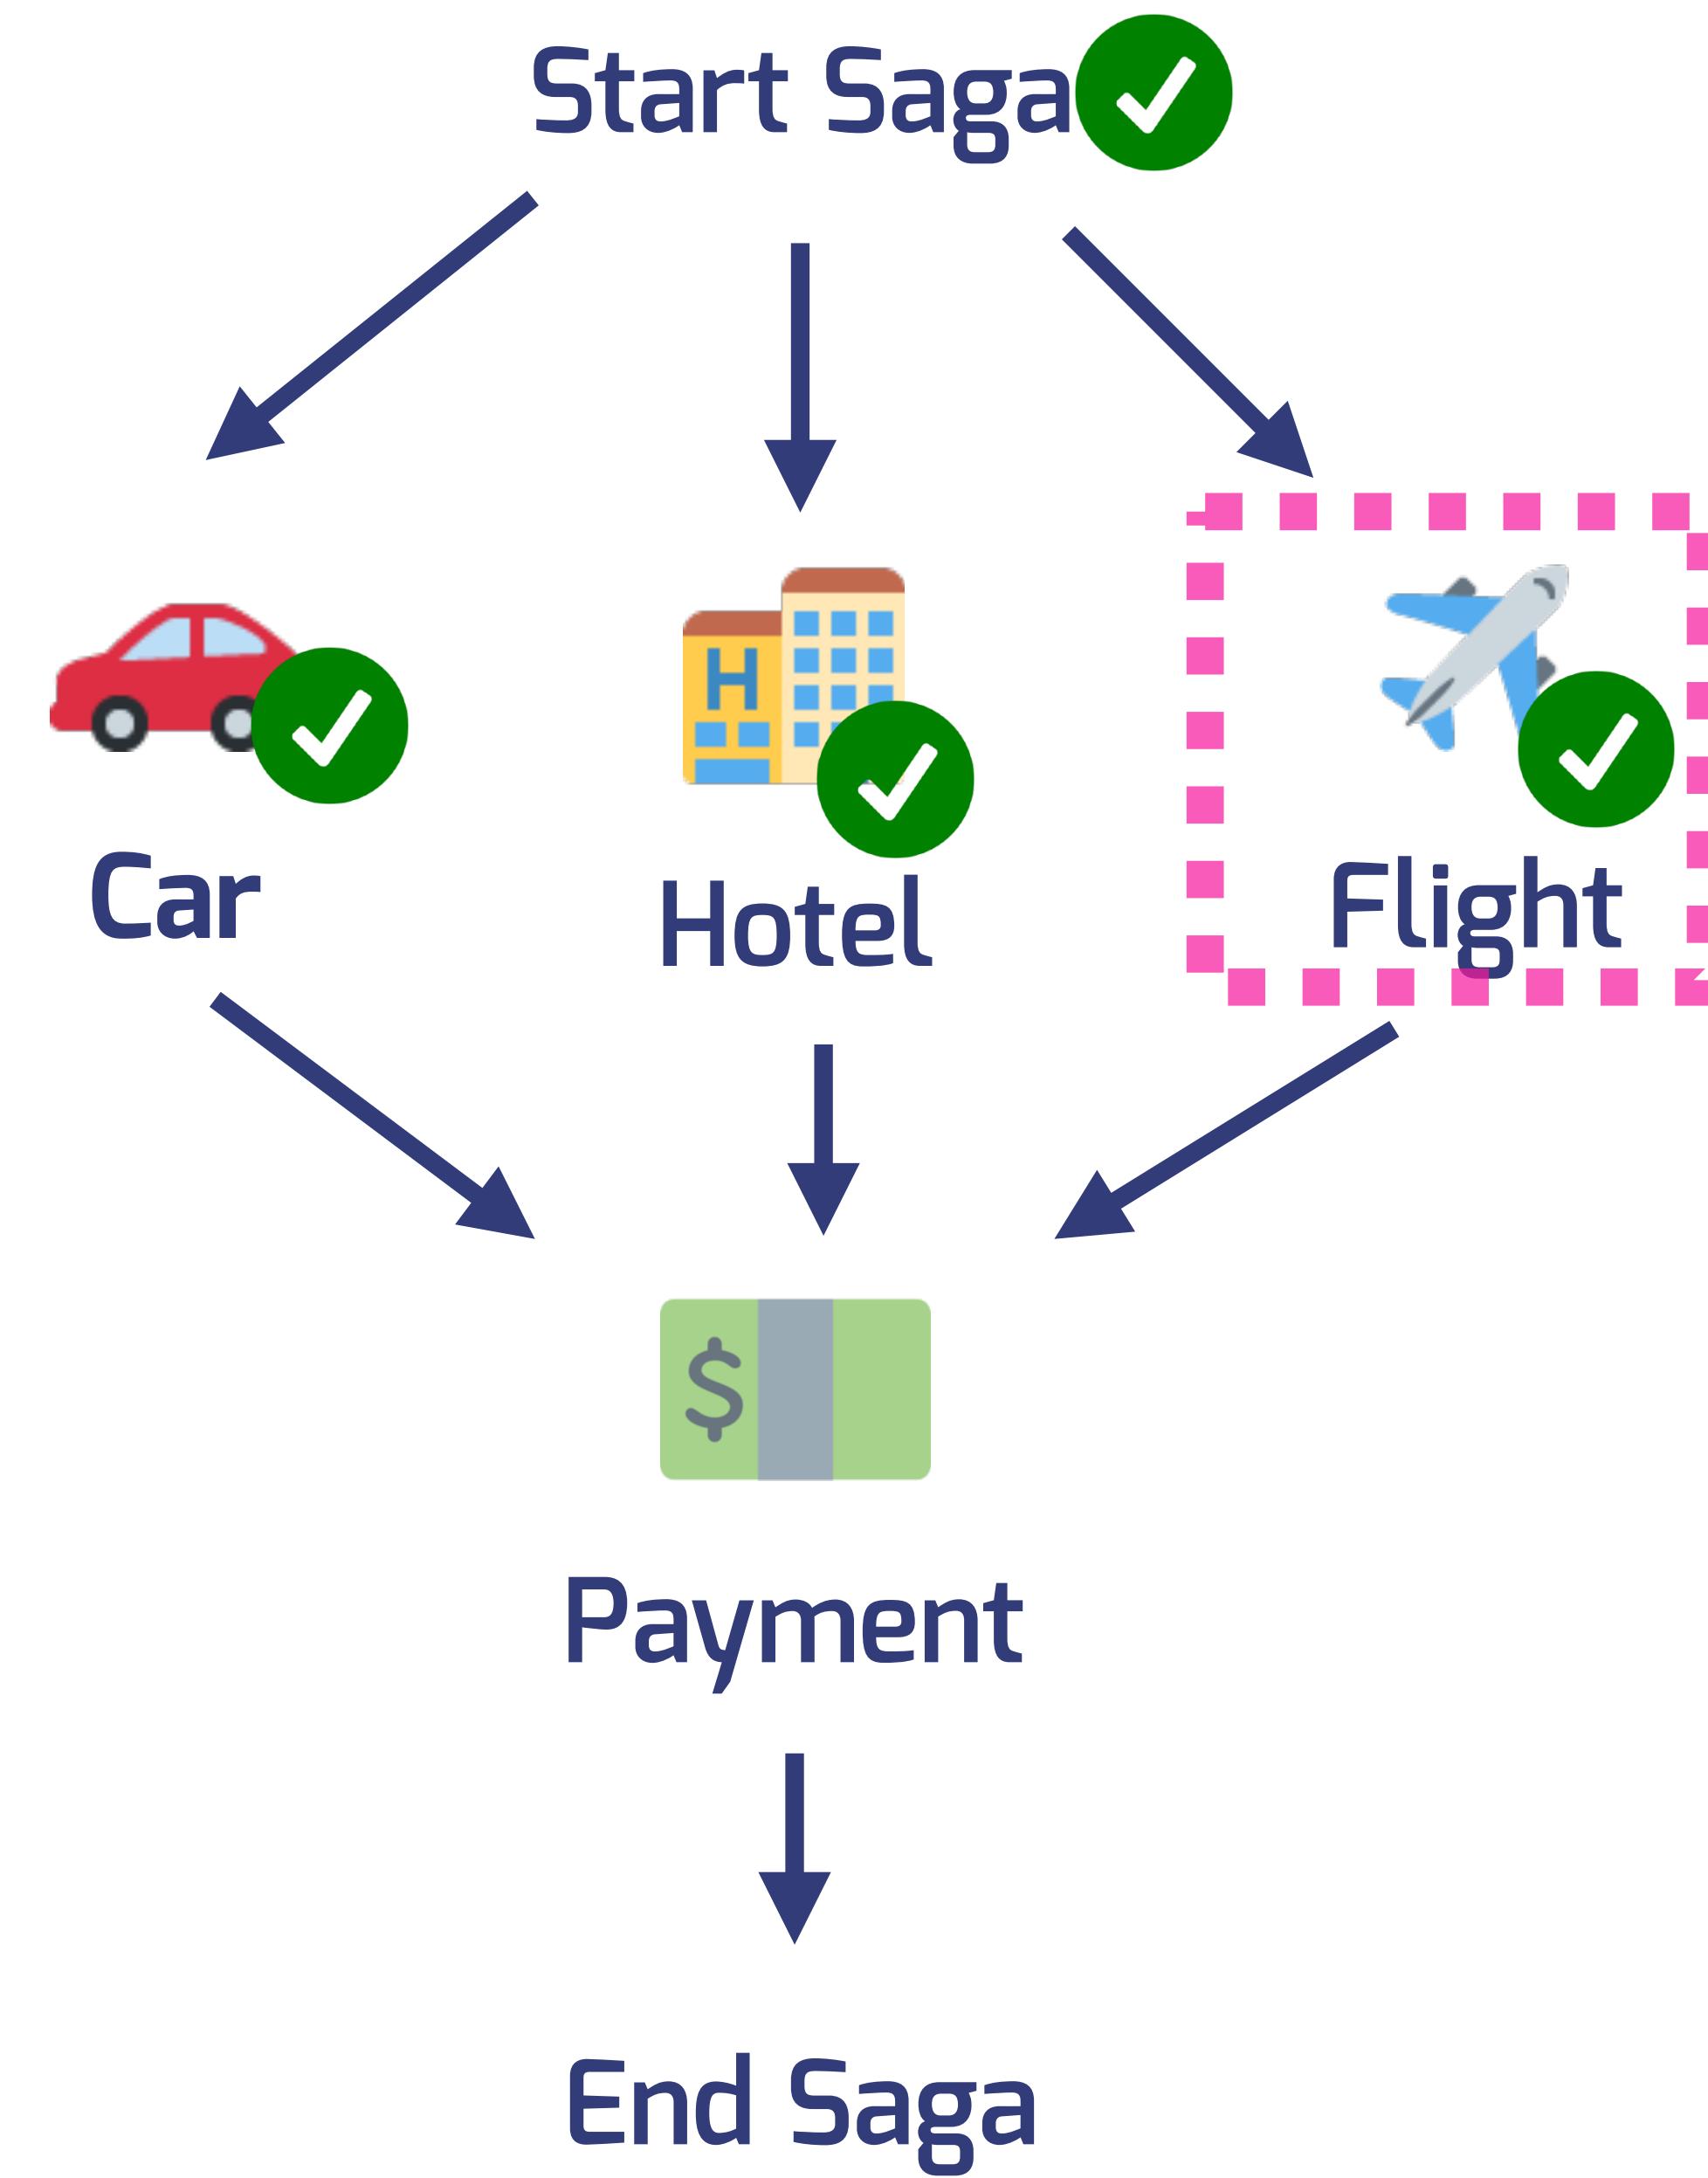
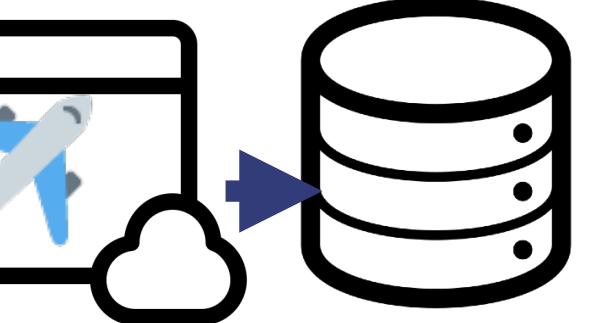


End Saga

Saga Log



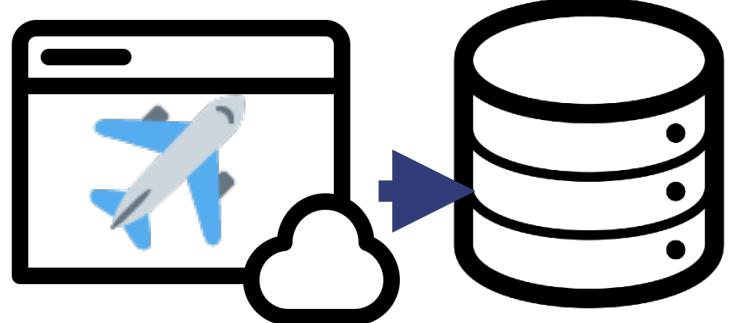
Done



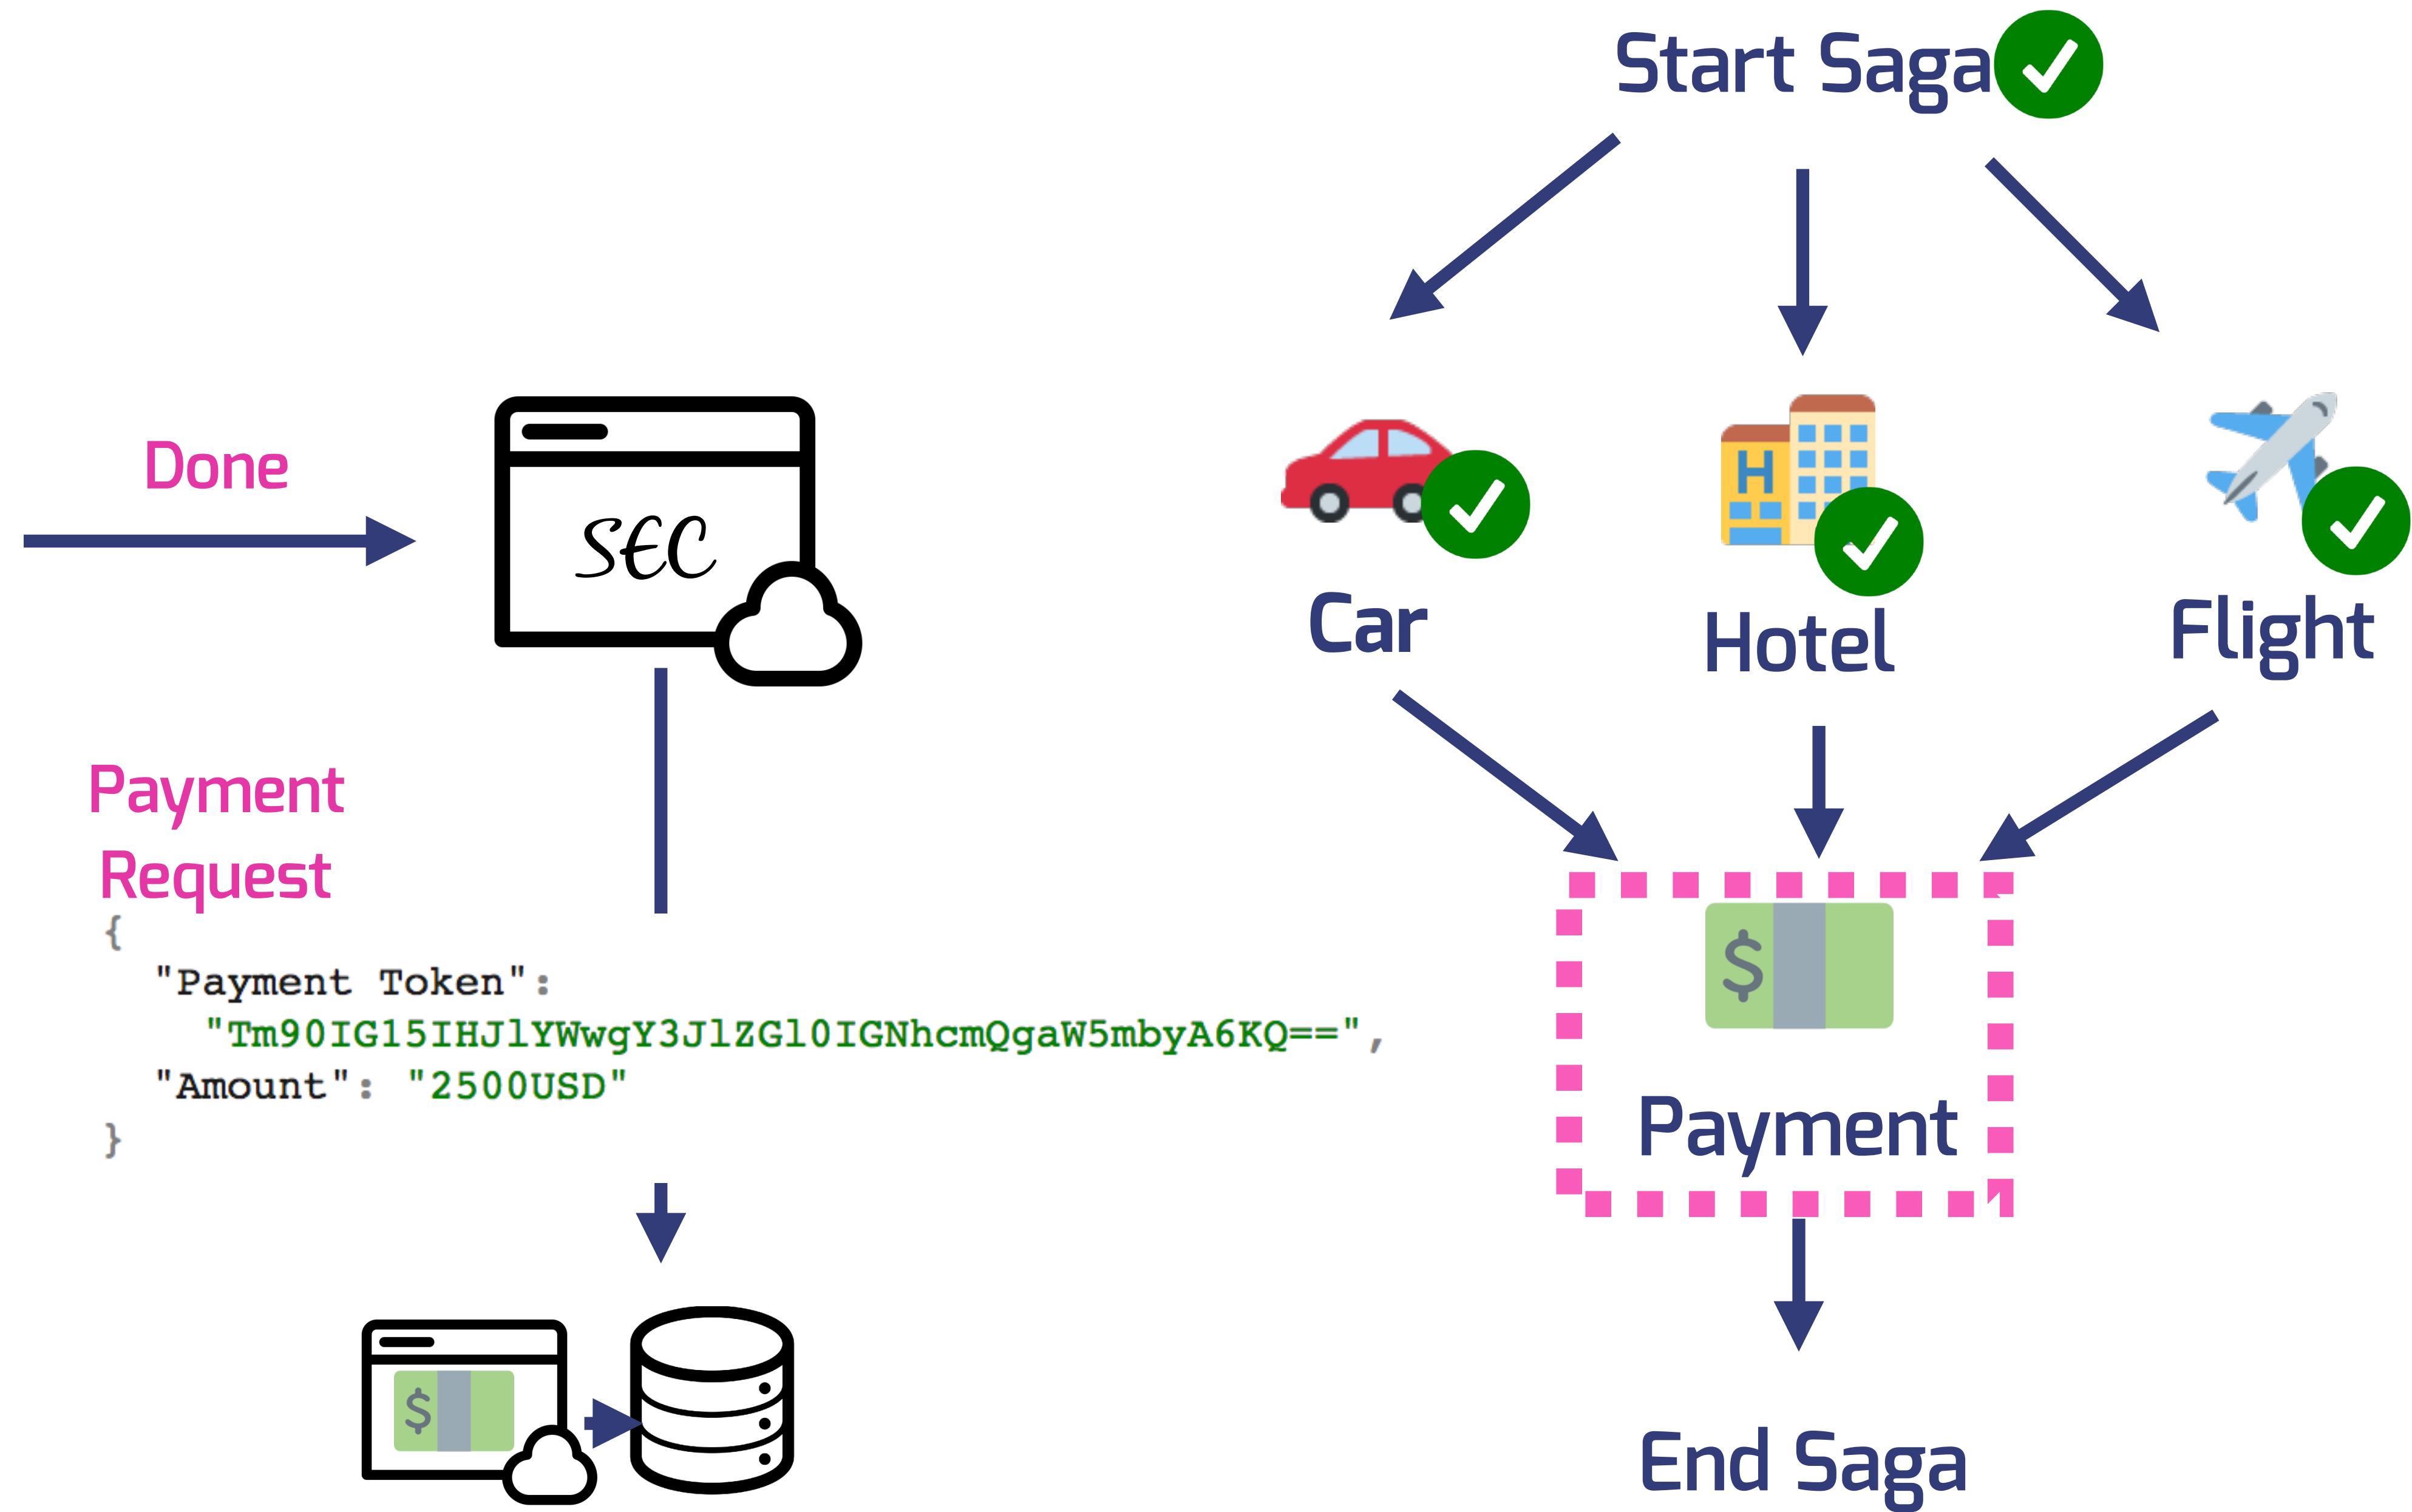
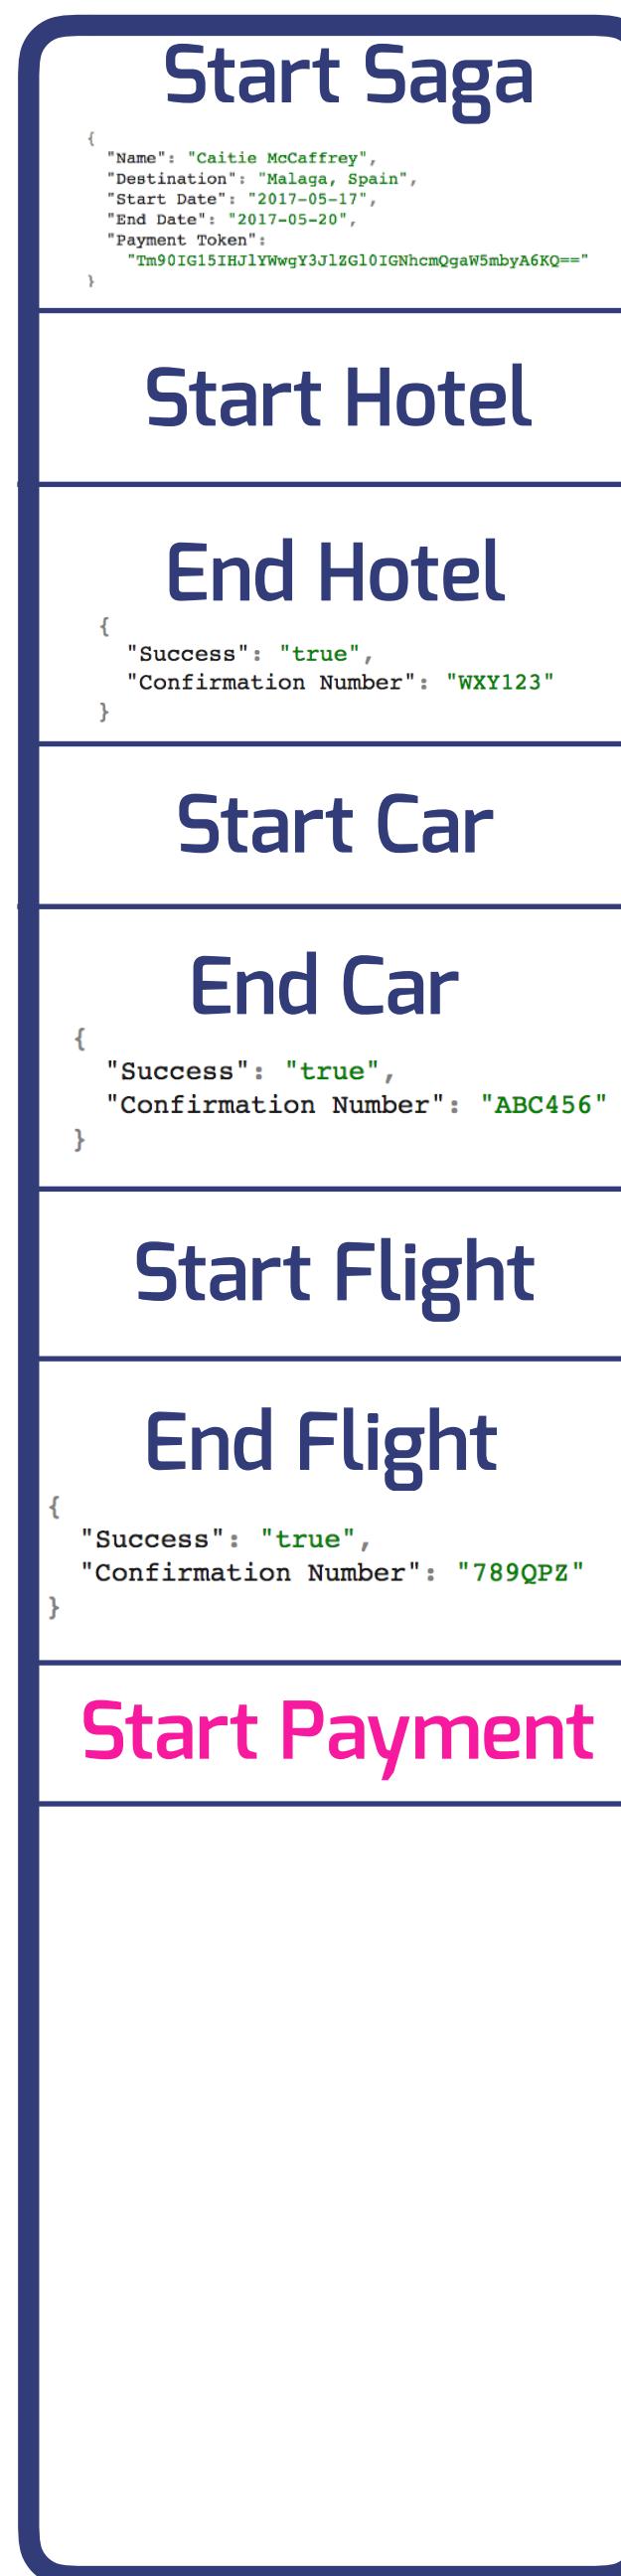
Saga Log

Start Saga
{ "Name": "Caitie McCaffrey", "Destination": "Malaga, Spain", "Start Date": "2017-05-17", "End Date": "2017-05-20", "Payment Token": "7m901G151HJ1YKwgy3J12G101GNhcmQgaW5mbA6KQ==" }
Start Hotel
End Hotel
{ "Success": "true", "Confirmation Number": "WXY123" }
Start Car
End Car
{ "Success": "true", "Confirmation Number": "ABC456" }
Start Flight
End Flight
{ "Success": "true", "Confirmation Number": "789QPZ" }

Start Payment



Saga Log



Saga Log

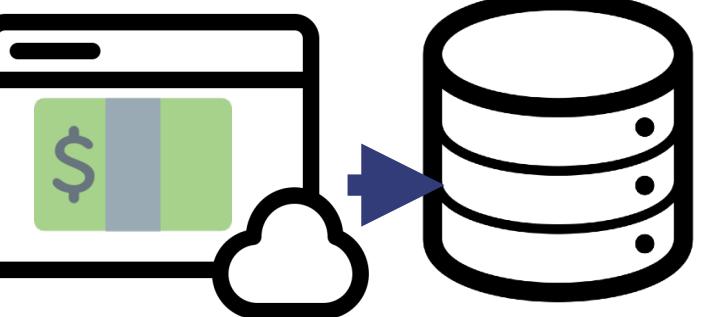
Start Saga
{ "Name": "Caitie McCaffrey", "Destination": "Malaga, Spain", "Start Date": "2017-05-17", "End Date": "2017-05-20", "Payment Token": "Um90IG15IHJ1YWwgY3J1ZG10IGNhcmQgaW5mbvA6KQ==" }
Start Hotel
End Hotel
{ "Success": "true", "Confirmation Number": "WXY123" }
Start Car
End Car
{ "Success": "true", "Confirmation Number": "ABC456" }
Start Flight
End Flight
{ "Success": "true", "Confirmation Number": "789QPZ" }
Start Payment

End Payment



Payment Response

```
{  
  "success": true,  
  "Invoice Number": 12345  
}
```



Start Saga 



Car



Hotel



Flight



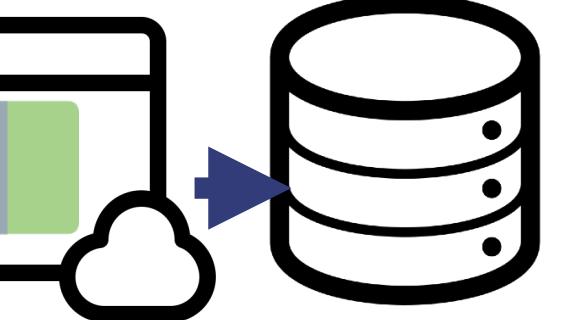
Payment

End Saga

Saga Log

Start Saga
{ "Name": "Caitie McCaffrey", "Destination": "Malaga, Spain", "Start Date": "2017-05-17", "End Date": "2017-05-20", "Payment Token": "Um90IG15IHJ1YWwgY3J1ZG10IGNhcmQgaW5mbjA6KQ==" }
Start Hotel
End Hotel
{ "Success": "true", "Confirmation Number": "WXY123" }
Start Car
End Car
{ "Success": "true", "Confirmation Number": "ABC456" }
Start Flight
End Flight
{ "Success": "true", "Confirmation Number": "789QPZ" }
Start Payment
End Payment
{ "success": true, "Invoice Number": 12345 }

Done



Saga Log

Start Saga
{ "Name": "Caitie McCaffrey", "Destination": "Malaga, Spain", "Start Date": "2017-05-17", "End Date": "2017-05-20", "Payment Token": "Um90IG15IHJ1YWwgY3J1ZG10IGNhcmQgaW5mbvA6KQ==" }
Start Hotel
End Hotel
{ "Success": "true", "Confirmation Number": "WXY123" }
Start Car
End Car
{ "Success": "true", "Confirmation Number": "ABC456" }
Start Flight
End Flight
{ "Success": "true", "Confirmation Number": "789QPZ" }
Start Payment
End Payment
{ "success": true, "Invoice Number": 12345 }

End Saga



Start Saga



Car



Hotel



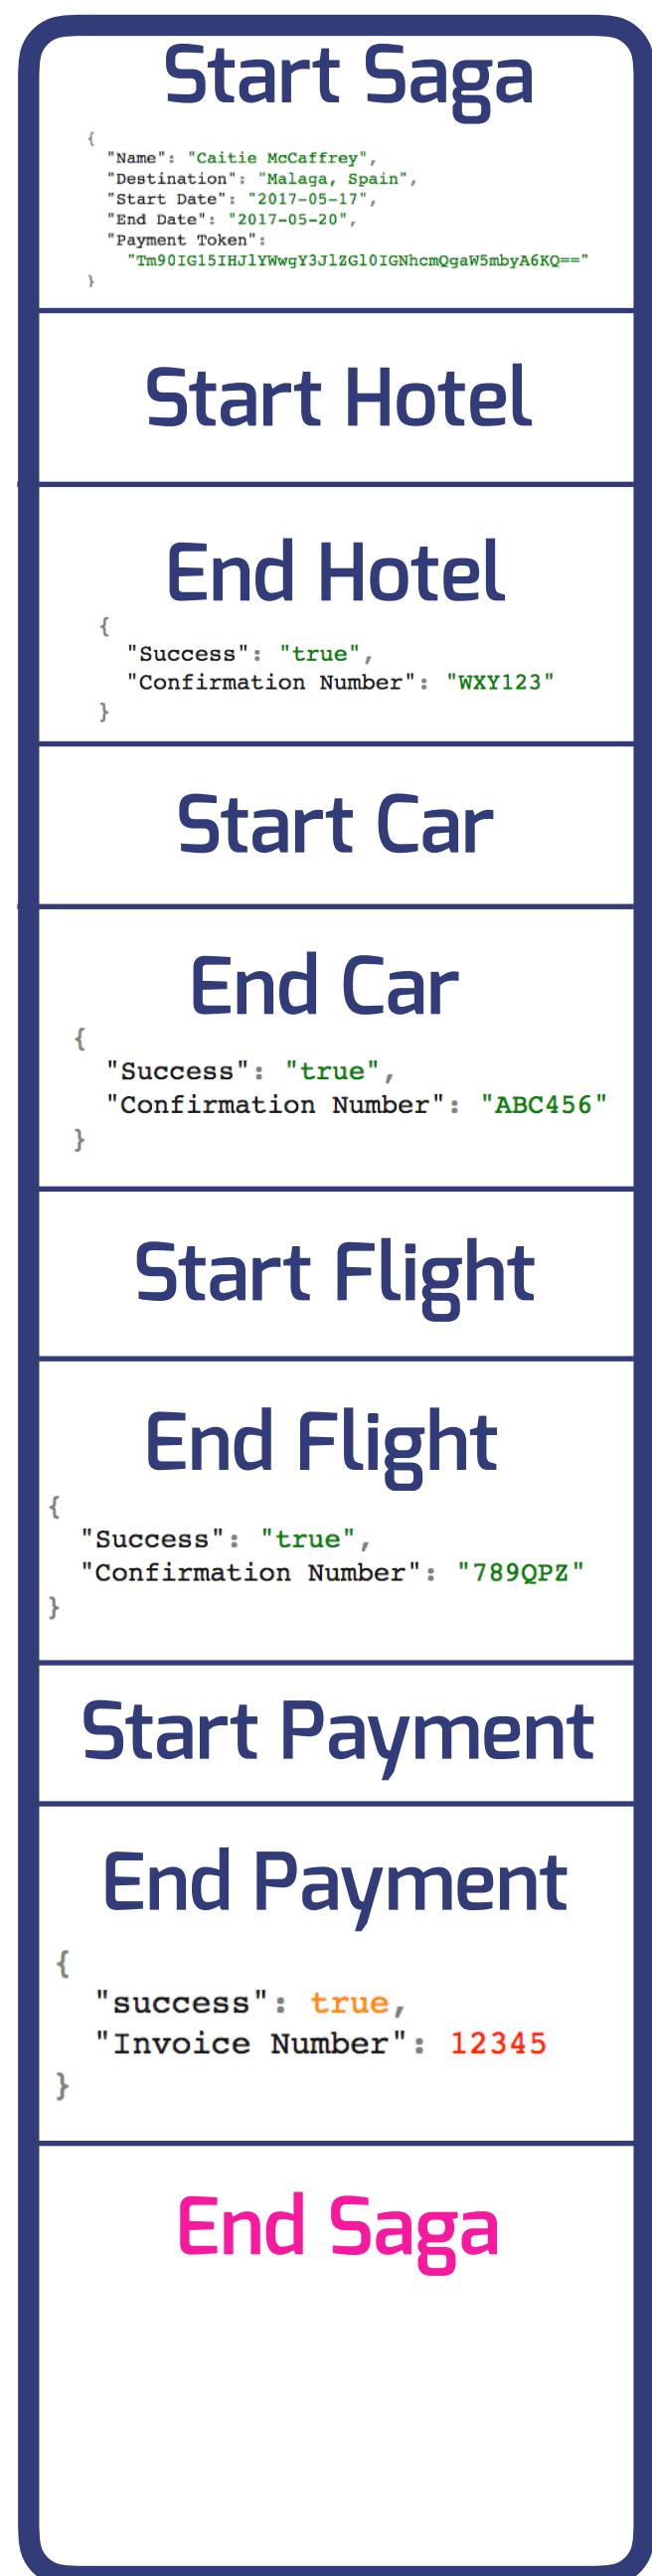
Flight



Payment

End Saga

Saga Log



Done



Saga Log

Start Saga
{ "Name": "Caitie McCaffrey", "Destination": "Malaga, Spain", "Start Date": "2017-05-17", "End Date": "2017-05-20", "Payment Token": "Um90IG15IHJ1YWwgY3J1ZG10IGNhcmQgaW5mbA6KQ==" }
Start Hotel
End Hotel
{ "Success": "true", "Confirmation Number": "WXY123" }
Start Car
End Car
{ "Success": "true", "Confirmation Number": "ABC456" }
Start Flight
End Flight
{ "Success": "true", "Confirmation Number": "789QPZ" }
Start Payment
End Payment
{ "success": true, "Invoice Number": 12345 }
End Saga



Start Saga



Car



Hotel



Flight

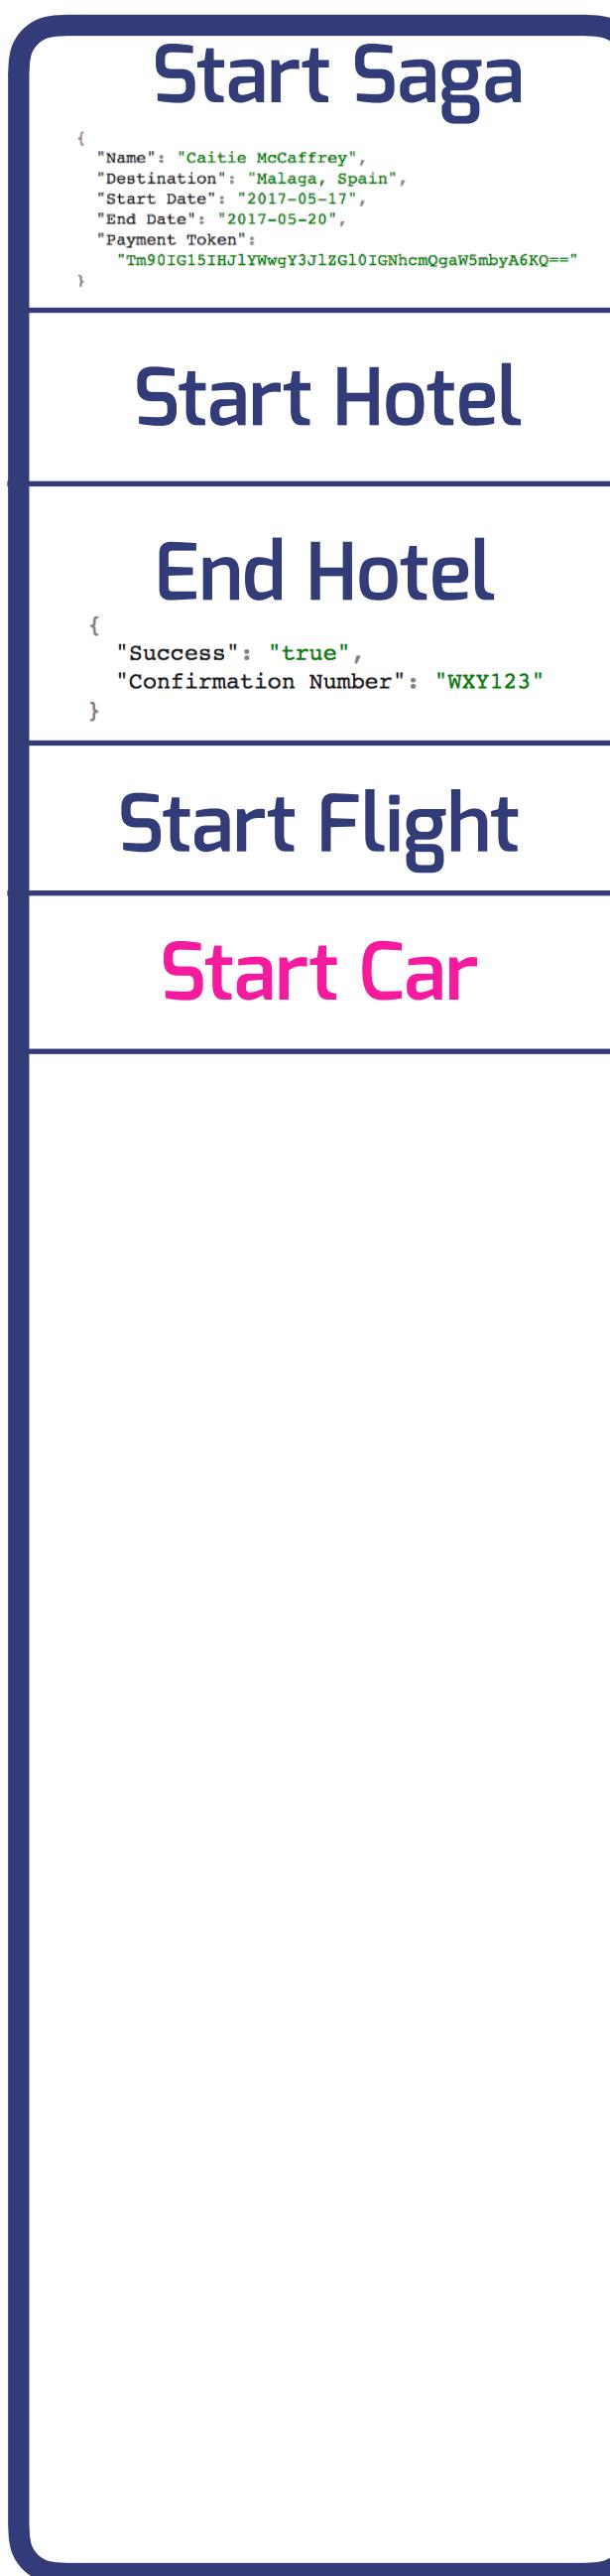


Payment

End Saga

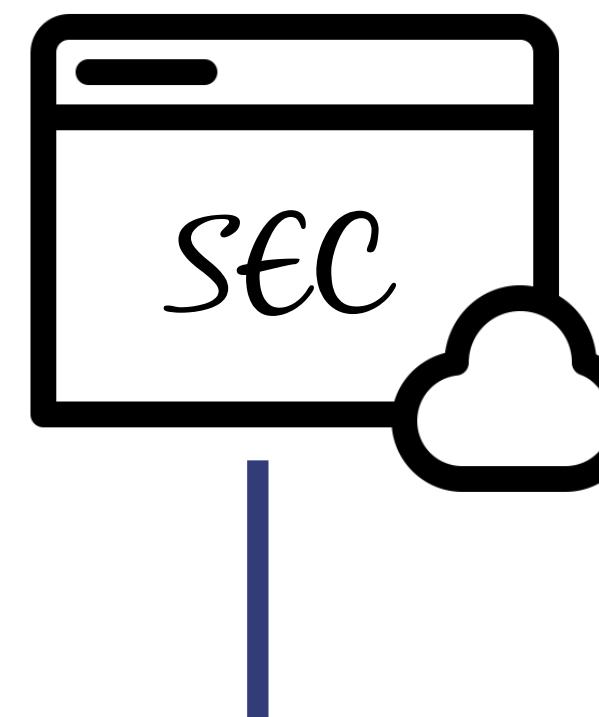
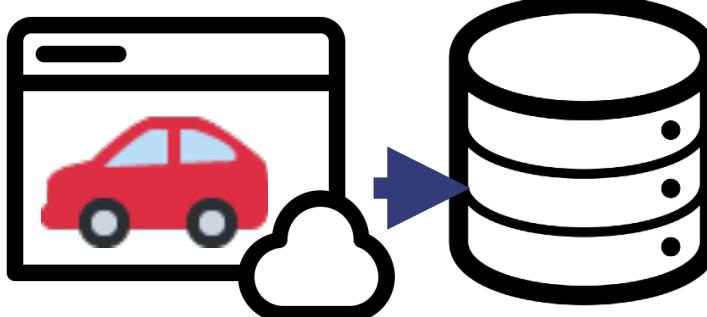
Failure of a
Distributed Saga

Saga Log



Book Car Request

```
{  
  "Name": "Caitie McCaffrey",  
  "Destination": "Malaga, Spain",  
  "Start Date": "2017-05-17",  
  "End Date": "2017-05-20"  
}
```



Start Saga



Hotel



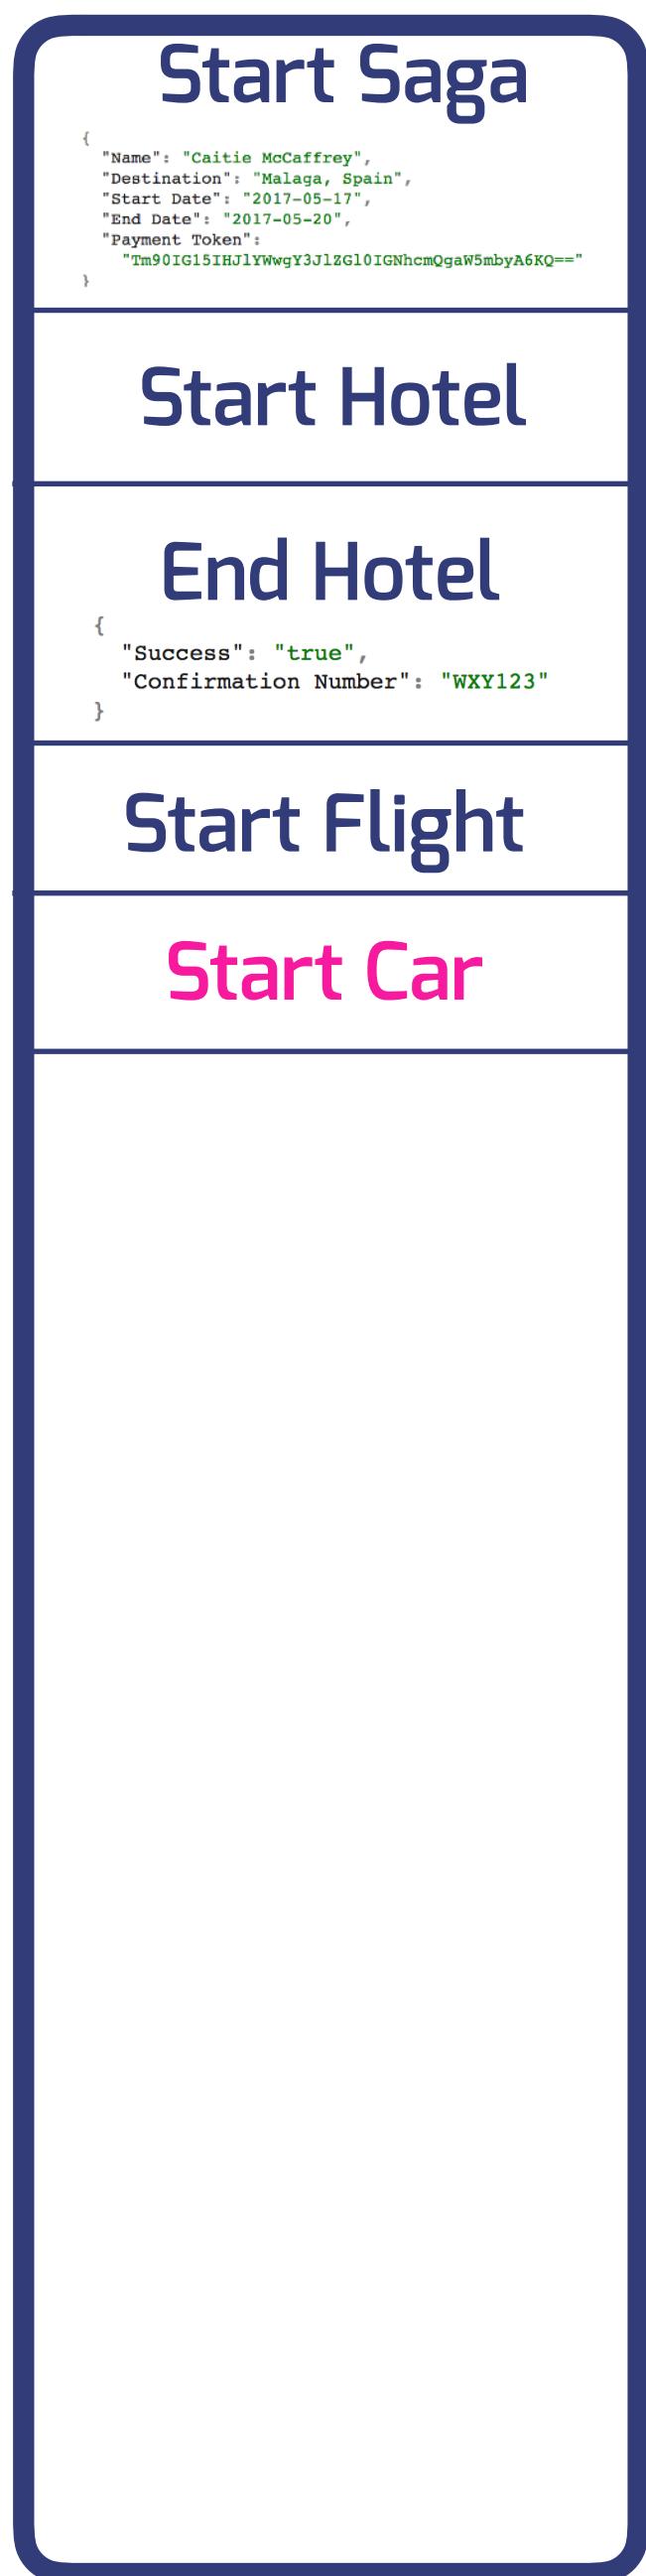
Flight



Payment

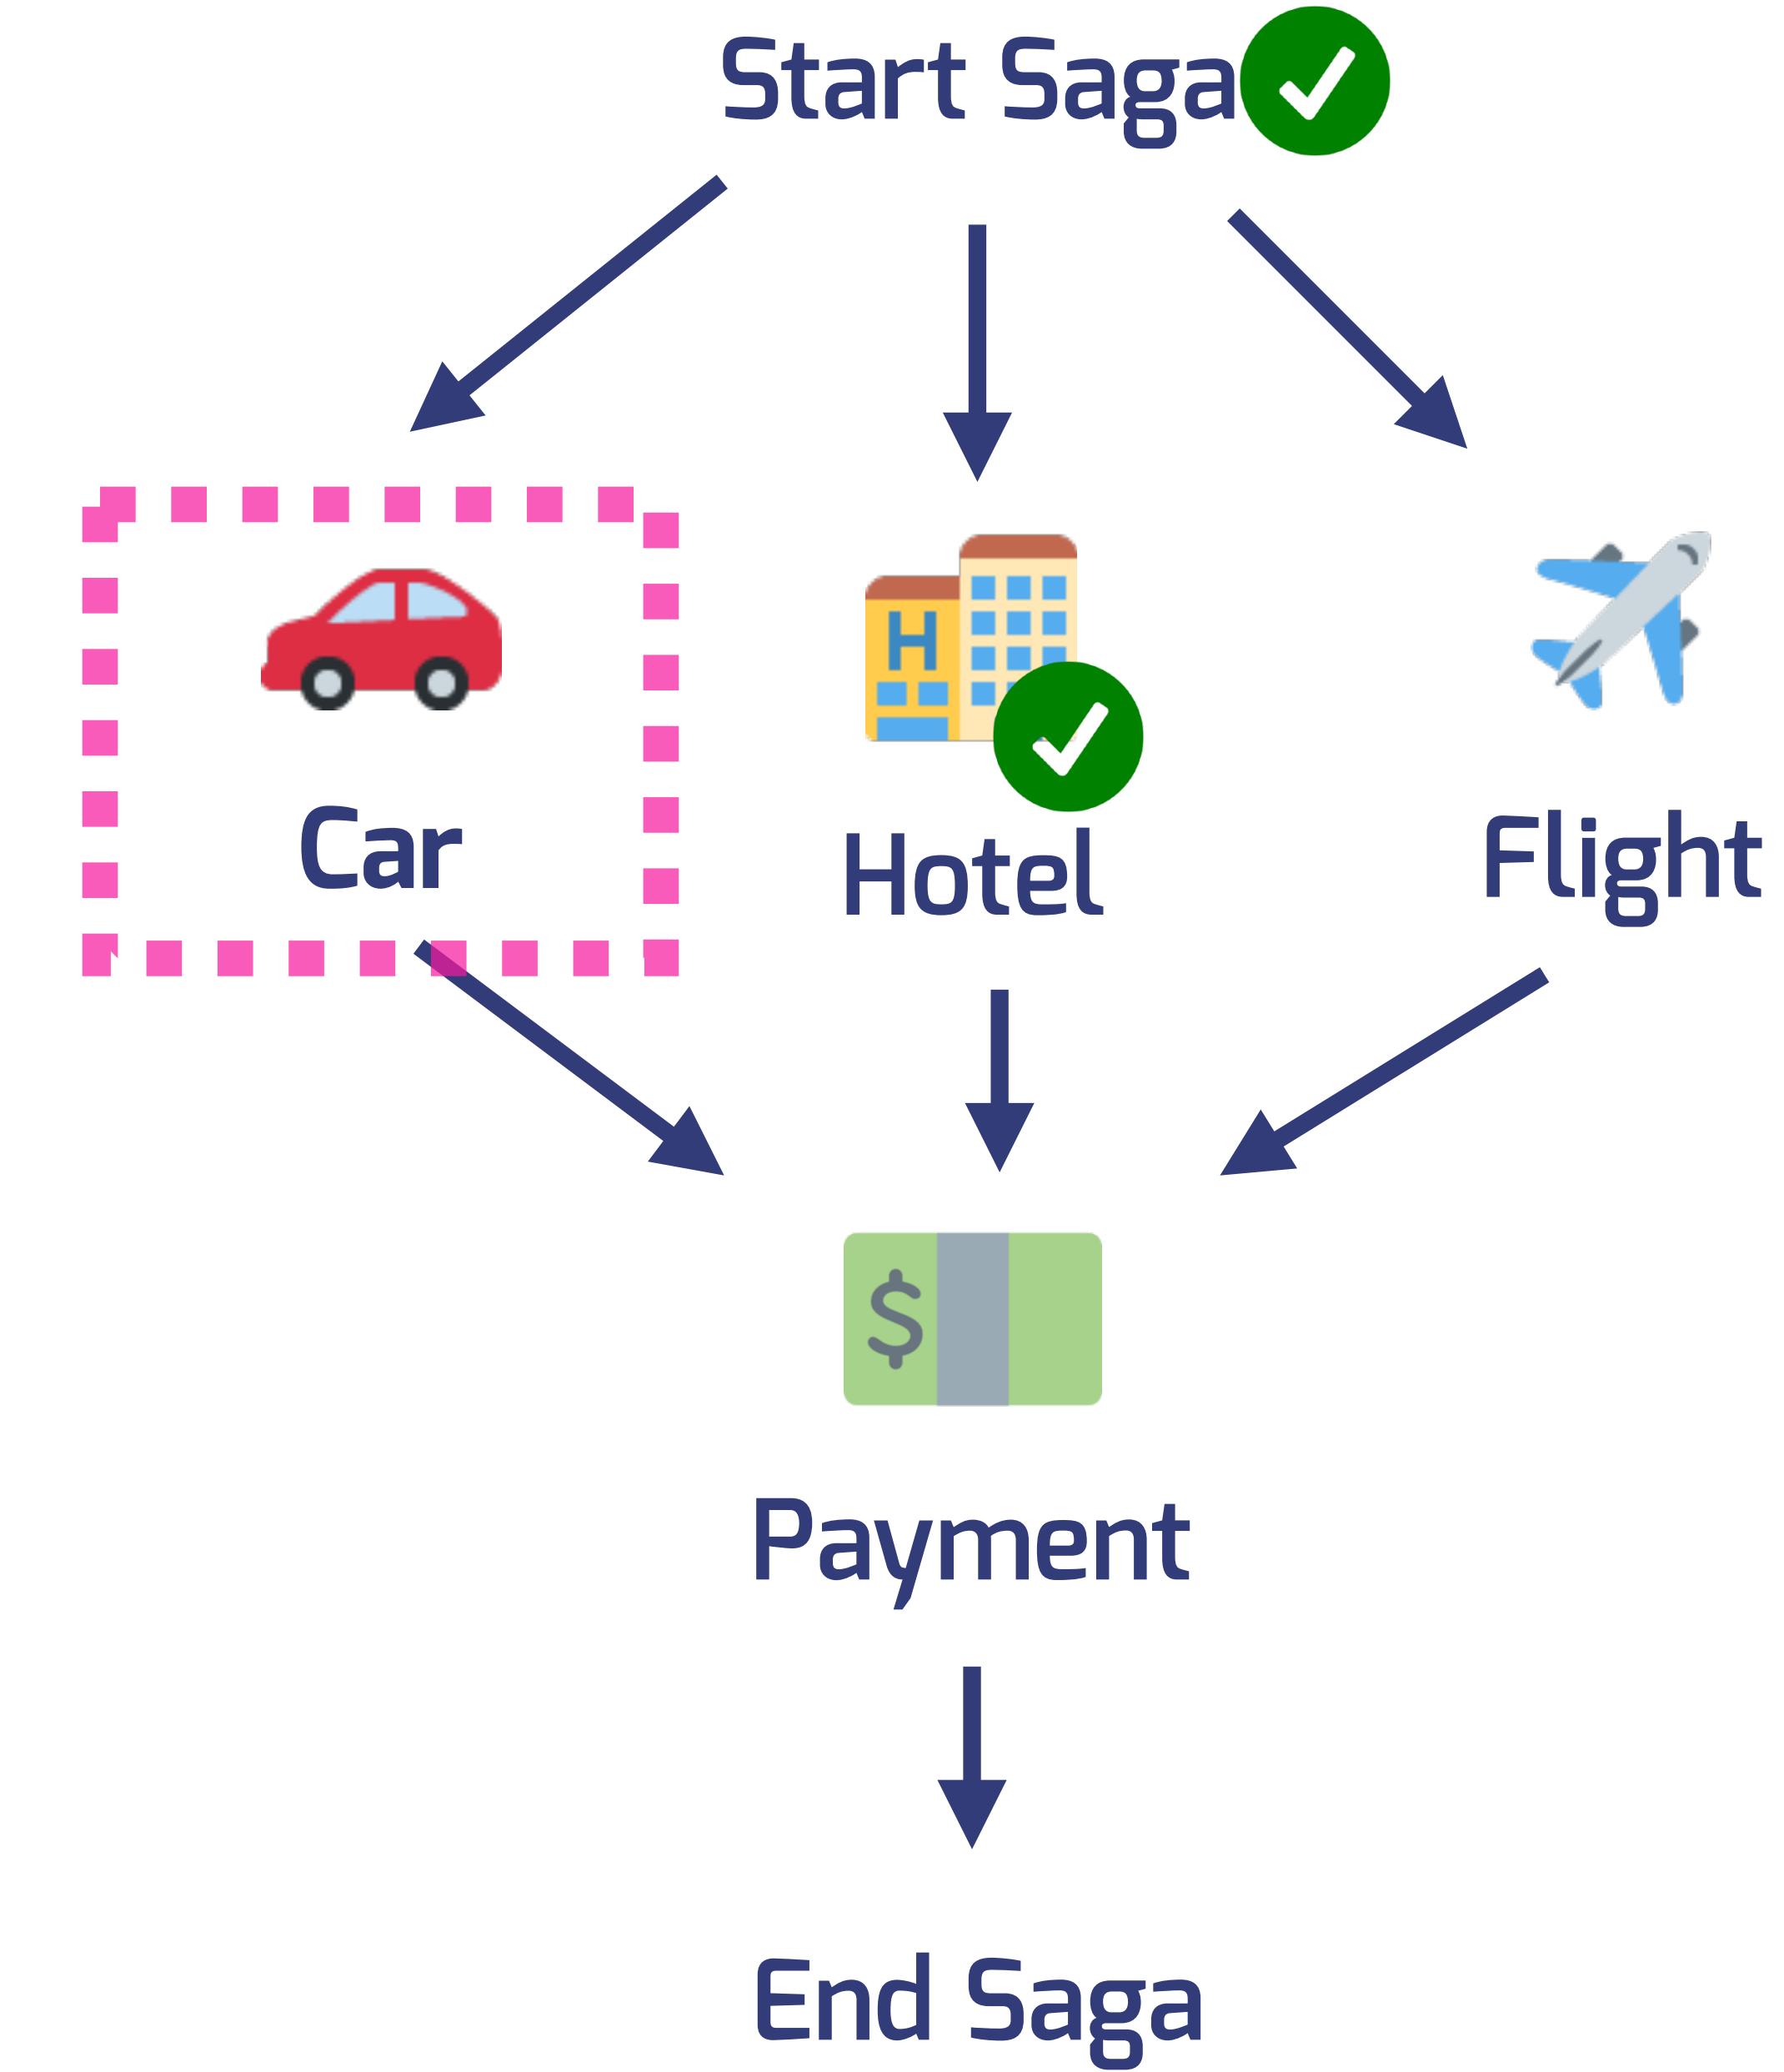
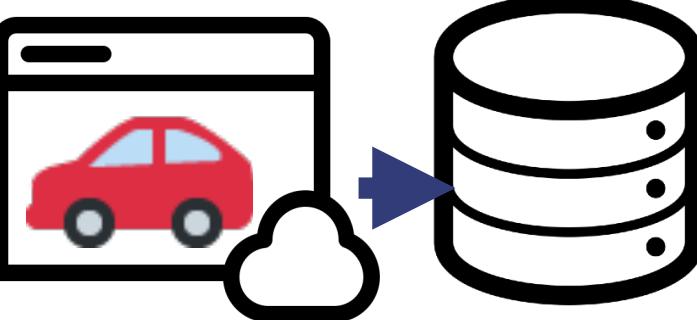
End Saga

Saga Log

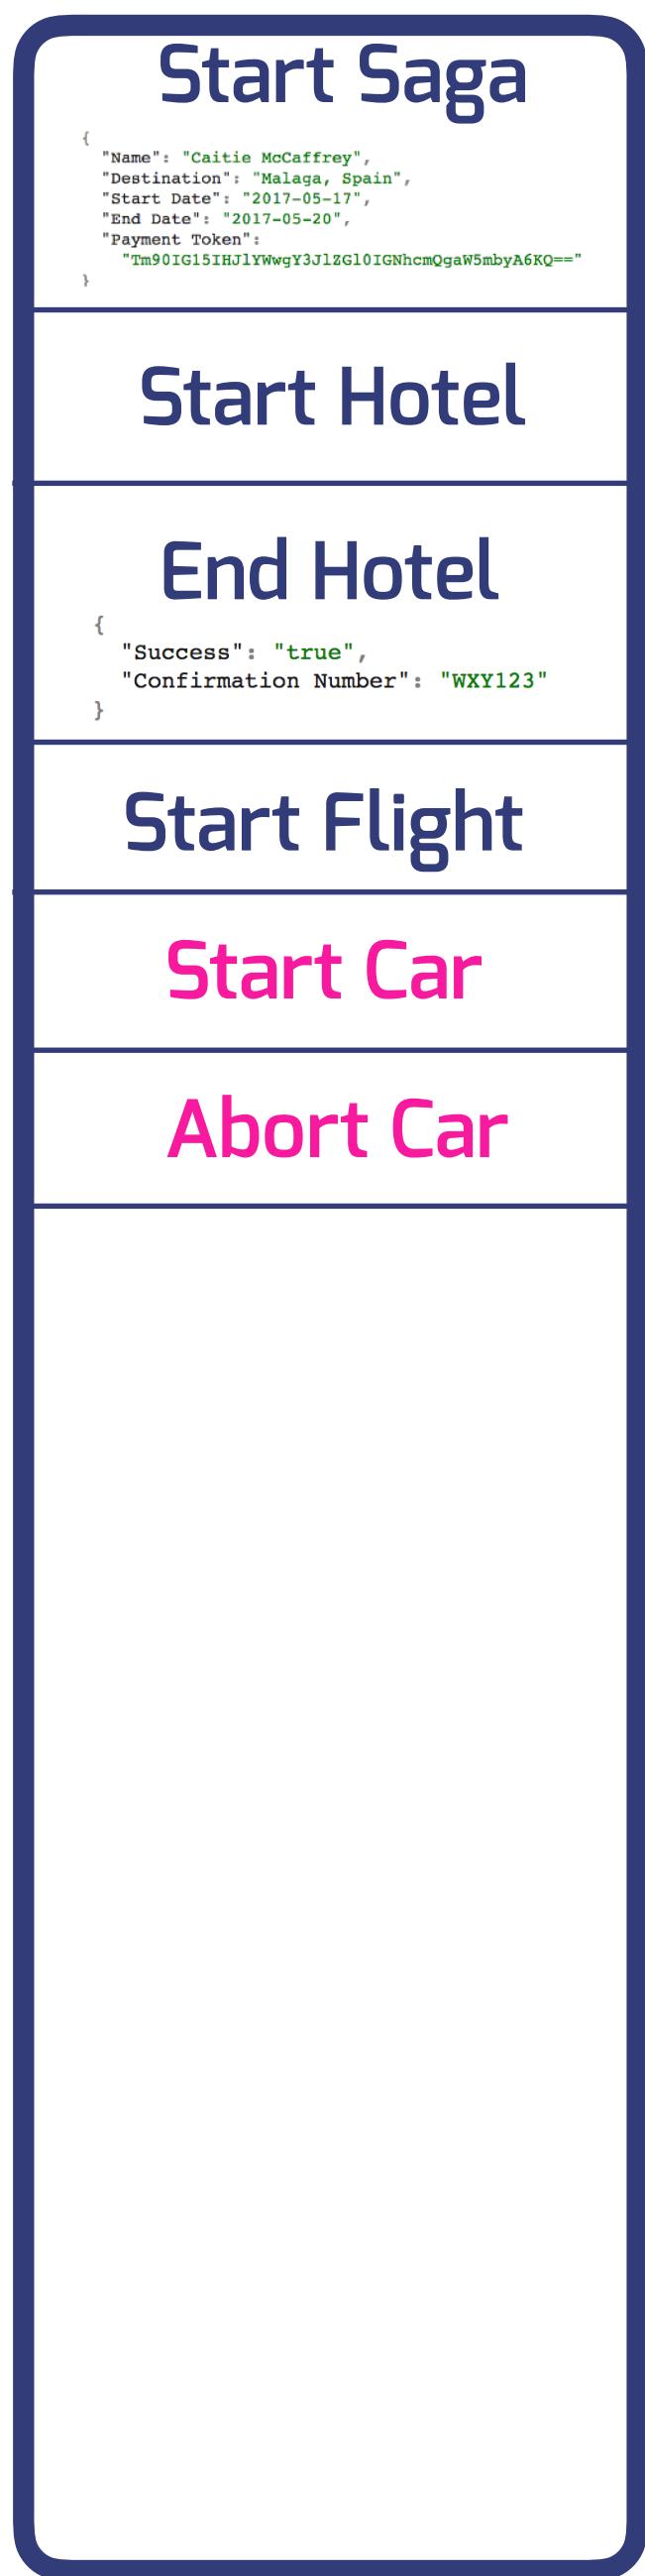


Abort Car

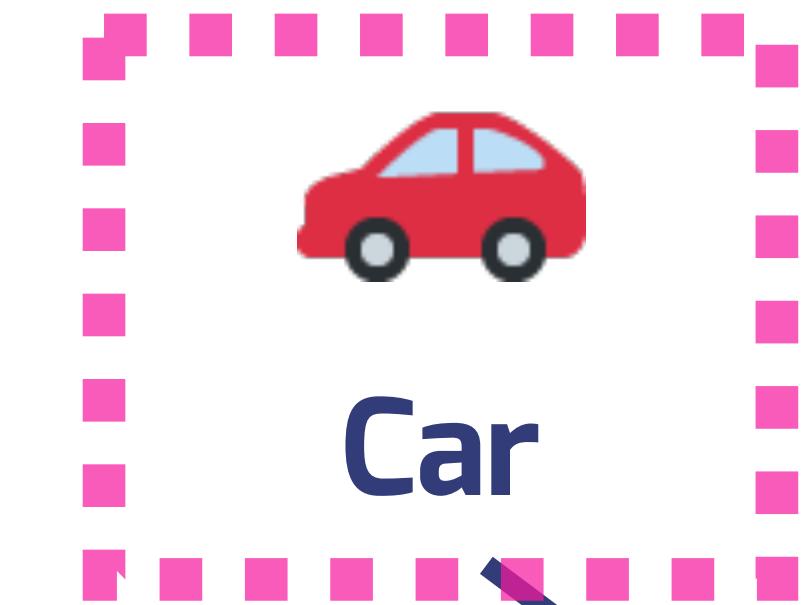
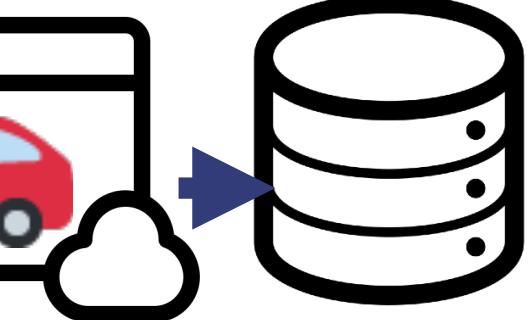
Book Car Response



Saga Log



Done



Car



Hotel



Flight

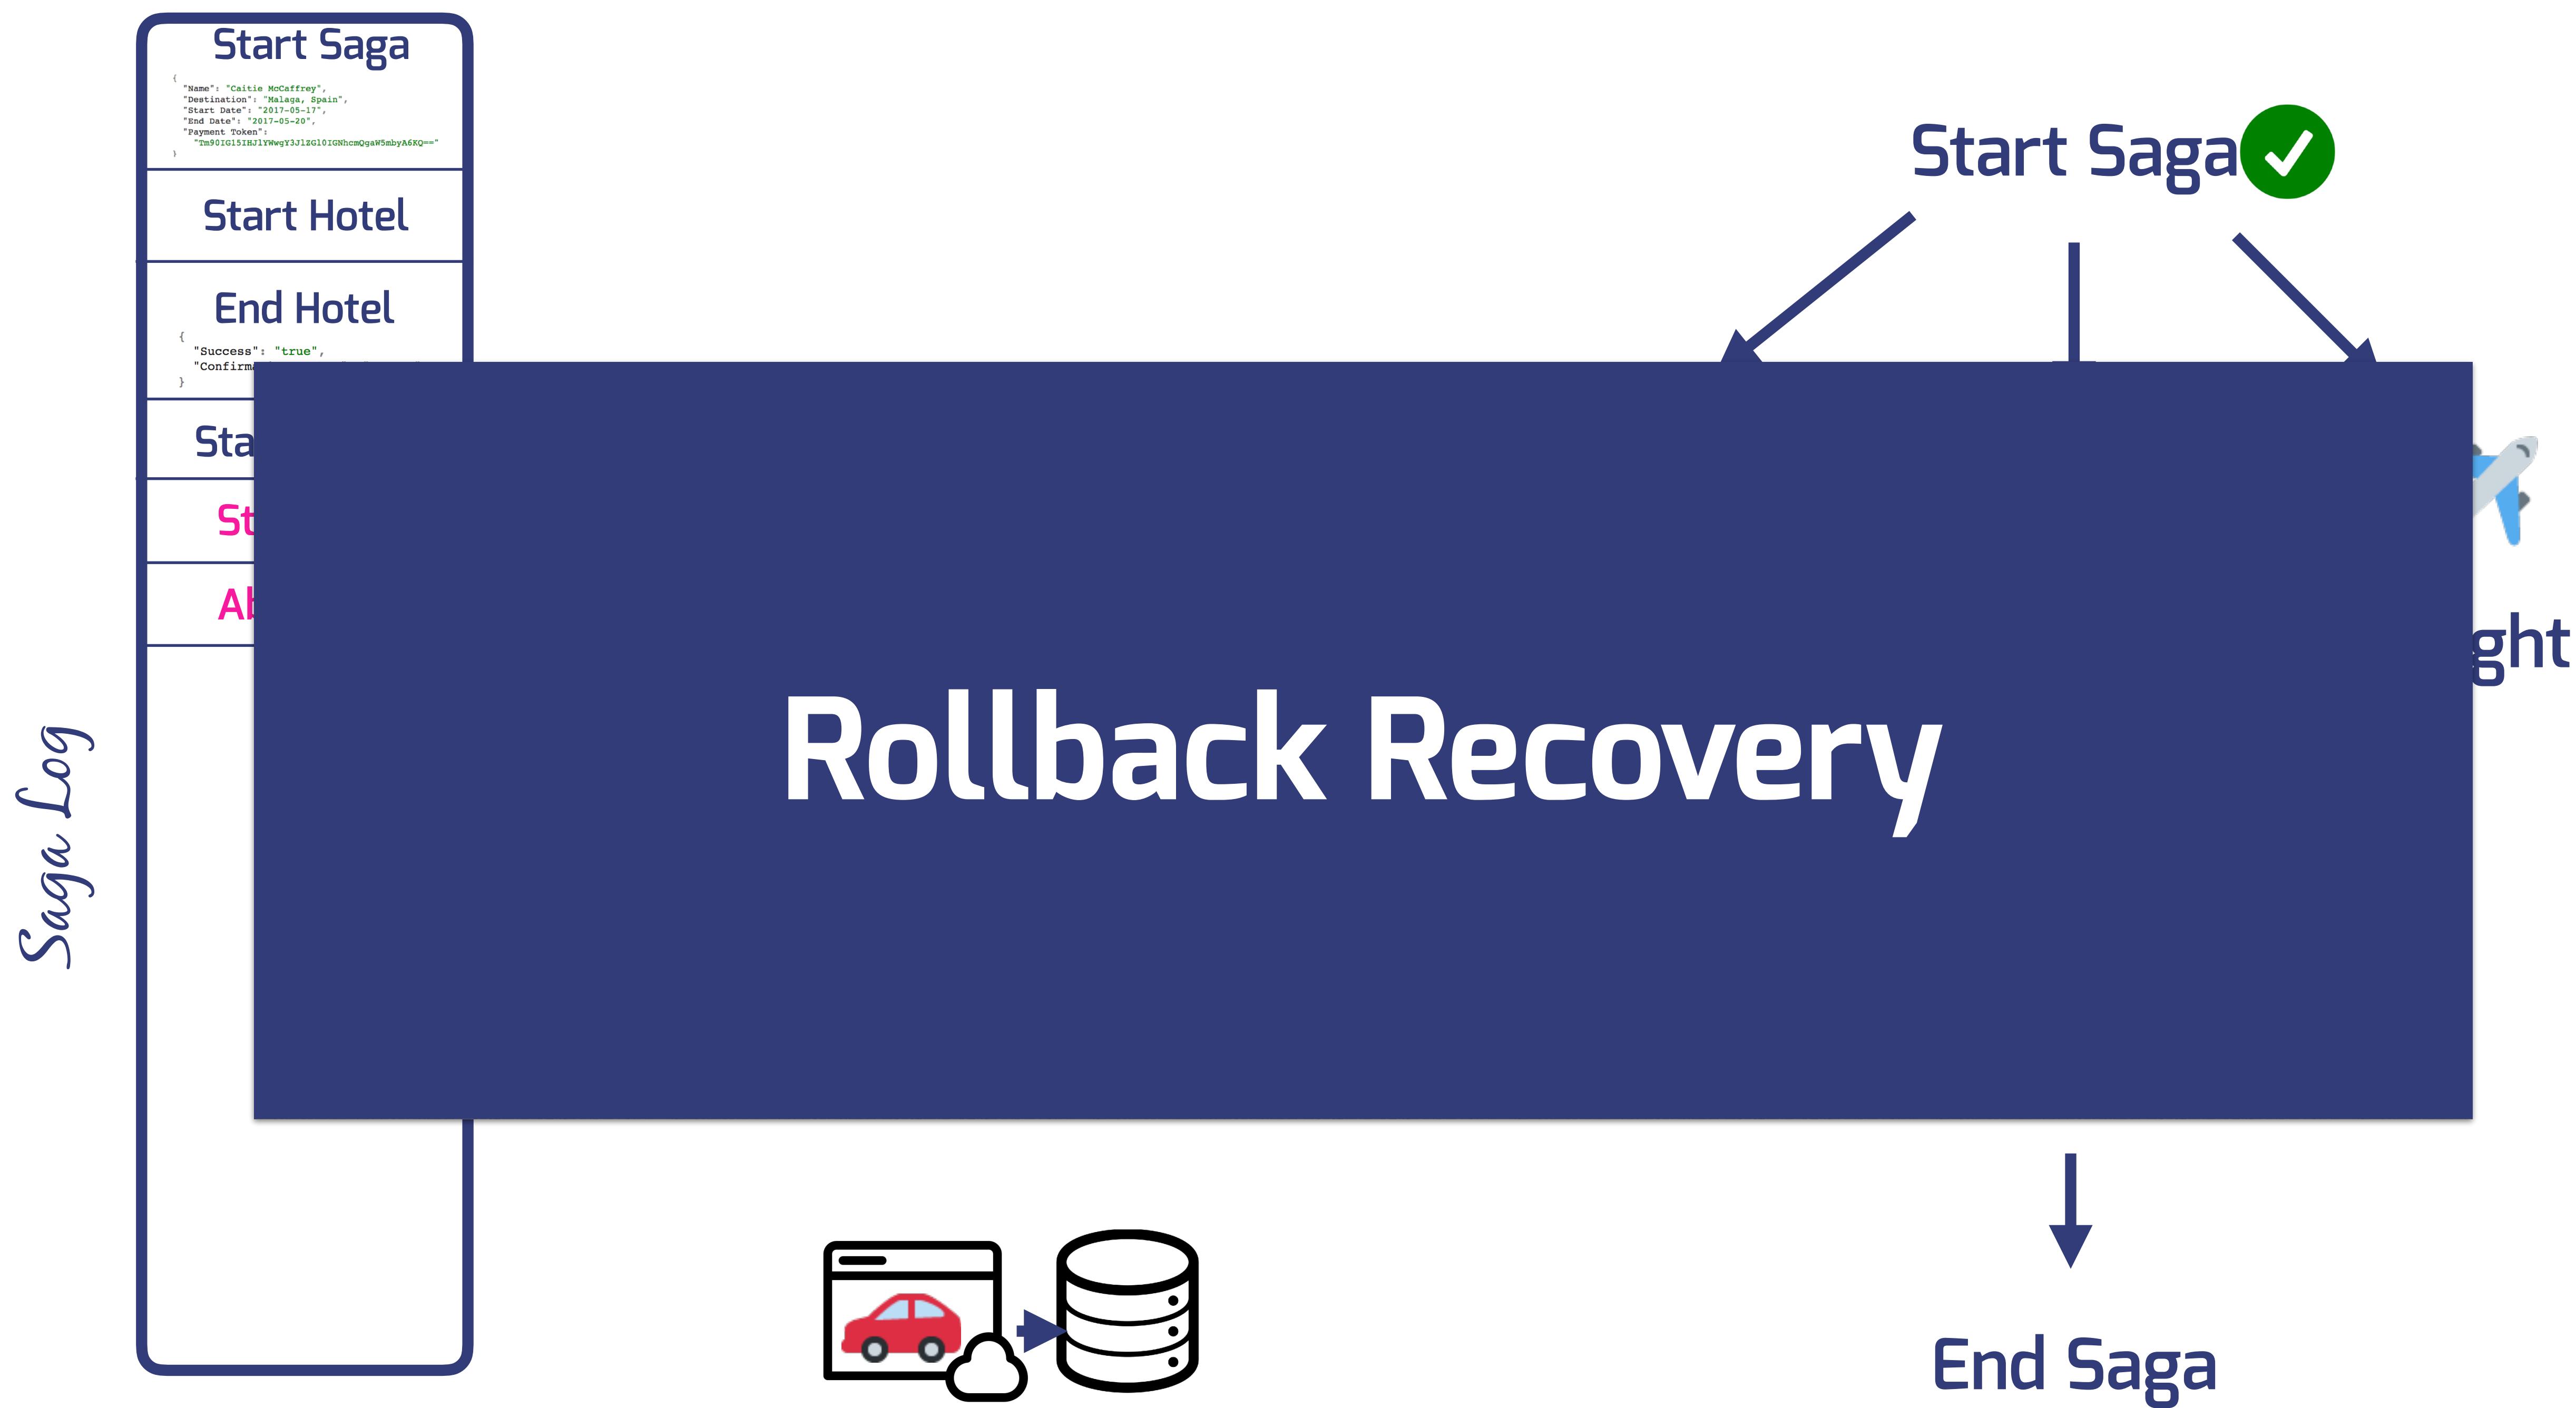


Payment

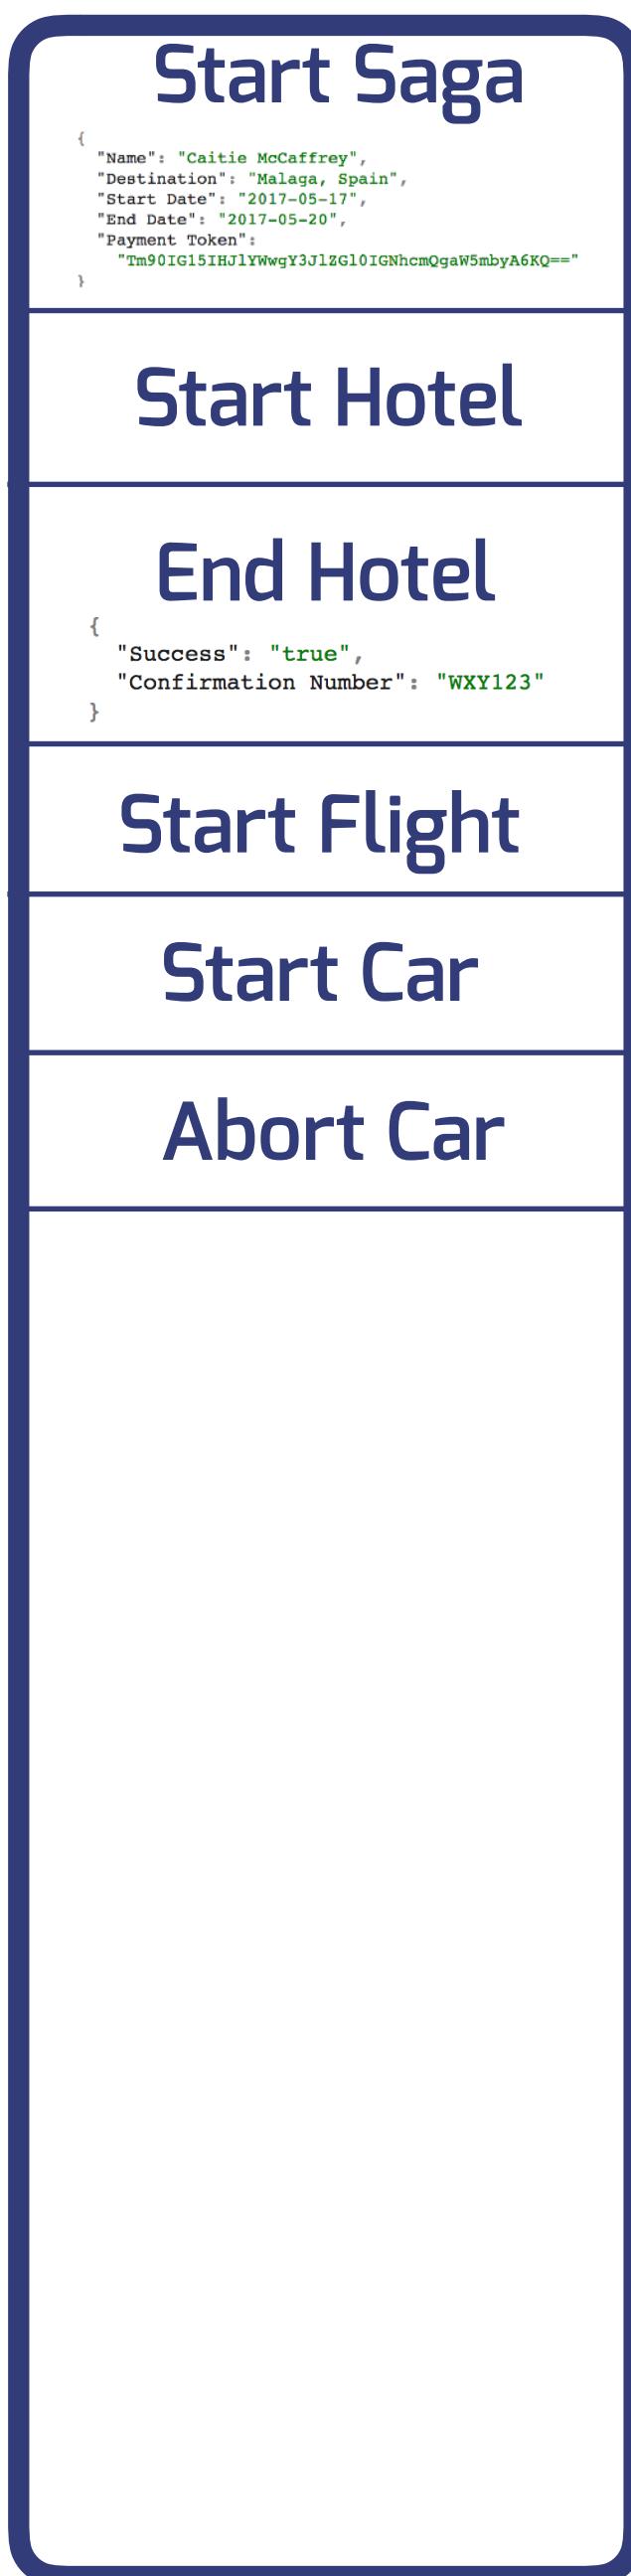
End Saga

Start Saga 





Saga Log



End Comp Saga



Car



Hotel



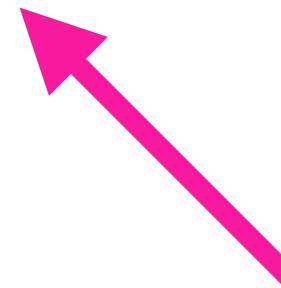
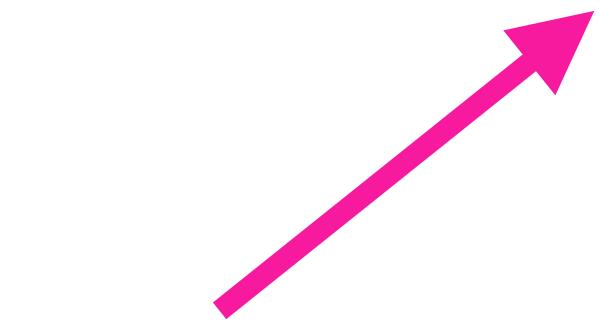
Flight



Payment



Start Comp Saga



Saga Log

Start Saga
{ "Name": "Caitie McCaffrey", "Destination": "Malaga, Spain", "Start Date": "2017-05-17", "End Date": "2017-05-20", "Payment Token": "7m901G151HJ1YWwgY3J12G101GNhcmQgaW5mbA6KQ==" }
Start Hotel
End Hotel
{ "Success": "true", "Confirmation Number": "WXY123" }
Start Flight
Start Car
Abort Car



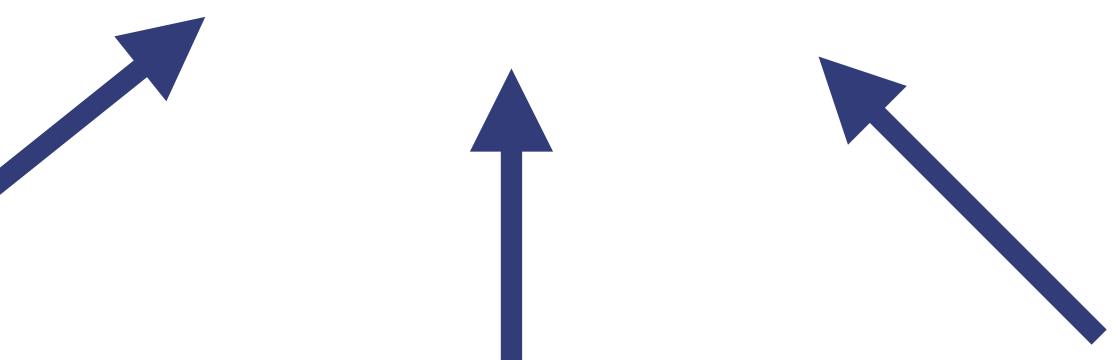
End Comp Saga

Payment

Hotel

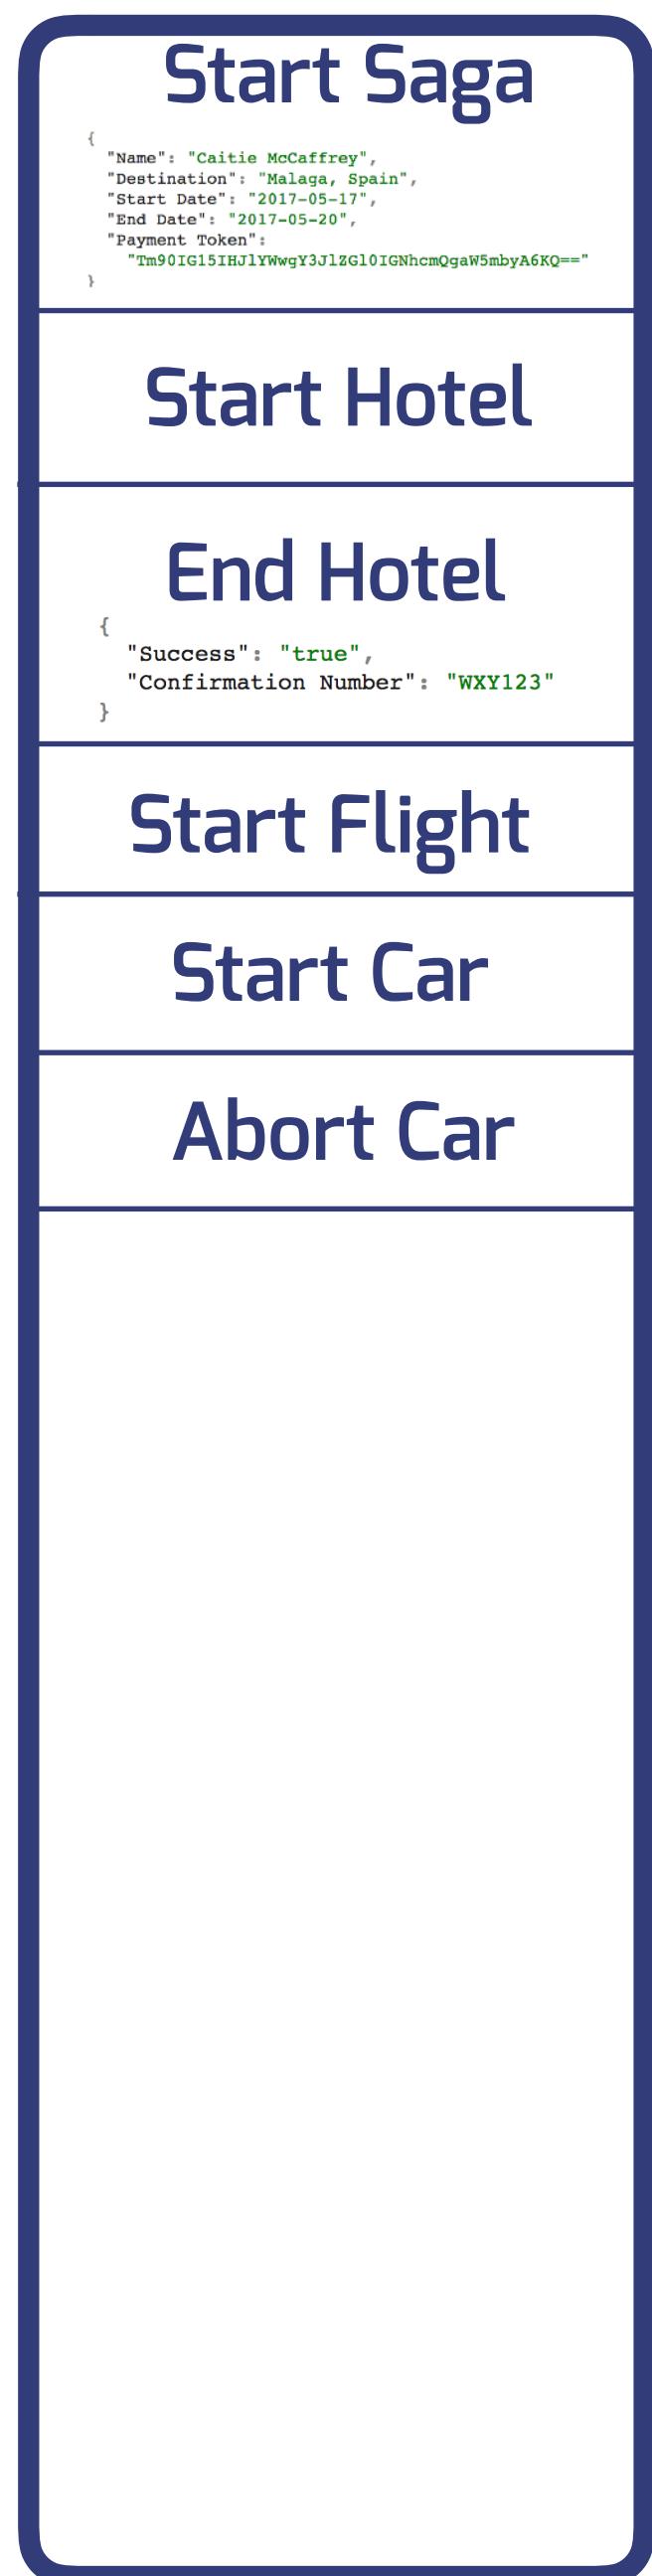


Flight



Start Comp Saga 

Saga Log



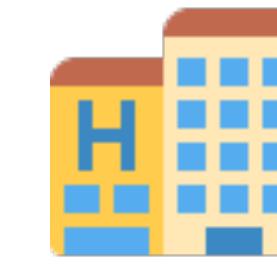
Payment Log
Entries?



End Comp Saga



Car



Hotel

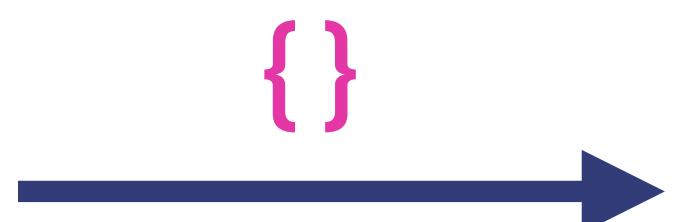
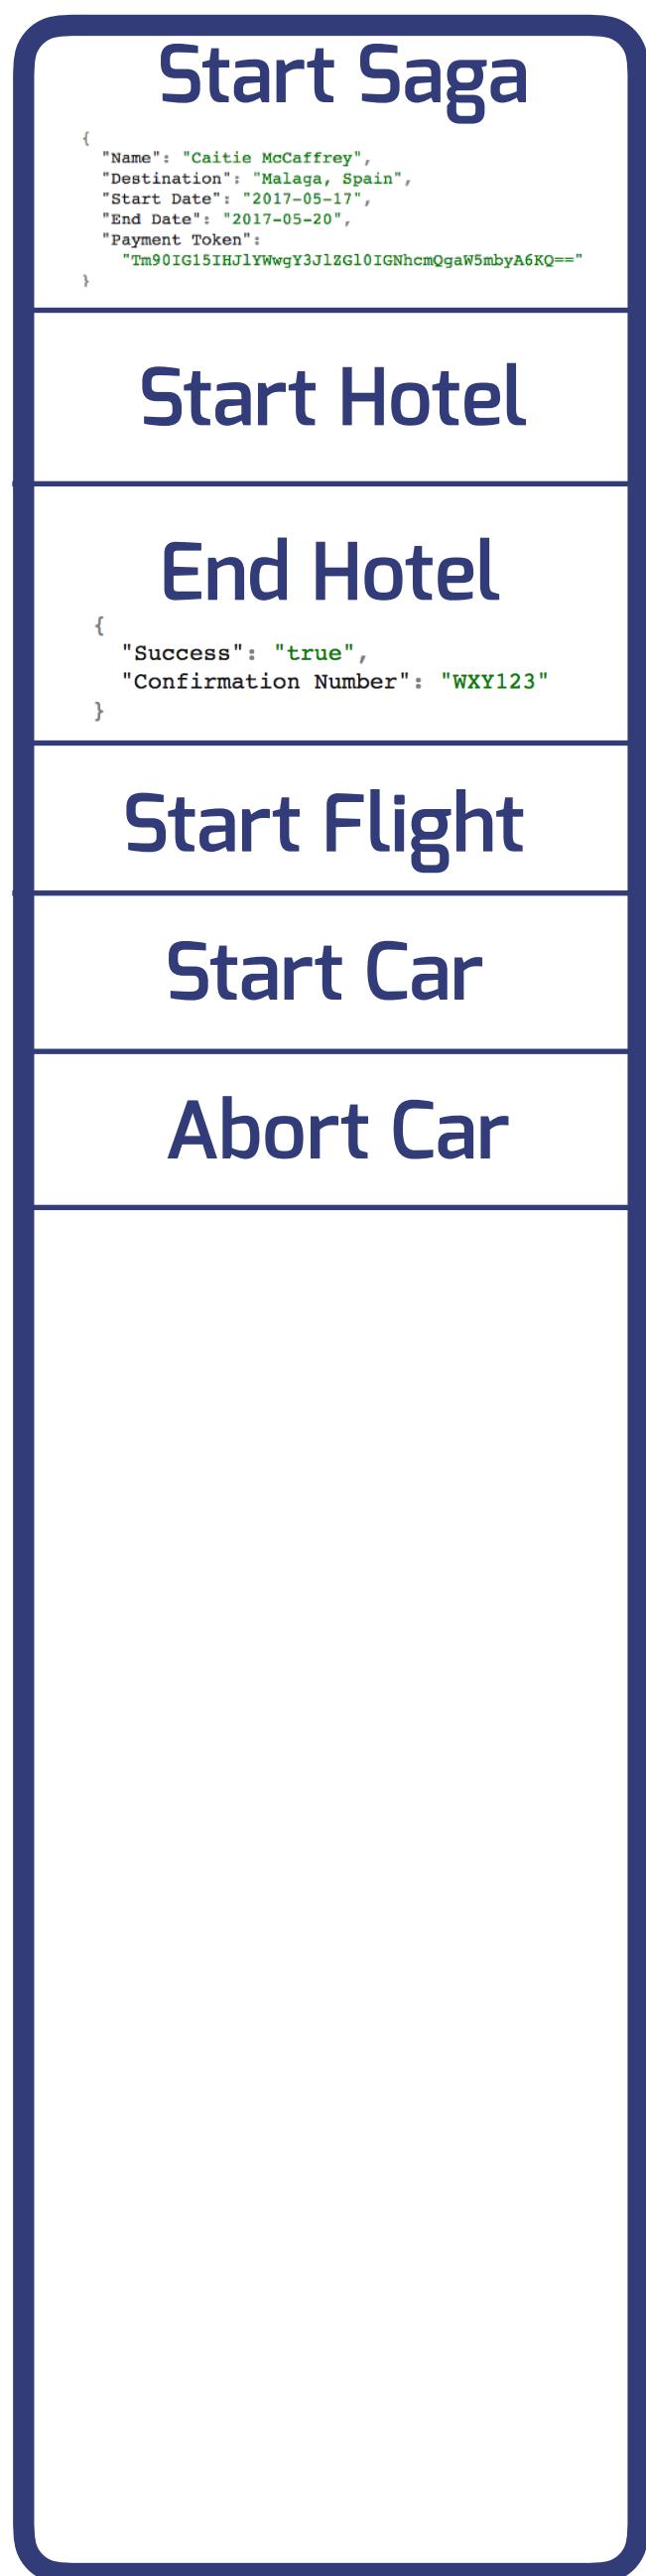


Flight

Payment

Start Comp Saga

Saga Log



End Comp Saga



Car



Hotel

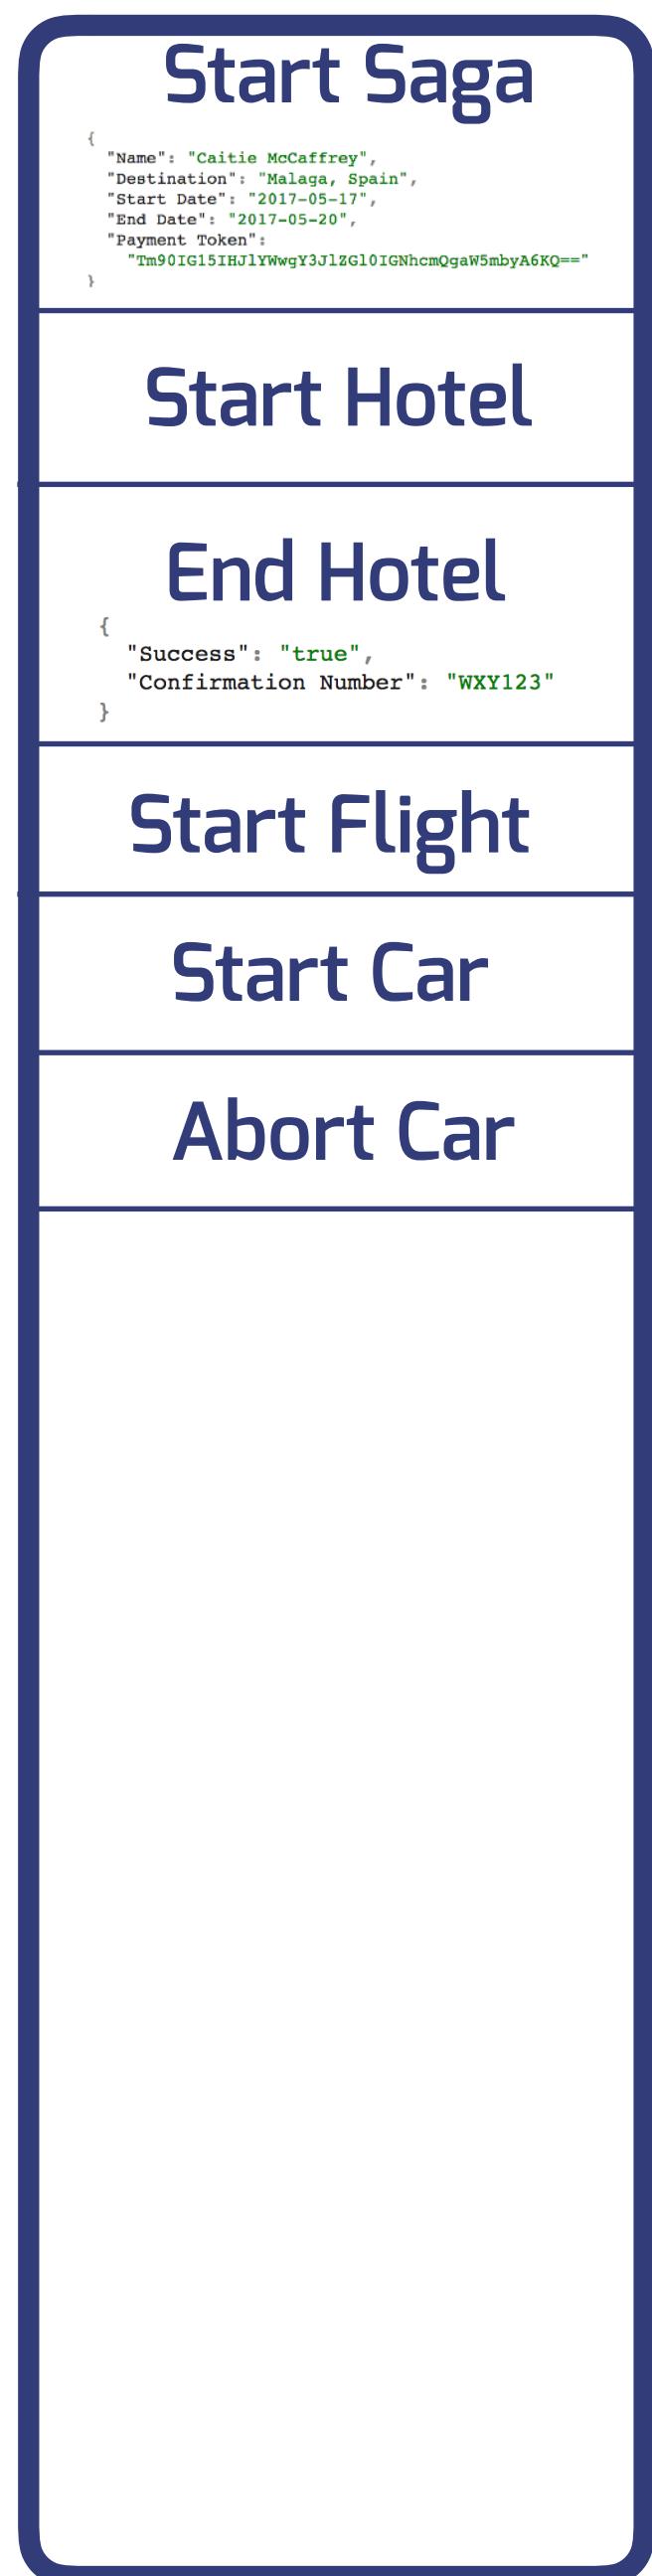


Flight

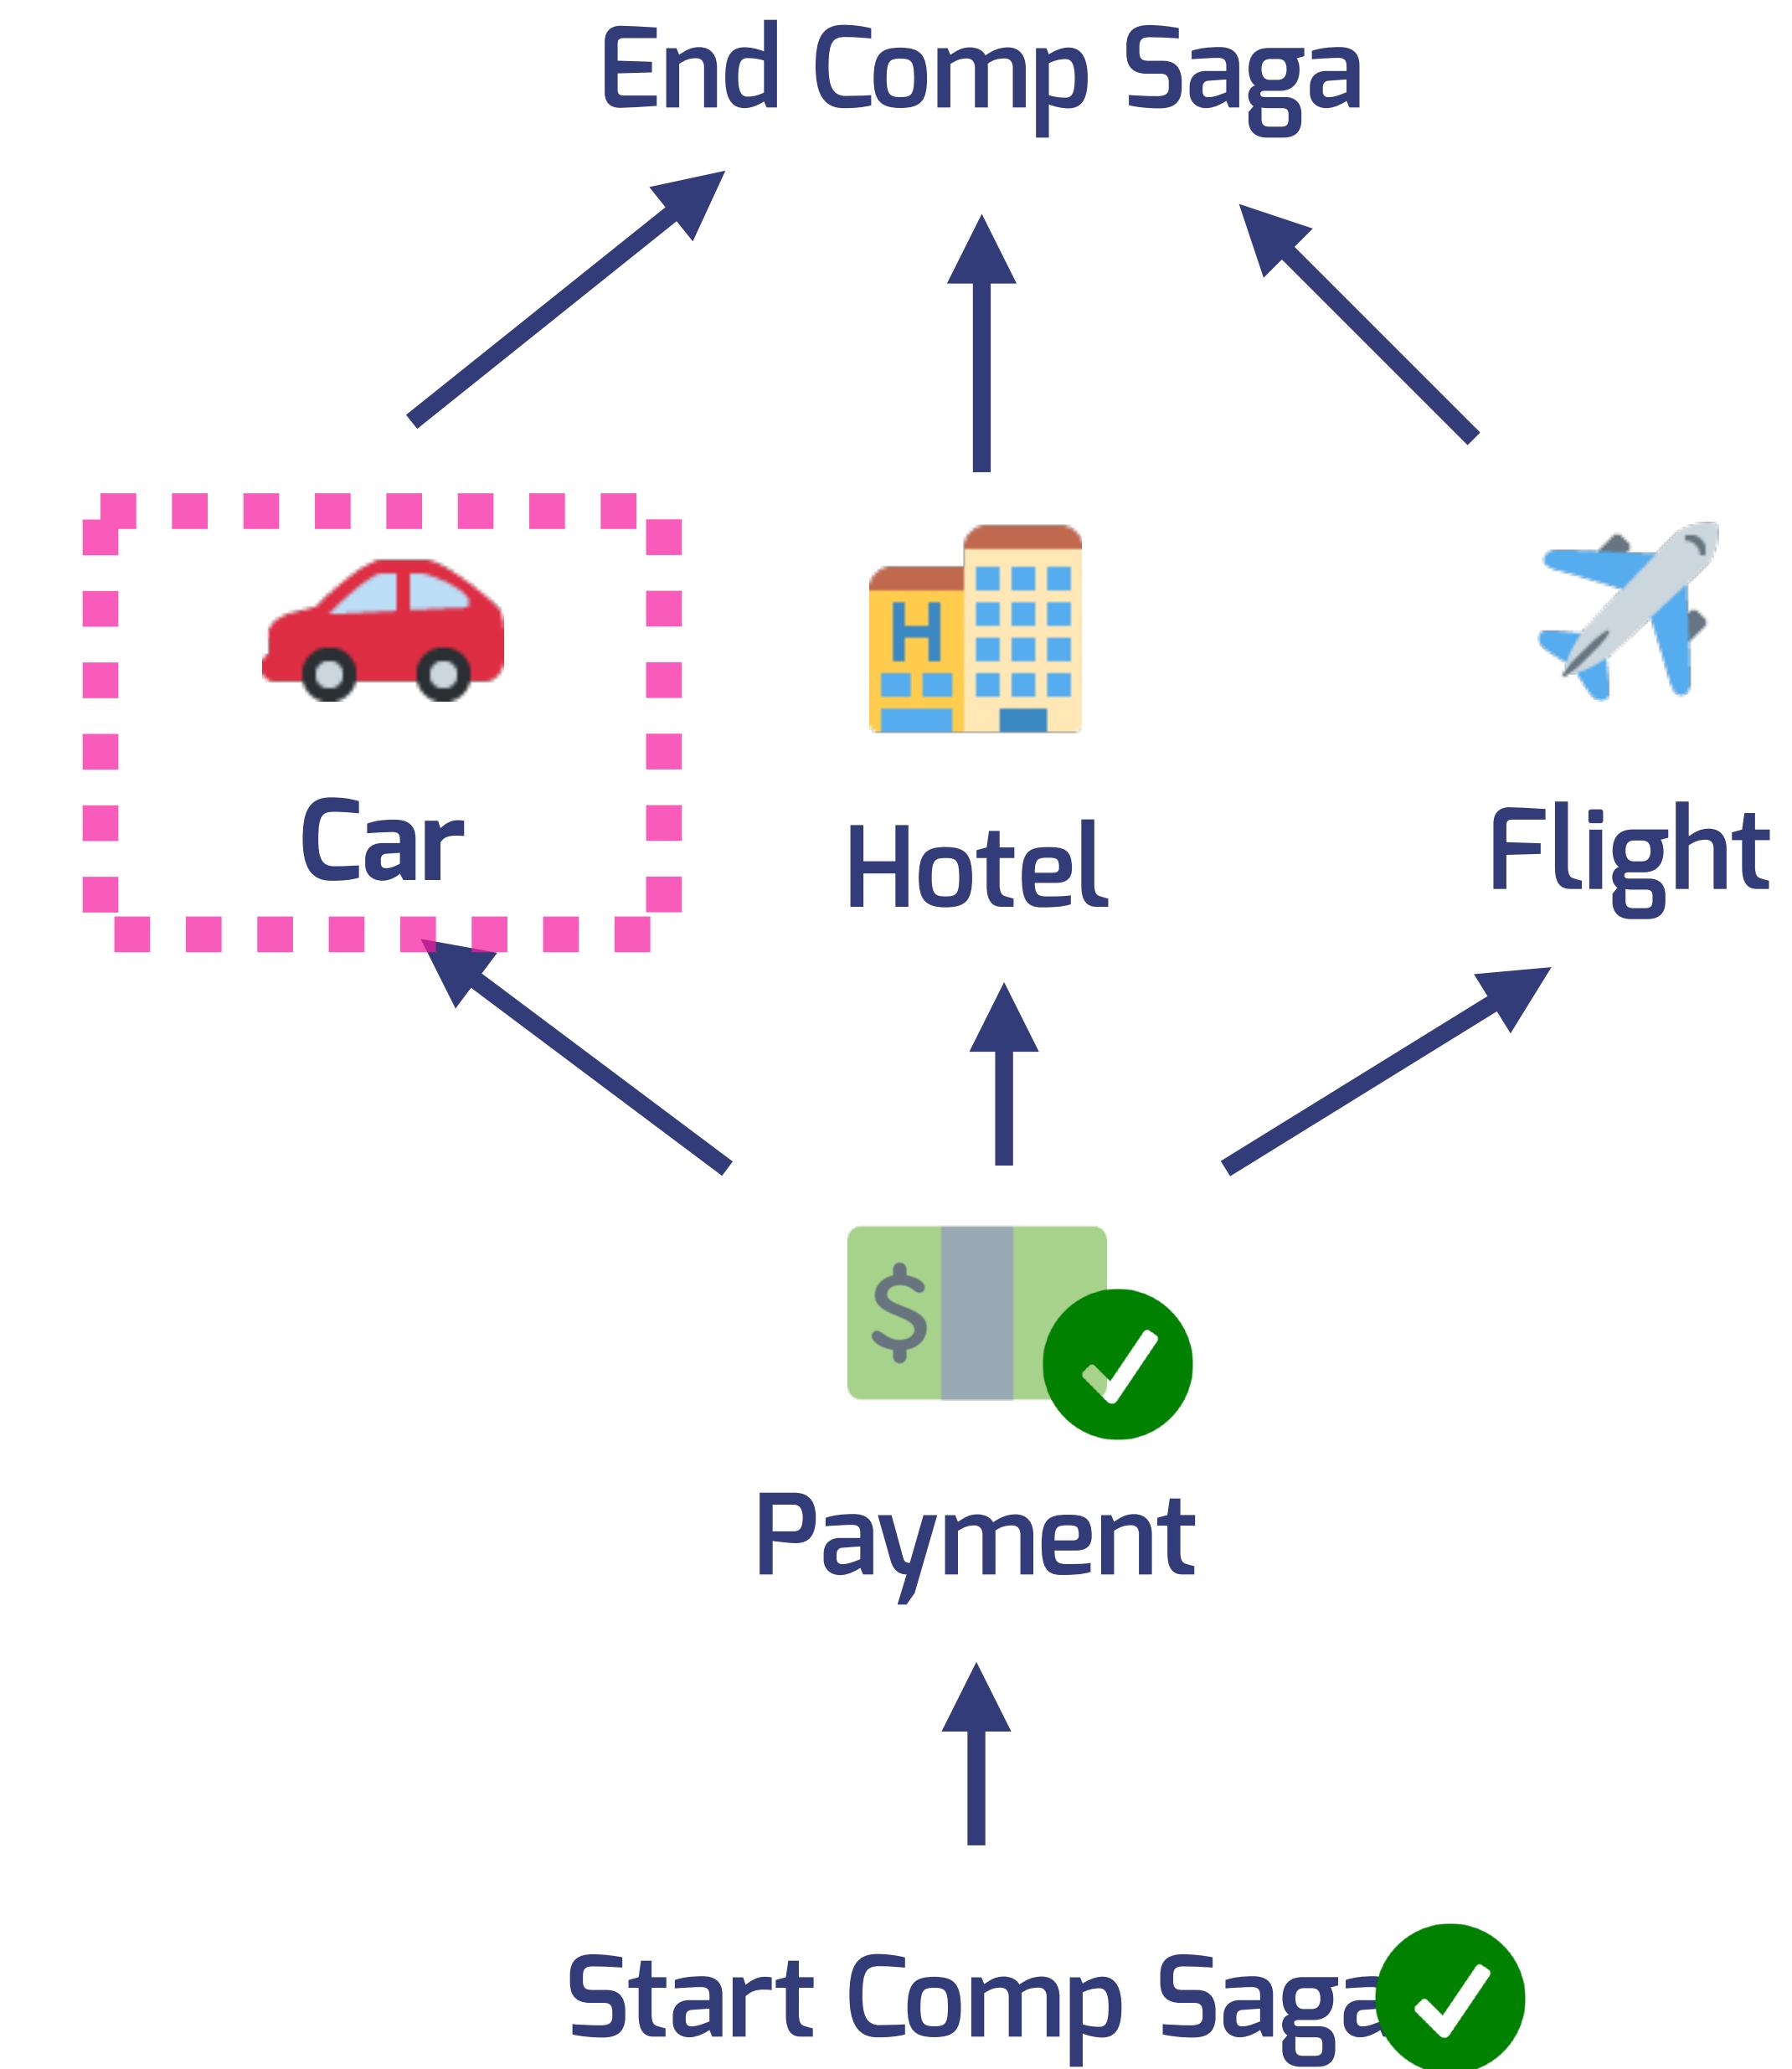
Start Comp Saga



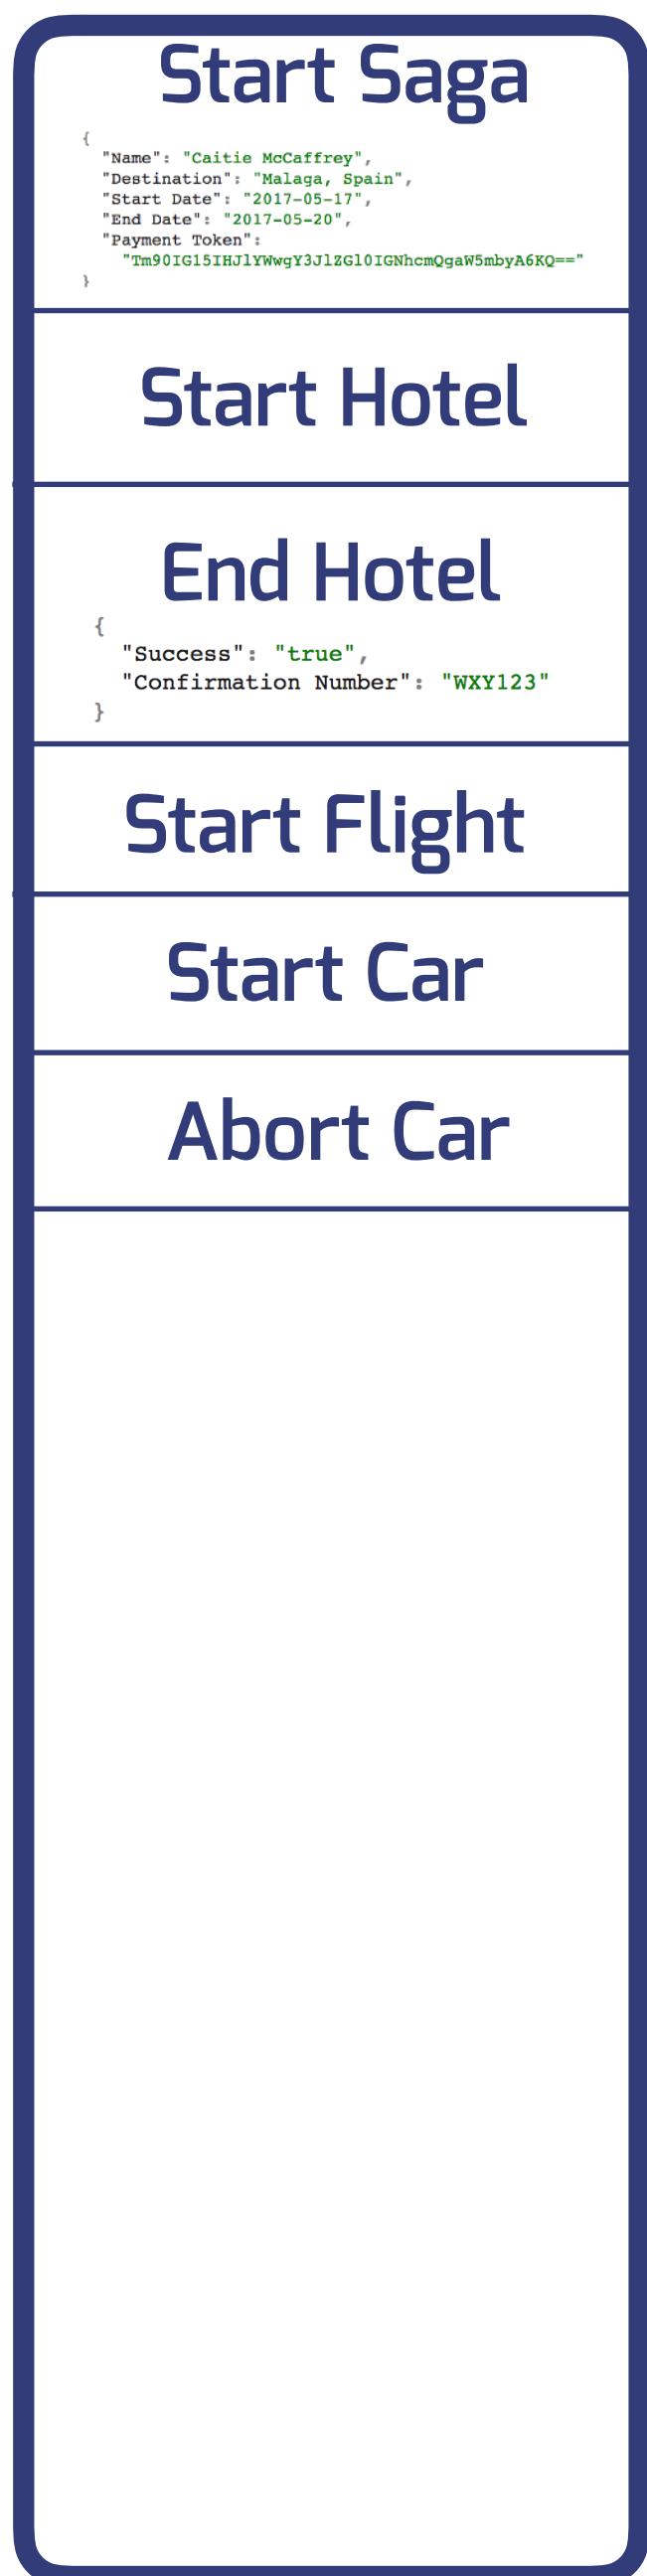
Saga Log



Car Log Entries?



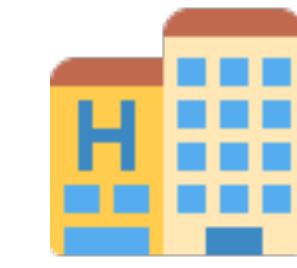
Saga Log



{ Start, Abort }



End Comp Saga



Hotel

Flight

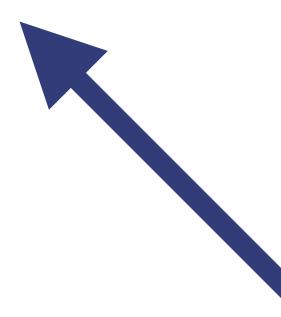
Payment



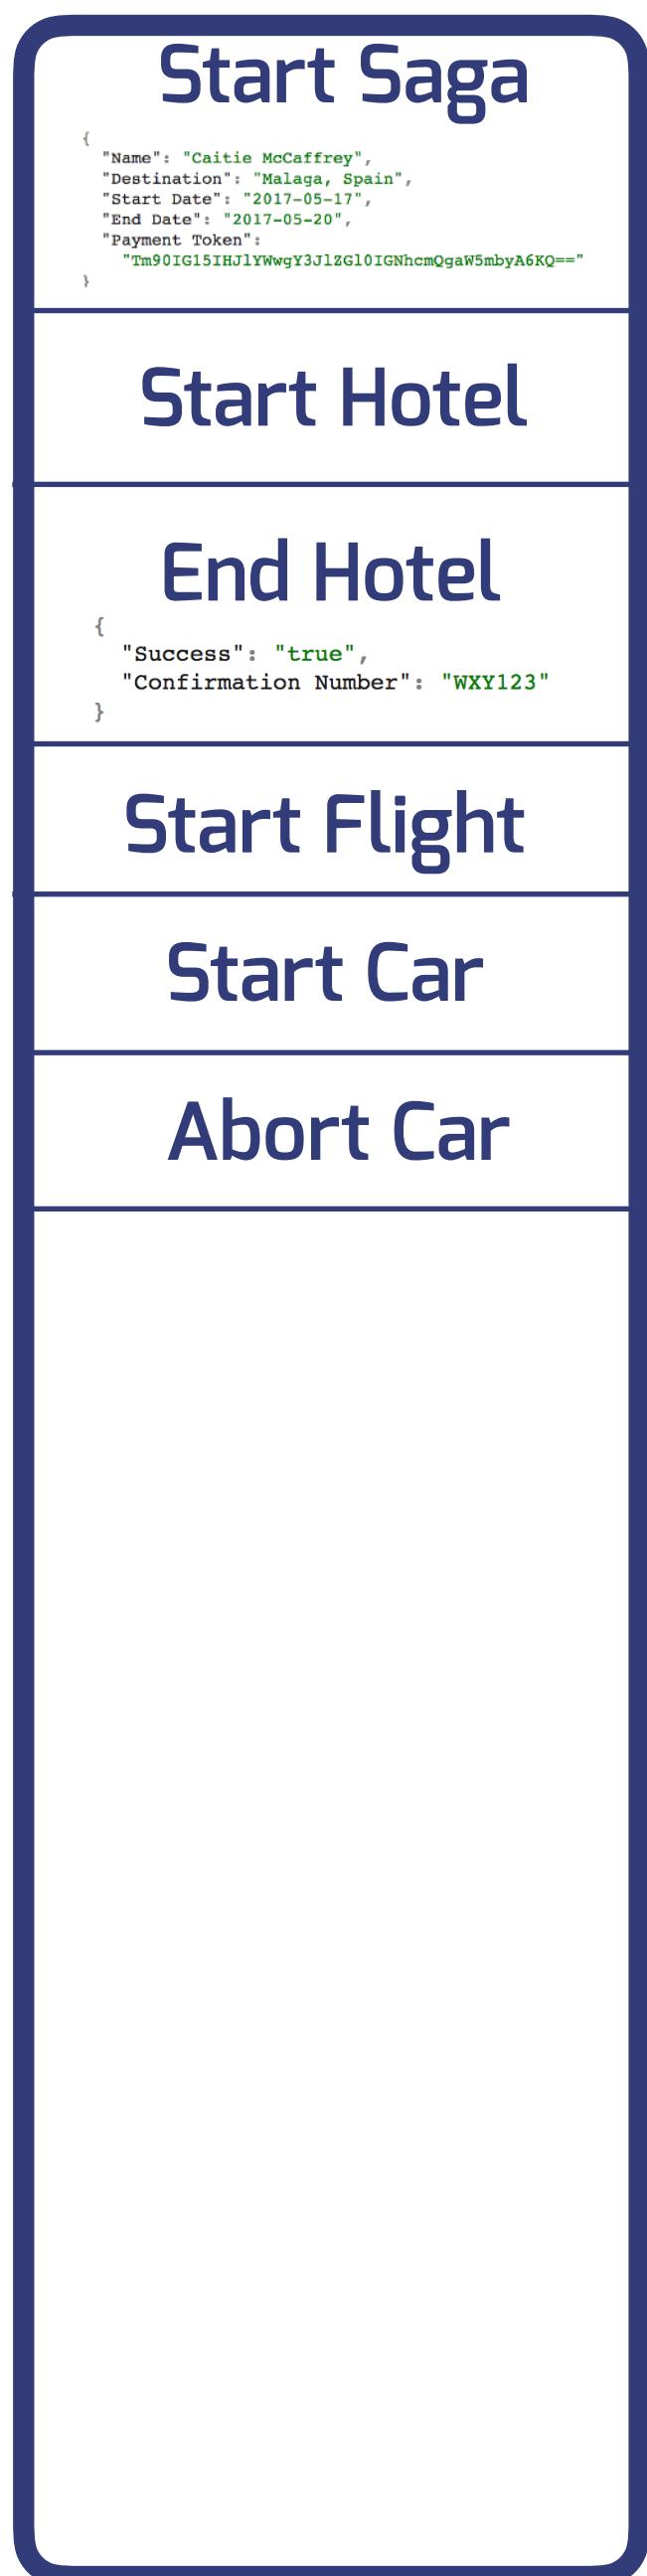
Start Comp Saga 



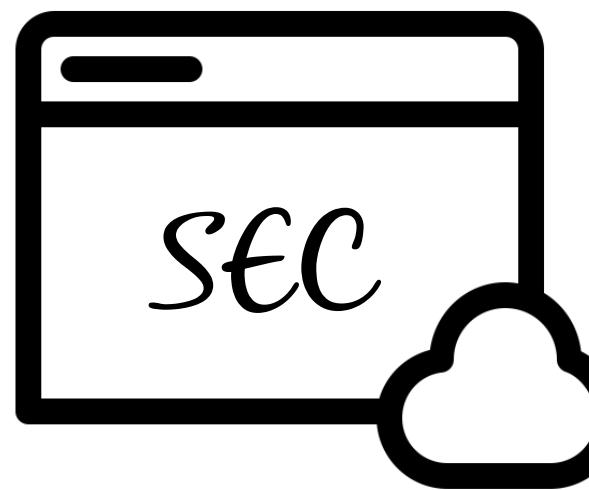
Car



Saga Log



Hotel Log
Entries?



End Comp Saga



Car



Hotel



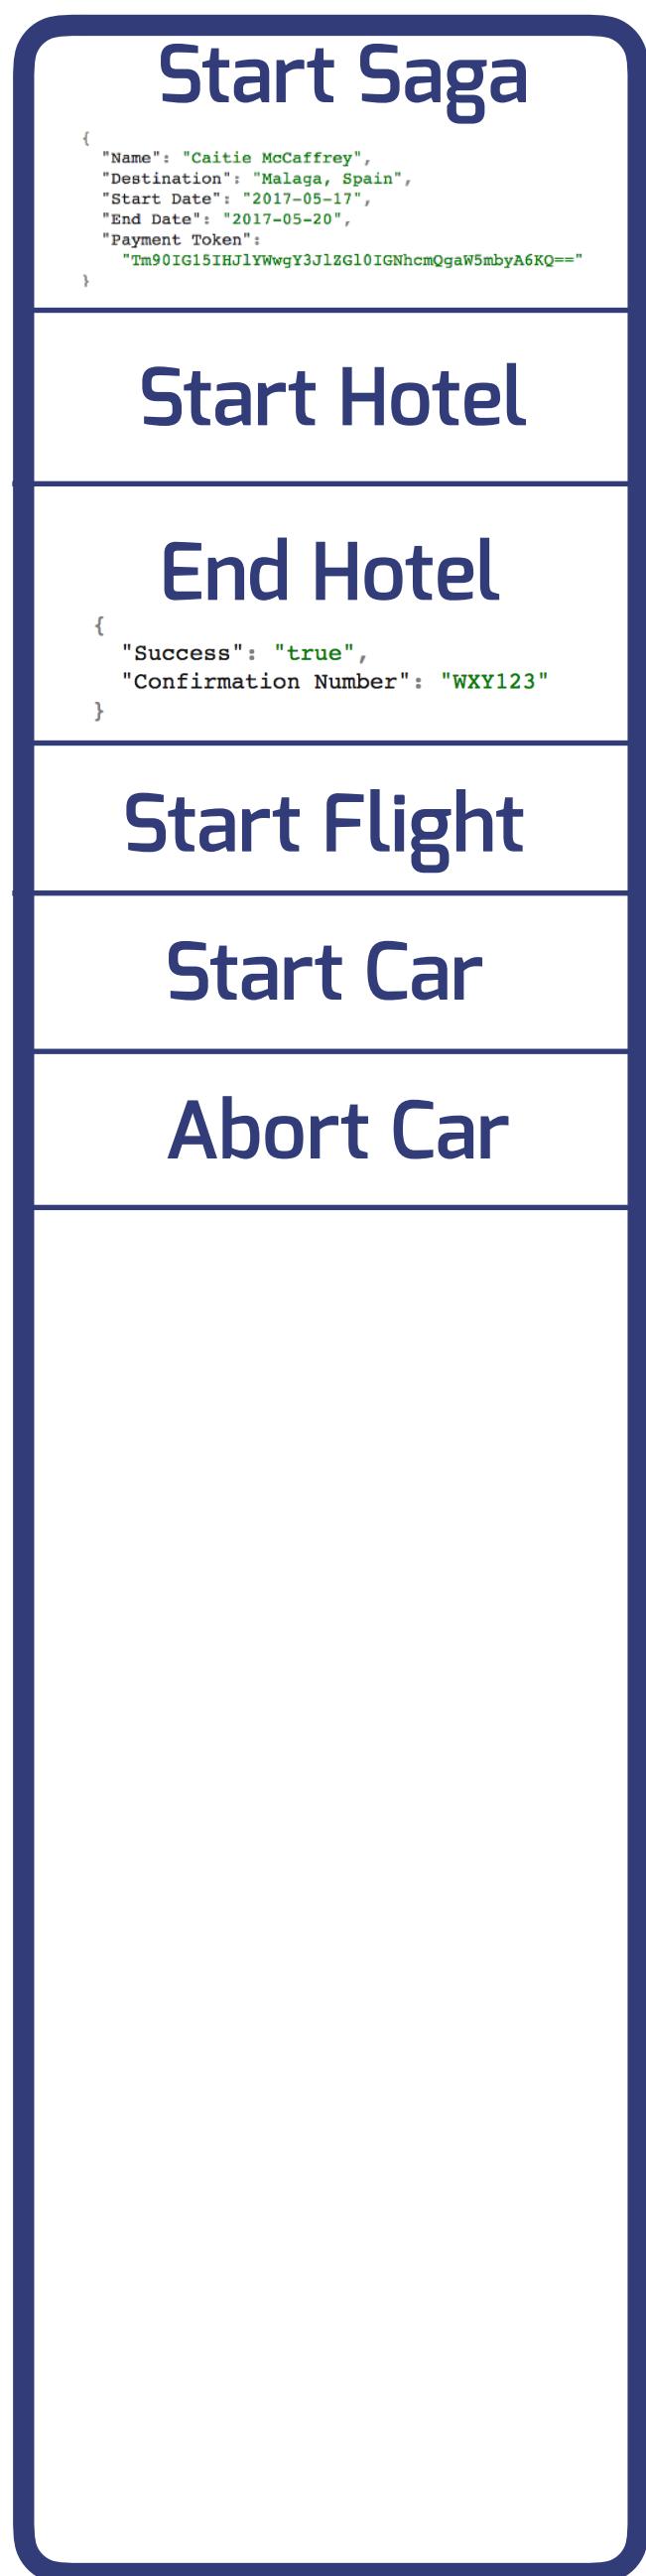
Flight



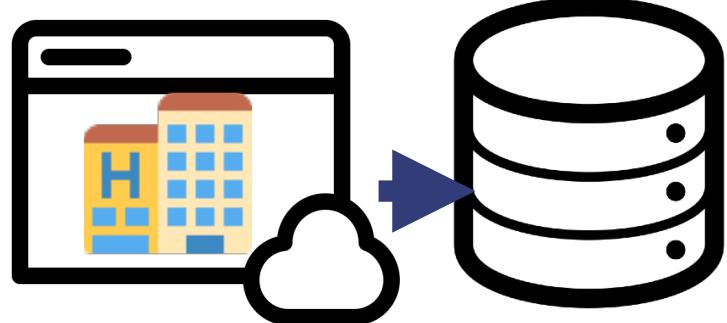
Payment

Start Comp Saga

Saga Log



{ Start, End }



End Comp Saga



Car



Hotel



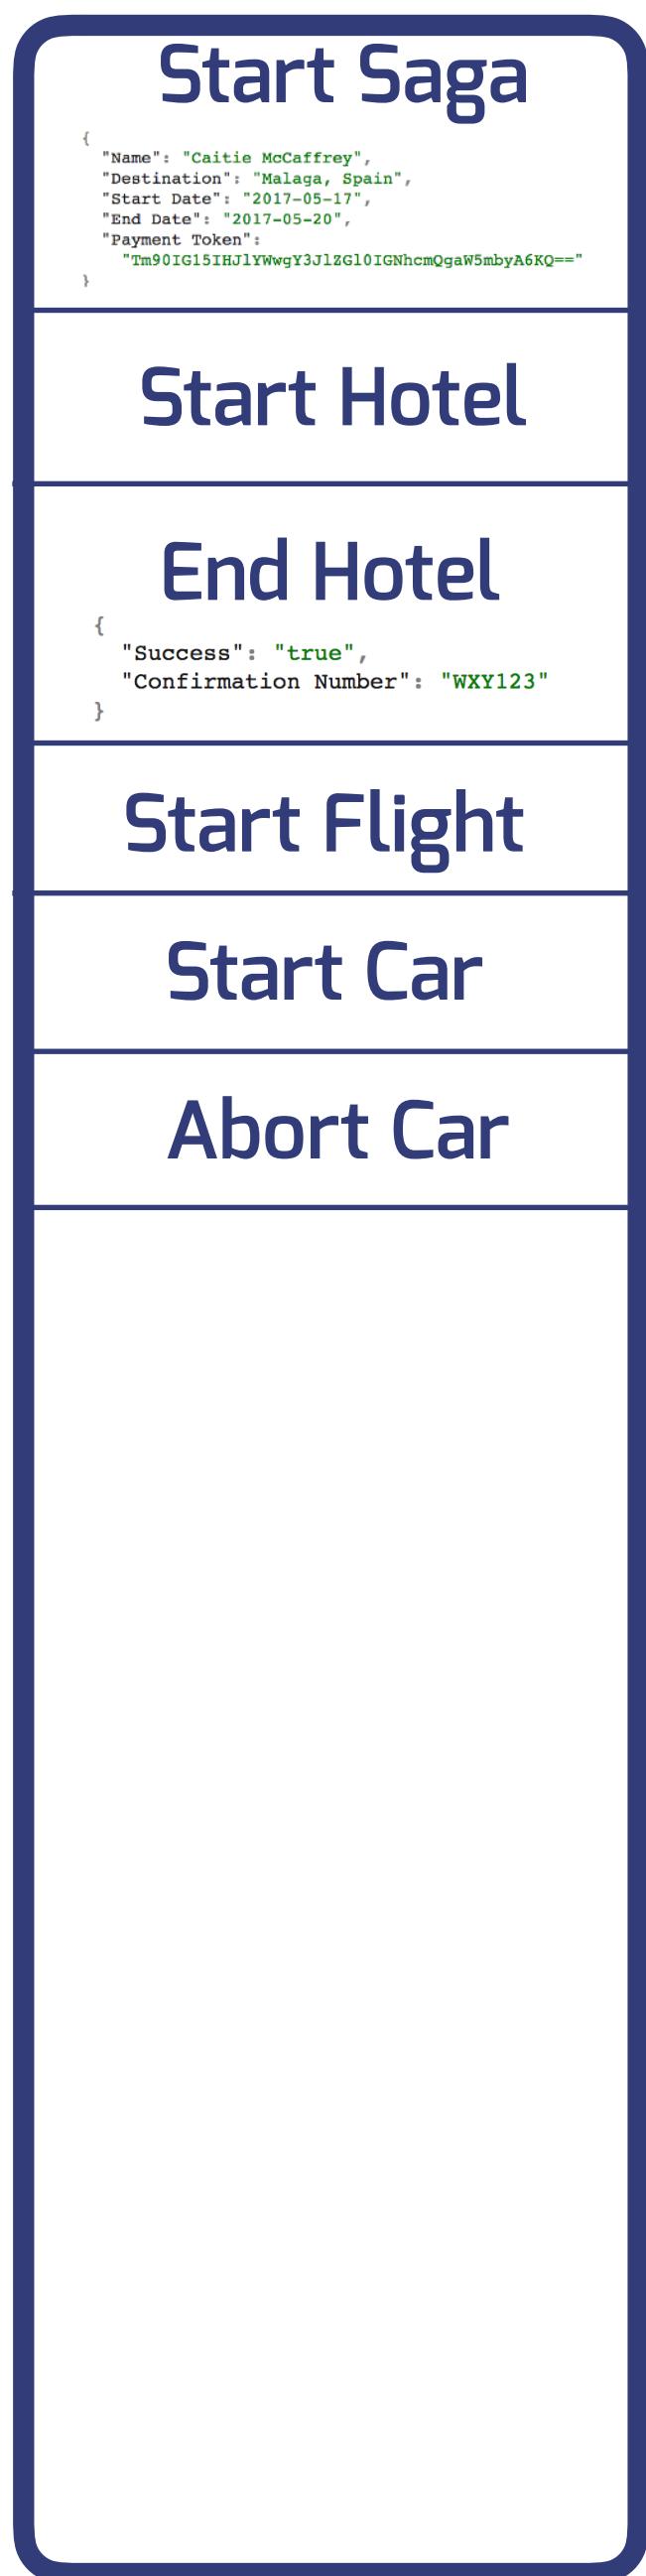
Flight



Payment

Start Comp Saga

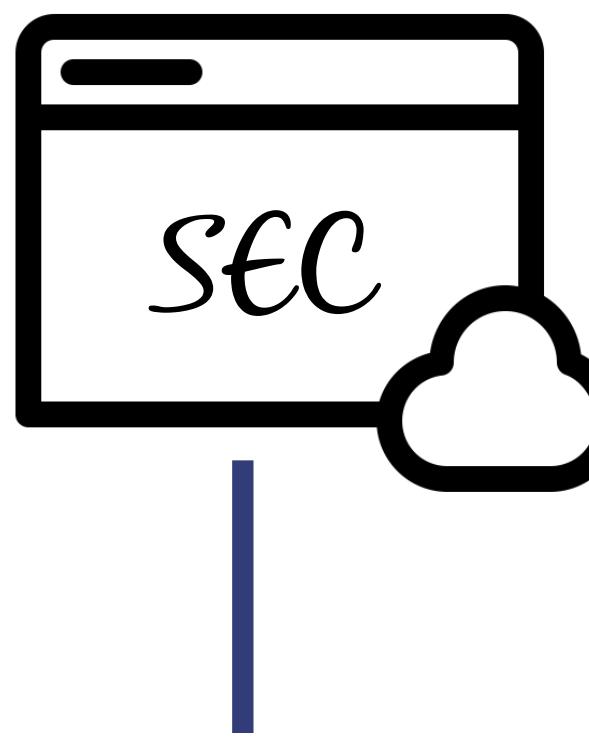
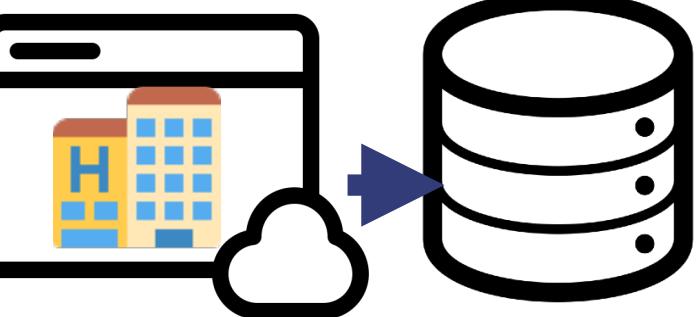
Saga Log



{ Start, End }

Cancel Hotel Request

```
{  
  "Name": "Caitie McCaffrey",  
  "Confirmation Number": "WXY123"  
}
```



End Comp Saga



Car



Hotel



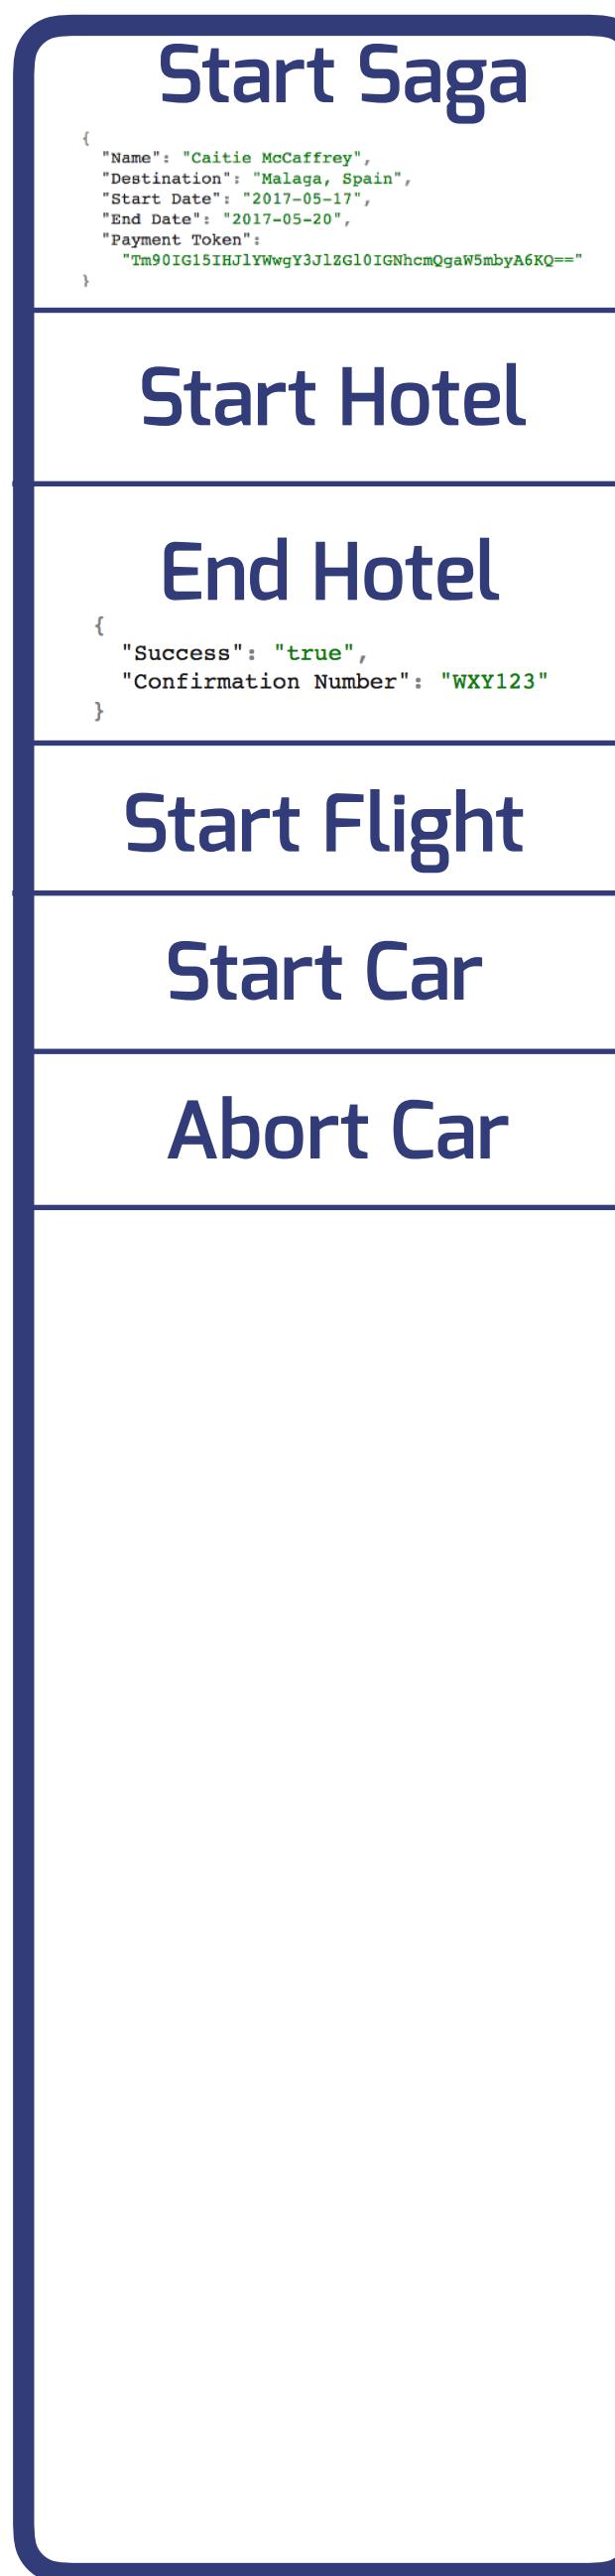
Flight



Payment

Start Comp Saga

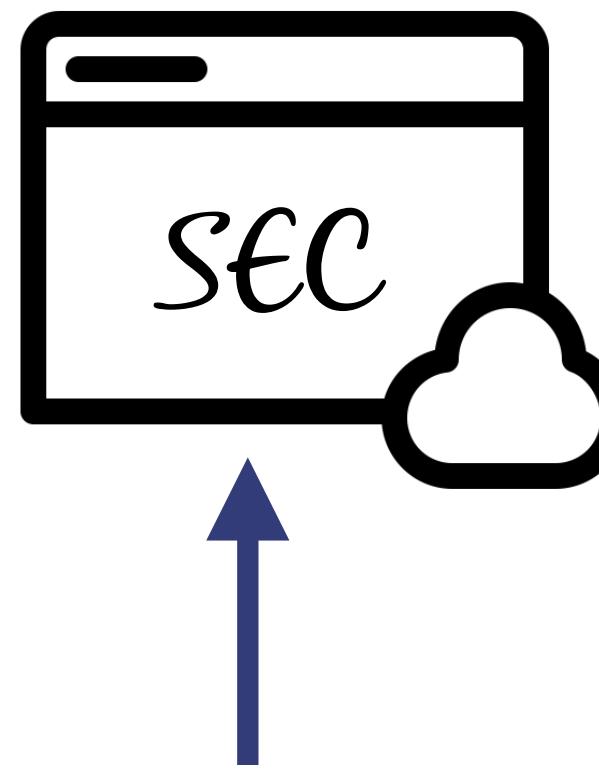
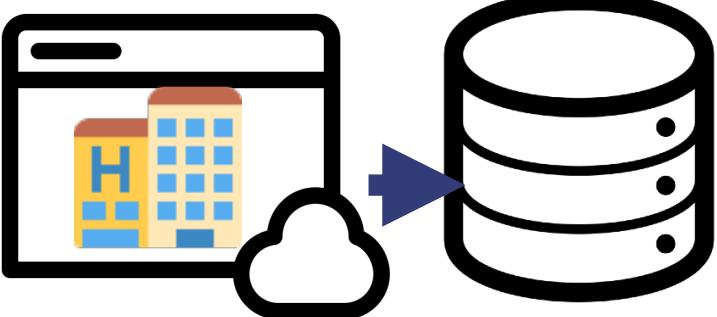
Saga Log



Comp Hotel

Cancel Hotel Response

```
{  
  "success": "true",  
  "Confirmation Number": "WXY123"  
}
```



End Comp Saga



Car

Hotel

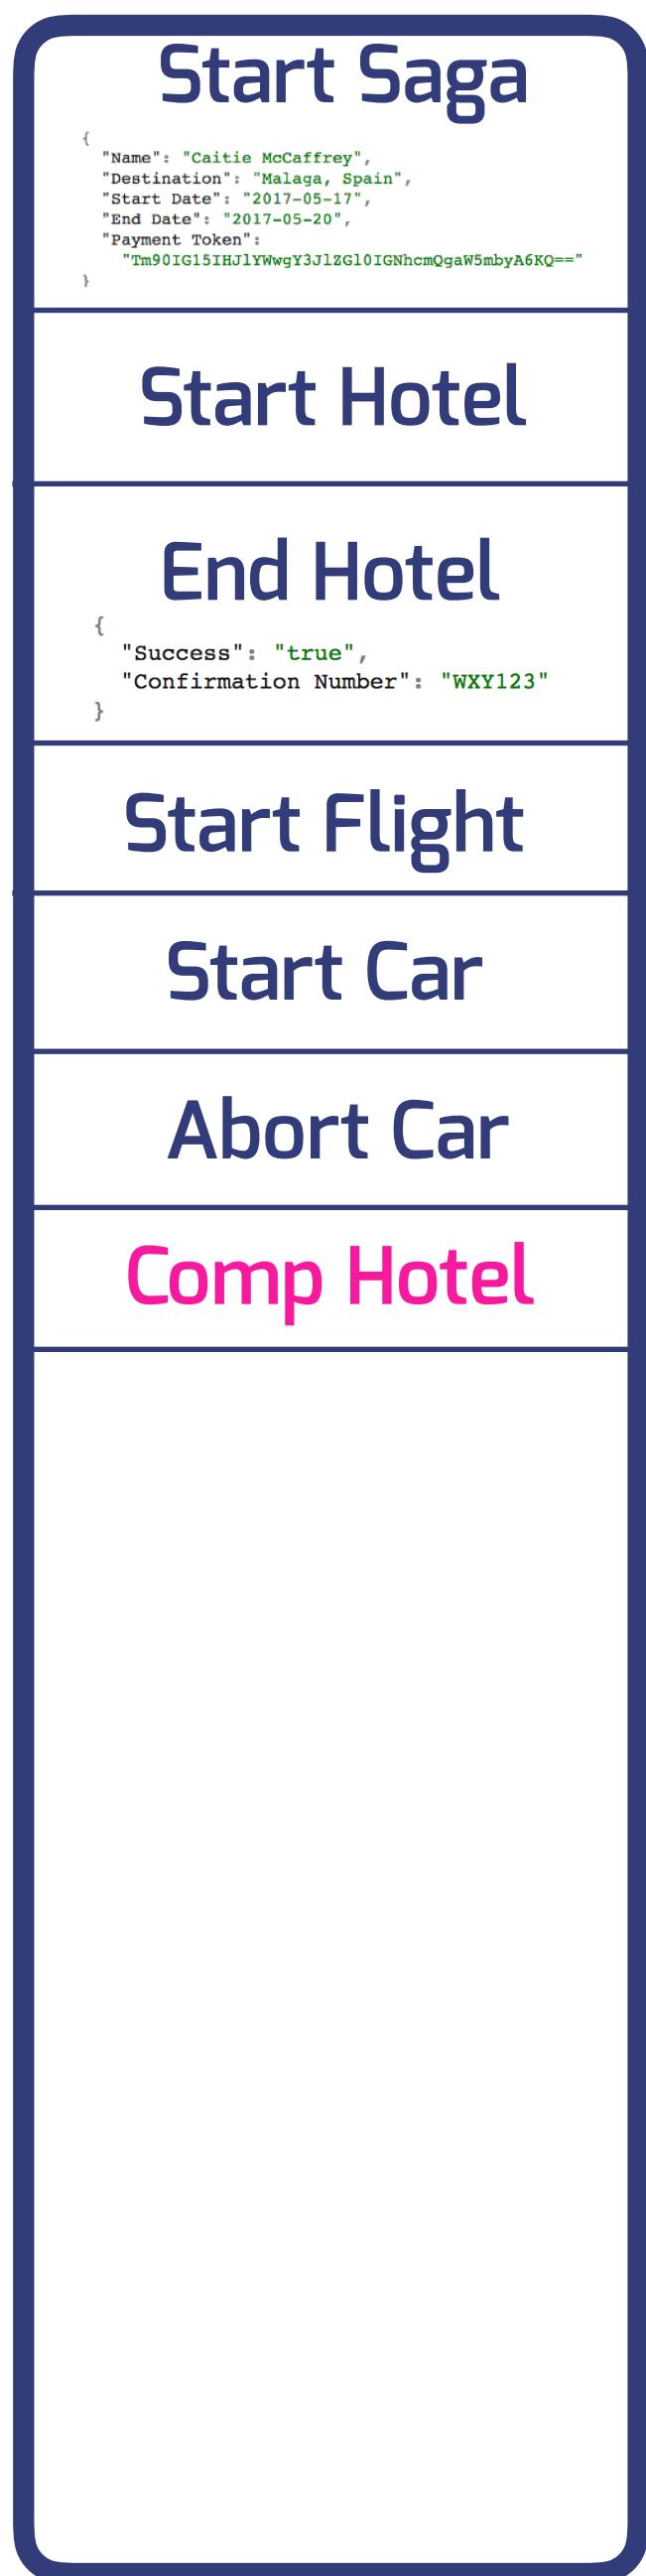
Flight



Payment

Start Comp Saga 

Saga Log



Done



End Comp Saga



Car



Hotel



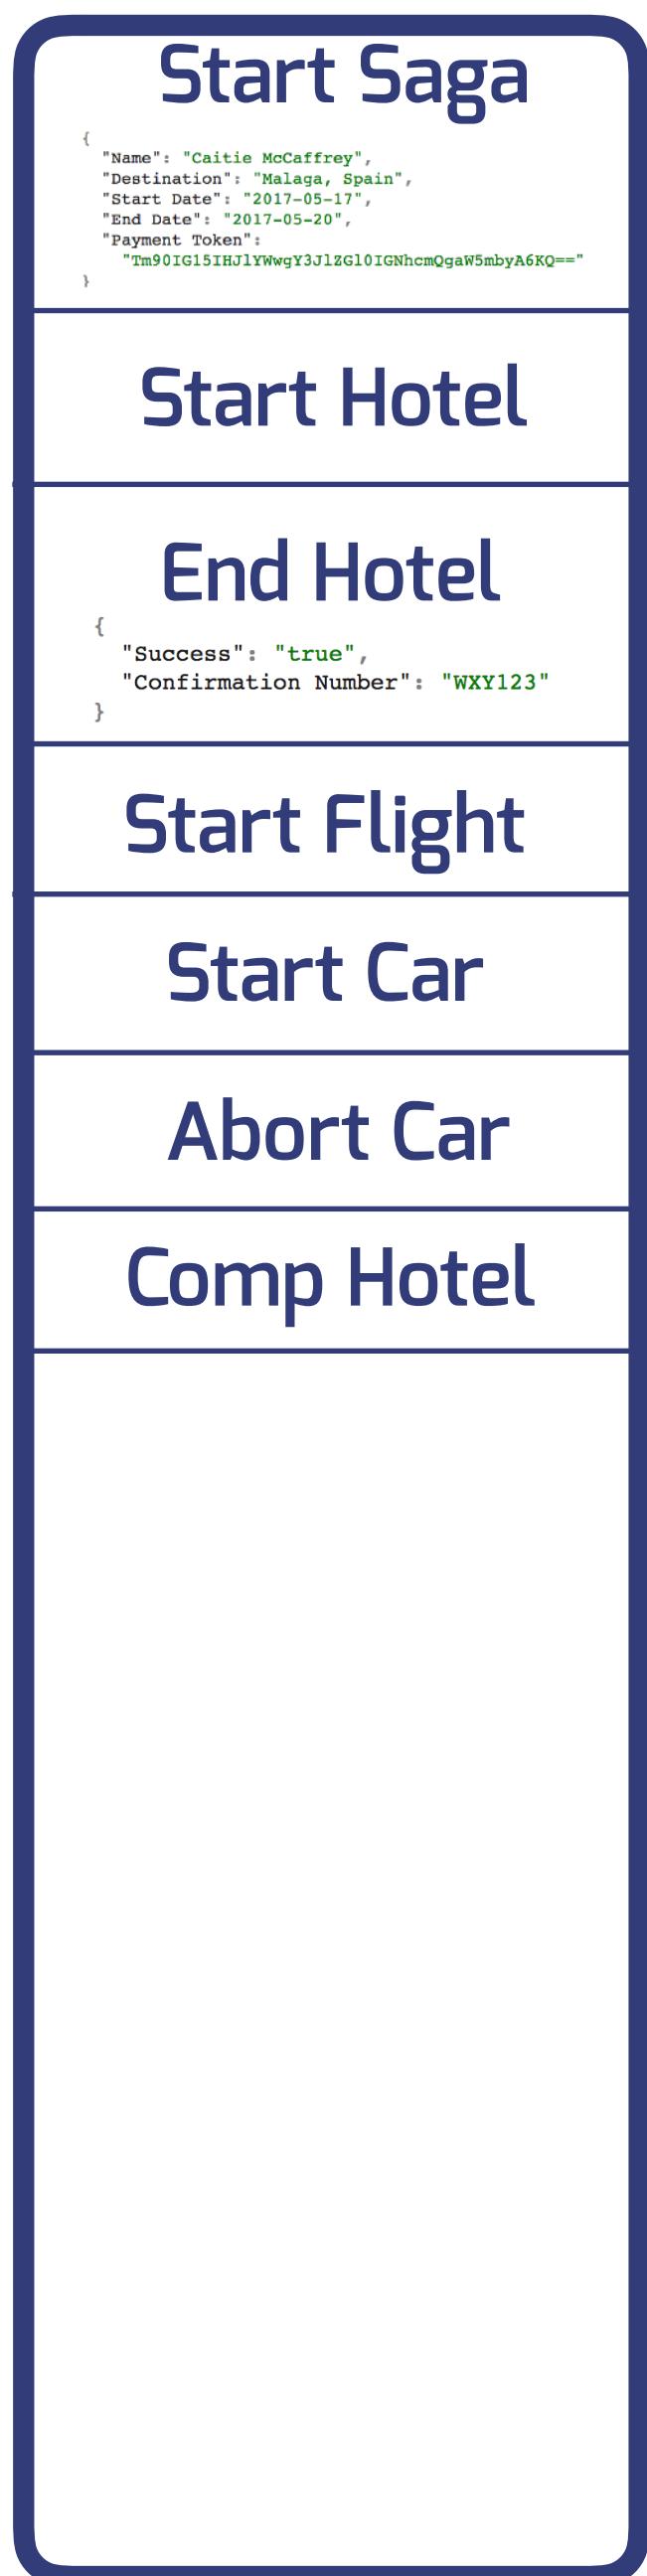
Flight



Payment

Start Comp Saga

Saga Log



Flight Log Entries?



End Comp Saga

Start Comp Saga 



Car



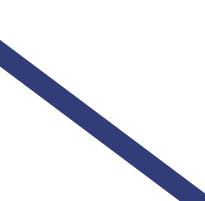
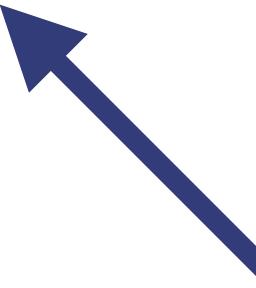
Hotel



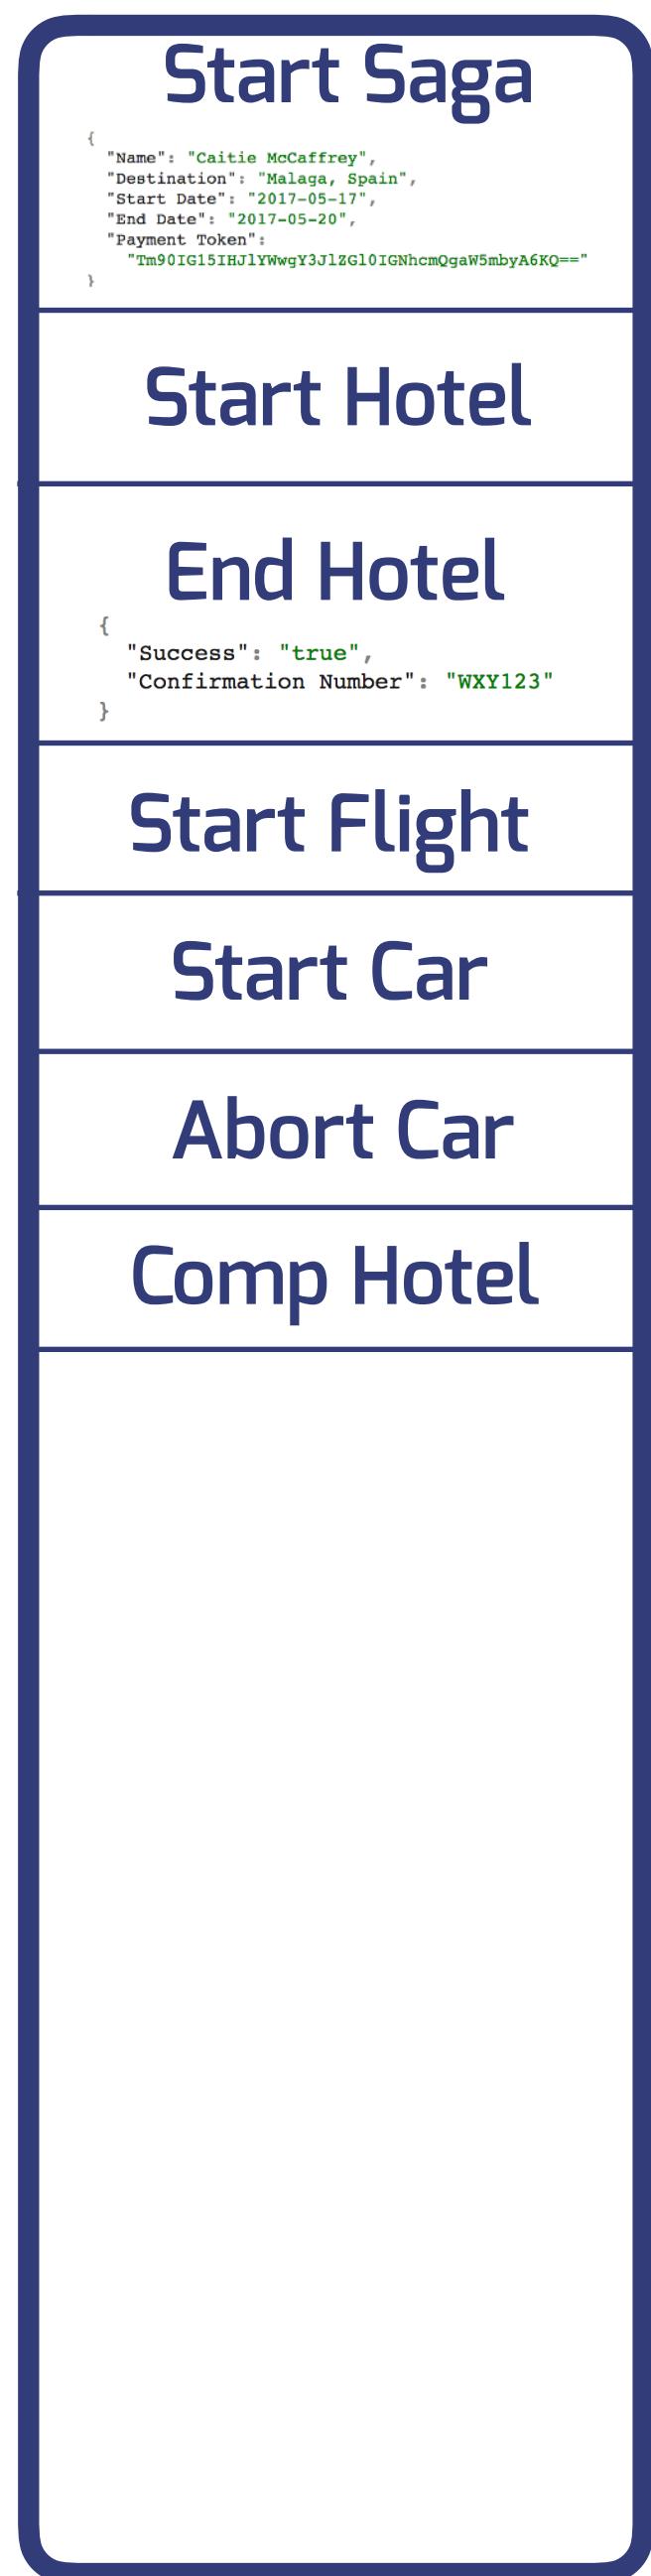
Flight



Payment



Saga Log



{ Start }



End Comp Saga

Start Comp Saga



Car



Hotel

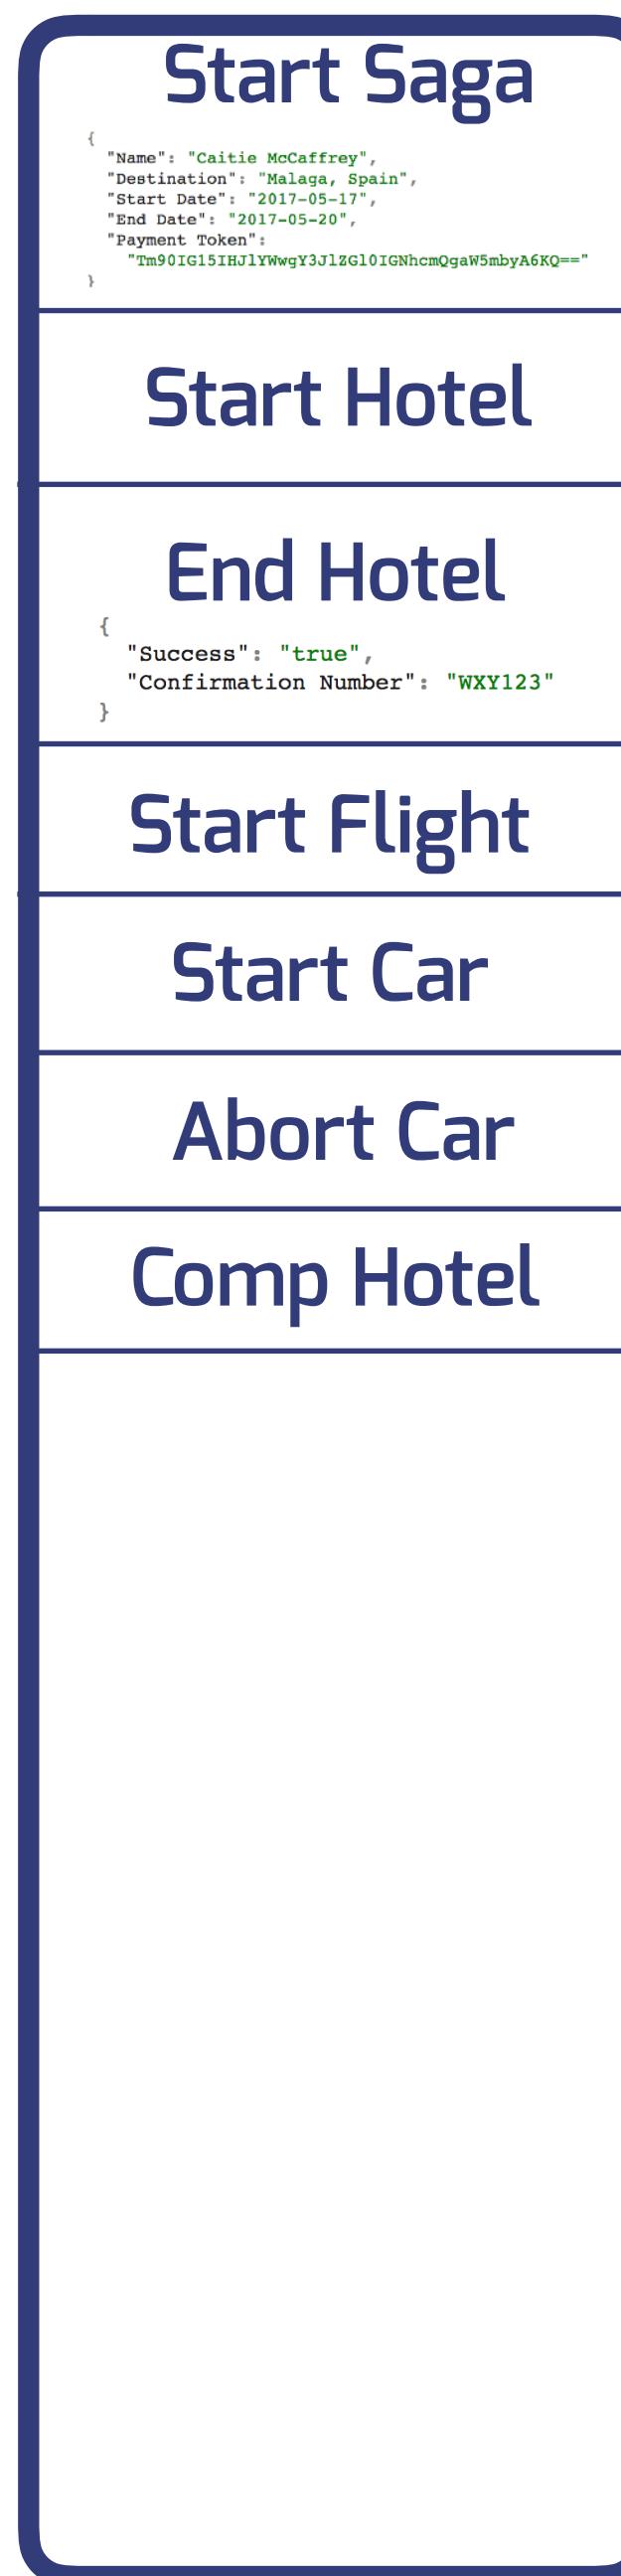


Payment



Flight

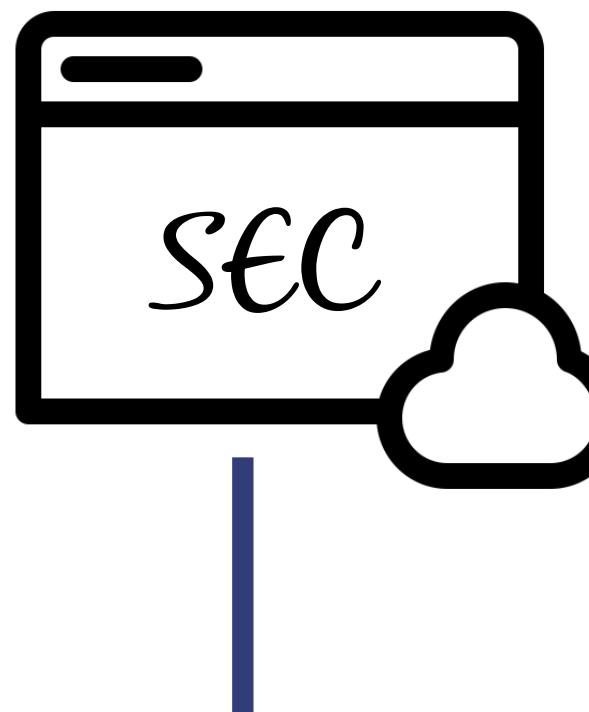
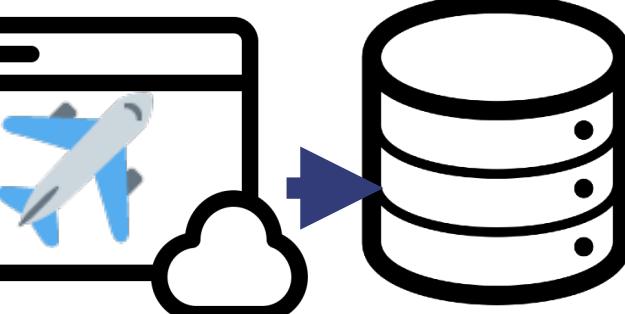
Saga Log



{ Start }

Book Flight Request

```
{  
  "Name": "Caitie McCaffrey",  
  "Destination": "Malaga, Spain",  
  "Start Date": "2017-05-17",  
  "End Date": "2017-05-20"  
}
```



End Comp Saga



Car



Hotel



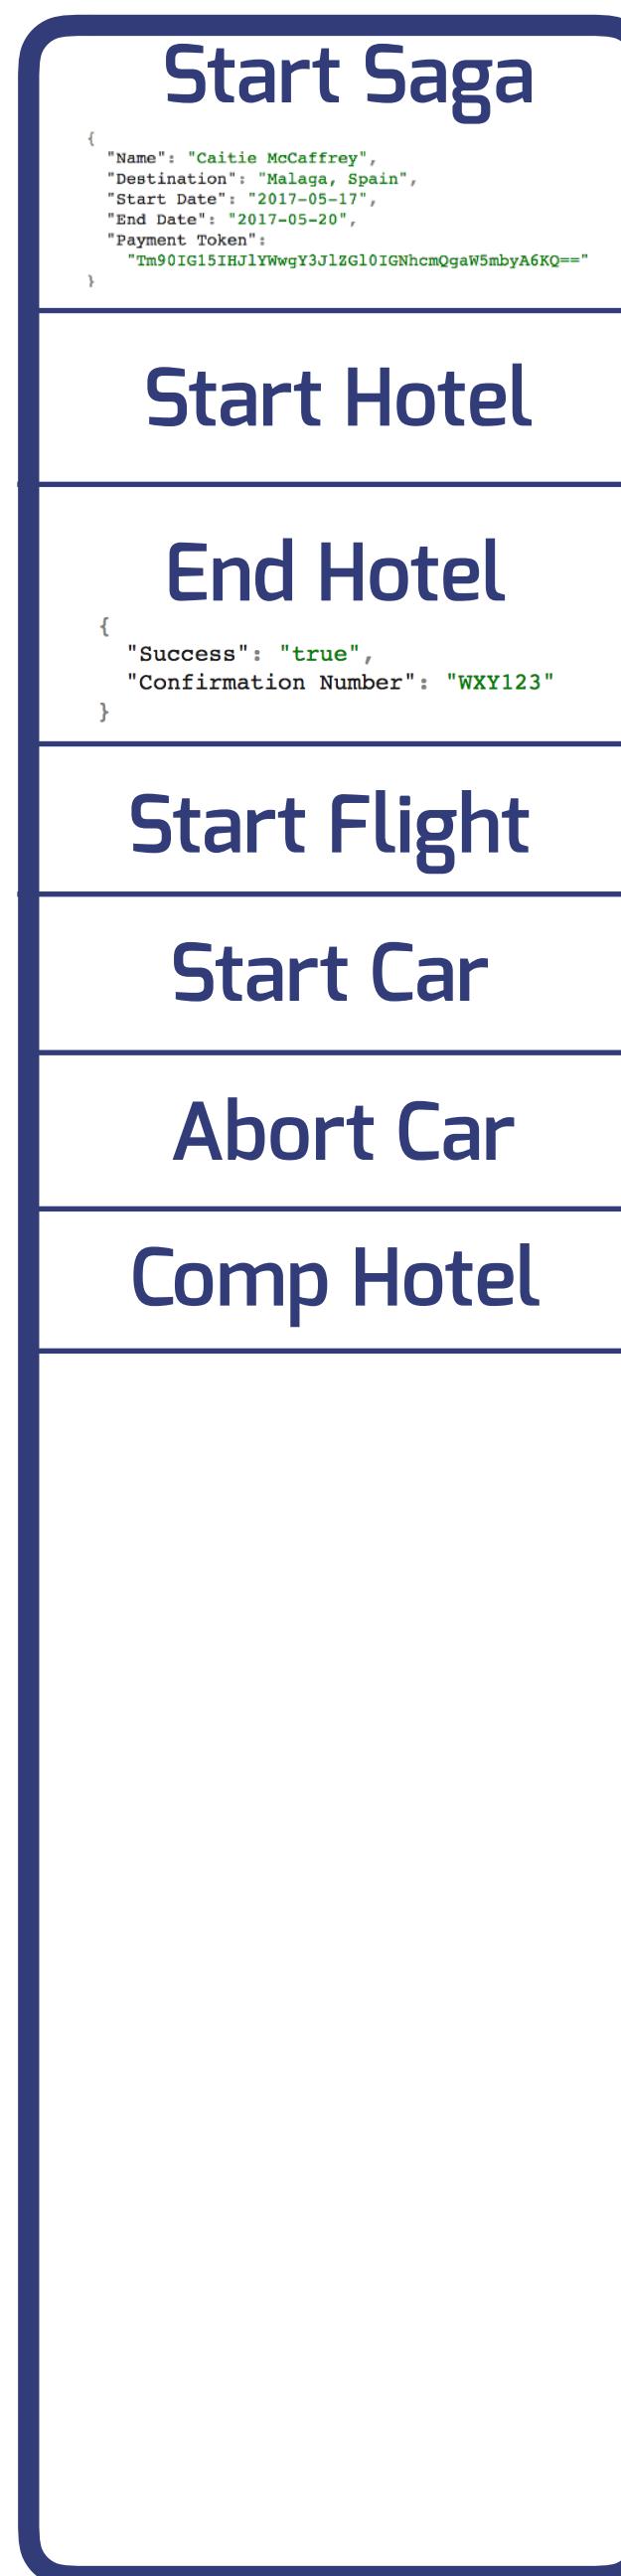
Flight



Payment

Start Comp Saga

Saga Log

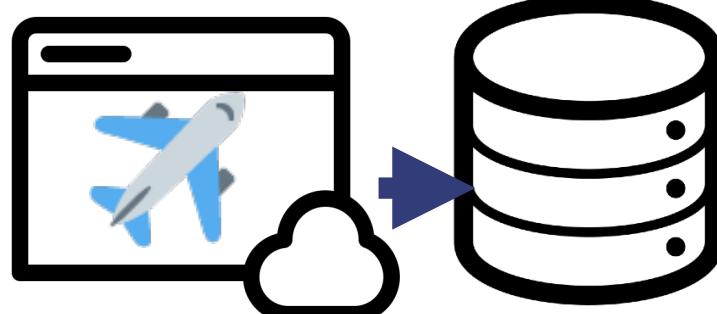


End Flight



Book Flight Response

```
{  
  "Success": "true",  
  "Confirmation Number": "789QPZ"  
}
```



End Comp Saga



Car



Hotel



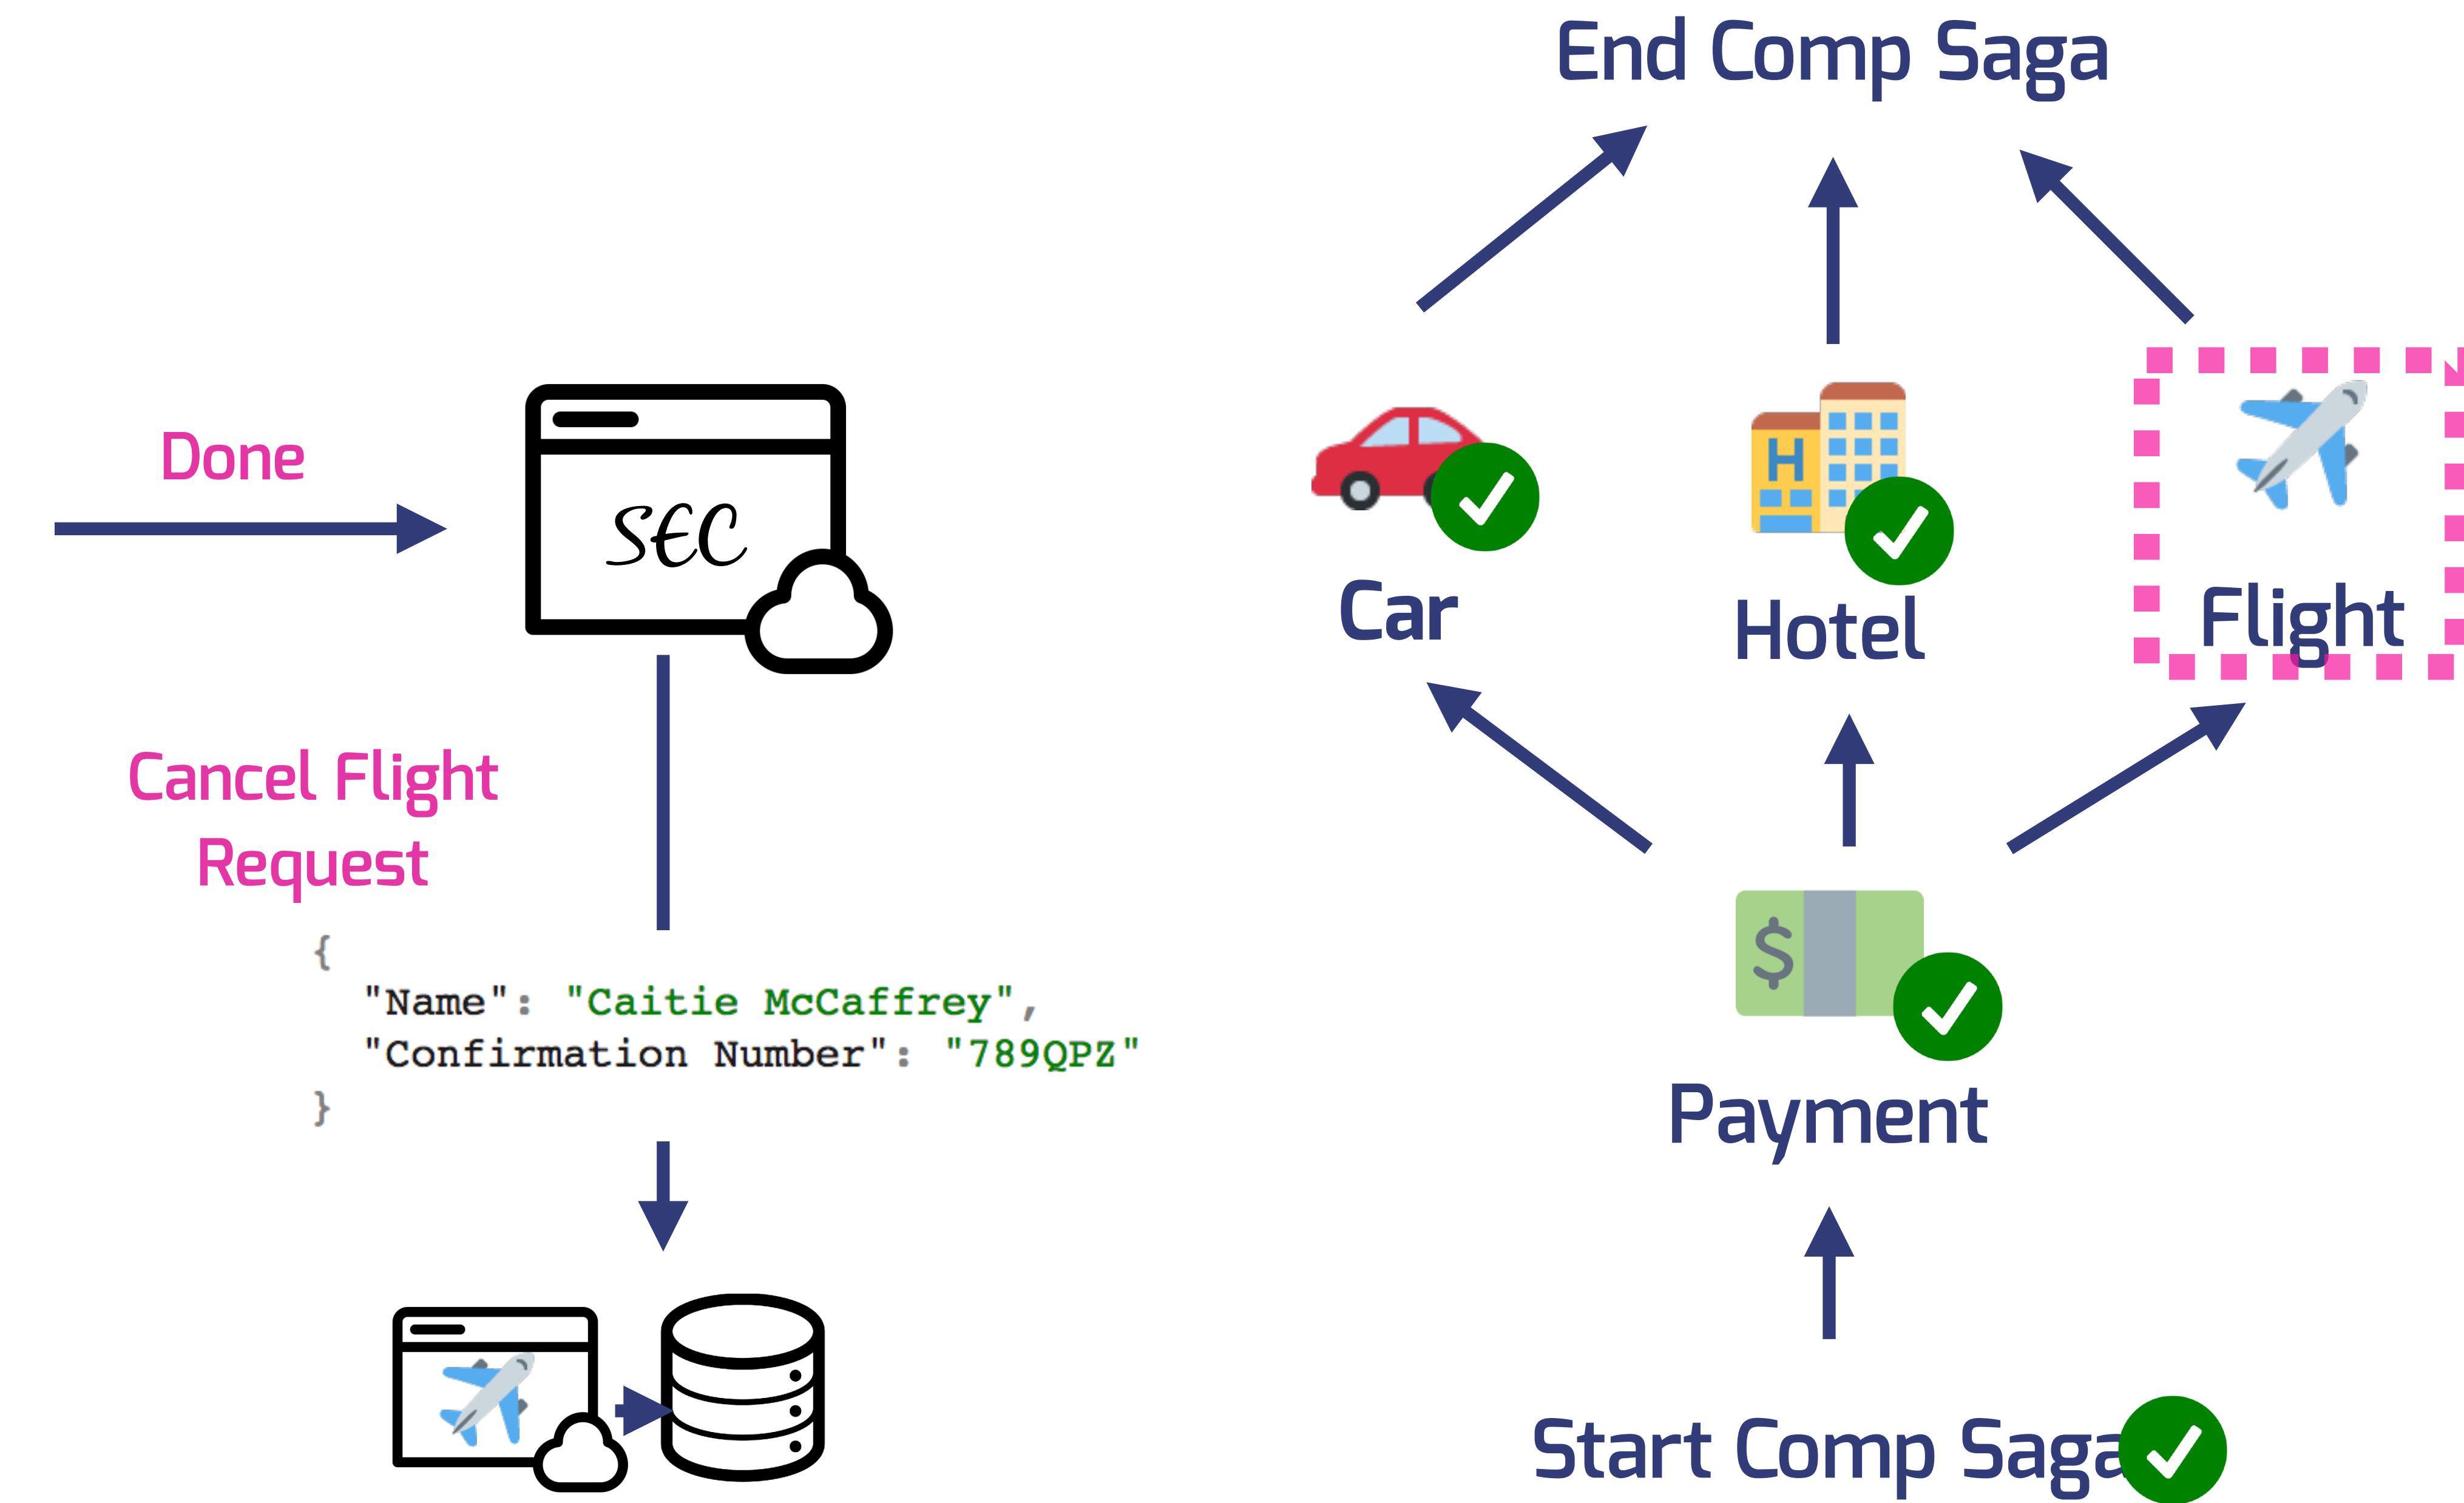
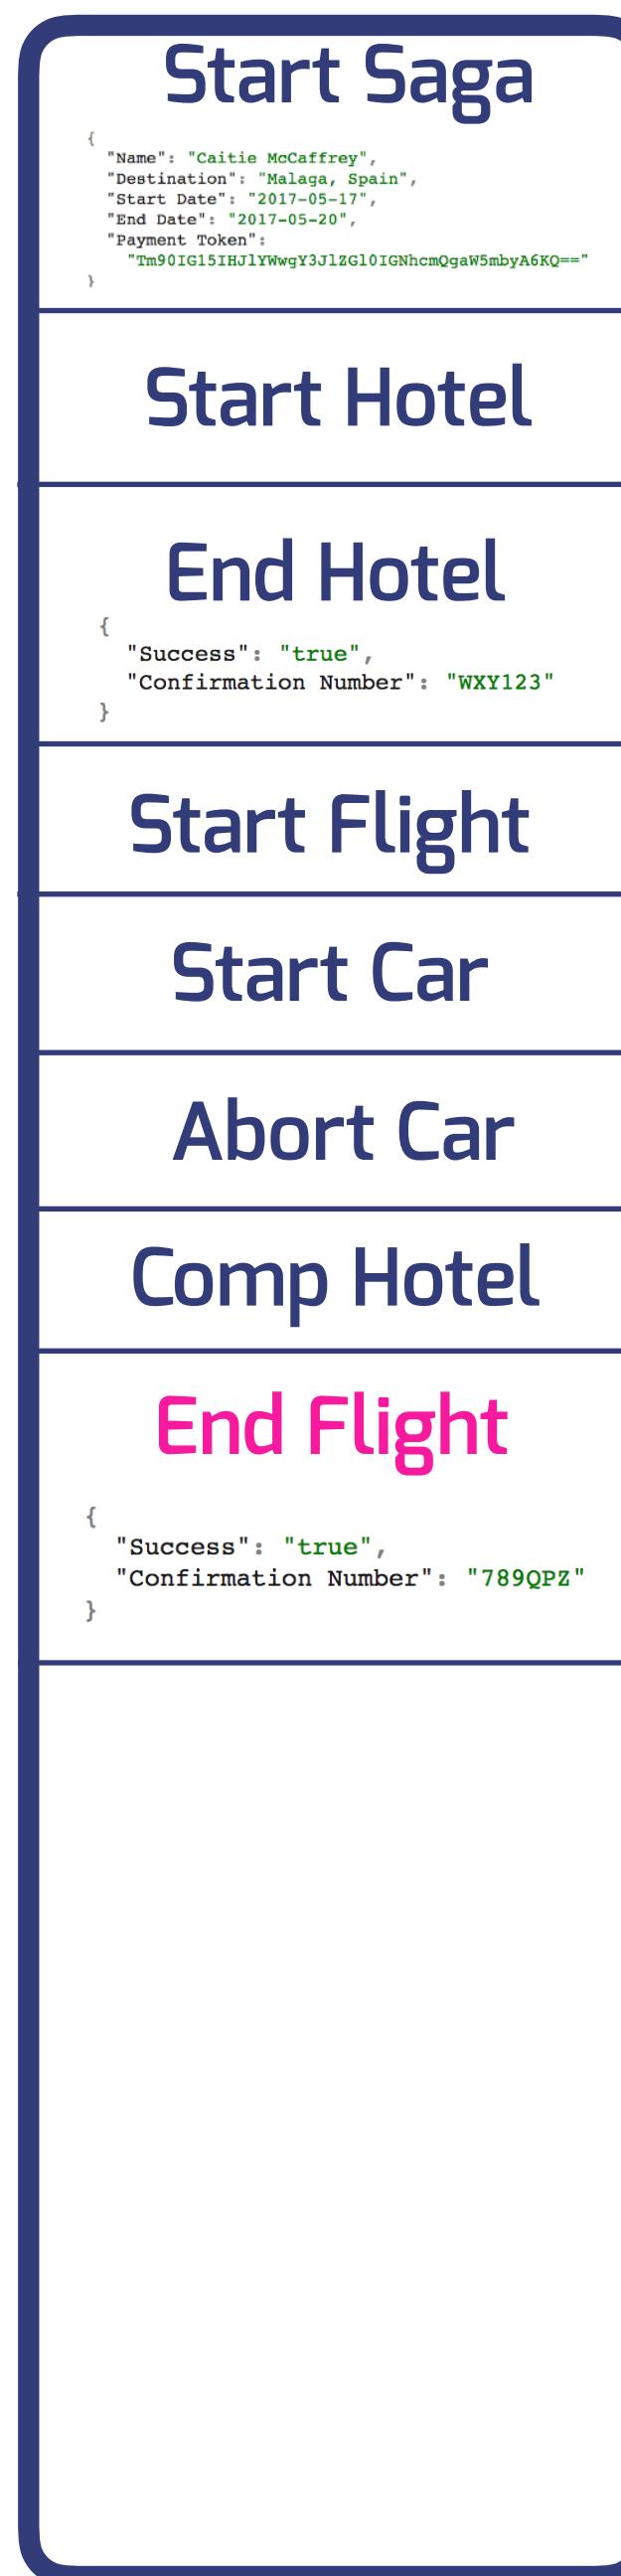
Flight



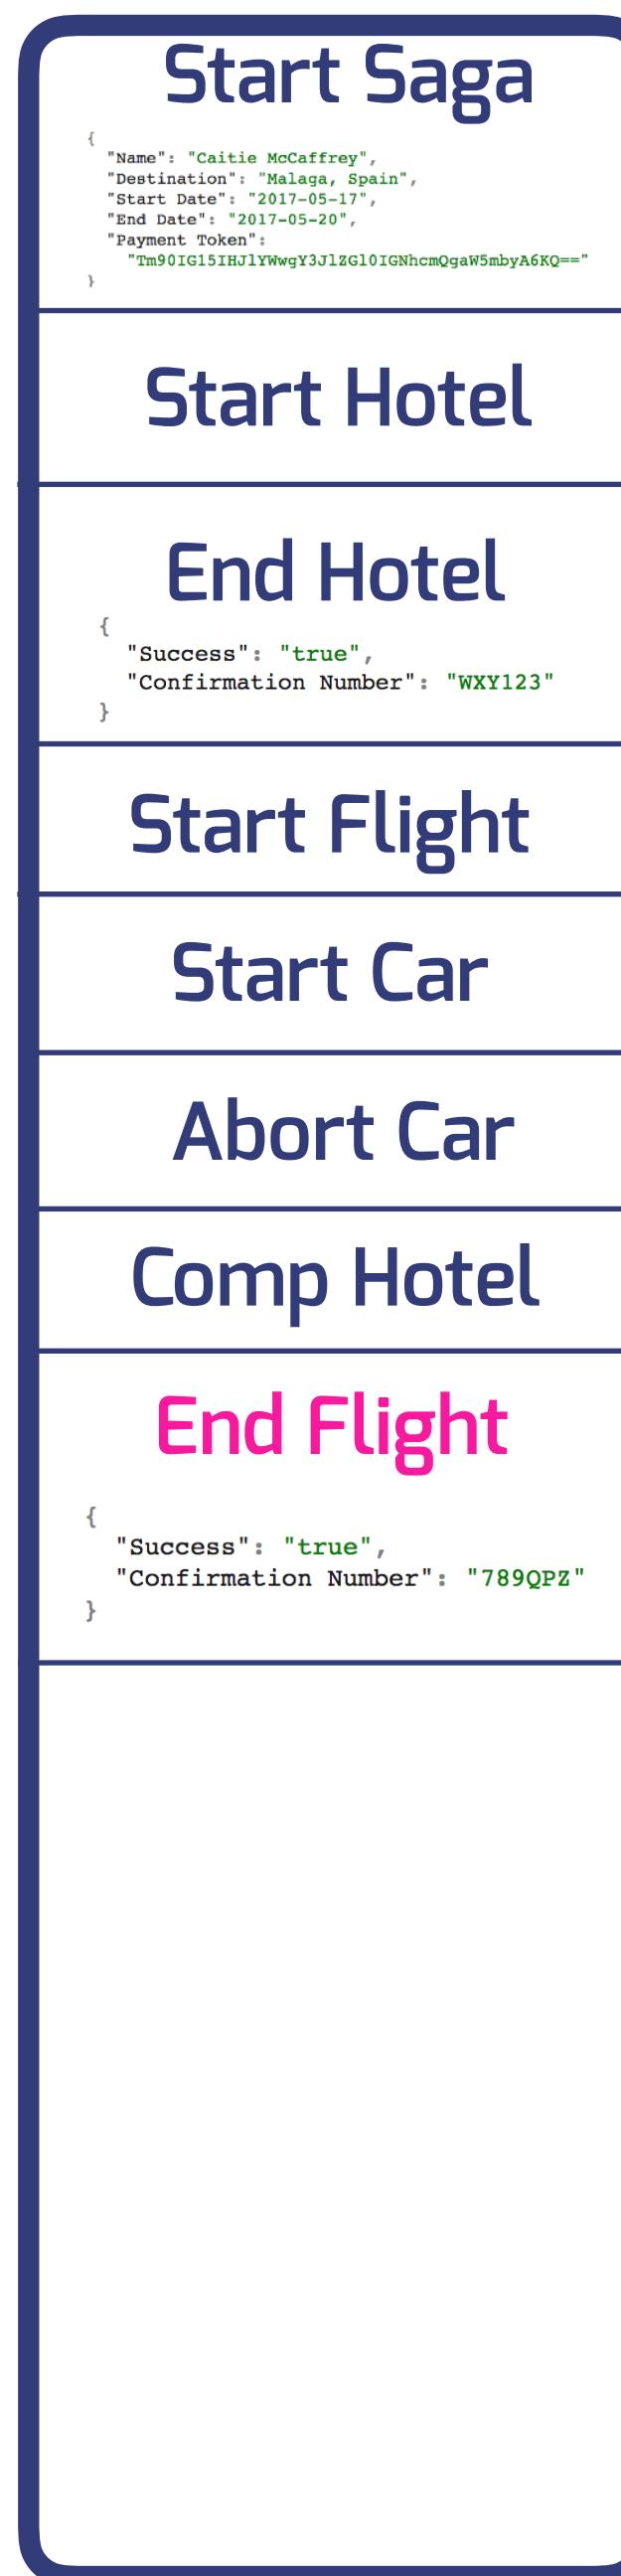
Payment

Start Comp Saga 

Saga Log



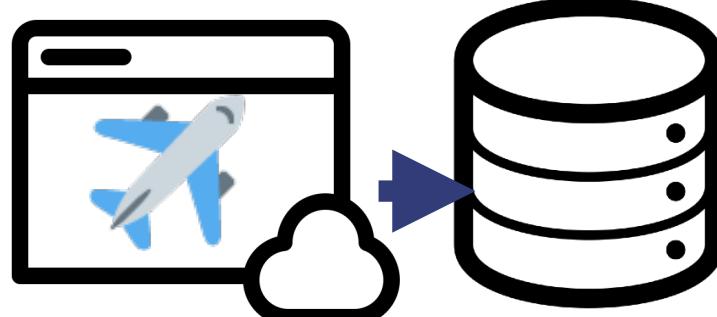
Saga Log



Comp Flight

Cancel Flight Response

```
{  
  "Success": "true",  
  "Confirmation Number": "789QPZ"  
}
```



End Comp Saga

Start Comp Saga 



Car



Hotel

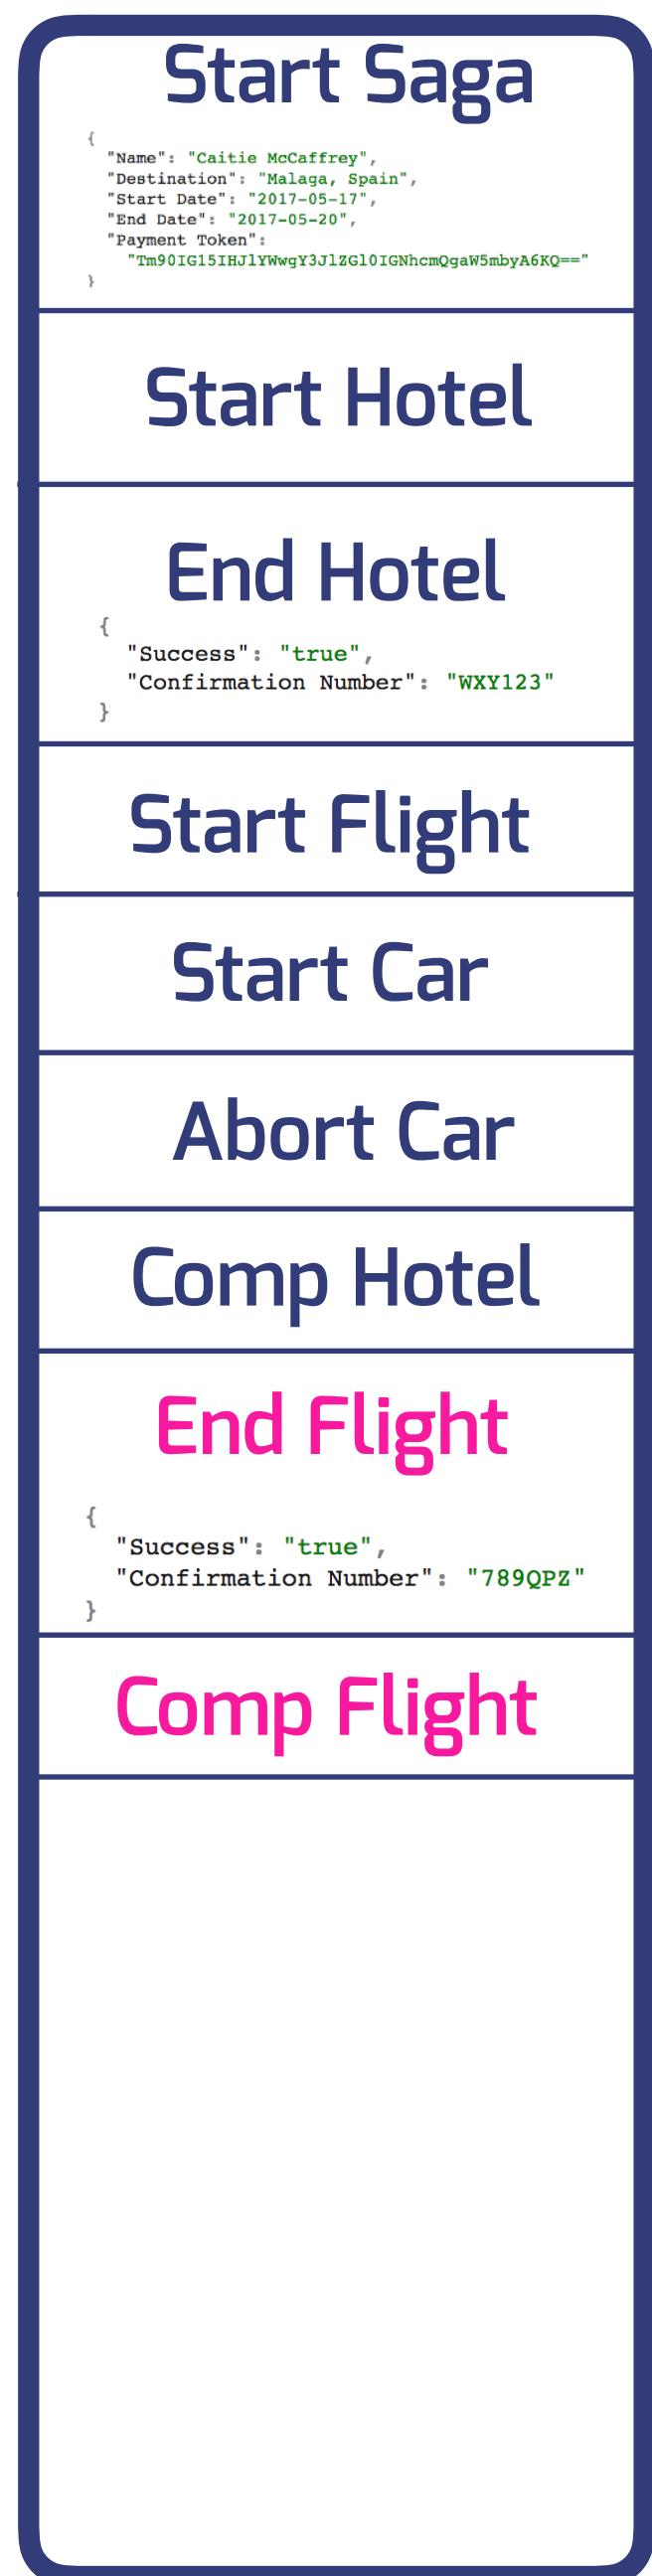


Payment

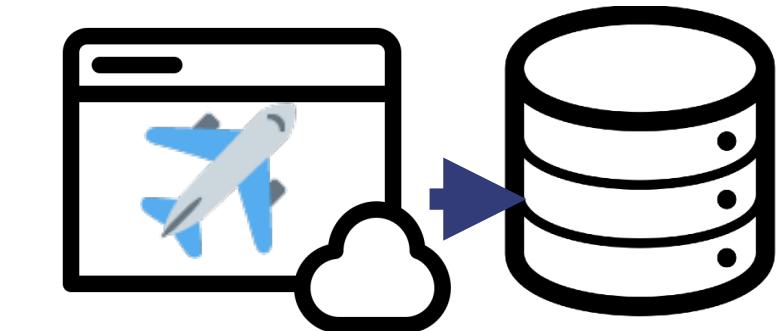


Flight

Saga Log



Done



End Comp Saga

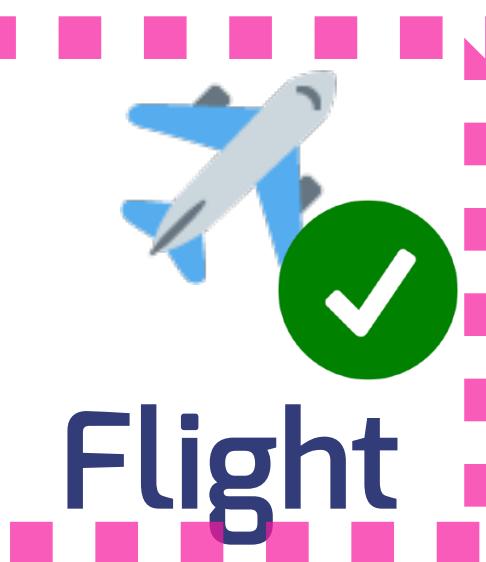
Start Comp Saga



Car



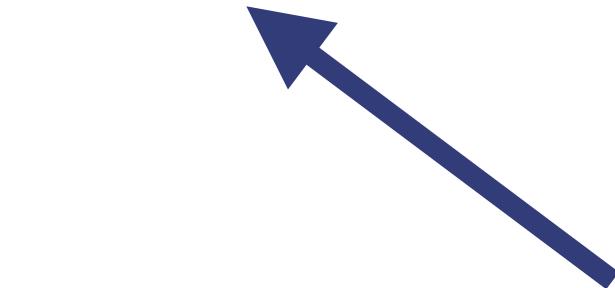
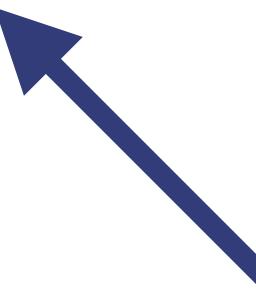
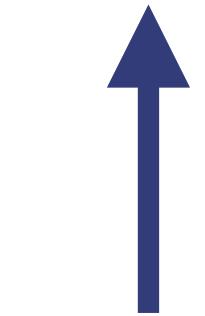
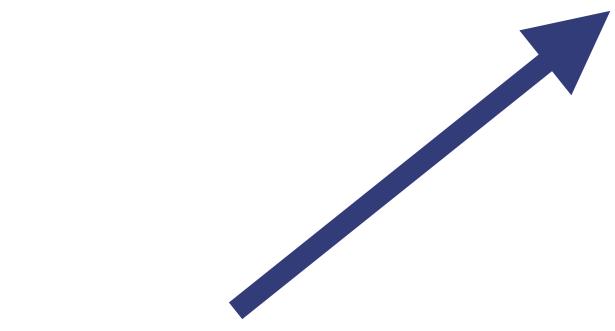
Hotel



Flight



Payment



Saga Log

Start Saga
{ "Name": "Caitie McCaffrey", "Destination": "Malaga, Spain", "Start Date": "2017-05-17", "End Date": "2017-05-20", "Payment Token": "Um90IG15IHJ1YWwgY3J1ZG10IGNhcmQgaW5mbvA6KQ==" }
Start Hotel
End Hotel
{ "Success": "true", "Confirmation Number": "WXY123" }
Start Flight
Start Car
Abort Car
Comp Hotel
End Flight
{ "Success": "true", "Confirmation Number": "789QPZ" }
Comp Flight

End Saga



Saga Log

Start Saga
{ "Name": "Caitie McCaffrey", "Destination": "Malaga, Spain", "Start Date": "2017-05-17", "End Date": "2017-05-20", "Payment Token": "Um90IG15IHJ1YWwgY3J1ZG10IGNhcmQgaW5mbvA6KQ==" }
Start Hotel
End Hotel
{ "Success": "true", "Confirmation Number": "WXY123" }
Start Flight
Start Car
Abort Car
Comp Hotel
End Flight
{ "Success": "true", "Confirmation Number": "789QPZ" }
Comp Flight
End Saga

Done



Distributed Saga Guarantee

All requests were completed successfully



Book Hotel



Book Car



Book Flight



Charge Money

Or a subset of requests and the corresponding compensating requests were executed



Book Hotel



Book Car

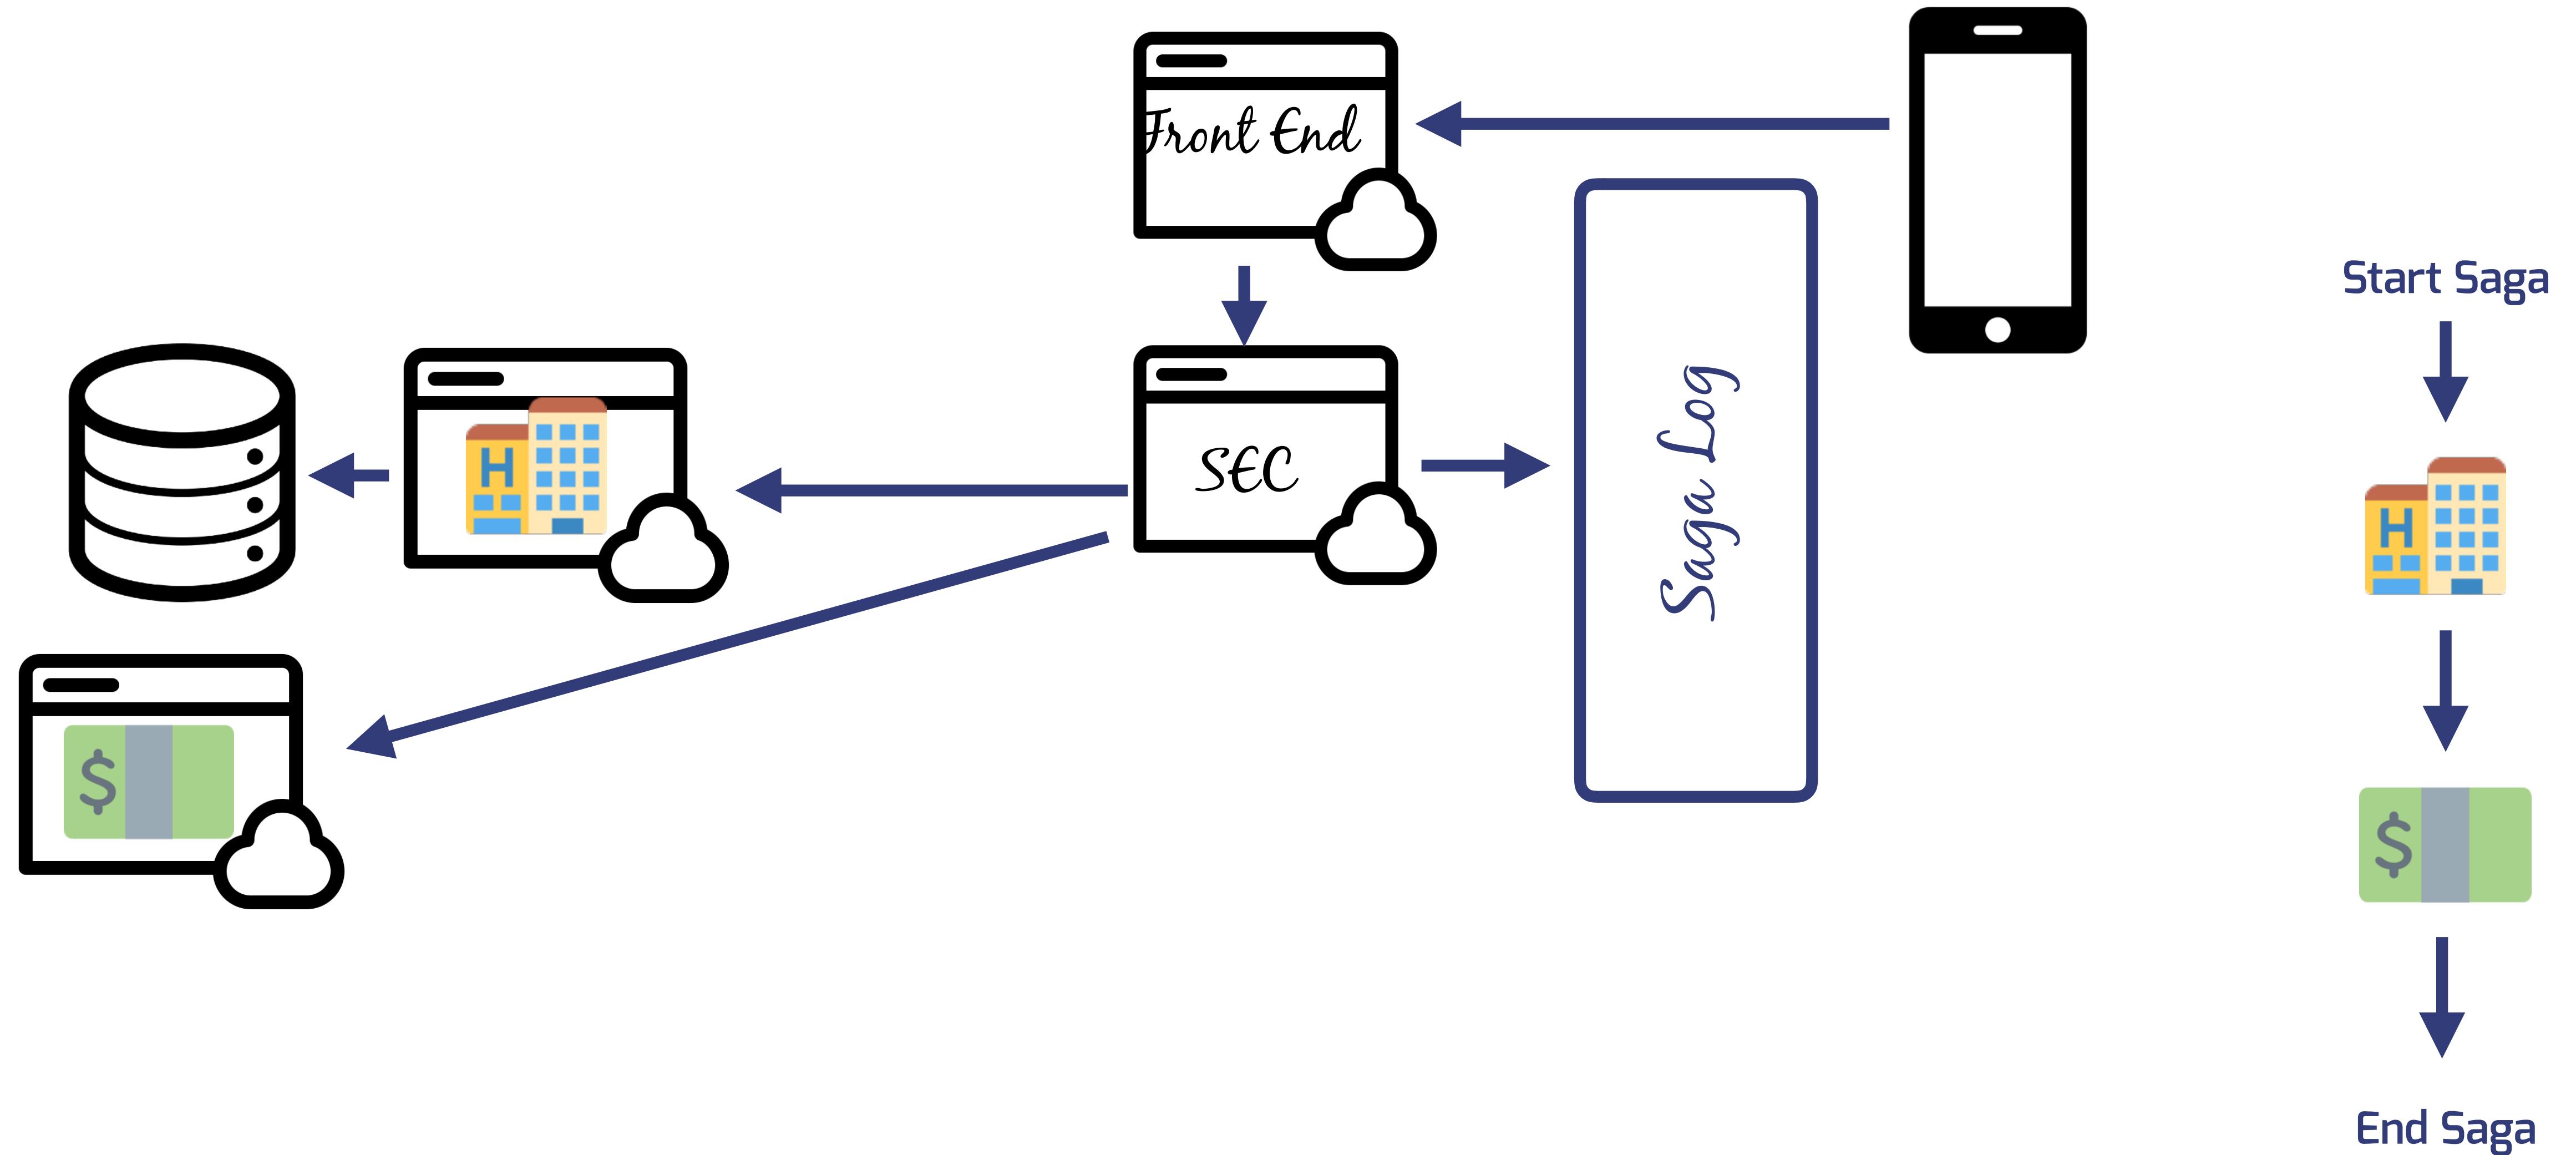


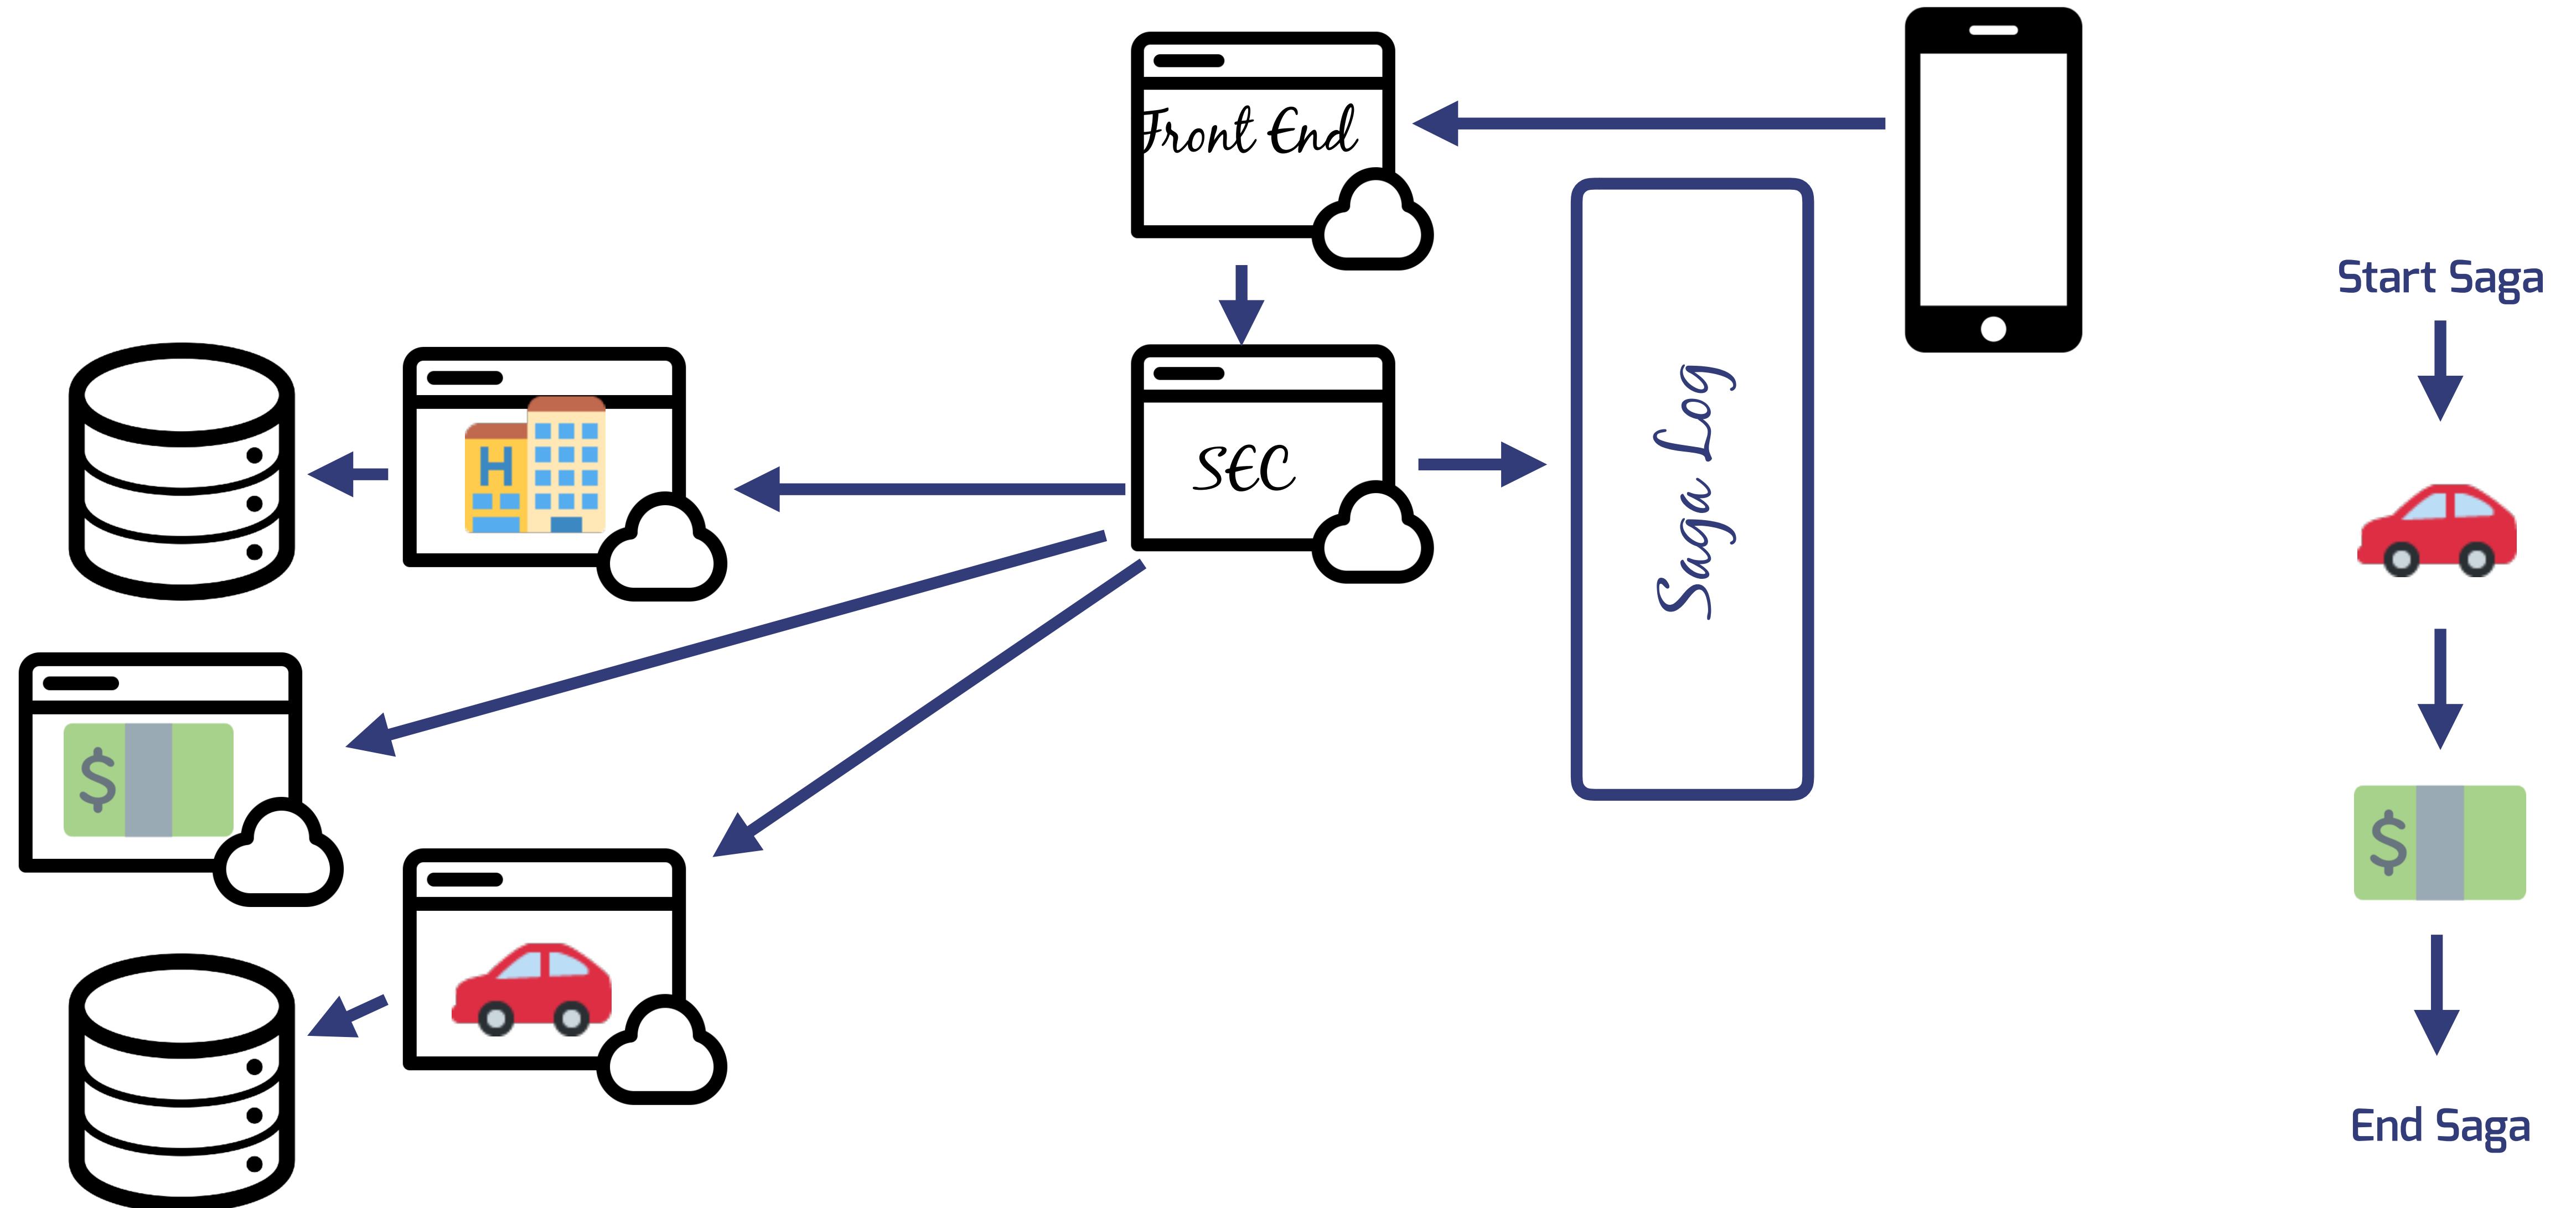
Cancel Hotel

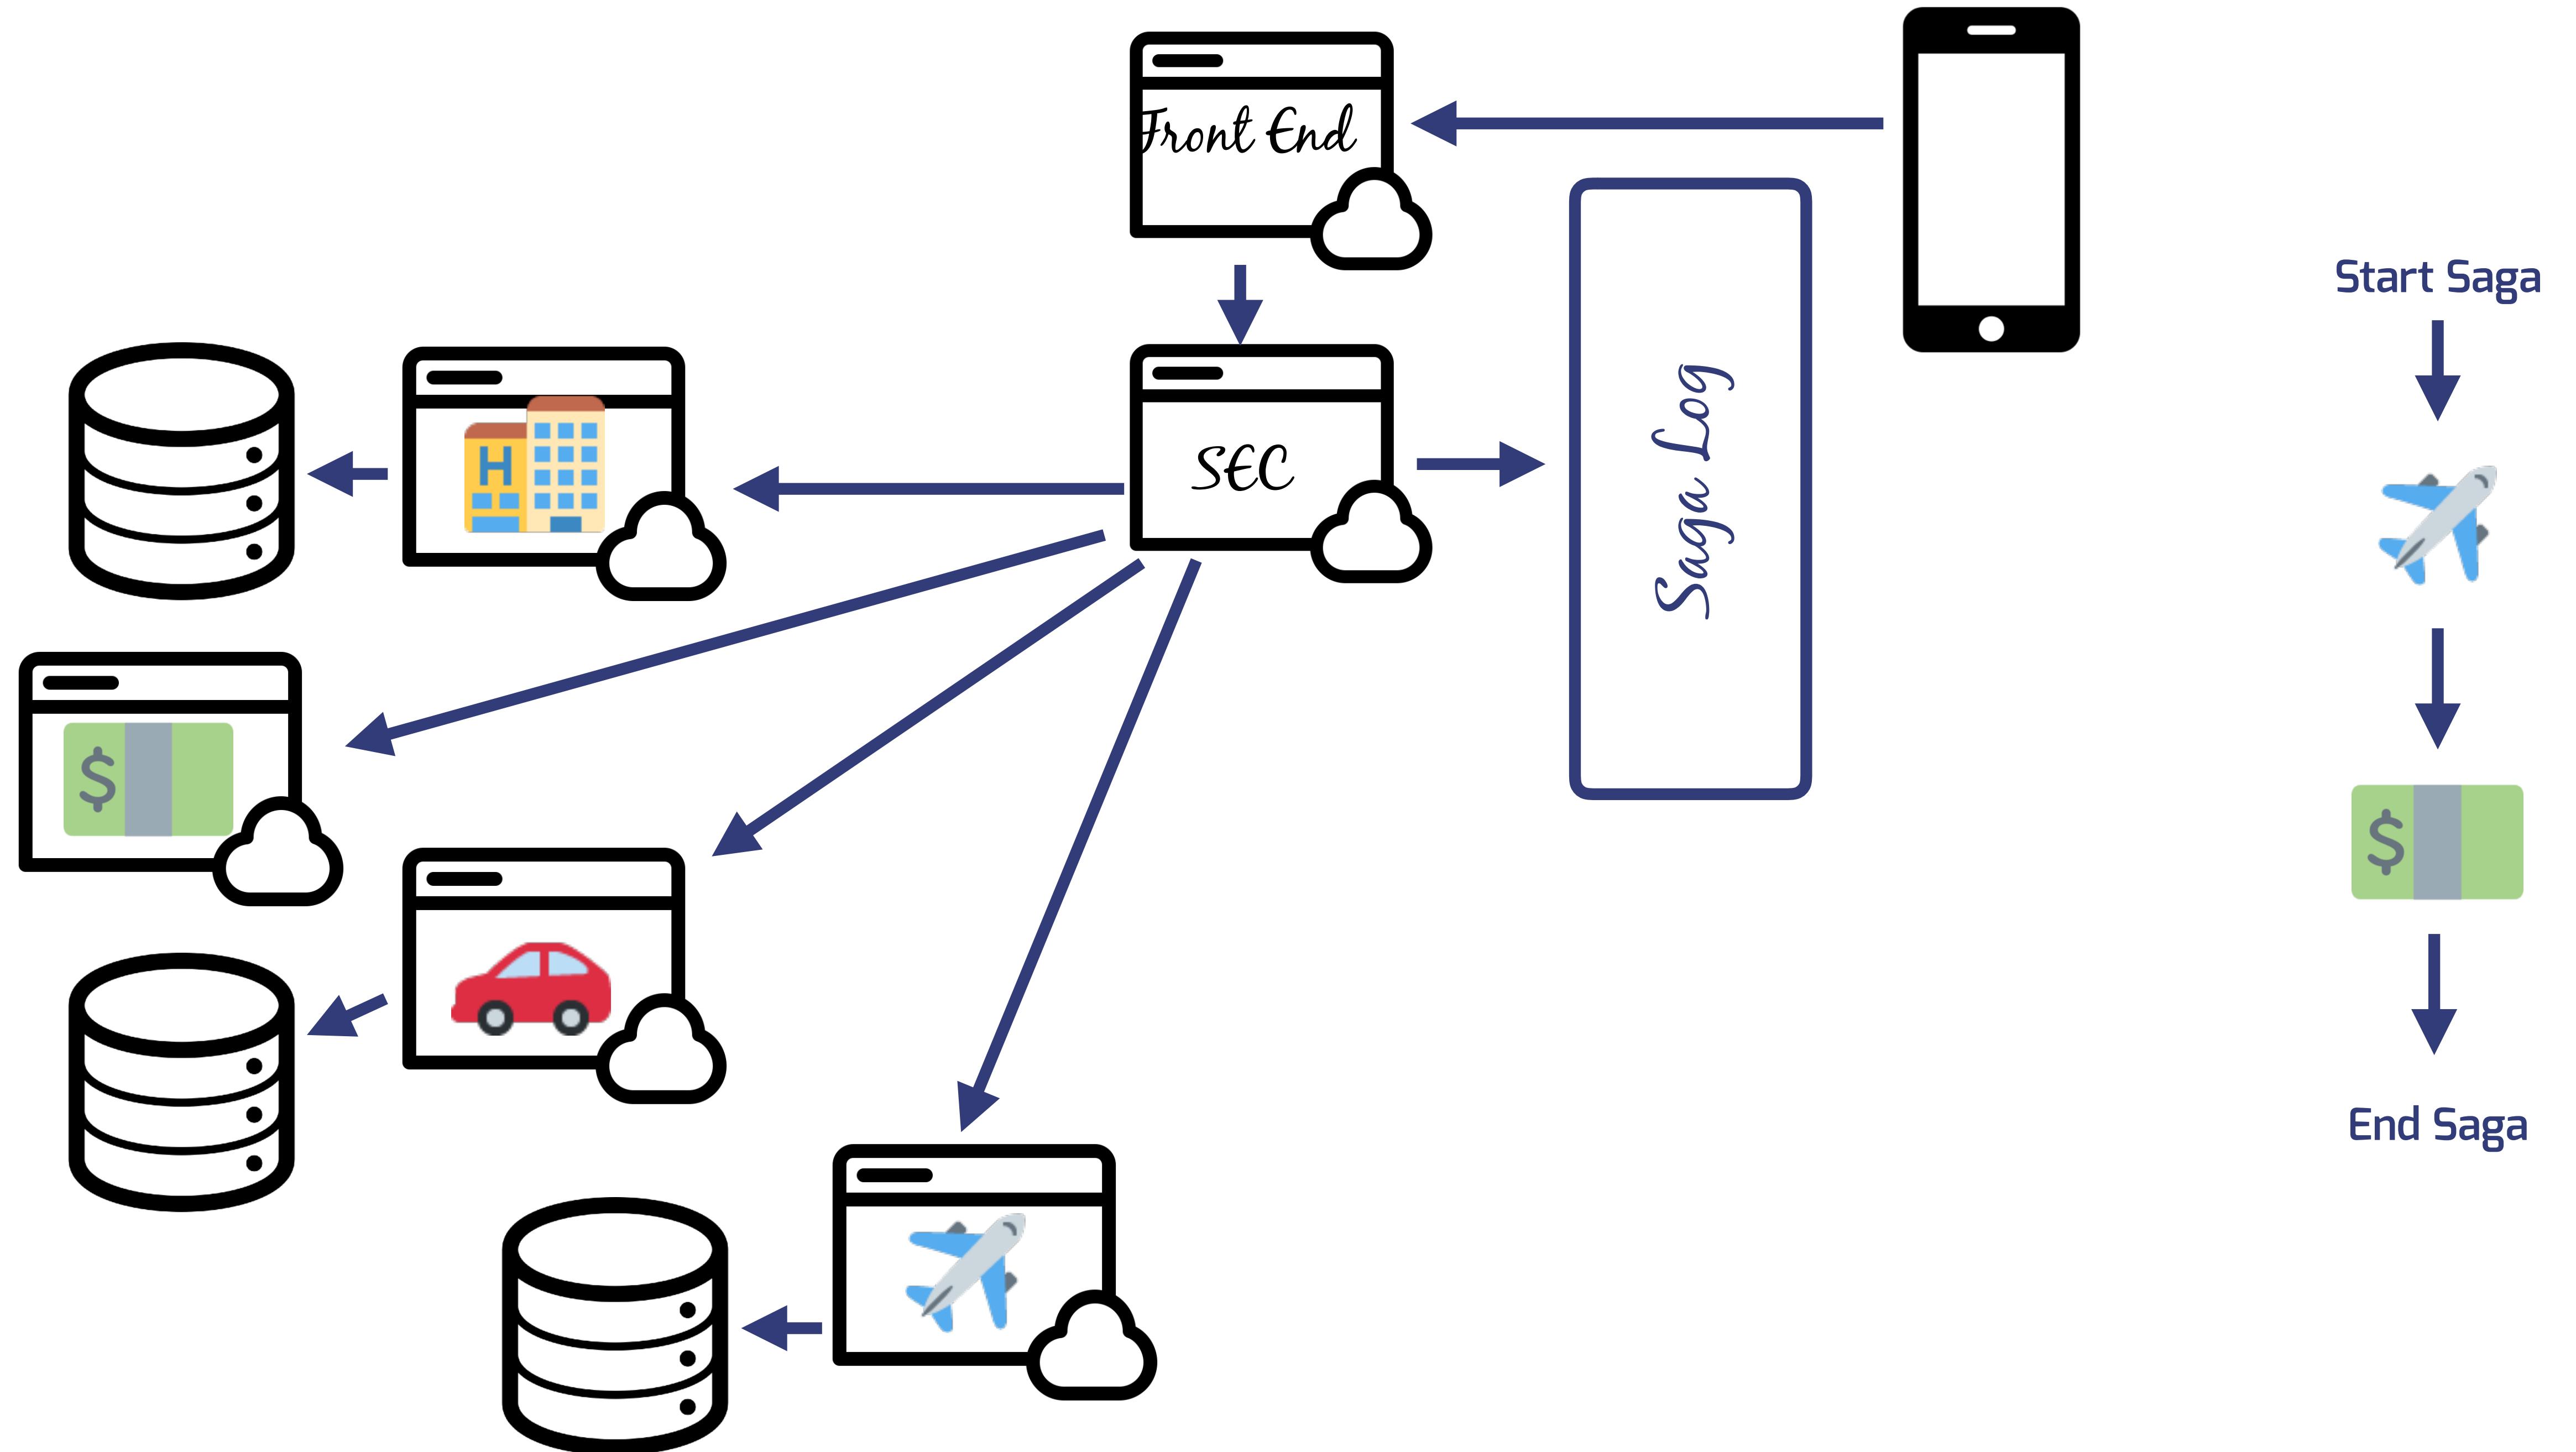


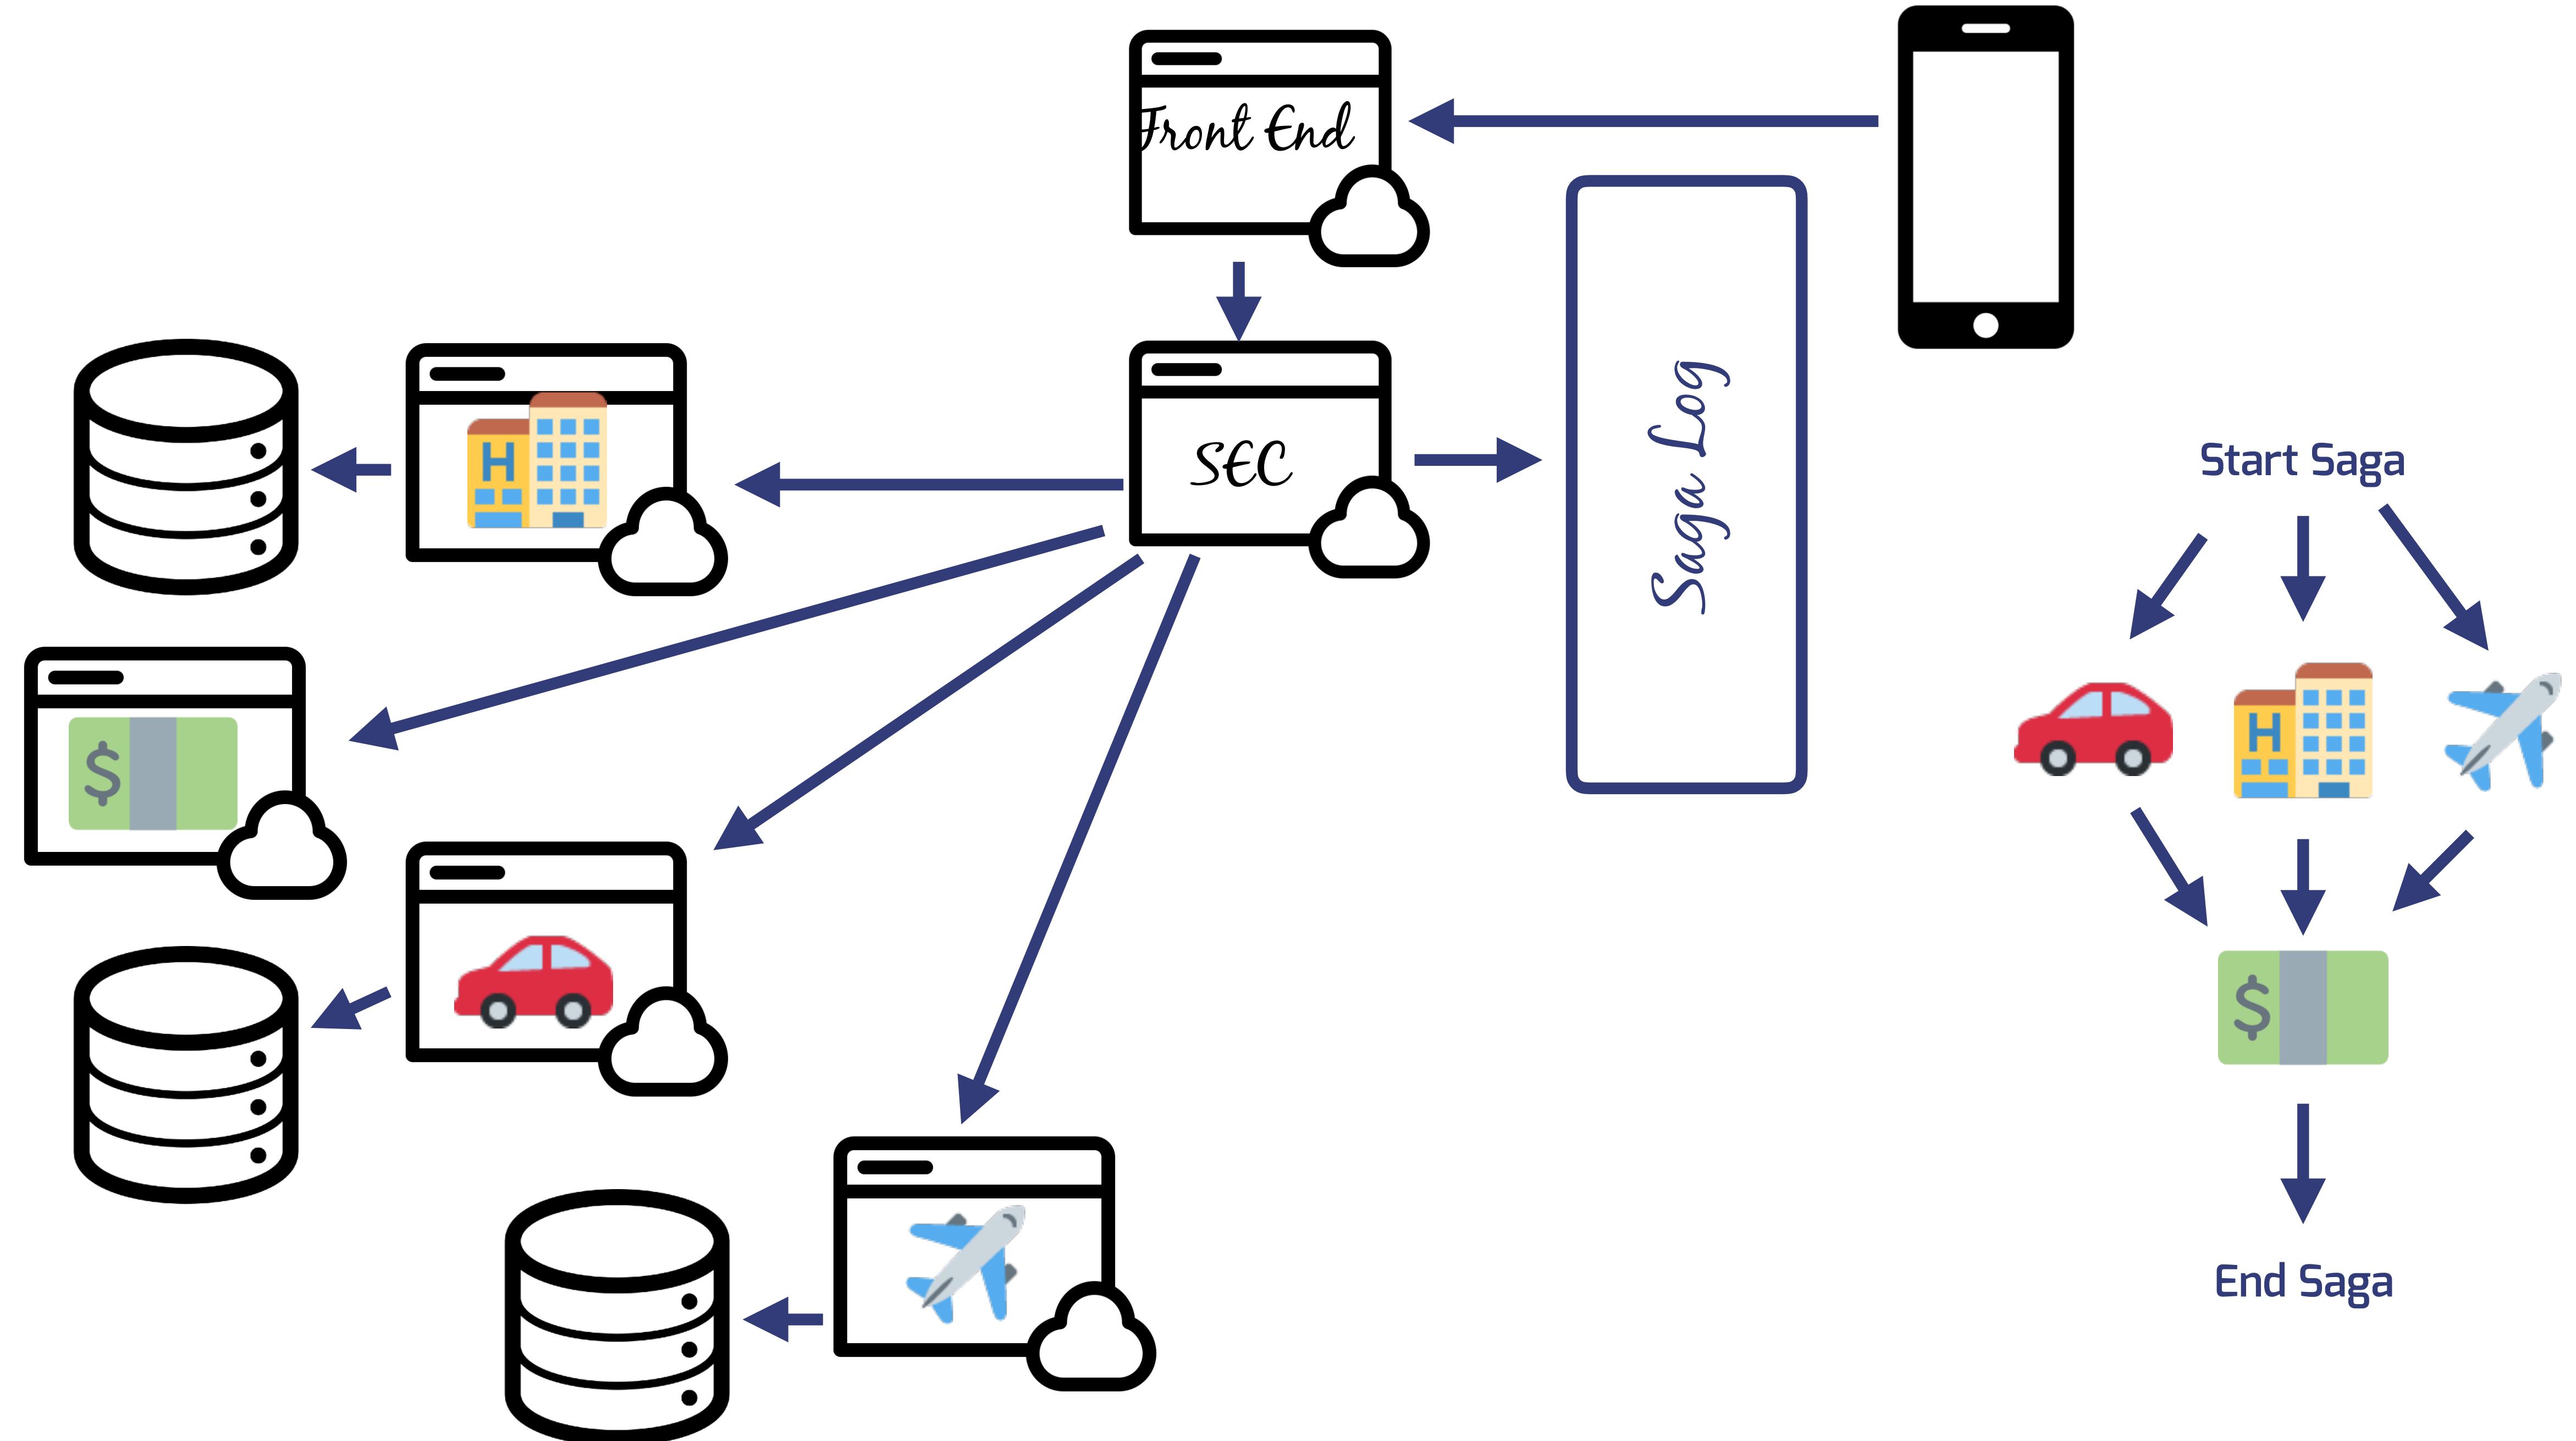
Cancel Car

Recovering from
Saga Execution Coordinator
Failure



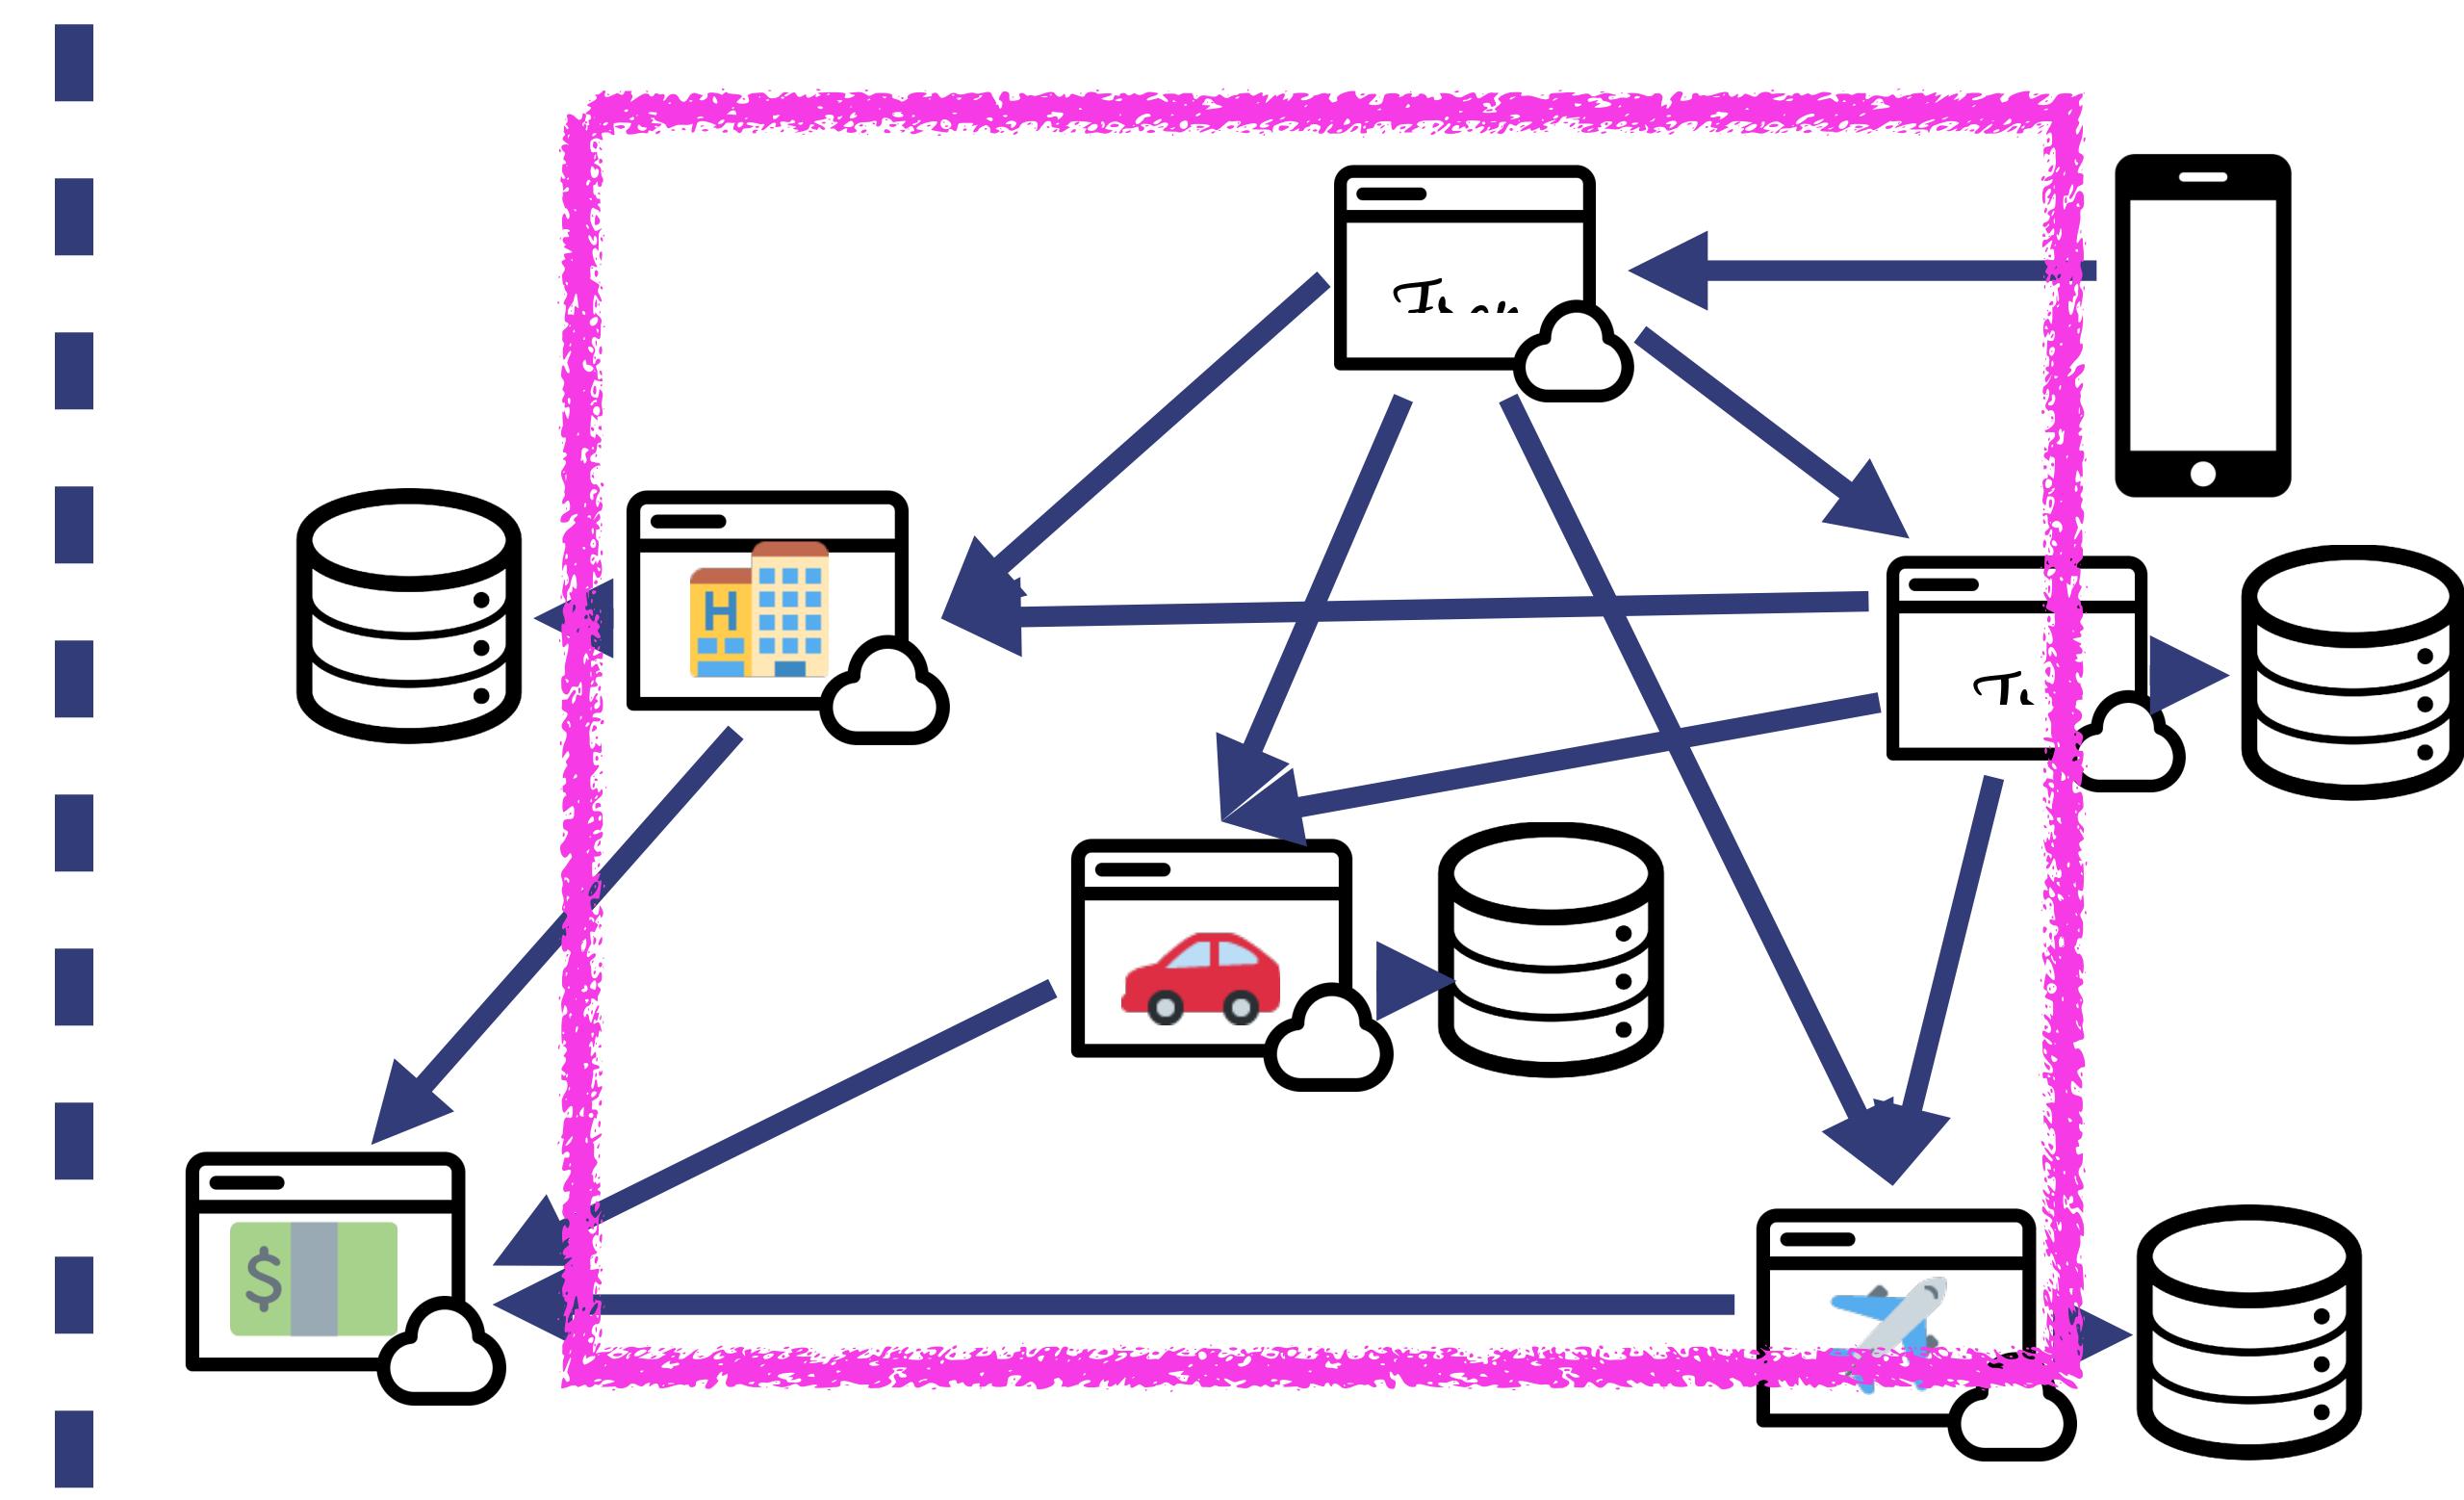
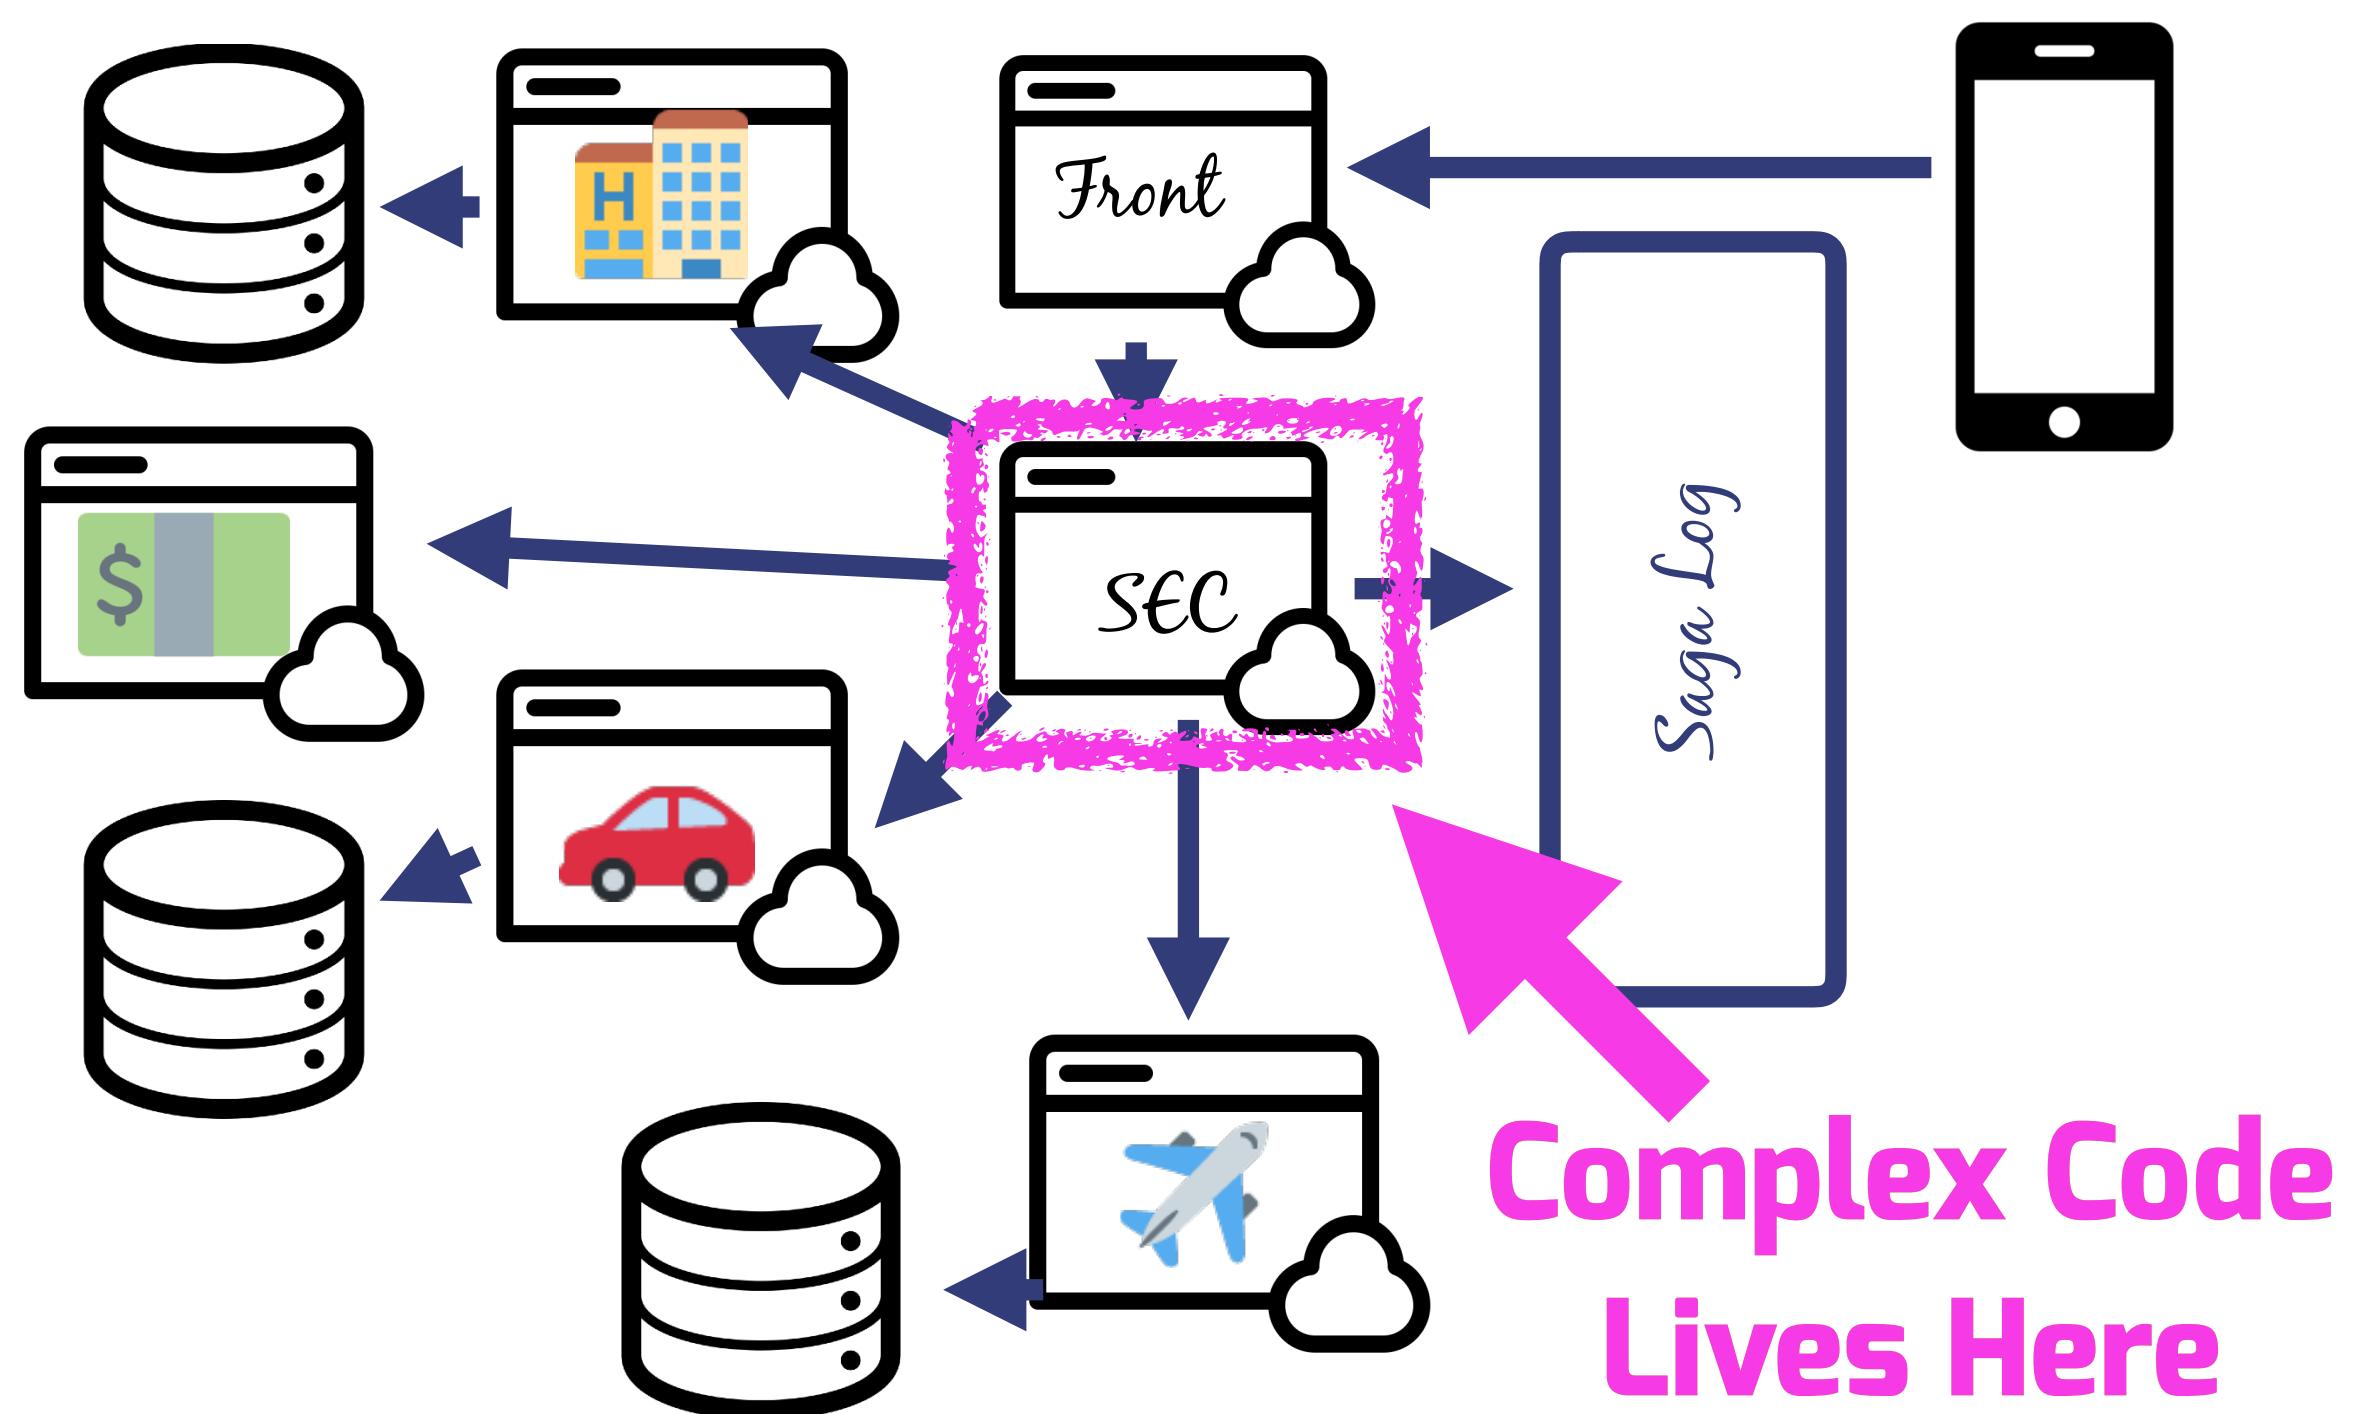






Isolation of Complex Code

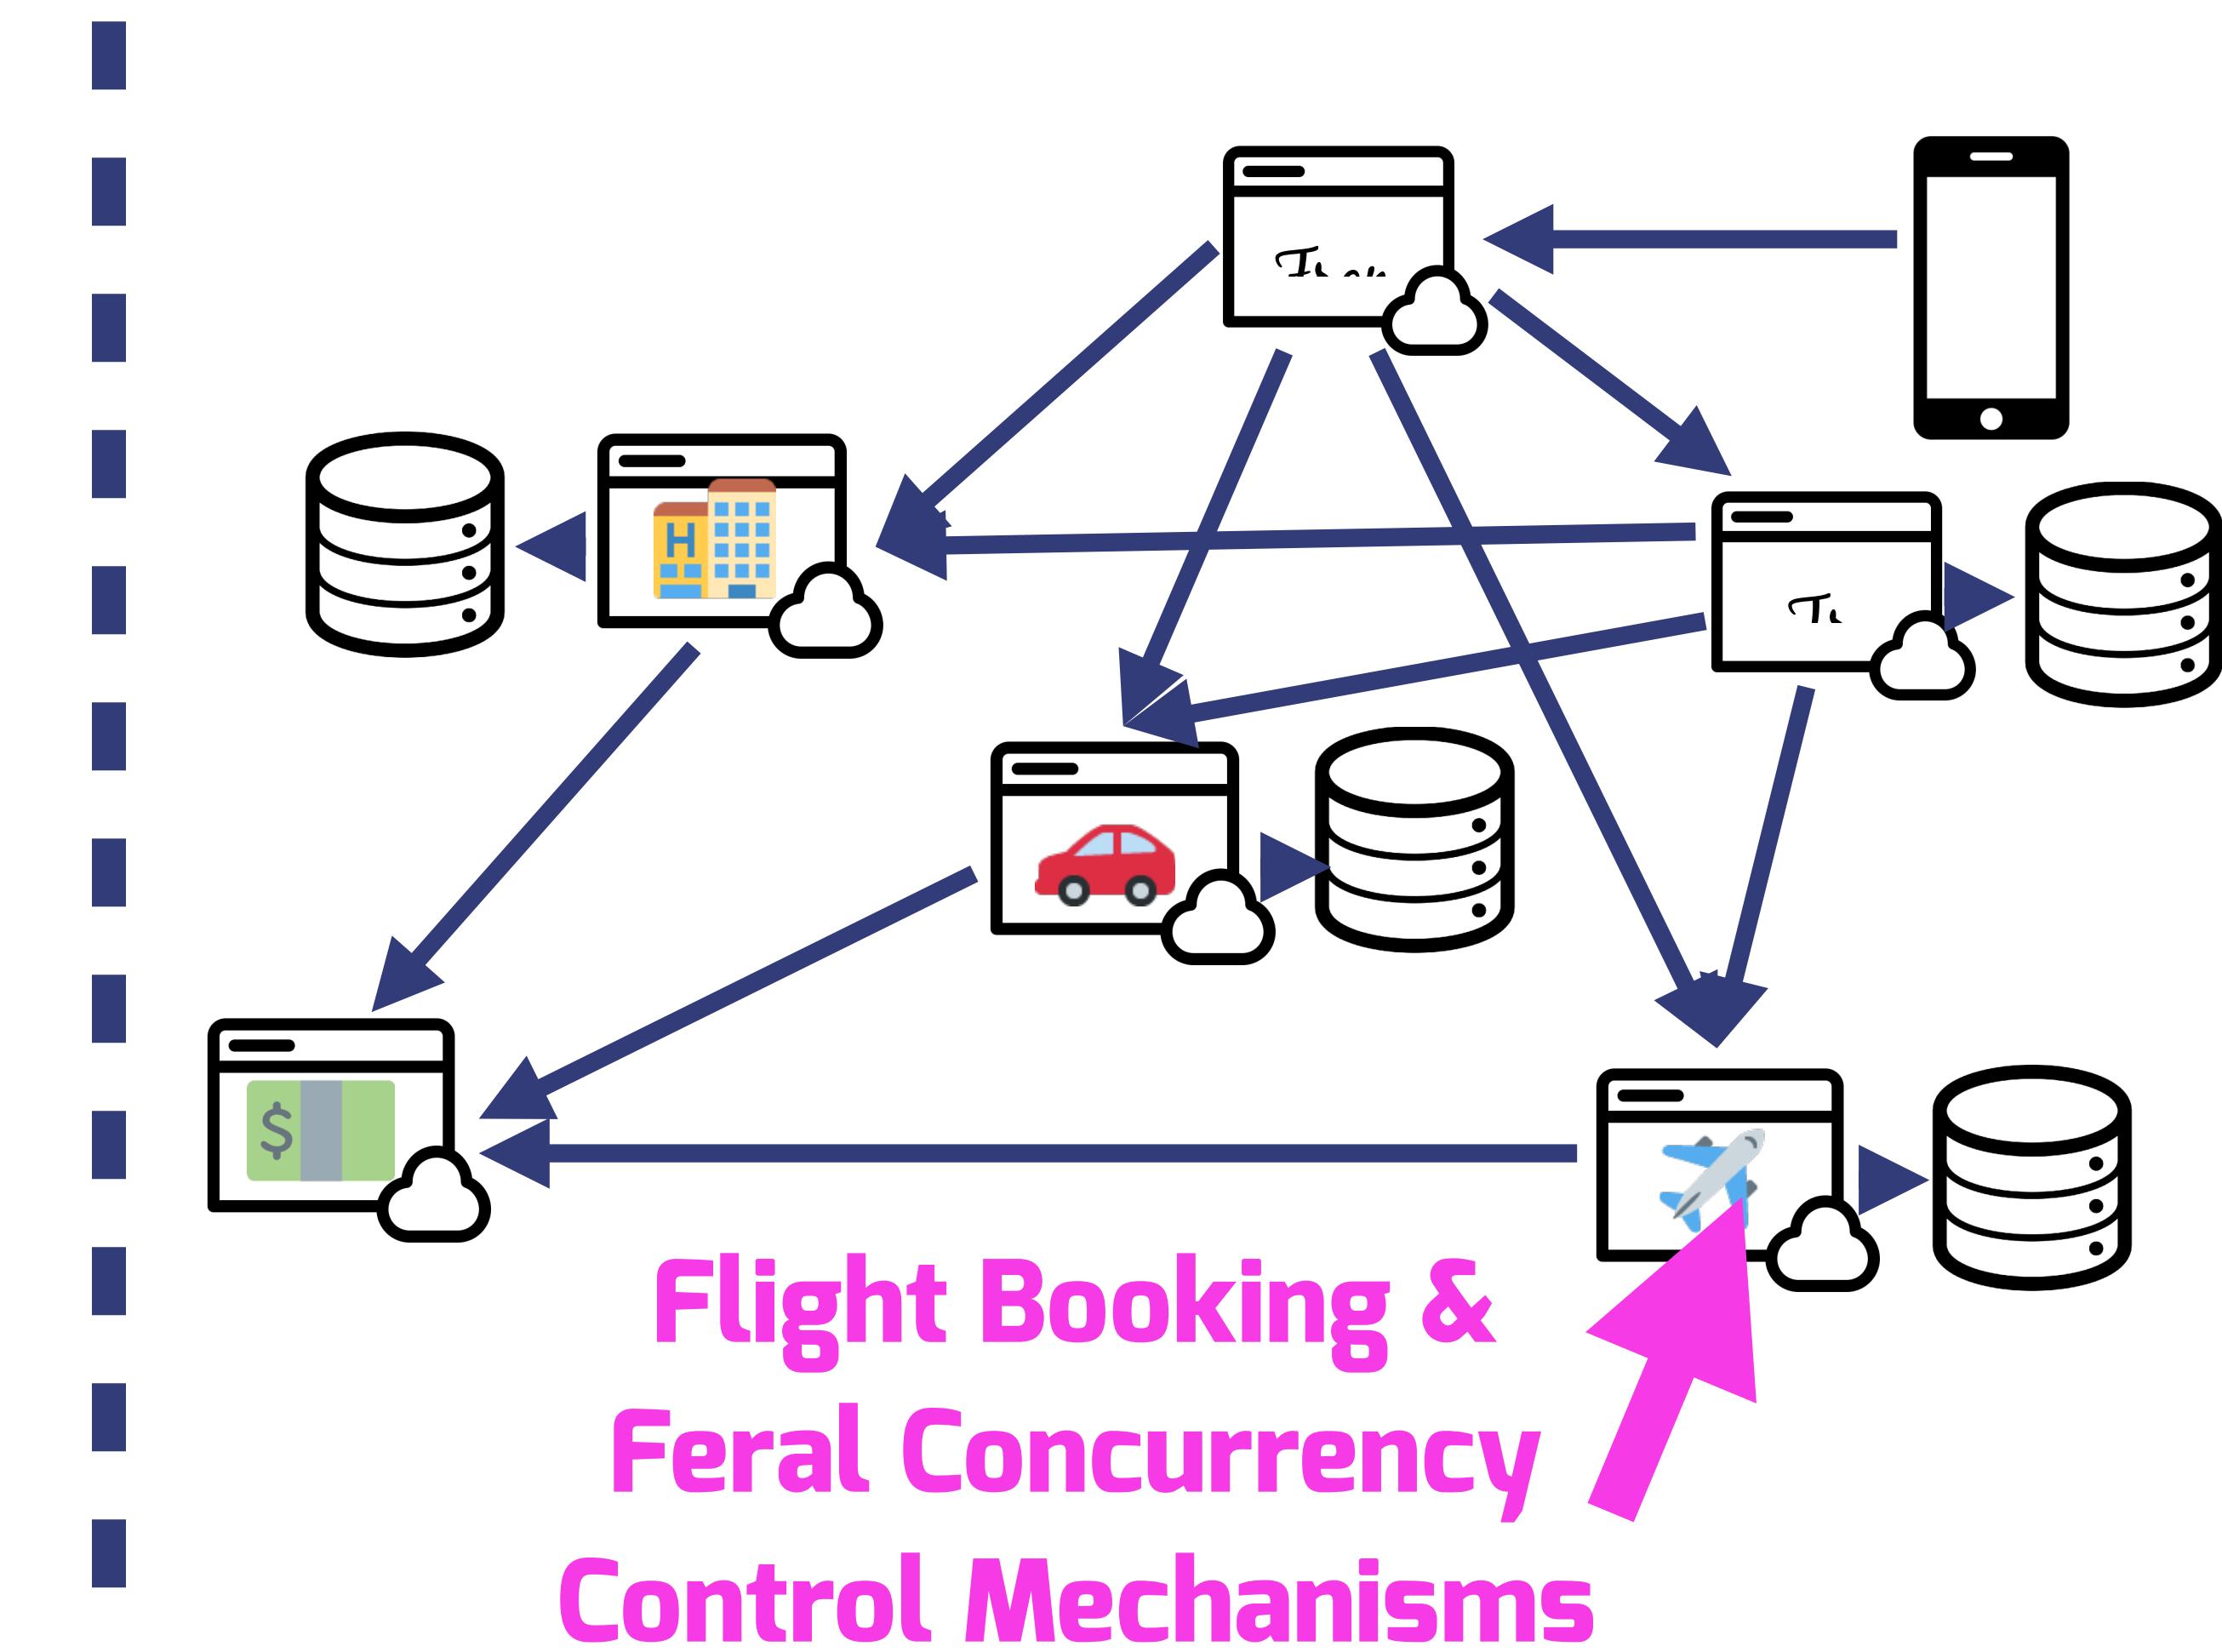
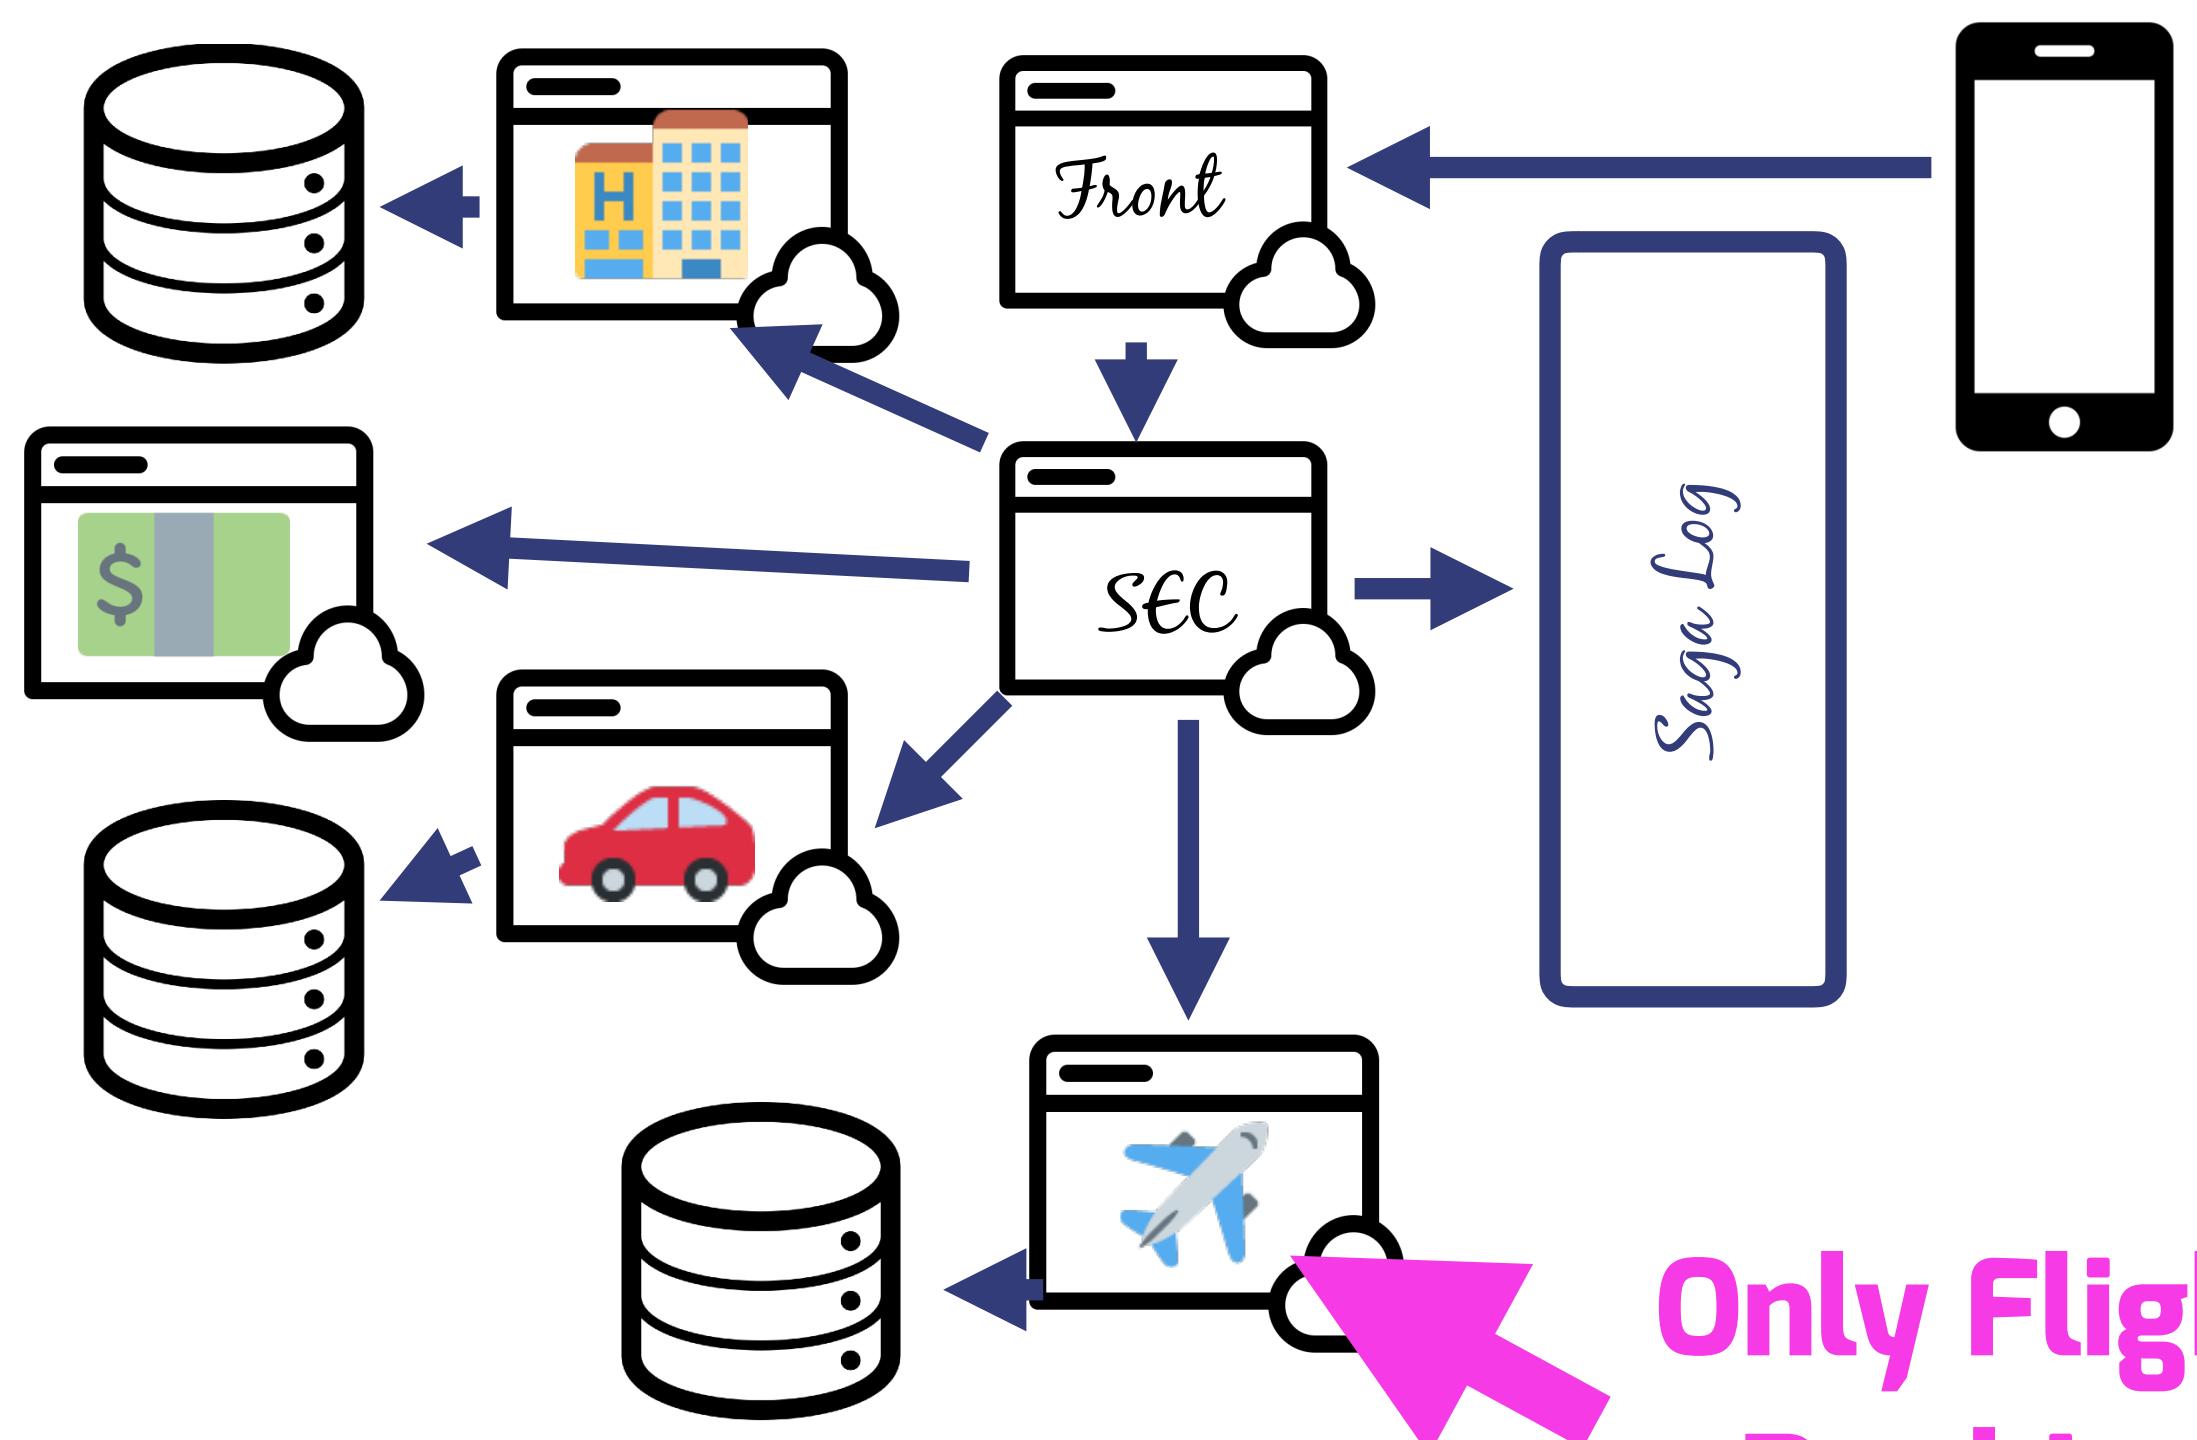
with Distributed Sagas



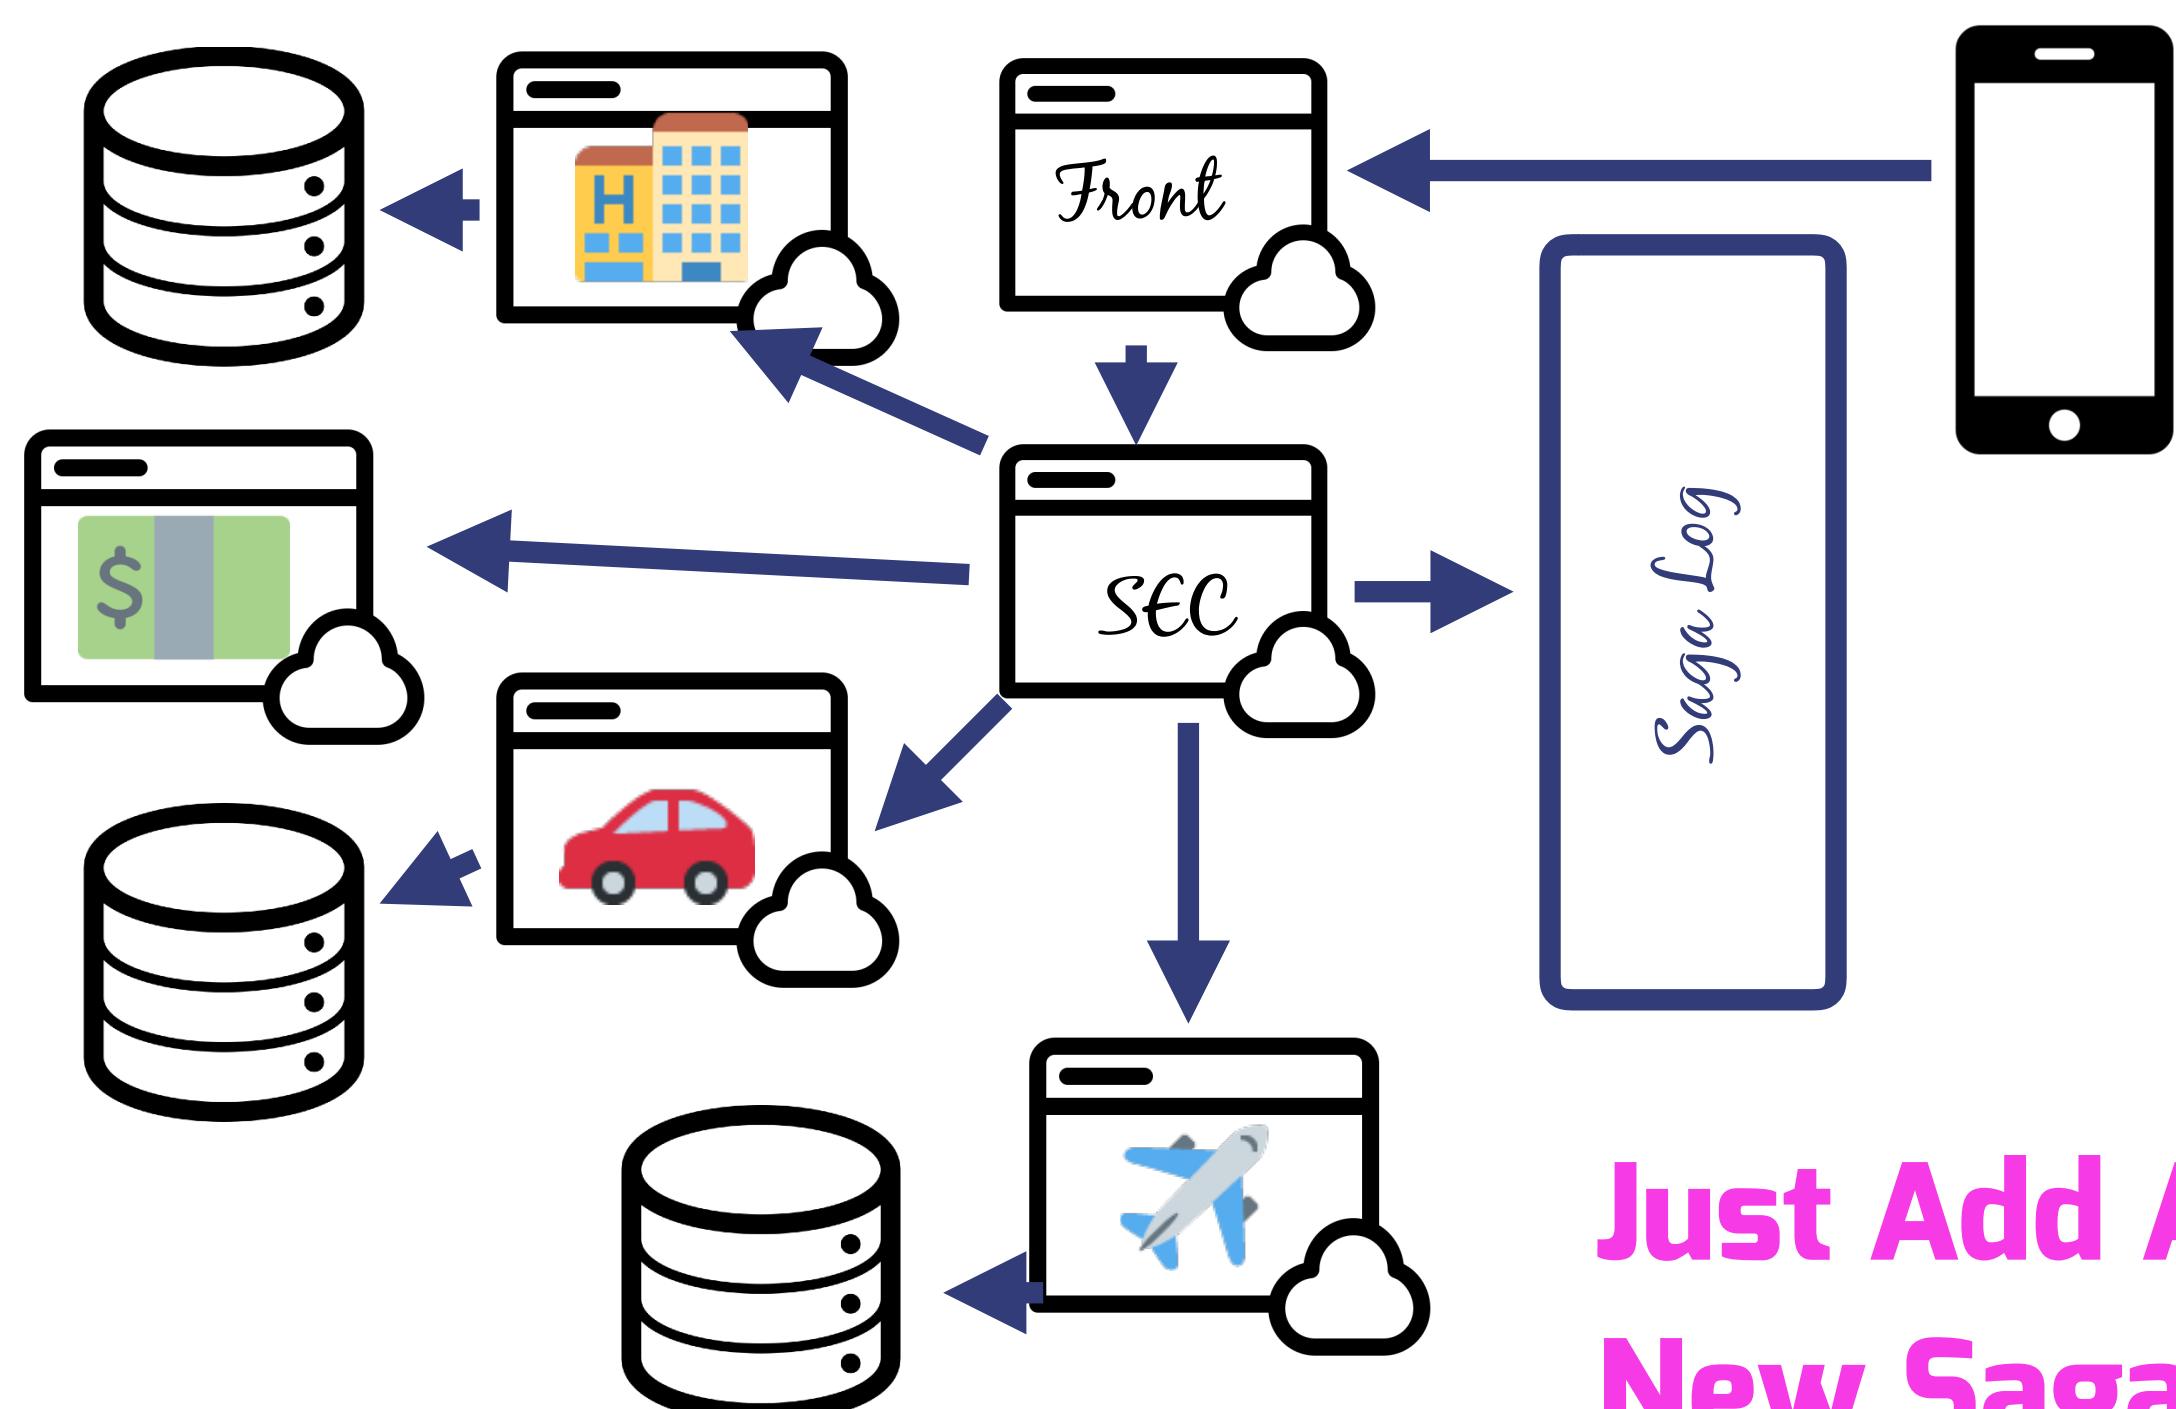
Complex Code
Lives Everywhere

Modular Services

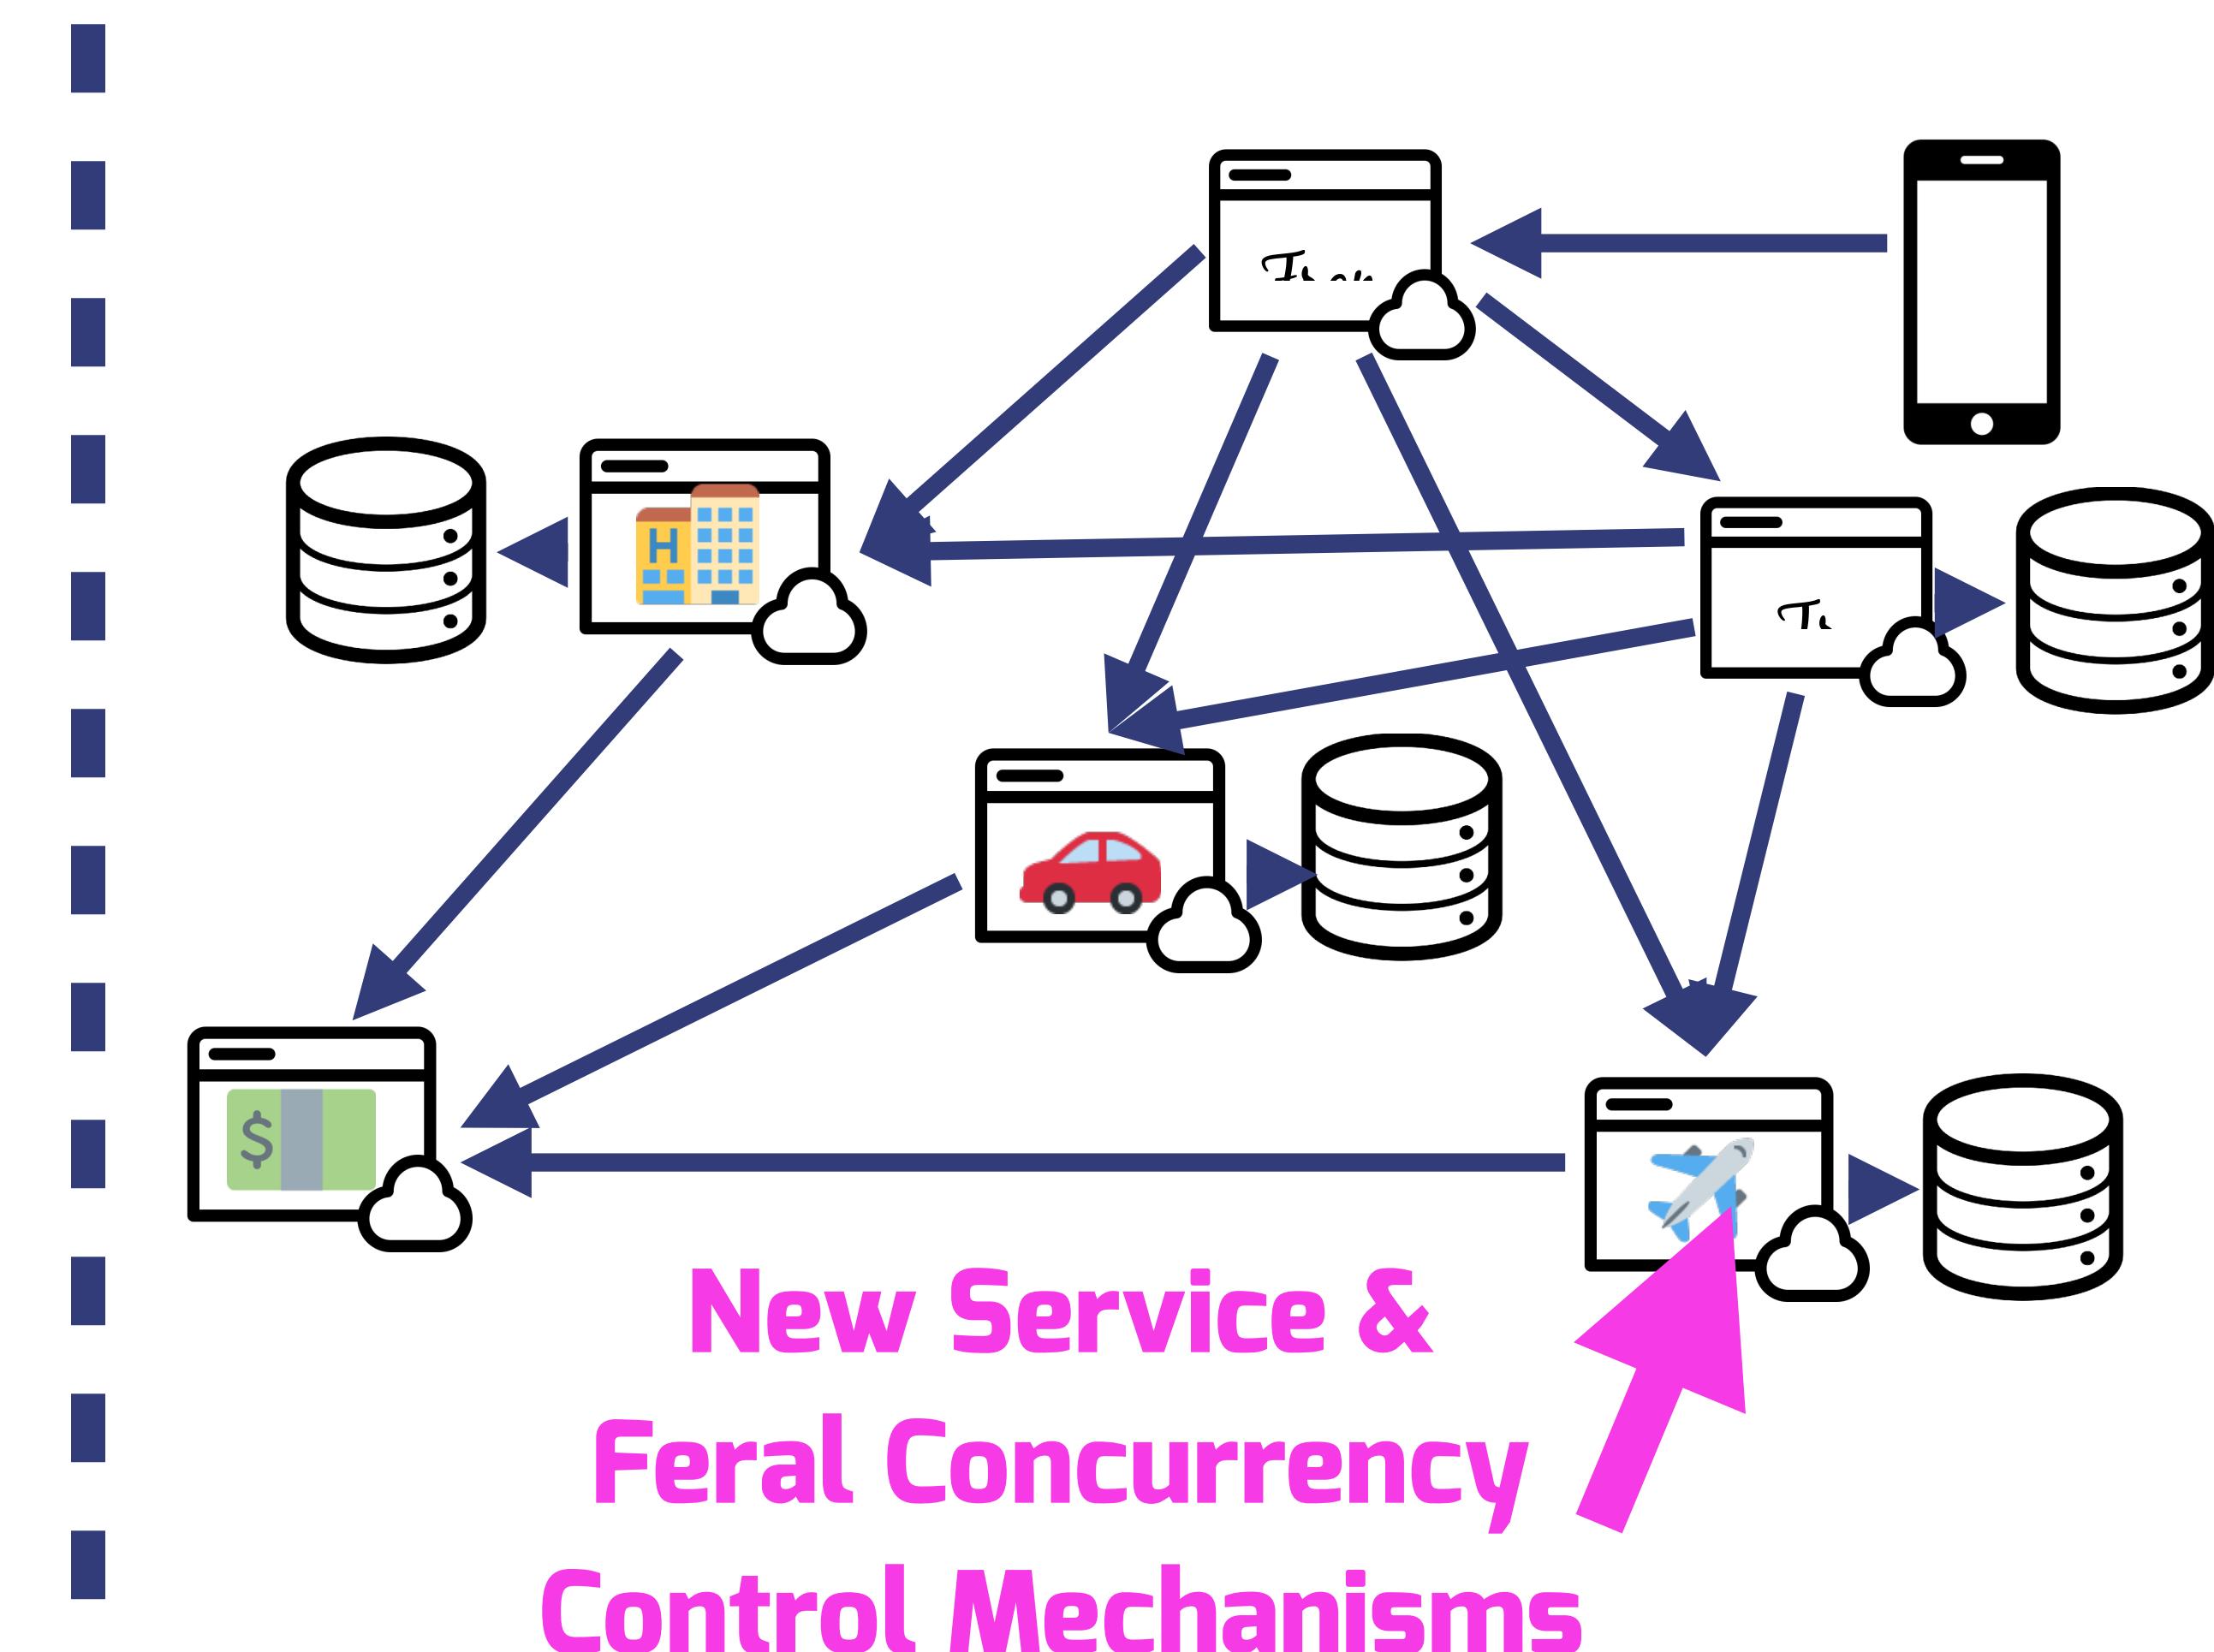
with Distributed Sagas



Service Composition with Distributed Sagas



Just Add A
New Saga!



New Service &
Feral Concurrency
Control Mechanisms

**Distributed Sagas Makes
Building & Modifying
Microservices Easier**

Thank you

@caitie